



Implementation Agreement (IA)

Common Management Interface Specification (CMIS)
Revision 5.3

IA Identification #
OIF-CMIS-05.3

September 4, 2024

Implementation Agreement created and approved
OIF
www.oiforum.com

The OIF is an international non-profit organization with over 150 member companies, including the world's leading carriers and vendors. Being an industry group uniting representatives of the data and optical worlds, OIF's purpose is to accelerate the deployment of interoperable, cost-effective, and robust optical internetworks and their associated technologies. Optical internetworks are data networks composed of routers and data switches interconnected by optical networking elements.

With the goal of promoting worldwide compatibility of optical internetworking products, the OIF actively supports and extends the work of national and international standards bodies. Working relationships or formal liaisons have been established with CFP-MSA, COBO, EA, ETSI NFV, IEEE 802.3, IETF, INCITS T11, ITU SG-15, MEF, ONF, Ethernet Alliance, IPEC, InfiniBand, SNIA SFF.

For additional information contact:

OIF

39221 Paseo Padre Pkwy, Suite J, Fremont, CA 94538 USA

+1.510.392.4902 – info@oiforum.com

www.oiforum.com

Physical and Link Layer (PLL) Working Group**Common Management Interface Specification (CMIS)****Rev. 5.3****Source****Technical Editor**

Stefan Langenbach
Cisco Systems, Inc.
Email: stefan.langenbach@cisco.com

Working Group Chair

David R. Stauffer, Ph.D.
Kandou Bus, S.A.
Email: david@kandou.com

PLL Working Group, Management Track, Vice Chairs

Ian Alderdice
Ciena Corp.
Email: ialderdi@ciena.com

Gary Nicholl
Cisco Systems, Inc.
Email: gary.nicholl@cisco.com

Abstract

This Implementation Agreement (IA) defines the Common Management Interface Specification (CMIS), which may be used by pluggable or on-board modules, such as QSFP Double Density (QSFP-DD), OSFP, COBO, or QSFP, as well as by existing or future module developments with host to module management communication based on a two-wire interface. This IA is targeted for systems manufacturers, system integrators, and suppliers of CMIS compliant modules.

Notice

This Technical Document has been created by the Optical Internetworking Forum (OIF). This document is offered to the OIF Membership solely as a basis for agreement and is not a binding proposal on the companies listed as resources above. The OIF reserves the rights to at any time to add, amend, or withdraw statements contained herein. Nothing in this document is in any way binding on the OIF or any of its members.

The user's attention is called to the possibility that implementation of the OIF implementation agreement contained herein may require the use of inventions covered by the patent rights held by third parties. By publication of this OIF implementation agreement, the OIF makes no representation or warranty whatsoever, whether expressed or implied, that implementation of the specification will not infringe any third party rights, nor does the OIF make any representation or warranty whatsoever, whether expressed or implied, with respect to any claim that has been or may be asserted by any third party, the validity of any patent rights related to any such claim, or the extent to which a license to use any such rights may or may not be available or the terms hereof.

Copyright © 2024 Optical Internetworking Forum

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction other than the following, (1) the above copyright notice and this paragraph must be included on all such copies and derivative works, and (2) this document itself may not be modified in any way, such as by removing the copyright notice or references to the OIF, except as needed for the purpose of developing OIF Implementation Agreements.

By downloading, copying, or using this document in any manner, the user consents to the terms and conditions of this notice. Unless the terms and conditions of this notice are breached by the user, the limited permissions granted above are perpetual and will not be revoked by the OIF or its successors or assigns.

This document and the information contained herein is provided on an "AS IS" basis and THE OIF DISCLAIMS ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY, TITLE OR FITNESS FOR A PARTICULAR PURPOSE.

Foreword

Please refer to chapters 1 – Introduction and 4 – General Concepts for orientation.

Revision History (Versions Produced and Published by the QSFP-DD MSA)

Please note that the revisions up to and including CMIS 5.2 were neither produced nor originally published by the OIF. These revisions are therefore not covered by the IP policy of the OIF. The OIF provides an archive of these revisions at <https://www.oiforum.com/documents/archived-non-oif-generated-specifications/>.

Rev 3.0 August 17, 2018

- Initial public release

Rev 4.0 May 8, 2019

WARNING: Implementations compliant to Revision 4.0 of this specification will not interoperate with Revision 3.0 implementations and vice-versa. The following specifications have changed:

- Converted InitMode signal to continuously sampled LPMode
- Associated Module and Data Path State Machine changes
- Added LPMode Override bit
- Data Path initialization control (Page 10h) meaning/polarity has changed
- Added LowPwr bit (Lower memory, Byte 26, bit 6) to switch power mode

Other changes:

- New format to conform to SFF
- Added Scope statement
- Added Conventions
- Added Definitions
- Updated Introduction
- Consolidation of chapter 5
- Rework and Partial Consolidation of chapter 6
- Added Pages 04h and 12h for tunable laser support
- Added diagnostic Pages 13h and 14h
- Added VDM Feature (Versatile Diagnostics Monitoring) (7.1)
- Added CDB Feature (Command/Reply Messaging) (7.2 and 9)
- Added Command Data Block (CDB) section (Pages 9Fh and A0h-Afh)
- Moved Examples to Appendix
- Change Upper Page Lower Page terminology (Upper and Lower Memory plus Pages)
- Removed Interface ID tables and replaced them with references to SFF-8024

Rev 5.0 May 8, 2021

This revision provides a significant number of changes and extensions, as well as comprehensive technical and editorial consolidation of CMIS 4.0. Specification weaknesses (gaps, bugs) have been fixed and some timing specifications have been adjusted, in order to improve host-module interaction performance in the future (see below for a detailed listing of changes). Resulting possibilities of cross-version incompatibilities have been deliberately accepted after careful consideration and broad industry survey.

Note: For instance, the meaning of the Flag Summary registers 00h:4-7 has changed. Careful study of the changes listed below is required to determine if any specific implementation may actually be unaffected.

WARNING: Depending on feature usage and implementation details, module implementations compliant to Revision 5.0 of this specification may or may not interoperate with Revision 4.0 host implementations, and host implementations of Revision 5.0 may need to be module version aware.

HINT: It is strongly recommended that all CMIS 4.0 implementations should be upgraded to CMIS 5.0.

The following list summarizes and classifies the changes made in this revision:

Register Map Changes

- 00h:2.6 SteppedConfigOnly (new advertisement bit)
- 00h:41 Fault Cause register (new, optional)
- 10h:232 OutputStatusChangedMaskRx (Mask for new Flag)
- 11h:235 DPInitPending (new status register)
- 11h:202-205 ConfigStatus (now required, new ConfigInProgress status code)
- 11h:132-133 OutputStatusRx, OutputStatusTx (new)
- 11h:153 OutputStatusChangedFlagRx (new)

- 12h:128-135.7-4 33GHz and 75GHz encoding swapped
- 15h New page for Timing Characteristics

CDB Message Changes

- CMD41h, 137.3: New scaling multiplier for maximum duration advertisements (backwards compatible)
- CMD00h, 138: Response message length shortened to 2, obsolete byte 138 removed from response.

Specification Bug Fixes

- TEC Current units of Aux2 Monitor
- LPL Message Length of some CDB command(s)
- Apply_DataPathInit clearing condition was not specified
- Apply_DataPathInit state versus trigger ambiguity resolved (it is a trigger register)
- MediaLaneSupported was mistakenly marked as optional (it is Rqd.)

Extensions (optional)

- Added Page 15h (Timing characteristics for PTP) in section 8.14 (optional)
- Added Fault Cause register in section 8.2.1 (optional)
- Added Firmware readback using CDB in section 7.3.1.3 (optional)
- Intervention-free reconfiguration on Apply* ¹may be advertised as not supported (SteppedConfigOnly)

Technical Requirements Added

- Added OutputStatusRx (with state change flag) and OutputStatusTx in section 8.10.2 (required)
- Added VDM Flag conformance per state in section 6.3.4.3 (new restrictions)
- Added clear specification of Bank and Page Select behavior (section 8.2.15) and timing (section 10.2)

Technical Requirements Changed

- Flag summaries now indicate bank and page (instead of bank and lane) (changed specification)
- Configuration command (Apply*) result status in ConfigStatus register now mandatory
- Configuration command (Apply*) now synchronized with host by ConfigStatus=ConfigInProgress
- ApplyImmediate in DPDeactivated now ignored. ApplyImmediate and ApplyDPInit now have a distinct and non-overlapping purpose; ApplyImmediate only for hot reconfiguration.
- Size matched read of scalar numerical RO data now guarantees data coherency (new restriction)
- Flags on optional pages are now masked by default (changed defaults)
- Generally revised password usage and support (standard CMIS provides only mechanism)
- No password protection allowed for CMIS standardized features
- Timing parameters harmonized with changes defined in QSFP-DD HW Spec revision 5.1.
- Frequency grid encoding in field (12h:128-135.7-4) has swapped the codes for 33GHz and 75 GHz
- Removed ill-defined peak detector from diagnostics features (configuration bits now Reserved)
- Apply* trigger registers allow single byte WRITE only (new restriction)
- The READ primitive allows to read N≤8 bytes, unless explicitly specified different (previous revisions did not have this restriction, allowing the host to read a full page in one READ transaction).

Technical Content Description

- Management interface timings now normatively defined in CMIS, initialized from QSFP-DD HW spec 5.1
- Refined description of conditionally required fields (section 8.1.3.7)
- TWI renamed as I2C-based MCI (I2CMCI) and described as the first and only instance of an MCI
- Clarification that READ/WRITE retry is an alternative to Acknowledge Polling
- Acknowledge polling is now abstracted as TEST access
- Added descriptive section 7.4: Tunable Lasers, Channel Selection and Center Frequency Setting
- Clarified Wavelength Information fields in Page 01h (section 8.4)
- Updated Appendix C: Example initialization flows
- Added sections 8.2.10.1: Password Type, and 8.2.10.2: Password entry and change
- Added sections on configuration and reconfiguration Command Handling

Changes of document structure

- Some table and figure numbers have changed due to insertion, splitting, or reordering
- Refactored chapter 5: Introduced register ACCESS primitives (READ, WRITE, TEST)
- Register ACCESS primitives implemented by MCI transactions described in new Appendix B.
- Reworked chapter 6 substructure for clarifications and additions
- Added chapter 10 for timing specifications (including MIS specs, migrating from QSFP-DD HW spec)
- Inserted Appendix B for Management Communication Interface (MCI) specifications
- Inserted Appendix E with illustrations of effects after ApplyDPInit
- Added Appendix G for CMIS evolution and maintenance guidelines

Changes of presentation and overall text consolidation

- Expanded Purpose and Scope chapter 1

¹ The asterisk '*' is used as a placeholder for arbitrary text, here e.g. for ApplyImmediate or ApplyDPInit.

- Updated document conventions for better text homogeneity in chapter 2
- Updated (extended and enhanced) term definitions in chapter 3
- Introduced arithmetic and other data type definitions in chapter 3
- Expanded description of CMIS modules and CMIS compliance, and document overview in chapter 4
- Updated text to comply to the documented conventions of chapter 2
- Updated text to use documented terminology consistently
- Chapter 6 text consolidated, sub structured, and revised
- Equations are now numbered
- Updated Figures 8-1 and 8-2 describing the management memory map
- Added section 8.1.3 with general specifications across the management memory map
- Added description of management memory access types in the Type column of tables in chapter 8
- Improved description of repetitive field array definitions. Register information in Description column
- Register description now often includes advertisement information
- Reduced verbosity in chapter 8
- Preparations for automated register map extraction
- Removed inverted logic bOOL data type (no longer used)

Renaming and Rewording

- Replaced socially inappropriate technical Master/Slave terminology with **Initiator/Target** terminology
- **Renamed memory map** fields systematically: informal name phrases converted to software-ready single-word field names with instance indexing suffixes (see temporary Appendix H)
- Systematic renaming of controls and signals contributing to a module **Low Power Mode request**:
 - ForceLowPwr → LowPwrRequestSW
 - LPMode → LowPwrRequestHW
 - LowPwr → LowPwrAllowRequestHW
- **Squelching** related terminology updated to better distinguish functionality and control of functionality
- Wording changes due to refactoring of chapter 5 (using WRITE, READ, TEST for register **ACCESS**)

Known Problems

- Unresolved inconsistencies in the description of Diagnostics (multi-bank control)
- Description of CDB in chapters 7, 8, 9 needs further consolidation (structure, dependency, redundancy)

Rev 5.1 November 2, 2021

This revision adds optional support for **client encapsulation modules** (see chapter 1) and applications, which can establish and maintain the media side transmission link independent of the presence of host side client signals, **with or without** client signal **multiplexing**.

Modules supported in CMIS 5.0 (**system interface modules**) are **not affected** by this revision at all.

Compatibility

- Existing CMIS 5.0 compliant module implementations are also CMIS 5.1 compliant
- Existing CMIS 5.0 compliant host implementations may not understand the content of new multiplex Application Descriptors advertised by CMIS 5.1 modules supporting client encapsulation applications

High Level Changes

- Added the concept of Network Path (NP) and Network Path State Machine (NPSM) to represent resources and control of the media side of uniplex or multiplex applications
- Added the concept of Host Paths (HP) as the host side segment of Data Paths (DP) feeding NPs in uniplex or multiplex applications, reusing DP registers and state machine (DPSM)
- Glossary adaptations

Register Map Changes

- 01h:142.7 advertises Network Path (NP) support
- 16h:128-255 new page for Network Path (NP) feature
- 17h:128-255 new general page for Flags and Masks

Bug Fixes

- Corrected CdbChkCode for CMD 0004h as FBh
- Changes of Document Structure**
- Added text section 7.6.5
 - Inserted memory map sections 8.14 and 8.15 for Pages 16h and 17h
 - Replaced temporary appendix H with examples for Network Path Applications

Editorial Bug Fixes

- Two tables following 8-88 were misnumbered and not in TOC (table numbering in chapter 8 changed)
- Captions of tables containing memory map definitions now all end with (Page #) (for Excel extraction)

Editorial Changes and Presentation Changes

- Change bars are retained to allow quick recognition of changes or additions.
- Slightly expanded description of the FarEndConfiguration of cable assemblies in section 8.3.8
- Expanded description of ModSelWaitTime in section 8.4.5 in response to a question
- Retired the locally defined term 'data link terminals' in favor of its alternative 'client encapsulation modules' that is now used consistently throughout the document (chapter 1)
- Line numbers for error reporting with higher contrast and visually less intrusive (smaller size)
- Advertisement registers cross references in section 8.9.2 corrected and some wording clarifications
- Clarification notes added in section 8.4.7
- A few strategically placed notes have been added to chapter 6, to help the reader understanding the role of this chapter (written for system interface applications) for multiplex applications

Revision History (Versions Produced and Published by the OIF)**Rev 5.2** April 27, 2022

This revision adds clarifications and optional features only.

Compatibility

- Existing CMIS 5.1 compliant (host and module) implementations are also CMIS 5.2 compliant

High Level Changes (triggered by contributions as indicated)

- Support of unidirectional Data Path configuration, in support of Fibre Channel needs (oif2021.530)
- Clarifications and optional bank broadcast feature for modules with 32 lanes (oif2021.525)
- Clarifications on sampling and sample statistics of VDM observables BER and FERC (oif2021.494)
- New VDM observables 25, 26 for accumulated rather than averaged FERC (oif2021.494)
- New optional Page 05h restricted to form-factor specific definition of meaning and management (advertisement, control, status) of form factor specific management signals, in separate OIF IA (oif2021.574)

Register Map Changes

- 00h:26.7 enable optionally supported lane-banked register WRITE broadcast across banks
- 01h:142.3 advertise support for new restricted page 05h (content defined elsewhere)
- 01h:156.7 advertise support for lane-banked register WRITE broadcast
- 01h:162.6 advertise support of unidirectional Data Path reconfiguration (for Fibre Channel)
- 05h:128-255 new optional page for management of form factor specific management signals
- 18h:128-255 new conditionally optional page reserved for additional lane specific controls
- 19h:128-255 new conditionally optional page for lane specific status information
- 19h:128-143 unidirectional Active Control Sets

Editorial and Presentation Changes

- New sections: 2.1.2, 7.7, 8.1.3.3, 8.8, 8.9.3.4, 8.9.4.4, 8.17, 8.18
- Refined text and added protocol stack clarifications to introductory sections 4.1.1 and 4.2
- Refined text of section 5.1 (management signaling layer, MSL extensions)
- Added clarification on 8 lane limit per Application in section 6.2.1
- Refined text of section 6.2.5 for clarification of explicit control
- Added note on significance of LowPwrRequestHW for startup in section 6.3.2.2
- Textual changes in sections 7.1.2 and 7.1.5 (importance of low sampling period variance)
- Updates and clarifications in sections 8.1.1, 8.1.2, 8.3.7, 8.9.2, 8.10.6
- Minor textual consolidation in informative section 8.12.11
- Clarification in Appendix A
- Reduced the number of orphan pages (by compaction of informal text)
- Removed typos and extraneous white space
- Fixed advertisement hint for 10h:137-139 in Table 8-69
- Fixed min/max mix-up in description of 01h:166.6-0 in Table 8-54

1 Rev 5.3 September 4, 2024

2 **Compatibility**

- 3 - Existing CMIS 5.2 compliant (host and module) implementations are also CMIS 5.3 compliant

4 **Project Based Extensions**

- 5 - Support for CPO (E01, E02, E16)
6 - Support for ELSFP (E01, E02, E15: Resource Module Concept)
7 - Support for Host Module Link Training (E03: Supplement)
8 - Resource Module Concept (E15, E30)
9 - SPI-based SPIMCI in support of CPO (E04, E05)

10 **Maintenance Changes¹**

- 11 - M11, M20, Revision History: added missing documentation about changes in Rev 5.0
12 - M21, 2.1.1: updated specification reference and table references in Table 6-1, Table 8-5, Table 8-18,
13 Table 8-33 to SFF-8024 rev 4.10
14 - M13, 3.3: added glossary entries on delay and hold-off a.k.a. reject for ACCESS and transaction, to
15 clarify the concepts of delay and rejection in both RAL and MCI layers.
16 - Editorial, 3.3: alphabetical sorting in the glossary (without change bars: Multiplex, Register)
17 - Editorial, beginning of chapter 5: Added a footnote on MSL signals.
18 - Editorial, 5.2.3: reformatted host options of how to react to ACCESS rejection as a list
19 - M4, 6.2.5.2: added notes to Equalization equations 6-1 and 6-2
20 - M23, 6.3.2.5.4, Table 6-14, glossary: clarified management functions available in ModuleLowPwrState
21 - E07: 6.3.4.2: added Flag Conformance for the Flags in Page 12h (see Table 6-21)
22 - M9, 7.1.2: added note on desirable saturation of VDM counters
23 - M26, 7.1.2: clarified limited scope of VDM advertised sampling period (FineIntervalLength)
24 - M22, 7.1.2, 8.7, and 8.22.3: clarified for the host that the module *may* update application dependent
25 alarm/warning thresholds, but only when commissioning a new Application
26 - M17, 7.1.4, 7.1.4.1, Table 8-129, and Table 8-170: updated definition of Host Side SNR
27 - M10, 7.3.1: added notes recommending use of a single firmware download file
28 - M27, 8.1.3.1: added notes on reading or writing reserved registers
29 - M18, 8.2.1: clarified SteppedConfigOnly, added note on prevalence of option 1 in Table 8-5
30 - M24, 0: added simplified note about writing to the Page Mapping registers
31 - M25, 8.9.3: 8.9.3: resolved contradicting specifications in section 8.9.3 (claiming that ApplyImmediate
32 triggers provisioning in DPDeactivated) and specifications in Table 6-3
33 - M8, 8.12.4: corrected text of note
34 - M5, 8.12.6 through 8.12.9: corrected textual references to advertisement bits in tables Table 8-109,
35 Table 8-111, Table 8-113, Table 8-115
36 - M1: 8.12.6 through 8.12.9, corrected pattern generator and checker field name component 'ByteSwap'
37 as 'SwapSymbolBits' for Bytes 13h:146, 13h:154, 13h:162, and 13h:170
38 - M16, 8.13.3: corrected addresses for MediaSideBERLane8 and GatedMediaSideBERLane8 in Table 8-129
39 - M14, 8.18: corrected size of reserved bytes on page 19h
40 - M15, 8.22.1.2: corrected VDM Observable Laser Age as media lane related in Table 8-170
41 - M12, 9.3.1: Clarified REPLY specification of CDB CMD 000h
42 - M6, 9.4.1: clarified advertisements for supported CDB commands in CDB CMD 0040h
43 - M7, 9.4.2 and 9.7: clarified description of advertisements for FW update commands in CDB CMD 0041h
44 - E17: 9.7.9, 9.7.10: clarified max duration for CDB CMD 109h (Run FW) and 10Ah (Commit FW)
45 - M2, B.2.5.1.2: fixed byte shown with 9 bits in Figure B-2

46 **Extension Changes²**

- 47 - E28: 2.1.2, 9.1, 9.4.6: added support of the new **Versatile Control Set** feature, with SI parameters
48 and their register locations advertised by the module (**CMIS-VCS** supplement [6])
49 - E23: 5.1, Appendix A: allows **modules without** dedicated **LowPwrRequestHw** pin, like SFP112
50 - E25: 5.2.2.1, 8.4.16: added advertisement for **full page READ**. May become mandatory in the future.
51 - E40: 6.1, 8.2.13, 8.3.8: add support **cables with more than 8 lanes** with uniform far end breakout
52 and **PCIe** support
53 - E26: 0, 8.4.13, 8.17.1, 8.18.2, 8.20: extended the **number of Applications** that can be advertised to
54 240, in groups of 15 **Normalized Application Descriptors (NAD)** in banked **page 1Ch** and identified

55 ¹ The M# identifiers refer to the relevant Maintenance Work List, for traceability

56 ² The E# identifiers refer to the relevant Objectives Work List, for traceability

- 1 by **Application Numbers** (the AppSel code is now a local identifier in a group).
2 - E27: 6.2.4, 8.2.1, 8.9.3: allow partial support of intervention-free **reconfiguration** advertised in Table
3 8-5, with associated **descriptions changed** for consistency in Table 6-3 and Table 6-4.
4 - E19: 7.1.2, 8.22.6: added optional (advertised) **VDM power saving mode**
5 - E42: 7.1.5, 8.22.1.2, 9.8.8, 9.8.9, 9.8.11: add CDB CMDs Control/Get **FEC SEW Histogram** and
6 Control/Get **Max FEC SEW**, and add VDM Statistics Observables for **SEWmax**
7 - E14: 7.8, 8.4.16, 8.21: added support for **Host Lane Switching** flexibility feature
8 - E13: 8.2.10: added **new fault codes** in 00h:41 (Table 8-16)
9 - E21: 8.2.11, 8.2.14, 8.4.16: added advertised **password entry result** register
10 - E10: 8.2.12: added **SFF8024FiberFaceType** field in 00h:42-45 (see Table 8-17)
11 - E11: 8.2.12: added **SFF8024ModuleSubType** field in 00h:42-45 (see Table 8-17)
12 - E09: 8.2.12: added **advertised** non-service-related **management limitations** in **ModuleLowPwr**
13 - E22: 8.2.14, 8.4.16: added password support advertisement
14 - E41: 8.3.5, 8.3.6, 8.3.9: extend scope of passive copper related data to **linear active copper**
15 - E37: 8.3.6: added a **new frequency** node to advertised **cable attenuation** in Table 8-34
16 - E36: 8.4.2: added **extension of reach advertisement** in Table 8-44
17 - E06: 8.4.4: added **C-CMIS support advertisement** in 01h:142.4 (see Table 8-46)
18 - E08: 8.4.12: added **advertised** max Bank/Page Change time **tBPC** in 01h:169 (Table 8-56)
19 - E33: 8.7: added **150GHz grid support** in Table 8-66, Table 8-99
20 - E33: 8.7, 7.5.1: added note on implicit **L-Band** support, renamed center frequency as reference
21 - E18: 8.12.13: added an 8-byte host **scratchpad register** 13h:184-191, advertised in 01h:251.7
22 - E32: 9.8.7: added CDB command for **RMON** statistics
23 - E38: 9.2.3, 9.2.4, 9.4.7: CDB command to query (mostly) string **descriptions** of an **Interface Code**
24 - E39: 9.2.3, 9.2.4, 9.4.8: CDB command to query **application** specific **attributes**
25 - E12: E31: 9.12: added new CDB Command Group for **security** features, module **authentication**

Structural Document Changes

- 27 - Demoted and combined Purpose and Scope chapter 1 with Document Overview in new intro chapter 1
28 - New section 1.1.2 on CMIS **supplement** specifications (extending the CMIS **base** specification)
29 - Added references to new CMIS supplement specifications in section 2.1.2
30 - Additions to Abbreviations and Glossary in chapter 3
31 - New subsection 4.1.2 on Resource Modules
32 - New subsection 4.4 on cross version interworking and compatibility
33 - Minor adaptations for Resource modules in chapter 6
34 - CPO Monitors explained in new subsection 7.1.7
35 - Updated memory map in Table 8-1
36 - New appendix for SPIMCI in section B.3
37

Register Map Changes

- 39 - **00h:2.5-2** added MciMaxSpeed Advertisements for SPIMCI and I2CMIC
40 - **00h:2.1-0** added AutoCommissioning advertisement
41 - **00h:27.3-0** added MciSpeedConfiguration
42 - **00h:41** added new Fault Codes (security or damage related)
43 - **00h:42.3-0** added PasswordCommandResult
44 - **00h:56** added CmisSmSupport for state machine support advertisement
45 - **00h:57** added ModuleFunctionType advertisement
46 - **00h:60** added SFF 8024 module sub type
47 - **00h:61** added SFF 8024 fiber face type
48 - **00h:62** added advertisement of implementation dependent LowPowerRestrictions
49 - **00h:208** added attenuation node AttenuationAt53p1GHz
50 - **00h:213** added MCciFlowControlDuration advertisement (SPIMCI)
51 - **01h:137.7-6** added additional fiber link length multipliers
52 - **01h:142.4** added advertisement for C-CMIS support
53 - **01h:142.3** added advertisement for CMIS-FF support
54 - **01h:162.7** added advertisement for CMIS-VCS support
55 - **01h:169** added advertisement MacDurationPBC for tBPC (bank/page change)
56 - **01h:175** added advertisement NADBanksSupported
57 - **01h:251.7-6** added advertisement for host scratch pad area
58 - **01h:251.5-2** added advertisement password entry and result
59 - **01h:251.1-0** added advertisement for full page read

- 01h:252.7 added advertisement for Host Lane Switching
- 01h:252.6 added advertisement LinkTrainingSupportedfor CMIS-LT
- **02h** added hint on possible Application dependency
- **04h:129.6** added GridSupported150GHz
- 04h:129.162-163 added GridLowChannel150GHz
- 04h:129.164-165 added GridHighChannel150GHz
- **06h-07h** marked pages restricted for use by Resource Modules
- **08h-0Bh** marked pages restricted for CMIS-LT
- **10h:156-159** renamed informative field name to HostControlledInputEqTarget (also elsewhere)
- **12h:128-135.7-4** added 150GHz grid spacing
- **13h:184-191** added host scratchpad register
- 13h:146 renamed informative field name to HostSideGeneratorSwapSymbolBits
- **18h:128-143** added Staged Control Set extensions for NAD Block Indices
- 18h:144-255 added CMIS-VCS parameter space
- **19h:144-151** added Active Control Set extensions for NAD Block Indices
- 19h:152-207 added CMIS-VCS parameter space
- **1Ah-1Bh** marked pages restricted for definition in Resource Module IAs
- **1Ch:128-247** added Banked Normalized Application Descriptors (**NAD**) page for NAD blocks
- **1Dh** added page for Host Lane Switching feature
- **50h-5Fh** pages restricted for CMIS-LT

VDM Observables Changes

- New Maximum FEC Corrected Symbols (Symbol Error Weight) Statistics (27-34) in subsection 8.22.1.2.
- New module power supply monitors (77-83) for CPO in subsection 8.22.1.2.
- New ELS optical input power monitors (84) for CPO in subsection 8.22.1.2.

CDB Messages Changes

- 0044h added Security Features and Capabilities query
- 0045h added Supported Externally Defined Features query
- 0050h added Get Application Attributes query
- 0051h added Get Interface ID Description query
- 0220h added Get Data Path RMON statistics query
- 0230h-0231h added FEC Symbol Error Correction statistics (symbol error weight histogram)
- 0232h-0233h added Maximum FEC Symbol Error Weight statistics
- 0400h-0406h added secure device authentication commands
- 4000h-40FFh restricted CMD IDs for commands defined in CMIS-VCS [6]

Editorial Changes

- Harmonized the CDB message description table format, generally separating REPLY status and data
- Bugfix of some RPLLength values incorrectly including the REPLY Header
- Extended lane related terminology for clarity of host lane switching in sections 2.3.4 and 3.3
- Editorial changes in chapters 1-6
- Hints that threshold values may change with Applications

Contents

Revision History (Versions Produced and Published by the QSFP-DD MSA)	4
Rev 3.0 August 17, 2018	4
Rev 4.0 May 8, 2019	4
Rev 5.0 May 8, 2021	4
Rev 5.1 November 2, 2021	6
Revision History (Versions Produced and Published by the OIF)	7
Rev 5.2 April 27, 2022	7
Rev 5.3 September 4, 2024	8
Contents	11
Figures	20
Tables	20
1 Introduction 26	
1.1 Purpose and Scope	26
1.1.1 CMIS Specification (Base)	26
1.1.2 CMIS Supplement Specifications	27
1.2 Document Overview	28
2 References and Conventions	29
2.1 Industry Documents	29
2.1.1 Interdependent Documents	29
2.1.2 CMIS Extensions and Supplements	29
2.1.3 General Background	29
2.2 Sources 29	
2.2.1 Standards and Specifications	29
2.2.2 Administrative Material	29
2.3 Conventions	30
2.3.1 General Conventions	30
2.3.2 Notational Conventions	30
2.3.3 Typographical Conventions	32
2.3.4 Addressing and Referencing Conventions	32
3 Definitions	34
3.1 Keywords	34
3.2 Abbreviations	35
3.3 Glossary	37
3.4 Data Types	46
3.5 Data Operations	46
4 General Concepts	47
4.1 CMIS Managed Module Classification	47
4.1.1 CMIS Compliant Transmission Modules	47
4.1.2 CMIS Compliant Resource Modules	48
4.2 CMIS Management Protocols and Layers	49
4.3 CMIS Compliance and Vendor Specific Extensions	50
4.4 CMIS Interworking and Compatibility	50
5 Management Interface	51
5.1 Management Signaling Layer (MSL) Control and Status Signals	51
5.2 Management Register Access Layer (RAL)	52
5.2.1 Registers in Addressable Memory	52
5.2.2 Register Access Methods for Addressable Memory	52
5.2.2.1 Read Access (READ)	52
5.2.2.2 Write Access (WRITE)	52
5.2.2.3 Test Access (TEST)	53
5.2.3 Register Access Rejection (Access Hold-Off)	53

1	5.2.4	Register Access Sequencing and Synchronization Requirements	53
2	5.2.5	Register Data Coherency	53
3	5.2.5.1	Coherency of Size-Matched Multi-Byte READ Results	54
4	5.2.5.2	Coherency of Multi-Byte WRITE Operations	54
5	6	Core Management Features	55
6	6.1	Transmission Module Management Basics	55
7	6.1.1	Host Interface	55
8	6.1.2	Media Interface	55
9	6.1.3	Memory Map Representations	56
10	6.2	Transmission Module Functional Model	56
11	6.2.1	Functional Transmission Module Capabilities – Applications	56
12	6.2.1.1	Applications and Application Instances	56
13	6.2.1.2	Data Paths for Application Instances	56
14	6.2.1.3	Multiple Application Instances and Multiple Applications	57
15	6.2.1.4	Application Advertising (Application Descriptors and AppSel Codes)	58
16	6.2.1.4.1	Basic Application Advertisement (up to 15 Applications)	58
17	6.2.1.4.2	Normalized Application Advertisement (up to 240 Applications)	60
18	6.2.1.5	Application Selection and Instantiation	60
19	6.2.2	Application Instances on Data Paths – Data Path Lane Assignment	60
20	6.2.3	Data Path Configuration – Control Sets	61
21	6.2.3.1	Control Sets Concept	61
22	6.2.3.2	Control Set Content	61
23	6.2.3.3	Control Set Usage (Applying a Staged Control Set)	63
24	6.2.4	Data Path Configuration – Procedures and Commands	65
25	6.2.4.1	Configuration Concept	65
26	6.2.4.2	Configuration Commands	67
27	6.2.4.3	Host Rules and Recommendations	68
28	6.2.4.4	Initialization Sequence Examples	69
29	6.2.5	Signal Integrity Related Controls	70
30	6.2.5.1	Tx Input Equalization Control	70
31	6.2.5.2	Rx Output Equalization Control	71
32	6.2.5.3	Rx Output Amplitude Control	71
33	6.2.5.4	Adaptive Tx Input Equalizer Store and Recall Mechanism	72
34	6.3	Module Behavioral Model	73
35	6.3.1	State Machine Concept	73
36	6.3.2	Module State Machine (MSM)	74
37	6.3.2.1	General Module Behavior	74
38	6.3.2.2	Module State Machine for Paged Memory Modules	75
39	6.3.2.3	Module State Machine for Flat Memory Modules (Simplified)	78
40	6.3.2.4	Module Power Mode (Module Characteristic)	79
41	6.3.2.5	Module State Machine States	80
42	6.3.2.5.1	Resetting State (Shutting Down)	80
43	6.3.2.5.2	Reset State (Ground State)	80
44	6.3.2.5.3	MgmtInit State (Initializing Management Interface)	80
45	6.3.2.5.4	ModuleLowPwr State (Basic Manageability)	81
46	6.3.2.5.5	ModulePwrUp state (Powering Up)	82
47	6.3.2.5.6	ModuleReady State (Fully Operational)	82
48	6.3.2.5.7	ModulePwrDn State (Powering Down)	82
49	6.3.2.5.8	ModuleFault State (Module Fault)	83
50	6.3.3	Data Path State Machines (DPSM)	84
51	6.3.3.1	DPSM State Transition Diagram and DPSM Specification	85
52	6.3.3.2	Data Path Control (Host)	88
53	6.3.3.3	Data Path Status (Module)	88
54	6.3.3.4	DPDeactivated State (Ground State)	89
55	6.3.3.5	DPIinit State (Initializing)	90
56	6.3.3.6	DPIinitialized State (Initialized)	91
57	6.3.3.7	DPEinited State (Deinitializing)	91
58	6.3.3.8	DPTxTurnOn State (Turning On)	92
59	6.3.3.9	DPAactivated State (Operational)	92

1	6.3.3.10 DPTxTurnOff State (Turning Off)	93
2	6.3.4 Flagging Conformance Rules per State	94
3	6.3.4.1 Module-Level Flagging Conformance Rules per Module State	94
4	6.3.4.2 Lane-Specific Flagging Conformance per Data Path State	94
5	6.3.4.3 VDM Flagging Conformance per State	96
6	7 Advanced Management Features	97
7	7.1 Versatile Diagnostics Monitoring (VDM)	97
8	7.1.1 Purpose and Background	97
9	7.1.2 Technical Overview	97
10	7.1.3 Technical Details	98
11	7.1.3.1 Observables	98
12	7.1.3.2 Thresholds	98
13	7.1.3.3 Structure	99
14	7.1.4 PAM4 Observables	100
15	7.1.4.1 Signal to Noise Ratio (SNR) Observable	101
16	7.1.4.2 PAM Level Transition Parameter (LTP) Observable	101
17	7.1.5 Error Performance Statistics (FEC)	102
18	7.1.5.1 FEC Related Primary Observables	102
19	7.1.5.2 FEC Related Statistics	102
20	7.1.6 Tunable Laser Monitors	104
21	7.1.6.1 TEC Current	104
22	7.1.6.2 Laser Frequency Deviation	104
23	7.1.6.3 Laser Temperature	104
24	7.1.7 Monitors for Co-Packaged Optics (CPO)	104
25	7.2 Command Data Block (CDB) Message Communication	105
26	7.2.1 Feature Overview	105
27	7.2.2 CDB Technical Basics	105
28	7.2.3 CDB Message Sending and Reply Receiving (Host)	105
29	7.2.4 CDB Message Receiving and Reply Sending (Module)	105
30	7.2.5 CDB Support Levels and Messaging Details	106
31	7.2.5.1 Foreground Mode CDB Messaging	106
32	7.2.5.2 Background Mode CDB Messaging	106
33	7.3 CDB Message Communication Based Features	108
34	7.3.1 Firmware Management using CDB	108
35	7.3.1.1 Firmware Management Concepts	108
36	7.3.1.2 Reference Firmware Download Procedure using CDB	110
37	7.3.1.3 Reference Firmware Upload Procedure using CDB	111
38	7.3.1.4 Firmware Administration and Status	113
39	7.3.2 Performance Monitoring (PM) using CDB	115
40	7.3.3 Security Features and Capabilities using CDB	115
41	7.4 Module Boot Record (MBR)	115
42	7.5 Tunable Lasers, Channel Selection and Center Frequency Setting	116
43	7.5.1 Concepts and Background	116
44	7.5.2 Programming Overview	116
45	7.6 Client Encapsulation Applications (Multiplexing)	118
46	7.6.1 Concepts and Technical Background	118
47	7.6.1.1 Multiplex or Uniplex Client Encapsulation Applications	118
48	7.6.1.2 Provisioned Host Replacement Signal Generation	118
49	7.6.1.3 Automatic Squelching	119
50	7.6.1.4 Tributary Assignment Flexibility Restrictions	119
51	7.6.1.5 Host Lane Granularity Restrictions	119
52	7.6.1.6 Clocking Aspects	120
53	7.6.2 Data Path, Network Path, and Host Paths	120
54	7.6.3 Network Path Applications	120
55	7.6.4 Network Path Application Advertisement	121
56	7.6.5 Network Path Application Instance Configuration	121
57	7.6.5.1 Configuration and Control	121
58	7.6.5.2 Reconfiguration	122
59	7.6.6 Modifications and Extensions of Prior Specifications for Data Path Applications	123

1	7.6.6.1	Host Lanes	123
2	7.6.6.2	Data Path State Machine (DPSM)	123
3	7.6.6.3	Media Lanes	123
4	7.6.6.4	Squelching	123
5	7.6.6.5	Configuration	123
6	7.6.7	Network Path State Machines (NPSM)	124
7	7.6.7.1	NPSM State Transition Diagram and NPSM Specification	125
8	7.6.7.2	Network Path Control (Host)	126
9	7.6.7.3	Network Path Status (Module)	126
10	7.6.7.4	Detailed State Descriptions	127
11	7.6.8	Network Path Dependent Flagging Conformance	128
12	7.6.8.1	Lane-Specific Flagging Conformance per NPSM State	128
13	7.6.8.2	VDM Flagging Conformance per NPSM State	128
14	7.7	Unidirectional Hot Data Path Reconfiguration	130
15	7.8	Host Lane Switching Capability	131
16	7.8.1	Connectivity Limitations	131
17	7.8.2	Host Lane Switch Model	132
18	7.8.3	Switch Configuration and Status	132
19	7.8.4	Host Lane Switch Operation	132
20	8	Module Management Memory Map	133
21	8.1	Overview and General Specifications	133
22	8.1.1	Management Memory Structure and Mapping	133
23	8.1.2	Map of Supported Pages and Banks	134
24	8.1.3	Specifications Conventions	137
25	8.1.3.1	Reserved Locations	137
26	8.1.3.2	Custom Locations	137
27	8.1.3.3	Bank-Dependent Lane Number Interpretation	137
28	8.1.3.4	Register Default Values	138
29	8.1.3.5	Byte Order (Endianness)	138
30	8.1.3.6	Access Types	138
31	8.1.3.7	Optionality Indications	138
32	8.1.4	General Specifications	139
33	8.1.4.1	Multi-Byte Reporting Registers	139
34	8.1.4.2	Flags, Masks, and Interrupts	139
35	8.1.4.3	Generic Checksums	140
36	8.2	Lower Memory (Control and Status Essentials)	141
37	8.2.1	Management Characteristics	142
38	8.2.2	Global Status Information	144
39	8.2.3	Flags Summary	145
40	8.2.4	Module-Level Flags	146
41	8.2.5	Module-Level Monitor Values	147
42	8.2.6	Module-Level Controls	148
43	8.2.7	Module-Level Masks	150
44	8.2.8	CDB Command Status	151
45	8.2.9	Module Active Firmware Version	153
46	8.2.10	Module Fault Information	153
47	8.2.11	Miscellaneous Status Information	154
48	8.2.12	Extended Module Information	154
49	8.2.13	Applications Advertising	156
50	8.2.14	Password Entry and Change	158
51	8.2.15	Page Mapping (Upper Memory Content Selection)	159
52	8.3	Page 00h (Administrative Information)	161
53	8.3.1	SFF-8024 Identifier Copy	161
54	8.3.2	Vendor Information	161
55	8.3.2.1	Vendor Name	162
56	8.3.2.2	Vendor Organizationally Unique Identifier	162
57	8.3.2.3	Vendor Part Number	162
58	8.3.2.4	Vendor Revision Number	162
59	8.3.2.5	Vendor Serial Number	162

1	8.3.2.6	Date Code	162
2	8.3.2.7	CLEI Code	162
3	8.3.3	Module Power Characteristics	163
4	8.3.4	Cable Assembly Link Length	163
5	8.3.5	Media Connector Type	163
6	8.3.6	Copper Cable Attenuation	164
7	8.3.7	Media Lane Information	164
8	8.3.8	Cable Assembly Lane Breakout Information	165
9	8.3.9	Media Interface Technology	166
10	8.3.10	MCI Related Advertisements	167
11	8.3.11	Page 00h Page Checksum (required)	167
12	8.3.12	Custom Info (non-volatile)	167
13	8.4	Page 01h (Advertising)	168
14	8.4.1	Inactive Firmware and Hardware Revisions	168
15	8.4.2	Supported Link Length Advertisement	169
16	8.4.3	Wavelength Information	170
17	8.4.4	Supported Pages Advertisement	171
18	8.4.5	Durations Advertisement	171
19	8.4.6	Module Characteristics Advertisement	173
20	8.4.7	Supported Controls Advertisement	175
21	8.4.8	Supported Flags Advertisement	176
22	8.4.9	Supported Monitors Advertisement	176
23	8.4.10	Supported Configuration and Signal Integrity Controls Advertisement	177
24	8.4.11	CDB Messaging Support Advertisement	178
25	8.4.12	Additional Durations Advertisement	181
26	8.4.13	Normalized Application Descriptor Support	182
27	8.4.14	Media Lane Assignment Options Advertisement	182
28	8.4.15	Additional Application Advertisement	183
29	8.4.16	Miscellaneous Advertisements	184
30	8.4.17	Page Checksum (Page 01h, Byte 255, RO RQD)	184
31	8.5	Page 02h (Module and Lane Thresholds)	185
32	8.5.1	Module-Level Monitor Thresholds	185
33	8.5.2	Lane-Related Monitor Thresholds	186
34	8.5.3	Page Checksum (Page 02h, Byte 255, RO RQD)	186
35	8.6	Page 03h (User EEPROM)	187
36	8.7	Page 04h (Laser Capabilities Advertising)	188
37	8.8	Page 05h (Form Factor Specific Management Signals Management)	191
38	8.9	Banked Page 10h (Lane Control and Data Path Control)	192
39	8.9.1	Data Path Initialization Control (DPDeinit Bits)	192
40	8.9.2	Lane-Specific Direct Effect Control Fields	193
41	8.9.2.1	Tx Output Muting Functions and Their Control	193
42	8.9.2.2	Rx Output Muting Functions and Their Control	193
43	8.9.2.3	Lane-specific Tx and Rx Control Fields	194
44	8.9.2.4	Register Lists	194
45	8.9.3	Staged Control Set 0	196
46	8.9.3.1	Apply Staged Control Set Triggers (Configuration Commands)	196
47	8.9.3.2	Data Path Configuration (Application Assignments)	197
48	8.9.3.3	Tx and Rx Signal Integrity Controls	198
49	8.9.3.4	Unidirectional Apply Staged Control Set Triggers	200
50	8.9.4	Staged Control Set 1	202
51	8.9.4.1	Apply Staged Control Set Triggers	202
52	8.9.4.2	Data Path Configuration (Application Assignment)	202
53	8.9.4.3	Tx and Rx Signal Integrity Controls	203
54	8.9.4.4	Unidirectional Apply Staged Control Set Triggers	204
55	8.9.5	Lane-Specific Masks	205
56	8.10	Banked Page 11h (Lane Status and Data Path Status)	208
57	8.10.1	Data Path States	208
58	8.10.2	Lane Output Status Indications	209
59	8.10.3	Lane-Specific Flags	209
60	8.10.4	Lane-Specific Monitors	213

1	8.10.5 Configuration Command Execution and Result Status (ConfigStatus)	214
2	8.10.6 Active Control Set	216
3	8.10.6.1 Provisioned Data Path Configuration (Application Assignment)	216
4	8.10.6.2 Provisioned Tx and Rx Signal Integrity Settings	217
5	8.10.7 Data Path Conditions	219
6	8.10.8 Media Lane to Media Wavelength and Fiber Mapping	219
7	8.11 Banked Page 12h (Tunable Laser Control and Status)	221
8	8.12 Banked Page 13h (Module Performance Diagnostics Control)	223
9	8.12.1 Loopback Capabilities Advertisement	224
10	8.12.2 Diagnostics Measurement Capabilities Advertisement	225
11	8.12.3 Diagnostic Reporting Capabilities Advertisement	226
12	8.12.4 Pattern Generation and Checking Location Advertisement	227
13	8.12.5 Pattern Generation and Checking Capabilities Advertisement	228
14	8.12.6 Host Side Pattern Generator Controls	231
15	8.12.7 Media Side Pattern Generator Controls	231
16	8.12.8 Host Side Pattern Checker Controls	233
17	8.12.9 Media Side Pattern Checker Controls	234
18	8.12.10 Clocking and Measurement Controls	235
19	8.12.11 Diagnostics Measurement Behavior	237
20	8.12.11.1 Un-Gated Measurements	237
21	8.12.11.2 Gated Measurements with Global Gate Timer	237
22	8.12.11.3 Gated Measurements with Per Lane Gate Timer	238
23	8.12.12 Loopback Controls	240
24	8.12.13 Host Scratchpad Area	241
25	8.12.14 Diagnostics Masks	242
26	8.12.15 User Pattern	243
27	8.13 Banked Page 14h (Module Performance Diagnostics Results)	244
28	8.13.1 Diagnostics Selection	244
29	8.13.2 Diagnostics Flags	245
30	8.13.3 Diagnostics Data	246
31	8.14 Banked Page 15h (Timing Characteristics)	249
32	8.15 Banked Page 16h (Network Path Functionality)	250
33	8.15.1 Network Path Provisioning	251
34	8.15.2 Network Path Control	253
35	8.15.3 Network Path Commands	254
36	8.15.4 Network Path Status	256
37	8.15.5 Network Path Related Advertisements (Capabilities and Restrictions)	258
38	8.15.5.1 Maximum Durations Advertisement	258
39	8.15.5.2 Miscellaneous Options	258
40	8.15.5.3 Application Advertisement Extensions	258
41	8.15.5.4 Multiplex and Uniplex Application Advertisement	259
42	8.15.5.5 Constraints and Advertisements for Parallel NP Applications	259
43	8.16 Banked Page 17h (Flags and Masks)	262
44	8.16.1 Flags	262
45	8.16.2 Masks	262
46	8.17 Banked Page 18h (Lane Control and Data Path Control Part 2)	263
47	8.17.1 Staged Control Set Extension – Application Descriptors Block Indices	263
48	8.18 Banked Page 19h (Lane Status and Data Path Status Part 2)	265
49	8.18.1 Direction-specific Provisioned Data Path Configuration	265
50	8.18.2 Active Control Set Extension – Applications Block Select	266
51	8.19 Banked Pages Range 1Ah-1Bh (Resource Modules)	267
52	8.20 Banked Page 1Ch (Normalized Application Descriptors)	268
53	8.20.1 Normalized Application Descriptor Format	268
54	8.20.2 Normalized Application Descriptor Block	268
55	8.21 Banked Page 1Dh (Host Lane Switching)	269
56	8.22 Banked Pages Range 20h-2Fh (VDM)	271
57	8.22.1 Pages 20h-23h (VDM Descriptor Groups Pages)	273
58	8.22.1.1 VDM Instance Descriptors	273
59	8.22.1.2 VDM Observable Types	274
60	8.22.2 Pages 24h-27h (VDM Sample Groups Pages)	275

1	8.22.3	Pages 28h-2Bh (VDM Threshold Set Groups Pages)	276
2	8.22.4	Page 2Ch (VDM Flags Page)	277
3	8.22.5	Page 2Dh (VDM Masks Page)	278
4	8.22.6	Page 2Fh (VDM Advertisement and Dynamic Controls)	279
5	8.23	Banked Page 9Fh (CDB Message)	282
6	8.23.1	Triggering CDB Command on WRITE ending at 9Fh:129	284
7	8.23.2	Triggering CDB Command on WRITE including 9Fh:129	284
8	8.24	Banked Pages Range A0h-AFh (CDB Extended Payload Pages)	286
9	9	CDB Command Reference	287
10	9.1	CDB Command Group Summary	287
11	9.2	General Messaging Rules	288
12	9.2.1	Command and Reply	288
13	9.2.2	Use of Multiple CDB Instances	288
14	9.2.3	Preparing for Command Extensions	288
15	9.2.4	Preparing for Query Reply Extensions	288
16	9.3	CDB Module Commands	289
17	9.3.1	CMD 0000h: Query Status	290
18	9.3.2	CMD 0001h: Enter Password	291
19	9.3.3	CMD 0002h: Change Password	292
20	9.3.4	CMD 0004h: Abort Processing	293
21	9.4	CDB Features and Capabilities Inquiry Commands	294
22	9.4.1	CMD 0040h: Module Features	295
23	9.4.2	CMD 0041h: Firmware Management Features	296
24	9.4.3	CMD 0042h: Performance Monitoring Features	298
25	9.4.4	CMD 0043h: BERT and Diagnostics Features	299
26	9.4.5	CMD 0044h: Security Features and Capabilities	300
27	9.4.6	CMD 0045h: Externally Defined Features	302
28	9.4.7	CMD 0050h: Get Application Attributes	303
29	9.4.8	CMD 0051h: Get Interface Code Description	305
30	9.5	CDB Bulk Read Commands	306
31	9.6	CDB Bulk Write Commands	306
32	9.7	CDB Firmware Management Commands	307
33	9.7.1	CMD 0100h: Get Firmware Info	309
34	9.7.2	CMD 0101h: Start Firmware Download	311
35	9.7.3	CMD 0102h: Abort Firmware Download	312
36	9.7.4	CMD 0103h: Write Firmware Block LPL	313
37	9.7.5	CMD 0104h: Write Firmware Block EPL	314
38	9.7.6	CMD 0105h: Read Firmware Block LPL	315
39	9.7.7	CMD 0106h: Read Firmware Block EPL	316
40	9.7.8	CMD 0107h: Complete Firmware Download	317
41	9.7.9	CMD 0108h: Copy Firmware Image	318
42	9.7.10	CMD 0109h: Run Firmware Image	319
43	9.7.11	CMD 010Ah: Commit Firmware Image	320
44	9.8	CDB Performance Monitoring Commands	321
45	9.8.1	CMD 0200h: Control PM	322
46	9.8.2	CMD 0201h: Get PM Feature Information	323
47	9.8.3	CMD 0210h/0211h: Get Module PM LPL/EPL	324
48	9.8.4	CMD 0212h/0213h: Get PM Host Side LPL/EPL	326
49	9.8.5	CMD 0214h/0215h: Get PM Media Side LPL/EPL	328
50	9.8.6	CMD 0216h/0217h: Get Data Path PM LPL/EPL	330
51	9.8.7	CMD 0220h: Get Data Path RMON Statistics	332
52	9.8.8	CMD 0230h: Control FEC Symbol Error Weight Histogram	333
53	9.8.9	CMD 0231h: Get FEC Symbol Error Weight Histogram	334
54	9.8.10	CMD 0232h: Control Max FEC Symbol Error Weight	336
55	9.8.11	CMD 0233h: Get Max FEC Symbol Error Weight	337
56	9.9	CDB Data Monitoring and Recording Commands	338
57	9.9.1	CMD 0280h: Data Monitoring and Recording Controls	338
58	9.9.2	CMD 0281h: Data Monitoring and Recording Advertisement	339
59	9.9.3	CMD 0290h: Temperature Histogram	339

1	9.10 CDB PRBS BERT Commands	341
2	9.11 CDB Diagnostics and Debug Commands	342
3	9.11.1 CMD 0380h: Loopbacks	342
4	9.12 CDB Security Related Commands	343
5	9.12.1 Module Authentication Overview	344
6	9.12.2 CMD 0400h: Get Initial Device ID Certificate in LPL	345
7	9.12.3 CMD 0401h: Get Initial Device ID Certificate in EPL	346
8	9.12.4 CMD 0402h: Set Digest To Sign given in LPL	347
9	9.12.5 CMD 0403h: Set Digest To Sign given in EPL	348
10	9.12.6 CMD 0404h: Get Digest Signature in LPL	349
11	9.12.7 CMD 0405h: Get Digest Signature in EPL	350
12	10 Management Timing Specifications	351
13	10.1 Timings for Management Control Signals (MSL)	351
14	10.2 Timings for Register Access (RAL)	352
15	10.2.1 Single ACCESS Timings	352
16	10.2.2 Consecutive ACCESS Timings	352
17	10.2.3 Register Content Dependencies	353
18	10.3 Timings between Conditions and Management Registers or Signals	354
19	10.3.1 Interrupt and Flag Related Timings	354
20	10.3.2 High Speed Signal Related Timings	354
21	10.4 Timings between High-Speed Signal Conditions	355
22	Appendix A Form Factor Specific MSL or MCI Signal Names	356
23	Appendix B Management Communication Interface Definitions	357
24	B.1 Generic Definitions	357
25	B.1.1 Communication Roles and Transactions	357
26	B.1.2 Current Byte Address	358
27	B.2 I2C-Based Management Communication Interface (I2CMCI)	359
28	B.2.1 Communication Topology	359
29	B.2.2 I2CMCI Control Signals	359
30	B.2.3 Physical Layer Signals	359
31	B.2.4 Serial Communication Protocol	359
32	B.2.4.1. Basic Definitions and Protocol Elements	360
33	B.2.4.1.1. Start Condition (START)	360
34	B.2.4.1.2. Stop Condition (STOP)	360
35	B.2.4.1.3. Word Size (Byte) and Bit Serial Transmission Order	360
36	B.2.4.1.4. Basic Operation Encoding (Control Byte)	360
37	B.2.4.1.5. Acknowledge (ACK and NACK)	360
38	B.2.4.2. Protocol Reset and Recovery	360
39	B.2.4.2.1. Power-On Reset	360
40	B.2.4.2.2. Protocol Reset and Recovery	360
41	B.2.5 Serial MCI Transactions for READ/WRITE/TEST Access	361
42	B.2.5.1. Transactions for a Byte Read Operation	362
43	B.2.5.1.1. Transaction for a Current Address Read Operation	362
44	B.2.5.1.2. Transaction for a Random Read Operation	363
45	B.2.5.2. Transaction for a Sequential Bytes Read Operation	364
46	B.2.5.2.1. Transaction for a Sequential Bytes Read from Current Start Address	364
47	B.2.5.2.2. Transaction for a Sequential Bytes Read from Random Start Address	365
48	B.2.5.3. Transaction for a Single Byte Write Operation	366
49	B.2.5.4. Transaction for a Sequential Bytes Write Operation	367
50	B.2.5.5. Transaction for a Test Operation	368
51	B.2.6 Transaction Flow Control Mechanisms	369
52	B.2.6.1. Delaying Current Transaction (Clock Stretching)	369
53	B.2.6.2. Rejecting Subsequent Transaction (Transaction Hold-Off)	369
54	B.2.7 Timing Specifications	370
55	B.2.7.1. Module Select Timings	370
56	B.2.7.2. Waveform Timings	370
57	B.2.7.3. Transaction Timings	371
58	B.2.7.3.1. tWR Timing	372

1	B.2.7.3.2. tNACK Timing	373
2	B.3 SPI-Based Management Communication Interface (SPIMCI)	374
3	B.3.1 Introduction	374
4	B.3.2 Communication Topology	374
5	B.3.3 SPIMCI Control Signals	374
6	B.3.4 Physical Layer Signals	375
7	B.3.5 Communication Speed	375
8	B.3.6 Physical Encoding	375
9	B.3.7 Serial Communication Protocol	376
10	B.3.7.1. Bus Transactions	376
11	B.3.7.1.1. Transaction Control Phase	377
12	B.3.7.1.2. Flow Control Phase (N Bytes)	377
13	B.3.7.1.3. Data Transfer Phase (M Bytes)	378
14	B.3.7.2. Basic Definitions and Protocol Elements	378
15	B.3.7.2.1. Start Condition (START):	378
16	B.3.7.2.2. Stop Condition (STOP):	378
17	B.3.7.2.3. Bus Word Size (Byte) and Bit Serial Transmission Order	378
18	B.3.7.2.4. Acknowledge (ACK and NACK)	378
19	B.3.7.2.5. TEST	378
20	B.3.7.3. Protocol Reset and Recovery	378
21	B.3.7.3.1. Power-On Reset	378
22	B.3.7.3.2. Protocol Violation	378
23	B.3.7.3.3. Protocol Reset and Recovery	378
24	B.3.8 SPI MCI Transactions for READ/WRITE/TEST Access	379
25	B.3.8.1. Read One Byte from Given Byte Address	379
26	B.3.8.2. Read n Bytes from Given Byte Address	380
27	B.3.8.3. Write 1 Byte to a Given Byte Address	381
28	B.3.8.4. Write n Bytes to a Given Byte Address	382
29	B.3.8.5. Test Command	383
30	B.3.9 Transaction Flow Control Mechanism	383
31	B.3.9.1. Stretching Current Transaction	383
32	B.3.9.2. Rejecting Subsequent Transaction (Transaction Hold-Off)	383
33	B.3.10 Timing Specification	383
34	Appendix C Examples of Application Advertisements	384
35	Appendix D Examples of Initialization and Deinitialization	388
36	D.1 Initialization Examples	388
37	D.1.1 Quick Hardware Initialization	388
38	D.1.2 Quick Software Initialization	389
39	D.1.3 Software Configuration and Initialization	390
40	D.2 Deinitialization Examples	394
41	D.2.1 Hardware Deinitialization	394
42	D.2.2 Software Deinitialization	395
43	Appendix E Illustration of Applying Control Sets	396
44	E.1 Default Behavior (SteppedConfigOnly = 0)	396
45	E.2 Restricted Behavior (SteppedConfigOnly = 1)	396
46	Appendix F Examples of Diagnostic Features Usage	397
47	F.1 Enabling and Disabling Pattern Generator (Host or Media Side)	397
48	F.2 Enabling and Disabling Pattern Checker (Host or Media Side)	397
49	F.3 Reading Pattern Checker Error Counters	397
50	F.3.1 Not Gated (Continuous) Error Counters, Individual Lanes	398
51	F.3.2 Not Gated (Continuous) Error Counters, Individual Lanes, Reset Error Counter	398
52	F.3.3 Not Gated (Continuous) Error Counters, All Lanes, all Banks	399
53	Appendix G Specification Evolution and Maintenance Notes	400
54	G.1 Definitions 400	400
55	G.2 Cross-Version Compatibility	400
56	G.3 Interpretation of CMIS Version Numbers	400

1	Appendix H Examples for Network Path Applications	401
2	H.1 Advertisement Examples	401
3	H.1.1 400G Module for 400ZR DP and NP Application supporting Homogeneous Multiplex	401
4	H.1.2 400G Module for 400ZR NP Application supporting Mixed Multiplex	402
5	H.1.3 400G Module with Alternative Support of 400G or 200G NP Application	403
6	H.1.4 800G Module with Parallel 400ZR NP or DP Applications	403
7	H.1.5 800G Module for 400ZR NP Application and Parallel DP Applications	404
8	H.2 Provisioning Examples	405
9	H.2.1 4x100G NP Application with Unused Multiplexing Slots	405
10	H.2.2 800G Module with Parallel 400ZR NP Applications	406
11	H.2.3 800G Module for 400ZR NP Application and Other Parallel DP Applications	406
12	Appendix I Companies Belonging to OIF at Time of Approval	407

Figures

15	Figure 4-1 CMIS Management Protocols and Layers	49
16	Figure 6-1 Lane Assignment Example (Case of Parallel DataPaths)	57
17	Figure 6-2 Control Set Data Flow Diagram	63
18	Figure 6-3 Paged Memory Module State Machine (MSM) State Transition Diagram	75
19	Figure 6-4 Flat Memory Module State Machine (MSM) State Transition Diagram	78
20	Figure 6-5 Data Path State Machine (DPSM) State Transition Diagram	85
21	Figure 7-1 Optical ingress path of Module	100
22	Figure 7-2 PAM4 amplitude histogram	100
23	Figure 7-3 BER short term measurements and interval statistics (example)	103
24	Figure 7-4 Firmware Update using CDB	110
25	Figure 7-5 Firmware Upload using CDB	112
26	Figure 7-6 Network Path State Machine (NPSM) State Transition Diagram	125
27	Figure 7-7 Host Lane Switching Model and Example	131
28	Figure 8-1 CMIS Module Memory Map (Conceptual View)	133
29	Figure 8-2 CMIS Bank and Page Group Iconic Memory Map Overview	135
30	Figure 8-3 Loopback Type Illustrations	224
31	Figure 8-4 PRBS Paths Reference Diagram	227
32	Figure 9-1 Module Authentication Flow	344
33	Figure B-1 Current Address Read	362
34	Figure B-2 Random Read	363
35	Figure B-3 Sequential Bytes Read Starting at Current Address	364
36	Figure B-4 Sequential Bytes Read Starting with Random Read	365
37	Figure B-5 Write Byte Transaction	366
38	Figure B-6 Sequential Bytes Write Transaction	367
39	Figure B-7 Test Readiness Transaction	368
40	Figure B-8 Test Readiness Transaction	369
41	Figure B-9 I2CMCI Waveform Timing Diagram	370
42	Figure B-10 tWR bus timing	372
43	Figure B-11 tNACK bus timing	373
44	Figure B-12 SPIMCI Physical Layer Signals	375
45	Figure B-13 SPIMCI Bus Transactions (general)	376
46	Figure B-14 SPIMCI READ Transaction	379
47	Figure B-15 SPIMCI READ N Transaction	380
48	Figure B-16 SPIMCI WRITE Transaction	381
49	Figure B-17 SPIMCI WRITE N Transaction	382
50	Figure B-18 SPIMCI TEST Transaction	383
51	Figure E-19 ApplyDPInit (default: SteppedConfigOnly=0)	396
52	Figure E-20 ApplyDPInit (restricted: SteppedConfigOnly=1)	396

Tables

56	Table 6-1 Application Descriptor structure	59
57	Table 6-2 Control fields activated by ExplicitControl bit	63
58	Table 6-3 Configuration Commands (Intervention-Free Reconfiguration Procedures Supported)	67

1	Table 6-4 Configuration Commands (Intervention-Free Reconfigurations Not Supported)	68
2	Table 6-5 Tx Input Eq control relationship to AdaptiveInputEqEnableTx	70
3	Table 6-6 Host-Controlled Tx Input Equalization Codes	70
4	Table 6-7 Rx Output Equalization Codes	71
5	Table 6-8 Rx Output Amplitude Codes	71
6	Table 6-9 ModuleStateChangedFlag behaviors	74
7	Table 6-10 Module State Machine exit condition priority	75
8	Table 6-11 ResetS transition signal truth table	75
9	Table 6-12 LowPwrS transition signal truth table	76
10	Table 6-13 LowPwrExS transition signal truth table	76
11	Table 6-14 Module state behaviors, paged memory modules	77
12	Table 6-15 Module state behaviors, flat memory modules	78
13	Table 6-16 DPDeinitS transition signal truth table	85
14	Table 6-17 DPReDeinitS transition signal truth table (default)	86
15	Table 6-18 Data Path state behaviors and Exit Conditions	87
16	Table 6-19 Data Path State Changed Flag behaviors	89
17	Table 6-20 Module Flag Conformance Rules	94
18	Table 6-21 Lane-Specific Flagging Conformance Rules	95
19	Table 6-22 VDM Flag Conformance Rules	96
20	Table 7-1 Status of an Image Storage Bank	113
21	Table 7-2 Typical FirmwareStatus Codes of Modules supporting image A and B.	113
22	Table 7-3 Typical FirmwareStatus Codes of Modules supporting image A	113
23	Table 7-4 Special FirmwareStatus Codes	114
24	Table 7-5 Network Path State Changed Flag behaviors	127
25	Table 7-6 Lane-Specific Flagging Conformance Rules per NPSM State	128
26	Table 7-7 VDM Flag Conformance Rules per NPSM State	129
27	Table 8-1 List of CMIS Pages	136
28	Table 8-2 Bank Dependent Lane Number Interpretation	137
29	Table 8-3 Access Types	138
30	Table 8-4 Lower Memory Overview	141
31	Table 8-5 Management Characteristics (Lower Memory)	142
32	Table 8-6 Global Status Information (Lower Memory)	144
33	Table 8-7 Module State Encodings	144
34	Table 8-8 Lane-Level Flags Summary (Lower Memory)	145
35	Table 8-9 Module Flags (paged memory modules only) (Lower Memory)	146
36	Table 8-10 Module-Level Monitor Values (paged memory modules only) (Lower Memory)	147
37	Table 8-11 Module Global Controls (paged memory modules only) (Lower Memory)	148
38	Table 8-12 Module Level Masks (paged memory modules only) (Lower Memory)	150
39	Table 8-13 CdbStatus fields (paged memory modules only) (Lower Memory)	151
40	Table 8-14 Bit definitions within CdbStatus fields	151
41	Table 8-15 Module Active Firmware Version (Lower Memory)	153
42	Table 8-16 Fault Information (paged memory modules only) (Lower Memory)	153
43	Table 8-17 Miscellaneous Status Information (Lower Memory)	154
44	Table 8-18 Extended Module Information (Lower Memory)	154
45	Table 8-19 Low Power Restrictions Byte	155
46	Table 8-20 Media Type Encodings	156
47	Table 8-21 Media Type Register (Lower Memory)	157
48	Table 8-22 Format of Application Descriptor Bytes 1-4	157
49	Table 8-23 Application Descriptor Registers Bytes 1-4 (Lower Memory)	157
50	Table 8-24 Password Change Entry (Lower Memory)	159
51	Table 8-25 Page Mapping Register Components (Lower Memory)	159
52	Table 8-26 Page 00h Overview	161
53	Table 8-27 SFF8024IdentifierCopy (Byte 00h:128)	161
54	Table 8-28 Vendor Information (Page 00h)	161
55	Table 8-29 Date Code (Page 00h)	162
56	Table 8-30 CLEI Code (Page 00h)	162
57	Table 8-31 Module Power Class and Max Power (Page 00h)	163
58	Table 8-32 Cable Assembly Link Length (Page 00h)	163
59	Table 8-33 Media Connector Type (Page 00h)	163
60	Table 8-34 Copper Cable Attenuation (Page 00h)	164

1	Table 8-35 Media Lane Information (Page 00h)	164
2	Table 8-36 Cable Assembly Information (Page 00h)	165
3	Table 8-37 Far End Configurations for Homogeneous Far End Breakout (Page 00h)	165
4	Table 8-38 Far End Configurations for up to 8 Near End Lanes (Page 00h)	165
5	Table 8-39 Media Connector Type (Page 00h)	166
6	Table 8-40 Media Interface Technology encodings	166
7	Table 8-41 MCI Related Advertisements (Page 00h)	167
8	Table 8-42 Page 01h Overview	168
9	Table 8-43 Module Inactive Firmware and Hardware Revisions (Page 01h)	169
10	Table 8-44 Supported Fiber Link Length (Page 01h)	169
11	Table 8-45 Wavelength Information (Page 01h)	170
12	Table 8-46 Supported Pages Advertising (Page 01h)	171
13	Table 8-47 Durations Advertising (Page 01h)	171
14	Table 8-48 State Duration Encoding (Page 01h)	172
15	Table 8-49 Module Characteristics Advertisement (Page 01h)	173
16	Table 8-50 Supported Controls Advertisement (Page 01h)	175
17	Table 8-51 Supported Flags Advertisement (Page 01h)	176
18	Table 8-52 Supported Monitors Advertisement (Page 01h)	176
19	Table 8-53 Supported Signal Integrity Controls Advertisement (Page 01h)	177
20	Table 8-54 CDB Advertisement (Page 01h)	178
21	Table 8-55 Overview of CDB advertising combinations	180
22	Table 8-56 Additional Durations Advertising (Page 01h)	181
23	Table 8-57 Normalized Application Descriptors Support (Page 01h)	182
24	Table 8-58 Media Lane Assignment Advertising (Page 01h)	182
25	Table 8-59 Additional Application Descriptor Registers (Page 01h)	183
26	Table 8-60 Miscellaneous Advertisements (Page 01h)	184
27	Table 8-61 Page 02h Overview	185
28	Table 8-62 Module-Level Monitor Thresholds (Page 02h)	185
29	Table 8-63 Lane-Related Monitor Thresholds (Page 02h)	186
30	Table 8-64 Page 03h Overview	187
31	Table 8-65 Page 04h Overview	188
32	Table 8-66 Laser capabilities for tunable lasers (Page 04h)	188
33	Table 8-67 Page 10h Overview	192
34	Table 8-68 Data Path initialization control (Page 10h:128)	192
35	Table 8-69 Lane-specific Direct Effect Control Fields (Page 10h)	194
36	Table 8-70 Staged Control Set 0, Apply Triggers (Page 10h)	197
37	Table 8-71 Data Path Configuration per Lane (DPConfigLane<i>)	197
38	Table 8-72 Staged Control Set 0, Data Path Configuration (Page 10h)	198
39	Table 8-73 Staged Control Set 0, Tx Controls (Page 10h)	199
40	Table 8-74 Staged Control Set 0, Rx Controls (Page 10h)	200
41	Table 8-75 Staged Control Set 0, Unidirectional Apply Triggers (Page 10h)	201
42	Table 8-76 Staged Control Set 1, Apply Triggers (Page 10h)	202
43	Table 8-77 Staged Control Set 1, Data Path Configuration (Page 10h)	202
44	Table 8-78 Staged Control Set 1, Tx Controls (Page 10h)	203
45	Table 8-79 Staged Control Set 1, Rx Controls (Page 10h)	203
46	Table 8-80 Staged Control Set 1, Unidirectional Apply Triggers (Page 10h)	204
47	Table 8-81 Lane-Specific Masks (Page 10h)	205
48	Table 8-82 Page 11h Overview	208
49	Table 8-83 Lane-associated Data Path States (Page 11h)	208
50	Table 8-84 Data Path State Encoding	208
51	Table 8-85 Lane-Specific Output Status (Page 11h)	209
52	Table 8-86 Lane-Specific State Changed Flags (Page 11h)	210
53	Table 8-87 Lane-Specific Tx Flags (Page 11h)	210
54	Table 8-88 Rx Flags (Page 11h)	212
55	Table 8-89 Media Lane-Specific Monitors (Page 11h)	213
56	Table 8-90 Configuration Command Status registers (Page 11h)	215
57	Table 8-91 Configuration Command Execution and Result Status Codes (Page 11h)	215
58	Table 8-92 Data Path Configuration per Lane (DPConfigLane<i> Field)	216
59	Table 8-93 Active Control Set, Provisioned Data Path Configuration (Page 11h)	216
60	Table 8-94 Active Control Set, Provisioned Tx Controls (Page 11h)	217

1	Table 8-95 Active Control Set, Provisioned Rx Controls (Page 11h)	218
2	Table 8-96 Data Path Conditions (Page 11h)	219
3	Table 8-97 Media Lane to Media Wavelength and Fiber mapping (Page 11h)	219
4	Table 8-98 Page 12h Overview	221
5	Table 8-99 Laser tuning, status, and Flags for tunable transmitters (Page 12h)	221
6	Table 8-100 Page 13h Overview	223
7	Table 8-101 Loopback Capabilities (Page 13h)	225
8	Table 8-102 Diagnostics Measurement Capabilities (Page 13h)	225
9	Table 8-103 Diagnostic Reporting Capabilities (Page 13h)	226
10	Table 8-104 Pattern Generation and Checking Location (Page 13h)	227
11	Table 8-105 Pattern IDs	228
12	Table 8-106 PRBS Pattern Generation Capabilities (Page 13h)	228
13	Table 8-107 Pattern Checking Capabilities (Page 13h)	229
14	Table 8-108 Pattern Generator and Checker swap and invert Capabilities (Page 13h)	229
15	Table 8-109 Host Side Pattern Generator Controls (Page 13h)	231
16	Table 8-110 Host Side Pattern Generator Pattern Select Controls (Page 13h)	231
17	Table 8-111 Media Side Pattern Generator Controls (Page 13h)	232
18	Table 8-112 Media Side Pattern Generator Pattern Select Controls (Page 13h)	232
19	Table 8-113 Host Side Pattern Checker Controls (Page 13h)	233
20	Table 8-114 Host Side Pattern Checker Pattern Select Controls (Page 13h)	233
21	Table 8-115 Media Side Pattern Checker Controls (Page 13h)	234
22	Table 8-116 Media Side Pattern Checker Select Controls (Page 13h)	234
23	Table 8-117 Clocking and Measurement Controls (Page 13h)	235
24	Table 8-118 PRBS Checker Behavior Un-Gated Mode	237
25	Table 8-119 PRBS Checker Behavior Single Gate Timer	238
26	Table 8-120 PRBS Checker Behavior Per Lane Gate Timer	239
27	Table 8-121 Loopback Controls (Page 13h)	240
28	Table 8-122 Host Scratchpad Area (Page 13h)	241
29	Table 8-123 Diagnostics Masks (Page 13h)	242
30	Table 8-124 User Pattern (Page 13h)	243
31	Table 8-125 Page 14h Overview	244
32	Table 8-126 Diagnostics Selection Register (Page 14h)	244
33	Table 8-127 Diagnostics Selector Options	244
34	Table 8-128 Latched Diagnostics Flags (Page 14h)	245
35	Table 8-129 Diagnostics Data (Bytes 192-255) Contents per Diagnostics Selector (Page 14h)	246
36	Table 8-130 Page 15h Overview	249
37	Table 8-131 Data Path Rx and Tx Latency, per lane (Page 15h)	249
38	Table 8-132 Page 16h Overview	250
39	Table 8-133 Network Path Provisioning per Lane (NPConfigLane<i>)	251
40	Table 8-134 Staged Control Set 0, Network Path Configuration (Page 16h)	252
41	Table 8-135 Staged Control Set 1, Network Path Configuration (Page 16h)	252
42	Table 8-136 Network Path Initialization Control (Page 16h)	253
43	Table 8-137 Network and Host Path Signal Source Selection (Page 16h)	253
44	Table 8-138 Staged Control Set 0, Apply Triggers (Page 16h)	254
45	Table 8-139 Staged Control Set 1, Apply Triggers (Page 16h)	254
46	Table 8-140 NP Configuration Command Status registers (Page 16h)	255
47	Table 8-141 NP Configuration Command Execution and Result Status Codes (Page 16h)	255
48	Table 8-142 NP Active Control Set, Network Path Configuration (Page 16h)	256
49	Table 8-143 Lane-associated Network Path States (Page 16h)	256
50	Table 8-144 Network Path State Encoding	256
51	Table 8-145 Network Path Conditions (Page 16h)	257
52	Table 8-146 NPSM Durations Advertising (Page 16h)	258
53	Table 8-147 Miscellaneous Options (Page 16h)	258
54	Table 8-148 NP Extended Application Advertisement (Page 16h)	259
55	Table 8-149 Multiplex Lane Grouping Advertisement	260
56	Table 8-150 Multiplex Granularities Advertisement (Page 16h)	261
57	Table 8-151 Global Multiplex Structures Advertisement (Page 16h)	261
58	Table 8-152 Page 17h Overview	262
59	Table 8-153 Network Path Related Flags (Page 17h)	262
60	Table 8-154 Network Path Related Masks (Page 17h)	262

1	Table 8-155 Page 18h Overview	263
2	Table 8-156 Staged Control Set 0 – Normalized Application Descriptor Block Indices (Page 18h)	263
3	Table 8-157 Staged Control Set 1 – Normalized Application Descriptor Block Indices (Page 18h)	263
4	Table 8-158 Page 19h Overview	265
5	Table 8-159 Active Control Set, Provisioned Tx Data Path Configuration (Page 19h)	265
6	Table 8-160 Active Control Set, Provisioned Rx Data Path Configuration (Page 19h)	266
7	Table 8-161 Active Control Set – Normalized Application Descriptor Block Indices (Page 19h)	266
8	Table 8-162 Page 1Ch Overview	268
9	Table 8-163 Normalized Application Descriptor (NAD) Structure (Page 1Ch)	268
10	Table 8-164 Normalized Application Descriptor Block (Page 1Ch)	268
11	Table 8-165 Page 1Dh Overview	269
12	Table 8-166 Host Lane Switching (Page 1Dh)	269
13	Table 8-167 Summary of Page Definitions for Page 20h-2Fh	271
14	Table 8-168 VDM Configuration (Page 20h-23h)	273
15	Table 8-169 Definition of 2-byte VDM Instance Descriptor	273
16	Table 8-170 VDM Observable Types (Type Coding)	274
17	Table 8-171 VDM Real-Time Values (Page 24h-27h)	275
18	Table 8-172 VDM Alarm/Warning Thresholds (Page 28h-2Bh)	276
19	Table 8-173 VDM Threshold Crossing (TC) Flags Byte	277
20	Table 8-174 VDM Alarm and Warning Configuration (Page 2Ch)	277
21	Table 8-175 VDM ThresholdSet0 to 15 Alarm and Warning Configuration (Page 2Dh)	278
22	Table 8-176 VDM Advertisement and Control Registers Summary (Page 2Fh)	280
23	Table 8-177 Page 9Fh Overview (CDB Message)	282
24	Table 8-178 CDB Command Message Header (Page 9Fh)	282
25	Table 8-179 CDB Reply Message Header (Page 9Fh)	283
26	Table 8-180 CDB Message Body (Page 9Fh)	283
27	Table 8-181 EPL Segments (Pages A0h-AFh)	286
28	Table 9-1 CDB Command Groups	287
29	Table 9-2 CDB Module Commands Summary	289
30	Table 9-3 CDB Command 0000h: Query Status	290
31	Table 9-4 CDB Command 0001h: Enter Password	291
32	Table 9-5 CDB Command 0002h: Change Password	292
33	Table 9-6 CDB Command 0004h: Abort	293
34	Table 9-7 CDB Feature and Capabilities Commands Overview	294
35	Table 9-8 CDB Command 0040h: Module Features	295
36	Table 9-9 CDB Command 0041h: Firmware Management Features	296
37	Table 9-10 CDB Command 0042h: Performance Monitoring Features	298
38	Table 9-11 CDB Command 0043h: BERT and Diagnostics Features	299
39	Table 9-12 CDB Command 0044h: Security Features and Capabilities	300
40	Table 9-13 CDB Command 0045h: Externally Defined Features	302
41	Table 9-14 CDB Command 0050h: Get Application Attributes	303
42	Table 9-15 CDB Command 0051h: Get Interface Code Description	305
43	Table 9-16 CDB Bulk Read Commands Overview	306
44	Table 9-17 CDB Bulk Write Commands Overview	306
45	Table 9-18 CDB Firmware Download Commands Overview	307
46	Table 9-19 CDB Command 0100h: Get Firmware Info	309
47	Table 9-20 CDB Command 0101h: Start Firmware Download	311
48	Table 9-21 CDB Command 0102h: Abort Firmware Download	312
49	Table 9-22 CDB Command 0103h: Write Firmware Block LPL	313
50	Table 9-23 CDB Command 0104h: Write Firmware Block EPL	314
51	Table 9-24 CDB Command 0105h: Read Firmware Block LPL	315
52	Table 9-25 CDB Command 0106h: Read Firmware Block EPL	316
53	Table 9-26 CDB Command 0107h: Complete Firmware Download	317
54	Table 9-27 CDB Command 0108h: Copy Firmware Image	318
55	Table 9-28 CDB Command 0109h: Run Firmware Image	319
56	Table 9-29 CDB Command 010Ah: Commit Image	320
57	Table 9-30 CDB Performance Monitoring Commands Overview	321
58	Table 9-31 CDB Performance Monitoring Observables Cross Reference	321
59	Table 9-32 CDB Command 0200h: Control PM	322
60	Table 9-33 CDB Command 0201h: Get PM Feature Information	323

1	Table 9-34 CDB Command 0210h/0211h: Get Module PM LPL/EPL	324
2	Table 9-35 CDB Command 0212h/0213h: Get PM Host Side LPL/EPL	326
3	Table 9-36 CDB Command 0214h/0215h: Get PM Media Side LPL/EPL	328
4	Table 9-37 CDB Command 0216/0217h: Get Data Path PM LPL/EPL	330
5	Table 9-38 CDB Command 0220h: Get Data Path RMON Statistics	332
6	Table 9-39 CDB Command 0230h: Control FEC Symbol Error Weight Histogram	333
7	Table 9-40 CDB Command 0231h: Get FEC Symbol Error Weight Histogram	334
8	Table 9-41 CDB Command 0232h: Control Max FEC Symbol Error Weight	336
9	Table 9-42 CDB Command 0233h: Get Max FEC Symbol Error Weight	337
10	Table 9-43 CDB Data Monitoring and Recording Commands Overview	338
11	Table 9-44 CDB Command 0280h: Data Monitoring and Recording Controls	338
12	Table 9-45 CDB Command 0281h: Data Monitoring and Recording Advertisements	339
13	Table 9-46 CDB Command 0290h: Temperature Histogram	339
14	Table 9-47 CDB BERT Commands Overview	341
15	Table 9-48 CDB Diagnostics and Debug Commands Overview	342
16	Table 9-49 CDB Command 0380h: Loopbacks	342
17	Table 9-50 CDB Diagnostics and Debug Commands Overview	343
18	Table 9-51 Detailed Return Status Codes for Security Commands	344
19	Table 9-52 CDB Command 0400h: Get Initial Device ID Certificate in LPL	345
20	Table 9-53 CDB Command 0401h: Get Initial Device ID Certificate in EPL	346
21	Table 9-54 CDB Command 0402h: Set Digest To Sign given in LPL	347
22	Table 9-55 CDB Command 0403h: Set Digest To Sign given in EPL	348
23	Table 9-56 CDB Command 0404h: Get Digest Signature in LPL	349
24	Table 9-57 CDB Command 0405h: Get Digest Signature in EPL	350
25	Table 10-1 Signal Waveform Timings	351
26	Table 10-2 Effect Latency Timings	351
27	Table 10-3 Timings for Register Access	352
28	Table 10-4 Maximum ACCESS Hold-Off Durations	352
29	Table 10-5 Content Dependency Timings	353
30	Table 10-6 Condition to Interrupt Timings	354
31	Table 10-7 Register to Interrupt Timings	354
32	Table 10-8 Register to High-Speed Signal Timings	354
33	Table 10-9 Signal Condition to Signal Condition Timings	355
34	Table 10-10 Module Select Timings	370
35	Table A-1 Form Factor Dependent Signal Name Associations	356
36	Table A-2 Symbolic Logical Signal Values	356
37	Table B-1 Management Communication Interface Variants	357
38	Table B-2 I2CMCI Transactions	361
39	Table B-3 I2CMCI Waveform Timing Parameters	371
40	Table B-4 Flow Control Timings	371
41	Table C-1 400GBASE-DR4 Transceiver with Dual Application Advertising	384
42	Table C-2 400GBASE-SR8 Fixed Transceiver Application Advertising	385
43	Table C-3 400GBASE-SR8 Transceiver supporting 200GBASE-SR4, 100GBASE-SR2 and 50GBASE-SR	386
44	Table C-4 8x50G AOC Application Advertising Example	387
45	Table H-1 400ZR NP Application Advertisement Example	401
46	Table H-2 400ZR NP Application Advertisement Mixed Multiplex Example	402
47	Table H-3 Global Multiplex Structure Advertisement	402
48	Table H-4 Multiple Multiplex Granularities Advertisement Example	403
49	Table H-5 2x400ZR NP Application Advertisement Example	403
50	Table H-6 400ZR + 400G-DR4 or 4x100G-DR1 Application Advertisement Example	404
51	Table H-7 400ZR NP Provisioning Example	405
52	Table H-8 2 x 400ZR NPs Provisioning Example	406
53	Table H-9 400ZR + 4x100G-DR1 NP Provisioning Example	406

1 Introduction

1.1 Purpose and Scope

This **Common Management Interface Specification** (CMIS) defines a generic management communication interface together with a generic management interaction protocol between hosts and managed modules.

The target audience of this specification includes suppliers of modules and transceivers, system manufacturers, and system integrators.

1.1.1 CMIS Specification (Base)

This specification, **CMIS**, has been developed to allow host and module software implementers to utilize a common code base across a variety of form factors and across a variety of transmission module¹ capabilities, and to foster the possibility of **vendor agnostic** management for **standardized** module functions.

To this end this document specifies a small **core** of management features that all modules must consider implementing (depending on module capabilities such as the level of programmability), and a larger evolving set of optional **advanced** features whose implementation is advertised in the so-called management memory map² of a module. This **advertisement approach** allows host software to adapt to optional module capabilities while ensuring interoperability with all modules at a basic level.

Characteristic and common to all CMIS compliant modules is that a well-defined set of management operations and associated data are transferred over a CMIS defined Management Communication Interface (MCI), e.g. a two-wire communication interface. The basic management operations are simple and allow the host to access a 256 byte addressable memory window, with mechanisms to dynamically switch 128 byte sized data pages of a much larger management memory space into the upper half of that host addressable memory window.

Note: This limited set of basic operations and the very small byte-oriented memory window are traced back to earlier SFF specification and allow also simple transducers or transceivers to be CMIS managed. For complex modules, extension mechanisms are implemented on top of these basic elements.

The **physical form factor scope** of CMIS includes pluggable or onboard form factors such as e.g. QSFP-DD, OSFP, or COBO. However, CMIS is developed as a generic management interface specification and can be implemented in a variety of existing form factors, such as QSFP, or also in future form factors. Generic advertisement fields in the management memory map inform the host about the form factor, the functional type of the module, and whether a module can be managed in a CMIS compliant fashion.

Note: Organizations working on new module developments are invited to use CMIS and contact the CMIS editors or the CMIS publishing organization for adding support and adapting CMIS incrementally, as needed.

The **functional scope** of CMIS includes module types which may range from electrical cable assemblies (hereafter also referred to as modules, unless cable assemblies are specifically mentioned) and active transceiver modules to versatile coherent DWDM modules with integrated framers.

The following classification distinguishes transmission modules or applications³ by the type of functionality they provide to the host system:

1. **data agnostic system interface modules** ("basic modules") map signals from host lanes to media lanes and vice versa, without knowledge of data formats and without participation in any communication protocol for that bit stream. Examples include cable assemblies and transceivers at not too high lane data rates, e.g. 100GBASE-SR4 modules
2. **data format aware system interface modules** ("complex modules") perform interface related single or multi-lane data processing (such as lane deskewing and FEC coding), e.g. 400ZR modules
3. **client encapsulation modules** ("multiplex modules") encapsulate one or more (single or multi-lane) host signals into a newly framed (single or multi-lane) network signal that may be transmitted and monitored independent of whether a host signal is present or whether a service is transported. Such modules employ framers with additional overhead for independent **media side data link termination**, encapsulating host signals as payload and comprising functionality like framing⁴, mapping, aggregation (multiplexing), switching, or inverse multiplex.

¹ In CMIS 5.3 so called **resource modules** have been added, which do not implement a transmission path.

² The management memory map defines registers and memory locations that are accessible to the host.

³ Versatile modules may be programmed to behave like modules of different functional types.

⁴ Note that system interfaces employing network side forward error correction (FEC) merely for media channel enhancement, not for independent network link operation, are not considered to be client encapsulating.

The *specification scope of this CMIS revision* covers both system interface modules and client encapsulation modules with at most (multiples of) eight host lanes and with management communication based on one of the Management Communication Interface definitions described in Appendix B.

The extensions for modules with more than 32 (4x8) lanes, or for more complex system interface or client encapsulation modules are left to future revisions of this specification, or to external extension specifications.

1.1.2 CMIS Supplement Specifications

Management specifications that are applicable to a significant but non-prevalent subset of modules are maintained as separate CMIS supplement specifications.

These CMIS **supplements rely on** the CMIS **base** specification and describe extensions or (sometimes) restrictions for the affected subclass of CMIS managed modules.

Due to the assumed pluggability of a CMIS managed module, a CMIS supplement cannot modify or restrict mandatory specifications from the CMIS base specification, by default, because a module in the scope of a CMIS supplement specification must still behave as a standard CMIS module, by default, when plugged into an arbitrary CMIS host¹.

The following CMIS Supplement Specifications exist, or are in preparation:

- **CMIS-VCS** [6] defines signal integrity (SI) parameters that can be part of a so-called Versatile Control Set, generalizing and extending the original, fixed set of signal integrity controls described in section 6.2.3. Note that this functionality, when enabled, may "override" certain fields of the standard memory map, which, if applicable, will be specified in [6].
- **C-CMIS** [7] provides specifications for modules with coherent transmission capabilities.
- **CMIS-FF** [8] provides specifications of form factor specific management facilities typically implemented by additional and sometimes programmable discrete value hardware signals² that are available as per individual module form factor hardware specification³.
- **CMIS-LT** [9] provides specifications in support of host to module electrical link training. More precisely, CMIS-LT specifies both data structures and mechanisms to emulate a bidirectional message exchange between the electrical link endpoints of a host lane for the purpose of link training (LT). Note that the activities and behaviors of the acting link endpoints that reside in both host and module and that actively use those facilities are specified elsewhere.
- **CMIS-ELSFP** [10] provides specifications for external laser source resource modules.

¹ It is, however, very well possible that a supplement allows the host to enable new behaviors that are in contradiction to standard default specifications. The key point is that a host must be aware of and willing to use any such modifications specified in a supplement.

² In the parlance of chapter 5, CMIS-FF is mostly concerned with form factor specific extensions to the CMIS management signaling layer (MSL), where certain management functions are associated with a small set of electrical 'low-speed' signals.

³ The industry trend moves away from management by non-scalable hardware signals, but widely deployed legacy form factors do offer such signals.

1.2 Document Overview

The remainder of this specification is organized as follows:

- Chapter 2 – References and Conventions provides references to other documents, sources, and describes conventions used in this document.
- Chapter 3 – Definitions introduces specific language, terminology, and vocabulary used in this document. It is highly recommended that readers familiarize themselves with this chapter.
- Chapter 4 – General Concepts introduces important concepts of be managed by a host via a management interface, and the notions of standardized or vendor-specific management facilities, as well as cross version interworking and compatibility aspects.
- Chapter 5 – Management Interface describes generic hardware signals for module management and the register access layer providing READ and WRITE access to an immediately host addressable space of 256 byte-sized registers or memory locations in the module (the lower layers of the management interface stack are described in Appendix B)
- Chapter 6 – Core Management Features describes the CMIS **core** features and behaviors that are required in each CMIS manageable module, including hardware only modules¹
- Chapter 7 – Advanced Management Features describes optional (advertised) CMIS **advanced** management features which require more complex module firmware
- Chapter 8 – Module Management Memory Map describes the management application layer of CMIS, which is based on a three-dimensionally addressed byte-organized management memory with register (or memory) read and write access primitives
- Chapter 9 – CDB Command Reference provides the message catalogue of a memory based messaging mechanism called Command Data Block (CDB)
- Chapter 10 – Management Timing Specifications collects specifications related to timing parameters and timing dependencies

Appendices provide further information and examples.

- Appendix A – Form Factor Specific MSL or MCI Signal Names describes the mapping of form factor specific signal names to CMIS generic signal names and the encoding of logical values
- Appendix B – Management Communication Interface Definitions describes data transfer and physical layers available to implement the essential register access operations of CMIS
- Appendix C – Examples of Application Advertisements provides examples of how module Applications are advertised
- Appendix D – Examples of Initialization and Deinitialization describes example sequences of events and interactions illustrating how a host may initialize or deinitialize a module
- Appendix E – Illustration of Applying Control Sets provides illustration of the command handling associated with configuration changes.
- Appendix F – Examples of Diagnostic Features Usage provides illustrative examples of how diagnostics features might be used
- Appendix G – Specification Evolution and Maintenance Notes documents guidelines for future evolution of this specification
- Appendix H – Examples for Network Path Applications provides examples for both advertising and provisioning of client encapsulation applications (a.k.a. multiplex applications)

¹ Only a small subset of chapters 6 and 8 are actually applicable to hardware only modules, conditional on advertisement bits in the core register map.

2 References and Conventions

2.1 Industry Documents

The following documents are relevant to this specification

2.1.1 Interdependent Documents

- [1] UM10204 I2C-bus specification and user manual, NXP Rev. 6.0, 2014, <https://www.nxp.com/docs/en/user-guide/UM10204.pdf>
- [2] QSFP-DD/QSFP-DD800/QSFP112 Hardware Specification for QSFP Double Density 8x and QSFP 4x Pluggable Transceivers, Rev. 6.01, QSFP-DD MSA, 2021, <http://www.qsfp-dd.com/specification>
- [3] OSFP MSA Specification for OSFP Octal Small Form Factor Pluggable Module, OSFP MSA, Rev. 4.1, 2021, <https://osfpmsa.org/specification.html>
- [4] COBO 8-Lane & 16-Lane On-Board Optics Specification, Consortium for On-Board Optics (COBO), Rev 1.1, 2018, <https://www.onboardoptics.org/specifications>
- [5] SFF-8024, SFF Module Management Reference Code Tables, SNIA: SFF TA TWG (Storage Networking Industry Association: Small-Form-Factor Technology Affiliate Technical Working Group) Rev 4.12, 2024, <https://www.snia.org/technology-communities/sff/specifications>

2.1.2 CMIS Extensions and Supplements

- [6] OIF-CMIS-VCS-01.0, CMIS Versatile Control Sets (to be published)
- [7] OIF-C-CMIS-01.3, Implementation Agreement for Coherent CMIS (October 2023)
- [8] OIF-CMIS-FF-01.0, CMIS Form Factor Specific Management Features (to be published)
- [9] OIF-CMIS-LT-01.0, CMIS Support for Link Training (to be published)
- [10] OIF-CMIS-ELSFP-01.0, CMIS for External Laser Small Formfactor Pluggable Modules (to be published)

2.1.3 General Background

- [11] CFP MSA Management Interface Specification Version 2.6, <http://www.cfp-msa.org/documents.html>
- [12] SFF-8074, SFP (Small Formfactor Pluggable) Transceiver, Rev 1.0, 2001
- [13] SFF-8636, Management Interface for 4-lane Modules and Cables, Rev. 2.10a, 2019
- [14] SFF-8679, QSFP+ 4X Hardware and Electrical Specification, Rev. 1.8, 2018
- [15] OIF Common Electrical Interface (CEI) Specifications
- [16] IEEE Std 802.3-2022, IEEE Standard for Ethernet, <https://standards.ieee.org/ieee/802.3/10422>
- [17] INCITS Fibre Channel Specification, https://standards.incits.org/apps/group_public/download.php
- [18] IEEE Organizationally Unique Identifiers (OUI)
- [19] PCI-SIG 5.0/6.0 CopperLink External Cable Specification, PCI-SIG (to be published)

2.2 Sources

2.2.1 Standards and Specifications

Copies of IEEE standards may be obtained from the Institute of Electrical and Electronics Engineers (IEEE) (<https://www.ieee.org>).

Copies of InfiniBand standards may be obtained from the InfiniBand Trade Association (IBTA) (<http://www.infinibandta.org>).

Copies of OIF Implementation Agreements may be obtained from the Optical Internetworking Forum (OIF) (<http://www.oiforum.com/technical-work/implementation-agreements-ias>).

Copies of small form factor (SFF) specifications may be obtained from the SNIA SFF Technology Affiliate site <https://www.snia.org/technology-communities/sff/specifications>

2.2.2 Administrative Material

Copies of Common Language Equipment Identification (CLEI) specifications may be obtained from <http://www.commonlanguage.com>.

Registration information for OUIs may be obtained from the IEEE Registration Authority <https://standards.ieee.org/products-services/regauth/index>

2.3 Conventions

The conventions defined in this section are used throughout this specification.

2.3.1 General Conventions

Definitions

In this specification, English words or general terms may be used as **technical terms** with well-defined and specific meaning that is either different from or narrower than the normal English meaning in a general context. Technical terms are preferably defined in chapter 3 (Definitions), or otherwise when they first appear in the main text flow.

Note: A technical term may, depending on context, be printed in bold to assist readers in recognizing that special meaning is attached. Typographical conventions are defined in subsection 2.3.3.

Order of Precedence

If a conflict arises between the interpretations of text, tables, or figures, the order of precedence to resolve the conflicts is text first, then tables, and finally figures. **Exceptions:** Tables that provide detailed and primary specification text (not just overview or illustration) take precedence over text referring to those tables. Diagrams used to convey exact graphical specifications (not just illustrations), such as state transition diagrams, take precedence over text and tables.

Figures and Tables

Figures and tables providing an overview of details specified elsewhere are illustrative. Tables originally specifying data, formats, and values are normative. Not all tables or figures are fully described in the main text.

Lists

Lists sequenced by lowercase or uppercase letters show no ordering relationship between the listed items. Lists sequenced by numbers show an ordering relationship between the listed items.

Names

Symbolic names assigned to registers, bits, fields, or constants have only local significance within this specification. While it is understood that software implementations may use similar or even identical names, this is neither assumed nor required. Names defined in this specification may change in exceptional cases.

Unnumbered Subheadings

Long text segments may be sub structured for readability by informal, unnumbered subheadings, including paragraph headings placed at the begin of a long paragraph. Such subheadings are set in bold font.

2.3.2 Notational Conventions

Field and Register Names

Field and register names are defined for specification convenience, allowing readers to refer to Memory Map locations by name rather than address.

Field and register names are recognized as one contiguous token (word without spaces) with capitalization of initials to indicate word boundaries. This notation helps to clearly distinguish field names from general text.

An asterisk * as part of a field name is used as a **wildcard** indicating that all fields with names matching the name pattern are referred to. A wildcard matching the instance identification at the end of a field array element name is often suppressed for readability.

The preferred name structure for configuration registers is <**noun**><**verb**>¹, where the noun denotes an object to which an action is applied, and where the <**verb**> denotes the action, as e.g., in OutputDisable.

For objects with several instances, a name part for instance identification is often appended, as in OutputDisableRx<i>, which allows arrays to be used in programming languages.

Register Structures and Containers

For repeated structures of registers, a **container** name may be prepended to names of these registers to achieve globally unique names. The container name and the register name are separated by a double colon, as in the syntax <**container name**>::<**register name**>.

¹ Angle brackets are used to enclose the name of a named placeholder.

1 Decimal Numbers

2 The American notation of decimal numbering is used, with a comma as thousands separator and a period for
3 the decimal point.

American	ISO
0.6	0.6
1,000.0	1 000
1,323,462.9	1 323 462.9

5 Non-Decimal Numbers

6 Hexadecimal numbers are marked with suffix **h** (e.g. 10h), often written with leading zeroes (0010h). Non-
7 numeric hexadecimal digits (ABCDEF) are capitalized (e.g. 0Fh).

8 Binary numbers are marked with suffix **b** (e.g. 10000b), often written with leading zeroes (00010000b)

9 Numerals without a base-indicating suffix are understood to be in decimal notation (e.g. 16)

10 Base-indicating suffixes may be omitted for unambiguous cases like 0=0b=0h and 1=1b=1h.

11 Spaces may be inserted to make long hexadecimal or binary digit strings readable (e.g. 0001 0000b).

12 Logical Operators and Expressions

13 Logical operators are written in uppercase (**OR**, **AND**, **NOT**) and parentheses are used to clarify precedence.

14 Logical Values

15 Logical values are TRUE and FALSE.

16 The default encodings of TRUE and FALSE are 1b and 0b, respectively.

17 Logical Hardware Signal Values

18 For generic HW signals, where physical signal levels depend on the signal logic encoding, the symbolic signal
19 levels ASSERTED and DEASSERTED are used, which correspond to TRUE and FALSE (or 1 and 0), respectively.

20 To allow using the symbolic signal levels ASSERTED and DEASSERTED in logical expressions of this specification,
21 two logical predicates (i.e. mappings into the value set {TRUE, FALSE} or {1,0}) are introduced:

- 22 • **ASSERTED**(signal) maps to a logic value of TRUE (or a binary value of 1) when the signal's symbolic
23 level is ASSERTED, and FALSE (or a binary value of 0) otherwise.
- 24 • **DEASSERTED**(signal) maps to a logic value of TRUE (or a binary value of 1) when the signal's symbolic
25 level is DEASSERTED, and FALSE (or a binary value of 0) otherwise.

26 Logical Expressions

27 Logical expressions can be defined by a mix of bits in the Memory Map, hardware signal levels, and logic values.
28 While Memory Map values are commonly represented as 1b or 0b, hardware signals are expressed as ASSERTED
29 or DEASSERTED, and Boolean logic values are expressed as TRUE or FALSE.

30 To simplify the truth tables and logic equations in this specification, the following equivalency is used.

CMIS	Bit	Boolean Logic	Hardware Signal
0	0b	FALSE	DEASSERTED
1	1b	TRUE	ASSERTED

31 A don't care term in a logical expression is denoted by X

32 Named Logical Expressions

33 Logical expressions representing exit conditions in state transition diagrams are called **transition signals**, by
34 analogy with state machine circuitry.

35 When symbolic names are assigned to transition signals, a name suffix **S** (e.g. ResetS) is used, as a mnemonic
36 of Transition **Signal**.

37 When symbolic names are assigned to general logical expressions, a name suffix **T** (e.g. ModuleReadyT) is
38 used, as a mnemonic of Logic **Term**.

2.3.3 Typographical Conventions

Technical Terms

This specification uses certain English words in a specific sense as technical terms. The specific meaning of these technical terms is defined in section 3.3. To avoid possible confusion with other meanings outside of this specification, words used as specific technical terms have capitalized initials.

Narrowed Meaning

This specification may use certain English words, composites, or general technical terms in a narrower sense than the traditional English meaning. The narrowed meaning of such words or technical term is also described in section 3.3, as a heads-up for the reader. Such words are usually not marked-up typographically.

Emphasized Text

Locally emphasized text (emphasis for better readability) is typeset in **bold font**.

Auxiliary Text

Non-normative, informative, and auxiliary text, such as examples, hints, notes, or explanations is printed in *italic font* to help in visual differentiation of auxiliary text versus normative specification text.

Syntactic Variables (Named Placeholders)

Syntactic variables (named placeholders) in register names or field names may be indicated by angle brackets or by italic or bold font, as e.g. *<n>*, *n*, or **n**. For example, the term OutputStatusRx*<n>* represents the 8 bits OutputStatusRx1, OutputStatusRx2, ..., OutputStatusRx8.

2.3.4 Addressing and Referencing Conventions

Note: The conventions in this subsection may require understanding of specific terminology defined in the glossary (see section 3.3).

Referencing Managed Module Resources (by Lanes)

In this specification, a **managed** transmission-related internal module **resource** is identified via the **lane number** (see next paragraph) of a **lane** (see glossary in section 3.3) that is associated with or connected to the managed resource in question.

In cases where a managed status or control aspect of a managed resource is related to lanes after multiplexing or demultiplexing has occurred, the status or control aspect is applicable to all lanes of the Data Path containing the lane number given as a reference, unless otherwise indicated.

Whenever the unspecific term 'lane' is used, instead of 'host lane' or 'media lane', and the context does not clearly allow to disambiguate its meaning, a host lane perspective is assumed.

Lane Numbers and Lane Numbering References

Lane numbers are used throughout CMIS to identify lanes and lane-related resources that are managed via named addressable registers (or fields) in the management memory map (see chapter 8).

The **lane numbers**, which are used as part of register names, are originally defined at physical **name binding reference points**¹. These name binding reference points are always located at a physical module boundary; examples include the geometric position of media contacts or receptacles in a connector, possibly including physical channel parameters (like, e.g., carrier wavelength) in case of multi-channel transmission schemes.

On the **host side**, the ascending host lane number sequence (e.g. 1 to 8) is always mapped one-to-one to the ascending numbering sequence in the names of the relevant electrical host interface connector contacts (such as Rx1 to Rx8), which are the physically identifiable name binding reference points as defined in the relevant hardware specification.

On the **media side** there is no universal mapping of media lane numbers to physically identifiable media side entities like fibers, wavelengths, or media side fiber connector positions. Instead, this mapping is advertised by the module and may depend on the relevant media interface standard (see section 8.10.8 for details).

Note: While the media lane definition may vary with interface standard, CMIS assumes that a media lane always comes with some associated optical attributes like transmitted power or received power.

¹ This description is true for the fixed name binding in CMIS core as described in chapter 6. As soon as flexible connectivity between lane segments is introduced, the description is true only for the trivial switch configuration where no switching is done. See e.g. section 7.8.

1 Lane Numbers and Lane Data Indexing

2 The host must locate (address) the specific register (or field) where the desired information (attributes) of a
3 lane with given lane number is presented to the host in the management memory map. Usually, lane-related
4 information (lane attributes) are arranged – at least conceptually – in the form of an **array** of a data structure,
5 such that an **address offset** (or more general: a zero-based **lane index**) can be used to access any desired
6 instance in the array. Generally, the lane index of a lane is the lane number of the lane minus one.

8 Referencing Elements of the Management Memory Map (Bytes, Bits, Fields)

9 Each host accessible Byte in the structured internal management Memory Map can be identified by an address
10 triple consisting of a Bank Index (0-255), a Page Index (00h-FFh), and a Byte Address (0-255). See chapter 7.8
11 for more information on this three-dimensional addressing structure.

12 Pages without Bank support implicitly have a Bank Index of 0 (which can be omitted).

13 The so-called Lower Memory implicitly has a Bank Index 0 and a Page Index 00h (both of which can be omitted).

14 The following **colon separated** Byte addressing syntax is used (*italics* denoting grammatical variables):

15 <i>Bank:Page:Byte</i>	(general Byte in the register Memory Map)
16 <i>Page:Byte</i>	(Byte in a page without Banking support)
17 <i>0:Page:Byte</i>	(fully specified alternate format)
18 <i>Byte</i>	(Byte in Lower Memory)
19 <i>00h:Byte</i>	(alternate format)
20 <i>0:00h:Byte</i>	(fully specified alternate format)

21 At each addressing level (i.e. at Bank, Page, or Byte addressing level), a **range** may be indicated by **dash separated** addresses *from-to*

22 Examples:

24 Banks 0-1
25 Pages A0h-AFh
26 Pages 0:9Fh-AFh
27 Bytes 1:9Fh:128-129
28 Bytes 0-3:9Fh:128

29 With Bits in a Byte indexed from 0 to 7, where index values denote arithmetic significance, the following **dot separated** notation may be used to identify a single **Bit** or a **Bit Field**:

32 <i>Bank:Page:Byte.Bit</i>	(single Bit)
33 <i>Bank:Page:Byte.Bit-Bit</i>	(Bit Field ranging from Bit to Bit)

34 Examples:

35 Field 2:10h:128.7-4	4-bit field of bits 7-4 of byte 128 in Upper Memory Bank 2 of page 10h
36 Bit 04h:128.0	least significant bit 0 of byte 128 in Page 04h (of unspecified Bank)
37 Byte 126	Byte at Byte address 126 in Lower Memory
38 Word 04h:128-129	A 16-bit word in Page 04h (see section 8.1.3.5 for Endianness conventions)

40 Referencing Field Values

41 When clear from context, the address of a field may also denote the value stored in the field. Otherwise square
42 brackets around a field address denote the value stored in the field.

43 <i>Address = Value</i>	An assignment
44 <i>[Address] = Value</i>	An expression, e.g. in a comparison

45 Examples:

46 When 10h:132 = 0 the module

3 Definitions

For the purposes of this specification, the following keywords, acronyms, and term definitions apply.

3.1 Keywords

This section lists verbs and adjectives intended to guide the interpretation of text in this specification in a formalized manner. Note, however, that this specification largely consists of detailed technical specifications which are presented in indicative and factual language, not in formalized requirements language.

Custom: Fields and formats described as **custom** are under control of each individual module vendor. The same custom resource may be used differently by different vendors or groups of vendors.

May: This verb indicates flexibility of choice with no implied preference, both for positive and for negative statements.

Obsolete: The adjective **obsolete** indicates that an item was defined in prior specifications but has been removed from this specification. The adjective obsolescent warns that an item is about to become obsolete.

Optional: The adjective **optional** describes features which are not required by the specification. However, if an optional feature is supported, the specifications in this specification do apply. Describing a feature as optional in the text is done to assist the reader.

Note: As specified in 8.1.3.1, when an optional feature is not supported the default values for bits and bytes associated with this feature shall be zero unless otherwise stated.

Note: Not supported optional registers or fields in a supported Page are accessible. Not supported optional Pages or Banks are not accessible.

Prohibited: The adjective **prohibited** describes a feature, function, or coded value that is defined in a referenced specification to which this specification makes a reference, where the use of said feature, function, or coded value is not allowed for implementations of this specification.

Reserved: The adjective **reserved** describes elements or resources set aside for future standardization. Such resources include Bits, Bytes, Fields, and coded values. A reserved element is not available for vendor specific use.

Note: As specified in 8.1.3.1 the default value of a reserved storage element on a supported Page shall be 0. The module is required to define a Reserved field or bit as 0, but, defensively, the host should not check Reserved fields or bits for 0.

Restricted: The adjective **restricted** describes features, bits, bytes, words, and fields that are set aside for other standardization purposes defined elsewhere, including CMIS supplement specifications. In contexts where these other standard specifications do not apply, the restricted bit, byte, word, or field is treated like a reserved bit, byte, word, or field (i.e. a **restricted** byte uses the same value as defined for a **reserved** byte).

Shall: This verb indicates a mandatory requirement of particular importance or of general nature. As is common in detailed technical specifications, simple descriptive statements of fact also express mandatory, or conditionally optional, technical requirements. Implementation of mandatory requirements is necessary to ensure interoperability with other products that conform to this specification. Note, however, that at this time no formally defined specification conformance criteria exist.

Should: This verb indicates flexibility of choice with a strongly preferred alternative.

Vendor specific: Indicates that something (e.g., bit, field, code value) is not defined by this specification. Specifications of vendor specific items are provided by the relevant vendor.

3.2 Abbreviations

This section lists abbreviations, most of which are acronyms, together with their full text expansion, without explanation of meaning (except for short hints). When abbreviated terms or acronyms need to be explained or used in a specialized sense, then the meaning of such abbreviations is defined in section 3.3.

5	ACC	Active Copper Cable (linear)
6	AEC	Active Electrical Cable (retimed)
7	ACK	Acknowledge
8	ACS	Active Control Set
9	Adv.	Advertised (a register classification tag)
10	AIS	Alarm Indication Signal
11	AOC	Active Optical Cable
12	ASCII	American Standard Code for Information Interchange (the numerical representation of a character)
13	ASN.1	Abstract Syntax Notation One
14	BER	Bit Error Ratio (dimensionless) or Bit Error Rate (per unit time)
15	BER	Basic Encoding Rules (for ASN.1)
16	BERT	BER Testing
17	BOL	Begin of Life
18	CA	Certification Authority
19	CDB	Command Data Block
20	CDM	Code Division Multiplex
21	CDR	Clock and Data Recovery
22	CLEI	Common Language Equipment Identification
23	CMIS	Common Management Interface Specification
24	Cnd.	Conditionally required (a register classification tag)
25	COR	Clear on Read (side effect)
26	CPO	Co-packaged Optics
27	CTLE	Continuous Time Linear Equalizer
28	CWDM	Coarse Wavelength Division Multiplexing
29	DAC	Direct Attach Cable (passive copper cable)
30	DER	Distinguished Encoding Rules (a restricted form of ASN.1 BER)
31	DFB	Distributed Feedback Laser
32	DPID	Data Path ID (the smallest host lane number of the Data Path)
33	DPSM	Data Path State Machine
34	DWDM	Dense Wavelength Division Multiplexing
35	EC	Explicit Control (name of a control bit)
36	ELS	External Laser Source
37	ELSFP	External Laser Source Formfactor Pluggable
38	EML	Externally Modulated Laser
39	EPL	Extended Payload
40	ePPS	enhanced Pulse Per Second (a timing signal)
41	EOL	End of Life
42	eSNR	SNR determined from electrical measurement
43	FC	Fibre Channel
44	FEC	Forward Error Correction
45	FERC	Frame Error Count
46	FP	Fabry-Perot (a laser type)
47	FSM	Finite State Machine
48	HP	Host Path
49	IA	Implementation Agreement
50	IB	InfiniBand
51	IC	Integrated Circuit
52	ID	Identifier
53	IDevID	Initial Device Identity (defined by the supplier of a device)
54	ISI	Inter-Symbol Interference
55	I2C	Inter IC (a two-wire bus specification)
56	I2CMCI	I2C-based MCI (a management communication interface specification defined here)
57	LF	Local Fault
58	LOL	Loss of Lock

1	LOS	Loss of Signal
2	LPL	Local Payload
3	LPO	Linear Pluggable Optics (linear, cable or transceiver)
4	LSB	Least Significant Bit (in a multiple bits context), Least Significant Byte (in a multi-byte context)
5	LT	Link Training
6	LTP	Level Transition Parameter
7	MBR	Module Boot Record
8	MCI	Management Communication Interface
9	MIS	Management Interface Specification
10	MSB	Most Significant Bit (in a multiple bits context), Most Significant Byte (in a multi-byte context)
11	MSL	Management Signaling Layer
12	MSM	Module State Machine
13	N/A	Not Applicable or Not Available
14	NACK	Not Acknowledged
15	NAD	Normalized Application Descriptor (a consolidated data structure describing an Application)
16	NP	Network Path
17	NPSM	Network Path State Machine
18	NV	Non-Volatile
19	OIF	Optical Internetworking Forum
20	OMA	Optical Modulation Amplitude
21	OOR	Out of Range
22	Opt.	Optional (register classification)
23	OTN	Optical Transport Network
24	OUI	Organizationally Unique Identifier (unique vendor code assigned by the IEEE)
25	PM	Performance Monitoring
26	PRBS	Pseudo Random Binary Sequence
27	PSL	Pattern Synchronization Loss
28	RO	Read-Only
29	RAL	Register Access Layer
30	RF	Remote Fault
31	Rqd.	Required (a register classification tag)
32	RS	Reed-Solomon (a FEC code)
33	RW	Readable and Writeable
34	Rx	Receiver function of module, receive direction (media to host)
35	SC	Self-Clearing (side effect)
36	SI	Signal Integrity
37	SCL	unidirectional Serial Clock
38	SCS	Staged Control Set
39	SDA	bidirectional Serial Data
40	SEW	Symbol Error Weight (number of symbol errors in a FEC frame)
41	SI	Signal Integrity
42	SFF	Small Form Factor
43	SM	State Machine
44	SPI	Serial Peripheral Interface
45	SNR	Signal power to Noise power Ratio
46	STD	State Transition Diagram
47	STT	State Transition Table
48	TC	Threshold Crossing
49	TDM	Time Division Multiplex
50	TEC	Thermoelectric Cooler
51	TWI	Two Wire Interface (an obsolescent alias for I2C)
52	Tx	Transmitter function of module, transmit direction (host to media)
53	VCS	Versatile Control Set (an optional CMIS feature defined in a supplement)
54	VCSEL	Vertical Cavity Surface Emitting Laser
55	VDM	Versatile Diagnostics Monitoring (a CMIS feature)
56	WDM	Wavelength Division Multiplexing
57	WO	Write-Only

3.3 Glossary

This Glossary defines the technical vocabulary used in this specification. The Glossary can be skimmed or skipped on first reading but should always be consulted when normative specification text is examined.

This specification uses certain English words as technical terms in a very specific sense. The specific meaning of these technical terms is defined here in this Glossary section. To avoid possible confusion with other meanings outside of this specification, words used as a technical term are usually capitalized (initials).

This specification also uses certain English words or general technical terms in a narrower sense than their general usage. The narrower meaning of such words or technical term is described in this section.

ACCESS: This all-caps spelling denotes a READ or WRITE access via management communication interface.

ACCESS Delay: A module may delay responding to an ACCESS by the host, with maximum delay durations specified in chapter 10. A module might need to slightly delay an incoming ACCESS to first complete some internal processing or synchronization for data consistency. See also Transaction delay.

ACCESS Hold-Off Period: is a period during which a module **rejects** any READ or WRITE access by the host.

ACCESS Rejection: A module may **reject** ACCESS by the host (see B.3.9.2). There are several reasons why a module might reject ACCESS, which are specified in the relevant section of this specification, with maximum durations specified in chapter 10. See also Transaction rejection a.k.a **Transaction Hold-Off**.

Advertisement: An advertisement provides in-advance information about a feature, characteristic, or property. An advertisement is available before use and does not change dynamically.

Alarm: An **alarm** informs the host about operationally undesired situations or about critical threshold crossings of monitored observables. The module raises an alarm by setting an associated **Flag** that represents the alarm.

Application: An **Application** is a well-defined transmission function provided by a CMIS managed module. For system interfaces, the application is usually defined by reference to a specific 1:1 combination of an industry standard host interface and an industry standard media interface. *See Section 6.2 for more information.* In contrast, multiplex applications are defined by an N:1 relationship between N host interfaces and one media interface. *See Section 7.6 for more information.*

Application Descriptor: The data describing the host and media side (multi-)lane interface characteristics of a particular Application to the host. Application Descriptors are presented to the host in a non-universal, module defined order. The position or sequence number of an Application Descriptor in this order is used as the key to identify the Application described by the Application Descriptor.

Array: Repeated data structures in Pages of the Management Memory Map are referred to as arrays. For example, the Applications supported by a module are described in arrays of Application Descriptor data structures.

Bank: In the management memory map, a Bank (capitalized) contains several Pages having the same Page Index distinguished by a Bank Index. In the context of firmware management, the term bank refers to a non-volatile storage location for a firmware image.

Bank Broadcast: An optional (advertised) feature where the module copies a value written to one Bank of a Page to all other Banks of the same Page.

Banking: A Memory Map architectural feature that allows modules to implement multiple Memory Map Pages that have the same Page address, effectively providing additional memory depth. For example, banking is used to provide additional lane related registers for more than eight lanes. Such Pages are called **lane banked**. *See Section 8.1 for more information.*

Bit (storage) versus **bit** (value): in common technical parlance the term bit may refer either to a binary value or it may refer to a binary storage element (a binary variable) that is usually part of a larger addressable storage element (e.g. part of a Byte). For clear and context-free differentiation in this specification, the spelling **Bit** (capitalized) is used to refer to a binary storage element, and the spelling **bit** (lowercase) refers to a binary value, which may be stored in a Bit.

Byte (storage) versus **byte** (value): in common technical parlance the term byte may refer either to a value or it may refer to a storage element (a variable). For clear and context-free differentiation in this specification, the spelling **Byte** (capitalized) is used to refer to an addressable storage element, whereas the spelling **byte** (lower case) refers to a value, which may be stored in a Byte or used in a computational expression.

Certificate: Cryptographically protected information proving integrity and authenticity of data or objects.

Channel: A (physical) channel is characterized by the resources (e.g., frequencies) within signal transmission media that are used to carry the modulated analog signal representing the data transmitted over lanes. Examples include baseband electrical and passband optical channels.

Checksum: A value derived from a block of digital data for the purpose of detecting errors in that block.

Client Encapsulation (Application or Module): Client encapsulation modules have been introduced in chapter 1. The associated transmission functionality is called a client encapsulation application. In practice, such modules are usually, but somewhat vaguely, referred to as transport, muxceiver, muxponder, or transponder modules. Compare with system interface modules.

Configuration (data): The term configuration is used in a broad sense to denote data which influence module behavior that can be modified by the host, i.e. excluding administrative data. See also Status and Trigger.

Conformance (Flagging with Interrupt generation rules): In this specification, the wording conformance, Interrupt conformance, Flag conformance, and **Flagging conformance** all relate to requirements, under which conditions the module may set Flags which potentially raise Interrupts.

Controls: Fields that are **writeable** and affect module operation are sometimes referred to as **controls**, in contrast to writeable fields which just store passive data. Typical examples of controls are **configuration** data fields and **trigger** fields where writing to the field triggers some processing in the module.

Data Path: A Data Path (**DP**) represents a contiguously numbered **group of host lanes** and an associated contiguously numbered **group of media lanes**, together with the associated internal resources, altogether carrying one module Application. For **system interface** applications, all parts of a (simple) Data Path are initialized, used, and deinitialized together, whereas for **client encapsulation** applications the overall Data Path is partitioned into one or more host side Host Paths and a single media side Network Path as the elements for a finer granularity of initialization, usage, and deinitialization. The Host Interface of a Data Path can carry a multi-lane signal (e.g. CAUI-4) or a single-lane signal (e.g. 25GAUI) in case of a non-multiplexing application, or groups of such signals in case of a genuine multiplexing application. The Media Interface of a Data Path can also consist of a single media lane or of multiple media lanes (in single-carrier or multi-carrier transmission schemes, respectively).

Default (value): The Default value of a management attribute is the value that a host would retrieve immediately after the management memory has become accessible (at some stage after power-on, or post-reset, when the management communication interface has been brought up)

Digest (message digest): The fixed size result of a cryptographic hash function applied to an arbitrary length "message" (data vector) which serves as a "fingerprint" of the message,

Dynamic (status): A dynamic status register displays status information that can change during operation.

Fault: In this specification the term fault is used exclusively for module detected conditions that incur a risk of physical damage to the module or its environment or that incur a safety risk. A partial or full loss of functionality or a deviation from desired state or behavior is called failure or error.

Field: In this specification, a **field** (lower case) is an informal notion referring to some scalar data element that is stored or accessed in management memory. It could be a portion of a Byte or a value larger than a Byte. When used with precise meaning, a **Field** (capitalized) is simply a part of a **Register** that is sub structured.

Firmware: The firmware of a module may be an aggregate or bundle consisting of multiple components such as executable code or data resources for multiple processors or storage locations. The individual components of the module firmware are not exposed to or visible via CMIS, albeit the firmware update mechanism allows vendors to support access to individual firmware components.

Firmware Update: This term refers to operations on firmware stores and to the function to pass control to another image. Operations on image stores include **download**, **upload**, and **copy** firmware images (to or from or between a module's internal firmware store). The operation to pass control is called **activate**. Firmware update can serve both purposes of **upgrade**, **repair**, and **downgrade** and may occur either immediately (**switch**) or on the next restart (**commit**).

Firmware Upgrade: This term is sometimes used as a feature name, but it may be misleading. This specification prefers the neutral term Firmware Update or Firmware Management.

Flag: A **Flag** (capitalized spelling) is a clear-on-read **Latch** with associated maskable **Interrupt** generation. Each Flag comes with an associated **Mask** bit. Interrupt generation from the Flag is suppressed when the associated Mask bit is set. Flags are used to reliably and (optionally) immediately inform the host about the onset of conditions or the occurrence of events. *Note: Using the word flag (lowercase) in a general sense to*

1 denote a simple binary indicator bit is discouraged in this specification. Such bits will be called status bits or
2 indicator bits instead (see Status).

3 **Flat Memory Module:** A Flat Memory module provides 256 Bytes of immediately addressable management
4 memory. Unlike a Paged Memory module, a Flat Memory module does not support dynamic Paging into the
5 Upper Memory. The simplest form of a Flat Memory module supports only constant read-only data.

6 *Note: In this version of CMIS, flat memory modules are assumed to support only read-only static data as*
7 *provided by an EEPROM. See sections 6.3.2.3. and 8.2.*

8 **Freezing** (Holding): Temporarily freezing dynamically updated registers for consistent readout. This is an
9 alternative to sampling the results into a separate result register.

10 *Note: In CMIS, freezing is used in the VDM feature, whereas sampling is used in Diagnostics Monitoring.*

11 **Gating:** A gated measurement or monitor runs for a specified amount of time, often triggered by an event. In
12 on-demand measurements, the gating start event is writing a trigger bit; there is then an unobserved time
13 between the end of the measurement and the start of the next one. In periodic measurements, the result of a
14 gating period is sampled in result registers while the next gating period is started without a gap. Ungated
15 measurements are started and stopped on host-command.

16 **Gray Coding:** used with PAM4 modulation. Defines the mapping of 2 binary bits into 4 levels.

17 (0,0) maps to level 0

18 (0,1) maps to level 1

19 (1,1) maps to level 2

20 (1,0) maps to level 3

21 **Host Interface:** The host interface is the high-speed electrical interface for connecting the module (or an
22 application on a module) to a host system via electrical high-speed lanes. The requirements of a specific host
23 interface are defined in the associated industry standard for that interface. See Section 6.1.

24 **Host Path:** A Host Path (**HP**) represents a **contiguously numbered** group of host lanes and associated
25 internal resources that will be initialized, used, and deinitialized together, essentially because the host lane
26 group carries one host signal in a **client encapsulation** application. In those applications, a Host Path (**HP**)
27 corresponds to a **host side segment** of a Data Path (**DP**), while the media side segment is modeled separately
28 as a connected-to Network Path (**NP**) serving one or more Host Paths, in uniplex or multiplex applications,
29 respectively. Hence a Host Path consists of the functional transmission resources between a Host Interface and
30 an internal multiplex connection point or plane (the set of multiplex connection points) within the module that
31 connects the host side receive (transmit) processing (the **HP**) and the media side transmit (receive) processing
32 (the **NP**).

33 **Interface:** The general term **interface** is ambiguous. In this specification three interpretations prevail:

- Interface as a **device-to-device** interconnection (e.g. Host Interface, Management Interface)
- Interface as a **device to media** attachment (e.g. Media Interface)
- Interface as an **adapter device function** processing signals carried on an interconnection
(e.g. when functions or registers of a module's interface are described).

38 **Interrupt:** The generic hardware signal that, when asserted, represents a pending interrupt request from
39 module to host is called **Interrupt**. In this specification, the Interrupt is asserted as long as any **Flag** is set
40 with its associated **Mask** cleared.

41 **Interrupt Flag Bit (obsolete):** The term interrupt flag or interrupt flag bit was used side by side with flag or
42 flag bit, in previous versions of this specification. These terms are now obsolete and replaced by the term **Flag**.

43 **Interval Statistics:** A raw performance monitor for an observable in the module provides a real-time value,
44 i.e. a current value **sample**, on request. Enhanced performance monitoring adds computation of statistics over
45 a periodic or on-demand **monitoring interval**. The **Interval Statistics** collected over a monitoring interval
46 typically include minimum value, maximum value, and sample mean (average), also known as min, max, avg.

47 **Lane** (general): Generally, a lane is a concatenation of passive pathways and active functional facilities inside
48 and around a module that are **traversed by a signal** being transmitted by the module. A lane has **external**
49 segments (within electrical interconnections to the host system or within optical interconnections to far end
50 modules), **internal** segments (within the module), and connection points (in module **ports** at physical module
51 boundaries) where external segments and internal segments meet.

52 *Note: Lanes should not be confused with signals: lanes carry signals like roads carry traffic.*

1 Note: A host lane consists of an external segment (host interconnection electrical lane) attached to a module-
2 internal segment (module host lane) at a host lane port of the module, and a media lane consists of an external
3 segment (optical interconnect media lane) attached to a module-internal segment (module media lane) at a
4 media lane port of the module.

5 **Lane** (segment, "pars pro toto"): For CMIS module management, most often the module-internal lane segments
6 ending at external module interfaces (lane ports) are of greatest interest. With moderate loss of language
7 accuracy, these lane segments of interest may be called just lanes as well, acting as a pars-pro-toto reference.

8 Such a module-internal **lane** (segment) encompasses the internal pathways and functional facilities used for
9 processing or propagating a serial high-speed signal entering or exiting the module via one of its physical
10 interfaces on host side or on media side.

11 Note: As an example for a functional facility, a Lane being squelched means that the signal normally carried on
12 the Lane is muted, which is achieved by activating a squelching facility in the module, i.e. by internal resource
13 configuration. This configurable squelching facility is considered part of the lane

14 Note: Module resources being associated with host interface lanes or media interface lanes are a specification
15 convenience. These resources may neither be dedicated to nor physically proximate to the associated connector.

16 **Lane** (port): The ending of a module internal lane segment, usually in a well defined physical connector
17 reference point at a physical module boundary (e.g. geometrically localized electrical contacts), is called a lane
18 port (of the module). Host side lane ports are the point of physical connection to host interconnection electrical
19 lanes (single-ended electrical wire or differential electrical wire pair). Media side lane ports are the points of
20 physical connection to optical lanes in optical fiber media¹.

21 Note: The location in physical connectors and other technical characteristics of lane ports (e.g. modulation
22 carriers and modulation formats) are specified in the relevant hardware and physical interface specifications.

23 **Lane** (external interconnect, electrical, optical): The module-external lane segment of a host lane is also called
24 a **host-interconnect** electrical lane, or simply **electrical lane**, when it must be distinguished from the
25 module-internal host lane (segment). The module-external lane segment of a media lane is also called a
26 **module-interconnect** optical lane or simply **optical lane** when it must be distinguished from the module-
27 internal media lane (segment). Host-interconnect lanes meet host lanes in host lane ports of the module.
28 Module-interconnect optical lanes meet (are attached to) media lanes in media lane ports of the module.

29 **Lane** (identifier): In principle any lane segment or any fixed segment-to-segment connection point of a lane
30 can be used to identify the entire lane unless there is some flexible connectivity between any lane segments.

31 Note: In the CMIS core (see chapter 6) there is only fixed connectivity between external and internal lane
32 segments, but optional host lane switching extensions allow flexibility (see section 7.8). Beware of confusion.

33 Note: From a module management point of view, the properly identified physical lane ports on one hand and
34 the lane identifiers used in management registers on the other hand are the parts of a lane that are relevant
35 for lane identification. When flexible connectivity is introduced (see section 7.8) location and configured state
36 of the "switch" must be considered.

37 **Lane group:** A lane group is a set of lanes with contiguous lane numbers. The lane group can be identified by
38 the smallest lane number in the set and the number of lanes in the set (which may be clear from context).

39 **Lane index:** When lane attributes (i.e. lane related data) of a lane with given lane number **n** need to be stored
40 in management register arrays, field arrays, or bit vectors, these attributes are usually stored in the array
41 element (register, field, or bit) with array index **n-1**. The index value **n-1** used for addressing lane data is called
42 a lane index for lane number **n**. Lane indexes are non-negative numbers; the lowest lane index is zero.

43 Note: Register or field names often contain lane numbers as a naming component, whereas lane references
44 stored in registers or fields are often lane indexes. See also the addressing conventions in section 2.3.4.

45 **Lane number:** A lane number is a positive number used to identify a lane segment for management purposes.

46 Note: When the optional host lane switching extensions are used, it becomes relevant at what side of a lane
47 switch the lane numbering is applied (see section 7.8). Beware of confusion.

48 **Latch:** A latched Bit or Latch, for short, is a read-only Bit in management memory with a sticky value (1b):
49 Once the module has written the sticky value to the Bit, the value is held until the Bit is eventually cleared as

¹ Note that a single optical fiber connector may host multiple lane connection points, one lane per independently receivable optical channel, when channelized transmission formats such as WDM are employed.

an implicit side effect of reading by the host (clear on read). Latches are used to reliably inform the host about events or possibly transient states. *Note: In this specification, Latches with maskable interrupt generation are called Flags.*

Latching (obsolete): This term has been used in CMIS 4.0 to describe storing a real-time value into a result register. It is now replaced by the term **sampling**. Latches for counters are now referred to as **results**.

Link Training (unidirectional, electrical): Algorithms aiming at achieving a desired level of data transmission performance over a host interface electrical link, usually by signal shaping, i.e. by tuning transmit signal characteristics based on information gathered in the receiving end.

Lower Memory: The 128 bytes of host addressable memory addressed by byte addresses 00h through 7Fh.

Low Power Condition: The term Low Power condition refers to the condition that the module must be and remain in Low Power Mode (see subsection 6.3.2.4).

Mask: In this specification a Mask bit or **Mask** is a host writable Bit in management memory which, when set, suppresses (or ceases) Interrupt generation by its associated **Flag** bit.

Management Communication Interface (MCI): The data communication interface and interconnection between host and module that is used to transport and implement READ and WRITE register access primitives across the interconnection.

Management Memory: The term **management memory** refers to the internal storage of all host **accessible** management data (**registers, memory**). Note that host accessible is not the same as host addressable because a host may need to switch accessible content into Upper Memory such that it becomes host addressable.

Management Memory, Addressable: The **addressable management memory** which is the memory that can be directly accessed by READ and WRITE primitives, via the MCI. In CMIS this a 256 Byte area consisting of Lower Memory and Upper Memory.

Media Interface: A media interface is defined as the high-speed interface for attaching the module to an interconnect medium, such as wires or optical fibers. The requirements of a specific media interface are defined in the associated industry standard for that interface. See Section 6.1.

Memory: In this specification the term memory is often used interchangeably with the term register. Moving forward, **memory** will preferably be used for a host accessible storage location when that storage location is used for temporary or permanent data storage and when there is no module functionality associated with the data stored in that location (cf. **register**, when functionality is associated).

Memory Map: In this specification the term management **Memory Map** refers to the conceptual organization of all host accessible management data, independent of whether the module stores these management data fully or partially in classical hardware **registers** (evaluated by module hardware) or in memory **locations** (virtual registers evaluated by module firmware).

Module: Pluggable transceivers and active or passive cable assembly terminations that plug into a host receptacle, or onboard transceivers, such as, but not limited to, those of QSFP-DD, OSFP, COBO, and QSFP form factors are hereafter referred to as modules unless cable assemblies are specifically mentioned. In CMIS, one classification of modules is based on the management memory map, where modules are distinguished by their manageability as **paged** memory modules or as **flat** memory modules (unpaged). Another classification is based on the types of dynamic state machines supported. Still another distinction is by functionality: **resource** modules and **transmission** modules.

Monitor: A monitor is the function to measure an observable or to obtain a value from an estimator and to make the results of the measurement or estimation available in a register, for consumption by the host. Many monitors do not just sample an observable or process but require measurement or estimation time until the monitor value can be updated. Many monitors in CMIS also provide additional functionality such as threshold crossing detection (for low/high warning/alarm indications) or statistics (cumulative or interval).

Multiplex (genuine): Applications, configurations, or topologies that implement a multiplexing (N:1) function with **encapsulation** of several (N>1) lower rate host side signals (**host signals**) into one higher rate media side signal (**network signal**). In a broad sense, multiplexing sometimes includes the degenerate multiplex case also called **uniplex** (1:1).

Multi-Wavelength Modules (Multi-Wavelength Applications): Modules running Applications using a fixed or tunable WDM grid or supporting parallel Applications using different wavelengths.

Muting: This generic term refers to any action with the intended effect that the output of a signal source facility becomes **quiescent**: i.e. a muted output is quiescent. A host mutes an output by management command or configuration, either by **disabling an output** or by **forcing** an output to be **squelched**. The module can mute an output when an auto-squelch controller is enabled. Note that unmuting an output is not the same as enabling an output: the inverse of output disable/squelch is therefore to un-disable/un-squelch an output, respectively.

Network Path: A Network Path (**NP**) represents a group of contiguously numbered media lanes together with the associated internal resources that will be initialized, used, and deinitialized together, essentially because the media lane group carries one **network signal** in a **client encapsulation** application. In those applications, a Network Path represents the **media side segment** of the client encapsulation application, while the host side segment is modeled separately as one or more connected Host Paths (**HP**), in uniplex or multiplex applications, respectively. Within the module, a Network Path offers one (uniplex NP) or more (multiplex NP) client connection and mapping points (multiplex connection points) to connect one or more Host Paths (each with its own Host Interface) and to carry the associated host signals over the Media Interface associated with the Network Path. The Media Interface of a Network Path can consist of a single media lane or of multiple media lanes, in single-carrier or multi-carrier transmission schemes, respectively.

*Note: In contrast, in **system interface** applications, the host side segment and the media side segment are not separated but modeled as one Data Path (**DP**) between Host and Media Interfaces.*

NV: Non-Volatile (NV) memory is persistent and maintains stored information permanently between writes, also when the module is not powered

Observable: An observable is something that can be measured, monitored, estimated, or otherwise observed. In this specification, observables are the inputs to performance monitors which in turn provide real time samples, threshold crossing supervision, and sometimes sample statistics. Previous versions of this specification have also used terms like parameter, attribute, or monitor.

OM2: Cabled optical fiber containing 50/125 um multimode fiber with a minimum overfilled launch bandwidth of 500 MHz-km at 850 nm and 500 MHz-km at 1300 nm as IEC 60793-2-10 Type A1a.1 fiber.

OM3: Cabled optical fiber containing 50/125 um laser optimized multimode fiber with a minimum overfilled launch bandwidth of 1500 MHz-km at 850 nm and 500 MHz-km at 1300 nm as well as an effective laser launch bandwidth of 2000 MHz-km at 850 nm in accordance with IEC 60793-2-10 Type A1a.2 fiber.

OM4: Cabled optical fiber containing 50/125 um laser optimized multimode fiber with a minimum overfilled launch bandwidth of 3500 MHz-km at 850 nm and 500 MHz-km at 1300 nm as well as an effective laser launch bandwidth of 4700 MHz-km at 850 nm in accordance with IEC 60793-2-10 Type A1a.3 fiber.

OM5: Cabled optical fiber containing 50/125 um laser optimized multimode fiber with a minimum overfilled launch bandwidth of 3500 MHz-km at 850 nm, 1850 MHz-km at 953 nm and 500 MHz-km at 1300 nm as well as an effective laser launch bandwidth of 4700 MHz-km at 850 nm and 2470 MHz-km at 953 nm in accordance with IEC 60793-2-10 Type A1a.4 fiber.

OMA: Optical Modulation Amplitude: The difference between two optical power levels of a digital signal generated by an optical source, e.g., a laser diode.

Operational: A functional resource is **operational** if it has been fully configured and initialized so that it is able to perform its function according to specification. Note that this term represents only a configuration and initialization view of the resource; the actual transmission behavior of an operational resource may depend on other conditions, such as valid input signals being available.

OSNR: Optical Signal to Noise Ratio: The ratio between the optical signal power in a given signal bandwidth and the noise power in a given noise reference bandwidth.

Page: A management memory segment of 128 bytes that can be mapped (paged) into the host addressable Upper Memory, which is host addressable by byte addresses 128 to 255 (80h to FFh)

Paged Memory Module: A module supporting a paging mechanism where parts of a module's Management Memory are dynamically mapped into a smaller host-addressable memory window. Usually a Paged Memory module runs firmware on a module internal controller. See also Flat Memory Module.

Parameter (controlled variable versus uncontrolled observable): The term **parameter** denotes a controlled or controllable property of an object or function, i.e. an attribute that is under control of the object's environment. Note: The term parameter is also used loosely to refer to an observable, i.e. to an emergent rather than controlled characteristic of an object.

P_{av}: Average Power: The average optical power P_{av}

1 **Performance Monitoring (PM):** This term is informally used to refer to the functionality of Interval Statistics
2 provided for an intensive observable and for associated threshold crossing Flags.

3 **Polling:** The repeated host-driven retrieval of module status data by the host. See also Reporting.

4 **Port:** A point of signal entry or signal exit at the module boundary. Sometimes short for Port Function.

5 **Port Function:** The functions and resources within the module that process a signal before leaving and after
6 entering the module through a module port. Colloquially this is often called an interface.

7 **Post-cursor equalization:** Module Rx output means used to reduce post-cursor ISI.

8 **Power Mode:** The Power Mode of a module is either **High Power** or **Low Power**. The maximum module
9 power consumption in Low Power Mode is defined in the relevant hardware specification. The maximum module
10 power consumption in High Power Mode is module implementation dependent and is advertised. The purpose
11 of the Low Power Mode is to ensure that a plugged-in module can be examined by the host and rejected (not
12 operated) if, e.g., the operational power consumption would exceed the host's capabilities. Power Mode, the
13 power level the module is permitted to consume, should not be confused with the Power Class which represents
14 an advertised range of power consumption limits defined in the relevant hardware specification.

15 **Pre-cursor equalization:** Module Rx output means used to reduce pre-cursor ISI.

16 **Pulse Amplitude Modulation, four levels (PAM4):** A modulation scheme where two bits are mapped into
17 four signal amplitude levels to enable transmission of two bits per symbol.

18 **Quiescent:** An output is **quiescent** when the signal being output **cannot be detected as a useful input**
19 signal by a remote input. An input receiving the signal from a quiescent output must detect a loss-of-signal
20 (LOS) condition, as specified in the relevant interface standard. Note that there may be multiple methods to
21 achieve quiescence (e.g. for modulated signals, either the carrier or the modulation or both can be removed).

22 **READ:** This all-capsitals spelling denotes a management register READ operation of one or more bytes, which
23 is executed via a management communication interface.

24 **Real Time Value (Sample):** A **real time value**, **current value**, or **sample** (preferred) is a value that is
25 sampled or evaluated (conceptually) at the time of a read request and returned with the read request.
26 Depending on the observable, sampling a new value may include some estimation or measurement time, usually
27 short, such that read operations faster than the sample measurement or estimation interval will yield the same
28 value. See also Statistics.

29 **Receiver, Rx:** An electronic or opto-electronic component (Rx) that converts an input signal (optical or
30 electrical) to an electrical (non-retimed, retimed, or processed and reformatted) output signal. In this
31 specification **Receiver (Rx)** refers to the module function as a whole, not to interface adapter circuitry. A
32 module's Rx converts media interface signals to host interface signals.

33 **Register:** In this specification a **register** denotes a host accessible management memory element with
34 associated management functionality (displaying status or exercising control) which can be accessed by read
35 or write primitives. In CMIS there are **single-Byte registers** (also referred to as Bytes) and **multi-Byte**
36 **registers**. Registers may be sub-structured into **fields**. Sometimes **location** is used as a synonym.

37 *Note: Effective addressability refers to the fact that memory may be mapped dynamically into a host addressable
38 memory window. Note that addressability has no implication of how a register is implemented, be it a true
39 register interpreted by module hardware or just a memory location interpreted by module firmware. Apart from
40 registers, there are also addressable memory elements that serve only for (temporary) storage without fixed
41 meaning aside (e.g. memory used for message transfer).*

42 *Note: CMIS Multi-Byte registers require a series of Byte accesses, which raises questions of data coherency for
43 the register value during a register modification (see section 5.2.3).*

44 **Reporting:** In this specification, setting a binary Flag (potentially) causing a host Interrupt has the purpose to
45 actively **report** some condition to the host. Unlike most other management interactions, reporting originates in
46 the module (or target). The term **indication** is also used, especially when conditions are reported.

47 **Reporting Register:** A reporting register is updated by the module to provide information to the host.

48 **Resource Module:** A module supporting the Module State Machine only, not providing transmission service,
49 with managed functionality described in a separate resource module supplement specification. Example: ELSFP
50 modules providing laser light to transmission modules, typically under control of a host system.

1 **Result Register:** A result register contains the results to be consumed by the host. For monitoring processes
2 with repeated result generation, the result updates occur either time-based or event-based (with handshake).

3 **Sample:** See Real Time Value and Result Register.

4 **Sampling:** Sampling data from a permanently running process copies an instantaneous current value (real
5 time value) into a stable sampling **result register** (sample register). The stored sample can be consumed by
6 the host while the internal process continues. In many cases a side effect of sampling is that some internal
7 process is reset and restarted.

8 *Note: Colloquially this sampling process is also called 'latching' of results. This "latching" terminology, however,
9 is discouraged because latched bits (used to record some binary state) often have clear-on-read access mode,
10 which is often not the case for numerical results.*

11 **Separable (interface):** In this specification, a physical interface or interconnection is called **separable**
12 (disconnectable) when the physically interconnected parts (e.g. module and optical fiber) can be disconnected
13 and detached from each other. Separable interfaces are connectorized and **disconnectable**.

14 **Signal:** The term signal is ambiguous. **Physical** signals are carried over physical interface **Lanes**. These signals
15 are sometimes called **single-lane** signals. Differential pairs of physical signals carrying one serial data stream
16 are usually considered to belong to the same Lane. **Logical** (data) signals may be carried over multiple physical
17 lanes and are sometimes called **multi-lane** signals.

18 **Signal Integrity (SI):** host side electrical layer equalization (or pre-equalization) **parameters** and
19 **functionality** to counteract signal distortions that come with high-speed electrical signals or long and
20 connectorized electrical signal paths. The SI related functionality ranges from fixed pre-defined equalization
21 using programmable equalizers configured at link startup time to run-time adaptive equalizers.

22 **SNR** (Signal to Noise Ratio): Here, the ratio of signal power to the noise power, always (implicitly) expressed
23 in decibels

24 **Squelching:** The act or function of intentionally muting an enabled module output to become quiescent such
25 that no signal is detectable at the remote end. Squelching is often an automatic module reaction, e.g. when no
26 module input is available for the output in question. This automatic control reaction (which can be disabled) is
27 also called squelching colloquially but should better be called **auto-squelching**. It is also possible to squelch
28 an output on host request. This is called **forced squelching**.

29 **State:** The term state is used in the context of conceptual "state machines" that are described by state
30 transition diagrams and indicate modal (i.e. state-dependent) behavior.

31 **State Machine:** In specifications, a (finite) **state machine** (FSM, SM) is a conceptual object obeying the state
32 change behavior specified in a state transition diagram or, equivalently, a state transition table.

33 **State Machine Modeling:** The syntactical and semantical rules for the elements allowed or required in a state
34 transition diagram establish a well-defined modeling language, which can be translated into an executable state
35 machine model. In CMIS, two semi-formal Moore state machine types are specified, while the actual behavior
36 in each state is specified in natural language.

37 **State Transition Diagram:** The states and state transitions of the most basic state machine (Moore SM) can
38 be described with a formal diagram of named nodes (representing states) and arrows (representing state
39 transitions) where the arrows are tagged with a predicate expressing a (possibly compound) logical condition
40 that must be true for the transition to occur. The semantics of what the states represent or the general behavior
41 of an object with state-dependent behavior described by the state transition diagram are not formalized in a
42 Moore SM and are therefore added as informal text.

43 **Static:** A field is **static** when the values in the field are not changed by the module except at power-up, or
44 because of a reset. Quasi static fields that are not changed by the module as long as a programmed Application
45 persists are also called static fields, for short.

46 **Statistics:** A **statistic** or **statistics** represent aggregated results computed over a set of **samples** that have
47 been collected in a currently running or past monitoring or **statistics interval**. Typical statistics include
48 minimum and maximum sample value observed in the monitoring interval, and the average over all sample
49 values. Taking samples for the statistics may require its own (shorter) estimation or **measurement interval**.

1 **Status** (data): The term status is used in a broad sense for any host **readable** (but usually not writeable) information. It includes especially variable binary indications and sample values, as well as constant data like advertisements¹. See also Controls, Configuration and Trigger.

4 **System Interface** (Application or Module): System interface modules have been introduced in chapter 1. The
5 associated functionality is called a system interface application. In practice, such modules are usually, but
6 somewhat vaguely, referred to as PHY, interface, or transceiver module. Compare with client encapsulation
7 modules.

8 **TEST**: This all-capsitals spelling denotes a test to determine if the module is ready for READ or WRITE access.

9 **Transaction**: The implementation of register ACCESS primitives across a Management Communication
10 Interface (MCI) consists of one or more acts of communication, also called bus transactions.

11 **Transaction Delay** Some Management Communication Interface (MCI) implementations have means for the
12 module to slow down or delay completion of bus transactions implementing register ACCESS, e.g., by **clock**
13 **stretching**. Delaying ACCESS completion is a means to accommodate module internal processing needs.

14 **Transaction Hold-Off (Rejection)**: The implementation of ACCESS Hold-off (a.k.a. ACCESS Rejection) in a
15 Management Communication Interface using bus transactions is called transaction hold-off (a.k.a. transaction
16 rejection, or negative acknowledgement, NACK).

17 **Transmitter, Tx**: an electronic or optoelectronic circuit (Tx) that converts an electrical input signal from a host
18 to a signal suitable for the communications media (optical or electrical). In this specification **Transmitter (Tx)**
19 refers to the module function as a whole, not to interface adapter circuitry. A module's Tx converts host interface
20 signals to media interface signals.

21 **Trigger**: The effect of writing a value of 1b to a (usually self-clearing) trigger bit in a register is that some
22 process is started (triggered). See also Controls, Configuration and Status.

23 **Uniplex**: Applications, configurations, or topologies that implement a uniplexing (1:1) function as a special
24 case of a (N:1) **multiplex** with **encapsulation** of one (N=1) host side host signal into one media side network
25 signal. Note that a system interface application is not a uniplex application (there is no encapsulation into an
26 independent media side network signal).

27 **Unit** (defined for a register): A specified **unit** of a numerical register defines what quantity is associated with
28 a numerical register value of 1. A unit is often specified as the product of a conventional measurement unit
29 (1m, 1kg, 1s, etc.) and a dimensionless scaling factor. Colloquially, this quantity is sometimes loosely referred
30 to as the "LSB". In other formulations the unit of a register value is expressed as "in increments of".

31 **Upper Memory**: The 128 Bytes of host addressable memory addressed by Byte addresses 80h through FFh.
32 The actual content of Upper Memory depends on the host-controlled selection of Bank and Page to be mapped
33 by the module into Upper Memory.

34 **User**: The word user refers to any party who has access to a CMIS module, including module makers, module
35 OEM users, host systems.

36 **Valid** (signal): A signal on a lane is called **valid** if its (electrical or optical) **physical signal characteristics**
37 are **stable** (i.e. not transient) and **proper** (i.e. compliant to the pertinent interface specification for a useful
38 state of transmission), and if the **data** carried on that physical signal is **correctly formatted** at the level of
39 data processing performed in the module. Signal validity plays a crucial role during initialization, because only
40 a valid input signal may allow a module or host to successfully and correctly initialize all resources required to
41 propagate or process the signal. Providing only valid output signals avoids scenarios where the receiving end is
42 misguided in its initialization, which otherwise may lead to undesired link flaps.

43 **Volatile** (register or memory location): A volatile register or memory location is not persistent and does not
44 preserve its value over a regular reset or a power cycle.

45 **Warning**: A warning is like an alarm. Only the implied severity is different. The reporting mechanism via Flags
46 is the same as for alarms. See Alarm.

47 **WRITE**: This all-capsitals spelling denotes a management register WRITE operation of one or more bytes, which
48 is executed via a management communication interface.

¹ The colloquial notion of status is usually assumed change dynamically.

3.4 Data Types

The following data types allow this specification to define the interpretation of data in multi-byte registers concisely.

The **default** Endianness (storage order) for numerical data types is defined in section 8.1.3.5 (**Big Endian**)¹.

In case of non-default storage order (Little Endian), the non-default Endianness must be explicitly specified.

ASCII ASCII character

ASCII[n] String of **n** ASCII characters (not 0 terminated)

Bool 1-bit Boolean value with encoding TRUE = 1 and FALSE = 0

Note: Boolean variables are often named such that the interpretation is self-explanatory (e.g. Bool Initialized instead of Bool InitializationStatus)

F16 A compact 16-bit floating point data type (originally defined in SFF-8636) used to represent non-negative real numbers. Representable values are expressed in the form $m \cdot 10^{s-24}$ where the mantissa m ranges from 0 to 2047 and the scaled exponent s ranges from 0 to 31. Hence, the smallest representable non-zero number is $1.0 \cdot 10^{-24}$ and the largest representable number is $2.047 \cdot 10^{10}$. An F16 value is stored in two bytes (where byte 1 is stored at the lower byte address) with the following representation (which is Big Endian for the 11-bit mantissa):

Byte Number	Bits	Description
1	7-3	Scaled Exponent (s)
	2-0	Mantissa (m), most significant bits 10-8
2	7-0	Mantissa (m), least significant bits 7-0

S8 8-bit signed integer (2's complement) with value range -128 to 127

S16 16-bit signed integer (2's complement) with value range -32768 to 32767

S32 32-bit signed integer (2's complement) with value range -2,147,483,648 to 2,147,483,647

S64 64-bit signed integer (2's complement) with approximate value range $-9.22 \cdot 10^{18}$ to $9.22 \cdot 10^{18}$

U8 8-bit unsigned integer with value range 0 to 255, stored in a single byte

U16 16-bit unsigned integer with value range 0 to 65535, stored in two bytes

U32 32-bit unsigned integer with value range 0 to 4 294 967 295, stored in four bytes

U48 48-bit unsigned integer with value range 0 to 281 474 976 710 655 stored in six bytes

U64 64-bit unsigned integer with value range 0 to 18,446,744,073,709,551,615, stored in eight Bytes

U<1> <1>-bit unsigned integer with value range 0 to $2^{<1>} - 1$, stored in a field

X16 Arbitrary 16-bit numerical data type (F16 or S18 or U16)

3.5 Data Operations

A Bit (variable) is said to be **set** when its current value is 1 and it is said to be **cleared** when its current value is 0. In other contexts, however, the same wording can have action semantics, where a Bit (variable) is said to be **set** when a value of 1 is written and **cleared** when a value of 0 is written. This ambiguity between state and action is best avoided by using active grammar (someone sets or clears a Bit) for activities and passive grammar for states (a bit is set or cleared).

To shorten specification text, the following verbs are used with specific meaning when describing management operations on binary data in registers or fields.

Set Write 1b (set X means X = 1b)

Clear Write 0b (clear X means X = 0b)

Raise Change value from 0b to 1b (raise X means X = 1b when X = 0b before the assignment)

Cease Change value from 1b to 0b (cease X means X = 0b when X = 1b before the assignment)

Assert Assign a numerical value representing the Boolean value TRUE

Deassert Assign a numerical value representing the Boolean value FALSE

Note that all these verbs except 'set' allow the reader to distinguish active operation activity from passive operation result (e.g. 'raise' versus 'raised').

¹ The most significant byte is at the lowest address, the least significant byte is at the highest address.

4 General Concepts

As a management interface specification CMIS defines the management interface and associated protocols for all required and certain optional management interactions between a CMIS aware host and a CMIS compliant module that are relevant for the host using the module in an application.

The remainder of this chapter is intended to provide high level orientation and informal overview. Normative specifications are found in the subsequent sections.

4.1 CMIS Managed Module Classification

Two main functional types of CMIS managed modules are distinguished

1. Transmission Modules providing data transmission service to a host system
 - a. Modules with detachable (connectorized) transmission media and unknown far end peer
 - b. Cables and cable assemblies with fixed media and fixed far end connectors
2. Resource Modules providing resources, usually to Transmission Modules, in a host system

4.1.1 CMIS Compliant Transmission Modules

A CMIS compliant **transmission module** has two main physical interfaces related to a module's mission of bidirectional signal or data transmission, and one CMIS compliant management interface for the host to manage module operation:

- The **host interface** connects a module and its host system for the purpose of exchanging high-speed electrical signals that the module transmits to and receives from a remote system. The host interface is decomposed into a set of host lanes. The host interface is often viewed as a bidirectional entity but sometimes it is viewed as a pair of two related unidirectional interfaces. This will be clear from context.
- A **host lane** carries a differential electrical high-speed signal over a named pair of electrical wires (or named contacts in a connector). Sometimes the pair of unidirectional host lanes used for bidirectional transmission are also referred to as a host lane. This will be clear from context.
- The **media interface** connects a module and the transmission media that interconnect the module to its far end peer¹. The media interface carries high-speed electrical or high-speed optical signals. The media interface is decomposed into a set of media lanes.
- A **media lane** carries one media dependent high-speed signal over a pair of copper wires, optical fibers, or optical wavelength channels (in DWDM or CWDM fiber links). The high-speed signal on a media lane may also comprise a multiplex of lower rate signals (TDM or CDM). Sometimes the pair of unidirectional media lanes used for bidirectional transmission are also referred to as a media lane. This will be clear from context.
- The **management interface** allows the host to manage the module's functionality and operation. It carries low-speed signals for management data communication and some discrete low-speed signals for low level control and status.
- Additional interfaces (interconnections) e.g. for power supply, reference clock signals, pulse per second signal, laser light, etc. will or may exist but are not considered here, unless there are management aspects to be considered.

A CMIS compliant module provides standardized signal or data transmission functions to its host system and each type of transmission function is referred to as an **Application**. In support of an Application the module propagates or processes signals between one or more host lanes and one or more media lanes.

A CMIS compliant module may also provide resources for parallel and mutually independent Applications. The lanes and module internal resources carrying one Application are called a **Data Path**. A Data Path may use only a subset of the host interface lanes or of the media interface lanes.

In CMIS compliant modules, unless advertised differently², symmetry is expected between the Tx hardware structure and the Rx hardware structure; for example, Tx lanes that are multiplexed in the Tx are likewise demultiplexed in the Rx. Each Application selected by the host is applied to the same Tx and Rx host lanes.

As described in section 7.8, the data structure of the CMIS management interface favors modules with 8, 16, or 32 lanes, because the core management features are prepared for modules with up to 4 x 8 lanes; future extensions, and applications using only a subset of lanes are, however, possible.

¹ For cable assemblies, the media interface and the media lanes are internal components

² This revision does not contain asymmetry advertisements.

1 Due to the way how support of Applications is advertised, the **Data Path** of a single **Application** is always
2 limited to **at most eight host lanes** and at most eight media lanes (even when a module supports 2x8 or 4x8
3 lanes).

4 In technical terms, explained in chapter 6, with the exception of passive cables or modules with very low power
5 consumption, a transmission module supports the Module State Machine for host-controlled power-up, and it
6 supports data transmission Applications on Data Paths, or Applications multiplexing host side Data Paths into
7 media side Network Paths, including the associated Data Path or Network Path State Machines.

9 **4.1.2 CMIS Compliant Resource Modules**

10 With the advent of co-packaged optics architectures, a new class of modules has emerged that do no provide
11 data transmission services themselves, but instead provide technical resources such as a bank of laser sources
12 to other modules that do. These modules are called **resource modules**.

13 Resource modules can also be managed by CMIS. They lack all aspects of Application management on Data
14 Paths, but they offer resource module type specific management facilities, and they are booted in a controlled
15 way, like transmission modules.

16 In technical terms, explained in chapter 6, a resource module supports the Module State Machine, but no Data
17 Path State Machines or Network Path State Machines.

18 For each fundamentally different type of resource module, a separate CMIS supplement specification is required,
19 each specifying the management memory structure and content of a management memory segment reserved
20 for resource modules.

4.2 CMIS Management Protocols and Layers

The CMIS management interface is best understood in terms of the management primitives exchanged across that interface and in terms of the protocol layers that govern the exchange of these management primitives (top to bottom):

- The **management application layer** at the top specifies **meaning** and **effects** of management operations on functions or behaviors of the module as well as any associated **behavioral rules** and **protocols**. This layer and its representation as a set of management **hardware signals**, management **register accesses**, or **messages exchanges** is described in chapters 6, 7, 7.8, 9, and 10.
- An optional **message exchange layer (CDB)** implemented on top of the essential register access layer provides primitives for a command-reply type of management message exchange between host and module. Implementation of this layer is described in sections 7.2 and 8.23.
- The essential **Register Access Layer (RAL)** provides the primitive management operations in CMIS which are **reading** and **writing** bytes from or to a 256-byte addressable memory window. The mechanisms and protocols pertinent to this layer are defined in chapter 5, whereas addressing extensions required to effectively access a larger management memory space are described in chapter 7.8.
- The **data transfer layer** underlying the register access layer provides the data transmission protocols required to let the host and the module exchange basic management operations (and associated data) via the physical interconnection layer. A set of data transfer stack variants is described under the name **Management Communication Interface (MCI)** in Appendix B.
- A small **management signaling layer (MSL)** exists in parallel to the MCI and provides electrical low-speed control and status signals with directly assigned management functionality, such as module reset or interrupt request. The MSL is defined in chapter 5.
- The **physical interconnection layer** is the lowest layer comprising the electrical and mechanical interface elements (that are connected to carry elementary signals) as well as the associated protocols i.e. rules for using those interface elements. Since CMIS is targeted to multiple form factors, the physical layer is defined in the relevant module form factor specification. In CMIS, generic signal names are used instead of form-factor specific signal names. Refer to Appendix A for the association between generic signal names and form-factor specific signal names. Electrical interfaces are also specified in Appendix B by reference to other standards [1] or form factor hardware specifications such as [2][3][4].

The following Figure 4-1 provides a pictorial illustration of these protocols, layers, and their relationships

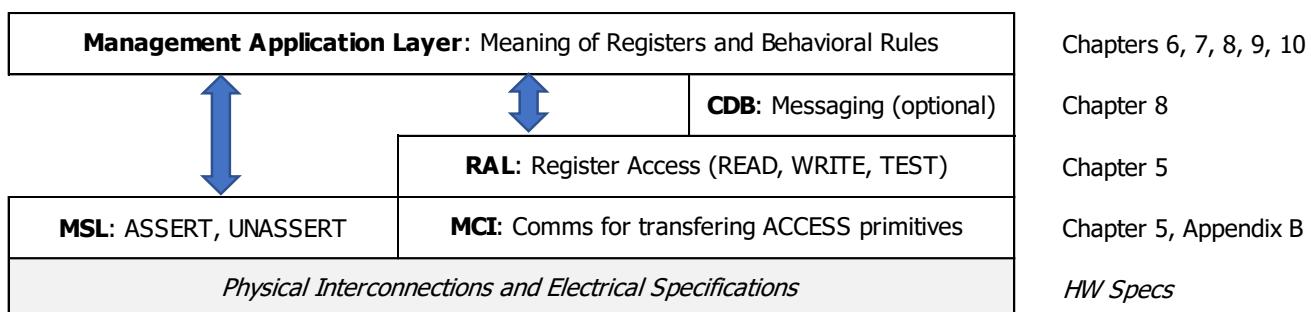


Figure 4-1 CMIS Management Protocols and Layers

The hardware specification of a form factor that supports or requires CMIS based module management contains mechanical and electrical specifications and defines which data transfer stack described in Appendix B is used.

4.3 CMIS Compliance and Vendor Specific Extensions

CMIS is designed to enable two important features

- **Vendor-agnostic standard management:** CMIS compliant hosts shall be able to manage optional or required standardized functionality of CMIS compliant modules using only CMIS standardized management functionality
- **Optional custom extensions:** CMIS compliant modules shall be able to provide optional enhanced custom functionality and optional custom management extensions

Note: Usually access to custom extensions may be limited or protected, but this is not technically assumed

Custom extensions require hosts that are entitled, aware, enabled, and willing to use those custom extensions. Any custom extension should therefore offer two important non-obvious properties:

- **Optionality:** A CMIS compliant host shall always be able to use a module with custom extensions as a plain CMIS compliant module, using only CMIS compliant management of standard functionality (as if there were no extensions).
- **Compliant default behavior:** A CMIS compliant module plugged into a CMIS compliant host slot shall initially provide its advertised standard functionality, independent of any vendor specific custom management extensions present. The module must not assume that its current host is capable to use or aware of its custom extensions.

Note: A host may be ignorant of custom management extensions, or may not have access, or may be not willing to use them in a given context, no specific scenario is assumed here.

Note: One subtle consequence to point out here is that all custom interrupt sources must remain silent until a host who is aware of and ready to use these extensions explicitly unmasks the custom interrupt sources.

4.4 CMIS Interworking and Compatibility

Further CMIS revisions with the same major revision number will be backwards specification-compatible extensions, allowing host to module interworking in cases when host and module implement different CMIS revisions¹.

¹ Note that hosts can learn the CMIS version supported by the module, which allows hosts to adapt to modules supporting earlier CMIS revisions. Working with newer modules may or may not be possible. See Appendix G.

5 Management Interface

The Management Interface may be understood as the means (mechanisms) required by, or available to, a management application on a CMIS host to manage a CMIS compliant module; its conceptual service interface for a host application can be horizontally partitioned as follows (see Figure 4-1 in section 4.2):

- A **Management Signaling Layer (MSL)** provides **management signals**, i.e. **electrical** control or status signals¹ with directly assigned (fixed or programmable) management functionality, such as module reset or interrupt request.
- A **Register Access Layer (RAL)** provides the elementary ACCESS mechanisms for **software** implementing the management application to READ from and WRITE to Bytes in the directly addressable management memory window of a module, where the assigned management functionality depends on the accessed address, i.e. on the **register map**, and on the value being read or written.

The side-by-side combination of RAL and MSL provides a set of basic **management mechanisms** (or primitives) available to a management application.

Due to the physical separation between a managing host and a managed module, data communication is needed to implement the basic READ and WRITE mechanisms of the RAL across the physical interconnection (which itself, somewhat confusingly, is also often referred to as the management interface). The data communication part of the vertical **protocol stack** (see Figure 4-1) used for this purpose is called the **Management Communication Interface (MCI)** and is specified in Appendix B.

Note: This vertical partitioning serves abstract specification purposes and does not constrain implementation. Especially there is no precise specification of the services and interactions at layer boundaries to which an implementation would have to adhere to.

5.1 Management Signaling Layer (MSL) Control and Status Signals

The **Management Signaling Layer (MSL)** consists of discrete control or status hardware signals which carry a module management function. These signals are called MSL signals, also known as low-speed signals.

The generic MSL signal names defined below are used in CMIS specification text and mapped to form factor dependent signal names in Appendix A.

In CMIS the binary values of MSL signals are denoted abstractly as **ASSERTED** or **DEASSERTED**.

The mapping of these abstract values to physical signal values or representations is defined in Appendix A.

The following MSL signals are **required** by CMIS

- a control signal allowing the host to request a module reset (**Reset**)
- a status signal allowing the module to assert an interrupt request (**Interrupt**) to the host

The following MSL signal is **optional**² in CMIS modules³

- a signal allowing the host to control full or partial power up of the module (**LowPwrRequestHW**)

When the module has no dedicated HW pin carrying an externally supplied LowPwrRequestHw signal it shall behave as if a LowPwrRequestHw signal was present and ASSERTED.

The usage of these management control signals in CMIS is described in chapters 6 and 7.8.

*Note: For different hardware form factors, there may be additional **form factor specific management signals** with fixed or programmable meaning. Such form factor specific **MSL extensions** may also need to be managed, in terms of advertisement, administration, and status reporting, i.e. they may require associated **register map extensions**. Such extensions are to be defined in one or more CMIS supplement specifications (see section 2.1.2), with any management registers allocated on a Page that CMIS has restricted for managing such MSL extensions (see section 8.8).*

¹ Sometimes referred to low-speed signals in hardware specifications.

² For optional MSL control signals, it is explicitly specified which of the signal values is assumed by the module, when there is no electrical pin or contact carrying the MSL control signal.

³ Signals to indicate module presence or signals to select the module as a target on a shared MCI (ModSel) may or may not be present. The latter is considered part of the MCI definition.

5.2 Management Register Access Layer (RAL)

Read and write access to the addressable memory of a module (as described in chapter 7.8) are the basic module management operations available to a host. These operations for memory or register access are implemented (emulated) by transactions over the Management Communication Interface (MCI) described in section 10.2 and Appendix B.

5.2.1 Registers in Addressable Memory

The directly host addressable management memory consists of 256 Bytes and is divided in two segments, called **Lower Memory** (Byte addresses 00h through 7Fh) and **Upper Memory** (Byte addresses 80h through FFh).

Note: The (indirectly) accessible management memory of a CMIS module is usually larger but requires a dynamic mapping mechanism for the Upper Memory which is described in chapter 7.8.

In the remainder of this section, read and write access to the directly addressable management memory of a module is simply referred to as register access (**ACCESS**).

5.2.2 Register Access Methods for Addressable Memory

Note: The following definitions of READ, WRITE, and TEST primitives are for specification purposes only and do not imply that software implementations must use drivers with interfaces of that signature.

For timing specification related to management register access see section 10.2.

5.2.2.1 Read Access (READ)

A read access can be modeled by the following abstract function:

VALUE = **READ**([ByteAddress,] N), $1 \leq N \leq N_{\max}$

A host may read N bytes of addressable module management memory in a READ access.

By default, $N_{\max} = 8$. When **full page read** is supported (as advertised in 01h:251.4) then $N_{\max} = 128$.

Note: Full page read may become a mandatory feature in the future. It is recommended to support it.

Note: A READ of more than 8 bytes may be allowed when specified explicitly in chapter 7.8.

The starting byte address for the read access is optionally specified by the host (ByteAddress). If omitted, the module uses an internally maintained **current byte address** as the starting byte address (see section B.1).

A successful READ access returns N bytes by VALUE.

Side effects of a READ (such as e.g. clear-on-read) may be specified for certain Bytes (see section 7.8). Such side effects are not guaranteed to be completed at the time when the READ is finished.

A READ access may temporarily be **rejected** by the module, for generic reasons specified in section 5.2.4 and specific reasons specified in chapter 7.8.

A rejected READ is modeled by returning an abstract REJECTED value that is distinguished from all possible VALUE byte strings returned by a successful READ.

Note: Syntactic or representational details are not important here, but the value REJECTED could be imagined, for example, by an X16 value outside of the X8 subrange representing possible byte values.

A rejected READ access can simply be retried until eventual success.

5.2.2.2 Write Access (WRITE)

A write access can be modeled by the following abstract function:

SUCCESS = **WRITE**(ByteAddress, VALUE, [VALUE, ...])

A successful WRITE writes a sequence of up to eight given byte values into the addressable memory of the module, beginning at the given ByteAddress and returns SUCCESS = TRUE.

Note: A WRITE with more than 8 bytes may be allowed when specified explicitly in chapter 7.8.

A WRITE access may temporarily be rejected by the module, for generic reasons specified in section 5.2.4 and specific reasons specified in chapter 7.8.

A rejected WRITE access has no effect in the target and is modeled by returning SUCCESS = FALSE.

A rejected WRITE access can simply be retried until eventual success.

1 The host shall not include a mixture of volatile and non-volatile registers in the same WRITE access.

2 **5.2.2.3 Test Access (TEST)**

3 Apart from certain advertised maximum durations, the host does not know in advance when exactly a module
4 is guaranteed to be ready and to accept the next READ or WRITE without rejection.

5 Apart from just trying and re-trying the next desired access, the host can explicitly test for module readiness.

6 This access readiness test can be modeled by the following pseudo-function

7 ACCESSIBLE = **TEST()**

8 A TEST access simply returns if the module is ready to accept a WRITE or READ access.

9 *Note: TEST is the register layer abstraction of the I2CMCI acknowledge polling primitive in section B.2.5.5*

10 **5.2.3 Register Access Rejection (Access Hold-Off)**

11 Any READ or WRITE register access may temporarily be **rejected** by the module, either for generic reasons
12 specified in section 5.2.4 or for specific reasons individually specified in chapter 8.

13 A host has several options to deal with temporary access rejection. A host may

- 14 • schedule access based on worst case timing specifications (see chapter 10)
- 15 • perform retries of a rejected access until eventual success
- 16 • TEST for readiness prior to a READ or WRITE access (a.k.a. acknowledge polling).

17 **5.2.4 Register Access Sequencing and Synchronization Requirements**

18 *Note: In this section, specifications are limited to the elementary READ and WRITE accesses to addressable
19 management memory. Sequencing and synchronization requirements for interactions at the management
20 application layer (e.g. Upper Memory Content Switch, CDB, VDM, etc.) are described in chapters 7 and 7.8.*

21 After a WRITE access the module ensures that the following **read-back conditions** are satisfied for all
22 registers with regular read-write (RW) access type.

23 When needed to fulfill these read-back conditions, the target temporarily **rejects** further READ or WRITE
24 accesses, for a maximum **ACCESS hold-off** duration specified in chapter 10.

25 *Note: The modification of a register by a WRITE access often leads to some intended **effect** in the module or
26 in its behavior. Like in most register-based interfaces, **completion of those effects is not synchronized** to
27 the end of a WRITE access, **nor confirmed** using general mechanisms. In those cases, where confirmation or
28 synchronization is required, handshaking or signaling mechanisms must be implemented using dedicated
29 elements of the register map, such as specific status indicators, Flags, handshaking bits, etc.*

30 **Single byte read-back condition:** After a properly terminated WRITE to a Byte address with RW access
31 type, an immediately following successful READ from that same Byte address yields the value just written.

32 **Multiple bytes read-back condition:** After a properly terminated WRITE to a Byte address range with RW
33 access type, an immediately following successful READ from any part of that Byte address range yields the
34 value of that part just written.

35 *Note: MCI mechanisms (see section B.2.6) can be used by the module to **reject** or **delay** a subsequent ACCESS,
36 either by rejecting or by internally delaying (clock stretch) a subsequent transaction while the read back
37 conditions are not yet satisfied.*

38 **5.2.5 Register Data Coherency**

39 Data coherency, of READ or WRITE accesses to a given byte range of interest, refers to module functionality
40 ensuring that concurrent accesses to that byte range, externally by the host and internally by the module, are
41 executed atomically (indivisible and under mutual exclusion), such that concurrent accesses to the byte range
42 as a whole are effectively serialized without interference. Neither the module nor the host will ever 'see' a
43 partially updated byte range that supports data coherency.

44 Both single-Byte READ accesses and single-Byte WRITE accesses guarantee data coherency.

5.2.5.1 Coherency of Size-Matched Multi-Byte READ Results

Size-matched READ access to one **scalar** multi-byte **read-only register** (as specified chapter 7.8) is atomic and delivers a coherent result, i.e. the module ensures not to update parts of a scalar read-only register during a size-matched READ of that multi-byte register.

This requirement applies in particular to any scalar 2-byte, 4-byte, or 8-byte status or monitoring register when read by the host with a single 2-byte, 4-byte, or 8-byte READ, respectively.

Note: See section 3.4 for a list of defined scalar multi-byte data types (e.g. F16, U16, S16, U64, X16, etc.).

READ access to other multi-byte registers, multiple registers or register arrays does not guarantee coherency.

If data coherency is not guaranteed but required, higher level access synchronization rules are explicitly specified, such as protocols using dedicated producer-consumer handshaking bits in the register map.

Note: Missing access synchronization rules in such cases would be considered a specification bug. Examples of explicit access synchronization protocols are given by the CDB command message send/receive protocol and the CDB reply message send/receive protocol described in sections 7.2 and 8.23.

5.2.5.2 Coherency of Multi-Byte WRITE Operations

Multi-Byte WRITE accesses (see section B.2.5.4) are generally **not** atomic, unless explicitly specified.

Note: Most writeable registers are byte sized. Examples of multi-byte registers with explicitly specified indivisible access include the PageMapping register and CMDID field in CDB messages.

If data coherency is not guaranteed but required, higher level access synchronization rules are explicitly specified, such as protocols using dedicated trigger or handshake fields in the register map.

Note: Missing access synchronization rules in such cases would be considered a specification bug. Examples of access synchronization protocol are the grouping of configuration data into Control Sets (see section 6.2.3) or the message composition phase before trigger to send a CDB command message (see section 8.23).

6 Core Management Features

The features, functions and behaviors described in this chapter usually assume the case of a **transmission module**. They are required for all CMIS managed modules, unless otherwise noted.

Note: Exceptions and simplifications exist for Resource Modules, Cable Assemblies, and Flat Memory modules. Resource Modules are programmable but do not provide data transmission service. Flat Memory modules support only (parts of) a constant 256 Byte read-only management memory map (i.e. they provide neither programmability nor dynamic status information). Cables Assemblies don't have an external media interface.

The letter of the text in this chapter applies only to **system interface** transmission modules and applications, while the extended management functionality of **client encapsulation** modules and applications is specified separately, in section 7.6, by defining changes and extensions to the management functionality of system interface modules described here. For these more complex but optional applications, this chapter must be read carefully and with a slightly different interpretation, as described in section 7.6.

Note: See the introduction in chapter 1 and the glossary definitions in section 3.3 for the distinction between system interface applications and client encapsulation applications, a.k.a. multiplex applications.

Note: The glossary defines some terms like Application and Data Path with a broad meaning that includes client encapsulation applications, while this chapter's text literally applies to system interface applications only. Some of the statements in this chapter are therefore true only for system interface applications, while the corresponding correct statement for client encapsulation modules becomes clear only with the interpretation put forward in the glossary and in section 7.6. Beware of confusion.

6.1 Transmission Module Management Basics

CMIS managed transmission modules have two physical interfaces that are related to the main purpose of signal transmission, a **Host Interface** and a **Media Interface**.

Note: The Host Interface is viewed as a device-to-device interconnection, whereas the Media Interface is viewed as a device to media attachment. For interface circuitry within the module, this difference is not important.

Note: This section does not apply to CMIS managed resource modules.

Note: For cable assemblies, the transmission media are fixed and not detachable from the connectors on both ends, hence the Media Interface may sometimes be considered unimportant.

6.1.1 Host Interface

The Host Interface is the high-speed electrical interface (interconnection) between a module and a host system.

The Host Interface carries signals travelling from the host into the module, referred to as **transmitter input** signals, and signals travelling from the module into the host, referred to as **receiver output** signals.

All electrical signals carried over the Host Interface are transmitted differentially over wire pairs, each of which is called a **host lane** (element of interconnection).

Note: Depending on context, a host lane may be viewed as unidirectional or as bidirectional.

The term **host lane** is also used to refer to the module internal resources (circuitry) associated with propagating or processing a host lane signal within the module.

6.1.2 Media Interface

The Media Interface is the high-speed electrical or optical interface between the module and the interconnection media (connecting the module to a remote peer). The interconnection media connecting to a remote end may consist of wires, of optical fibers, or of optical channels, a.k.a. wavelengths (in DWDM or CWDM links).

The Media Interface carries signals that travel from the module into the media, referred to as **transmitter output** signals, and signals that travel from the media into the module, referred to as **receiver input** signals.

All Media Interface signals are ultimately carried between the module and a remote peer. From a near-end viewpoint the Media Interface signals are carried either over electrical differential pairs (on a copper cable) or over optical wavelengths on physical fibers, generically called **media lanes**.

Note: Depending on context, a media lane may be viewed as unidirectional or as bidirectional.

The term **media lane** is also used to refer to the module internal resources (optics and circuitry) associated with propagating or processing a media lane signal within the module.

6.1.3 Memory Map Representations

A set of registers is associated with each lane, both of the host interface and of the media interface, to control processing and report status for signals at that interface. In this specification all **lane numbers** or other references to host lanes or media lanes refer to the respective registers or fields that control or describe those lanes. When the term 'lane' is used without reference to 'host' or 'media', a host lane perspective is assumed.

6.2 Transmission Module Functional Model

The following subsections define functional aspects that are common to all CMIS managed transmission modules, unless otherwise noted.

Note: The functional and behavioral management requirements related to module reconfiguration have grown to some complexity, at least in certain details. Significant simplifications apply, implicitly, to modules with limited or no programmability.

Note: CMIS 5.0 introduced an option for functional simplification with subtle behavioral variations for modules that support only so-called step-by-step reconfiguration, but do not support so-called intervention-free reconfiguration. For historical reasons, the text treats this simpler and often sufficient case as a restriction of the more general default specification; it is not optimized for readers interested only in those simplifications.

Note: Recall that the text in this section literally applies to so-called system interface applications. For the more complex but optional case of client encapsulation applications (a.k.a. multiplex applications), it must be read with a slightly different interpretation, as described in section 7.6.

6.2.1 Functional Transmission Module Capabilities – Applications

6.2.1.1 Applications and Application Instances

An **Application** is a type of functional transmission configuration that is characterized by specific signal propagation or signal processing between one or more host lanes and one or more media lanes, overall providing a well-defined signal or data transmission function to the host.

An **Application instance** is one implementation of an Application by a module

An Application is typically characterized and **specified by reference** to a **pair of industry standards**, one for the host interface and one for the media interface, each comprising one or more lanes. These interface standards define all necessary attributes for the respective interface, including the signaling rate (Baud rate), the signaling modulation format, the number of lanes, as well as any digital processing (such as lane alignment or FEC encoding and decoding) of the data stream (if applicable).

In this version of CMIS, Applications are limited to at most eight host lanes and at most eight media lanes.

A simple module may support only a fixed single Application, or several parallel Application **instances** (for Applications using fewer lanes than the module provides), while a versatile programmable module may support variant Applications or even multiple instances of one or more Applications operated in parallel.

*Note: Modules providing parallel Application instances of one homogeneous Application type are sometimes called **breakout** modules.*

6.2.1.2 Data Paths for Application Instances

The specific host and media lanes of a module that are used to implement one instance of an Application, together with all required internal module resources, is called the **Data Path** of that **Application instance**.

The host and media lanes of the Data Path **allocated** to an Application instance implement the host interface instance and the media interface instance of that Application, respectively.

Note: Here it is important to distinguish the host and media interface of a module and those of an Application. Only for modules specifically implementing only a single Application instance there is no difference.

The Data Path allocated to an Application instance always uses host and media lanes that are **numbered consecutively** on both module interfaces, starting with the lowest lane number, respectively.

Note: Due to this restriction, both the host interface and the media interface of a Data Path allocated to an Application instance are uniquely identified by the lowest lane number of the allocated host interface lanes and of the allocated media interface lanes, respectively.

Given the possible lane groups, i.e. the options for allocating host and media interface instance of an Application type, the Data Paths available for allocation to an Application instance are limited by the **restriction** that the n'th host interface instance option must be connected straight to the n'th media interface instance option.¹

Note: This rule assumes transceiver Applications with identical information bit rates on host and media side.

Note: This honors the fact that an Application is defined by a pair of interfaces. The restriction above means that the pairing is always fixed and straight without an option to switch the host to media side connectivity.

6.2.1.3 Multiple Application Instances and Multiple Applications

A module may support one or more instances of a single Application. Each configured instance of an Application is independent, from a host viewpoint, of all other concurrently configured Application instances (if any).

For example, a module may support an Application with a CAUI-4 host interface and a 100GBASE-SR4 media interface. In this example, a module could advertise support for one or two instances of the Application.

A module may also support multiple Applications.

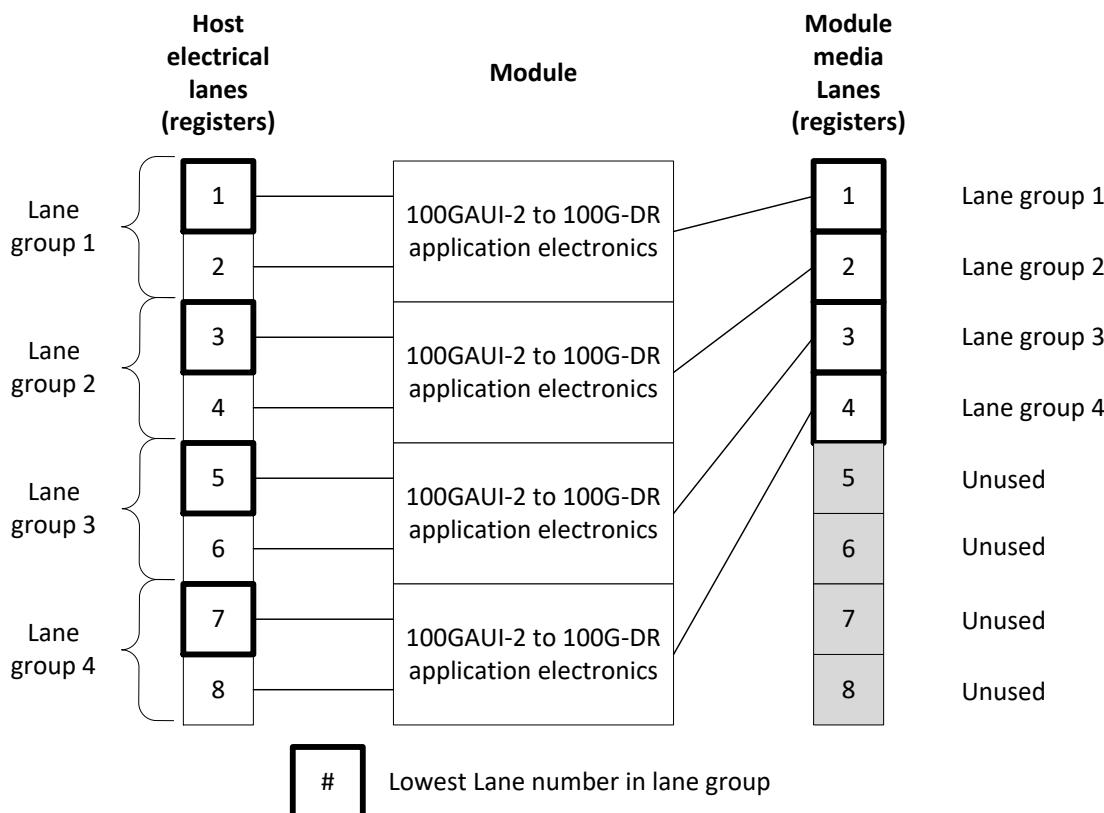
For example, a module may support one 400Gbps Application that is characterized by a 400GAUI-8 host interface and a 400GBASE-DR4 media interface combination, and a second 100Gbps Application that is characterized by a 100GAUI-2 host interface and a 100GBASE-DR media interface combination. The module may be programmable to work as one instance of the first Application or to work as up to four instances of the second Application.

Each instance of each active Application is independent of all other Application instances, from a host viewpoint.

For each Application where the module supports multiple Application instances, a host needs to be able to determine which host lane group corresponds to which media lane group for each possible Application instance in the module. As specified in the previous section, modules associate the nth host lane group of the Application with the nth media lane group for the same Application.

This rule is best illustrated by means of an example. A module advertising support for four instances of a 100GAUI-2 to 100GBASE-DR Application must identify the lowest lane number for each instance on both interfaces. As shown in Figure 6-1 for this example, the Application can be assigned starting on host lane 1, 3, 5, or 7.

Figure 6-1 Lane Assignment Example (Case of Parallel DataPaths)



¹ See section 7.8 for an optional extension that introduces a host lane switching capability.

These possibilities are the first, second, third, and fourth lane groups. By rule, the first lane group on the media interface in this example, which starts at media lane 1, must correspond to the first lane group on the host interface, which starts on host lane 1. The second lane group on the media interface, which starts on media lane 2, must correspond to the second lane group on the host interface, which starts on host lane 3, and so on. Therefore, in this example, the module would advertise host lanes 1, 3, 5, and 7 are supported as the lowest numbered host lane for each instance of the Application, and media lanes 1, 2, 3, and 4 are supported as the corresponding lowest numbered media lanes.

Note: For modules with more than eight lanes, lane counting starts fresh in each group of eight lanes. For instance, the next host and media lane numbers in Figure 6-1 would both be 9. See section 8.1.3.3 for details.

6.2.1.4 Application Advertising (Application Descriptors and AppSel Codes)

The module uses **Application Descriptor** data structures to advertise the set of supported instantiations of the set of supported Applications: For each supported Application, a module advertises the lanes that can carry a host interface instance and a media interface instance of a supported instance of that Application, respectively.

6.2.1.4.1 Basic Application Advertisement (up to 15 Applications)

The Application Descriptors are stored by the module in three Application Advertising areas of the Memory Map (see Table 8-23, Table 8-58, Table 8-59), but conceptually they form a contiguous array of Application Descriptors, each one structured as shown in Table 6-1.

Each Application Descriptor is identified by a unique **AppSel** code (short for Application Selection code).

The AppSel code of an Application Descriptor is simply the **sequence number** of that Application Descriptor in the list of basic Application Descriptors provided by the module. Hence, with a bit of simplification, an **AppSel** code identifies or selects an **Application**, albeit in a **module dependent** way.

Note: The AppSel code identifying a specific type of Application is module dependent because the sequence in which the module advertises its supported Applications in the Application Descriptor list is not specified.

Basic advertisement provides space for advertisement of up to fifteen Applications. Applications identified by AppSel codes 1-8 are advertised in registers described in Table 8-23 and Table 8-58, and Applications identified by AppSel codes 9-15 are advertised in registers described in Table 8-59 and Table 8-58.

A module advertises at least one supported Application as its power-up default Application in the first Application Descriptor (identified by AppSel code 1), and it advertises any additionally supported types of Applications consecutively, in a module defined arbitrary order, in a contiguous array of Application Descriptors (starting from the descriptor identified by AppSel code 2).

A module advertises only supported Applications (which may include custom Applications).

Note: Examples of Application advertising scenarios are shown in Appendix C.

An **Application Descriptor** consists of five bytes and describes all possible instantiations of one Application.

The first byte (**HostInterfaceID**) identifies the industry standard for the host interface. The list of defined HostInterfaceID values can be found in [5] where each Host Interface ID is associated with an industry standard, host interface signaling rate, modulation format, lane count(s), and, implicitly, with signal or data processing characteristics defined in the referenced industry standard. The module defines the number of supported host interface lanes in the HostLaneCount field, described below. The module encodes the first unused entry in the list of Application Descriptors as FFh to indicate the end of the list of supported Applications.

The second byte (**MediaInterfaceID**) identifies the industry standard for the media interface. The list of defined MediaInterfaceID values can be found in [5] in one of the media interface code tables. The module identifies the applicable media interface code table in the **MediaType** advertising field in Byte 00h:85 (see Table 8-21). The media interface code tables implicitly identify the media interface signaling rate, modulation format, and standard-defined lane count(s) for each MediaInterfaceID. The module defines the number of supported media interface lanes in the MediaLaneCount field, described below.

Note: The MediaInterfaceID may be undefined for certain cables with non-detachable media.

Note: The combination of HostInterfaceID and MediaInterfaceID identifies one type of Application, whereas the following fields describe the possible instantiations of that Application on the lanes of the module.

The third byte (**Lane Counts**) defines the number of lanes for the host interface and for the media interface. These lane counts are derived from the standards identified in the first and second bytes.

The fourth byte (**HostLaneAssignmentOptions**) identifies the lanes where the Application is supported on the module's host interface. The module may support multiple instances of a given Application, so each Lane Assignment Options field identifies the lowest numbered lane in a consecutive group of lanes to which the Application can be assigned.

For example, a module supporting two instances of an Application with CAUI-4 host interface that can be assigned to host interface lanes 1-4 and 5-8 would advertise HostLaneAssignmentOptions as 00010001b, to indicate that the lowest numbered lane for assignment of an instance of the Application can be lane 1 or lane 5.

The fifth byte (**MediaLaneAssignmentOptions**) identifies where the Application instance is supported on the module's media interface. Note that the MediaLaneAssignmentOptions registers are located on Memory Map Page 01h as described in Table 8-58, separated from the first four bytes.

Note: The MediaLaneAssignmentOptions information is not required for flat Memory Map modules.

Table 6-1 Application Descriptor structure

Byte	Bits	Name	Description
First	7-0	HostInterfaceID	ID from Table 4-5 "Host Electrical Interface IDs" in [5]. FFh marks an unused Application Descriptor. <i>Note: The end of the list of supported Applications is indicated by an unused Application Descriptor with HostInterfaceID FFh.</i>
Second	7-0	MediaInterfaceID	ID from one of several " media interface IDs " tables in [5], where the relevant table is selected by the MediaType field (see Table 8-23) containing a ModuleTypeEncoding value as described in Table 8-20.
Third	7-4	HostLaneCount	Number of host lanes 0000b: lane count defined by interface ID (see [5]) 0001b: 1 lane, 0010b: 2 lanes, ..., 1000b: 8 lanes 1001b-1111b: reserved
	3-0	MediaLaneCount	Number of media lanes. For cable assemblies, this is the number of lanes in the cable. 0000b: lane count defined by interface ID (see [5]) 0001b: 1 lane, 0010b: 2 lanes, ..., 1000b: 8 lanes 1001b-1111b: reserved
Fourth	7-0	HostLaneAssignmentOptions	Bits 0-7 form a bit map corresponding to Host Lanes 1-8. A bit value of 1b indicates that the Application can begin on the corresponding host lane. The host lane numbers of a multi-lane Application are contiguous. If multiple instances of an Application are allowed each supported starting point is identified. If multiple instances are advertised, all possible combination of Application instances are supported concurrently.
Fifth	7-0	MediaLaneAssignmentOptions	Bits 0-7 form a bit map corresponding to Media Lanes 1-8. A bit value 1b indicates that the Application can begin on the corresponding media lane. The media lane numbers of multi-lane Applications are contiguous. If multiple instances of an Application are allowed each supported starting point is identified. If multiple instances are advertised, all possible combination of Application instances are supported concurrently. This field is not required for flat Memory Map modules.

*Note: Since an Application describes a **combination** of a host interface and a media interface, a particular HostInterfaceID or a particular MediaInterfaceID may occur in multiple Application Descriptors.*

Custom Applications

For cases where a module supports a standardized Host Interface with a HostInterfaceID that is included in [5] but uses a media interface that is proprietary or not yet listed in the media interface advertising ID tables in [5], the module can use a null ID (00h=undefined) or a custom ID for the media interface that the module supplier has established, combined with the standardized HostInterfaceID. The HostInterfaceID and the lane

counts provide the required information for the host to interoperate with modules that have unfamiliar or vendor-specific media types and future technologies.

Likewise, when a module supports a host interface specification for which no HostInterfaceID has been allocated in [5], a custom HostInterfaceID can be used, which, however, requires the host to be aware of its meaning.

Note: Timely allocation of a standard MediaInterfaceID or HostInterfaceID in [5] is strongly recommended when a new type of interface is targeted for standardized use.

6.2.1.4.2 Normalized Application Advertisement (up to 240 Applications)

To overcome the limited number of 15 Applications that can be advertised in the basic Application Advertisement registers, a module can advertise up to $16 \cdot 15 = 240$ Applications in so called Normalized Application Descriptors.

Support of Normalized Application Descriptors is optional and hence advertised (see Table 8-57).

A **Normalized Application Descriptor** is a contiguous 8-byte data structure (see Table 8-163) providing all basic Application attributes, including all information contained in the basic Application Advertisement registers.

To allow interworking with legacy host, the content of the first 15 Normalized Application Descriptor (**NAD**) instances must be mirrored into the basic Application Advertisement registers.

The list of NAD instances is stored systematically in up to 16 segments, each segment providing a contiguous array of 15 NAD data structures (see section 8.20). The segments are implemented as Banks.

Each NAD instance is identified by a unique Application Number (**AN**), which determines the Byte, Page, and Bank addresses of that instance in management memory.

Note: Recall that the sequence of Application advertisements is module defined; the same Application in two different modules will usually have different Application Numbers.

The NAD instances located in one segment (i.e. in one Bank) are referred to as an **NAD Block**, which is identified by the segment index, alias the Bank index (0 – 15).

Combining the **NAD Block** index (**B** = 0–15) and the basic **AppSel** code (**c** = 1–15) identifying the instances within the block results in the globally unique **Application Number N** = **B**·15 + **c**.

A module advertising n Applications uses Application Numbers from 1 to n in a module-defined order and marks the first unused NAD (with AN = n+1) by setting its first byte to FFh (unless n=240).

6.2.1.5 Application Selection and Instantiation

A host refers to an Application described in one of the (Basic) Application Descriptors by means of the **AppSel** code of that Application Descriptor. When Normalized Application Descriptors are used, a particular **Normalized Application Descriptor** is referred to by its **Application Number**.

Note: When reference to Application selection via its AppSel code is made in this specification, this implicitly includes the alternative of Application selection via its Application Number AN. The basic data structures using just an AppSel code are extended to allow storing the NAD Block index (see sections 8.17.1 and 8.18.2)

The actual configuration mechanisms and procedures to instantiate ("provisioning") and implement ("commissioning") one or more Applications is described in the following subsections 6.2.3 and 6.2.4.

Note: The module behavior when a host mistakenly selects an unused Application Descriptor is not specified.

6.2.2 Application Instances on Data Paths – Data Path Lane Assignment

As introduced in section 6.2.1.2, the lanes that are associated with one Application instance, together with all associated module internal resources, are collectively referred to as the **Data Path** carrying that Application instance.

In module configurations with multiple Application instances operating independently, separately, and in parallel, the Data Paths assigned to these instances do not share any lanes.

The host configures the module for its intended use by selecting one or more non-conflicting Applications, from the list of supported Applications in the Application Advertising registers, and by assigning or allocating a unique set of lanes as a Data Path to each of the intended non-conflicting Application instances.

The assigned Data Path must be advertised as being supported for this Application, and the lanes used by the Data Path must not create resource conflicts with the lanes of other currently configured Data Paths.

1 *Note: Most Data Path related attributes are duplicated as lane attributes on the lanes of the Data Path. The host configures these attributes on all lanes of the Data Path identically, and likewise the module reports Data Path status attributes on all lanes of a Data Path.*

4 All module internal resources associated with a Data Path are initialized and deinitialized as a group.

5 *Note: In case of resources internally shared between Data Paths, this statement is a simplification.*

6 Separate Data Paths operating in parallel on disjoint (non-overlapping) lane groups are initialized, operated,
7 and deinitialized independently, from a host viewpoint. In particular, the host can issue management operations
8 on separate Data Paths concurrently and asynchronously.

9 *Note: This means that the module must provide functionally independent control for each data path to the host.
10 However, module implementations are permitted to utilize shared physical resources for items such as lasers,
11 TECs, processors, etc., as long as the control mechanisms presented to the host for individual data paths and
12 their associated behaviors remain independent of other data paths.*

13 *In one example, a module advertises an Application that consists of a CAUI-4 host interface and a 100GBASE-
14 SR4 media interface combination. In this example, an eight-lane module could advertise support for one or two
15 instances of the Application. Each instance of the Application would use a separate Data Path, where each Data
16 Path includes four host lanes and four media lanes and is independent of the other Data Path.*

17 *In another example, a module advertises an Application that consists of a 400GAUI-8 host interface and a
18 400GBASE-DR4 media interface combination, and a second Application that consists of a 100GAUI-2 host
19 interface and a 100GBASE-DR media interface combination. The host may select one instance of the first
20 Application or up to four instances of the second Application. Each instance of each selected Application
21 becomes a separate, independent Data Path.*

22 **6.2.3 Data Path Configuration – Control Sets**

23 The fundamental configuration of a module's Data Paths is defined in so called Control Sets.

24 **6.2.3.1 Control Sets Concept**

25 A **Control Set** is a group of registers containing per lane configuration settings that are fundamental for the
26 transmission functions to be provided by Data Paths.

27 The contents of Control Sets are eventually used by the module during actual Data Path initialization or
28 reinitialization, i.e. when a provisioned Data Path configuration is commissioned.

29 There are two types of control sets. The **Active Control Set** (see section 8.10.6) **reports** the currently
30 provisioned configuration settings that are used (or to be used) by the module to control its hardware, whereas
31 a **Staged Control Set** (see sections 8.9.3 and 8.9.4) is used by the host to **define** new configuration settings
32 for future use, without immediate effect in the module.

33 *Note: This staging mechanism decouples the timing and sequence of host writes to the Staged Control Set from
34 the module actions to configure the module hardware once the Active Control Set is evaluated.*

35 Each programmable module implements the Active Control Set and at least one Staged Control Set.

36 The Active Control Set is read-only, and it normally reports settings representing the current hardware
37 configuration of a Data Path, except transiently between a nominal **provisioning** update of the Active Control
38 Set and the completion of a subsequent configuration **commissioning** procedure, in which the Data Path
39 hardware is actually initialized or reinitialized from the Active Control Set.

40 To **realize** a desired configuration that has been **defined** in a Staged Control Set, the host first **requests** the
41 module to replace the current configuration in the Active Control Set with the content of that Staged Control
42 Set (**provisioning**), and afterwards the module commits the new settings into hardware (**commissioning**).

43 Requesting a configuration provisioning procedure (or a combined provisioning and commissioning procedure)
44 is also called **Applying** the Staged Control Set, which is described in more detail, including validation and error
45 checking, in sections 6.2.3.3. and 6.2.4, and in sections 8.9.3.1 and 8.10.5.

46 **6.2.3.2 Control Set Content**

47 Each Control Set contains a **Data Path Configuration** register array consisting of **DPConfigLane<i>**
48 registers (see Table 8-72, Table 8-77, and Table 8-93) defining two configuration settings for each host lane

- 49 • **Application Instance** related settings (described below)
- 50 • **Signal Integrity** (SI) related settings (described mainly in section 6.2.5)

The host defines an instance of an Application by assigning the relevant Application Descriptor "number" (AppSel code, see below) and its Data Path "location" (i.e. the Data Path lanes allocated to the Application instance, identified by the first host lane) to all lanes carrying that Application instance.

These Application and Data Path related settings must be **identical on all lanes** of a Data Path, while the Signal Integrity related settings may be chosen individually per lane.

For this purpose, each **DPConfigLane<i>** register includes an **AppSelCode** field and a **DataPathID** field (together defining the Application and the Data Path which lane <i> is part of), and a signal integrity related **ExplicitControl** bit (allowing host-defined SI settings to be used for lane<i>, instead of SI settings derived by the module for the Application).

AppSel Code Field (Application Descriptor Number)

As described in section 0, each supported Application is advertised by an Application Descriptor and each Application Descriptor is identified by a so called **AppSel** code (an Application Descriptor sequence number).

The host uses the AppSel code of the relevant Application Descriptor to assign the Application identified in that Application Descriptor to one or more host lanes, in the array of DPConfigLane<i> registers.

For a multi-lane Data Path, i.e. when an Application requires multiple lanes, the host writes the same AppSel code into the AppSelCode field of each lane of that Data Path.

The special AppSel code value **0000b** in the Data Path Configuration register of a host lane indicates that the lane (together with its associated resources) is **unused** and not part of a Data Path. The DataPathID and ExplicitControl fields of unused host lanes are irrelevant and may be ignored by the module.

Note: This zero valued AppSel code corresponds to a NIL or NULL pointer in programming languages.

In per-lane state reporting, the module always reports a DPDeactivated state for unused lanes (see Table 8-84).

DataPathID Field

The value of the DataPathID field in a DPConfigLane<i> register localizes and identifies a specific Data Path. The DataPathID of a Data Path is defined as the **lowest lane number** of all host lanes in that Data Path.

Note: See section 6.2.1.2 for rules to determine the media lanes of the Data Path.

Note: The host lane numbers of a Data Path are ultimately referenced to physical lane reference points. However, this does not limit the scope of internal module resources associated with the Data Path: The module resources associated with the Data Path may be located anywhere in the module and may be shared with other host lanes or associated with a media interface in the Memory Map. Any sharing of resources in the module is hidden in the sense that it does not affect independent operation of separate Data Paths in normal operation.

For a multi-lane Data Path, the host writes the same DataPathID in all lanes of that Data Path.

The host must assign lanes to Data Paths in accordance with the Lane Assignment Options field advertised by the module for that Application.

ExplicitControl Field

The **ExplicitControl** bit (EC) determines if the **signal integrity (SI) settings** for the functional resources associated with a lane are programmed by the host (EC=1) or by the module (EC=0).

The **effective** signal integrity settings resulting from host input (EC=1) or determined by the module (EC=0) are applied by the module during Data Path initialization.

With **host-controlled** settings (EC=1) the host programs a signal integrity configuration using the controls listed in Table 6-2 (which are described in section 6.2.5).

With **module-controlled** settings (EC=0) the module internally programs a signal integrity configuration that is determined to comply with the (targets of the) interface specifications of the relevant Application.

Note: In the EC=0 case, the signal integrity settings to be used may not be fully and uniquely derivable from the pertinent interface specification of an Application. In that case the module implements settings that accomplish the signal integrity targets of the Application in a best effort manner.

Table 6-2 lists all SI controls together with the values of the ExplicitControl bit required to enable the control. Otherwise the SI control field value is ignored, and the module instead uses SI settings that are compliant with the interface industry standard associated with the selected Application.

Table 6-2 Control fields activated by ExplicitControl bit

Signal Integrity Control Field (see e.g. section 8.9.3.3)	ExplicitControl (value to enable)
AdaptiveInputEqEnableTx	1
AdaptiveInputEqRecallTx	X
HostControlledInputEqTargetTx	1
CDREnableTx	1
CDREnableRx	1
OutputEqPreCursorTargetRx	1
OutputEqPostCursorTargetRx	1
OutputAmplitudeTargetRx	1

Note: The intent of the AdaptiveInputEqRecallTx feature is to quickly restore previously stored settings of adaptive equalization. This function is always available and may be used for that purpose also when EC=0.

The usage of these signal integrity control fields is defined in section 6.2.5.

6.2.3.3 Control Set Usage (Applying a Staged Control Set)

Initially, the module populates both Staged Control Set 0 and the Active Control Set registers with the module-defined default Application and signal integrity settings before exiting the MgmtInit state (see section 6.3.2.5.3).

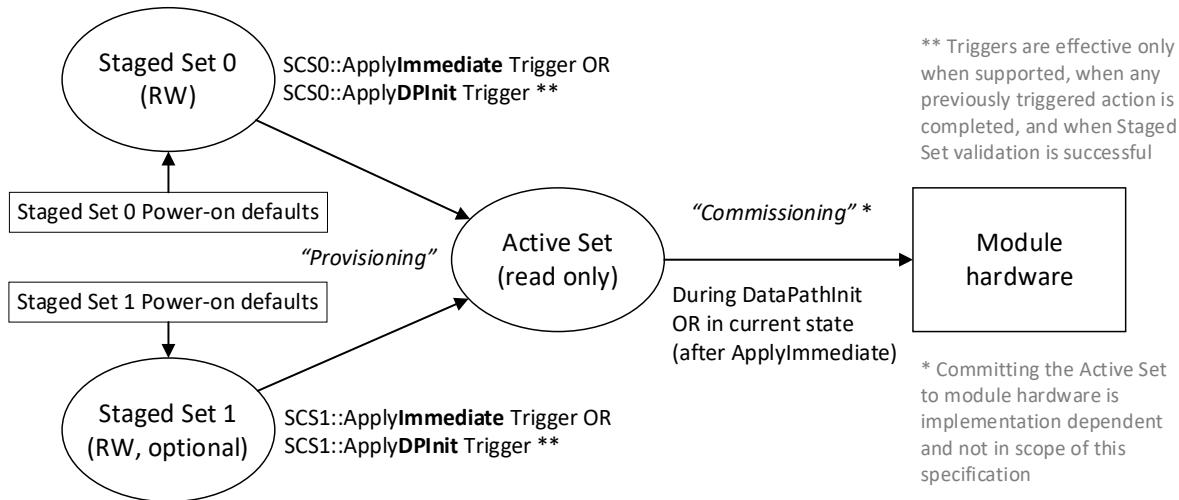
The host can request **provisioning** of one or more Data Paths by asking the module to **validate** and then **copy** the settings of the relevant lanes from the Staged Control Set into the Active Control Set, as illustrated in Figure 6-2.

The resulting lane configuration must consist only of completely valid Data Paths and unused lanes (if any).

Note: Recall that an AppSel value of 0000b in the Data Path Configuration register of a lane denotes an unused lane (not part of a Data Path) for which the module reports the state DPDeactivated (see section 6.3.3).

When the host requests **commissioning** of a Data Path provisioned into the Active Control Set, the module evaluates the settings in the Active Control Set to initialize all resources associated with that Data Path.

Note: Depending on past host provisioning activities, current settings in the Active Control Set may be module-initialized, host-provisioned, or host-commissioned values.

**Figure 6-2 Control Set Data Flow Diagram**

Apply Trigger Registers

For each Staged Control Set there are two **Apply** controls (trigger registers) available for the host to request copying **Data Path** settings from a Staged Control Set to the Active Control Set (using a bit mask to indicate the lanes of selected Data Paths): **SCS<k>::ApplyDPIInit** and **SCS<k>::ApplyImmediate**.

Note: Typical use cases for these Apply triggers are reconfiguration of the Application to be carried over a Data Path, modification of Data Path lane widths, modification of signal integrity settings on lanes of a Data Path, or changing interface speed by selecting another Application associated with that speed.

Note: ApplyImmediate may be unsupported, depending on module advertisement (see section 6.2.4).

The operations requested by Apply triggers are **Data Path operations**: A host must trigger **Apply** on **all** lanes of a Data Path simultaneously (even if only lane specific attributes of individual lanes are changed), with **one exception**: ApplyImmediate, if supported, can be used on a per lane basis, but only when only lane specific signal integrity settings are changed and only on lanes of Data Paths in state DPActivated or DPInitialized.

A host can trigger **Apply** on multiple Data Paths simultaneously in one WRITE access.

The effects of the host writing a lane selection bit mask to one of those trigger registers depend both on the trigger register instance used and on the states of the **Data Path State Machine (DPSM)** instances associated with the lanes selected in the bit mask.

Note: See section 6.3.3 for information on DPSM instances and states, and see section 6.2.4 and Table 6-3 for a command handling oriented description of the module procedures associated with Apply triggers.

Stepwise Provisioning and Commissioning

The host requests **provisioning** of prepared staged settings by invoking any Apply trigger on lanes indicating the **DPDeactivated** state, which causes the module to (1) **validate** and (2) **copy** the staged settings into the Active Control Set, without effect on module hardware.

The subsequent **commissioning** of the Active Control Set occurs separately, after a **host-initiated transition to the DPInit state**, where the module will eventually (3) **initialize** the associated Data Path resources based on the configuration then found in the Active Control Set, regardless of which sequence or type of Apply trigger commands had been used to define the current content of the Active Control Set.

Note: For historical reasons, the following intervention-free reconfiguration procedures are supported by the default advertisement, but a module may advertise that only stepwise configuration is supported.

Combined Provisioning and Commissioning (intervention-free)

When invoked on lanes indicating an **initialized** Data Path that is currently in use, i.e. on lanes indicating the **DPInitialized** or **DPActivated** state, both Apply triggers cause the module to perform the provisioning and commissioning steps sequentially, albeit in two different ways, with or without DPSM state changes, described as **Regular Reconfiguration** or **Hot Reconfiguration** as follows:

Regular Reconfiguration (Provisioning Procedure with Commissioning in DPSM State Cycle)

After an **ApplyDPInit** trigger in an **initialized** state the module will first (1) **validate** and (2) **copy** the staged settings into the Active Control Set, then (3) **deinitialize** the Data Path by **auto-transitioning** the DPSM to the **DPDeactivated** state, and finally (4) **initialize** the Data Path to **realize** the new Active Control Set content.

This **reinitialization** (3+4) occurs **automatically** as a sequence of **DPSM state transitions** (by internal module procedures executed in the states), eventually moving back to the state where the **ApplyDPInit** was triggered (unless DPSM state transition relevant configurations have meanwhile changed).

Note: An intervention-free procedure executed internally on host command does not require READ or WRITE access for each step and hence can execute faster, as is required in Applications like InfiniBand or Fibre Channel.

Note: This intervention-free reconfiguration function is provided for cases where new settings need to be applied faster than possible with the stepwise management interactions that are necessary when the host fully controls the reconfiguration or reinitialization loop.

Hot Reconfiguration (Provisioning and Commissioning Procedure without DPSM State Change)

After an **ApplyImmediate** trigger in an **initialized** state the module performs the provisioning and commissioning actions, i.e. (1) **validate**, (2) **copy**, and (3) **re-initialize**, in one optimized internal procedure, **without changing** the DPSM state.

Note: This "hot reconfiguration" function is provided for cases where new settings need to be applied quickly without disabling the media lane transmitters, such as in Fibre Channel Link Speed Negotiation (LSN).

*Note: Using **ApplyImmediate** in the **DPInitialized** or **DPActivated** state allows still faster completion of a desired configuration change in cases where this is essential. The intent of this specification is for the module to perform the actions associated with **ApplyImmediate** as quickly and efficiently as possible in a minimally disruptive manner. However, in some cases such as changing speeds, disruption is unavoidable.*

Silent Rejection

When invoked on lanes indicating a transient state (**DPInit**, **DPDeinit**, **DPTxTurnOn**, **DPTxTurnOff**) the module may ignore Apply triggers; the result is then undetermined in general.

1 Note: The detailed behavior depends on the support of intervention-free reconfiguration, as described below.

2 Note: Modules should preferably provide rejection feedback (ConfigRejected), when possible, see below.

3 Error Handling and Result Feedback

4 In all cases, if a requested Data Path configuration is not accepted during validation, the module skips the
5 remaining steps of the procedure without copying any settings for that Data Path into the Active Control Set.

6 When terminating the procedure the module always reports a result status code (see Table 8-91) informing
7 about successful execution (ConfigSuccess) or about rejection (ConfigRejected*) for each affected lane, using
8 the Configuration Command Execution and Result Status fields (**ConfigStatus**, see Table 8-90).

9 Realization (Commissioning)

10 The translation of the Memory Map settings in the Active Control Set to module hardware is implementation-
11 specific and outside the scope of this specification. Modules should implement a best effort approach when
12 trying to translate Memory Map settings to module hardware settings.

13 As a visual summary, Figure 6-2 illustrates the data flow of configuration settings from a Staged Control Set
14 into the Active Control Set and further into module hardware.

15 6.2.4 Data Path Configuration – Procedures and Commands

16 This section describes the process of changing the Data Path configuration of a module at the level of procedures
17 and commands.

18 6.2.4.1 Configuration Concept

19 Bringing a Data Path into service, such that it provides the Application-defined transmission function, requires
20 three steps, herein called **definition**, **provisioning**, and **commissioning**.

21 *Note: It is customary to call the combination of all three steps a configuration activity, but sometimes
22 configuration refers only to the first two steps. For accuracy one would have to distinguish these cases as actual
23 configuration and as nominal configuration, respectively, but we hope that the meaning is clear from context.*

24 *Note: Recall that configuration changes are requested at Data Path granularity even if only settings affecting
25 individual lanes are changed (e.g. lane-specific SI settings).*

26 Configuration Steps

27 The first step, **definition**, is a host-executed procedure defining a desired Data Path configuration in a selected
28 Staged Control Set, using an arbitrary sequence of register WRITE accesses, without module reaction. Definition
29 is a pure **host procedure**.

30 The second step, **provisioning**, is a host-triggered module procedure that **validates** the settings of a Data
31 Path defined in a given Staged Control Set and, if successful, **copies** them into the Active Control Set, for all
32 lanes of a Data Path. This step is a **module procedure**, triggered by the host.

33 Commissioning

34 The third step, **commissioning**, is either a host-initiated procedure or an automatic follow-up of a previously
35 host-triggered provisioning procedure: in both cases the module commits the nominal configuration of selected
36 lanes in the Active Control Set to hardware such that the Active Control Set eventually matches the actual
37 hardware configuration again. This step is a **module procedure** that is either directly **initiated** or indirectly
38 **triggered** by the host.

39 *Note: **Triggering** and **initiating** are deliberately distinguished here: triggering has an event nature and is
40 implemented as a host writing to a stateless trigger register while initiating is a causal consequence of changing
41 a setting in a configuration register. For triggered procedures, the module provides feedback, while initiated
42 procedures provide feedback only indirectly, by their observable effects.*

43 The host selects all host lanes of a Data Path by writing to lane-associated registers or fields (in definition or
44 for initiating a module procedure), or by using a lane selection bit mask (when triggering a module procedure).

45 Host Procedures

46 This specification describes two types of configuration procedures.

47 In a **stepwise** procedure, the host performs all configuration actions step-by-step, according to the **Data Path**
48 **State Machine (DPSM)** initialization and configuration model described in section 6.3.3.

Characteristic features of this procedure are that the host provisions only lanes that are currently either **unused** or part of a deactivated Data Path (lanes in **DPA deactivated** state), and that the commissioning procedure is first explicitly initiated by the host and then performed by the module while initializing a Data Path (in **DPI init** state).

Note: For Applications where the stepwise procedure is fast enough, for example Ethernet or Transport, it is recommended that the host uses only a stepwise procedure and avoids the intervention-free procedures described in the following.

In **intervention-free** (or **automatic**) procedures, the host triggers only the provisioning step on initialized lanes that are currently in use (in **DPI initialized** or **DPA activated** state) and the module subsequently performs the commissioning step automatically, without further host intervention.

However, a module may advertise that it does not support, or only partially support, intervention-free reconfigurations (see fields SteppedConfigOnly and AutoConfigSupport in Table 8-5).

In the intervention-free **regular reconfiguration** procedure (triggered via an **ApplyDPI init** register) the module automatically **cycles through the states** of a Data Path (via DPA deactivated and DPI init) to perform the commissioning step when the triggered provisioning step was successfully executed.

Note: This method is targeted for Applications where the stepwise procedure is not fast enough, for example Fibre Channel and InfiniBand.

In the intervention-free **hot reconfiguration** procedure (triggered via an **ApplyImmediate** register), the module performs all configuration and commissioning actions automatically while **staying in the current** Data Path **state**.

Note: This method is targeted for Applications where neither the stepwise nor the intervention-free regular reconfiguration procedure is fast enough, for example Fibre Channel.

Note: The option SteppedConfigOnly has been introduced in CMIS 5.0 to simplify implementation and test, for modules supporting only Applications that can always be reconfigured step-by-step. Later, advertisements have been added to support of only one of the two intervention-free reconfiguration procedures. See Table 8-5.

Module Procedures

Unlike configuration definition, the provisioning and commissioning steps are **module procedures** executed by the module after a host action (see the following section 6.2.3.3). See also Table 6-3 below.

Module procedures may take time to complete and may fail. Hence the module provides **feedback** to the host, both about an execution-in-progress status and, eventually, about success or failure of the triggered procedure.

The host **triggers** the beginning of a module procedure by writing a lane selection bit mask to an **Apply** register in the relevant Staged Control Set (see Table 8-70 and Table 8-76), and the module provides feedback about the execution and about the result status of the procedure via per-lane **ConfigStatus** registers (see Table 8-90 and Table 8-91).

Module Procedure Execution Protocol (Command Handling)

The execution of a host triggered Data Path module procedure consists of four steps, which may be understood, conceptually, as a **command handling** sequence.

First, in **command acceptance**, the module determines if the requested procedure can currently be accepted at all: here, a new trigger cannot be accepted when a previously triggered configuration procedure is still executing. When accepted, the module immediately sets the execution status for this lane (**ConfigInProgress**); else the command is not executed and silently **ignored, without feedback** to the host.

*Note: It is therefore host responsibility to check the **ConfigStatus** of the lanes of a Data Path in advance, or to otherwise ensure that no procedure is still executing, prior to triggering a configuration procedure.*

Second, in the **command validation** step, the module validates the parameters of the command (content of Staged Control Set) and skips the next step on those Data Paths for which parameter validation has failed.

Third, in the actual **command execution** step, the desired procedure is executed.

Last, the module reports the **command result** status in the **ConfigStatus** registers of the lanes of the Data Path. The result status reports either successful execution (**ConfigSuccess**) or rejection without execution due to validation failure prior to command execution (**ConfigRejected***).

Note: Unexpected failures during command execution are not reported.

The host can always query the **ConfigStatus** of a lane to see if a configuration command is currently still being processed or otherwise what the result of the previously executed command was (see Table 8-91).

6.2.4.2 Configuration Commands

The overall response of the module to a host WRITE of a lane selection bit mask to a trigger register is best understood as per Data Path **command handling**, which consists of (1) command handling readiness check, (2) input parameter validation, (3) actual command execution, and (4) positive or negative result status feedback. This is now described in more detail.

Note: Recall that configuration commands are applied to entire Data Paths (with one exception), while the trigger and result reporting mechanism is defined per lane (all lanes reporting the same value). Also note that the desired reaction to a trigger possibly depends on several factors: DPSM state, Apply trigger register type and instance used, nature of changes, and on restriction advertisement.

Configuration Command Handling

Writing a lane selection bit mask to an **Apply** register represents a **(re-)configuration command** and **triggers** execution of one or more parallel **module procedures**, depending both on the trigger register used and on the current DPSM state(s) associated with each affected lane, as indicated in Table 6-3 and Table 6-4.

Writing a set lane selection bit to an Apply register is only allowed when the selected lane indicates a stable and steady DPSM state and when no previously triggered configuration command is still being executed on the selected lane.

Note: This precondition can be checked by the host by reading the Configuration Command Execution and Result Status registers described below.

Note: Writing to an apply register in transient states or in transitory steady states may have unpredictable effects.

When a previously triggered **Provision** or **Provision-and-Commission** command is still being processed for lanes of a Data Path, the module **ignores** new triggers for those lanes, and does not execute the command on the Data Paths those lanes belong to.

*Note: Ignoring lane selection or trigger bits due to configuration being in progress is **not** indicated to the host.*

Table 6-3 Configuration Commands (Intervention-Free Reconfiguration Procedures Supported)

Control (Trigger Reg.)	Data Path State *	Host Intent (configuration command type)	Triggered Module Procedure ** (when trigger is accepted)
ApplyDPInit	DPDeactivated	Stepwise (host-controlled) Regular *** Configuration ("copy only") (provisioning)	Provision 0 ConfigStatus = ConfigInProgress 1 Validate content of Staged Control Set 2 Copy content to Active Control Set 3 Set DPInitPending **** 4 Report result status in ConfigStatus <i>Note: commissioning is done via DPSM</i> <i>Note: DPInitPending is cleared in DPInit</i>
	DPActivated or DPInitialized	Intervention-free (host-trigg'd) Regular *** Reconfiguration ("copy and cycle") (provisioning + commissioning)	
ApplyImmediate		Intervention-free (host-trigg'd) Hot Reconfiguration ("copy and commit") (provisioning + commissioning)	Provision-and-Commission 0 ConfigStatus = ConfigInProgress 1 Validate content of Staged Control Set 2 Copy content to Active Control Set 3 Commit Active Control Set to HW 4 Report result status in ConfigStatus

* Triggers are ignored in Data Path states not explicitly mentioned

** Steps 2 and 3 are conditional on successful validation in step 1. Note that the module procedure completes only the provisioning step when the host intention is provisioning and commissioning, while the commissioning step is handled by triggering or unblocking DPSM behavior.

*** Reconfiguration represented by DPSM state transitions is called "regular reconfiguration", as opposed to "hot reconfiguration" which is performed without DPSM state transitions.

**** These bits cause the Data Path State Machine to automatically execute the commissioning

Table 6-4 Configuration Commands (Intervention-Free Reconfigurations Not Supported)

Control (Trigger Reg.)	Data Path State *	Host Intent	Triggered Module Procedure **
ApplyDPInit	Any DPSM state	Provisioning (part of stepwise procedures)	Provision 0 ConfigStatus = ConfigInProgress 1 Validate content of Staged Control Set 2 Copy content to Active Control Set 3 Set DPInitPending *** 4 Report result status in ConfigStatus <i>Note: DPInitPending is cleared in DPInit</i>
ApplyImmediate	Any DPSM state	None	None (no reaction)

* Triggering recommended only in DPDeactivated (avoid difference of Hardware and Active Control Set)

** Steps 2 and 3 are conditional on successful validation in step 1

*** The DPInitPending bits are only indicators in this case, because the Data Path State Machine reacts only when Intervention-Free Regular Reconfiguration is supported

Configuration Command Execution Status and Result Status

The module reports both the transitory execution-in-progress status and the final result status of the configuration command using the Configuration Command Execution and Result Status (**ConfigStatus**) fields described in Table 8-90. The status reporting for Data Path configurations is per lane (as always in CMIS).

When accepting an Apply trigger, the module **immediately** updates the ConfigStatus field of all affected lanes to indicate **ConfigInProgress**.

When a Module Procedure triggered on a Data Path has been completed, the module updates the ConfigStatus field of all triggered lanes of that Data Path with the appropriate result: **ConfigSuccess** or **ConfigRejected***. Only one Data Path result status is reported on all triggered lanes of a Data Path.

6.2.4.3 Host Rules and Recommendations

This subsection provides a host-oriented compilation of important rules and recommendations that follow from module-oriented specifications elsewhere, without a claim of completeness:

- The host must assign a valid **AppSel** code to all host lanes in the Active Control Set
- All lanes of the application identified by an AppSel code must be assigned the same AppSel code
- The host must assign **AppSel** = 0000b to each unused host lane (which is not part of a Data Path)
- When a lane is used in a Data Path, the host can reconfigure it to become unused only when the Data Path is in the DPDeactivated state, i.e. the host must not invoke an Apply trigger on lanes marked as unused in the Staged Control unless the Data Path currently associated with those lanes is in the **DPDeactivated** state
- Each used lane in the Active Control Set must be part of a completely valid Application instance (i.e. a valid Application completely allocated on lanes supported for that Application).
- The host can change the width of a Data Path only while in the **DPDeactivated** state, i.e. the host must always transition an existing Data Path to **DPDeactivated** before selecting an Application with a different lane count. Any lane that becomes unused must be marked as such (AppSel = 0000b) or it must be assigned to a new valid Data Path (remaining in DPDeactivated state until eventually used)
- It is recommended that hosts **use only** the fully host-controlled **stepwise reconfiguration** procedure and use ApplyDPInit only for Data Paths in DPDeactivated state – this is the cleanest way to perform a Data Path configuration change, more likely to behave exactly like initial power up configuration
- It is recommended that hosts use the **intervention-free** reconfiguration procedures using ApplyDPInit or ApplyImmediate on Data Paths in steady states DPInitialized or DPActivated **only when required** to meet protocol performance requirements of the Application
- When **intervention-free** reconfiguration procedures are supported, hosts are advised not to invoke an Apply trigger on the lanes of a Data Path in a transient state (**DPInit**, **DPDeinit**, **DPTxTurnOn**, or **DPTxTurnOff**), as the module silently ignores requests received while still being in a transient state.
- Hosts are advised to use ApplyImmediate only when intervention-free reconfiguration is supported, and only in initialized states (DPInitialized and DPActivated), because the module silently ignores ApplyImmediate in all other cases.

- 1 • Before involving any Apply trigger, hosts are advised to ensure that the **ConfigStatus** of all affected
2 lanes is different from **ConfigInProgress**, as the module may otherwise silently ignore the trigger

3 **6.2.4.4 Initialization Sequence Examples**

4 The Data Path architecture described above is intentionally designed to support a broad array of
5 implementations while ensuring compatibility across hosts and modules. Some Applications may not use all of
6 the features provided in the architecture.

7 Appendix D contains some example host-module initialization flows that can be used for popular Applications.

6.2.5 Signal Integrity Related Controls

An explicit control option and signal integrity control fields in the Staged Control Sets provide a per lane mechanism for the host to **override** signal integrity (**SI**) settings that are otherwise determined by the module.¹

If the **ExplicitControl** bit is set in the applicable Staged Control Set (see Table 8-72 and Table 8-77), the module copies all **host provisioned** signal integrity control values from the Staged Control Set to the Active Control Set (subsequently committed to module hardware, see Figure 6-2) when an **Apply** trigger is used.

*Note: In this case the host is responsible to provide meaningful values for **all** signal integrity controls.*

If the **ExplicitControl** bit is cleared, the module ignores the Staged Control Set and writes suitable Application-dependent default signal integrity control values into the Active Control Set when an **Apply** trigger is used.

Note: In this case the module decides, in a best effort manner, which signal integrity control settings accomplish the targets specified in the pertinent interface standard of an Application.

Note: Adaptive or Non-Adaptive equalization control is only available to the host when the ExplicitControl bit is set. While the AdaptiveInputEqRecallTx recall works also when the ExplicitControl bit is not set.

For all signal integrity controls, the control values represent desired behavior in terms of reference receiver or transmitter metrics and the module makes a best effort to achieve the desired effect in its implementation.

6.2.5.1 Tx Input Equalization Control

The controls for Tx input equalization can be grouped by equalization type, as shown in Table 6-5.

Table 6-5 Tx Input Eq control relationship to AdaptiveInputEqEnableTx

Equalization Type	Control	AdaptiveInputEqEnableTx
Adaptive	AdaptiveInputEqFreezeTx	1
	AdaptiveInputEqStoreTx	
	AdaptiveInputEqRecallTx	
Non-Adaptive	HostControlledInputEqTargetTx	0

The controls relevant for **adaptive** Tx input equalization are described in section 6.2.5.4.

The controls relevant for **non-adaptive** Tx input equalization, when Tx input equalization settings are either **pre-determined** by the module or **host-controlled**, are described below.

Note: The meaning of these input equalization values may be specified elsewhere, outside of this specification.

The module **ignores** control field values that are not relevant for the current AdaptiveInputEqEnableTx setting.

Host-Controlled Equalization

Tx input equalization values in dB are based on a reference CTLE and may not directly apply to the equalizer implemented in the module.

SCS<k>::**HostControlledInputEqTargetTx**<i> is a four-bit control field for lane <i> and encoded as shown in Table 6-6. This field allows the host to program a Tx input equalization target, and it is ignored by the module if AdaptiveInputEqEnableTx<i> is set for that lane.

The module advertises support of host-controlled Tx input equalization as described in Table 8-53.

The module advertises the maximum supported Tx input equalization values as described in Table 8-49.

Table 6-6 Host-Controlled Tx Input Equalization Codes

Value	Bit pattern	Input Equalization
0	0000b	No Equalization
1 – 12	0001b ... 1100b	1 dB ... 12 dB
13-15	1101b ... 1111b	Custom

¹ Here, CMIS defines a fixed structure for the SI parameters in a control set. See CMIS supplement [6] for the description of a versatile and module-defined configuration set.

6.2.5.2 Rx Output Equalization Control

Rx output equalization is defined at an appropriate test point defined by the relevant standard.¹

`SCS<k>::OutputEqPreCursorTargetRx<i>` and `SCS<k>::OutputEqPostCursorTargetRx<i>` are four-bit control fields for lane <i> defined in Table 8-74 and encoded as shown in Table 6-7.

The module advertises support of Rx output equalization control as described in Table 8-53.

The module advertises the maximum supported Rx output equalization values as described in Table 8-49.

Modules may advertise (in RxOutputEqControlSupported) that they support only a single output emphasis setting (Pre-Cursor only or, more likely, Post-Cursor only). In both these cases and despite the apparent name mismatch in the Pre-Cursor only case, the `SCS<k>::OutputEqPostCursorTargetRx<i>` fields are used to control the single emphasis settings (as they offer a larger dB range), and the module ignores the `SCS<k>::OutputEqPreCursorTargetRx<i>` fields.

Table 6-7 Rx Output Equalization Codes

Code Value	Bit pattern	Post-Cursor Equalization	Pre-Cursor Equalization
0	0000b	0dB (No Equalization)	0dB (No Equalization)
1	0001b	1 dB	0.5 dB
2	0010b	2 dB	1.0 dB
3	0011b	3 dB	1.5 dB
4	0100b	4 dB	2.0 dB
5	0101b	5 dB	2.5 dB
6	0110b	6 dB	3.0 dB
7	0111b	7 dB	3.5 dB
8-10	1000b-1010b	Reserved	Reserved
11-15	1011b-1111b	Custom	Custom

Note: The pre-cursor equalizer settings in dB approximates to

$$\text{Pre EQ (dB)} = -20 \cdot \log_{10}(1 - C_1 / (C_1 + C_0 + C_1)) \quad (\text{Eq. 6-1})$$

The post-cursor equalizer settings in dB approximates to

$$\text{Post EQ (dB)} = -20 \cdot \log_{10}(1 - C_1 / (C_1 + C_0 + C_1)) \quad (\text{Eq. 6-2})$$

Equalizer coefficients C_n are pre-cursor for $n < 0$ and post-cursor when $n > 0$.

6.2.5.3 Rx Output Amplitude Control

The Rx output amplitude is measured without Rx output equalization and defined at an appropriate test point defined by the relevant standard.

`SCS<k>::OutputAmplitudeTargetRx<i>` is a four-bit field for lane <i> to specify the Rx output signal level being in a particular amplitude range, with amplitude range encoding as described in Table 6-8.

The module advertises support of Rx output amplitude control as described in Table 8-53.

The module advertises the maximum supported Rx output amplitudes as described in Table 8-49.

Table 6-8 Rx Output Amplitude Codes

Value	Bit pattern	Output Amplitude
0	0000b	100-400 mV (P-P)
1	0001b	300-600 mV (P-P)
2	0010b	400-800 mV (P-P)
3	0011b	600-1200 mV (P-P)
4-14	0100b-1110b	Reserved
15	1111b	Custom

¹ Note, however, that the equations below Table 6-7 approximate pre-cursor and post-cursor equalization in terms of the normalized tap weights of the FFE (in the device), C_0 , C_1 , and C_2 . These approximated equalization values are not measurable; measurements at external test points will be lower due to module internal losses and board losses.

6.2.5.4 Adaptive Tx Input Equalizer Store and Recall Mechanism

Note: Equalizer adaptation can be time-consuming. In some applications, the available time for a speed change (which incurs selecting a new Application and hence a Data Path reconfiguration) does not include equalizer adaptation time. To better support such applications, an optional host-controlled Store and Recall mechanism is specified for storing adapted equalizer settings for later recall and use.

Note: Only modules supporting applications that require fast application change (e.g. for speed negotiation) with critical time budgets are expected to support this mechanism.

The **TxInputEqRecallBuffersSupported** field (see Table 8-53) advertises support of the adaptive Tx input equalizer store and recall mechanism by advertising the number of recall buffers supported by the module.

Recall buffers are numbered and used independently of Staged Control Set instances.

The module provides enough storage in each recall buffer to store adapted equalizer settings for each lane of the module. The storage mechanism is implementation specific and not defined in this specification.

The **AdaptiveInputEqStoreTx** trigger field is described in Table 8-69. To store the most recent adapted Tx input equalizer settings for lane <i>, the host writes the desired target recall buffer number into **AdaptiveInputEqStoreTx<i>**. Equalizer adaptation is then stopped until the settings have been stored and continues afterwards, unless adaptation is frozen (i.e. unless **AdaptiveInputEqFreezeTx<i>** is set).

Host requests to store equalizer settings while adaptation is disabled (i.e. **AdaptiveInputEqEnableTx<i>** is cleared for lane <i>) are ignored by the module.

The host may trigger storage by writing to **AdaptiveInputEqStoreTx<i>** at any time while the Data Path state is DPInitialized or DPActivated, and a requested storage occurs then immediately.

The **AdaptiveInputEqRecallTx** control field is described in Table 8-73 for Staged Control Set 0 and in Table 8-78 for Staged Control Set 1.

The Active Control Set provides a read-only indication of the current **AdaptiveInputEqRecalledTx** status for each lane (see Table 8-94).

To recall a stored Tx input equalizer adaptation setting into a Staged Control Set, the host writes the storage buffer number to be recalled to the applicable **AdaptiveInputEqRecallTx** lane controls.

Note: These settings are not recalled into the active equalizer until the host has triggered ApplyDPInit or ApplyImmediate for that Staged Control Set.

The **AdaptiveInputEqRecallTx** field is used independent of the **ExplicitControl** field settings for that lane.

If **AdaptiveInputEqFreezeTx** is cleared, the recalled Tx input Eq adaptation setting is used as the starting point for continuous adaptation for the applicable lanes. Otherwise, if **AdaptiveInputEqFreezeTx** is set, the recalled Tx Input Eq Adaptation is used as the frozen Tx Input Eq value for the applicable lanes.

6.3 Module Behavioral Model

6.3.1 State Machine Concept

Fundamental power up, initialization, and reinitialization interactions between host and module are governed and described by state machine based behavioral models. Conceptually, these state machines are considered parts of the module.

Note: The state machines that are actually supported by the module depend on the type of module (such as transmission module, resource module, etc.). This support is advertised in the CmisSmSupported register (see section 8.2.12).

Note: These state machines are purely conceptual, precise models to specify required host and module interactions and behaviors associated with those interactions; they do not constrain software implementation.

State machines describe both autonomous behavior (i.e. what the module does) and reactive behavior (i.e. how the module reacts to events caused by the host). In certain situations, the states of the state machines in the module cannot be observed by the host, but they still govern behavior and reactions of the module in these situations.

Module and Data Path State Machines

Two types of state machine are distinguished: Module State Machine and Data Path State Machine.

A **Module State Machine (MSM)** defines host-module interactions and behavioral characteristics of the module as a whole, such as the initialization of the management interface and the Module Power Mode.

A **Data Path State Machine (DPSM)** defines host-module interactions and behavioral characteristics needed for the initialization of one particular **Data Path**, which represents signal flow and signal processing of one instance of one type of **Application**.

Note: For information on Applications and Data Paths see sections 6.2.1 and 6.2.2, respectively. For more information on Module State Machine and Data Path State Machines see sections 6.3.2. and 6.3.3.

A module has a single Module State Machine, but the number of Data Path State Machines depends on module configuration: A simple module with a single Data Path for a single Application, has one DPSM. A complex module supporting parallel Data Paths for multiple instances of one Application or instances of multiple Applications has several DPSM instances, one for each provisioned Application instance.

Note: As defined elsewhere in this specification, parallel Data Path State Machines are controlled independently, and they operate concurrently and mutually independent. Parallel Data Paths behave largely like independent virtual modules within a physical module.

Transient and Steady States

Two kinds of states are distinguished (in both types of state machines), steady states and transient states.

In **steady states** the module is essentially waiting for the host to initiate action.

Note: some exceptional module internal events may also cause exit from steady states.

In **transient states** the module performs dedicated activities and automatically performs a state transition upon completion (unless events cause other state transitions before completion).

In the state machine illustrations, such as Figure 6-3, Figure 6-4, and Figure 6-5, steady states are displayed with a rectangular outline and transient states are displayed with an oval outline.

The durations of steady states are unbounded, from a module's viewpoint, as the exit conditions depend on events not controlled by the module.

The durations of transient states are bounded because they are associated with module-controlled activities. The duration bounds may be implementation dependent. For most transient states a maximum duration is either advertised or bounded by a specified limit (see chapter 10).

Form Factor Independent Specification

CMIS is designed to be applied to multiple form factors. Therefore, actions and properties are described using generic signal names instead of using form factor specific signal names. Appendix A describes the association of generic and form factor specific signal names, and the implementation of abstract signal levels.

6.3.2 Module State Machine (MSM)

The **Module State Machine (MSM)** defines host-module interactions and behavioral characteristics of the module as a whole. The MSM state represents a module's readiness to be managed and used by its host.

The Module State Machine describes module-wide behaviors and characteristics.

Note: For Data Path-specific behaviors and characteristics, refer to the Data Path State Machine in section 6.3.3.

Host View

By observing the MSM state, the host can determine when the management interface has completed initialization after power on or on exit from reset. The host can also determine from the Module State Machine state when the high-speed circuitry in **paged memory modules** may be fully powered such that the host can begin to initialize Data Paths through the Data Path State Machine (described in section 6.3.3).

Module View

The MSM is engaged immediately after module insertion and power on.

The MSM is applicable both to modules and to cable assemblies, whether passive or active.

- The MSM for paged memory modules is described in section 6.3.2.2.
- The simplified MSM for flat memory modules is described in section 6.3.2.3.

6.3.2.1 General Module Behavior

This section describes general module behavior and requirements associated with the Module State Machine.

Management Data Relevance, Validity and Accuracy

Prior to proper initialization, i.e. before the state of basic manageability (ModuleLowPwr) has been reached, the contents of the management Memory Map are neither evaluated nor maintained by the module.

During transient states, the contents of reporting registers that are usually dynamically changed by the module may generally be invalid or inaccurate until a fully managed operational state is reached again.

Note: Dynamic reporting registers are read-only updated by the module to report status or performance data.

Flag Setting and Clearing

The set of Flags (possibly causing Interrupts) that a module may set is defined specifically for each state. The state dependent Flag setting rules are defined in section 6.3.4.

Unless specified otherwise, the module shall not clear any Flags on a state transition, and Interrupts are only cleared when the host has read or masked all set Flags.

State Change Indication

The **ModuleStateChangedFlag** is set after entering a new state, but only for the subset of module-initiated state transitions that are defined in Table 6-9.

In addition, setting the **ModuleStateChangedFlag** is generally suppressed when the new state is exited immediately because its exit conditions are already fulfilled on entry.

Note: The rationale for this suppression rule is that intermediate state change notifications are not interesting and should not cause Interrupt; only the last transition into a new stable state should be notified.

Table 6-9 ModuleStateChangedFlag behaviors

New State	ModuleStateChangedFlag flagged on state entry? *
Resetting	No
Reset	No
MgmtInit	No
ModuleFault	Yes
ModuleLowPwr	Yes
ModulePwrUp	No
ModuleReady	Yes
ModulePwrDn	No

*unless the exit condition was already satisfied on entry into the entered state

6.3.2.2 Module State Machine for Paged Memory Modules

The behavior of Paged Memory Modules is described by the Module State Machine shown in Figure 6-3

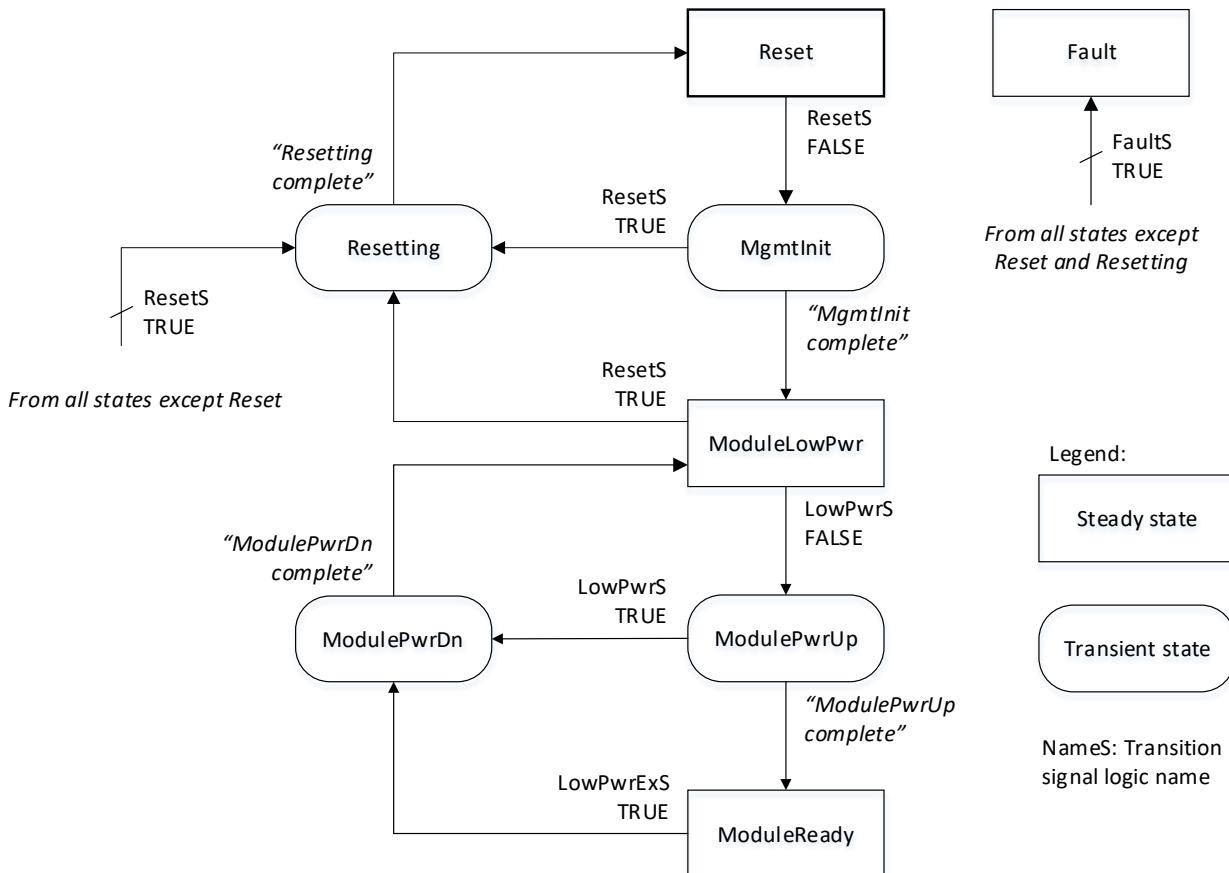


Figure 6-3 Paged Memory Module State Machine (MSM) State Transition Diagram

On module power-up, the initial Module State Machine state is the Reset state, and the State Machine remains in this state while ResetS is TRUE. Otherwise, the Module State Machine transitions to the MgmtInit state.

The state machine exits a given state when specific exit conditions are satisfied. So called **transition signals** (names ending in **S**) are used to describe the logic conditions governing a state transition. Other compound logical terms may be symbolically named with names ending in **T**.

The following table describes the priority of exit conditions, if more than one exit condition is satisfied at the same time. Note that not all exit conditions are applicable to all states.

Table 6-10 Module State Machine exit condition priority

Priority	Exit Condition
1	ResetS
2	FaultS
3	All other exit conditions

Signals and Conditions

The **ResetS** transition signal is described using the truth table shown in Table 6-11, below.

Table 6-11 ResetS transition signal truth table

VccReset (due to low Vcc)	Reset (hardware signal)	SoftwareReset (see Table 8-11)	ResetS (transition signal)
ASSERTED	X	X	1
DEASSERTED	ASSERTED	X	1
DEASSERTED	DEASSERTED	1	1
DEASSERTED	DEASSERTED	0	0

The ResetS transition signal can also be represented by the logic equation

$$\text{ResetS} = \text{ASSERTED(VccReset)} \text{ OR ASSERTED(Reset)} \text{ OR SoftwareReset} \quad (\text{Eq. 6-3})$$

The internal **VccReset** generic hardware signal is asserted within the module when the voltage of one or more of the Vcc power rails as observed at the module input drops below an implementation-defined value. Implementation of VccReset is optional.

The **Reset** generic hardware signal must be asserted by the host for longer than the minimum reset pulse duration to trigger a module reset. Refer to form factor-specific documentation for the minimum reset pulse duration. See Appendix A for a mapping of the generic signal name to form factor specific signals.

The **FaultS** transition signal truth table and logic equation are module implementation specific.

The **LowPwrS** transition signal is described by the truth table shown in Table 6-12, below.

Table 6-12 LowPwrS transition signal truth table

LowPwrRequestSW see Table 8-11	LowPwrAllowRequestHW see Table 8-11	LowPwrRequestHW hardware signal	LowPwrS transition signal
1	X	X	1
0	1	ASSERTED	1
0	1	DEASSERTED	0
0	0	X	0

The **LowPwrS** transition signal can also be represented by the logic equation

$$\text{LowPwrS} = \text{LowPwrRequestSW OR (LowPwrAllowRequestHW AND ASSERTED(LowPwrRequestHW))}$$

Note: The power up default values LowPwrRequestSW=0 and LowPwrAllowRequestHW=1 (see Table 8-1) imply that a module powers up under hardware control, i.e. by the LowPwrRequestHW control input signal. To pause module startup in the ModuleLowPwr state, LowPwrRequestHW needs to be asserted prior to startup.

The **LowPwrExS** transition signal is described by the truth table shown in Table 6-13, below. It represents the exit conditions from the ModuleReady state, when the LowPwrS transition signal is TRUE AND all Data Paths have reached the DPDeactivated state (represented by the logical term **ModuleDeactivatedT**).

Table 6-13 LowPwrExS transition signal truth table

LowPwrS transition signal	ModuleDeactivatedT	LowPwrExS transition signal
1	1	1
1	0	0
0	1	0
0	0	0

The **LowPwrExS** transition signal can also be represented by the logic equation

$$\text{LowPwrExS} = \text{LowPwrS AND ModuleDeactivatedT} \quad (\text{Eq. 6-4})$$

where

$$\begin{aligned} \text{ModuleDeactivatedT} = & (\text{DPStateHostLane1} = \text{DPDeactivated}) \text{ AND} \\ & (\text{DPStateHostLane2} = \text{DPDeactivated}) \text{ AND ...} \\ & \dots \text{ AND ...} \\ & (\text{DPStateHostLaneN} = \text{DPDeactivated}) \end{aligned} \quad (\text{Eq. 6-5})$$

N = number of host lanes in the module

State Transition Table

Table 6-14 provides a summary of the high-level behaviors and properties of each module state for paged memory module implementations. Refer to sections 6.3.2.5.1-6.3.2.5.8 for detailed requirements for each state.

Table 6-14 Module state behaviors, paged memory modules

State	Power Mode	Behavior in state	Exit condition	Next state	Req./Opt.
Resetting	High/ Low Power	Management interface and all module electronics transition to reset	Resetting completed	Reset	Rqd.
Reset	Low Power	Management interface and all module electronics in reset	ResetS transition signal is FALSE	MgmtInit	Rqd.
MgmtInit	Low Power	Management interface powering up and initializing	ResetS transition signal is TRUE	Resetting	Rqd.
			FaultS transition signal is TRUE	ModuleFault	
			Module management interface ready AND Interrupt by ModuleStateChangedFlag set. Note: must occur within a duration of tMgmtInit (see section 10.1)	ModuleLowPwr	
ModuleLowPwr	Low Power	Management interface available, entire paged management memory accessible for reading and writing. Module supports all management functions that are consistent with the Low Power condition	ResetS transition signal is TRUE	Resetting	Rqd.
			FaultS transition signal is TRUE	ModuleFault	
			LowPwrS transition signal is FALSE	ModulePwrUp	
ModulePwrUp	High Power	Module transitioning to high power mode	ResetS transition signal is TRUE	Resetting	Rqd.
			FaultS transition signal is TRUE	ModuleFault	
			LowPwrS transition signal is TRUE	ModulePwrDn	
			Power up activities are complete	ModuleReady	
ModuleReady	High Power	Module may be consuming power up to the level defined in the fields in Table 8-31	ResetS transition signal is TRUE	Resetting	Rqd.
			FaultS transition signal is TRUE	ModuleFault	
			LowPwrExS transition signal is TRUE	ModulePwrDn	
ModulePwrDn	High Power	Module transitioning to Low Power mode	ResetS transition signal is TRUE	Resetting	Rqd.
			FaultS transition signal is TRUE	ModuleFault	
			Module has returned to Low Power mode	ModuleLowPwr	
ModuleFault	*Low Power	Module is waiting for host action	Module power down	N/A	Opt.
			ResetS transition signal is TRUE	Resetting	

*It is suggested, if possible, that the ModuleFault state be in Low Power mode.

2

3

6.3.2.3 Module State Machine for Flat Memory Modules (Simplified)

Flat Memory Modules shall adhere to the behaviors described by the simplified Module State Machine shown in Figure 6-4.

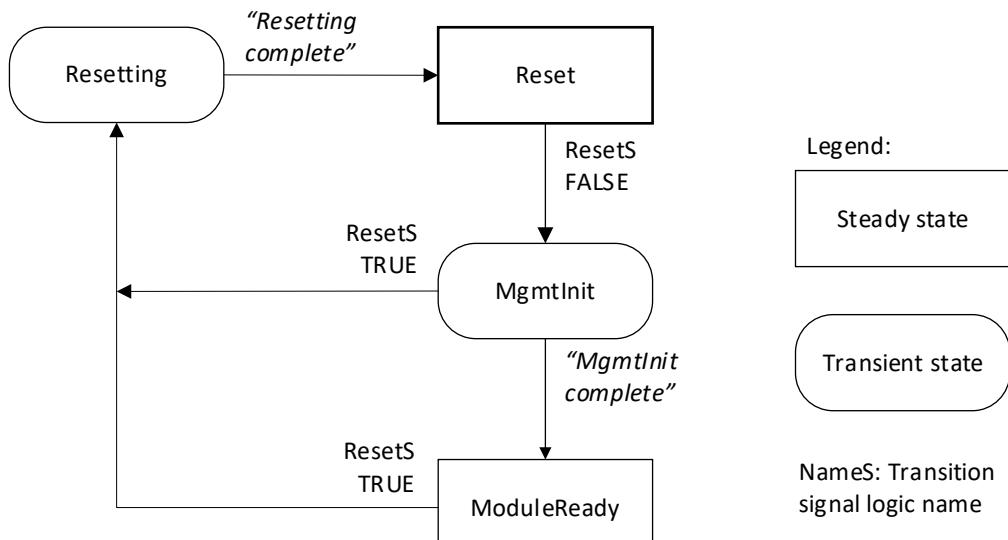


Figure 6-4 Flat Memory Module State Machine (MSM) State Transition Diagram

As shown in Figure 6-4, flat memory modules transition to the **ModuleReady** state without host interaction.

Module state transitions for flat memory modules are just a **specification formalism** since the contents of the static EEPROM-based Memory Map does not change. Although the behaviors of the **Resetting**, **Reset**, and **MgmtInit** states apply to such modules, those **states are not reported** to the host through the Memory Map.

Flat memory modules shall statically report a module state of **ModuleReady** (see Table 8-5).

The **ResetS** transition signal is described in Table 6-11. For flat memory modules, assertion of the **Reset** signal may optionally hold the EEPROM in reset, however the Data Path remains active.

Flat memory modules are not required to support **SoftwareReset** or **VccResetL**.

At initial module insertion or application of power, the (conceptual) Module State Machine initializes to the **Reset** state. Table 6-15 provides a summary of the high-level behaviors and properties of each module state for flat memory modules. Refer to sections 6.3.2.5.1-6.3.2.5.8 for detailed requirements for each state.

Table 6-15 Module state behaviors, flat memory modules

State	Power Mode	Behavior in state	Exit condition	Next state	Required/Optional
Resetting (not reported)	Low Power	Management interface transitions to reset	Resetting completed	Reset	Optional
Reset (not reported)	Low Power	Management interface in reset	ResetS transition signal becomes TRUE	MgmtInit	Optional
MgmtInit (not reported)	Low Power	Management interface powering up and initializing until host can access the flat management memory	ResetS transition signal becomes TRUE	Resetting	Required
			Module Management Interface ready is simply assumed after a duration tMgmtInit (see section 10.1), so hosts should just wait for that period.	ModuleReady	
ModuleReady (reported)	Low Power	Management interface available for host to access the flat management memory	ResetS transition signal becomes TRUE	Resetting	Required

1 6.3.2.4 Module Power Mode (Module Characteristic)

2 The **Module Power Mode** dictates the maximum electrical power that the module is permitted to consume
3 while operating in that Module Power Mode.

4 The Module Power Mode is a function of the state of the Module State Machine.

5 Two Module Power Modes are defined:

- 6 • In **Low Power Mode** (characteristic of all MSM steady states except ModuleReady) the maximum
7 module power consumption is defined in the form factor-specific hardware specification.
- 8 • In **High Power Mode** (characteristic of the MSM state ModuleReady) the implementation dependent
9 maximum module power consumption is advertised in the **MaxPower** Byte 00h:201 (see Table 8-31).

10 All modules initially boot in Low Power Mode. After the management interface has been initialized and the MSM
11 has reached the steady ModuleLowPwr state the host may eventually transition paged memory modules to High
12 Power Mode using the conditions defined by the LowPwrS transition signal (see Table 6-12) provided that the
13 advertised MaxPower value is supported in the host system.

14 If LowPwrS is FALSE when the module is in the ModuleLowPwr state, the module begins to enable High Power
15 Mode operation, using the power up procedures defined for the ModulePwrUp state (see section 6.3.2.5.5).

16 Conversely, whenever LowPwrS or LowPwrExS (as applicable) is TRUE while the module is in or moving towards
17 High Power Mode, the module begins to return to the ModuleLowPwr state and hence to Low Power Mode
18 operation, using the power down procedures defined for the ModulePwrDn state (see section 6.3.2.5.7).

19 *Note: The LowPwrS transition signal only controls the Module Power Mode but it does not control Data Path
20 initialization. Refer to section 6.3.3 for more information on Data Path initialization.*

1 **6.3.2.5 Module State Machine States**

2 **6.3.2.5.1 Resetting State (Shutting Down)**

3 The **Resetting** state is a transient state.

4 The Resetting state is entered from any state except the Reset state when the ResetS transition signal is TRUE.
5 The ResetS transition signal is defined in Table 6-11.

6 **On Entry**

7 When a paged memory module enters the Resetting state, all Data Path State Machines cease to exist. Refer
8 to section 6.3.3 for Data Path State Machine behaviors.

9 **Autonomous Behavior**

10 The module tries to gracefully power down module optics and electronics before entering the Reset state. The
11 Resetting state initiates a complete module reset. The shutdown procedure used by the module for a reset
12 event is implementation dependent.

13 *Note: The module may be in High Power Mode during portions of the Resetting state.*

14 **Reactive Behavior**

15 Management interactions initiated by the host during the Resetting state may be ignored by the module.
16 Transactions in progress may be aborted when entering the Resetting state. Note: While the ResetS transition
17 signal is TRUE, the management communication interface may be held in reset and may not respond (NACK).

18 **Exit**

19 When all module electronics have been powered down and are in reset, the module state transitions
20 automatically to the Reset state.

22 **6.3.2.5.2 Reset State (Ground State)**

23 The **Reset** state is a steady state.

24 The Reset state is the initial Module State Machine state on module insertion or power-up.

25 **On Entry**

26 The self-clearing SoftwareReset bit (see Table 8-11) effectively returns to its default value. All other bits are
27 not affected.

28 *Note that this means the module must have some mechanism to clear this bit when in the Reset State or upon
29 exiting the Reset State. Other bits will be set to their power-up default values later, in the MgmtInit State,
30 regardless of their value when exiting the Reset State.*

31 **Autonomous Behavior**

32 All internal module electronics are held in reset.

33 The module remains in Low Power mode.

34 All Interrupts to the host are suppressed.

35 **Reactive Behavior**

36 The module may abort any MCI transactions in progress on entry to the Reset state.

37 The module may ignore all MCI transactions while in the Reset state.

38 *Note: While the ResetS transition signal is TRUE, the management interface may be held in reset and may not
39 respond (NACK); when the ResetS transition signal has become FALSE, the module may still not respond until
40 the subsequent MgmtInit state is exited.*

41 **Exit**

42 The module exits the Reset state when the ResetS transition signal is found FALSE (see Table 6-11).

43 The Reset state can only be exited when the ResetS transition signal is FALSE and power is applied.

44 On exit from the Reset state, the Module State Machine enters the MgmtInit state.

46 **6.3.2.5.3 MgmtInit State (Initializing Management Interface)**

47 The **MgmtInit** state is a transient state.

1 The MgmtInit state is applicable to both paged memory modules and flat memory modules.

2 **Entry**

3 The MgmtInit state is entered any time the module comes out of the Reset state.

4 **Autonomous Behavior**

5 During the MgmtInit state, the module initializes the Memory Map to default values and initializes the
6 management communication interface allowing the host to eventually manage the module.

7 The module may perform limited power-up of Data Path circuitry provided the module remains in Low Power
8 Mode throughout this state.

9 *Note that, for paged memory modules, all Data Path States are DPDeactivated in MgmtInit.*

10 **Reactive Behavior**

11 The module may ignore all MCI transactions while in the MgmtInit state.

12 **Exit**

13 The regular exit occurs automatically when the specified autonomous behavior has run to completion.

14 The next state after regular exit is ModuleLowPwr.

15 On regular exit to ModuleLowPwr state the module has ensured that all Memory Map register locations have
16 been set to their power-on default values.

17 **6.3.2.5.4 ModuleLowPwr State (Basic Manageability)**

18 The **ModuleLowPwr** state is a steady state.

19 It represents the situation where the **management interface** is fully initialized and operational while the
20 device maintains its Low Power Mode.

22 *Note: The Data Path state of all lanes is still DPDeactivated in the ModuleLowPwr state.*

23 During this state, the host may query and configure the module using the management interface to read from
24 and write to the management Memory Map.

25 *Note: Some examples of query and configuration activities include reading the ID and device property fields,
26 reading capability advertisements including, if supported, CDB advertisement messages, setting CDR and other
27 lane attributes, and configuration of Interrupt related Masks.*

28 Details of host behavior and of resulting host-module interactions in the ModuleLowPwr state are outside the
29 scope of this specification.

30 *Note: This is typical for a host-controlled steady state, where interactions are initiated by the host. For instance,
31 a host may choose to advance to the ModuleReady state before reading all advertisements, or they may do
32 that in ModuleLowPwr.*

33 **On Entry**

34 On entry into the ModuleLowPwr state, the module sets the Module State register (Table 8-5) to the
35 ModuleLowPwr state and it sets the ModuleStateChangedFlag (Table 8-9) unless the ModuleLowPwr exit criteria
36 are already met upon entry into the state.

37 **Autonomous Module Behavior**

38 The module waits for exit conditions to become true.

39 **Reactive Behavior**

40 The module reacts to all management operations but is allowed to ignore or reject commands that are **not**
41 **consistent with the Low Power** condition¹.

42 *Note: In other words, the module supports all management functions (generally required by CMIS or advertised
43 by the module as being supported) that are consistent with the characteristic Low Power condition of the
44 ModuleLowPwr state.*

¹ Unfortunately, lacking corresponding advertisement, a CMIS5.2 host is unaware exactly which management functions would violate the Low Power condition in a specific module implementation. Since CMIS 5.3, modules advertise their LowPowerRestrictions as described in Table 8-60.

1 Note: Management functions which require that a Data Path has been initialized are, by definition, considered
2 inconsistent with the Low Power condition.

3 **Exit**

4 The module state transitions to ModulePwrUp when the LowPwrS transition signal is FALSE (see Table 6-12).

5 Note: This transition can occur at any time during ModuleLowPwr, and so modules shall sense LowPwrS
6 throughout the ModuleLowPwr state. This behavior was defined differently in CMIS 3.0

7 Note: The LowPwrS transition signal may evaluate to 0 the first time it is sampled in ModuleLowPwr. In such
8 circumstances, the transition to ModulePwrUp may be too fast for the host to detect that the module was
9 transiently in the ModuleLowPwr state.

10 **6.3.2.5.5 ModulePwrUp state (Powering Up)**

11 The **ModulePwrUp** state is a transient state. In this state the host is informed that the module is in the process
12 of powering up to High Power Mode.

13 Note: The module is expected to power up module components as needed to expedite later module and Data
14 Path reconfigurations for operational use.

15 Note that the Data Path State of all lanes is DPDeactivated in the ModulePwrUp state.

16 **On Entry**

17 On entry the module shall set the Module State register (Table 8-5) to ModulePwrUp.

18 **Autonomous Behavior**

19 The module may be in High Power mode at any time during the ModulePwrUp state.

20 **Reactive Behavior**

21 The module reacts to all management operations.

22 **Exit**

23 When the **LowPwrS** transition signal is TRUE at any time during the **ModulePwrUp** state, the module state
24 immediately transitions to **ModulePwrDn**.

25 When the module power up sequence has completed, the module state transitions to the **ModuleReady** state.

26 **6.3.2.5.6 ModuleReady State (Fully Operational)**

27 The **ModuleReady** state is a steady state.

28 In this state the module is in High Power mode and the host may initialize or deinitialize Data Paths.

29 **On Entry**

30 On entry into the **ModuleReady** state, the module shall set the Module State register (Table 8-5) to the
31 ModuleReady state and set the ModuleStateChangedFlag (Table 8-9).

32 **Autonomous Behavior**

33 The module operates as configured by the host.

34 **Reactive Behavior**

35 The module reacts to all management operations and to all relevant input signal changes.

36 **Exit**

37 Except for Reset or Fault, the action that results in an exit from ModuleReady is if the **LowPwrExS** transition
38 signal is TRUE (see Table 6-13), which causes the module state to transition to **ModulePwrDn**.

41 **6.3.2.5.7 ModulePwrDn State (Powering Down)**

42 The **ModulePwrDn** state is a transient state.

43 In this state the host is informed that the module is in the process of returning to Low Power mode.

44 Note that the Data Path State for all lanes is DPDeactivated in the ModulePwrDn state.

45 **On Entry**

46 On entry into the ModulePwrDn state, the module shall set the Module State register (Table 8-5) to the
47 ModulePwrDn state.

Autonomous Behavior

The module reduces the power consumption of module electronics such that the module power consumption is less than the Low Power Mode threshold. The electronics associated with the management interface remains powered and available.

The module may still be in High Power mode at any time during the ModulePwrDn state.

Reactive Behavior

Modules in ModulePwrDn shall ignore the LowPwrS transition signal, so if this signal is FALSE during ModulePwrDn, the module will complete the power-down sequence and transition to ModuleLowPwr before sampling LowPwrS again.

Exit

When the module has arrived in Low Power Mode, the module state transitions to the ModuleLowPwr state.

6.3.2.5.8 ModuleFault State (Module Fault)

The **ModuleFault** state is a steady state.

The ModuleFault state is provided for notification to the host that a module **fault** has occurred after which the module aims to prevent physical damage and to avoid safety risks for the module or for its environment (e.g. host).

The exact ModuleFault state entry conditions are implementation dependent. The ModuleFault state is only entered when module detects a condition (e.g. TEC runaway, memory corruption) that could compromise safety or cause damage. The specification intent of the ModuleFault state is to put the module in a condition that does not compromise safety or create further equipment failures.

On Entry

On entry the module shall set the ModuleState register (see Table 8-6) to ModuleFault.

The module may write information about the fault cause into the ModuleFaultCause register (see Table 8-16)

Autonomous Behavior

It is strongly recommended that the module enters Low Power mode during the ModuleFault state but the response to a Fault condition is implementation specific.

Reactive Behavior

The module reacts only to events that cause the ResetS transition signal to become TRUE.

The module may respond to READ access if that is still possible.

Exit

Except for a power cycle, the only exit path from the ModuleFault state is to perform a module reset by taking an action that causes the ResetS transition signal to become TRUE (see Table 6-11).

6.3.3 Data Path State Machines (DPSM)

A **Data Path State Machine (DPSM)** instance describes Data Path-specific behaviors and properties that are related to the configuration of the Data Path, as managed by the host.

Note: Recall that the text in this section literally applies to system interface applications only. For the more complex but optional case of client encapsulation applications (a.k.a. multiplex applications), it must be read with a slightly different interpretation, as described in section 7.6. For instance, the DPSM described here applies to the entire Data Path of a system interface application, whereas in client encapsulation applications a separate DPSM is used for each Host Path (which is then only a host side segment of the overall multiplex Data Path).

Note: For configuration dependent behaviors and properties of the module as a whole, refer to the Module State Machine described in section 6.3.2.

Note: As described in section 6.2.2, a Data Path is a bidirectional combination of one or more host lanes, one or more media lanes, and a set of internal module resources implementing the Application that is described in an associated Application Descriptor.

*Note: The DPSM state represents a **management** or **configuration realization status** of a Data Path, representing the effects of certain host configuration commands and of module reactions to those commands. It **does not** necessarily represent other behavioral or operational aspects of a bidirectional Data Path, e.g. in terms of current input or output signal conditions or in terms of transmission service being provided.*

*Note: The DPSM state should neither be confused with the **operational status** of the functional resources of a Data Path in Tx direction or in Rx direction, nor with the resulting signal **output status** of Rx host lane outputs or of Tx media lane outputs, which are reported independently in separate Output Status registers (see section 8.10.2).*

Data Paths (and Data Path State Machines) are only applicable to **paged memory modules**.

Module State and DPSM Life Cycle

All **DPSM** instances required to represent the power-up default Application defined in the Data Path Configuration field values of the Active Control Set are initially created and set-up during the **MgmtInit** state.

After its creation, a DPSM remains in the DPDeactivated State until the Module State Machine is in the **ModuleReady** state and an exit condition from the DPDeactivated state is met.

When the host updates the Data Path Configuration fields in the Active Control Set, in either the **ModuleLowPwr** or **ModuleReady** states, the module tears down any previous DPSM that is no longer defined and then creates and sets up any newly defined DPSM.

All Data Path State Machines are torn down in the **Resetting** state.

Note: Refer to section 6.3.2 for an overview of the Module State Machine, section 6.2.1 for an overview of Applications, section 6.2.2 for an overview of Data Paths, and section 6.2.3 for an overview of Control Sets.

DPSM Purpose

A Data Path State Machine is used by the module to represent the initialization status of the resources associated with a Data Path in response to certain host configuration settings or commands. Although individual resources within a Data Path may complete initialization activities at different times, the module waits to report the updated DPSM state until all resources associated with the Data Path have completed the requested configuration or reconfiguration action. This synchronized status reporting across all lanes and resources in a Data Path reflects the fact that there is only one Data Path State Machine per Data Path.

Note: The DPSM specification model does not imply any specific way of implementation.

Note: Some example Data Path initialization sequences are provided in Appendix D.

Note: Modules identify supported Data Path configurations through the Application advertisement fields.

DPSMs for parallel Data Paths of Multiple Application Instances

Each Data Path in a module is required to operate independently of other Data Paths: if the host changes the Data Path State of one Data Path, the other Data Paths in the module shall be unaffected and uninterrupted.

Note: Independent operation of Data Paths may require independent clocking per Data Path, from a recovered clock within that Data Path. See the applicable hardware specification for further information. The module shall only advertise Applications and lane configurations that are supported by the implemented clocking scheme.

6.3.3.1 DPSM State Transition Diagram and DPSM Specification

Figure 6-5 shows the state transition diagram (STD) of a DPSM representing the Data Path configuration related state of one Data Path instance¹.

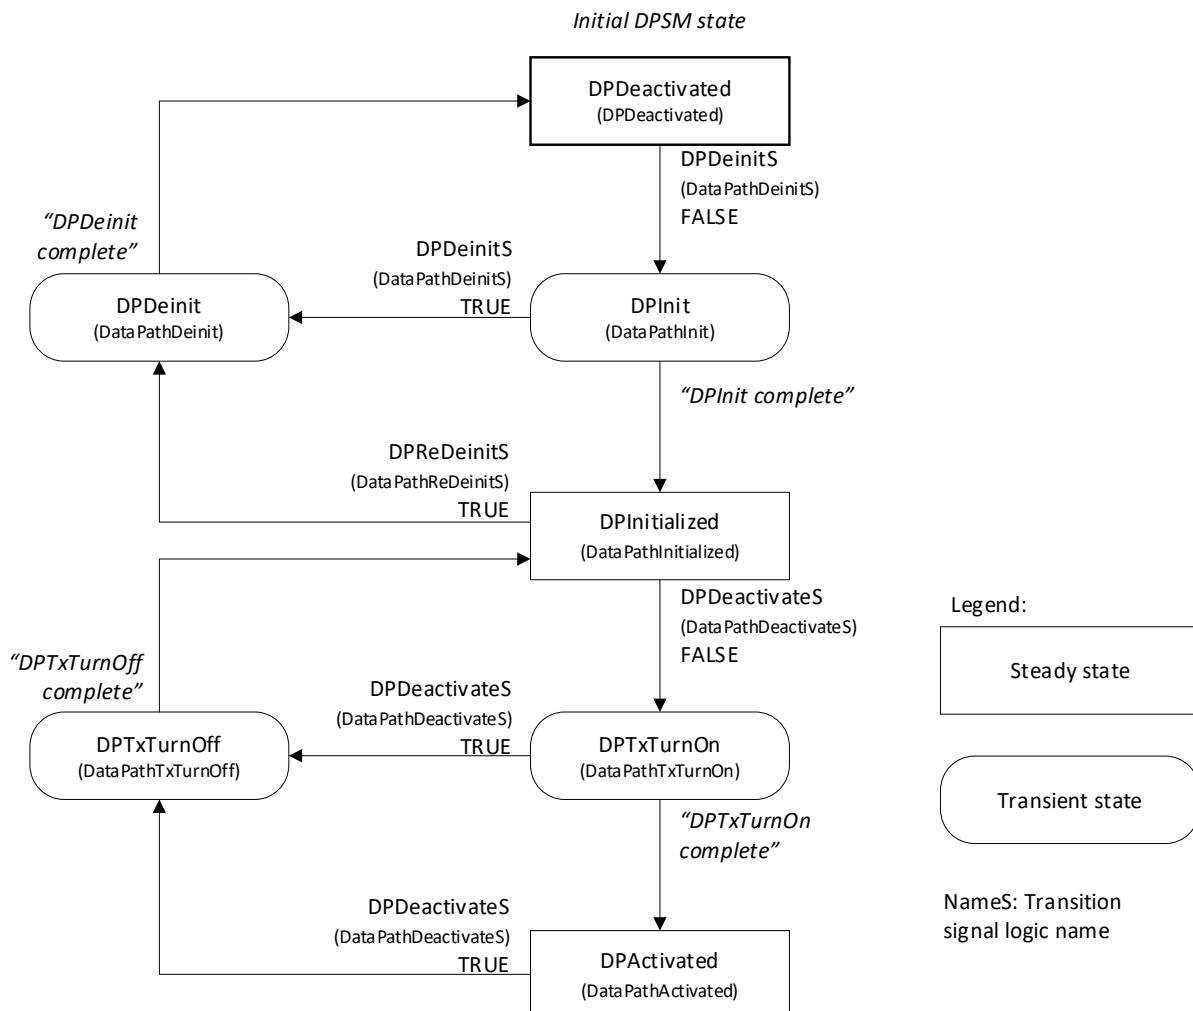


Figure 6-5 Data Path State Machine (DPSM) State Transition Diagram

Note: Prior to exit from the MgmtInit module state, all Data Paths initialize to the DPDeactivated state.

State Exit Conditions and Transition Signals

The state machine exits a given state when specific exit conditions are satisfied. So called **Transition Signals** (recognized by name suffix S) represent these exit conditions for steady states.

DPDeinitS (DataPathDeinitS)

The DPDeinitS (DataPathDeinitS) transition signal is defined using the truth table shown in Table 6-16.

Table 6-16 DPDeinitS transition signal truth table

ModuleReadyT (MSM term)	LowPwrS (MSM, Table 6-12)	DPDeinitT (term)	DPDeinitS (transition signal)
0	X	X	1
1	1	X	1
1	0	1	1
1	0	0	0

¹ For brevity, a 'DataPath' prefix in the name of a state, a transition signal, or a named logical predicate may always be replaced by 'DP'. The shortened names are fully equivalent synonyms to the full names.

The DPDeinitS transition signal can also be represented by the logic equation

$$\text{DPDeinitS} = (\text{NOT ModuleReadyT}) \text{ OR LowPwrS OR DPDeinitT} \quad (\text{Eq. 6-6})$$

where

$$\text{ModuleReadyT} = (\text{ModuleState} = \text{ModuleReady}) \quad (\text{Eq. 6-7})$$

$$\text{DPDeinitT} = \begin{aligned} & \text{DPDeinitLane<N>} \\ & \text{OR DPDeinitLane<N+1>} \\ & \dots \\ & \text{OR DPDeinitLane<N+M-1>} \end{aligned} \quad (\text{Eq. 6-8})$$

N = first host lane in the Data Path

M = number of host lanes in the Data Path

DPReDeinitS (DataPathReDeinitS)

The definition of the DPReDeinitS transition signal depends on (advertised) module capabilities. Specifically, the DPReinitT term is included only when intervention free regular reconfiguration is supported (see Table 8-5).

Case 1: Regular Intervention-Free Reconfiguration is Supported

The DPReDeinitS (DataPathReDeinitS) transition signal is defined using the truth table shown in Table 6-17 when the module does support intervention-free reconfiguration (for advertisement see Table 8-5).

Table 6-17 DPReDeinitS transition signal truth table (default)

DPDeinitS (transition signal)	DPReinitT (term)	DPReDeinitS (transition signal)
1	1	1
1	0	1
0	1	1
0	0	0

The DPReDeinitS transition signal can also be represented by the logic equation

$$\text{DPReDeinitS} = \text{DPDeinitS OR DPReinitT} \quad (\text{Eq. 6-9})$$

where

$$\text{DPReinitT} = \begin{aligned} & \text{DPInitPendingLane<N>} \\ & \text{OR DPInitPendingLane<N+1>} \\ & \dots \\ & \text{OR DPInitPendingLane<N+M-1>} \end{aligned} \quad (\text{Eq. 6-10})$$

N = first host lane in the Data Path

M = number of host lanes in the Data Path

The **DPInitPending** status register is described in section 8.10.7. Each **DPInitPendingLane<i>** bit represents the condition following a successful **ApplyDPInit** trigger that a transit through state **DPInit** is pending.

The module simultaneously **sets** the DPInitPendingLane<i> bits for all triggered lanes <i> of a Data Path after it has successfully copied all settings for all those triggered lanes <i> from a selected Staged Control Set to the Active Control Set during a **Provision** procedure that was triggered by a host WRITE to one of the two SCS<k>::**ApplyDPInit** trigger registers (see section 8.9.3.1, 8.9.4.1).

*Note: DPInitPending bits are **not** set in response to **ApplyImmediate** triggers.*

The module **clears** all DPInitPendingLane<i> bits of a Data Path while in DPSM state **DPInit**.

*Note: The desired effect of the DPInitPending condition variable is to **report a discrepancy** between Active Control Set and hardware until it is resolved in the DPInit state. When intervention-free reconfiguration is supported (by default: SteppedConfigOnly = 0), this discrepancy causes an intervention-free DPSM state cycle through DPDeactivated and DPInit after ApplyDPInit was successfully triggered in DPInitialized or DPActivated.*

Case 2: Regular Intervention-Free Reconfiguration is Not Supported

When the module does not support for intervention-free reconfiguration (see Table 8-5), then the DPReDeinitS transition signal is simplified as follows (the DPReinitT term is omitted):

$$\mathbf{DPReDeinitS} = \mathbf{DPDeinitS} \quad (\text{Eq. 6-11})$$

Note: The intentional change in behavior, when only stepwise configuration is supported, is that ApplyDPIInit has no impact on the DPSM dynamics at all and is therefore fully independent of the DPSM.

DPDeactivateS (DataPathDeactivateS)

The DPDeactivateS (DataPathDeactivateS) transition signal is defined by the logic equation

$$\text{DPDeactivateS} = \text{DPReDeinitS} \text{ OR } \text{DPTxDisableT} \text{ OR } \text{DPTxForceSquelchT} \quad (\text{Eq. 6-12})$$

where

$$\begin{aligned} \mathbf{DPTxDisableT} = & \quad \text{OutputDisableTx}\langle N \rangle \\ & \text{OR OutputDisableTx}\langle N+1 \rangle \\ & \dots \\ & \text{OR OutputDisableTx}\langle N+M-1 \rangle \end{aligned} \quad (\text{Eq. 6-13})$$

$$\begin{aligned} \mathbf{DPTxForceSquelchT} = & \quad \text{OutputSquelchForceTx}\langle N \rangle \\ & \text{OR OutputSquelchForceTx}\langle N+1 \rangle \\ & \dots \\ & \text{OR OutputSquelchForceTx}\langle N+M-1 \rangle \end{aligned} \quad (\text{Eq. 6-14})$$

N = first media lane in the Data Path

M = number of media lanes in the Data Path

Note: When in **DPActivated** or **DPTxTurnOn** states, setting *OutputDisableTx* or *OutputSquelchForceTx* on one media lane of a Data Path intentionally causes that entire Data Path to transition to **DPIinitialized** via *DPTxTurnOff*. Although some media lanes of the Data Path may continue to be operational (i.e. media lanes with output neither disabled nor squelched), as long as some media lanes are not operational, the Data Path as a whole is considered not activated.

Reaction to Module Reset

When the MSM **ResetS** transition signal (see Table 6-11) becomes TRUE, any Data Path related power down activities are performed in the **Resetting** module state. The DPSM state machines cease to exist in this case.

Note: Module dependent pre-reset clean up and power down activities may be implemented, possibly depending also on the reset trigger in either hardware or software.

Reaction to Module Fault

When Module State Machine transitions to **ModuleFault** state, the DPSM behavior is not defined formally, but governed by the behavioral requirements of the ModuleFault state.

DPSM State and Tx Output Status

Table 6-18 provides a high-level summary of the Tx output behaviors and characteristics of each Data Path state and outlines the meaning of exit conditions of transient states that are not captured in formal transition signal definitions. Refer to sections 6.3.3.4-6.3.3.10 for detailed requirements for each state.

Table 6-18 Data Path state behaviors and Exit Conditions

State	Output Status Tx	Exit condition	Next State
DPDeactivated	Quiescent	DPDeinitS is FALSE	DPIInit
DPInit	Quiescent	DPDeinitS is TRUE DPInit complete: Data Path resource initialization completed	DPDeinit DPInitialized
DPInitialized	Depends on per-lane OutputDisableTx and OutputSquelchForceTx	DPReDeinitS is TRUE DPDeactivateS is FALSE	DPDeinit DPTxTurnOn
DPDeinit	In transition	DPDeinit complete: Module dependent deinitialization completed	DPDeactivated
DPTxTurnOn	In transition	DPDeactivateS is TRUE DPTxTurnOn complete: All Data Path Tx outputs are operational	DPTxTurnOff DPActivated
DPActivated	Operational	DPDeactivateS is TRUE	DPTxTurnOff
DPTxTurnOff	In transition	DPTxTurnOff complete: Physical status of some Tx outputs reflects the nominal	DPInitialized

State	Output Status Tx	Exit condition	Next State
		configuration of being squelched or disabled on host command	

1 DPSM State and Rx Output Status

3 Except for an implicit dependency on the initialization of internal functional resources of the Rx Data Path on
4 host and media side, which occurs in the first DPInit state traversal, the Rx output state is **not further**
5 **controlled** by the Data Path State Machine.

6 After Rx Data Path resource initialization in the DPInit state, a module always forwards a **valid** Rx input signal
7 to a valid Rx output signal in the DPInitialized, DPTxTurnOn, DPTxTurnOff, DPActivated states (unless
8 overridden by Rx output muting host commands).

9 When **returning** to the other DPSM states (DPDeinit, DPDeactivated, DPInit), a module may or may not mute
10 the Rx output signal; especially it may transmit a **valid** Rx output signal when a valid Rx input signal is available.

11 *Note: See section 3.3 for the definition of a **valid** signal.*

12 The host can always ensure a muted Rx output using the OutputDisableRx control (Table 8-69).

13 The module mutes the Rx output when supported automatic squelching functionality is enabled and when the
14 relevant squelch conditions are present.

15 The module always reports the **resulting actual status** of the Rx output signal in the OutputStatusRx register,
16 independent of the Data Path State Machine state.

17 6.3.3.2 Data Path Control (Host)

18 A single main configuration register is provided for the host to control initialization and deinitialization of all
19 Data Paths represented in a given Bank. This **DPDeinit** register (see Table 8-68) defines per host lane if lane
20 or the associated Data Path resources are determined to be unused for functional operation (and hence can be
21 deinitialized) or if they are determined for functional operation (and hence need to be initialized).

22 *Note: Per-lane configuration has been chosen to allow a variety of Data Path configurations from a single
23 Memory Map specification.*

24 *Note: Initialization status and behavior of Tx media lane outputs are further controlled by the host using the
25 media lane specific control bits **OutputDisableTx<i>** and **OutputSquelchForceTx<i>**.*

26 A host requesting initialization or deinitialization of a Data Path ensures that the Active Control set contains the
27 desired configuration settings and then writes the value 0 or 1, respectively, to the DPDeinit bits associated
28 with the lanes of that Data Path.

29 The host may request initialization or deinitialization of multiple Data Paths with one register access.

30 Some informative Data Path initialization flow examples are provided in Appendix D to facilitate understanding
31 of the relationship between the initialization of physical structures in the module and Data Path-level reporting
32 in the Memory Map.

33 6.3.3.3 Data Path Status (Module)

34 The module provides information on the current state of the Data Path (DPSM **current state reporting**) and
35 on entry to certain DPSM states (DPSM **state change indication**).

36 DPSM Current State Reporting

37 On entry to a DPSM state the module reports the DPSM state entered as the current DPSM state in the **DPState**
38 status register (see Table 8-83), on all lanes of the Data Path, with optional exceptions specified below.

39 *Note: Due to the identical behavior of all lanes of a Data Path the host needs to read only the first lane of the
40 Data Path to determine the Data Path state.*

41 *Note: The DPSM model in this specification describes a single state machine per Data Path, with Data Path
42 attributes replicated on all lanes of the Data Path. This does not limit software implementations.*

43 DPSM Current State Reporting Exceptions

44 The module **may** suppress reporting the current DPSM state in the DPStateHostLane<i> registers when that
45 state is known to be transitional, i.e. when it is exited immediately because its exit conditions are fulfilled on
46 entry, or when the duration of staying in that state is known to be in the order of 1 ms or less.

1 Note: The duration specification is intentionally vague. The intention for allowing exceptions in state reporting
 2 is to avoid reporting short-lived status data which the host is unlikely to read and react upon.

3 DPSM State Change Indication (Flag)

4 A DPSM State Change Indication consist of the module setting a **DPStateChangedFlag** for each lane of the
 5 Data Path associated with the relevant DPSM instance.

6 Note: The intention of the following specification is that the module indicates a state change only on entry to a
 7 lasting steady state and only when the transition time since the previous lasting steady state was significant.

8 The maximum duration of a transient state is advertised in a MaxDuration* field (see Table 8-47, Table 8-56)
 9 and is considered insignificant when the coded MaxDuration* field value is 0000b (see Table 8-48).

10 The module performs a DPSM State Change Indication on entry to a **steady** DPSM state when the following
 11 two Flag setting conditions are fulfilled:

- 12 • No exit condition of the entered **steady** state is fulfilled on state entry (state is not visited transitorily)
- 13 • The advertised maximum state **duration** is **significant** for at least one **transient** state passed through
 14 since the previous State Change Indication for this Data Path, or since module reset if there was no such
 15 State Change Indication yet since reset.

16 The module does **not** perform a DPSM State Change Indication on entry to a **transient** DPSM state.

17 Table 6-19 defines the Flag behavior for each DPSM state entry.

18 **Table 6-19 Data Path State Changed Flag behaviors**

Entered state	DPStateChangedFlag may be set *
DPIInit	No
DPIInitialized	Yes
DPDeinit	No
DPDeinit	No
DPDeactivated	Yes
DPTxTurnOn	No
DPActivated	Yes
DPTxTurnOff	No
DPTxTurnOff	No
DPIInitialized	Yes

19 * Note: The Flag setting conditions are described in the main text.

20
 21 Note: Steady state exit conditions may already be met upon entry into the steady state and lead to immediate
 22 transition to the next state (after state entry or state exit activities, if defined).

23 Flag Related Behavior

24 The module does not clear any Flag due to a state change of a Data Path State Machine.

25 The module raises Flags only according the DPSM state-specific conformance rules defined in section 6.3.4.

26 6.3.3.4 DPDeactivated State (Ground State)

27 The **DPDeactivated** (DataPathDeactivated) state is a steady state.

28 This per-lane state indicates that no Data Path is initialized on the indicated lane(s).

29 The host may configure or reconfigure Data Paths on lanes reporting DPDeactivated.

30 Note: Reconfiguration begins with update of the Active Control Set.

31 Note: When SteppedConfigOnly=1, the host may reconfigure Data Paths in other states as well, but this is not
 32 recommended, as described in section 6.2.4.3.

33 Autonomous Behavior

34 On entry to this state, the module updates the Data Path state register (see Table 8-83) and the Data Path
 35 State Changed Flag (Table 8-87) for all lanes in the Data Path according to the rules described in section 6.3.3.3.

36 While in DPDeactivated state, all Tx media lane outputs of the Data Path shall be quiescent while the status of
 37 all Rx host lane outputs of the Data Path are undefined.

1 Note: The Rx output status of the Rx host lanes of a Data Path depends on host controlled configuration history
2 and on unspecified module controlled deinitialization behaviors (see sections 6.3.3.7 and 6.3.3.10).

3 Reactive Behavior (on Host Actions)

4 Changes to the OutputDisableTx or OutputSquelchForceTx register values for Data Paths in DPDeactivated shall
5 have no impact on the output quiescence of those Data Paths.

6 Exit

7 The Data Path remains in DPDeactivated as long as ResetS is TRUE (see Table 6-11).

8 Otherwise, the DPSM transitions to DPInit when the DPDeinitS transition signal is FALSE (see Table 6-16).

9 The Host shall provide a **valid** high-speed Tx input signal at the required signaling rate and encoding type prior
10 to causing a DPSM to exit the DPDeactivated state.

11 Note: The module must receive a valid input signal while performing initialization activities in the subsequent
12 DPInit state, such as adaptation of signal integrity equalizer settings. Otherwise, if no valid Tx input signal is
13 available, the resulting initialization behavior is not predictable (the module may or may not wait for a valid
14 signal and may or may not initialize to a possibly inadequate condition).

15 6.3.3.5 DPInit State (Initializing)

16 The **DPInit** (DataPathInit) state is a transient state.

17 In this state the module performs all initialization activities on the internal resources of the Data Path that are
18 necessary to make the Data Path operational.

19 Note: Initialization activities include the realization of the selected Application properties and/or adaptation of
20 signal integrity settings, and possibly power up and initialization of Tx and Rx Data Path electronics if
21 opportunistic power savings were employed by the module in DPDeactivated.

22 Note: The module assumes a valid high-speed Tx input signal when entering the DPInit state.

23 Autonomous Behavior

24 On entry to this state, the module updates the Data Path state register (see Table 8-83) for all lanes in the Data
25 Path according to the rules described in section 6.3.3.3.

26 Within the DPInit state, the module performs any necessary power-up and initialization activities for module
27 electronics associated with the Data Path.

28 Note: In some cases, these electronics may be shared between multiple Data Paths. Depending on prior power
29 up and power down actions, some or all of these electronics may already be powered; in such cases, the power
30 up sequence is bypassed. The details of the power up sequence are implementation-dependent and outside the
31 scope of this specification.

32 During DPInit, the module realizes the Application properties and signal integrity settings as defined in the
33 Active Control Set (see section 6.2.3) by configuring the relevant internal resources of the Data Path.

34 Note: The details of how the module applies Application settings is implementation-dependent and outside the
35 scope of this specification.

36 Note: Attributes that require adaptation, such as CTLE settings, are adapted at the appropriate time during
37 DPInit. The order in which signal integrity settings are applied and adapted is implementation-dependent and
38 outside the scope of this specification.

39 The module clears the DPInitPending bits of all lanes in the Data Path.

40 While in DPInit, all Tx media lane outputs of the Data Path shall be quiescent, while the status of all Rx host
41 lane outputs associated with the Data Path are undefined.

42 Note: The actual Rx output status of the host lanes of the Data Path depends on configuration history and on
43 unspecified (i.e. module or vendor specific) deinitialization behaviors (see sections 6.3.3.7 and 6.3.3.10).

44 When no Rx input signal is present at the module Rx input at the time of initialization, the module configures
45 its electronics such that any required adaptation or CDR locking occurs automatically at a later point in time
46 when an input signal is provided, without host intervention.

1 Reactive Behavior (on Host Actions)

2 Changes to OutputDisableTx<i> or to OutputSquelchForceTx<i> controls shall have no impact on the media
3 lane output quiescence of the Data Path.

4 *Note: It is recommended that hosts minimize management operations while in this state. Dynamic Memory Map*
5 *content may be unreliable while in this state and should not be read or written.*

6 Exit

7 If the DPDeinitS logic signal is TRUE at any time during DPInit, the Data Path State Machine transitions to
8 DPDeinit.

9 Otherwise, once the module has completed power-up and initialization of all Tx and Rx resources associated
10 with the Data Path, and all associated Tx and Rx Flags and status registers are valid, the Data Path State
11 Machine state transitions to DPInitialized.

12 *Note: If the module fails to complete DPInit, the host may determine this by observing that the maximum*
13 *advertised duration of DPInit has been exceeded (see MaxDurationDPInit in Table 8-47). Actions taken by the*
14 *host in response to such a failure are outside the scope of this specification.*

15 6.3.3.6 DPInitialized State (Initialized)

16 The **DPInitialized** (DataPathInitialized) state is a steady state.

17 In this state all functional resources of the Data Path are fully initialized. However, the output of one or more
18 Tx media lanes whose Data Path stays in DPInitialized is either squelched by the host or disabled, and so the
19 Data Path is not ready to transmit traffic.

20 Autonomous Behavior

21 On entry to this state, the module updates the Data Path state register (see Table 8-83) and the Data Path
22 State Changed Flag (Table 8-87) for all lanes in the Data Path according to the rules described in section 6.3.3.3.

23 Reactive Behavior (on Host Actions)

24 Transmitter output quiescence for Data Paths in DPInitialized is configured per media lane by the setting in the
25 OutputDisableTx<i> and OutputSquelchForceTx<i> controls.

26 *Note: A media lane output may also be quiescent when a module internal auto-squelching controller reacts to*
27 *a Tx LOS condition on a host lane input and decides to squelch the media lane output.*

28 Exit

29 If the DPReDeinitS transition signal is TRUE at any time during DPInitialized, the Data Path State Machine
30 transitions to DPDeinit.

31 Otherwise, if the DPDeactivateS signal is FALSE at any time during DPInitialized, the Data Path State Machine
32 transitions to DPTxTurnOn. Either of these conditions may be met upon entry into DPInitialized.

33 6.3.3.7 DPDeinit State (Deinitializing)

34 The **DPDeinit** (DataPathDeinit) state is a transient state.

35 In this state the module may deinitialize the module internal resources associated with a Data Path.

36 *Note: Deinitialization tasks are implementation dependent but can include tasks such as opportunistic power*
37 *savings, hardware reconfiguration, or module software variable clean-up.*

38 Autonomous Behavior

39 On entry to this state, the module updates the Data Path state register (see Table 8-83) for all lanes in the Data
40 Path according to the rules described in section 6.3.3.3.

41 During DPDeinit, the module may power down applicable Data Path electronics for opportunistic power savings.
42 In some cases, electronics may be shared with other Data Paths that are not in DPDeinit or DPDeactivated. In
43 such cases, these electronics remain powered. Similarly, modules may keep electronics that require significant
44 power up times powered even when the host requests Data Path deinitialization.

45 *Note: If a host wants to ensure maximum power savings, the host should initiate a module transition to Low*
46 *Power Mode by causing LowPwrS to become TRUE, but this will deactivate all Data Paths in the module.*

1 In DPDeinit, all Tx media lane outputs of the Data Path shall eventually become quiescent while the status of
2 all Rx host lane outputs associated with the Data Path are undefined.

3 *Note: The actual Rx output status of the Rx host lanes of the Data Path depends on configuration history and*
4 *on unspecified (i.e. module or vendor specific) deinitialization behaviors (see also section 6.3.3.10).*

5 **Reactive Behavior (on Host Actions)**

6 Changes to OutputDisableTx or OutputSquelchForceTx for Data Paths in DPDeinit shall have no impact on the
7 output quiescence of those Data Path outputs.

8 *Note: It is recommended the host minimize management operations while in this state. Dynamic Memory Map*
9 *content may be unreliable for lanes in this state and should not be read or written.*

10 **Exit**

11 When the module has completed deinitialization activities on all resources associated with the Data Path, the
12 Data Path State Machine transitions to DPDeactivated.

13 **6.3.3.8 DPTxTurnOn State (Turning On)**

14 The **DPTxTurnOn** (DataPathTxTurnOn) state is a transient state.

15 In this state the module unmutes the Tx output for all media lanes associated with the Data Path.

16 *Note: In somewhat abnormal operational conditions, when auto-squelched by the module, a lane output may*
17 *actually remain or become quiescent, overriding the host configured output status.*

18 **Autonomous Behavior**

19 On entry to this state, the module updates the Data Path state register (see Table 8-83) for all lanes in the Data
20 Path according to the rules described in section 6.3.3.3.

21 While in DPTxTurnOn, all Tx outputs associated with the Data Path shall be in transition and the status of all
22 Rx outputs associated with the Data Path are as reported in the Rx Output Status indicator (see Table 8-85).

23 **Reactive Behavior (on Host Actions)**

24 *Note: It is recommended the host minimize management operations while in this state.*

25 **Exit**

26 If the DPDeactivateS transition signal becomes TRUE at any time during DPTxTurnOn, the Data Path State
27 Machine transitions to DPTxTurnOff.

28 The Data Path state advances to DPActivated once all Tx outputs associated with the Data Path are enabled,
29 have stabilized, and are ready to transmit live traffic.

30 **6.3.3.9 DPActivated State (Operational)**

31 The **DPActivated** (DataPathActivated) state is a steady state.

32 In this state Data Paths are fully operational (initialized and ready to transmit traffic).

33 **Autonomous Behavior**

34 On entry to this state, the module updates the Data Path state register (see Table 8-83) and the Data Path
35 State Changed Flag (Table 8-87) for all lanes in the Data Path according to the rules described in section 6.3.3.3.

36 All Tx outputs associated with the Data Path in DPActivated are unmuted and operational throughout the state.

37 *Note: While operational, Tx outputs may still be auto squelched, overriding the host configured output status.*

38 **Reactive Behavior (on Host Actions)**

39 The module reacts to all events at all interfaces in this DPSM state of providing regular transmission service for
40 the fully operational Data Path

41 **Exit**

42 The Data Path state transitions to DPTxTurnOff if the host causes the DPDeactivateS transition signal to become
43 TRUE for that Data Path.

44 One way for the DPDeactivateS transition signal to become TRUE is if the host triggers the ApplyDPIInit bits
45 associated with the Data Path.

1 The host may reconfigure one or more Data Paths while in DPActivated by defining a new Application in one of
2 the Staged Control Sets and then using ApplyDPInit.

3 The host shall set ApplyDPInit to the same value for all lanes in the Data Path being reinitialized.

4 *Note: When ApplyDPInit bits are triggered, the Data Path State Machine will transition through the DPTxTurnOff*
5 → *DPIInitialized → DPDeinit → DPDeactivated → DPInit → DPInitialized → DPTxTurnOn → DPActivated state*
6 *sequence, reinitializing the new Data Path configuration in DPInit.*

7 Prior to triggering the ApplyDPInit bits for applicable lanes, the host shall provide a valid high-speed input signal
8 at the required signaling rate and encoding type.

9 The ApplyDPInit bits for all lanes in the Data Path shall be triggered with one register access.

10 The host may request reinitialization of multiple Data Paths in the same register access.

11 Data Paths excluded from the ApplyDPInit selector are not affected.

12 *Note: This selective control allows host reconfiguration of individual Data Paths without affecting the operation*
13 *of other Data Paths in the module.*

14 The DPDeactivateS transition signal will also become TRUE if the host sets Tx Output Disable or Tx Force Output
15 Squelch for any lane of the Data Path.

16 **6.3.3.10 DPTxTurnOff State (Turning Off)**

17 The **DPTxTurnOff** (DataPathTxTurnOff) state is a transient state.

18 In this state the module performs the programmed OutputDisableTx and/or OutputSquelchForceTx action on
19 applicable media lanes in the Data Path.

20 **Autonomous Behavior**

21 On entry to this state, the module updates the Data Path state register (see Table 8-83) for all lanes in the Data
22 Path according to the rules described in section 6.3.3.3.

23 Tx media lane output quiescence for Data Paths in DPTxTurnOff is eventually determined per media lane by
24 the programmed setting in the OutputDisableTx and OutputSquelchForceTx controls.

25 *Note: An externally caused Tx LOS condition of the lane also mutes the Tx output.*

26 **Reactive Behavior (on Host Actions)**

27 *Note: It is recommended that a host minimizes register accesses while in this state.*

28 **Exit**

29 The Data Path state advances to DPInitialized after the OutputDisableTx or OutputSquelchForceTx configuration
30 actions (causing entry to the DPTxTurnOff state) have been realized and the lanes are quiescent and stable.

6.3.4 Flagging Conformance Rules per State

Some Flags (causing a host Interrupt unless masked) are raised by virtue of the Module State Machine (MSM) or by virtue of an active Data Path State Machine (DPSM), but the majority of Flags are raised by other sources.

In some states, certain Flags are not applicable and shall not be raised by the module.

The following sections define the conformance rules for all Flags, for each state of MSM and of DPSM.

Note: The presence of a flagging conformance rule for a particular Flag does not imply that that Flag is required. A conformance rule for a Flag is applicable only when the (required or optional) Flag is actually supported.

6.3.4.1 Module-Level Flagging Conformance Rules per Module State

Table 6-20 describes the Flagging conformance rules for all module-level Flags, per MSM state.

Flag Setting Restrictions

While in an MSM state where a Flag is indicated as **N/A** (Not Allowed), the module shall not set that Flag.

All module-level Flags are generally N/A throughout the **Resetting**, **Reset**, and **MgmtInit** MSM states.

Module-level Flagging conformance is independent of Data Path State.

The meaning of some module-level Flags is configuration dependent; Table 6-20 defines the Flag conformance for each configuration option for those Flags.

Note: The host can suppress undesirable Interrupts caused by known Flags by setting the corresponding known Mask bit at any time after the management interface is initialized.

Table 6-20 Module Flag Conformance Rules

Flag / Flag Group	Page	Byte	Resetting Reset MgmtInit	ModuleLowPwr ModuleFault	ModulePwrUp ModulePwrDn ModuleReady
ModuleStateChangedFlag	00h	8	N/A	allowed	allowed
ModuleFirmwareErrorFlag	00h	8	N/A	allowed	allowed
DataPathFirmwareErrorFlag	00h	8	N/A	allowed	allowed
CdbCmdCompleteFlag*	00h	8	N/A	allowed	allowed
TempMon*Flag	00h	9	N/A	allowed	allowed
VccMon*Flag	00h	9	N/A	allowed	allowed
Aux1Mon*Flag (see Table 8-10)	00h	10	N/A	N/A	allowed
Aux2Mon*Flag (see Table 8-10)	00h	10	N/A	N/A	allowed
Aux2Mon*Flag (see Table 8-10)	00h	10	N/A	N/A	allowed
Aux3Mon*Flag (see Table 8-10)	00h	11	N/A	N/A	allowed
Aux3Mon*Flag (see Table 8-10)	00h	11	N/A	allowed	allowed
CustomMon*Flag	00h	11	N/A	1	allowed

Note 1: The CustomMon*Flag of a vendor-defined monitor is allowed in the ModuleLowPwr and ModuleFault states only if the monitor relates to functionality available in Low Power Mode.

Flag Specification Conformance

Setting allowed alarm and warning Flags of module-level monitors including associated Interrupt generation as per specified Flag semantics is only assured in the **ModuleReady** MSM state.

6.3.4.2 Lane-Specific Flagging Conformance per Data Path State

Table 6-21 and Table 6-22 describe the Flagging conformance for all lane-specific Flags, per DPSM state.

Flag Setting Restrictions

While in a DPSM state where a Flag is indicated as **N/A** (not allowed), the module shall not set that Flag.

All lane-specific Flags are N/A throughout the **Reset** and **MgmtInit** MSM states.

For all other MSM states, the DPSM State determines lane-specific Flagging conformance.

1 Note: For Flags allowed in a DPSM state, additional and more specific rules may exist

2 Note: The host can suppress undesirable Interrupts by setting the corresponding Mask bit at any time after the
3 management interface is initialized.

4 Flag Specification Conformance

5 Setting permitted ('allowed') alarm and warning Flags of Data Path related monitors, including associated
6 Interrupt generation, is only assured in the DPInitialized and DPActivated states.

7 **Table 6-21 Lane-Specific Flagging Conformance Rules**

Flag / Flag Group ¹	Page	Byte	DPAcknowledged	DPIInit	DPAcknowledged	DPDeinit	DPInitialized	DPTxTurnOn	DPTxTurnOff	DPActivated
Data Path Related Flags										
DPStateChangedFlag*	11h	134	allowed	N/A	N/A	allowed	N/A	N/A	N/A	allowed
Tx Related Flags										
FailureFlagTx*	11h	135	allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
LOSFlagTx*	11h	136	N/A	N/A	N/A	allowed	allowed	allowed	allowed	allowed
CDRLOLFlagTx*	11h	137	N/A	N/A	N/A	allowed	allowed	allowed	allowed	allowed
AdaptiveInputEqFailFlagTx*	11h	138	N/A	allowed	N/A	allowed	allowed	allowed	allowed	allowed
OpticalPowerHighAlarmFlagTx*	11h	139	allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
OpticalPowerLowAlarmFlagTx*	11h	140	N/A	N/A	N/A	allowed ¹	allowed	allowed	allowed	allowed
OpticalPowerHighWarningFlagTx*	11h	141	allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
OpticalPowerLowWarningFlagTx*	11h	142	N/A	N/A	N/A	allowed ¹	allowed	allowed	allowed	allowed
LaserBiasHighAlarmFlagTx*	11h	143	allowed	allowed	allowed	allowed	allowed ¹	allowed	allowed	allowed
LaserBiasLowAlarmFlagTx*	11h	144	N/A	N/A	N/A	allowed ¹	allowed	allowed	allowed	allowed
LaserBiasHighWarningFlagTx*	11h	145	allowed	allowed	allowed	allowed	allowed ¹	allowed	allowed	allowed
LaserBiasLowWarningFlagTx*	11h	146	N/A	N/A	N/A	allowed ¹	allowed	allowed	allowed	allowed
TargetOutputPowerOORFlagTx*	12h	231- 238	allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
FineTuningOutOfRangeFlagTx*	12h		allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
TuningNotAcceptedFlagTx*	12h		allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
InvalidChannelNumberFlagTx*	12h		allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
WavelengthUnlockedFlagTx*	12h		N/A	allowed	N/A	allowed	allowed	allowed	allowed	allowed
TuningCompleteFlagTx*	12h		N/A	N/A	N/A	allowed	allowed	allowed	allowed	allowed
Rx Related Flags										
LOSFlagRx*	11h	147	allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
CDRLOLFlagRx*	11h	148	N/A	N/A	N/A	allowed	allowed	allowed	allowed	allowed
OpticalPowerHighAlarmFlagRx*	11h	149	allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
OpticalPowerLowAlarmFlagRx*	11h	150	N/A	N/A	N/A	allowed	allowed	allowed	allowed	allowed
OpticalPowerHighWarningFlagRx*	11h	151	allowed	allowed	allowed	allowed	allowed	allowed	allowed	allowed
OpticalPowerLowWarningFlagRx*	11h	152	N/A	N/A	N/A	allowed	allowed	allowed	allowed	allowed
OutputStatusChangedFlagRx* ²	11h	153	N/A	N/A	N/A	allowed	allowed	allowed	allowed	allowed

8 Note 1: For instance, Tx output power and Tx bias Flags are not allowed (N/A) in the DPInitialized state
9 for media lanes where the Tx output is squelched or disabled by the host.

10 Note 2: While the OutputStatusChangedFlagRx is N/A in certain states, the associated OutputStatusRx*
11 status fields are valid in all DPSM states. An OutputStatusChangedFlagTx is deliberately not available.

¹ An asterisk '*' in a name is a wildcard: All Flags matching the name pattern are referred to.

6.3.4.3 VDM Flagging Conformance per State

Table 6-22 describes the Flag conformance for all Flags related to the (optional) Versatile Diagnostic Monitoring Observables, per Data Path State. See section 7.1 and section 8.19 for information on the optional VDM feature.

In Data Path States where a Flag is indicated as 'Not Allowed', the module shall not set the associated Flag bit while the Data Path is in that state.

All VDM Flags are 'Not Allowed' throughout the Reset and MgmtInit module states. For all other module states, implementers should refer to the Data Path State to determine lane-specific Flag conformance.

Note: The host can suppress undesirable Interrupts by setting the corresponding Mask bit at any time after the management interface is initialized.

Table 6-22 VDM Flag Conformance Rules

VDM Observable Type	DPDeactivated	DPIinitialized	DPIInit	DPDeinit	DPTxTurnOn DPTxTurnOff DPActivated
Laser Age	N/A	allowed	allowed	allowed	allowed
TEC Current	allowed	allowed	allowed	allowed	allowed
Laser Frequency Error	N/A	allowed	allowed	allowed	allowed
Laser Temperature	N/A	allowed	allowed	allowed	allowed
SNR Media Input	allowed	allowed	allowed	allowed	allowed
SNR Host Input	N/A	allowed	allowed	N/A	allowed
PAM4 Level Transition Parameter (LTP) Media Input	N/A	allowed	N/A	allowed	allowed
PAM4 Level Transition Parameter (LTP) Host Input	N/A	allowed	allowed	N/A	allowed
Pre-FEC BER Minimum Sample Media Input (DP)	N/A	allowed	N/A	allowed	allowed
Pre-FEC BER Minimum Sample Host Input	N/A	allowed	N/A	N/A	allowed
Pre-FEC BER Maximum Sample Media Input	N/A	allowed	N/A	allowed	allowed
Pre-FEC BER Maximum Sample Host Input	N/A	allowed	N/A	N/A	allowed
Pre-FEC BER Sample Average Media Input	N/A	allowed	N/A	allowed	allowed
Pre-FEC BER Sample Average Host Input	N/A	allowed	N/A	N/A	allowed
Pre-FEC BER Current Sample Value Media Input	N/A	allowed	N/A	allowed	allowed
Pre-FEC BER Current Sample Value Host Input	N/A	allowed	N/A	N/A	allowed
FERC Minimum Sample Media Input	N/A	allowed	N/A	allowed	allowed
FERC Minimum Sample Host Input	N/A	allowed	N/A	N/A	allowed
FERC Maximum Sample Media Input	N/A	allowed	N/A	allowed	allowed
FERC Maximum Host Input	N/A	allowed	N/A	N/A	allowed
FERC Sample Average Media Input	N/A	allowed	N/A	allowed	allowed
FERC Sample Average Host Input	N/A	allowed	N/A	N/A	allowed
FERC Current Sample Value Media Input	N/A	allowed	N/A	allowed	allowed
FERC Current Sample Value Host Input	N/A	allowed	N/A	N/A	allowed
FERC Total Accumulated Media Input	N/A	allowed	N/A	allowed	allowed
FERC Total Accumulated Host Input	N/A	allowed	N/A	N/A	allowed

7 Advanced Management Features

This chapter describes advanced management features and capabilities, which all are optional.

A module advertises the actually supported advanced features in the module management Memory Map.

7.1 Versatile Diagnostics Monitoring (VDM)

Versatile diagnostics monitoring (VDM) is an optional feature, extending performance and diagnostics monitoring capabilities both in terms of new functionality (statistics) and in terms of new observables that can possibly be monitored (when supported by a module).

The module advertises general VDM support in Bit 01h:142.6. When supported in general, VDM Page 2Fh provides more details as well as descriptor arrays for the set of VDM observables actually supported.

VDM functionality is represented in Pages 20h-2Fh as described in detail in section 8.19, where also a summary of these VDM Pages is shown in Table 8-167.

7.1.1 Purpose and Background

Advanced modules come with new types of observables. For all observables, regular performance monitoring functionality is desired: real-time **samples** and **threshold crossing notifications** for warning and alarm thresholds.

For instance, for modules with tunable lasers, implementing a Dense Wavelength Division Multiplexing (DWDM) optical interface, laser wavelength (or frequency) monitoring is important because it allows estimating the health of the laser. When direct measurement of the laser frequency error is not available, laser temperature deviation from target is often used instead. In addition, DWDM modules typically use a thermo-electric cooler (TEC) to control the laser temperature, and the current flowing through the TEC is a module health indicator. Threshold crossing warnings or alarms for such observables are early indications of pending module failure.

For modules using PAM4 signaling several additional observables are useful to determine the health of the module and the line signal quality. These include bit error ratio (BER) and frame error ratio or frame error count, signal-to-noise ratio estimation, and a level transition (decision threshold) measurement that characterizes the received PAM eye.

Modules of all kinds may contain forward error correction (FEC), which allows estimating digital transmission quality and link margins: Counters of corrections performed, and counters of uncorrectable frames are used to derive frame, symbol, or bit error counts, and associated error rates (per unit time) or ratios (per received data volume). The short-term fluctuations of the error rates or ratios over time are best observed by **statistics**, keeping track of minimum, maximum, and average values. The statistics interval is under host control.

7.1.2 Technical Overview

VDM first extends the list of observables that can be monitored and alarmed beyond those defined as part of core management functionality in Lower Memory and in Banked Page 11h. The list of additional observables that a module may offer for monitoring is provided in Table 8-170.

VDM also extends the monitoring functionality of modules by adding support for **interval statistics** for basic observables with fast internal sampling rates.

The basic functionality of any VDM monitor is the same as the functionality of the basic monitors:

- Providing a **current value** (a.k.a. **sample**, **real time value**) of the monitored observable
- **Flags** indicating high and low **threshold crossings**, both for warning and alarm threshold levels
- **Masks** to suppress Interrupt generation associated with the threshold crossing Flags
- Threshold crossing level information (advertised Flag setting criteria)

The functional extensions of VDM provide host-controlled **interval statistics** for a subset of observables, i.e.

- Minimum
- Maximum
- Average

*Note: Statistics are currently supported for **error performance observables only**, because a fine sampling time resolution is used, intentionally, to capture fluctuations for these performance metrics. Other fluctuating observables of more general interest are usually sampled or updated at much slower time scales, which enables hosts to build any desired statistics themselves.*

1 Note: It is important to distinguish two types of intervals for observables with statistics support: the (shorter)
2 sampling interval and the (longer) statistics collection interval (sometimes just called a statistics interval).

- 3 • A **sampling interval** (or measurement interval) is the module internal interval used to periodically
4 estimate or measure **one sample** of an observable. The durations of the sampling intervals are not
5 specified (**vendor defined**) but typically short (sub-second) because one is interested in performance
6 relevant fluctuations at a short time scale. The series of **samples** are used to **feed the statistics** for
7 the observable, such as min/max/average values of the observable samples within a statistics interval.
- 8 • A **statistic collection interval** is the duration between two "VDM Freeze" requests issued by the host
9 to get the current statistics results and simultaneously restart statistics collection for the next statistics
10 interval. The typical duration of a statistics collection interval is in the order of one or a few seconds.
11 Note that the statistics interval is entirely **host defined** and may vary intentionally or unintentionally.
12 This is a difference compared to a module generated periodic "PM tick" that is known from other
13 systems. Otherwise the VDM Freeze command acts like a host controlled "PM tick".

14 Note: The module should minimize implementation dependent variations of the sampling interval duration in a
15 best effort manner.

16 Note: Since counters are not specified as saturating, hosts are advised to use a collection interval duration that
17 ensures that counter type monitors do not exceed the maximum value that the register can contain. To avoid
18 confusion when a host exceeds this interval, PM counters should be implemented with saturation. See also
19 section 7.1.5 for a description of error performance statistics and section 8.19 for more detailed description of
20 the VDM Freeze/Unfreeze mechanism.

21 Note: The advertised sampling interval (sample measurement interval) of VDM error performance statistics
22 (FineIntervalLength, 2Fh:129-139) does not specify sample measurement durations, or value update periods,
23 for other observables, no matter if with or without statistics. Notably the sampling interval for error performance
24 observables is a CMIS module characteristic, whereas a "PM interval" for readout is a CMIS host characteristic.

25 In some applications, VDM monitoring is not used for performance supervision but only for debugging and
26 diagnostics. To allow computational power saving during normal operation, the VDM functionality can be
27 disabled in a Power Saving Mode (when advertised as being supported).

28 7.1.3 Technical Details

29 See also section 8.19 for more detailed descriptions

30 VDM functionality is represented in the following Banked Pages:

- 31 • Pages 20h-23h: Descriptors of supported VDM observable (64 per Page).
- 32 • Pages 24h-27h: Real-time samples (current values) of the supported VDM observables (64 per Page).
- 33 • Pages 28h-2Bh: Quad Thresholds used to generate TC Alarms or Warnings (16 Quads per Page).
- 34 • Page 2Ch: Up to 256 Flag registers, 4 Flags per observable (alarm/warning-high/low).
- 35 • Page 2Dh: Up to 256 Mask Bits associated with the Flags.
- 36 • Page 2Fh: Advertisement registers and handshake controls.

37 7.1.3.1 Observables

38 VDM allows a module to offer up to 256 observables for 8 lanes in Bank 0. Up to 3 additional Banks can provide
39 extensions for modules with more than 8 lanes, where each Bank provides observables for 8 lanes.

40 VDM observables may be related to the module as a whole, to a Data Path, or they may be lane specific.

41 VDM observables are identified by a numerical Observable ID.

42 It is expected that not all possible observables will be supported by all modules or on all lanes. To address this
43 situation, modules advertise which observables are being monitored.

44 Note: This allows adding new observables without major specification changes.

45 A module advertises supported observables by a 2-byte descriptor register per observable instance (slots).

46 Unused slots are indicated by a value 00h in both bytes of the slot configuration register.

47 7.1.3.2 Thresholds

48 Up to 64 threshold value sets of four thresholds (quads) each are provided, and each of the 256 observables is
49 associated (by the module) with one of the threshold value sets from a threshold group.

1 *Note: Some performance related observables may not have four thresholds implemented. For example, BER*
2 *does not need low alarm or warning thresholds and SNR does not need high alarm or warning thresholds*
3 *implemented.*

4 To indicate a missing high (low) threshold, the module shall advertise the maximum (minimum) representable
5 threshold value, because such a threshold cannot be crossed outwards.

6 *Note: It is not specified if crossing (preferred) or reaching a threshold causes Flag setting. For full accuracy, it*
7 *is left for vendors to specify in data sheets which observables have less than all four thresholds implemented.*

8 The actual threshold values can depend on the commissioned set of Applications and therefore may change
9 whenever a new Application is commissioned.

10 *Note: The module will update threshold values only for that reason. For Data Path Applications, the thresholds*
11 *will be updated when the relevant Data Path reaches DPInitialized. For Network Path Applications, the Media*
12 *Side thresholds will be updated when the relevant Network Path reaches NPInitialized, and the Host side*
13 *thresholds will be updated when the relevant Host Path reaches DPInitialized.*

14 Note: It is recommended that hosts interested in the thresholds read them after Data Path Initialization and/or
15 Network Path Initialization.

16 The observable descriptor registers also indicate the threshold set used for each observable.

17 *Note: The VDM mechanism assumes that several observables may share the same threshold set, e.g. when the*
18 *same observable is measured on multiple lanes.*

19 **7.1.3.3 Structure**

20 The 256 observable slots are split into 4 groups. Each group consists of one Page of descriptors, one Page of
21 real-time values (samples), one Page of threshold sets and 1/4 Page each of Flags and Masks.

7.1.4 PAM4 Observables

Two VDM observables, Signal to Noise Ratio (SNR) and Level Transition Parameter (LTP), are determined from an electrical PAM4 amplitude histogram, from which the receiver may also determine the best decision thresholds for the slicer.

Note: Both SNR and LTP are measures of signal-to-noise ratio, but LTP is more sensitive to a noise floor.

Figure 7-1 below shows a typical block diagram of the optical ingress path of a module showing the location where SNR and level transition (decision threshold) observables are measured.

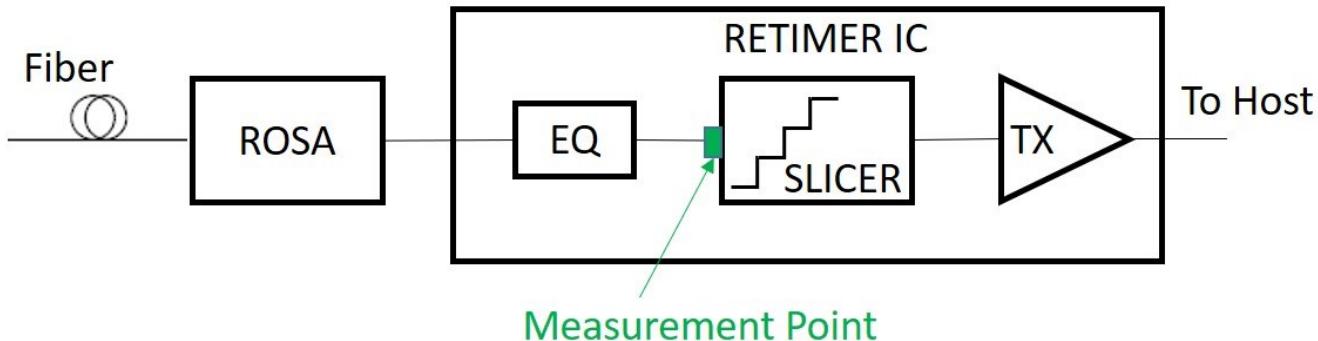


Figure 7-1 Optical ingress path of Module

Figure 7-2 below shows a PAM4 amplitude histogram collected by sampling the received signal in the horizontal center of the eye, just before deciding the data in the slicer. The histogram x-axis is in bins and the y-axis is in number of bin hits. The number of bins and the hit count magnitude is vendor specific.

Note: The same observables can be defined for PAM2, with obvious modifications.

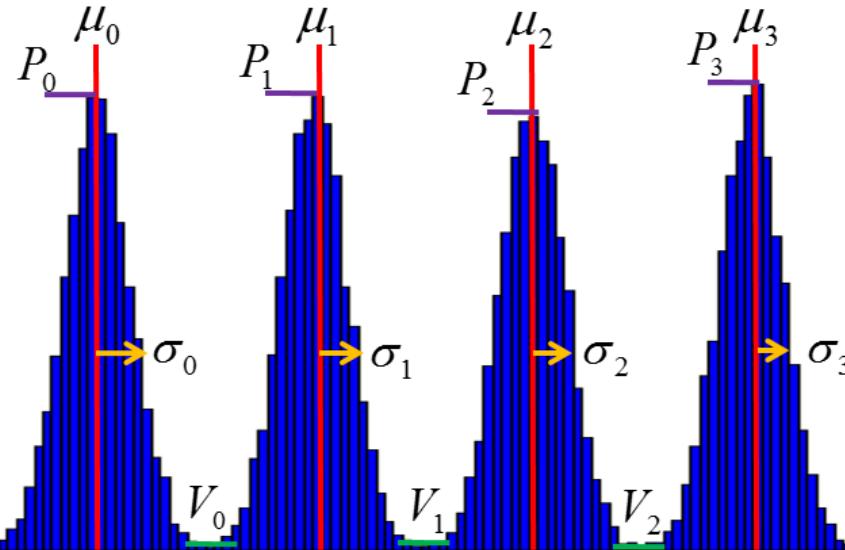


Figure 7-2 PAM4 amplitude histogram

Legend:

- μ_i: location of ith peak, optionally determined as histogram mean over neighboring bins
- σ_i: std deviation of ith peak, estimated from histogram standard deviation over neighboring bins
- P_i: height of ith peak, optionally averaged over neighboring bins
- V_i: height of ith valley, optionally averaged over neighboring bins

The peak locations are the bins with the largest number of hits between any two valleys (or below valley 1 and above valley 3 for the first and last peaks). The valleys are the bins with lowest hit count between two peaks, and the height of a valley is the value of this lowest hit count.

1 Note: The peak locations can be computed more accurately as the histogram mean over neighboring bins,
2 preferably over the same bins used to compute the sample standard deviation.

3 The calculations for the reported PAM4 eye parameters SNR and LTP are

4 **SNR** $\text{SNR} (\text{dB}) := 10 * \log_{10}(\min\{\text{SNR}_0, \text{SNR}_1, \text{SNR}_2\})$

5 where for an **optical** lane $\text{SNR}_i := (\mu_{i+1} - \mu_i) / (\sigma_{i+1} + \sigma_i)$

6 and for an **electrical** lane $\text{SNR}_i := \frac{1}{2} \cdot (\mu_{i+1} - \mu_i)^2 / (\sigma_{i+1}^2 + \sigma_i^2)$

7 **LTP** $\text{LTP} (\text{dB}) := 10 * \log_{10}(\min\{\text{LTP}_0, \text{LTP}_1, \text{LTP}_2\})$

8 where $\text{LTP}_i := (P_{i+1} + P_i) / (2V_i)$

9 Note: In the management register map both $\text{SNR}(\text{dB})$ and $\text{LTP}(\text{dB})$ are expressed numerically in 1/256 dB units.

10 Note: These formulas are also recommended to be used outside of VDM, in module performance diagnostics
11 (see sections 8.12 and 8.13).

12 For the vendor specified wavelength, the accuracy of the reported SNR and LTP observables shall be better
13 than +/-3 dB over specified temperature and voltage.

14 **7.1.4.1 Signal to Noise Ratio (SNR) Observable**

15 For an ingress **optical** lane this observable represents an optical signal-to-noise ratio estimated from electrical
16 amplitude statistics. It is defined as the minimum of the three individual eye SNR_i values, where the SNR_i for
17 each of the three eyes is defined as the ratio of the difference of the mean voltage between neighboring levels
divided by the sum of the standard deviations of the two neighboring levels.

18 For an ingress **electrical** lane this observable represents the electrical signal-to-noise ratio estimated from
19 electrical amplitude statistics. It is defined as the minimum of the three individual eye SNR_i values, where the
20 SNR_i for each of the three eyes is defined as half of the ratio of the squared difference of the mean voltage
between neighboring levels divided by the sum of the variances of the two neighboring levels.

21 SNR is measured in dB and numerically represented as U16 in units of 1/256 dB.

22 Note: For example, a value of 1380h corresponds to an SNR of 19.5 dB.

23 **7.1.4.2 PAM Level Transition Parameter (LTP) Observable**

24 This observable represents electrical noise around slicing signal levels (decision thresholds), which ideally would
25 be located at the signal levels with minimum bin counts (valleys), by considering the ratio between the height
26 of the histogram peaks (at or near the constellation levels) and the height of the valleys between them.

27 The overall LTP is defined as the minimum of three individual LTP_i values, where LTP_i for each PAM level is
28 defined as the average of the peak histogram height of neighboring PAM levels divided by the minimum
29 histogram height between them.

30 LTP is measured in dB and numerically represented as U16 in units of 1/256 dB. When the minimum histogram
31 height between PAM levels is zero the LTP value is encoded as FFFFh. Finite LTP values greater than 255.996
32 dB are encoded as FFFEh.

1 7.1.5 Error Performance Statistics (FEC)

2 7.1.5.1 FEC Related Primary Observables

3 Three detection error metrics may be supported as sampled VDM observables:

- 4 • Frame Errors Counted¹ (**FERC**) in a sampling interval: This observable represents the number of uncorrectable FEC frames that occurred in a sampling interval, measured as RS(544,514) **equivalent** frames. This is a post-FEC decoder error metric, which appears to be an extensive (growing) observable, but implicitly is an intensive (fluctuating) observable, by virtue of the (implicit but fixed) duration of the sampling interval.
- 5 • Pre-FEC Bit Error Ratio (**BER**) in a sampling interval: This is an intensive (fluctuating) observable representing the total number of bits that were corrected by FEC decoding, during a sampling interval, divided by the total number of bits received in that sampling interval. This is an **estimate** of the average pre-FEC bit error ratio or bit error probability.
- 6 • Maximum FEC Symbol Error Weight² (**SEW_{max}**) in a sampling interval: This observable represents the highest number of correctable FEC symbols in a FEC frame that have been detected by a FEC decoder during a sample measurement interval. This is a fluctuating pre-FEC performance metric.

16 **Frame errors** are reported in units of RS(544,514)³ **equivalent frames** per sampling interval: When the FEC actually used is different, then the measured frame error count is scaled proportional to the binary FEC frame size divided by 5440.

19 *Note: For example, if the FEC frame size is 10% larger than the RS(544,514) FEC frame, then the reported equivalent frame error count will be 10% higher than the actually measured true frame error count.*

21 *Note: This is to allow comparing frame error counts regardless of the FEC encoding actually employed.*

22 For algebraic decoders the symbol error weight (**SEW**) of a received FEC frame is an integer number between 0 and **t**, where **t** denotes the symbol error correction capability⁴ of the FEC decoder. The SEW value of **t+1** is used for uncorrectable FEC frames, while an SEW value of 0 implies that the FEC frame appeared error free.

25 *Note: Unlike for frame error counting, the SEW_{max} can not be scaled to an equivalent frame size of 5440 bits.*

26 For other FEC schemes, e.g. with iterative binary decoding, the definition of SEW remains valid (as the maximum number of bit errors corrected within some reference frame) and useful as a relative performance measure.

28 *Note: Different FEC schemes have different maximum pre-FEC BER requirements for achieving the same maximum post-FEC BER target.*

30 *Note: The **vendor defined** error performance **sampling interval** ideally has a constant duration (e.g. a constant number of bits received in the sampling interval), and this **nominally constant** duration is advertised in a VDM field (**FineIntervalLength**). However, the **actual durations** of the sampling intervals will vary slightly for implementation reasons, such as e.g. indeterministic task scheduling jitter. The allowable **variability** of the actual sampling interval duration is **not specified**.*

35 *Note: For the sample statistics to be most meaningful and most useful as a performance indicator, module implementations **should** ensure that the **relative** variation of the sampling interval duration and its deviation from the nominal sampling interval length remains **as small as reasonable possible**. Implementations with larger sampling interval variation call for longer sampling interval durations.*

39 7.1.5.2 FEC Related Statistics

40 Current Sample

41 If supported, the host can always read a most recent BER, FERC, or SEW_{max} sample ("real-time-value" or "current value") measured in the last completed sampling interval.

43 Sample Statistics (Minimum, Maximum, Average, Total)

44 A module supporting statistics takes short-term measurements (**samples**) of BER, FERC, or SEW_{max} over a vendor-specific fine measurement time interval (e.g. 1 ms or 10 ms) and then updates internal sample statistics

¹ Also known as uncorrectable blocks (UCB) or uncorrectable words (UCW).

² Also known as maximum FEC bin or as maximum FEC symbol errors, colloquially alluding to the fact that the symbol errors per FEC frame are often collected into a histogram.

³ see IEEE 802.3 Clause 91.5, a.k.a. KP-FEC

⁴ For the two most used Reed-Solomon FEC schemes, the value of **t** is 7 (KR-FEC) and 15 (KP-FEC), respectively, while a FEC symbol consists of 10 bits.

variables (min, max, and average or more likely: cumulative sum plus sample count). For FERC, the cumulative sum of all FERC samples in the statistics collection period is also computed.

Note: The sample average for BER is best computed from internally accumulated bit error counts and total number of bits received, because this method is mathematically equivalent to averaging the BER samples, but more robust against sampling interval length variations.

Note: The sample average for FERC per sampling interval is best computed from internally accumulated frame error counts and total number of sampling intervals processed, because this method is mathematically equivalent to averaging the FERC samples, but more robust against sampling interval length variations.

Statistics Collection Behavior

The cumulative statistics in a current **statistics collection interval** are updated with each new sample until the host eventually terminates the current statistics collection interval by requesting to freeze the statistics reporting registers and to internally restart the statistics collection for a new interval at the same time.

While no freeze is active, the host may read out the "live" statistics at a slower pace at any time during the currently running statistics collection interval.

Note: While not frozen, the FERC total statistics is an extensive (growing) observable.

This mechanism of closing a statistics interval and reading stable results is described in section 8.19.

Example

Figure 7-3 BER short term measurements and interval statistics (example)

Figure 7-3 shows an *example* of a time series of fine interval measurements (i.e. of samples determined in periodic sampling intervals) punctuated by statistics Freeze events that demarcate the host statistics monitoring interval (sample statistics collection interval). While no Freeze request is pending, after an Unfreeze request, the host can still read calculated statistics values, which then cover the current monitoring period between the most recent Freeze request and the latest completed sampling interval.

Note: Recall that the set of statistics available depends on observable type.

The module periodically calculates a BER/FERC sample for each fine interval (sampling interval), indicated with light borders. The host-determined monitoring interval (statistics collection interval) is indicated with dark borders. The top row in the diagram shows the samples measured internally in each fine interval (sampling interval). The bottom four rows show the statistics reported by the module on request.

The Min, Max, and Avg rows show that after an Unfreeze request, the module continuously updates the statistics reporting values. Then, upon the next Freeze request, the statistics reporting values stop changing until the following Unfreeze request, whereupon it resumes reporting the updating statistics of the new yet unfinished host monitoring interval. Current sample values are always updated, independent of statistics freezing.

Note: Threshold crossing detection is available also for the BER and FERC observables. Modules should use low thresholds of 0. To alarm post-FEC errors the high threshold for FERC should also be 0 unless other error correcting schemes are present.

7.1.6 Tunable Laser Monitors

7.1.6.1 TEC Current

If supported, this monitor observes the amount of current flowing to the TC of a cooled laser.

A single observable for a whole-module TC is available in the Auxiliary monitoring capabilities. Modules should use the VDM system to provide for a per-lane TEC controller.

See Table 8-170 for encoding, units, and scaling information for the monitored observable.

7.1.6.2 Laser Frequency Deviation

If supported, this monitor observes the difference (in frequency units) between the target center frequency and the actual current center frequency. It is a similar measurement to the Laser Temperature except expressed as a frequency difference instead of a temperature difference, and vendors may support one or the other measurement, or both.

See Table 8-170 for encoding, units, and scaling information

7.1.6.3 Laser Temperature

If supported, this monitor observes the laser temperature difference between the target laser temperature for a cooled laser, and the actual current temperature.

It is a similar measurement to the Frequency Error except expressed as a temperature difference instead of a frequency difference, and vendors may support one or the other measurement, or both.

A single observable for a whole-module laser frequency is available in the Auxiliary monitoring capabilities. Modules should use the VDM system to provide for a per-lane laser temperature.

See Table 8-170 for encoding, units, and scaling information for the monitored observable.

7.1.7 Monitors for Co-Packaged Optics (CPO)

In many ways, Co-Packaged Optics (CPO) appears to a host as a multi-lane transmission module with parallel applications over many lanes.

Note: The host of a CPO subsystem is assumed to have prior knowledge about the presence of a CPO subsystem and of the CPO architecture. Therefore, there is no specific advertisement provided.

However, a CPO module comes with a small set of new monitored observables that are not present in a normal transmission module: additional voltages and optical input power levels delivered by external lasers (if any).

These additional observables are implemented as VDM observables. See section 8.22.1.2.

7.2 Command Data Block (CDB) Message Communication

7.2.1 Feature Overview

The optional Command Data Block (CDB) feature allows host-originated **functional interactions** between host and module beyond the basic CMIS core interactions, which allow only to READ or WRITE registers.

CDB is best understood as an asymmetric **message exchange mechanism** for implementing a **command and reply** type of interaction, not much different from invoking a function and waiting for a result being returned. There are **message sending** facilities for the host (initiator) and **message reply** facilities for the module (target) which include status information and completion notification using a Flag.

Two command **processing variants** exist. In both cases the host may send a command and then go off doing other things, coming back later when the reply has returned a result, as determined either by polling or by receiving an Interrupt. With **foreground processing**, the module rejects any register ACCESS during command processing, while with **background processing** register ACCESS is possible.

7.2.2 CDB Technical Basics

The module advertises support for CDB in several fields 01h:163-166 as described below and in section 8.4.11.

The host sends a CDB Command (**CMD**) message which is identified by a **CMD ID** and the module responds with a CDB Reply (**REPLY**) message without changing the CMD ID.

The CMD and REPLY message exchange is performed by reading and writing **message header** fields located on **CDB Message Page** 9Fh (section 8.23) and **message body** data (also called **payload**) either locally on Page 9Fh or on the **Extended Payload Pages** A0h-Afh (section 8.24).

A **CDB Status Register** (00h:37 or 00h:38) informs the host about the status of a CDB facility and the status of the most recent message exchange, as described in section 8.2.8.

A **CDB Command Completion Flag** (00h:8.6 or 00h:8.7) informs the host when the REPLY is available, as described in section 8.2.4.

The **CDB Message Page** 9Fh contains the CMD **message header** of the command to be executed and provides an on-page area for a **message body** of up to 120 bytes, called the **local payload (LPL)** area.

The **CDB EPL Pages** A0h-Afh provide an area of up to 2048 bytes of **extended payload (EPL)** for messages with a large message body. Support for EPL Pages is advertised as described in section 8.4.11.

The same LPL and EPL areas are used for both CMD parameters and REPLY data, depending on command specifications in chapter 9. While the message REPLY does not overwrite the CMD header, the REPLY message body may overwrite the CMD message body. This implies that at most one message exchange is performed at any one time in a CDB instance (i.e. queuing of messages is not supported).

7.2.3 CDB Message Sending and Reply Receiving (Host)

The host triggers sending a message either when writing to a particular Byte address (9Fh:129) or when a multi-byte write to Page 9Fh that includes Byte 9Fh:129 finishes with the STOP condition (see Table 8-54). The method to be used is advertised by the module.

When the module performs CDB command processing in the background, the host can poll the status of CDB command execution in byte **CdbStatus1** (00h:37) or **CdbStatus2** (00h:38). See section 8.2.8 for details.

When the processing of a CDB command has been completed, the module notifies the host using a **CdbCommandComplete** Flag (with associated maskable interrupt) in Byte 00h:8 (see Table 8-9).

A module not responding to a supported CDB command is considered defective and the host may need to initiate a hardware Reset or issue CMD 0004h (Abort Processing) to recover the module.

7.2.4 CDB Message Receiving and Reply Sending (Module)

If a module receives a CDB command, it may perform the following steps, typically:

1. Capture CMD header and body (into a buffer or queue)
2. Recompute the integrity check code (CdbChkCode) for verification
3. Validate the CMD ID and the command parameters
4. Check if the requested command can be executed in the current state
5. Schedule the command for execution and update the command status register
6. Execute the command to completion
7. Determine if the command was executed successfully or if it has failed

- 1 8. Compose the REPLY header and body from execution results
- 2 9. Update the command status register with command result codes (indicating success or failure)
- 3 10. Send the REPLY by setting the command completion Flag, which causes the Interrupt signal to be asserted (unless the command completion Mask is set).

6 The steps described above may vary by module vendor but the overall behavior of the CDB message exchange
7 is identical: a CDB command sent will result in a CDB reply received.

8 7.2.5 CDB Support Levels and Messaging Details

9 Modules may provide different levels of CDB support, e.g. based on processing capabilities.

10 Support level options include:

- 11 • **CDB foreground** processing (ACCESS hold-off until command completion)
- 12 • **CDB background** processing (ACCESS during command processing)
- 13 • Number of CDB instances supported for parallel CDB command execution in background mode

15 *Note: Modules supporting CDB background processing typically require a more powerful CPU and more memory
16 resources, while modules supporting CDB foreground processing may be more cost or power effective.*

17 This CMIS revision specifies support of up to two CDB instances (with foreground or background processing).

18 Each CDB instance uses its own Bank. CDB instance 1 resides in Bank 0 and CDB instance 2 resides in Bank 1.

19 All supported CDB instances support the same set of advertised CDB commands.

20 When background processing of two CDB instances is supported, up to two CDB commands may be executed
21 in parallel, while register access to the memory mapped registers is also possible at the same time.

22 When foreground processing is used, even when two CDB instances are supported, at most one CDB command
23 may be executed at any one time and ACCESS to registers is impossible until the CDB command completes its
24 execution.

25 On completion of a command the module updates the relevant CdbStatus register (00h:37 or 00h:38 for Bank
26 0 or 1) and then sets the relevant command completion Flag in Byte 00h:8:7-6.

27 *Note: Setting a Flag causes an Interrupt unless the Flag is masked by its associated Mask in Byte 00h:31:7-6.*

28 Command completion with status update and Flag setting occurs in both background and foreground mode.

29 7.2.5.1 Foreground Mode CDB Messaging

30 A module may support CDB command execution in foreground mode only. This is indicated by Bit 01h:163.5
31 cleared (see Table 8-54).

32 In foreground mode the module **rejects** any register ACCESS until a currently executing CDB command
33 execution has completed.

34 *Note: READs of the CdbStatus registers 00h:37 or 00h:38 (see Table 8-13) will also be rejected by the module.*

35 Once a command has eventually completed execution, the module sets the relevant CdbStatus register 00h:37
36 or 00h:38 (for Bank 0 and 1), then sets the relevant command completion Flag 00h:8:7 or 00h:8:6, respectively,
37 and resumes accepting register ACCESS by the host (e.g. to read the reply).

38 *Note: While a foreground mode CDB command execution is in progress, a host may use either register ACCESS
39 with retry, or TEST access, a.k.a. Acknowledge Polling (see section 5.2.2.3).*

40 *Note: In the error event that the module does not stop to NACK, the host might need to eventually reset the
41 module via the hardware Reset signal, or toggle the power supply to the module.*

42 7.2.5.2 Background Mode CDB Messaging

43 One CDB Instances

44 A module advertising [01h:163.7-5] = 011b implements one CDB instance in background mode.

45 In background mode the module will generally allow register ACCESS while CDB commands are being
46 processed: it will block ACCESS only for a short time until the command has been captured. Register access is
47 possible as soon as the command is captured.

48 When only one CDB is supported, the module responds only to CDB commands in Bank 0. Reading the CDB
49 status Bit 00h:37.7 will return 1b (CdbIsBusy=1).

1 The host shall not use the same CDB instance, until the current CDB command execution is completed or
2 aborted.

3 *Note: If the host attempts to send another CDB command while the CDB instance is busy (indicated by the*
4 *CdbIsBusy status bit) the command may be ignored. Under race conditions, the command may also be*
5 *processed if CdbIsBusy is cleared in background during the MCI transaction time. This makes the behavior of*
6 *the CDB command or response unpredictable.*

7 The only CDB command that is allowed to be sent while a CDB instance is busy (CdbIsBusy), is CMD 0004h
8 (CDB Abort Processing), which is designed to abort a command that has not completed within expected
9 command execution time and to clear the CdbIsBusy status.

10 The module indicates command completion by setting Flag 00h:8.6 (CdbCmdCompleteFlag1). An Interrupt
11 request will be asserted, unless masked by the associated Mask bit 00h:31.6 (M-Cdb1CommandComplete).

12 *Note: Modules with one CDB instance and background processing capability may e.g. support firmware update*
13 *using CDB while the module is in service. Depending on host software implementation, the module may also*
14 *support running Performance Monitors (PM) at the same time.*

15 Two CDB Instances

16 A module may advertise with [01h:163.7-5] = 101b that it supports two CDB instances in background mode.

17 A module supporting two CDB instances in background mode can run two CDB commands simultaneously.

18 Otherwise, the rules described in the previous section apply to each CDB instance, individually.

19 *Note: This type of module allows the host to have separate threads or tasks for dedicated purposes; for*
20 *example, having one thread specifically for firmware update and another thread for regular performance*
21 *monitoring, diagnostics, and other features.*

22 *Note: When using multiple threads, the host must implement proper interlocks for Bank and Page registers and*
23 *MCI writes/reads. This is one of the main reasons that the status registers 37 and 38 have been allocated in*
24 *Lower Memory where these registers are always available.*

7.3 CDB Message Communication Based Features

7.3.1 Firmware Management using CDB

This section describes concepts and procedures related to firmware management, especially firmware update, using the set of CDB commands defined in section 9.7.

Firmware update based on CDB messaging is an advertised optional feature.

7.3.1.1 Firmware Management Concepts

The **generic** CDB based firmware update facilities allow hosts to update the module firmware completely or partially in a **vendor supported** depth and granularity.

Firmware

Implementation dependent examples of firmware items that might be updated include main controller firmware, the parts needed to support another Memory Map revision, modified administrative or operational data resources, microcode in the module, code for internal devices such as converters or DSPs, or a bundle of such items, in any granularity as defined and supported by the module vendor.

The definition of what updateable items exist and if they may be included in a complete or partial firmware update is module implementation dependent and outside the scope of the specification. This is possible because the CDB based firmware update mechanism is not dependent on the content of a firmware update.

Note: Host suppliers are encouraged to coordinate with module suppliers to ensure firmware update packages meet their requirements.

Firmware Revisions

In the CMIS core, firmware revisions are identified by a **major** version number and a **minor** version number. With CDB firmware management, a finer distinction using a **build number** and an extra **string** is possible.

Note: As the aggregate firmware of a module may consist of multiple components, it is strongly recommended that the CMIS firmware revision numbers refer to that complete aggregate (bundle) of all firmware components, not to a particular firmware component.

Note: It is also strongly recommended that revision numbers, at least when complemented with build number and extra string, uniquely specify exactly one aggregate firmware configuration.

Note: The recommendation for firmware revisions to refer to the entire aggregate module firmware incurs the recommendation to deliver a single firmware download file for a firmware revision.

Note: When piecewise or incremental firmware updates are used or supported by a vendor, it is strongly recommended that transiently invalid or unsupported combinations of firmware components in the module are considered invalid and not declared runnable until eventually a supported and well defined aggregate firmware configuration status is achieved, which is identified by a unique aggregate version identification.

Firmware Update

Firmware update includes operations on module-internal firmware stores (called firmware banks), and the function of passing control to firmware in a firmware store.

- Operations on internal firmware stores include **query**, **download**, **upload**, and **copy** of stored firmware (to, or from, or between firmware stores), which can serve the purposes of **upgrade**, **repair**, and **downgrade**.
- The operation to eventually pass control to downloaded firmware is called **activation** and occurs on demand either immediately (**run**) or on the next restart (**commit**)

Note: If a module contains firmware but does not support firmware update, the module may still allow to query firmware information using CMD 0100h (Get Firmware Info). This is especially recommended when vendor specific firmware update mechanisms exist.

Firmware Storage and Administration

Modules may have up to two host updateable images stored in up to two storage banks, called A and B; and modules may also have an internal factory image.

Each Bank (or its content) is classified by its current content as follows:

- The **Administrative Status** indicates if the firmware in the bank is run upon resets (**committed**) or not (**uncommitted**). Only one bank is committed at a time. The administrative status can be changed

1 with CMD 010Ah (Commit Firmware Image).

- 2 • The **Operational Status** of a bank indicates if the image in the bank is currently **running** or **not running**. Only one image can be running at a time. In special conditions, the module may indicate that it is running a factory image which is neither A nor B. The Operational Status may be changed with CMD 0109h (Run Firmware Image), or with a module reset (triggered by software or hardware signal), or with power cycling. Upon a reset or power cycle the module will run the "committed" image as indicated by the Administrative Status.
- 3 • The **Bank Validity Status** indicates if the firmware in the bank is runnable (**valid**) or not (**invalid**). A valid firmware is runnable and persistently stored completely undamaged. The Bank validity of images A and B may be affected by the CDB firmware **update** commands with IDs 0101h (Start), 0102h (Abort), 0103h and 0104h (Write), 0107h (Complete) and 0108h (Copy). Only a proper sequence will be able to restore the Bank Validity Status from "Invalid" to "Valid".

4
5
6
7
8 *Note: Firmware Update using a single storage bank is possible as well. In this case, there is no option to fall back to a previous firmware during trial phase for a downloaded but not yet committed firmware image.*

9 **Firmware Update Media**

10 For any specific firmware update purpose, the module vendor can provide a specific **firmware download file** which includes vendor specific instructions for updating firmware in a module, completely or partially.

11 A firmware download file consists of a vendor specific and host transparent **header** section and a **firmware** section containing a representation of the actual executable firmware image bundle. This representation may be coded (ASCII) or uncoded (binary).

12 *Note: Supplying a single firmware download file for the complete module firmware bundle is recommended to ease generic firmware administration for hosts.*

13 If firmware readback is supported, the module vendor may also provide a **firmware readback file** which serves to control reading back the most recently downloaded firmware from the module. A delivered firmware readback file consists of a header and of dummy data with the expected length of the firmware to be read back.

14 *Note: The purpose of firmware readback is vendor or host specific. CMIS makes no assumptions here.*

15 The vendor specific **header** section in both types of firmware files is initially downloaded to and interpreted by the module. Since size, content, and interpretation of these headers are vendor specific, the module only advertises the number of header bytes at the beginning of a firmware file in the field **StartCmdPayloadSize** (9Fh:138) returned by CMD 0041h (Firmware Management Features).

16 The file format of a firmware file may be either binary or ASCII (for example Intel hex file format). If the file is ASCII encoded, the host first decodes the file contents into binary format before downloading or after uploading.

17 *Note: The method for a host to determine the firmware file format are not specified in CMIS.*

7.3.1.2 Reference Firmware Download Procedure using CDB

The host begins by reading module capabilities using CMD 0041h (Firmware Management Features).

The host then reads (and possibly converts to binary format) the vendor specific and vendor provided firmware download file into a contiguous addressable byte array defined as the binary **download image**.

To start the firmware download, the host sends the download image header consisting of the first **StartCmdPayloadSize** bytes from the download image using CMD 0101h (Start Firmware Download). This header instructs the module in a vendor specific way about the full or partial download content to be expected.

Before the module updates an image bank in a download procedure the module ensures that the bank is marked as empty or corrupt until the download has eventually finished successfully.

When CMD 0101h was successful, the module is ready to accept data from the host using the advertised method, either CMD 0103h (Write Firmware Block LPL) or CMD 0104h (Write Firmware Block EPL). These command differ only in the allowed size of a download image block, which is advertised for the EPL variant.

The host zero-initializes a variable containing the address of the next data block to be sent or received in subsequent Write Firmware Block or Read Firmware Block commands, respectively.

In a loop the host then reads subsequent bytes of the download image in blocks not exceeding the allowed block size and sends it using CMD 0103h or 0104h.

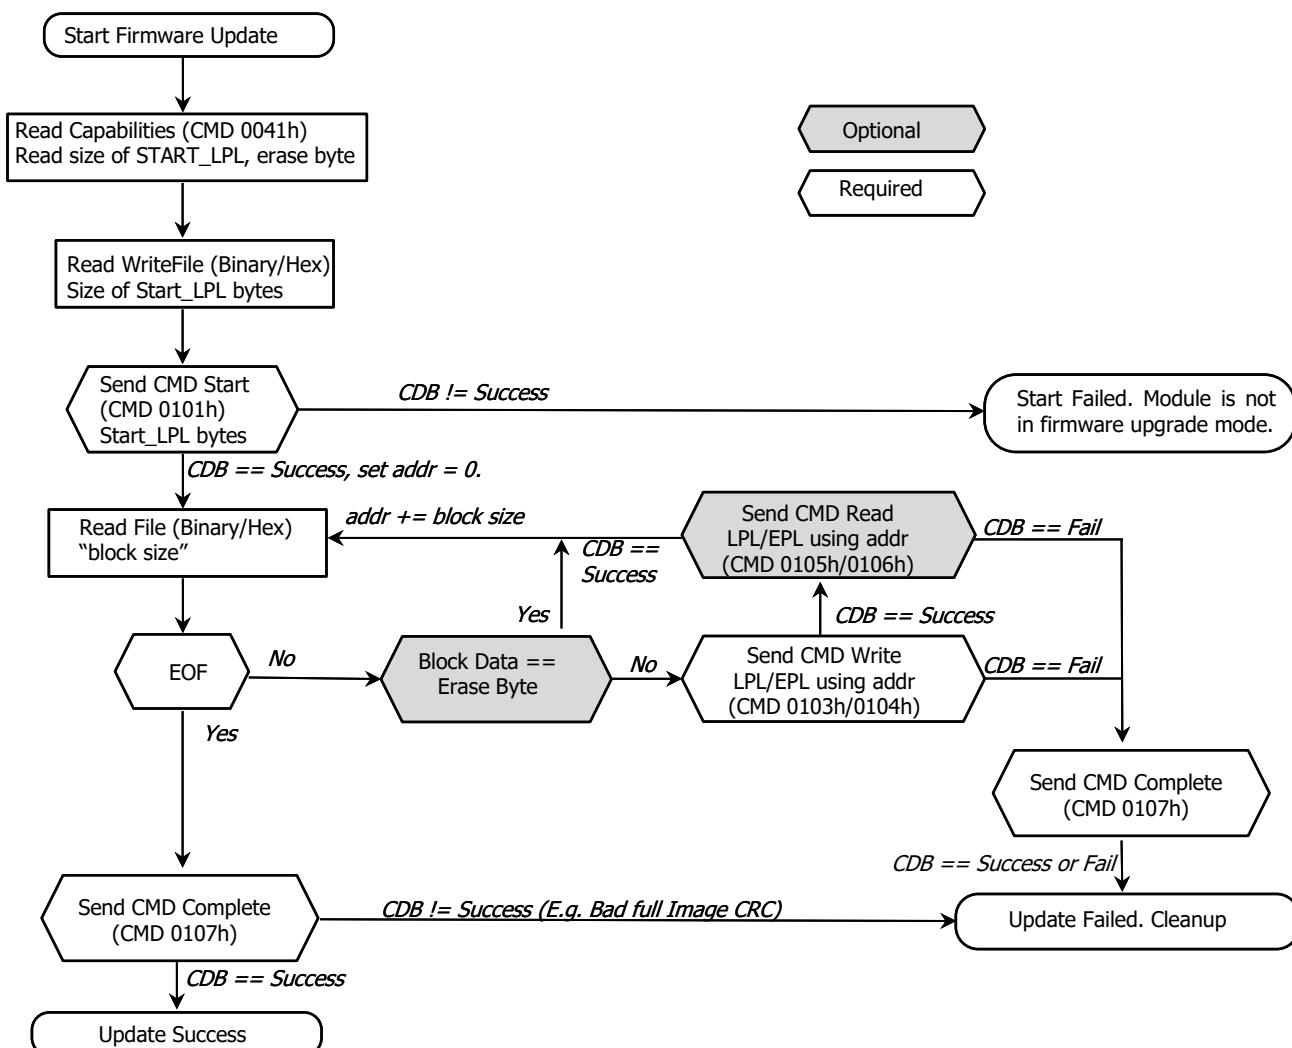


Figure 7-4 Firmware Update using CDB

The flow diagram in Figure 7-4 shows two optional blocks:

The first optional block describes that the host may optionally skip sending such blocks of data that consist entirely of bytes marking an erased state (this special byte value is advertised). Skipping the download of such blocks can speed up the firmware download process.

1 The second optional block describes that the host may optionally read back the most recently sent block. This
2 slows down the firmware download process and is only recommended when needed.

3 A successful firmware update process ends when the host has downloaded all data from the download image
4 and no unrecovered errors have occurred in the associated message exchanges. The host then completes the
5 firmware update by sending the CMD 0107h (Complete Firmware Download). Receiving this command will
6 trigger the module to perform any appropriate validation and stop the current firmware download process. If
7 the validation fails, this command returns a failure code in the **CdbStatus** register.

8 *Note: A defensive host should not assume that all modules will perform a full validation.*

9 When the firmware download was unsuccessful, e.g. due to commands returning failure, or when it needs to
10 be aborted for any reason, the host must also terminate the ongoing firmware update process either by sending
11 a CMD 0107h (returning failure but terminating the download process) or, if supported, by using CMD 0102h
12 (Abort Firmware Download). A new firmware update process may then be started to retry updating the firmware
13 in the module.

15 **7.3.1.3 Reference Firmware Upload Procedure using CDB**

16 Reading back a complete or partial image, as controlled by the header in the readback firmware file, may be
17 achieved using the flow chart shown in Figure 7-5.

18 *Note: The upload retrieves data from the most recently updated image store.*

19 *Note: Modules may not support reading back firmware from the module. If readback is supported, a host may
20 generate a file from the readback data such that the data of the file matches that of the download file used. At
21 least one bit will be different as the headers of the readback and download files are different.*

22 Readback is controlled by the header portion of a vendor-supplied **firmware readback file**, which serves two
23 purposes: it contains a vendor specific header that controls which parts of the overall firmware of the module
24 will be uploaded and its size defines the size of the expected readback data.

25 The readback file header has the same number of (StartCmdPayloadSize) bytes like the downloaded file. Within
26 the vendor-specific header bytes, a field must indicate to the module that the operation is a readback, in which
27 case CMD 0101h (Start Firmware Download) does not perform an image erasure.

28 To start the readback, the host reads the first number of StartCmdPayloadSize bytes from the binary readback
29 image and sends these bytes using CMD 0101h (Start Download). The content of the readback image after the
30 first "Start Command Payload Size" bytes is irrelevant and will not be sent to the module.

31 The host zero-initializes a variable containing the address of the next data block to be received in subsequent
32 Write Firmware Block or Read Firmware Block commands, respectively. The number of bytes to readback as
33 shown in Figure 7-5 is the size of the binary readback image (excluding the header).

34 When CMD 0101h was successful, the module is in a state ready to return data to the host using the advertised
35 method, CMD 0105h (Read Firmware Block LPL) or CMD 0106h (Read Firmware Block EPL). These commands
36 differ mainly in the allowed size of an upload image block, which is advertised for the EPL variant

37 After completion of the readback, but also on readback errors, the host sends CMD 0107h (Complete Firmware
38 Download). In the reply, the module may return either nothing or an appropriate error code in subsequent Read
39 LPL or Read EPL commands.

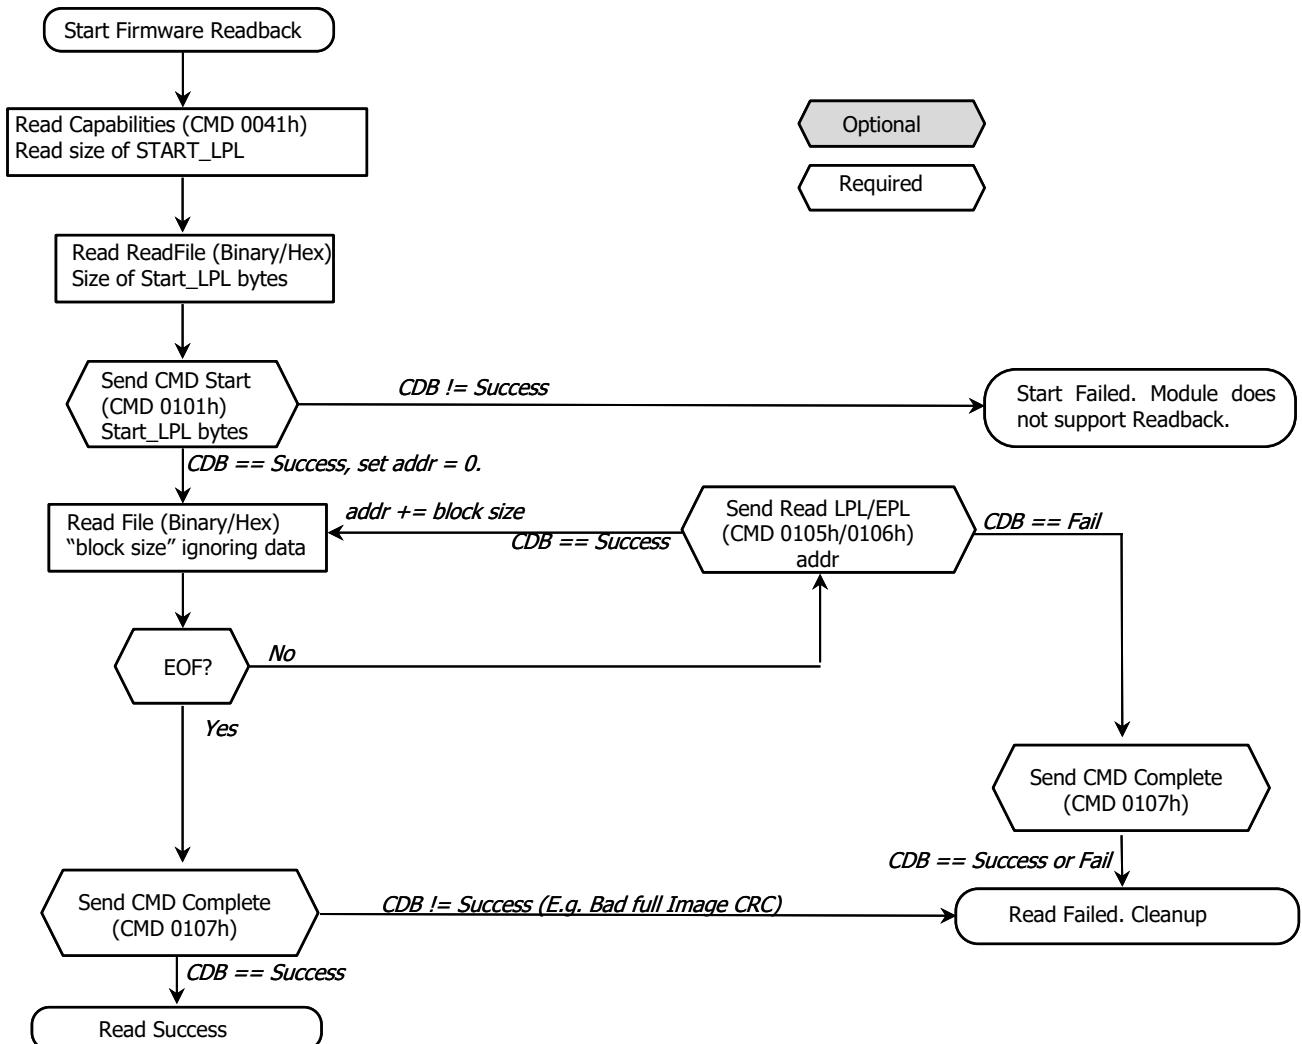


Figure 7-5 Firmware Upload using CDB

1

2

7.3.1.4 Firmware Administration and Status

This section describes the **FirmwareStatus** Byte 9Fh:136 returned by CMD 0100h (Get Firmware Info). This byte represents status information about the firmware images stored in the module. See section 9.7.1.

The Status of (the content of) an image storage bank is encoded in the bits of a nibble (half byte)

Table 7-1 Status of an Image Storage Bank

Nibble Bit	Status Aspect	Encoding
0	Operational Status	1: running 0: not running
1	Administrative Status	1: committed 0: uncommitted
2	Validity Status	1: invalid (not runnable) 0: valid (runnable)
3	-	0: reserved

Table 7-2 and Table 7-4 below show a non-exhaustive list of common FirmwareStatus codes. Table 7-2 shows typical codes returned by modules that supports two image banks A and B. The presence of a bootloader or Factory image is not important nor directly indicated here. In normal operating situations, there will not be a way for the bootloader to continue running.

Table 7-2 Typical FirmwareStatus Codes of Modules supporting image A and B.

Firmware Status BA	Firmware Running	Committed Firmware	Bank Validity Status	Description
03h	A	A	A,B	Normal Operation. A running and A committed. Image B valid, can be activated by Run B command.
43h	A	A	A,-	Normal Operation. A running and A committed. Image B update is in progress, has failed or has been aborted. There is no valid image in Bank B.
12h	B	A	A,B	Trial run of image in Bank B.
16h	B	A	-,B	Errored Mode (Bank A found invalid). Image in Bank A is committed but invalid. The module is running the valid image in Bank B.
30h	B	B	A,B	Normal Operation. B running and B committed. Image A valid, can be activated by Run A command.
34h	B	B	-,B	Normal Operation. B running and B committed. Image A update in progress, failed or aborted. There is no valid image in Bank A.
21h	A	B	A,B	Trial run of image in Bank A.
61h	A	B	A,-	Errored Mode (Bank B found invalid). Image in Bank B is committed but invalid. The module is running the valid image in Bank A.

Table 7-3 Typical FirmwareStatus Codes of Modules supporting image A

Firmware Status BA	Running	Committed	Bank Validity Status	Description
44h	-	-	-,-	Factory image or boot loader running, no valid image available
40h	-	-	A,-	Factory image or boot loader running, valid uncommitted image available
42h	-	A	A,-	Factory image or boot loader running, valid committed image available
41h	A	-	A,-	Running an uncommitted image
43h	A	A	A,-	Running a committed image

Table 7-4 shows FirmwareStatus codes that may be returned by a module in special situations, when neither image A nor B is running and neither image A nor B is committed. However, the firmware is still responding correctly to this CDB command. In this situation the host should interpret the running firmware as an internal factory image that cannot be updated via the CDB command set. The vendor may have other vendor specific methods to update the firmware and should be contacted for additional details if these codes are returned.

In some situations, these may be error conditions, for example, if the module supports image A and B as well as a backup factory image, these are error codes that may be returned by the module if an unexpected event has occurred that caused the factory image to be running.

Table 7-4 Special FirmwareStatus Codes

Firmware Status BA	Running	Committed	Bank Validity Status	Description
00h	Factory (none of A,B)	Factory (none of A,B)	A,B	This may be returned by a module that does not support firmware update and only has the factory image. A module that supports image A and B may return these codes when the module's boot control block that defines which image Bank to boot is damaged or when a default factory image is running.
40h				A,-
04h				-,B
44h				-,-

7.3.2 Performance Monitoring (PM) using CDB

The optional CDB implementation of performance monitoring supports **statistics** such as minimum, average, maximum and current values for a subset of **monitored observables** (see Table 9-31 in section 9.8).

A host can query the list of supported CDB performance monitoring features using CMD 0042h and CMD 0201h (see sections 9.4.3 and 9.8.2).

CDB based performance monitoring is based on host polling and allows the host to read many observables in a single CDB message transaction.

Alarm indications are not supported by the polling based CDB mechanism. If autonomous alarm reporting via Flags is required, Versatile Diagnostics Monitoring (VDM) implemented in Pages 20h-2Fh may be used.

CDB and VDM may both be supported and then used simultaneously. Both CDB and VDM can collect statistics (such as minimum, average, and maximum) for an overlapping subset of observables.

The CDB and VDM statistics may be collected either independently (**independent statistics mode**) or synchronously (**linked statistics mode**). The statistics collection mode defaults to independent mode.

A CDB command (CMD 0200h) is available for selecting the statistics collection mode, e.g. to switch from independent tracking to synchronous tracking.

Whenever the module restarts (including restart due to a Run Image CDB command), the performance monitoring mode reverts to independent tracking of all supported and advertised observables.

All statistics results are cleared and reset with a mode change.

In **independent statistics mode** statistics read using CDB and statistics read using VDM Pages 20h-2Fh will generally differ because start time and duration of the data collection periods are not synchronized.

Clearing and resetting of statistics result registers in CDB and VDM Page 2Fh will be completely independent, and the module maintains two separate internal data sets for statistics collecting (minimum, average and maximum values).

In **linked statistics mode**, the module maintains one common minimum, average, and maximum value for each observable that supports statistics and then generally returns the same information via CDB or VDM.

Furthermore, in a linked mode a statistics interval result will be cleared whether it is cleared by a CDB command or by writing to a register in VDM Page 2Fh.

For both performance monitoring methods, VDM or CDB, the time at which the module originally starts collecting statistics and the method or internal time granularity at which these data are processed remains undefined.

It is assumed that a host will initially setup Data Paths and later, when the operational conditions are satisfactory for the data of interest, read and discard the first set of statistics results, before actually (re-)starting statistics collection. This is recommended because the first statistics results will likely be abnormal due to initialization transients in the system or non-operational conditions.

7.3.3 Security Features and Capabilities using CDB

CDB currently supports one optional security feature: Module **Device Authentication**.

Module authentication is based on a cryptographically protected Initial Device Identity (IDevID) certificate which is securely installed during module manufacturing. Pluggable modules manufactured with a secure identity can provide assurance of module integrity and authenticity to the host, based on standard public key cryptography.

Secure, trustworthy modules are built upon hardware-based roots of trust. Without these roots of trust there can be no assurance of integrity or authenticity. The root of trust provides secure key storage and uses a private key for digital signatures of host message digests in a challenge-response authentication mechanism. A module certificate is accessible to the host to obtain the corresponding public key.

An IDevID certificate contains the module's unique serial number (SN), product ID (PID), public key, and a signature generated by a certificate authority (CA). Both SN and PID values should be covered by the signature.

7.4 Module Boot Record (MBR)

This section is reserved for a future feature extension.

7.5 Tunable Lasers, Channel Selection and Center Frequency Setting

7.5.1 Concepts and Background

An optical carrier or optical channel is usually characterized by its optical center frequency and the optical channel bandwidth available to carry a modulated signal.

Note: Instead of specifying the optical frequency, a (signed) channel number on a given channel grid is often used, representing the grid offset against a channel number 0, which is defined as the grid's anchor or origin. Instead of specifying the channel bandwidth, the frequency grid spacing between optical carriers on a grid is often used. See below for more information.

For management purposes an optical channel is represented in CMIS by a Media Lane together with associated Media Lane attributes (registers). CWDM module transmit over multiple channels and therefore have multiple Media Lanes. DWDM modules often use only a single optical channel and have only a single Media Lane.

From a management viewpoint, optical channels may be either **pre-defined** by an interface standard (often in single-channel and CWDM applications) or **host-programmable** (often in DWDM applications).

For **host-programmable** channels, conceptually there are two provisioning schemes:

- indirectly **selecting** a channel from a predefined set of channels (in **grid-based** applications), where a channel and its associated center frequency is selected by a channel number
- directly **setting** a center frequency arbitrarily within an overall supported band (in **grid-less** applications), i.e. without the help of a channel number

Note: In both cases, the module must ensure, for operation, that the relevant lasers are tuned to the center frequencies of the provisioned channels. When a tunable module wants to implement an Application with pre-defined channels, it will perform the tuning based on the selected Application, not based on the channel programming registers.

In CMIS, a particular **channel grid** is defined by four attributes

- a grid spacing bandwidth (the frequency difference between neighbors)
- a channel numbering scheme (mapping channel numbers to carrier frequencies)
- a **numbering origin** or **grid anchor** frequency (193.1THz)
- a **range of channel numbers** supported on the grid

Note: The CMIS anchor frequency is not necessarily at the center of the supported band.

7.5.2 Programming Overview

A module advertises the **grid types** it supports on Page 04h. The grid type definitions include grid spacing, channel numbering scheme including numbering origin, and the supported range of channel numbers.

The registers for programming a single optical channel (for one Media Lane) are located on Page 12h:

- For **grid-based** applications, the grid to be used and a desired channel number on that grid are set.
- For **grid-less** applications, for historical reasons¹ the center frequency is not programmed directly, but in a base-plus-offset approach: The host selects and tunes to a nearby channel from a coarse-grained grid and then add a suitable fine-grained frequency offset (called a fine tuning offset) to achieve any desired off-grid center frequency.

To perform this function, so-called fine-tuning support (adding a small frequency offset to a grid frequency) must be available and enabled. *For instance, selecting a grid with 50 GHz channel spacing (or less) and offset tuning with 1MHz resolution over a +/-32 GHz range allows to program arbitrary grid-less channels. Note that tuning is performed as an initial setup step.*

When selecting another optical channel (grid spacing or channel number), the module must be in the **DPDeactivated** state. Attempts to change the optical channel (grid selection or channel selection) while the corresponding Data Path is in any other Data Path state may result in unspecified behavior, as the module may tune the laser before or after the change has been made.

A **Tuning Not Accepted** Flag is raised when the module is currently unable to serve the tuning request following a host programming grid or channel. The Flag indicates that the nominally programmed channel or grid values do no longer reflect the actual laser condition.

¹ The absence of a direct frequency setting method has historical reasons: grid-less applications are a relatively recent development and modules with tunable lasers have traditionally been programmed via channel numbers. Future versions of CMIS may include a direct frequency setting method.

1 Unlike channel selection, laser frequency fine-tuning offsets, and target output power settings may be
2 programmed if the corresponding Data Path is not in a transient state.

3 The module is expected to align the laser to the fine-tuning offset value regardless of whether the value has
4 been changed before, during or after laser tuning.

5 Modules that support programmable output power are also expected to align the output power to the target
6 output power value regardless of when the change was made.

7.6 Client Encapsulation Applications (Multiplexing)

This section specifies management functionality and essential behavior of **client encapsulation** applications (a.k.a. **multiplex** applications) in the form of differences and extensions to the baseline of **system interface** applications described in chapter 6.

Note: System interface and client encapsulation modules are introduced in chapter 1, while Chapter 6 describes the concepts of Applications and of Application Instances allocated to Data Paths (groups of lanes) in modules. However, the Data Paths described in chapter 6 were limited to system interface applications with tightly coupled host and media side initialization and operation. For client encapsulation applications, new concepts will be introduced here, to allow loose coupling between host side and media side initialization and operation.

7.6.1 Concepts and Technical Background

To ease conceptual understanding, this subsection introduces client encapsulation a.k.a. multiplex applications in general terms. The specific technical language representing the same subject matter in CMIS data structures is described in the following subsections.

7.6.1.1 Multiplex or Uniplex Client Encapsulation Applications

In client encapsulation applications, the data stream carried by one or more host side **host signals** (sometimes also referred to as client signals or tributaries) is digitally mapped into (and demapped from) an independently maintained media side **network signal**, which carries the data of the host signals as encapsulated payload.

Usually the data rate of the network signal is slightly higher than the combined rate of the host signals, and the signal mapping often includes some form of frequency adaptation (stuffing) in order to allow mutually asynchronous host signals and network signals.

Both the host signals and the network signals can be single-lane or multi-lane signals, independently.

Note: As an important and initially surprising restriction, the host lane data rates of all host signals that are encapsulated into a (i.e. into any) network signal must be identical! See subsection 7.6.1.5.

In a (genuine) **multiplex** application, several ($N > 1$) lower bandwidth host signals fill the higher available bandwidth of a network signal in a particular order, also called the (N:1) multiplex structure.

In a **uniplex** application, a single host signal fills the available bandwidth of a network signal; a uniplex application can also be viewed as a degenerate or trivial case ($N = 1$) of (N:1) multiplexing.

Note: Unlike for a system interface application, the media side of a uniplex application can be initialized and operated also when the host signal is absent.

A module may offer a set of supported **multiplex structures** (supported sequences of host signals with possibly different bandwidths, multiplexed in a particular order into supported network signals).

A multiplex structure can be viewed as an ordered list of module-internal **multiplex connection points**, to which host signals of **suitable type** and **bandwidth** can be connected (within the module).

Once a specific client encapsulation application (and its multiplex structure) is provisioned, the network signal may be commissioned (initialized, and remain operational) totally independent of the presence of any of the multiplexed host signals.

The multiplex structure can be changed without loss of continuity both for the network signal and for those host signals that are not affected by the multiplex structure change.

Note: A typical use case would be to take tributary signals out of service, or to bring a new tributary signal into service without disturbing the existing tributary transmission services.

7.6.1.2 Provisioned Host Replacement Signal Generation

Modules may optionally include facilities to transmit an internally generated host signal (actually generating a replacement data stream) instead of forwarding a received host signal or data stream.

Note: Unlike the provisioning of the mechanisms causing the insertion of replacement signals (if supported), the overall scenarios for using host replacement signals are not in the scope of CMIS.

Note: Examples of standardized host replacement signals include SONET unequipped signaling, or insertion of Ethernet Idle block streams (without LF or RF failure indications)

*Note: Host configured replacement signals should not be confused with the data and signals that a module reactively transmits instead of a failed host signal, as an automatic **consequential action** in response to certain*

failure conditions, as governed by the network signal processing specifications that is associated with the MediaInterfaceID. Consequent actions are **automatic module behaviors** defined for an application, and as such not in the scope of CMIS. Examples of consequent actions include AIS maintenance signal insertion in OTN applications, or LF or RF insertion in Ethernet applications.

Provisioned Tx Host Signal Replacement

Client encapsulation applications may optionally include facilities to insert (map) a host replacement signal or data stream in **Tx** direction that can be enabled by the host, typically when (intentionally) no host signal is connected to feed a certain multiplex connection point.

Note: A configurable host replacement signal is especially useful when it actually conveys the unambiguous meaning that a multiplex connection point and its associated payload capacity is intentionally 'unused' or 'unequipped' rather than absent due to an undesired host signal fail situation.

Data content and digital format of the replacement signal in the Tx direction depend on the transmission standard underlying the advertised MediaInterfaceID, and often are such that the received signal is recognizable as a client replacement data at the remote end, in Rx direction.

Provisioned Rx Host Signal Replacement

Client encapsulation applications may optionally include facilities to generate and transmit a host replacement signal in **Rx** direction, which might be used for fully independent link bring-up on host side and media side.

Data content and physical format of the replacement signal in the Rx direction depends on the transmission standard underlying the advertised HostInterfaceID.

Note: The OutputStatusRx indicators of the relevant lanes are set when an Rx host path replacement signal is transmitted, because the Rx host path replacement signal is a valid physical layer signal.

7.6.1.3 Automatic Squelching

Automatic squelching of the network signal outputs in **Tx** direction based on host side signal defects is not meaningful for multiplex applications. It is therefore not supported for client encapsulation applications.

Note: Support of forced squelching on host demand is maintained for client encapsulation applications.

Automatic squelching host signal outputs in **Rx** direction is assumed to be supported both for physical network signal fails (LOS), as well as for digital received host signal fails (Rx Host Path input fails), whereby the conditions for presence and meaning of digital received host signal fails are derived from the pertinent interface standards.

Note: Rx auto-squelching is a mandatory module function. It is recommended that modules allow it to be disabled via the AutoSquelchDisableRx register (as advertised in the AutoSquelchDisableRxSupported register)

Transmitting a provisioned Rx host replacement signal has priority over automatic Rx squelching.

7.6.1.4 Tributary Assignment Flexibility Restrictions

Previous versions of the specification did not provide internal routing flexibility for host signals (tributaries) from host lane ports to module internal multiplex connection points of a multiplex application. Instead, the assignment of host signals to multiplex connection points was based on the sequence of lane groups on the module's host interface: the host signal had to appear on the host lanes in the same order and with the same signal types as the multiplex connection points in the multiplex structure of the Application.

Note: Near end and far end systems therefore had to support the same multiplex structure not only from a network signal processing viewpoint, but also in relation to the host signal perspective.

In this version of the specification a module may support host lane switching as described in section 7.8 to allow the host to redirect host lanes to different internal connection points in a multiplex application.

7.6.1.5 Host Lane Granularity Restrictions

This version of the specification requires that all lanes of all (single or multilane) host side signals that are encapsulated into any of one or more network signals must **all** carry the same data rate per lane, i.e. within a mixed multiplex application and even across any parallel multiplex applications.

Note: This restriction is established only to limit the complexity of advertising capabilities and restrictions for parallel multiplex applications and for mixed multiplex applications (see section 8.15.5.4 for details). Note that no such restriction exists for parallel system interface applications as described in chapter 6.

7.6.1.6 Clocking Aspects

This version of the specification does not include clocking or clock dependency aspects of multiplex schemes.

Note: Especially, no assumptions are made on mutually synchronous or asynchronous clocking of the host signals nor on the network signal being asynchronous to the host signals.

7.6.2 Data Path, Network Path, and Host Paths

The **Data Path** of a client encapsulation application, like the Data Path of a system interface application, consists of all involved host and media lanes, including all related module internal resources.

However, unlike the Data Path of system interface applications described in chapter 6, where the Data Path is also a unit of independent initialization, usage, and deinitialization, the Data Path of a N:1 multiplexing client encapsulation application is partitioned into N host side Host Paths (see below) and a single media side Network Path (see below), with finer granularity of initialization, usage, and deinitialization governed by a set of N+1 parallel state machines (one per Host Path and one for the Network Path), rather than by a single DPSM.

Network Path

The media side network signal related transmit-and-receive functions of a **multiplex** or **uniplex** client encapsulation application are realized by a **Network Path (NP)** that spans the segment between the internal multiplex connection points to the media lane physical interfaces.

Like the media side of a system interface Data Path, an NP can use one or more media lanes, employing single-carrier or multi-carrier transmission.

A Network Path is controlled using a **separate** set of Network Path related **registers** and its dynamic state of configuration is represented by a Network Path State Machine (**NPSM**)

Host Path

The host side host signal related transmit-and-receive functions of one host signal in a client encapsulation application are realized by a **Host Path (HP)** that spans the segment between internal multiplex connection points and the host lane physical interfaces.

Each **Host Path** is controlled using the host side related **Data Path registers** for its constituent lanes and its dynamic state of configuration is represented by a DPSM that is effectively restricted to act on the relevant Host Path segment of the Data Path only (by limiting the achievable states to DPDeactivated and DPInitialized).

Note: In a certain sense, all literal, textual references to host side aspects of Data Paths in chapter 6, actually refer only to a Host Path segment of the Data Path of an NP Application, when the relevant lane is part of a NP Application and hence a HP lane. Beware of confusion. See also section 7.6.6.

Where distinction is needed, the combination of one HP and its associated NP is called a **partial** Data Path, because the host lanes of the other HPs in the NP Application are not included.

7.6.3 Network Path Applications

Due to the characteristic presence of an independently operated Network Path, client encapsulation applications are also referred to as **NP Applications**. Likewise, for ease of distinction, system interface applications (without client encapsulation) are now also referred to as **DP Applications**.

A **uniplex** NP Application includes only one Host Path.

A **multiplex** NP Application includes as many Host Paths as there are multiplexed host signals.

A **homogeneous** (simple) NP Application encapsulates one type of host signals (characterized by a single HostInterfaceID).

A **mixed** (heterogeneous) NP Application encapsulates different types of host signals (each type characterized by a different HostInterfaceID).

Note: Recall that this specification restricts the lane data rate to be identical for all Host Path lanes, not for essential reasons, but to simplify the advertisement of mixed multiplex capabilities and restrictions.

One NP Application comprises one media side Network Path that is connected (internally, in the module) to one or more host side Host Paths (HPs) such that the information bandwidths of the Host Paths add up to the information bandwidth of the Network Path.

One NP Application instance is therefore defined as a specific association of one or more Host Interfaces comprising a group of host lanes (each associated with a separate Host Path) and a single Media Interface (associated with a single Network Path) comprising a group of media lanes.

A module may support several **alternative** NP applications, such as e.g. rate scaled multiplex applications, which are structurally identical but use a scaled host signal data rate.

A module may also support several NP Applications operating in **parallel** where each NP Application requires only a subset of the lanes available in the module.

7.6.4 Network Path Application Advertisement

As with Data Path Applications, Application Descriptors (see section 0) are used to describe all possible instantiations of an NP Application, expressed from a host side perspective in terms of host lanes.

An Application Descriptor where the information bandwidth associated with the MediaInterfaceID is a multiple of the information bandwidth associated with the HostInterfaceID is a **partial** Application Descriptor.

Using an application descriptor extension, the module explicitly advertises for each Application Descriptor, if it advertises a DP Application or an NP Application (see section 8.15.5.3).

Note: The partial Application Descriptors required to advertise genuine multiplex applications were semantically invalid in CMIS 5.0. CMIS 5.0 hosts may therefore be confused by CMIS 5.1 modules, if they sanity test for equal data rate on host and media side, or if they do not fully match desired versus advertised Applications.

Uniplex NP Applications

Structurally, the Application Descriptor for a uniplex NP Application cannot be distinguished from a DP Application Descriptor; since the meaning associated with the MediaInterfaceID does not always allow the distinction, application advertisement extensions are used to disambiguate.

Homogeneous Multiplexing NP Applications

The Application Descriptor for a homogeneous multiplexing NP Application can be recognized by the fact that the information bandwidth of the host side signal (identified by HostInterfaceID) is lower than the information bandwidth of the media side signal (identified by MediaInterfaceID), i.e. it is a partial Application Descriptor

Note: Normally, the MediaInterfaceID information bandwidth is an integer multiple of the HostInterfaceID information bandwidth. When this is not the case, the host side signal cannot be used in a homogeneous multiplexing NP Application but possibly only in a mixed multiplexing NP Application.

Mixed Multiplexing NP Applications

A mixed multiplexing NP Application is described by a **set** of mutually **consistent** partial Application Descriptors, one for each type of host interface being part of the multiplex structure.

Two partial Application Descriptors are **consistent** when they use the same MediaInterfaceID and when they offer at least one identical option in the MediaLaneAssignmentOptions, as well as at least one set of conflict-free HostLaneAssignmentOptions.

The module advertises separately (see section 8.15.5.5) which lane groups and bandwidth granularities can be mixed to form the Host Paths of the mixed multiplexing NP Application, subject to restrictions described in section 7.6.1.5.

7.6.5 Network Path Application Instance Configuration

Application Descriptors and associated extensions (see section 8.15.5.3) advertise all potential instantiations of DP or NP Applications that are supported by a module supporting NP Applications.

The configuration of a particular DP Application instance is described in sections 6.2.1 through 6.2.4, whereas the configuration and control of a particular NP Application instance is described in the following.

7.6.5.1 Configuration and Control

The actual configuration of an NP Application Instance in terms of its host lanes, media lanes, and internal resources is determined by first **provisioning** (into Active Control Sets) and then **commissioning** (into hardware) a configuration previously **defined** in Staged Control Sets.

This procedure allocates the host lanes (of one or more Host Paths) and the media lanes (of a Network Path) to the NP Application Instance and internally connects the HP instances and the multiplex connection points of the NP, within the module.

1 The host lanes are assigned to Host Paths by using the regular Data Path related Staged Control Set registers
2 (see section 6.2.3), and these host lanes are also assigned to the Network Path using a dedicated NP Staged
3 Control Set (see section 8.15.2), i.e. the Host Paths and the Network Paths are associated via host lanes.

4 Like Data Path instances, also Host Path and Network Path instances are identified by reference to the lowest
5 lane number of all assigned host lanes, called HP DPID and NPID, respectively.

6 The outputs (inputs) of a HP become associated to NP inputs (outputs) at internal multiplex connection points,
7 in order of increasing host lane numbers (HP DPIDs), when the lanes of the HP (in the provisioned DPConfig)
8 are also assigned to the NP (in the provisioned NPConfig).

9 In the overall configuration procedure, the group of Host Paths (HP) and the Network Path (NP) participating
10 in an N:1 NP Application instance are defined, provisioned, and commissioned sequentially (i.e. in separate
11 configuration commands), and the dynamically controlled states of all commissioned HPs (as represented in
12 DSPM states) and of the NP (as represented in an NPSM state) are also observed and controlled separately.

13 The **definition, provisioning, commissioning, and control** of NP Application instances works as follows:

14 The HPs are **defined**, individually or as a group, in a Staged Control Set. Each HP is identified by reference to
15 its first associated host lane number, called HP DPID, and all lanes of a HP must be assigned this HP DPID.

16 The NP is **defined** in an NP Staged Control Set. The NP is identified by reference to its first associated host
17 lane number (of all tributary HPs), called NPID, and all lanes of NP must be assigned this NPID.

18 HPs and NP are then (requested to be) **provisioned** into the relevant Active Control Set or NP Active Control
19 Set in separate steps, by writing to the associated ApplyDPInit trigger register or ApplyNPInit register,
20 respectively. As usual, the provisioning procedure includes (partial) validation prior to copying a Staged Control
21 Set into the associated Active Control Set.

22 *Note: Since the order of HP and NP provisioning is unspecified, and since the time of completing the provisioning
23 is unknown, the module can only validate certain aspects during each of the provisioning commands.*

24 *Note: A future version might add a command to validate the currently provisioned configuration to allow on
25 demand validation of a completely provisioned NP Application prior to commissioning at a time when the host
26 knows that its provisioning is complete.*

27 The sequential **commissioning** of provisioned HPs and provisioned NP is **controlled** via the **DPDeinit** and
28 **NPDeinit** registers, causing DPInitPending and NPInitPending indicators to be set, respectively.

29 *Note: Prior to the commissioning a configured NP Application Instance a suitable multiplex structure must be
30 fully provisioned into the DP and NP Active Control Sets. However, some of the multiplex connection points may
31 actually remain unused (i.e. the relevant host signals of the HP connected to the multiplex connection point are
32 intentionally not present) by keeping the relevant DPSMs in the DPDeactivated state.*

33 The actually commissioned **configuration** of HPs and NP can eventually be **observed** in the Active Control
34 Set and in the NP Active Control Set after the DPInitPending and NPInitPending indicators are cleared, while
35 current **states** can be **observed** in the DPState and NPState registers, respectively.

36 *Note: See also Appendix H for configuration and control examples.*

37 7.6.5.2 Reconfiguration

38 The multiplex structure of a NP can be changed by selective HP reconfiguration and, optionally, NP source
39 reconfiguration, without disrupting the transmission of unchanged Host Paths carrying traffic.

40 *Note: This can be achieved because the NPSM and the HP DPSMs are decoupled.*

1 7.6.6 Modifications and Extensions of Prior Specifications for Data Path Applications

2 The textual specifications for DP Applications in chapter 6 are generally applicable to the Host Paths of NP
3 Applications, with the following modifications assumed but not spelled out in the text.

4 Register specifications in chapter 7.8 referring to the host side aspects of Data Paths apply also to Host Paths.

5 7.6.6.1 Host Lanes

6 All statements about the host side aspects of Data Paths apply also to Host Paths.

7 All register specifications referring to the host side aspects of Data Paths apply also to Host Paths.

8 7.6.6.2 Data Path State Machine (DPSM)

9 Each Host Path is controlled by a slightly modified Data Path State Machine (DPSM) in which the transition
10 signal **DPDeactivateS** is **forced to TRUE**, limiting the steady states achievable to DPDeactivated and
11 DPInitialized.

12 *Note: This reflects the fact that the media side functions of a NP Application are represented by the Network
13 Path, not by the Host Path, and therefore controlled by the NPSM.*

14 Formally this is achieved by adding a configuration dependent term **NPIInUseT** to the DPDeactivateS equation

$$15 \quad \text{DPDeactivateS} = \text{DPReDeinitS} \text{ OR } \text{DPTxDisableT} \text{ OR } \text{DPTxForceSquelchT} \text{ OR } \text{NPIInUseT} \text{ (Eq. 7-1)}$$

16 where:

17 **NPIInUseT** is defined to be TRUE for a Host Path lane and FALSE otherwise.

18 Formally, the presence of a Host Path (instead of a Data Path) is derived from **NPIInUseLane< i >**, where **< i >** is
19 the ID of the Data Path or Host Path (see Table 8-133).

20 The media side Network Path is not controlled by a DPSM, but instead by a Network Path State Machine (NPSM)
21 as described in section 7.6.7.

22 7.6.6.3 Media Lanes

23 All statements made about Media Lanes of Data Paths need to be reinterpreted as statements about Media
24 Lanes of Network Paths.

25 Unlike for DP Applications, in NP Applications the conditions for Tx output signals to be valid do not depend on
26 host signal conditions in the module.

27 7.6.6.4 Squelching

28 Squelching related functions for the HPs of a NP application in the **Rx** direction are supported when they are
29 supported for DP applications.

30 In the CMIS core specification only physical signal defects (LOS) trigger automatic squelching in the Rx direction.
31 However, depending on the MediaInterfaceID, Network Paths may come with additional conditions detected in
32 digital signal processing layers that signify the absence of a genuine host signal to be forwarded. Such conditions
33 are not specified here but are assumed to contribute to Automatic Squelching if available, and if enabled.

34 Squelching related functions for the HPs of a NP application in the **Tx** direction are supported when they are
35 supported for DP applications, with the strict exception of auto-squelching.

36 *Note: In an NP application, consequent actions in case of a multiplexed host signal being unavailable are
37 assumed to be specified in the pertinent transmission standards.*

38 7.6.6.5 Configuration

39 There is no intervention-free reconfiguration for Network Paths nor for Host Paths, i.e. only stepped
40 reconfiguration functionality is available for Host Paths of multiplex Data Paths, independent of the
41 SteppedConfigOnly setting for regular Data Paths.

7.6.7 Network Path State Machines (NPSM)

A **Network Path State Machine (NPSM)** instance describes Network Path-specific behaviors and properties that are related to the configuration of the media side **Network Path**, as managed by the host.

*Note: The NPSM state represents a **management** or **configuration realization status** of a Network Path, representing the effects of certain host configuration commands and of module reactions to those commands. It **does not** necessarily represent other behavioral or operational aspects of a Network Path, e.g. in terms of current input or output signal conditions or in terms of transmission service being provided.*

*Note: The NPSM state should neither be confused with the **operational status** of the functional resources of a Network Path (in Tx direction or in Rx direction) nor with the resulting signal **output status** of Rx host lane outputs or of Tx media lane outputs.*

Module State and NPSM Life Cycle

All **NPSM** instances required to represent the power-up default Application defined in the Network Path Configuration field values of the Active Control Set are initially created and set-up during the **MgmtInit** state.

After its creation, a NPSM remains in the NPDeactivated State until the Module State Machine is in the **ModuleReady** state and an exit condition from the NPDeactivated state is met.

When the host updates the Network Path Configuration fields in the Active Control Set, in either the **ModuleLowPwr** or **ModuleReady** states, the module tears down any previous NPSM that is no longer defined and then creates and sets up any newly defined NPSM.

All Network Path State Machines are torn down in the **Resetting** state.

NPSM Purpose

A Network Path State Machine is used by the module to represent the initialization status of the resources associated with a Network Path in response to certain host configuration settings or commands.

Although individual resources within a Network Path may complete initialization activities at different times, the module waits to report the updated NPSM state until all resources associated with the Network Path have completed the requested configuration or reconfiguration action. This synchronized status reporting across all lanes and resources in a Network Path reflects the fact that there is only one NPSM per Network Path.

NPSMs for parallel Network Paths of Multiple Application Set Instances

Each Network Path in a module is required to operate independently of other Network Paths: if the host changes the Network Path State of one Network Path, the other Network Paths in the module shall be unaffected and uninterrupted.

7.6.7.1 NPSM State Transition Diagram and NPSM Specification

Figure 7-6 shows the state transition diagram (STD) of a NPSM representing the Network Path configuration related state of one Network Path instance.

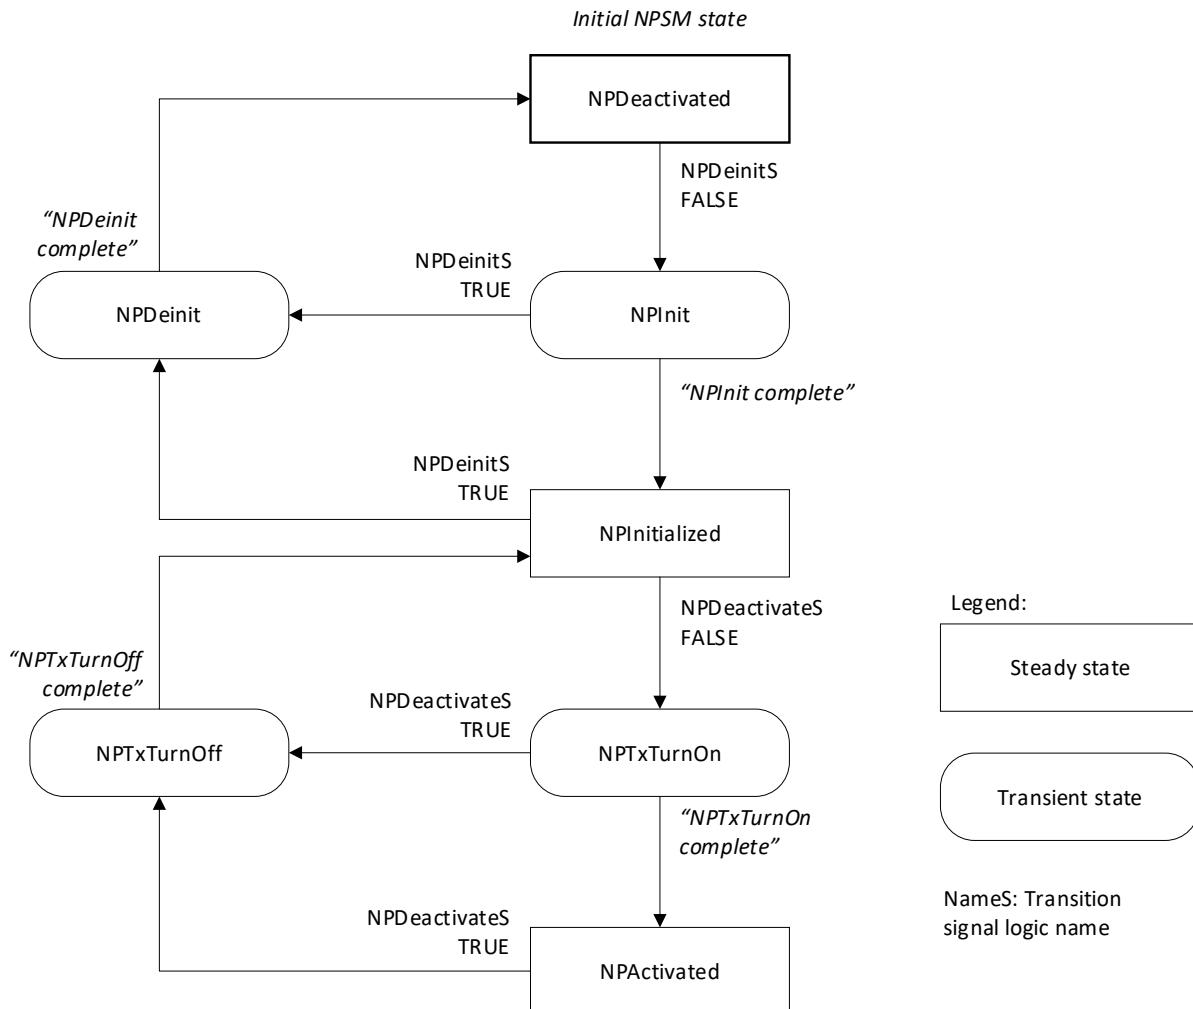


Figure 7-6 Network Path State Machine (NPSM) State Transition Diagram

Note: Prior to exit from the MgmtInit module state, all Network Paths initialize to the NPDeactivated state.

State Exit Conditions and Transition Signals

The state machine exits a given state when specific exit conditions are satisfied. So called **Transition Signals** (recognized by name suffix S) represent these exit conditions for steady states.

NPDeinitS

The NPDeinitS transition signal can also be represented by the logic equation

$$\text{NPDeinitS} = (\text{NOT NPInUseT}) \text{ OR } (\text{NOT ModuleReadyT}) \text{ OR LowPwrS OR NPDeinitT} \quad (\text{Eq. 7-2})$$

where:

$$\text{ModuleReadyT} = (\text{ModuleState} = \text{ModuleReady}) \quad (\text{Eq. 7-3})$$

$$\text{NPDeinitT} = \begin{aligned} & \text{NPDeinitLane}<\text{N}> \\ & \text{OR NPDeinitLane}<\text{N+1}> \end{aligned} \quad (\text{Eq. 7-4})$$

$$\dots$$

$$\text{OR NPDeinitLane}<\text{N+M-1}>$$

$$\text{NPInUseT} = \begin{aligned} & \text{NPInUseLane}<\text{N}> \\ & \text{OR NPInUseLane}<\text{N+1}> \end{aligned} \quad (\text{Eq. 7-5})$$

1 ...

2 OR NPInUseLane<N+M-1>

3 N = first host lane in the Network Path

4 M = number of host lanes in the Network Path

5 **NPDeactivateS**

6 The NPDeactivateS transition signal is defined by the logic equation

$$7 \quad \text{NPDeactivateS} = \text{NPDeinitS} \text{ OR } \text{NPTxDisableT} \text{ OR } \text{NPTxForceSquelchT} \quad (\text{Eq. 7-6})$$

8 where:

$$9 \quad \begin{aligned} \text{NPTxDisableT} = & \text{ OutputDisableTx}<\text{N}> \\ & \text{OR OutputDisableTx}<\text{N+1}> \\ & \dots \\ & \text{OR OutputDisableTx}<\text{N+M-1}> \end{aligned} \quad (\text{Eq. 7-7})$$

$$10 \quad \begin{aligned} \text{NPTxForceSquelchT} = & \text{ OutputSquelchForceTx}<\text{N}> \\ & \text{OR OutputSquelchForceTx}<\text{N+1}> \\ & \dots \\ & \text{OR OutputSquelchForceTx}<\text{N+M-1}> \end{aligned} \quad (\text{Eq. 7-8})$$

11 N = first media lane in the Network Path

12 M = number of media lanes in the Network Path

20 **Reaction to Module Reset**

22 When the MSM **ResetS** transition signal (see Table 6-11) becomes TRUE, any Network Path activities related
23 to power down are performed in the **Resetting** module state. The NPSM state machines then ceases to exist.

24 *Note: Module dependent pre-reset clean up and power down activities may be implemented, possibly depending
25 also on the reset trigger in either hardware or software.*

26 **Reaction to Module Fault**

27 When Module State Machine transitions to **ModuleFault** state, the NPSM behavior is not defined formally, but
28 governed by the behavioral requirements of the ModuleFault state.

29 **7.6.7.2 Network Path Control (Host)**

30 A single main configuration register is provided for the host to control initialization and deinitialization of all
31 Network Paths represented in a given Bank. This **NPDeinit** register (see Table 8-136) defines per host lane if
32 the associated Network Path resources are determined to be unused for functional operation (and hence can
33 be deinitialized) or if they are determined for functional operation (and hence need to be initialized).

34 *Note: Initialization status and behavior of Tx media lane outputs are further controlled by the host using the
35 media lane specific control bits **OutputDisableTx**<*i*> and **OutputSquelchForceTx**<*i*>.*

36 A host requesting initialization or deinitialization of a Network Path ensures that the Active Control set contains
37 the desired configuration settings and then writes the value 0 or 1, respectively, to the DPDeinit bits associated
38 with the host lanes of that Network Path.

39 The host may request initialization or deinitialization of multiple Network Paths with one register access.

40 **7.6.7.3 Network Path Status (Module)**

41 The module provides information on the current state of the Network Path (NPSM **current state reporting**)
42 and on entry to certain NPSM states (NPSM **state change indication**).

43 **NPSM Current State Reporting**

44 On entry to a NPSM state the module reports the NPSM state entered as the current NPSM state in the **NPState**
45 status register (see Table 8-144

46 Table 8-143), on all lanes of the Network Path, with optional exceptions specified below.

47 *Note: Due to identical behavior of all lanes feeding a Network Path the host needs to read only the first lane of
48 the Network Path to determine the Network Path state.*

1 NPSM Current State Reporting Exceptions

2 The module **may** suppress reporting the current NPSM state in the NPStateHostLane<i> registers when that
 3 state is known to be transitional, i.e. when it is exited immediately because its exit conditions are fulfilled on
 4 entry, or when the duration of staying in that state is known to be in the order of 1 ms or less.

5 *Note: The duration specification is intentionally vague. The intention for allowing exceptions in state reporting
 6 is to avoid reporting short-lived status data which the host is unlikely to read and react upon.*

7 NPSM State Change Indication (Flag)

8 A NPSM State Change Indication consist of the module setting a **NPStateChangedFlag** for each lane of the
 9 Network Path associated with the relevant NPSM instance.

10 *Note: The intention of the following specification is that the module indicates a state change only on entry to a
 11 lasting steady state and only when the transition time since the previous lasting steady state was significant.*

12 The maximum duration of a transient state is advertised in a MaxDuration* field (see Table 8-146) and is
 13 considered **insignificant** when the coded MaxDuration* field value is 0000b (see Table 8-48), and **significant**
 14 otherwise.

15 The module issues a NPSM State Change Indication on entry to a **steady** NPSM state when no exit condition
 16 of the entered **steady** state is fulfilled on state entry (state is not visited transitorily)¹.

17 The module does **not** perform a NPSM State Change Indication on entry to a **transient** NPSM state.

18 Table 6-19 defines the Flag behavior for each NPSM state entry.

19 **Table 7-5 Network Path State Changed Flag behaviors**

Entered state	NPStateChangedFlag may be set *
NPInit	No
NPIinitialized	Yes
NPDeinit	No
NPDeinit	No
NPDeactivated	Yes
NPTxTurnOn	No
NPActivated	Yes
NPTxTurnOff	No
NPTxTurnOff	No
NPIinitialized	Yes

20 * Note: The Flag setting conditions are described in the main text.

21 *Note: Steady state exit conditions may already be met upon entry into the steady state and lead to immediate
 22 transition to the next state (after state entry or state exit activities, if defined).*

23 Flag Related Behavior

24 The module does not clear any Flag due to a state change of a Network Path State Machine.

25 The module raises Flags only according the NPSM state-specific conformance rules defined in section 7.6.8.

26 7.6.7.4 Detailed State Descriptions

27 The DPSM state descriptions in the subsections of section 6.3.3 apply similarly to the NPSM states, with the
 28 following exceptions:

29 Intervention-free reconfiguration of the NP is not supported.

30 ¹ Note the deliberate difference to the DPSM behavior specification in section 6.3.3.3.

7.6.8 Network Path Dependent Flagging Conformance

7.6.8.1 Lane-Specific Flagging Conformance per NPSM State

Network Path(s) adhere to the following NPSM flagging conformance rules.

Table 7-6 and Table 7-7 describe the Flagging conformance for media lane-specific Flags, per NPSM state.

Note: The Network Path Flagging Conformance tables are limited to flags or flag groups and VDM observables that are related to the media interface.

Flag Setting Restrictions

While in an NPSM state where a Flag is indicated as **N/A** (not allowed), the module shall not set that Flag.

All media lane-specific Flags are generally N/A throughout the **Reset** and **MgmtInit** MSM states. For all other MSM states, the NPSM State determines media lane-specific Flagging conformance.

Note: For Flags allowed in a NPSM state, additional and more specific rules may exist

Note: The host can suppress undesirable Interrupts by setting the corresponding Mask bit at any time after the management interface is initialized.

Flag Specification Conformance

The setting of allowed alarm and warning Flags of Network Path related monitors including associated Interrupt generation is only assured in the NPInitialized and NPActivated states.

Table 7-6 Lane-Specific Flagging Conformance Rules per NPSM State

Flag / Flag Group ¹	Page	Byte	NPDeactivated	NPInit	NPDeinit	NPInitialized	NPTxTurnOn	NPTxTurnOff	NPActivated
Network Path Related Flags									
NPStateChangedFlag*	17h	128	allowed	N/A	N/A	allowed	N/A	N/A	allowed
Tx Media Related Flags									
OpticalPowerHighAlarmFlagTx*	11h	139	allowed	allowed	allowed	allowed	allowed	allowed	allowed
OpticalPowerLowAlarmFlagTx*	11h	140	N/A	N/A	N/A	allowed	allowed	allowed	allowed
OpticalPowerHighWarningFlagTx*	11h	141	allowed	allowed	allowed	allowed	allowed	allowed	allowed
OpticalPowerLowWarningFlagTx*	11h	142	N/A	N/A	N/A	allowed	allowed	allowed	allowed
LaserBiasHighAlarmFlagTx*	11h	143	allowed	allowed	allowed	allowed	allowed	allowed	allowed
LaserBiasLowAlarmFlagTx*	11h	144	N/A	N/A	N/A	allowed	allowed	allowed	allowed
LaserBiasHighWarningFlagTx*	11h	145	allowed	allowed	allowed	allowed	allowed	allowed	allowed
LaserBiasLowWarningFlagTx*	11h	146	N/A	N/A	N/A	allowed	allowed	allowed	allowed
Rx Media Related Flags									
LOSFlagRx*	11h	147	allowed	allowed	allowed	allowed	allowed	allowed	allowed
CDRLOLFlagRx*	11h	148	N/A	N/A	N/A	allowed	allowed	allowed	allowed
OpticalPowerHighAlarmFlagRx*	11h	149	allowed	allowed	allowed	allowed	allowed	allowed	allowed
OpticalPowerLowAlarmFlagRx*	11h	150	N/A	N/A	N/A	allowed	allowed	allowed	allowed
OpticalPowerHighWarningFlagRx*	11h	151	allowed	allowed	allowed	allowed	allowed	allowed	allowed
OpticalPowerLowWarningFlagRx*	11h	152	N/A	N/A	N/A	allowed	allowed	allowed	allowed

7.6.8.2 VDM Flagging Conformance per NPSM State

Table 7-7 describes the Flag conformance for all Flags related to the (optional) Versatile Diagnostic Monitoring (VDM) Observables, per NPSM State. See section 7.1 and section 8.17 for information on the VDM feature.

In NPSM States where a Flag is indicated as 'Not Allowed', the module shall not set the associated Flag bit while the Network Path is in that state.

All VDM Flags are 'Not Allowed' throughout the Reset and MgmtInit module states. For all other module states where a NPSM is in use, implementers should refer to the Network Path State to determine lane-specific Flag conformance.

¹ An asterisk '*' in a name is a wildcard: All Flags matching the name pattern are referred to.

Table 7-7 VDM Flag Conformance Rules per NPSM State

VDM Observable Type	NPDeactivated	NPIInit	NPDeinit	NPIInitialized	NPTxTurnOn NPTxTurnOff NPActivated
Laser Age	N/A	allowed	allowed	allowed	allowed
TEC Current	allowed	allowed	allowed	allowed	allowed
Laser Frequency Error	N/A	allowed	allowed	allowed	allowed
Laser Temperature	N/A	allowed	allowed	allowed	allowed
eSNR Media Input	allowed	allowed	allowed	allowed	allowed
PAM4 Level Transition Parameter (LTP) Media Input	N/A	allowed	N/A	allowed	allowed
Pre-FEC BER Minimum Media Input (data-path)	N/A	allowed	N/A	allowed	allowed
Pre-FEC BER Maximum Media Input	N/A	allowed	N/A	allowed	allowed
Pre-FEC BER Average Media Input	N/A	allowed	N/A	allowed	allowed
Pre-FEC BER Current Value Media Input	N/A	allowed	N/A	allowed	allowed
FERC Minimum Sample Media Input	N/A	allowed	N/A	allowed	allowed
FERC Maximum Sample Media Input	N/A	allowed	N/A	allowed	allowed
FERC Sample Average Media Input	N/A	allowed	N/A	allowed	allowed
FERC Current Sample Value Media Input	N/A	allowed	N/A	allowed	allowed
FERC Total Accumulated Media Input	N/A	allowed	N/A	allowed	allowed

7.7 Unidirectional Hot Data Path Reconfiguration

Support for fast, intervention-free, **unidirectional hot reconfiguration** of certain Data Path attributes in the current DPSM state (see section 6.2.4) is motivated by requirements of certain link speed negotiation protocols, as used, e.g., in Fibre Channel applications, which require to temporarily switch unidirectionally between similar Applications at different speeds, while retaining the lane allocation of the Data Path.

Note: Colloquially, this unidirectional reconfiguration feature is also known as Fibre Channel support.

Advertisement

Support for unidirectional reconfiguration is advertised by the UnidirReconfigSupported bit (see Table 8-53).

Commands

The fastest intervention-free hot reconfiguration employs switching between two prepared Staged Control Sets.

Note: Support for the 2nd Staged Control Set is advertised independently of UnidirReconfigSupported, but typical applications of unidirectional control will require two Staged Control Sets, for speed purposes.

For the fastest intervention-free **bidirectional** hot reconfiguration, the basic ApplyImmediate commands (Table 8-70 and Table 8-76) can be used.

For **direction-specific** intervention-free hot reconfiguration, two optional **unidirectional** hot reconfiguration commands are introduced, **ApplyImmediateTx** and **ApplyImmediateRx** (see Table 8-75 and Table 8-80), that restrict the effect of a hot reconfiguration command (including provisioning and commissioning) that the host has initiated by applying a Staged Control Set, to the Tx or Rx direction, respectively.

Note: To further speed-up the overall reconfiguration during time-critical link speed negotiations, some module form factors may even offer hardware control signals to trigger unidirectional reconfiguration commands, instead of writing to the CMIS trigger registers. The management of such 'rate or speed select Rx/Tx' functions (advertisement, administration, reporting) as well as the encoding and behavior of such control signals is outside of the scope of CMIS and is therefore left to form factor specific CMIS adjoint specifications. See section 8.8.

Unidirectional hot reconfiguration command triggers are effective only in the **DPIInitialized** or in the **DPActivated** steady states of the relevant DPSM. The module (silently) ignores them in all other states.

Note: This implies that the original Data Path on first entry to DPInitialized is Rx/Tx symmetrical.

Unidirectional hot reconfiguration of a Data Path Application is **restricted** to changing the AppSel code consistently on all lanes of the Data Path, and to optionally changing Explicit Control parameters (SI attributes) on some or all lanes of the Data Path. The lane allocation to the Data Path must not be modified.

Unidirectional hot reconfiguration of NP Applications is not supported.

Command Parameters

The command parameters for a unidirectional target reconfiguration are defined as usual in one of the supported Staged Control Sets and are selected by using the Apply registers associated with that Staged Control Set.

Command Validation, Execution, and Results

The positive or negative results of a unidirectional reconfiguration command and the reconfiguration command execution status (in progress) is found in the Configuration Command Status register (see Table 8-90).

In case of command parameter validation failure, the module shall not update the Active Control Set and not commit any changes to hardware.

A host shall allow a triggered immediate reconfiguration command to complete in the current state before triggering any DPSM state transition.

Committed Status Information

To represent the provisioned (and eventually also commissioned) configuration per direction, two new data structures are introduced **ACS::DPConfigTx** (see Table 8-160) and **ACS::DPConfigRx** (see Table 8-159), in addition to the SI settings which are already direction specific (see Table 8-94 and Table 8-95).

The classical Active Control Set (see Table 8-93) always represents the Tx configuration (which is always identical to the Rx configuration if unidirectional reconfiguration is not used).

Note: Hosts not using unidirectional reconfiguration facilities supported by a module may safely ignore the directional Active Control Sets and use the classical Active Control Set instead.

7.8 Host Lane Switching Capability

In the CMIS core specifications (see chapter 6) there is a fixed association between a physical electrical lane (identified by geometric and other physical characteristics of its physical connection point) and its lane number (or lane index) used in management registers, such as, e.g., the Data Path configuration register array. See also section 2.3.4.

The optional host lane switching functionality introduced here in this section breaks this simplicity and makes it necessary to distinguish the **external lane segment** and the **internal lane segment** that are **connected by** a variable **switching** function in between.

Likewise, we will need to carefully distinguish if a given **lane number** (or index) refers to the fixed **external lane segment** or to the **internal lane segment** that the external lane segment is switched to, because the same lane has different lane numbers, depending on which lane segment is used to identify the lane in the current switch configuration.

Note: Please refer to the glossary in section 3.3 for background of lane related terminology.

We will use the phrases **electrical** lane (number) and **nominal** lane (number) to distinguish the lane numbers of external host lane segments on one hand and flexibly connected internal lane segments on the other hand. These numbers are identical when switching is configured for pass-through (or not present).

7.8.1 Connectivity Limitations

The host interconnection electrical lane segments attached to the module's internal host lane segments of the n^{th} of M parallel Data Paths, either in a breakout module with parallel Data Paths (each with their own media lanes), or in a multiplexing module with several Host Paths attached to a Network Path, are determined by the lane-number-based rules described in 6.2.1.3 and depicted in Figure 6-1.

According to these rules, host interconnection electrical lanes are attached to Data Paths contiguously and sequentially in terms of their lane numbers. Since the association between lane numbers and physical lane ports is fixed in the CMIS core specifications, a given lane number always refers to the same host lane port of the module, and hence to the same host interconnection electrical lane attached to that port.

For modules supporting multiple media or host lane data rates and a variable number of host lanes for a given Data Path bandwidth, this can become a nuisance, because the host lanes assigned to a Data Path depend on the number of host lanes used for other Data Paths which are carried over lanes with smaller lane numbers.

To overcome these limitations a module can advertise support of an eight-by-eight host lane switching function.

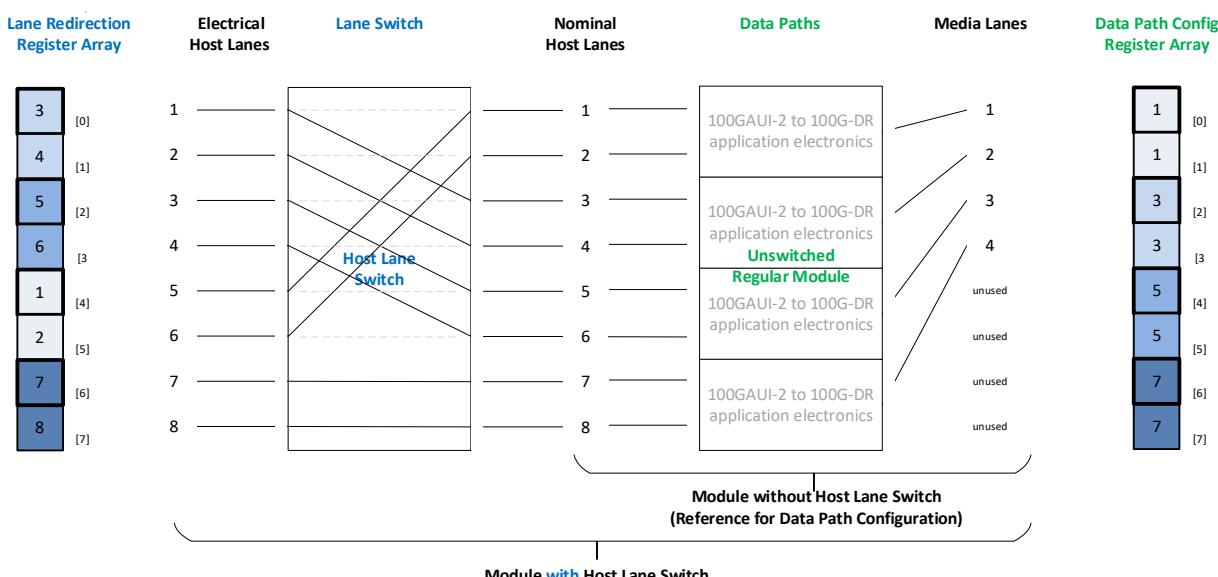


Figure 7-7 Host Lane Switching Model and Example

This lane switching function allows the host to freely define which module-external host electrical lane segments are connected to the module-internal nominal host lane segment of a Data Path, within each group of eight lanes.

7.8.2 Host Lane Switch Model

The 8x8 host lane switching function shown in Figure 7-7 allows the host to modify the normally fixed connections between module-external host electrical lanes and module-internal host lanes that are part of module-internal Data Paths, within each group of eight lanes.

This lane switch is modeled by (conceptually) inserting a configurable 8x8 host lane switch for each group of eight lanes between the **physical lane ports** of a module (with **fixed physical attachment** of the host electrical lanes) and the newly introduced module-internal **nominal host lanes** (ports) (with **fixed nominal association** to the management register view, with lane numbering that is unaffected by lane switching).

Note: The (ports of) these nominal lanes are only a modeling construct to allow the Application Description and Data Path definitions to remain unaffected by host side lane switching. For modules not supporting or not using lane switching, the nominal lane ports and the physical lane ports are indistinguishable.

7.8.3 Switch Configuration and Status

To represent the configuration of this host lane switch, the viewpoint of the physical lane ports (or external host interconnection electrical lanes) is taken, by introducing an array of eight lane redirection registers **RedirectionOfLane**<i>, with i in {1, ..., 8}. See also section 8.21.

Each redirection register **RedirectionOfLane**<i> contains the **nominal host lane** number j to which the host interconnection **electrical lane** i (physically attached to external lane port i) will be switched (connected to).

Elsewhere in this document every reference k to a host lane (via a lane number or lane index), e.g. in an Application Descriptor or in a Data Path configuration, refers to the **nominal host lane** k (shown on the right-hand side of the host lane switch in Figure 7-7). When the switch is unused, this distinction between physical and nominal lanes is not necessary.

Note: In a simplified view, the switch just rewrites the electrical host lanes connected to a "regular" CMIS module without a switch. It does not allow forking or joining data flows. Applications are described, configured, and supervised in terms of the nominal host lane number, not in terms of the reordered electrical lane numbers.

The current switch configuration status is always reported, which is especially important in case of failure. The switch configuration status is represented by a **RedirectStatusOfLane**<i>status register array that represents the actual state of the switch. Each **RedirectStatusOfLane**<i> register with i in {1, ..., 8} contains the **nominal host lane** j to which host interconnection **electrical lane** i is currently switched (connected to).

7.8.4 Host Lane Switch Operation

The usage model of the host lane switch consists of a **provisioning** area and a command with both in-progress and result feedback to **commit** the provisioned configuration. See section 8.21.

For simplicity, there is no advertisement of any 8x8 switching restrictions, nor of potential switching needs due to module limitations in certain Applications, even if implementations might have some. Therefore, the reconfiguration commit command returns if it is in progress, has succeeded, or has failed.

*Note: Recall that all registers or fields with a host lane association (like **ApplyDPIInit**, **DPCConfig**, **DPStatus**) continue to refer to nominal host lanes (shown on the right-hand side of the host lane switch in Figure 7-7). Only the actual electrical lane carrying the data of a given nominal host lane changes when the host lane switching function is used.*

Note: Changing the switch configuration does not affect the DPSM, nor the host lane numbering elsewhere, e.g., in registers with bits assigned to lanes.

Committing a new host lane switch configuration is allowed at any time in ModuleLowPwr and ModuleReady states, and independent of the Data Paths (and their DPSM states) that are associated with lanes.

The host must ensure that any initialized or activated Data Paths are either affected as a whole or not at all by a change in the Host Lane Switch configuration.

The module should validate that the Host Switch describes a true lane permutation and that all lanes of an existing Data Path are switched together.

Traffic integrity of lanes with unaffected switching configuration is expected to be unaffected. The traffic integrity of lanes with modified switching configuration is expected to be affected during the reconfiguration.

Note: The host should avoid interlacing Data Path commands and Host Lane Switching commands.

8 Module Management Memory Map

This chapter defines the structure and the meaning of the registers and fields in the Management Memory Map of a CMIS compliant module.

8.1 Overview and General Specifications

8.1.1 Management Memory Structure and Mapping

The register access protocol defined in section 5.2 only supports eight-bit Byte addresses. This limits the size of the **host addressable memory** that can be immediately accessed by the host to 256 bytes.

The host addressable memory is divided in two segments, **Lower Memory** (addresses 00h through 7Fh) and **Upper Memory** (addresses 80h through FFh).

This is illustrated in the upper part of Figure 8-1:

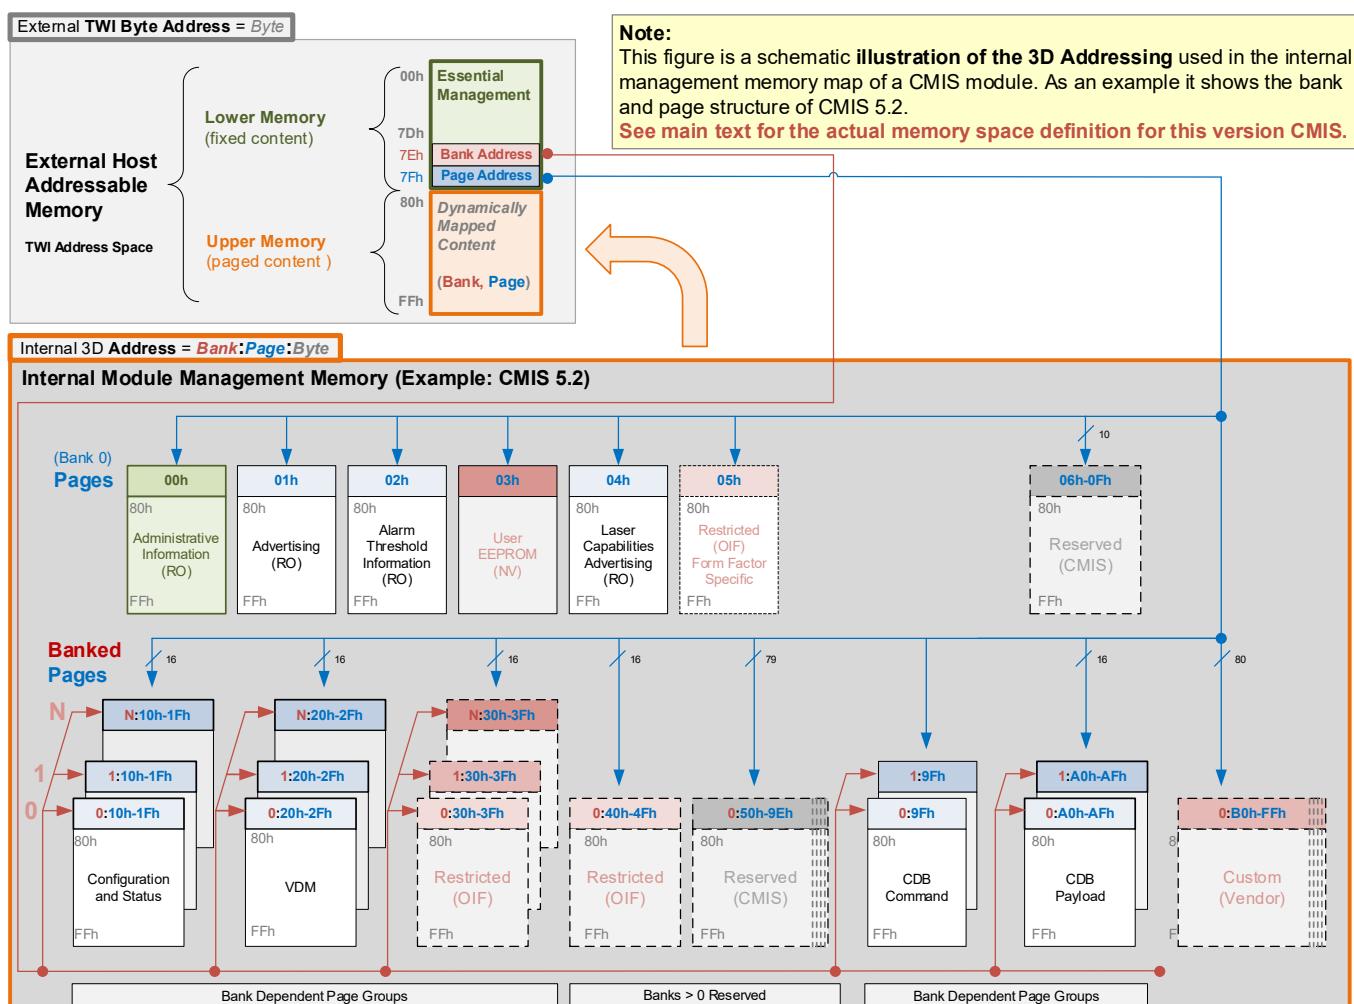


Figure 8-1 CMIS Module Memory Map (Conceptual View)

A larger **accessible** management memory is required for all but the most basic modules. This additional memory is supported by a structure of 128-byte **Pages** and by a mechanism for dynamically mapping any of these 128-byte Pages from a larger internal management memory space into the Upper Memory of the host addressable space.

Basic modules that support only a fixed address space without Upper Memory content switching are called **flat memory modules**, whereas modules that support a larger internal management memory with dynamic mapping of Pages into Upper Memory are called **paged memory modules**.

Note: In this version of CMIS, flat memory modules are assumed to support only read-only static data as provided by an EEPROM. See sections 6.3.2.3. and 8.2.

The conceptual structure of the additional internal management memory¹ and the dynamic mapping into Upper Memory is shown in the lower part of Figure 8-1. The management memory inside the module is arranged as a unique and always immediately host accessible address space of 128 bytes (Lower Memory) and as multiple upper address subspaces of 128 bytes each (**Pages**), only one of which is selected at any one time to be host visible in Upper Memory.

A second level of page selection is available for Pages for which several instances exist (e.g. where a **Bank** of Pages with the same Page number is supported).

This structure supports the flat 256-byte memory for **flat memory** modules like passive copper modules and permits timely access to addresses in the Lower Memory, e.g. Flags and Monitors, also for **paged memory modules**.

Less time critical entries, e.g. serial ID information and threshold settings, are available using the Page Select function in Lower Memory. For more complex modules which require a larger amount of management memory the host needs to use dynamic mapping of the various Pages into the host addressable Upper Memory address space, whenever needed.

Note: The management Memory Map has been designed largely after the QSFP Memory Map. This Memory Map has been changed in order to accommodate 8 electrical lanes and to limit the required memory space. The single address approach is used as found in QSFP. Paging is used in order to enable time critical interactions between host and module.

8.1.2 Map of Supported Pages and Banks

A basic 256-byte subset of the Management Memory Map is mandatory for all CMIS compliant devices, flat memory modules and paged memory modules.

Other parts are only available for paged memory modules, where a few basic Pages are required for all paged memory modules, while most other pages are supported as explicitly advertised by the module. See Table 8-46 for details regarding the advertisement of supported management memory spaces (Banks, Pages).

In particular, Lower Memory and Page 00h in Upper Memory is supported by all modules.

Flat memory modules support only Lower Memory and Page 00h fixed mapped into Upper Memory.

Paged memory modules must additionally support Pages **01h**, **02h** and Bank 0 of Pages **10h** and **11h**.

Bank 0 of Pages 10h-1Fh and 20h-2Fh provides lane-specific registers for the first 8 lanes, and each additional Bank provides support for an **additional set of 8 lanes**. These Pages are therefore called **lane banked**.

Note: Generally the allocation of information over the Banks may be Page specific and may not be related to grouping data for a group of 8 lanes, i.e. not all banked Pages are lane banked. For instance, additional Banks of Pages can also be supported for modules which require larger volume of management data.

For lane banked pages, a module may optionally support a **bank broadcast** feature, where writing to a particular current Page and Bank writes, virtually simultaneously, to all supported Banks of the current Page.

The following Figure 8-2 and Table 8-1 provide an overview and a detailed map of all Pages in the Management Memory Map in the two-dimensional space spanned by Page Address and Bank Address.

¹ The actual physical storage structure of these memory pages is not in the scope of this specification. This specification only defines how the available memory pages can be addressed.

1

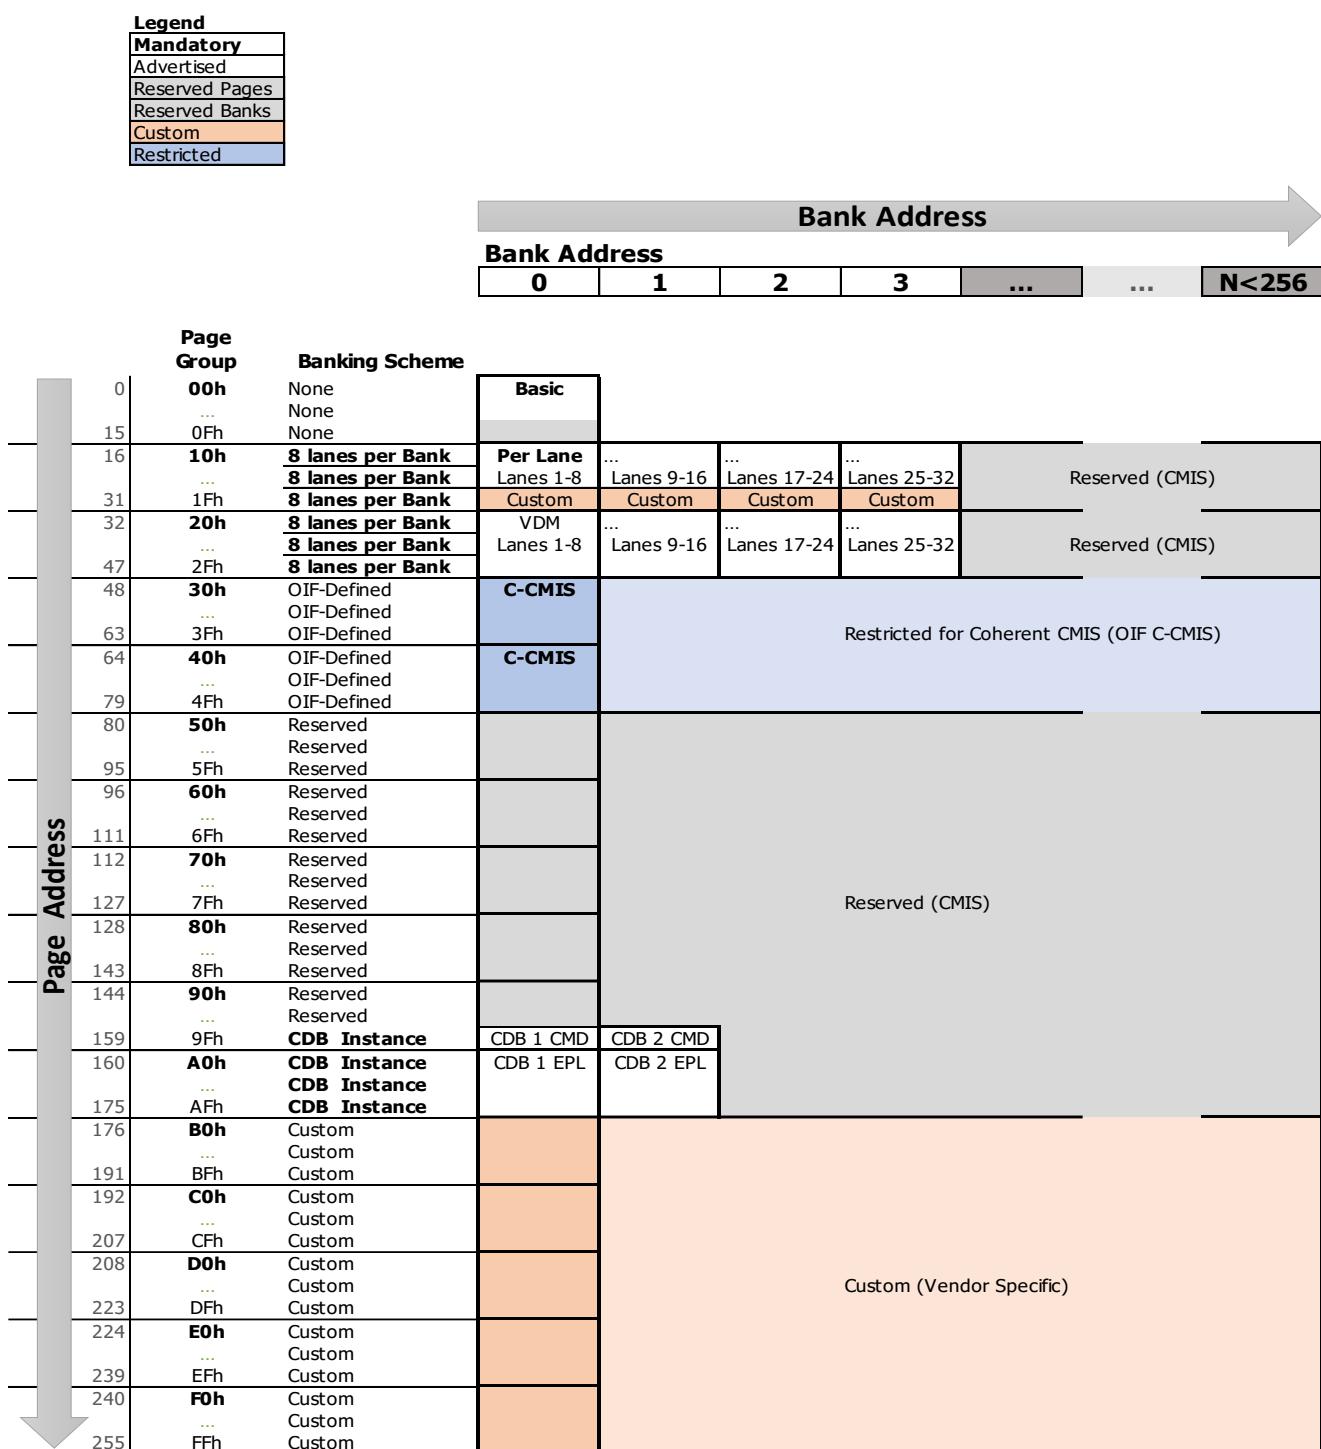


Figure 8-2 CMIS Bank and Page Group Iconic Memory Map Overview

2

3

4

Table 8-1 List of CMIS Pages

Page	#	Page Description	Data per Bank	# Banks	Type	See	
00h	1	Administrative Information	n/a	1 of 1	RO	8.3	
01h	1	Advertising	n/a	1 of 1	RO	8.4	
02h	1	Thresholds Information	n/a	1 of 1	RO	8.5	
03h	1	User NV RAM	n/a	1 of 1	RW	8.6	
04h	1	Laser Capabilities Advertising	n/a	1 of 1	RO	8.7	
05h	1	Restricted for OIF (CMIS-FF [8])	n/a	1 of 1	-	8.8	
06h-07h	2	Restricted for OIF (Resource Modules)	n/a	1 of 1	-		
08h-0Bh	4	Restricted for OIF (CMIS-LT [9])	n/a	1 of 1	-		
0Ch-0Fh	3	Reserved Pages (CMIS)	n/a	1 of 1	-	-	
10h	1	Lane and Data Path Configuration	8 lanes group	4 of 256	RW	8.9	
11h	1	Lane and Data Path Status	8 lanes group	4 of 256	RO	8.10	
12h	1	Tunable Laser Control and Status	8 lanes group	4 of 256	mixed	8.11	
13h	1	Module Performance Diagnostics Control	8 lanes group	4 of 256	RW	8.12	
14h	1	Module Performance Diagnostics Results	8 lanes group	4 of 256	RO	8.13	
15h	1	Timing Characteristics	8 lanes group	4 of 256	RO	8.14	
16h	1	Network Path Control and Status	8 lanes group	4 of 256	RW	8.15	
17h	1	Flags and Masks	8 lanes group	4 of 256	RO/COR	8.16	
18h	1	Lane and Data Path Configuration (Extensions)	8 lanes group	4 of 256	RW	8.17	
19h	1	Lane and Data Path Status (Extensions)	8 lanes group	4 of 256	RO	8.18	
1Ah-1Bh	2	Restricted for OIF (Resource Modules)	-	256	-	8.19	
1Ch	1	Normalized Applications Advertising	15 Applications	adv.	RO	8.20	
1Dh	1	Host Lane Switching	8 lanes group	4 of 256	mixed	8.21	
1Eh-1Fh	2	Custom Pages	8 lanes group	4 of 256	-	-	
20h	1	VDM Observable Descriptors[1][64]	8 lanes group	4 of 256	RO	8.22.1	
21h	1	VDM Observable Descriptors[2][64]	8 lanes group	4 of 256	RO		
22h	1	VDM Observable Descriptors[3][64]	8 lanes group	4 of 256	RO		
23h	1	VDM Observable Descriptors[4][64]	8 lanes group	4 of 256	RO		
24h	1	VDM Samples[1][64]	8 lanes group	4 of 256	RO	8.22.2	
25h	1	VDM Samples[2][64]	8 lanes group	4 of 256	RO		
26h	1	VDM Samples[3][64]	8 lanes group	4 of 256	RO		
27h	1	VDM Samples[4][64]	8 lanes group	4 of 256	RO		
28h	1	VDM TC Flagging ThresholdQuads[1][16] ¹	8 lanes group	4 of 256	RO	8.22.3	
29h	1	VDM TC Flagging ThresholdQuads[2][16]	8 lanes group	4 of 256	RO		
2Ah	1	VDM TC Flagging ThresholdQuads[3][16]	8 lanes group	4 of 256	RO		
2Bh	1	VDM TC Flagging ThresholdQuads[4][16]	8 lanes group	4 of 256	RO		
2Ch	1	VDM TC FlagQuads[256] (threshold crossing)	8 lanes group	4 of 256	RO/COR	8.22.4	
2Dh	1	VDM MaskQuads[256]	8 lanes group	4 of 256	RW	8.22.5	
2Eh	1	Reserved Pages (CMIS)	8 lanes group	4 of 256	-	-	
2Fh	1	VDM Advertisements and Dynamic Control	8 lanes group	4 of 256	RW	8.22.6	
30h-3Fh	16	Restricted for OIF (C-CMIS [7])	-	-	-	-	
40h-4Fh	16	Restricted for OIF (C-CMIS [7])	-	-	-	-	
50h-5Fh	16	Restricted for OIF (CMIS-LT [9])	-	-	-	-	
60h-9Eh	63	Reserved Pages (CMIS)	-	-	-	-	
9Fh	1	CDB Command/Response with Local Payload	CDB Instance	2 of 256	RW	8.23	
A0h	1	CDB EPL[1][128] Extended Payload Segments	CDB Instance	2 of 256	RW	8.24	
A1h	1	CDB EPL[2][128]	CDB Instance	2 of 256	RW		
A2h	1	CDB EPL[3][128]	CDB Instance	2 of 256	RW		
A3h	1	CDB EPL[4][128]	CDB Instance	2 of 256	RW		
A4h	11	CDB EPL[5][128]	CDB Instance	2 of 256	RW		
AEh		...					
AFh	1	CDB EPL[15][128]	CDB Instance	2 of 256	RW		
B0h-FFh	80	Custom Pages	Vendor specific	-	-	-	

¹ Note that one ThresholdQuad set of thresholds serves several Observables

8.1.3 Specifications Conventions

This section presents specifications conventions that are used in the Management Memory Map definitions.

8.1.3.1 Reserved Locations

Pages or Banks marked as **reserved** are reserved for future use in this specification. Reserved Pages or Banks are not accessible, i.e. neither functionality nor locations are supported (see section 8.2.15).

Locations within a Page (Bytes, Fields, or Bits) marked as **reserved** are also reserved for future use in this specification (see section 3.1). However, such locations are **accessible**, albeit without associated functionality.

The module shall zero-initialize accessible reserved locations. There are no other obligations for the module.

Both host and module shall neither use (evaluate) nor modify data from reserved locations.

Note: A subtle exception is hosts reading and testing for module-initialized zero, which allows newer hosts to work with older modules, when new registers are encoded with backwards interworking in mind.

Note: Another subtle exception is for a defensive host to read (but not use) locations reserved for Flags.

The effects of a forbidden host WRITE to a reserved location is undefined.

Note: As a module is therefore not obliged to store a WRITE in a reserved location, introducing new writeable registers even with encoding for backwards interworking still always requires explicit advertisement.

Note: Other organizations must contact the managing organization or the editor of this specification to request allocations of registers, fields, or bits.

8.1.3.2 Custom Locations

Pages or Banks marked as **custom** are for vendor defined use. Access to custom Banks or Pages that are not functionally supported by a module shall not be supported, i.e. no access shall be possible (see section 8.2.15).

Locations within a Page (Bytes, Fields, or Bits) marked as **custom** are not governed by this specification and may be vendor defined. The use of registers defined as custom may be subject to additional agreements between module users and vendors.

Note: The requirement to generally enable vendor-agnostic management of fully CMIS compliant modules implies that the use of custom managed elements must be optional.

Note: The requirement to allow vendor-agnostic management of CMIS compliant modules implies that custom interrupt sources (Flags) must remain silent unless the host is known to be aware of the custom extensions.

8.1.3.3 Bank-Dependent Lane Number Interpretation

In this subsection, the term lane number refers to **host** lane numbers and **media** lane numbers.

Banking is used on certain Pages (such as 10h-1Fh and 20h-2Fh) to extend the number of supported lanes in groups of eight lanes. Such Pages are also called **lane banked**.

On these lane-banked pages, the reader shall interpret references to lane numbers $\langle i \rangle$ in $\{1, 2, \dots, 8\}$ in the register name or register description on a banked page as the “representation” of a unique bank-dependent lane number $\langle j \rangle$ in $\{1, 2, \dots, 32\}$, as follows:

The unique lane numbers $\langle j \rangle$ on bank $\langle k \rangle$ in $\{0, 1, 2, 3\}$ are $\langle j \rangle = \langle k \rangle \cdot 8 + \langle i \rangle$.

Vice versa, the lane number $\langle i \rangle$ in $\{1, \dots, 8\}$ that “represents” a unique lane number $\langle j \rangle$ in $\{1, 2, \dots, 32\}$ in the text is found as follows:

The bank 0 lane number $\langle i \rangle$ corresponding to $\langle j \rangle$ is determined as $\langle i \rangle = (\langle j \rangle - 1) \bmod 8 + 1$

These relations are shown in the following table

Table 8-2 Bank Dependent Lane Number Interpretation

BankSelect \ Lane	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
1	9	10	11	12	13	14	15	16
2	17	18	19	20	21	22	23	24
3	25	26	27	28	29	30	31	32

8.1.3.4 Register Default Values

Default values for all **control** registers are 0 unless otherwise specified.

Note: Hosts are encouraged to review or explicitly set critical registers and not rely on module default values.

Note: For Control Set registers the default settings may be Application dependent, see section 6.2.3.

The default value of **Masks** for Flags residing on **optional** pages is 1 (Interrupt is masked by default).

*Note: A default value of 1 is also recommended for all Masks of **custom** Interrupt sources.*

8.1.3.5 Byte Order (Endianness)

The default byte order of multi-byte registers representing numerical types is **Big Endian**, i.e. the lowest byte address contains the most significant byte of the multi-byte value. Exceptions are stated explicitly

Note: Independent of proper representation, atomic accesses to multi-byte registers may require access synchronization protocols (see also section 5.2.3 for a discussion of data coherency)

8.1.3.6 Access Types

The following Field access types are distinguished and indicated (together with an optionality indication) in the Type column of the field definition tables defining the Management Memory Map later in this chapter. The word "element" is used here to refer to an entire Byte, a Field, or a Bit:

Table 8-3 Access Types

Access Type	Description
RW	A readable and writeable element.
RWW	A readable and writeable element that can also be modified by the module. <i>Note: This access type should be used only for interactions governed by a protocol.</i>
RO	A read-only element. A WRITE of a value to a read-only element is allowed but has no effect.
WO	A write-only element. A READ from a write-only element is allowed but delivers unpredictable values.
WO/SC	A write-only element with self-clearing side effect. A READ from a WO/SC element is allowed and delivers a zero value, except transiently when reading before the module has evaluated and cleared the non-zero bits written.
RO/COR	A read-only element with clear-on-read (clear-after-read) side effect. All bits in a RO/COR Byte are cleared by the module after the Byte value has been read. <i>Note: Modules may also combine clearing and update (clear-or-update-on-read).</i>

Note: The clear-on-read (COR) access type is used for latched Flags that indicate that an event has occurred, or a condition was present since the last read. Note that a Flag cleared by reading will be set again if the module evaluates the flagged condition as still being present. The timing of this Flag update after a clear-on-read is specified in section 9.12.7.

Note: Bits with WO or WO/SC access types are often used as command or trigger bits that start the execution of something. A WO/SC specification is mainly useful when privacy protection of written data is to be specified.

Note: The ACCESS protocol allows to read from or write to any addressable register; the Access Type is therefore only a specification of value related behavior, not an actual access restriction. Especially there are neither protocol violations nor protocol exceptions when the host, e.g., writes to a RO element.

8.1.3.7 Optionality Indications

Identified elements of the Memory Map described in this chapter are conditionally or unconditionally optional. An optional element is either supported or not-supported by a specific module.

Three levels of optionality are distinguished: Optionality of Pages, of Banks, and of Fields on (banked) Pages.

For Pages, only two cases are distinguished, as documented in the text describing the Page

- Required: All paged memory modules must support required Pages.
- Optional: All other Pages are optional. The host can always determine, directly or indirectly from advertisements or other field values, if an optional Page is supported.

For Banks, the host can always determine, directly or indirectly from advertisements or other field values, if an optional Bank is supported.

1 For Fields, the following classification is documented in the Type column of the Field definition tables:

Required	Rqd.	Always supported when the relevant Page is supported.
Advertised	Adv.	Supported when indicated in a clearly associated feature advertisement
Conditional	Cnd.	Supported when some (documented) run-time evaluated condition is present
Optional	Opt.	optionally supported, unknown conditions, inferred only from feature behavior

2
3 Note: Ideally the host should be able to find out if an optional field is actually supported. However, for historical
4 reasons, there are some optional fields for which no advertisement exists. These tend to be features which
5 have a meaningful default value where absence of support can be inferred from non-changing default values.

7 8.1.4 General Specifications

8 This section presents specification that are applicable to wide parts of the Management Memory Map.

9 8.1.4.1 Multi-Byte Reporting Registers

10 Scalar multi-byte reporting registers (e.g. monitored numerical values) with explicitly specified data type (i.e.
11 specified as e.g. F16, S32, S64) are coherently and atomically readable in a multibyte READ over the length of
12 the multi-byte reporting register. See section 3.4 for defined data types and section 5.2.5 for data coherency
13 in multi-byte ACCESS operations.

14 Note: Unlike in previous revisions a multi-byte register is now described as one multi-byte register covering a
15 byte address range, not as individual bytes with individual byte names.

16 Other multi-byte registers do not support coherent and atomic ACCESS unless explicitly specified.

17 8.1.4.2 Flags, Masks, and Interrupts

18 Flags

19 In this specification, a **Flag** is a latched indicator bit in management memory with associated maskable Interrupt
20 request generation.

21 While a Flag is set, an **Interrupt** request is generated unless an associated **Mask** bit is set.

22 Once asserted, a Flag bit remains set (latched) until cleared by a READ of the Byte containing the Flag (or
23 cleared in the course of the reaction to a Reset signal).

24 A Flag has **clear on read** access type.

25 After being read and cleared, a Flag indicating presence of a condition is raised again if the flagged condition
26 persists; this will cause the Interrupt request to be asserted again (unless masked). Flags indicating the
27 occurrence of an event will only be raised again when the next flagged event occurs.

28 Note: In a situation where a flagged condition itself is permanent, the module may first clear (on read) and
29 then re-assert the cleared Flag after a finite condition re-evaluation time interval. A fast sequence of Flag
30 readings may therefore result in toggling Flag values while the flagged condition continues to be present.

31 Note: Generally, the timing limits, sequencing relationships, and interdependencies with regard to state changes
32 of Flags, Mask, and Interrupt Requests is not specified in this version of CMIS (see chapter 9.12.7). Hosts should
33 therefore be aware that the synchronization between changes of a flagged condition or event, Flag or Mask
34 state changes, and Interrupt request assertion changes (and hence the transient behavior between stable
35 conditions) may vary significantly across modules.

36 Note: A module raises Flags only when allowed by Flagging Conformance rules (see section 6.3.4) which depend
37 on the states of the Module State Machine and the Data Path State Machine.

38 Flag Summaries

39 To allow quick scanning of a Flag register tree, a bit in a Flag summary register indicates if a Flag in an
40 associated underlying Flag register (or in a next level Flag summary) is set. Note that a Flag summary bit itself
41 is not a Flag and reading a Flags summary register does not clear any underlying (summarized) Flag.

42 Interrupt Request

43 The Interrupt request hardware signal is asserted as long as there is any Flag set with its corresponding Mask
44 bits cleared.

An Interrupt is raised at the onset of an unmasked flagged condition or at the occurrence of an unmasked flagged event and remains asserted until all asserted unmasked Flags (both module-level and lane-specific) either have been cleared by a host read or have been masked by the host (by setting the relevant Mask bits as described below).

Masks

In this specification a **Mask** bit is a host writable bit in management memory which, when set, suppresses future Interrupt generation, and ceases current Interrupt generation caused by its associated Flag bit.

Note: Mask bits may be used to prevent an Interrupt request from remaining active while the host performs actions to acknowledge and handle the condition causing the Flag to be set.

For each Flag there is a corresponding Mask bit, usually at the same Bit index in a Mask configuration Byte. A set Mask bit suppresses the contribution of its corresponding Flag to the hardware Interrupt signal generation. A cleared Mask bit enables the contribution of its corresponding Flag to the hardware Interrupt signal generation. Setting a Mask on a set Flag deasserts the hardware Interrupt unless there are other unmasked Flags set.

Mask bits are volatile: During module initialization all CMIS defined Mask bits are set to default values (at exit from **MgmtInit** state to **ModuleReady** state, see section 6.3.2.5.3).

Note: This choice exists for historical reasons. Normally default Interrupt silence would be preferable.

Note: The host may use known Mask bits to temporarily suppress a continued hardware Interrupt assertion due to stable conditions, i.e. from not cleared Flags, which would otherwise continually reassert the hardware Interrupt signal.

Note: Custom masks for custom functions should be masked by default.

8.1.4.3 Generic Checksums

Page Checksum

Unless specified differently, a Page Checksum value is the arithmetic sum of a given Byte range, modulo 256.

CDB Checksum

A CDB message checksum value is the one's complement of the arithmetic sum of a given Byte range, modulo 256.

Note: Both checksums may be applied to non-contiguous ranges (where Bytes are excluded or zeroed).

8.2 Lower Memory (Control and Status Essentials)

Lower Memory consists of the “lower” 128 bytes of the 256-byte host-addressable memory.

Lower Memory is addressed with byte addresses in the range 00h to 7Fh. The byte addresses of Lower Memory are statically interpreted: a given byte address always refers to the same physical register in Lower Memory.

Note: In contrast, Upper Memory byte addresses in the range 80h to FFh are dynamically interpreted: the actually addressed register depends on the current page mapping as described in section 8.1.1 and below.

Note: Lower Memory content is always immediately accessible and therefore contains mainly registers that are deemed to be most frequently accessed by the host, such as module state and interrupt status, module level monitors, module level Flags and Masks, lane level Flags summary, and global control functions.

Flat memory modules that support only constant read-only data are not required to support any writeable control fields or readable dynamic status reporting fields defined in Lower Memory.

Note: This enables ROM-based implementation of the addressable memory in flat memory modules.

Lower Memory is subdivided into subject areas as shown in the following table:

Table 8-4 Lower Memory Overview

Address	Size	Subject Area	Description	Type
0-2	3	Management Characteristics	Basic advertisement about how this module is managed	Adv.
3	1	Global Status Information	Current state of Module, Interrupt signal status	Status
4-7	4	Flags Summary	Summary of Flags set on specific Pages (and Banks)	
8-13	6	Module-Level Flags	Flags that are not lane or Data Path specific	
14-25	12	Module-Level Monitors	Monitors that are not lane or Data Path specific	
26-30	5	Module-Level Controls	Controls applicable to the module as a whole	Control
31-36	6	Module-Level Masks	Mask bits for the Module-Level Flags	
37-38	2	CDB Command Status	Status of current CDB command	
39-40	2	Module Active Firmware Version	Module Active Firmware Version	
41	1	Module Fault Information	Fault cause for entering ModuleFault state	Status
42-45	4	Miscellaneous Status Information	Feedback indicators (e.g. password entry results)	
46-55	10	-	Reserved[10]	-
56-63	48	Extended Module Information	Additional information how module is to be managed	Adv.
64-84	21	-	Custom[21]	-
85-117	33	Applications Advertising	Applications supported by module Data Path(s)	Adv.
118-125	8	Password Facilities	Password Entry and Change	Control
126-127	2	Page Mapping	Page mapping into host addressable Upper Memory	

15

Note: PCI-SIG has coordinated to use the custom fields 00h:64-84 for PCI-specific feature advertisement for PCIe applications, with specifications in [19].

8.2.1 Management Characteristics

The Management Characteristics fields described in Table 8-5 provide fundamental management characteristics of the module that allow hosts to verify if, and to determine how, they can operate and manage the module.

The Management Characteristics fields include an SFF-8024 module type identifier (this is a value defined in the Identifier Values table in [5]) which implicitly contains information about the management protocol and the management Memory Map offered by the module.

Note: Since the management protocol is not systematically encoded in the SFF 8024 identifiers, hosts will have to test against a list of supported SFF 8024 module type identifiers.

The other fields indicating management characteristics like the CMIS version number, memory model (flat memory or paged memory), support for intervention-free reconfiguration, and management interface speed, can be interpreted once the module has been recognized as a CMIS compliant module.

*Note: Modules may also be classified by their functional "module type", depending on the unique or main Application that they support. This kind of "module type" is **not** encoded in the SFF8024Identifier field, but can instead be derived from other advertisements.*

Table 8-5 Management Characteristics (Lower Memory)

Byte	Bits	Field Name	Field Description	Type
0	7-0	SFF8024Identifier	SFF8024Identifier is an SFF-8024 module type Identifier from Table 4-1 “Identifier Values” in [5] which allows to infer both physical form factor and management protocol of the module. <i>Note: The CMIS interpretation of all other registers or fields is valid only when the fundamental SFF8024Identifier indicates that the module uses the CMIS management protocol.</i>	RO Rqd.
1	7-0	CmisRevision	CMIS revision number (decimal): The upper nibble (bits 7-4) is the integer part (major number) The lower nibble (bits 3-0) is the decimal part (minor number) <i>Example: 01h indicates version 0.1, 21h indicates version 2.1.</i> <i>Note: See Appendix G.3 for interoperability implications of the major revision number (integer part).</i>	RO Rqd.
2	7	MemoryModel	Indicator of the memory model of the module: 0b: Paged memory (Pages 00h-02h, 10h-11h supported) 1b: Flat memory (Page 00h supported only)	RO Rqd.
	6	SteppedConfigOnly	Indicates which reconfiguration procedures are supported: 0b: Module supports all types of reconfigurations: essential step-by-step reconfiguration as well as both types of intervention-free reconfiguration: hot (in state) and regular (with automatic state change). <i>(legacy default, not meant to be pre-valent, see notes below)</i> 1b: Module supports the essential step-by-step reconfiguration but none or (optionally) only one of the two types of intervention-free reconfiguration. See the AutoCommissioning field below for more details. <i>Note: The baseline option 1b is assumed to be prevalent for modules that do not require hot reconfiguration, such as, e.g., Ethernet transceivers.</i> <i>Note: The extended option 0b with additional support for intervention-free reconfiguration is intended only for module Applications such as InfiniBand or Fibre Channel that may require time critical speed negotiations when tight timing is not achievable with basic step-by-step configuration.</i> <i>Note: See section 6.2.4 for more information.</i>	

	5-2	MciMaxSpeed	<p>Indicates maximum supported clock speed of Management Communication Interface (MCI):</p> <p>I2CMCI: 0: Module supports up to 400 kHz 1: Module supports up to 1 MHz 2: Module supports up to 3.4 MHz 3-15: Reserved</p> <p>SPI MCI: 0: Module supports up to 1 MHz 1: Module supports up to 2 MHz 2: Module supports up to 4 MHz 3: Module supports up to 8 MHz 4: Module supports up to 12 MHz 5: Module supports up to 16 MHz 6: Module supports up to 20 MHz 7: Module supports up to 30 MHz 8: Module supports up to 40 MHz 9: Module supports up to 50 MHz 10-15: Reserved</p>	
	1-0	AutoCommissioning	<p>This field allows modules to support just one of the two intervention-free reconfiguration procedures, either the regular (automatic DPSM state changing) one or the hot one (staying in DPSM state):</p> <p>SteppedConfigOnly = 0: xx: both regular and hot supported (legacy default)</p> <p>SteppedConfigOnly = 1: 00: none, neither regular nor hot supported 01: only regular supported (affects ApplyDPIInit) 10: only hot supported (affects ApplyImmediate) 11: reserved</p> <p>Effect on ApplyImmediate: When hot intervention-free reconfiguration is unsupported, the module ignores any WRITE to ApplyImmediate registers</p> <p>Effect on ApplyDPIInit: When regular intervention-free reconfiguration is not supported, the module accepts ApplyDPIInit in all states as a Provision command (see Table 6-4), but the DPSM excludes the DPReinitT term in the DPDeReinitS transition signal (see Case 2 in section 6.3.3.1)</p>	

8.2.2 Global Status Information

The fields described in Table 8-6 provide fundamental module status indicators.

See section 6.3 for information on the Module State Machine

Table 8-6 Global Status Information (Lower Memory)

Byte	Bits	Field Name	Field Description	Type
3	7-4	-	Reserved	RO Rqd.
	3-1	ModuleState	Current Module State (see Table 8-7 for encoding and section 6.3.2 for a description of the meaning of ModuleState) <i>Notes:</i> - Flat memory modules always report ModuleReady. - Not all states of the Module State Machine are observable.	
	0	InterruptDeasserted	Status of Interrupt output signal 1b: Interrupt not asserted (default) 0b: Interrupt asserted	

Table 8-7 Module State Encodings

Code	Module State	Description
000b	-	Reserved
001b	ModuleLowPwr	
010b	ModulePwrUp	
011b	ModuleReady	This is the only state reported by flat memory modules
100b	ModulePwrDn	
101b	ModuleFault	
110b	-	Reserved
111b	-	Reserved

8.2.3 Flags Summary

The Flags Summary bits indicate when any Flags are asserted on specific pages, for up to 4 Banks.

Note: To clear a summarized Flag, the Flag itself must be read from the relevant Page on the appropriate Bank.

Table 8-8 Lane-Level Flags Summary (Lower Memory)

Byte	Bit	Field Name	Field Description	Type
4	7	-	Reserved	
	6	-	Reserved	
	5	-	Reserved	
	4	-	Reserved	
	3	FlagsSummaryBank0Page2Ch	1b: at least one Flag is set on Bank 0, Page 2Ch	RO Adv.
	2	FlagsSummaryBank0Page14h	1b: at least one Flag is set on Bank 0, Page 14h	RO Adv.
	1	FlagsSummaryBank0Page12h	1b: at least one Flag is set on Bank 0, Page 12h	RO Adv.
	0	FlagsSummaryBank0Page11h	1b: at least one Flag is set on Bank 0, Page 11h	RO Rqd.
	7	-	Reserved	
5	6	-	Reserved	
	5	-	Reserved	
	4	-	Reserved	
	3	FlagsSummaryBank1Page2Ch	1b: at least one Flag is set on Bank 1, Page 2Ch	RO Adv.
	2	FlagsSummaryBank1Page14h	1b: at least one Flag is set on Bank 1, Page 14h	RO Adv.
	1	FlagsSummaryBank1Page12h	1b: at least one Flag is set on Bank 1, Page 12h	RO Adv.
	0	FlagsSummaryBank1Page11h	1b: at least one Flag is set on Bank 1, Page 11h	RO Adv.
	7	-	Reserved	
	6	-	Reserved	
6	5	-	Reserved	
	4	-	Reserved	
	3	FlagsSummaryBank2Page2Ch	1b: at least one Flag is set on Bank 2, Page 2Ch	RO Adv.
	2	FlagsSummaryBank2Page14h	1b: at least one Flag is set on Bank 2, Page 14h	RO Adv.
	1	FlagsSummaryBank2Page12h	1b: at least one Flag is set on Bank 2, Page 12h	RO Adv.
	0	FlagsSummaryBank2Page11h	1b: at least one Flag is set on Bank 2, Page 11h	RO Adv.
	7	-	Reserved	
	6	-	Reserved	
	5	-	Reserved	
7	4	-	Reserved	
	3	FlagsSummaryBank3Page2Ch	1b: at least one Flag is set on Bank 3, Page 2Ch	RO Adv.
	2	FlagsSummaryBank3Page14h	1b: at least one Flag is set on Bank 3, Page 14h	RO Adv.
	1	FlagsSummaryBank3Page12h	1b: at least one Flag is set on Bank 3, Page 12h	RO Adv.
	0	FlagsSummaryBank3Page11h	1b: at least one Flag is set on Bank 3, Page 11h	RO Adv.

8.2.4 Module-Level Flags

The registers described in Table 8-9 contain module-level (global) Flags, with Masks described in section 8.2.7.

Note: The general behavior of Flags, Masks, and Interrupts is specified in section 8.1.4.2.

Module-level Flags are used to report module-level status changes, operating failures, as well as threshold crossing alarms and warnings for monitored observables. Monitors with associated alarm and/or warning thresholds have associated alarm Flags, warning Flags. For normal operation and in default state, these Flags are cleared. Refer to section 6.3.4 for Flagging rules depending on ModuleState.

Byte 13 is provided for Custom Module Level Flags.

Table 8-9 Module Flags (paged memory modules only) (Lower Memory)

Byte	Bit	Name	Field Description	Type
8	7	CdbCmdCompleteFlag2	Latched Flag to indicate completion of a CDB command for CDB instance 2. Support is advertised in field 01h:163.7-6	RO/COR Adv.
	6	CdbCmdCompleteFlag1	Latched Flag to indicate completion of a CDB command for CDB instance 1. Support is advertised in field 01h:163.7-6	RO/COR Adv.
	5-3	-	Reserved	Rqd.
	2	DataPathFirmwareErrorFlag	Latched Flag to indicate that subordinated firmware in an auxiliary device for processing transmitted or received signals (e.g. a DSP) has failed.	RO/COR Adv.
	1	ModuleFirmwareErrorFlag	Latched Flag to indicate that self-supervision of the main module firmware has detected a failure in the main module firmware itself. There are several possible causes of the error such as program memory becoming corrupted and incomplete firmware loading.	RO/COR Adv.
	0	ModuleStateChangedFlag	Latched Flag to indicate a Module State Change	RO/COR Rqd.
9	7	VccMonLowWarningFlag	Latched Flag for low supply voltage warning	RO/COR Adv.
	6	VccMonHighWarningFlag	Latched Flag for high supply voltage warning	
	5	VccMonLowAlarmFlag	Latched Flag for low supply voltage alarm	
	4	VccMonHighAlarmFlag	Latched Flag for high supply voltage alarm	
	3	TempMonLowWarningFlag	Latched Flag for low temperature warning	
	2	TempMonHighWarningFlag	Latched Flag for high temperature warning	
	1	TempMonLowAlarmFlag	Latched Flag for low temperature alarm	
	0	TempMonHighAlarmFlag	Latched Flag for high temperature alarm	
10	7	Aux2MonLowWarningFlag	Latched Flag for low Aux 2 monitor warning	RO/COR Adv.
	6	Aux2MonHighWarningFlag	Latched Flag for high Aux 2 monitor warning	
	5	Aux2MonLowAlarmFlag	Latched Flag for low Aux 2 monitor alarm	
	4	Aux2MonHighAlarmFlag	Latched Flag for high Aux 2 monitor alarm	
	3	Aux1MonLowWarningFlag	Latched Flag for low Aux 1 monitor warning	
	2	Aux1MonHighWarningFlag	Latched Flag for high Aux 1 monitor warning	
	1	Aux1MonLowAlarmFlag	Latched Flag for low Aux 1 monitor alarm	
	0	Aux1MonHighAlarmFlag	Latched Flag for high Aux 1 monitor alarm	
11	7	CustomMonLowWarningFlag	Latched Flag for low Vendor Defined Monitor warning	RO/COR Adv.
	6	CustomMonHighWarningFlag	Latched Flag for high Vendor Defined Monitor warning	
	5	CustomMonLowAlarmFlag	Latched Flag for low Vendor Defined Monitor alarm	
	4	CustomMonHighAlarmFlag	Latched Flag for high Vendor Defined Monitor alarm	
	3	Aux3MonLowWarningFlag	Latched Flag for low Aux 3 monitor warning	
	2	Aux3MonHighWarningFlag	Latched Flag for high Aux 3 monitor warning	
	1	Aux3MonLowAlarmFlag	Latched Flag for low Aux 3 monitor alarm	
	0	Aux3MonHighAlarmFlag	Latched Flag for high Aux 3 monitor alarm	
12	7-0	-	Reserved[1]	
13	7-0	-	Custom[1]	

8.2.5 Module-Level Monitor Values

Real time monitoring for module-level observables includes two monitors with fixed observables (temperature and supply voltage) and four monitors with selectable function (3 auxiliary and 1 vendor defined) as shown in Table 8-10.

Note: The data format of a monitored value may have greater resolution and range than required.

Measurement accuracy is defined by the relevant interface standard or module product specification.

The reported monitoring results of supported module level monitors shall be within the relevant accuracy requirements when the module is in the ModuleReady state.

Table 8-10 Module-Level Monitor Values (paged memory modules only) (Lower Memory)

Byte	Bit	Register Name	Register Description	Type
14-15	7-0	TempMonValue	S16 Module Temperature Monitor (Current Value) internally measured temperature in 1/256 degree Celsius increments	RO Adv.
16-17	7-0	VccMonVoltage	U16 Supply Voltage Monitor (Current Value) internally measured input supply voltage in 100 μ V increments	RO Adv.
18-19	7-0	Aux1MonValue	S16 Aux1 Monitor (see Table 8-49) (Current Value) The monitored observable is advertised in 01h:145.0: 0b: Custom 1b: TEC Current in 100%/32767 increments of maximum TEC current magnitude, i.e. of the larger of the heating or cooling current magnitudes +32767 (100%) of the max current magnitude when heating -32767 is -100% of the max current magnitude when cooling	RO Adv.
20-21	7-0	Aux2MonValue	S16 Aux2 Monitor (see Table 8-49) (Current Value) The monitored observable is advertised in 01h:145.1: 0b: Laser Temperature : in 1/256 degree Celsius increments 1b: TEC Current in 100%/32767 increments of maximum TEC current magnitude, i.e. of the larger of the heating or cooling current magnitudes +32767 (100%) of the max current magnitude when heating -32767 is -100% of the max current magnitude when cooling	RO Adv.
22-23	7-0	Aux3MonValue	S16 Aux3 Monitor (see Table 8-49) (Current Value) The monitored observable is advertised in 01h:145.2: 0b: Laser Temperature : in 1/256 degree Celsius increments 1b: Additional Supply Voltage : in 100 μ V increments	RO Adv.
24-25	7-0	CustomMonValue	S16 or U16: Custom monitor (Current Value)	RO Adv.

8.2.6 Module-Level Controls

Module-Level (global) controls are applicable to the entire module or to all Lanes or Data Paths in the module.

Note: Lane-specific controls are located in Page 10h (see section 8.9).

Table 8-11 Module Global Controls (paged memory modules only) (Lower Memory)

Byte	Bit	Field Name	Field Description	Type
26	7	BankBroadcastEnable	<p>0b: Bank broadcast for lane-banked pages disabled 1b: Bank broadcast for lane-banked pages enabled</p> <p>When BankBroadcastEnable is set, a WRITE to a control register (i.e. to a register with RW or WO access) in any bank of a lane-banked page is executed as a bank broadcast. Recall that a banked page is lane-banked if banking is used to add support for additional lanes. A bank broadcast is a virtually simultaneous and atomic WRITE of the same value to the same register and the same page, in all supported banks. The module ensures a generalized broadcast register readback condition (see section 5.2.4), such that a READ from any supported bank for the same page and register always yields the value written in the broadcasted WRITE.</p> <p>Advertisement: 01h:156.7</p>	RW Adv.
6	LowPwrAllowRequestHW		<p>Enables evaluation of the LowPwrRequestHW hardware signal 0b: Module ignores the LowPwrRequestHW signal 1b: Module evaluates the LowPwrRequestHW signal (default)</p> <p><i>Note: See Table 6-12 and section 6.3.2.2 for more information Note: As LowPwrRequestSW is cleared by default, evaluation of LowPwrRequestHW is enabled by default, allowing the host to request start-up to halt in Low Power mode.</i></p>	RW Rqd.
5	SquelchMethodSelect		<p>0b: Squelching of Tx output reduces OMA 1b: Squelching of Tx output reduces P_{av}</p> <p>Advertisement: 00h:156.5-4 <i>Note: Method to choose depends on interface standard used. See Table 8-50 for SquelchMethodSelect capability advertising.</i></p>	RW Adv.
4	LowPwrRequestSW		<p>0b: No request 1b: Request for the module to stay in, or to return into, Low Power mode</p> <p><i>Note: See Table 6-12 and section 6.3.2.4 for more information</i></p>	RW Rqd.
3	SoftwareReset		<p>Self-clearing trigger bit that causes the module to be reset when 1b is written to it.</p> <p>The effect of a SoftwareReset trigger is the same as asserting the Reset hardware signal for the appropriate hold time, followed by its de-assertion.</p> <p>0b: No action 1b: Software reset</p>	WO/SC Rqd.
2-0	-		Custom	

27	7-4	-	Reserved	
	3-0	MciSpeedConfiguration	<p>For SPIMCI:</p> <p>The host configures the speed used for SPI bus transactions.</p> <p>The module uses this MciSpeedConfiguration field to compute the number of fill bytes if it advertises its AddressDecodingDurationFormat as "time" (0h:213.7 = 1).</p> <p>0: 1MHz 1: 2MHz 2: 4MHz 3: 8MHz 4: 12MHz 5: 16MHz 6: 20MHz 7: 30MHz 8: 40MHz 9: 50MHz</p> <p>After a change, the next bus transaction uses the new speed.</p> <p>The initial speed setting of the module on exit from the MgmtInit state is always 1MHz (MciSpeedConfiguration = 0).</p>	
28	All	-	Reserved[1]	
29-30	All	-	Custom [2]	

1

8.2.7 Module-Level Masks

The host can control which Flags may cause a hardware Interrupt by setting Mask bits described in Table 8-12.

Note: The general behavior of Flags and Masks is specified in section 8.1.4.2.

Table 8-12 Module Level Masks (paged memory modules only) (Lower Memory)

Byte	Bits	Field Name	Field Description	Type
31	7	CdbCmdCompleteMask2	Mask bit for CdbCmdCompleteFlag2	RW Adv.
	6	CdbCmdCompleteMask1	Mask bit for CdbCmdCompleteFlag1	RW Adv.
	5-3	-	Reserved	
	2	DataPathFirmwareErrorMask	Mask bit for DataPathFirmwareErrorFlag	RW Adv.
	1	ModuleFirmwareErrorMask	Mask bit for ModuleFirmwareErrorFlag	RW Adv.
	0	ModuleStateChangedMask	Mask bit for ModuleStateChangedFlag	RW Rqd.
	32	VccMonLowWarningMask	Mask bit for VccMonLowWarningFlag	RW Adv.
	6	VccMonHighWarningMask	Mask bit for VccMonHighWarningFlag	
	5	VccMonLowAlarmMask	Mask bit for VccMonLowAlarmFlag	
	4	VccMonHighAlarmMask	Mask bit for VccMonHighAlarmFlag	
	3	TempMonLowWarningMask	Mask bit for TempMonLowWarningFlag	
	2	TempMonHighWarningMask	Mask bit for TempMonHighWarningFlag	
	1	TempMonLowAlarmMask	Mask bit for TempMonLowAlarmFlag	
	0	TempMonHighAlarmMask	Mask bit for TempMonHighAlarmFlag	
33	7	Aux2MonLowWarningMask	Mask bit for Aux2MonLowWarningFlag	RW Adv.
	6	Aux2MonHighWarningMask	Mask bit for Aux2MonHighWarningFlag	
	5	Aux2MonLowAlarmMask	Mask bit for Aux2MonLowAlarmFlag	
	4	Aux2MonHighAlarmMask	Mask bit for Aux2MonHighAlarmFlag	
	3	Aux1MonLowWarningMask	Mask bit for Aux1MonLowWarningFlag	
	2	Aux1MonHighWarningMask	Mask bit for Aux1MonHighWarningFlag	
	1	Aux1MonLowAlarmMask	Mask bit for Aux1MonLowAlarmFlag	
	0	Aux1MonHighAlarmMask	Mask bit for Aux1MonHighAlarmFlag	
34	7	CustomMonLowWarningMask	Mask bit for CustomMonLowWarningFlag	RW Adv.
	6	CustomMonHighWarningMask	Mask bit for CustomMonHighWarningFlag	
	5	CustomMonLowAlarmMask	Mask bit for CustomMonLowAlarmFlag	
	4	CustomMonHighAlarmMask	Mask bit for CustomMonHighAlarmFlag	
	3	Aux3MonLowWarningMask	Mask bit for Aux3MonLowWarningFlag	
	2	Aux3MonHighWarningMask	Mask bit for Aux3MonHighWarningFlag	
	1	Aux3MonLowAlarmMask	Mask bit for Aux3MonLowAlarmFlag	
	0	Aux3MonHighAlarmMask	Mask bit for Aux3MonHighAlarmFlag	
35	7-0	-	Reserved[1] for Masks	
36	7-0	-	Custom[1] Module level Masks	

8.2.8 CDB Command Status

The CDB command status fields **CdbStatus**<i> provide the status of the most recently triggered CDB command execution or its result, separately for each CDB Instance <i>.

See section 7.2 for a description of the optional CDB feature, and sections 8.23 and 9 for more technical details

When CDB is used in background operation mode (see Table 8-54), the host can read the relevant CDB command status field while a CDB command is executing to obtain its current status.

When CDB is used in foreground operation mode, the host can read the relevant CDB command status field only after the command has completed, as can be determined by Acknowledge polling.

Table 8-13 CdbStatus fields (paged memory modules only) (Lower Memory)

Byte	Bits	Register Name	Register Description	Type
37	7-0	CdbStatus1	Status of the most recent CDB command in CDB instance 1	RO Adv.
38	7-0	CdbStatus2	Status of the most recent CDB command in CDB instance 2	RO Adv.

A CdbStatus field has the following format:

Table 8-14 Bit definitions within CdbStatus fields

Bits	Field Name	Field Description												
7	CdbIsBusy	Bool: CdbIsBusy status bit indicates whether the module is still busy, or idle and ready to accept a new CDB command 0b: Module idle, host can write 1b: Module busy, host needs to wait												
6	CdbHasFailed	Bool: CdbHasFailed bit indicates if there was a failure, after the module has completed execution of the last CDB command 0b: Last triggered CDB command completed successfully 1b: Last triggered CDB command failed												
5-0	CdbCommandResult	<p>The CdbCommandResult field provides more detailed classification for each of the three coarse query results that are encoded by the pair of bit 7 (CdbIsBusy) and bit 6 (CdbHasFailed)</p> <table> <thead> <tr> <th>Coarse Status</th> <th>CdbIsBusy</th> <th>CdbHasFailed</th> </tr> </thead> <tbody> <tr> <td>IN PROGRESS</td> <td>1</td> <td>X (don't care)</td> </tr> <tr> <td>SUCCESS</td> <td>0</td> <td>0</td> </tr> <tr> <td>FAILED:</td> <td>0</td> <td>1</td> </tr> </tbody> </table> <p>The interpretation of CdbCommandResult therefore depends on the coarse status as follows:</p> <p>IN PROGRESS</p> <ul style="list-style-type: none"> 00h=Reserved 01h=Command is captured but not processed 02h=Command checking is in progress 03h=Command execution is in progress 04h-2Fh=Reserved 30h-3Fh=Custom <p>SUCCESS</p> <ul style="list-style-type: none"> 00h=Reserved 01h=Command completed successfully 02h=Reserved 03h=Previous CMD was ABORTED by CMD Abort 04h-1Fh=Reserved 20h-2Fh=Reserved 30h-3Fh=Custom <p>FAILED</p> <ul style="list-style-type: none"> 00h=Reserved 01h=CMDID unknown 02h=Parameter range error or parameter not supported 03h=Previous CMD was not properly ABORTED (by CMD Abort) 	Coarse Status	CdbIsBusy	CdbHasFailed	IN PROGRESS	1	X (don't care)	SUCCESS	0	0	FAILED:	0	1
Coarse Status	CdbIsBusy	CdbHasFailed												
IN PROGRESS	1	X (don't care)												
SUCCESS	0	0												
FAILED:	0	1												

Bits	Field Name	Field Description
		04h=Command checking time out 05h=CdbChkCode Error 06h=Password related error (command specific meaning) 07h=Command not compatible with operating status 08h-0Fh=Reserved for STS command checking error 10h-1Fh=Reserved 20h-2Fh=For individual STS command or task error 30h-3Fh=Custom

1

2 The module sets the **CdbIsBusy** bit when a CDB command is triggered on the respective CDB instance.

3 The module clears the **CdbIsBusy** bit on successful or unsuccessful completion of a CDB command, after
4 updating the fields describing the result of the completed command

- 5 • The **CdbHasFailed** bit reports if the command has failed
6 • The **CdbCommandResult** field provides a detailed classification

7 After command completion, the module does not change the **CdbStatus** byte of a CDB Instance, until the next
8 CDB command is triggered for that CDB instance.

9
10 *Note: CdbStatus fields are cleared on exit from module state MgmtInit (see section 6.3.2.5.3).*

11

8.2.9 Module Active Firmware Version

The Bytes described in Table 8-15 allow a module to report the firmware major and minor revision for the active (i.e. currently running) firmware, or to indicate that no firmware is running.

Note: Module firmware may consist of multiple firmware components (program images, non-volatile data). It is strongly recommended that the firmware version identifies the aggregate (or bundle) of all firmware elements that can be updated by the vendor or by the host.

Note: For modules supporting firmware update, it is also strongly recommended that the firmware major and minor revision numbers together with build number (available by CDB query) uniquely specifies exactly one aggregate firmware configuration. Firmware aggregates that differ in only a single component should never have the same version identification (major, minor, build).

Note: The identification of temporary firmware component versions for lab test or debug, or identification of individual firmware components is not in the scope of this specification. However, the CDB firmware query command supports additional fields which could be used to mark such temporary firmware aggregates.

Reporting firmware version information is generally required and independent of whether a module supports firmware update by any method.

Content and meaning of firmware major and minor revision information reported are vendor dependent, but the format of both fields is defined to be integer. The encoding of the major and minor revision fields is:

- Major Revision = 0 and Minor Revision = 0 indicates that a module does not have any firmware.
- Major Revision = FFh and Minor Revision = FFh indicates that the active firmware image is invalid¹.
- All other Major and Minor Revision combinations indicate the active firmware version.

Table 8-15 Module Active Firmware Version (Lower Memory)

Byte	Bits	Register Name	Register Description	Type
39	7-0	ModuleActiveFirmwareMajorRevision	U8 Numeric representation of the module's active firmware major revision	RO Rqd.
40	7-0	ModuleActiveFirmwareMinorRevision	U8 Numeric representation of the module's active firmware minor revision	RO Rqd.

Note: A module supporting a second image reports version information for the inactive image in the Bytes described in section 8.4.1 with the same encoding.

Note: A module may or may not support a second (inactive) firmware image.

Note: Flat memory modules with firmware may also report their active firmware version in these Bytes.

Note: Modules that support a proprietary vendor specific firmware update method (instead of CDB as described later) shall still report the active firmware version in these bytes.

8.2.10 Module Fault Information

The optional ModuleFaultCause Byte describes the reason for the module having entered the ModuleFault state.

Note: In certain fault reaction scenarios the management communication interface may not be available.

Table 8-16 Fault Information (paged memory modules only) (Lower Memory)

Byte	Bits	Register Name	Register Description	Type
41	7-0	ModuleFaultCause	Reason of entering the ModuleFault state 0: No Fault detected (or field not supported) 1: TEC runaway 2: Data memory corrupted 3: Program memory corrupted 4: Transmitter fault 5: Receiver fault 6: Temperature related fault 7-31: Reserved (fault codes) 32-63: Custom (fault codes) 64-255: Reserved (general)	RO Opt.

¹ Such an indication could e.g. be given by a fixed bootloader that is independent of loaded firmware.

1 **8.2.11 Miscellaneous Status Information**

2 **Table 8-17 Miscellaneous Status Information (Lower Memory)**

Byte	Bits	Field Name	Field Description	Type
42	7-4	-	Reserved	RO Rqd.
	3-0	PasswordCmdResult	Results status for the most recent password entry or password change entry in 00h:118-125 0000b: not supported (legacy before CMIS 5.3) 0001b: module password entry or change has been accepted 0010b: host password entry or change has been accepted 0011b: password entry not accepted 1000b: password validation in progress Advertisement: 01h:251.3-2	
43	7-0	-	Reserved	RO Rqd.
44	7-0	-	Reserved	
45	7-0	-	Reserved	

4 **8.2.12 Extended Module Information**

5 **Table 8-18 Extended Module Information (Lower Memory)**

Byte	Bits	Field Name	Field Description	Type
56	7-0	CmisSmSupport	CMIS State Machines (SM) dynamics supported: 0: undefined transceiver or muxceiver (legacy) 1: no state machine supported (e.g. passive cable) 2: MSM only (Resource Module or fixed transceiver) 3: MSM + DPSM (programmable transceiver) 4: MSM + DPSM + NPSM (Muxceiver) 5-FF: Reserved <i>Note: When undefined (prior to CMIS 5.3), the type of module is implicit but can usually be determined from MemoryModel (00h:2) and from other advertisements.</i>	RO Rqd.
57	7-0	ModuleFunctionType	0: Transmission Module 1: ELSFP Resource Module [10] 2-127: Reserved 128-255: Custom	RO Rqd.
58-59	7-0	-	Reserved [2]	RO Opt.
60	7-4	-	Reserved <i>Note: This reservation is conservative since SFF-8024 theoretically allows for 255 transceiver sub-type codes (while only a few are expected to be used).</i>	RO Rqd.
	3-0	SFF8024ModuleSubtype	SFF8024ModuleSubtype is an SFF-8024 module subtype identifier, with meaning defined in the Transceiver Sub-type codes table in [5]. It allows the reader to infer information regarding variants of physical form factor or thermal characteristics, depending on the SFF8024Identifier value.	
61	7-3	-	Reserved	RO Rqd.
	1-0	SFF8024FiberFaceType	Physical contact specification at optical connectors (assumed to be the same for all optical interfaces) 00b: Unknown, unspecified, or not applicable 01b: PC/UPC (Physical/Ultra Physical contact) 10b: APC (Angled Physical Contact) 11b: Reserved <i>Note: See 00h:203 for connector type information</i>	

62	7-0	LowPowerRestrictions	<p>The module advertises if certain management functions, despite being supported in ModuleReady state, are not fully available on entry to ModuleLowPwr state, due to LowPower mode implementation restrictions.</p> <p>See Table 8-19 for the management function groups that may be restricted for that reason.</p> <p><i>Note: Because restrictions are advertised here, the encoding is somewhat unusual as follows:</i></p> <p>0: management functionality fully supported (unrestricted) 1: management functionality not fully supported (restricted)</p> <p><i>Note: The purpose of this advertisement is to provide full transparency to the host if a module can or cannot fully support certain non-traffic related but possibly power-hungry management functions, because the implementation would violate the power dissipation constraint of Low Power mode. For management functions with result feedback (such as CDB commands) a host may also use a try-and-err approach to find out if the function is supported.</i></p>
63	7-0	-	Reserved

Table 8-19 Low Power Restrictions Byte

Bits	Field Name	Field Description
7	ValidityIndication	0b: Advertisements in this byte are uncertain (legacy modules only) 1b: Advertisements in this byte are valid
6-4	-	Reserved
3	CdbQueryUnsupported	1b: CDB query commands are not generally supported in ModuleLowPwr state
2	CdbFwQueryUnsupported	1b: CDB query commands for FW management are not generally supported in ModuleLowPwr state
1	CdbCmdsUnsupported	1b: CDB commands that can cause state changes in the module are not generally supported in ModuleLowPwr state
0	CdbFwCmdsUnsupported	1b: CDB commands that can cause FW state changes in the module are not generally supported in ModuleLowPwr state

8.2.13 Applications Advertising

Bytes 00h:86-117 (see Table 8-21) provide space for an array of the first four of five bytes (see Table 8-22) of eight **Application Descriptors** (see section 6.2.1.4).

Note: The module advertises supported applications in Application Descriptors. See sections 6.2.1.4 and 7.6.4 for more information about advertising system interface applications and client encapsulation, respectively. Extensions for advertising client encapsulation applications are described in section 8.15.5

Note: The fifth byte (MediaLaneAssignmentOptions) of each Application Descriptor is stored separately in Byte array 01h:176-190 (see Table 8-58).

Note: When more than eight Applications need to be advertised, additional Application Descriptors can be stored in memory described in Table 8-59 and in the second half of memory described in Table 8-58. Even more Applications can be advertised by using so called Normalized Application Descriptors (see section 6.2.1.4.2)

All modules advertise **at least one** Application (so the first Application Descriptor is always used).

*Note: For passive copper cables it is recommended that modules advertise **only one Application Descriptor**, indicating the Host Interface ID for the fastest, single lane data rate supported by the assembly.*

The module indicates the end of the list of valid Application Descriptors by setting the HostInterfaceID field of the first unused Application Descriptor to a value of FFh.

Application Descriptor Fields

The **HostInterfaceID** indicates the interface standard describing the Host Interface of the Application. HostInterfaceID values and their meaning are specified in [5].

The **MediaInterfaceID** indicates the interface standard describing the Media Interface of the Application. MediaInterfaceID values and their meaning are specified in [5], depending on the module's Media Type as indicated in the **MediaType** field (Byte 00h:85).

The **MediaType** indicates the particular Media Interface Type table in [5] that applies to the module and hence the interpretation of MediaInterfaceID values. MediaType values are specified in Table 8-20.

The **HostLaneCount** and **MediaLaneCount** fields specify the number of lanes either explicitly (nonzero value) or by implicit reference, via the Interface ID, to the relevant interface specification (zero value).

The **HostLaneAssignmentOptions** register specifies which lane groups can be used for a Data Path carrying the advertised Application. Bits 0-7 form a bit map corresponding to Host Lanes 1-8. A bit value of 1 indicates that the lane group of the advertised Application can begin on the corresponding host lane. See section 6.2.1.4 for a more detailed description.

Note: The MediaLaneAssignmentOptions parts of the Application Descriptor are stored separately.

All parts of the Application Descriptor are linked by a number known as the AppSel Code, which is simply the sequential position number of the Application Descriptor in any of the memory locations storing (parts of) the Application Descriptor array.

Table 8-20 Media Type Encodings

Code	Media Type	Associated Interface ID Table
00h	Undefined	None, not applicable
01h	Optical Interfaces: MMF	[5] Table 4-6 "MMF media interface IDs"
02h	Optical Interfaces: SMF	[5] Table 4-7 "SMF media interface IDs"
03h	Passive and Linear Active Copper Cables	[5] Table 4-8 "Passive and Linear Active Copper Cable and Passive Loopback media interface codes"
04h	Active Cables	[5] Table 4-9 "Limiting and Retimed Active Cable assembly and Active Loopback media interface codes"
05h	BASE-T	[5] Table 4-10 "BASE-T media interface codes"
06h-3Fh	-	Reserved
40h-8Fh	-	Custom
90h-FFh	-	Reserved

Table 8-21 Media Type Register (Lower Memory)

Byte	Bits	Register Name	Register Description	Type
85	7-0	MediaType	The MediaType field defines the interpretation of MediaInterfaceID values in the following Application Descriptors. See Table 8-20 for the MediaType encoding.	RO Rqd.

2

3

Table 8-22 Format of Application Descriptor Bytes 1-4

Byte Offset	Bits	Field Name	Field Description	Type
0	7-0	HostInterfaceID	ID from [5] or FFh to mark as unused (empty Application Descriptor)	RO Rqd.
1	7-0	MediaInterfaceID	ID from a suitable table of IDs in [5] that is identified by the MediaType Byte 00h:85 (see Table 8-20)	
2	7-4	HostLaneCount	0000b: lane count defined by interface ID 0001b: 1 lane, 0010b: 2 lanes 1000b: 8 lanes. 1001b-1111b: reserved	
	3-0	MediaLaneCount		
3	7-0	HostLaneAssignmentOptions	Bits 0-7 form a bit map corresponding to Host Lanes 1-8. A set bit indicates that the Application may begin on the corresponding host lane. <i>Note: Refer to section 6.2.1.4 for more information.</i>	

4 Note: The fifth bytes of these Application Descriptors are stored elsewhere (see Table 8-58)

5 Note: For passive copper cables (MediaType= 03h) it is recommended to set all Application Descriptor fields
6 except the HostInterfaceID field to zero. This informs the host that the HostLaneCount is derived from the
7 HostInterfaceID value and that the MediaInterfaceID is undefined (see [5])

8

9

Table 8-23 Application Descriptor Registers Bytes 1-4 (Lower Memory)

Byte	Bits	Field Name	Register Description	Type
86	7-0	HostInterfaceIDApp1	AppDescriptor1 AppSel 1 (0001b) See Table 8-22	RO Rqd.
87	7-0	MediaInterfaceIDApp1		
88	7-4	HostLaneCountApp1		
	3-0	MediaLaneCountApp1		
89	7-0	HostLaneAssignmentOptionsApp1	AppDescriptor2 AppSel 2 (0010b) See Table 8-22	RO Rqd.
90	7-0	HostInterfaceIDApp2		
91	7-0	MediaInterfaceIDApp2		
92	7-4	HostLaneCountApp2		
	3-0	MediaLaneCountApp2	AppDescriptor3 AppSel 3 (0011b) See Table 8-22	RO Cnd.
93	7-0	HostLaneAssignmentOptionsApp2		
94	7-0	HostInterfaceIDApp3		
95	7-0	MediaInterfaceIDApp3		
96	7-4	HostLaneCountApp3	AppDescriptor4 AppSel 4 (0100b) See Table 8-22	RO Cnd.
	3-0	MediaLaneCountApp3		
97	7-0	HostLaneAssignmentOptionsApp3		
98	7-0	HostInterfaceIDApp4		
99	7-0	MediaInterfaceIDApp4	AppDescriptor5 AppSel 5 (0101b) See Table 8-22	RO Cnd.
100	7-4	HostLaneCountApp4		
	3-0	MediaLaneCountApp4		
101	7-0	HostLaneAssignmentOptionsApp4		
102	7-0	HostInterfaceIDApp5	AppDescriptor5 AppSel 5 (0101b) See Table 8-22	RO Cnd.
103	7-0	MediaInterfaceIDApp5		
104	7-4	HostLaneCountApp5		
	3-0	MediaLaneCountApp5		
105	7-0	HostLaneAssignmentOptionsApp5		

106	7-0	HostInterfaceIDApp6	AppDescriptor6 AppSel 6 (0110b) See Table 8-22	RO Cnd.
107	7-0	MediaInterfaceIDApp6		
108	7-4	HostLaneCountApp6		
	3-0	MediaLaneCountApp6		
109	7-0	HostLaneAssignmentOptionsApp6	AppDescriptor7 AppSel 7 (0111b) See Table 8-22	RO Cnd.
110	7-0	HostInterfaceIDApp7		
111	7-0	MediaInterfaceIDApp7		
112	7-4	HostLaneCountApp7		
	3-0	MediaLaneCountApp7		
113	7-0	HostLaneAssignmentOptionsApp7	AppDescriptor8 AppSel 8 (1000b) See Table 8-22	RO Cnd.
114	7-0	HostInterfaceIDApp8		
115	7-0	MediaInterfaceIDApp8		
116	7-4	HostLaneCountApp8		
	3-0	MediaLaneCountApp8		
117	7-0	HostLaneAssignmentOptionsApp8		

1

8.2.14 Password Entry and Change

The password entry and change facility described in this section provides a standardized **mechanism** to allow password protection of custom data or functionality, which itself are outside of the scope of this specification.

Note: An alternative password entry and password change mechanism providing feedback about success or failure is available via CDB commands CMD 0001h and CMD 0002h (see section 9.3).

Password protected custom functionality and password entry mechanism support is optional (advertised in 01h:151.5-4) but if it is supported, it shall conform to the following specifications.

Password protection of CMIS specified features, data or functionality is **prohibited** for CMIS compliant module, unless explicitly specified¹ or advertised².

The password entry and change facility uses 32-bit passwords (U32 values in Big-Endian format).

A password entry register must be written using a size-matched four-byte WRITE access

Password entry registers are self-clearing; the module maintains the state of password protection internally.

In the password value range, two types of passwords are distinguished by the most significant bit:

A **Host Password** value is in the range of 0000 0000h to 7FFF FFFFh. It can be changed by the host.

A **Module Password** value is in the range of 80000000h to FFFF FFFFh. It **cannot be changed** by the host.

Note: A host password may e.g. be used to protect module instance-specific data (e.g. inventory data) that are under control of the module user.

Note: A module password may e.g. be used by module vendors to grant access to custom features only to those hosts who have been given a certain module password.

The factory default of a Host Password is 0000 1011h.

The factory default of a Module Password is defined by the module vendor.

The host can change a current Host Password value by writing a new Host Password value into the Password Change Entry Area (Bytes 00h:118-121) after the correct current Host Password value has been entered into the Password Entry Area (Bytes 00h:122-125). The effect of writing a Module Password value is undefined.

The result of the last password entry or of the last password change by a WRITE to one of the following registers is eventually indicated in the PasswordCmdResult register 00h:42.3-0. While password validation is still in progress (indicated in PasswordCmdResult) and until the validation result is available, the host should refrain from using password entry registers. The module may also **reject** READ access to the PasswordCmdResult register until the result has been determined and updated by the module within a period not exceeding tWRITE.

Password-protected features are unlocked when a valid password is entered and remain unlocked until either an invalid password is entered or until the module is reinitialized.

¹ The non-volatile module data in Page 03h may be password protected

² In this version of CMIS, there are no such advertisements.

The initial internal state of password protection when the module exits the MgmtInit state shall be locked.

Table 8-24 Password Change Entry (Lower Memory)

Byte	Bits	Register Name	Register Description	Type
118-121	7-0	PasswordChangeEntryArea	U32: new password value	WO/SC Opt.
122-125	7-0	PasswordEntryArea	U32: password value	WO/SC Opt.

8.2.15 Page Mapping (Upper Memory Content Selection)

The **PageMapping** register (00h:126-127) is a **two-byte** register containing a **Page Address** value that determines which register in the internal Management Memory Map is **actually** accessed when the host performs an ACCESS addressing a Byte in Upper Memory (Byte address range 128-255).

Note: This programmable redirection from a fixed Upper Memory address window to a selected Page in the Management Memory Map is referred to as Page Mapping. Note that CMIS is agnostic of how a module actually implements or emulates page mapping.

The **PageMapping** register has two components storing the two components of a **Page Address**:

- The **BankSelect** Byte (00h:126) storing a **Bank Index**
- The **PageSelect** Byte (00h:127) storing a **Page Index**

Note: Beware of confusion – a Page Address is a two-dimensional value that uniquely identifies a Page, whereas a Page Index is a one-dimensional value identifying a set of Pages with the same Bank Index.

The Management Memory Map (see Table 8-1) defines **banked** Pages (with the potential of Bank Support) and **unbanked** Pages (without Bank Support), whereby the Bank Index of unbanked Pages is irrelevant (albeit using a nominal value of 0 is recommended).

Note: See section 8.1 for background on the host accessible Management Memory Map, on the split of host addressable memory into LowerMemory and UpperMemory, and on Banking.

For an arbitrary Page Address change or for just a Bank Index change, a host must write both BankSelect and PageSelect in one WRITE access, even if the PageSelect value does not change. The module does not begin processing the BankSelect value until after the PageSelect register has been written.

For just a Page Index change (mapping another Page in the current Bank), or for mapping an arbitrary unbanked Page to Upper Memory, a host may WRITE only the PageSelect Byte.

Note: In other words, a host needs to write both the BankSelect and the PageSelect register in a single WRITE transaction, except for the case when BankIndex value in the BankSelect register does not change. In this case it is sufficient to WRITE to the PageSelect register.

Note: However, modules may also choose to accept the two registers written in two subsequent WRITE transactions, to work with non-compliant hosts. Note that rainy day scenarios remain unspecified in this case, and the PageMapping Validity assertion described below is compromised.

Table 8-25 Page Mapping Register Components (Lower Memory)

Byte	Bits	Register Name	Field Description	Type
126	7-0	BankSelect	Bank Index of Page mapped to Upper Memory (if applicable). The current BankSelect value determines which Bank of a Page is accessed (for Pages with Banking) when a host ACCESS addresses a Byte in Upper Memory (address 128 through 255). Ignored when the Page indexed by PageSelect is unbanked.	RW Cnd.
127	7-0	PageSelect	Page Index of Page mapped to Upper Memory The current PageSelect value determines which Page (or Page in a Bank) is accessed when a host ACCESS addresses a Byte in Upper Memory (address 128 through 255) <i>Note: The module will clear the PageSelect when needed to prevent mapping an unsupported page.</i>	RWW Cnd.

1 PageMapping Validity

2 The module ensures that the **PageMapping** register always contains a **Page Address** that is **actually**
3 supported by the module: When a host write would result in a not supported Page Address in the PageMapping
4 register, the module **clears the PageSelect** Byte (without clearing the BankSelect Byte), such that the
5 resulting **PageMapping** register selects Page 00h (as the BankSelect value is ignored unbanked Page 00h).

6 *Note: This is a deliberate exception from the usual rule that either host or module modify a field, but not both.*

7 Stale Data Access Prevention during Page Remapping

8 The module may need time to effectively switch the content of Upper Memory (i.e. to effectively map a new
9 Page) after a change of the PageMapping register.

10 While performing the page switch, when an ACCESS to Upper Memory would or could access stale data, the
11 module **rejects** such ACCESS, as described in section 5.2.3, using methods described in Appendix B.

12 For simplicity, the module may at the same time also **reject** ACCESS to Lower Memory.

13 The maximum time allowed for the module to complete a page switch is specified in section 10.2.2 and applies
14 both to supported and to non-supported values, as well as to the case where the register values do not change.

8.3 Page 00h (Administrative Information)

Page 00h contains static read-only module characteristic information.

Page 00h is supported by all modules (including cable assemblies).

Table 8-26 Page 00h Overview

Address	Size (bytes)	Subject Area or Field	Description
128	1	SFF8024IdentifierCopy	Copy of Byte 00h:0
129-144	16	VendorName	Vendor name (ASCII)
145-147	3	VendorOUI	Vendor IEEE company ID
148-163	16	VendorPN	Part number provided by vendor (ASCII)
164-165	2	VendorRev	Revision level for part number provided by vendor (ASCII)
166-181	16	VendorSN	Vendor Serial Number (ASCII)
182-189	8	DateCode	Manufacturing Date Code (ASCII)
190-199	10	CLEICode	Common Language Equipment Identification Code (ASCII)
200-201	2	ModulePowerCharacteristics	Module power characteristics
202	1	CableAssemblyLinkLength	Cable length (for cable assembly modules only)
203	1	ConnectorType	Connector type of the media interface
204-209	6	Copper Cable Attenuation	Attenuation characteristics (copper ¹ cables only)
210	1	MediaLaneInformation	Supported near end media lanes (all modules)
211	1	Cable Assembly Information	Far end module breakout information (cable assemblies)
212	1	MediaInterfaceTechnology	Information on media side device or cable technology
213-214	2	MCI Related Advertisement	MCI protocol parameters (SPIMCI)
215-220	6	-	Reserved [6]
221	1	-	Custom [1]
222	1	PageChecksum	Page Checksum over bytes 128-221
223-255	33	-	Custom [33] Information (non-volatile)

8.3.1 SFF-8024 Identifier Copy

This field shall contain the same value as Byte 00h:0.

Note: This duplication requirement is maintained for historical reasons and for similarity with predecessor SFF MIS specifications.

Table 8-27 SFF8024IdentifierCopy (Byte 00h:128)

Byte	Bits	Name	Description	Type
128	7-0	SFF8024Identifier	This Byte shall contain the same value as the SFF8024 Identifier Byte 00h:0. See Byte 00h:0 for a description of its meaning.	RO Rqd.

8.3.2 Vendor Information

Table 8-28 Vendor Information (Page 00h)

Bytes	Length (bytes)	Register Name	Register Description	Type
129-144	16	VendorName	Vendor name (ASCII)	RO Rqd.
145-147	3	VendorOUI	Vendor IEEE company ID	RO Rqd.
148-163	16	VendorPN	Part number provided by vendor (ASCII)	RO Rqd.
164-165	2	VendorRev	Revision level for part number provided by vendor (ASCII)	RO Rqd.
166-181	16	VendorSN	Vendor Serial Number (ASCII)	RO Rqd.
182-189	8	DateCode	Manufacturing Date Code (ASCII)	RO Rqd.
190-199	10	CLEICode	Common Language Equipment Identification Code (ASCII)	RO Rqd.

¹ Up to CMIS 5.2 this field was for passive copper cables only.

8.3.2.1 Vendor Name

The **VendorName** is a 16 character read-only field that contains ASCII characters, left aligned and padded on the right with ASCII spaces (20h).

The VendorName shall contain the full name of the corporation, a commonly accepted abbreviation of the name of the corporation, the SCSI company code for the corporation, or the stock exchange code for the corporation. The VendorName may be the original manufacturer of the module or the name of the module reseller. In both cases, the VendorName and VendorOUI (if specified) shall correlate to the same company. At least one of the VendorName or the VendorOUI fields shall contain valid serial number manufacturing data.

8.3.2.2 Vendor Organizationally Unique Identifier

The vendor organizationally unique identifier field (**VendorOUI**) is a 3-byte field that contains the IEEE Company Identifier for the vendor. A value of all zero in the 3-byte field indicates that the vendor's OUI is unspecified.

8.3.2.3 Vendor Part Number

The vendor part number (**VendorPN**) is a 16-byte field that contains ASCII characters, left aligned and padded on the right with ASCII spaces (20h), defining the vendor part number or product name. A value of all zero in the 16-byte field indicates that the vendor part number is unspecified.

8.3.2.4 Vendor Revision Number

The vendor revision number (**VendorRev**) is a 2-byte field that contains ASCII characters, left aligned and padded on the right with ASCII spaces (20h), defining the vendor's product revision number. A value of all zero in the field indicates that the vendor revision number is unspecified.

8.3.2.5 Vendor Serial Number

The vendor serial number (**VendorSN**) is a 16-character field that contains ASCII characters, left aligned and padded on the right with ASCII spaces (20h), defining the vendor's serial number for the Product. A value of all zero in the 16-byte field indicates that the vendor serial number is unspecified.

8.3.2.6 Date Code

The **DateCode** is an 8-byte field that contains the vendor's date code in ASCII characters. The date code is mandatory. The date code shall be in the following format:

Table 8-29 Date Code (Page 00h)

Byte	Bits	Register Name	Description	Type
182-183	All	Year	ASCII two low order digits of year (00=2000)	RO Rqd.
184-185	All	Month	ASCII digits of month (01=Jan through 12=Dec)	RO Rqd.
186-187	All	DayOfMonth	ASCII day of month (01-31)	RO Rqd.
188-189	All	LotCode	ASCII custom lot code, may be blank	RO Opt.

8.3.2.7 CLEI Code

The **CLEI** (Common Language Equipment Identification) code is a 10-byte field that contains the vendor's CLEI code in ASCII characters.

The CLEI code value is optional. If CLEI code value is not supported, a value of all ASCII 20h (spaces) shall be entered.

Table 8-30 CLEI Code (Page 00h)

Byte	Bits	Register Name	Description	Type
190-199	All	CLEICode	Vendor's CLEI Code (ASCII)	RO Opt.

8.3.3 Module Power Characteristics

Module power characteristics are advertised in Bytes 00h:200 and 00h:201 (see Table 8-31).

The **MaxPower** field specifies worst case maximum power consumption over operating conditions and lifetime.

The **ModulePowerClass** field provides a form factor specific classification of the MaxPower value.

Note: The preferred method for the host to verify if the maximum power consumption is acceptable in the host system, before allowing the module to enter High Power Mode, is to use the MaxPower field. See section 6.3.2.4.

Note: The host may also use the ModulePowerClass field to determine if the maximum power consumption is acceptable in the host system, but this indirect and form factor dependent method is not preferred.

Table 8-31 Module Power Class and Max Power (Page 00h)

Byte	Bits	Field Name	Field Description	Type
200	7-5	ModulePowerClass ¹	000: Power class 1 001: Power class 2 010: Power class 3 011: Power class 4 100: Power class 5 101: Power class 6 110: Power class 7 111: Power class 8	RO Rqd.
	4-0	-	Reserved	
201	7-0	MaxPower	Maximum power consumption in multiples of 0.25 W rounded up to the next whole multiple of 0.25 W	RO Rqd.

Note 1: See relevant hardware specification for maximum power allowed in each Power class

8.3.4 Cable Assembly Link Length

The **CableAssemblyLinkLength** Byte 00h:202 advertises the physical interconnect length of cable assemblies (including both passive copper cables and active optical or electrical cables).

Modules with separable optical media shall set the CableAssemblyLinkLength value to 0000 0000b.

A CableAssemblyLinkLength value of 1111 1111b indicates a link length greater than 6300 m.

Table 8-32 Cable Assembly Link Length (Page 00h)

Byte	Bits	Field Name	Field Description	Type
202	7-6	LengthMultiplier	Multiplier for value in bits 5-0: 00: multiplier 0.1 01: multiplier 1 10: multiplier 10 11: multiplier 100	RO Rqd.
	5-0	BaseLength	Link length base value in meters. To calculate actual link length use multiplier in bits 7-6. A value of 0 indicates an undefined Link Length, e.g. when the physical media can be disconnected from the module.	

8.3.5 Media Connector Type

The **ConnectorType** field indicates the connector type for the media side of the module, as defined and maintained in the Connector References section of [5].

Note: The list of connector types includes a special connector type code for copper and optical cables or cable assemblies with non-detachable media, which do not have a media connector.

Table 8-33 Media Connector Type (Page 00h)

Byte	Bits	Register Name	Description	Type
203	7-0	ConnectorType	Type of connector present in the module. See Table 4-3 in [5] for Connector Type codes.	RO Rqd.

8.3.6 Copper Cable Attenuation

These Bytes advertise the externally measurable cable attenuation characteristics for copper cables. For active linear copper cables with host-programmable gain, the characteristics are reported for the 0dB gain setting.

For other modules bytes 204-209 are **reserved**.

Table 8-34 Copper Cable Attenuation (Page 00h)

Byte	Bits	Register Name	Register Description	Type
204	7-0	AttenuationAt5GHz	U8 cable attenuation at 5 GHz in 1 dB increments	RO Cnd.
205	7-0	AttenuationAt7GHz	U8 cable attenuation at 7 GHz in 1 dB increments	RO Cnd.
206	7-0	AttenuationAt12p9GHz	U8 cable attenuation at 12.9 GHz in 1 dB increments	RO Cnd.
207	7-0	AttenuationAt25p8GHz	U8 cable attenuation at 25.8 GHz in 1 dB increments	RO Cnd.
208	7-0	AttenuationAt53p1GHz	U8 cable attenuation at 53.125 GHz in 1 dB increments. A value of 0 dB indicates that this characteristic is not available (not relevant or otherwise unknown)	RO Cnd.
209	All	-	Reserved	RO

Note: As per [19], when the module advertises itself as a PCIe application, the cable attenuation fields above are reported for frequencies 2.5, 4.0, 8.0, 16.0, 32.0 GHz, instead of 5.0, 7.0, 2.9, 25.8 and 53.125 GHz.

8.3.7 Media Lane Information

The **MediaLaneUnsupported** Byte (see Table 8-35) indicates which Media Lanes are not supported.

Note: An optical media lane may represent fibers or WDM wavelengths.

Note: This lane related register is on a non-banked page. Modules with more than 8 host lanes can therefore not unambiguously advertise unsupported media lanes.

Note: The MediaLaneUnsupported Byte (see Table 8-35) is especially important for cable assemblies, where it describes the near end media lanes of the assembly, while the FarEndConfiguration Byte(see Table 8-36) describes how these lanes belong to different modules at the far end (in breakout cables).

Table 8-35 Media Lane Information (Page 00h)

Byte	Bits	Field Name	Field Description	Type
210	7	MediaLaneUnsupportedLane8	Bool: MediaLaneUnsupportedLane<i>	RO Cnd.
	6	MediaLaneUnsupportedLane7	0b: Media Lane <i> supported 1b: Media Lane <i> not supported	
	5	MediaLaneUnsupportedLane6		
	4	MediaLaneUnsupportedLane5		
	3	MediaLaneUnsupportedLane4		
	2	MediaLaneUnsupportedLane3		
	1	MediaLaneUnsupportedLane2		
	0	MediaLaneUnsupportedLane1	Condition: Module supports at most 8 host lanes. Otherwise, byte is reserved (0 valued). Hosts should ignore this byte when the module supports more than 8 host lanes.	

8.3.8 Cable Assembly Lane Breakout Information

For Cable Assemblies, the interconnect media are not detachable, i.e. they are fixed attached to both near end and far end modules. With breakout topologies, the far end module granularity may be finer than that of the host managed near end module: lane groups of the near end module may break out to several far end modules. By far the most common breakout type is a homogeneous 1:N breakout where all N far end modules are of the same type, but for 8 lane modules also heterogeneous breakouts can be specified (see below).

The **FarEndConfiguration** Byte (see Table 8-38) indicates a cable assembly's fixed grouping of lanes as connected to discrete far end modules.

For modules with detachable media connectors the FarEndConfiguration byte is cleared.

Table 8-36 Cable Assembly Information (Page 00h)

Byte	Bits	Field Name	Field Description	Type
211	7-5	-	Reserved	RO
	4-0	FarEndConfiguration	Configuration of the far end module breakout. See Table 8-38 for configuration codes	RO Cnd.

Table 8-37 defines the generic FarEndConfiguration values for each possible far end **homogeneous** lane group breakout configuration into 1-lane, 2-lane, 4-lane, 8-lane, and 16-lane modules.

Table 8-37 Far End Configurations for Homogeneous Far End Breakout (Page 00h)

Code ¹	Binary	Homogeneous Far End Breakout into
0	00000b	Undefined. Module with detachable media.
1	00001b	1-lane modules (connectors)
12	01100b	2-lane modules (connectors)
3	00011b	4-lane modules (connectors)
2	00010b	8-lane modules (connectors)
27	11011b	16-lane modules (connectors)
28	11100b	Reserved
29	11101b	
30	11110b	Custom
31	11111b	

Table 8-38 defines a configuration code for each possible partitioning of 8 near end lanes into lane groups (of various group sizes) identified by letters a to h (depending on the lowest lane number in the group), each of which might be connected to a discrete far end module.

Unique letters distinguish both the possible lane groups and the connected discrete far end modules. Note that the discrete far end modules may or may not be of the same module type.

For modules with less than 8 lanes, the remaining columns are ignored. The lane groups must not straddle into the ignored columns.

For modules with more than 8 lanes, the topology defined for 8 lanes **repeats** in each group of 8 lanes².

Table 8-38 Far End Configurations for up to 8 Near End Lanes (Page 00h)

Far End Cable Lane Groups Advertising Codes								
Config Code		Near End Host Lane Number						
Decimal	Binary	1	2	3	4	5	6	7
0	00000b	Undefined. Module with detachable media.						
1	00001b	a	b	c	d	e	f	g
2	00010b	a	a	a	a	a	a	a
3	00011b	a	a	a	a	e	e	e
4	00100b	a	b	c	d	e	e	e

¹ These codes representing homogeneous breakout are found boldfaced in Table 8-38

² This is also the background for Table 8-37

5	00101b	a	b	c	c	e	e	e	e
6	00110b	a	a	c	d	e	e	e	e
7	00111b	a	a	c	c	e	e	e	e
8	01000b	a	a	a	a	e	f	g	h
9	01001b	a	a	a	a	e	f	g	g
10	01010b	a	a	a	a	e	e	g	h
11	01011b	a	a	a	a	e	e	g	g
12	01100b	a	a	c	c	e	e	g	g
13	01101b	a	b	c	c	e	e	g	g
14	01110b	a	a	c	d	e	e	g	g
15	01111b	a	b	c	d	e	e	g	g
16	10000b	a	a	c	c	e	f	g	g
17	10001b	a	b	c	c	e	f	g	g
18	10010b	a	a	c	d	e	f	g	g
19	10011b	a	b	c	d	e	f	g	g
20	10100b	a	a	c	c	e	e	g	h
21	10101b	a	b	c	c	e	e	g	h
22	10110b	a	a	c	d	e	e	g	h
23	10111b	a	b	c	d	e	e	g	h
24	11000b	a	a	c	c	e	f	g	h
25	11001b	a	b	c	c	e	f	g	h
26	11010b	a	a	c	d	e	f	g	h
27	11011b	Far End break out with 16-lane connector(s)							
28	11100b	Reserved							
29	11101b								
30	11110b								
31	11111b	Custom							

8.3.9 Media Interface Technology

The **MediaInterfaceTechnology** Byte 00h:212 classifies the media interface device or cable technology, using the encodings in Table 8-40.

An active optical cable may be distinguished from a separable module with detachable media by examining the **ConnectorType** field Byte 00h:203 (see section 8.3.5), with values defined in the Connector References section of [5].

Table 8-39 Media Connector Type (Page 00h)

Byte	Bits	Register Name	Register Description	Type
212	7-0	MediaInterfaceTechnology	Media Interface Technology as per Table 8-40	RO Rqd.

Table 8-40 Media Interface Technology encodings

Code	Description of physical device
00h	850 nm VCSEL
01h	1310 nm VCSEL
02h	1550 nm VCSEL
03h	1310 nm FP laser
04h	1310 nm DFB laser
05h	1550 nm DFB laser
06h	1310 nm EML
07h	1550 nm EML
08h	Others

09h	1490 nm DFB laser
0Ah	Copper cable, passive, unequalized
0Bh	Copper cable, passive, equalized
0Ch	Copper cable with near and far end limiting active equalizers
0Dh	Copper cable with far end limiting active equalizers
0Eh	Copper cable with near end limiting active equalizers
0Fh	Copper cable with linear active equalizers (deprecated, do not use for new designs)
10h	C-band tunable laser
11h	L-band tunable laser
12h	Copper cable with near end and far end linear active equalizers
13h	Copper cable with far end linear active equalizers
14h	Copper cable with near end linear active equalizers
15h-FFh	Reserved

8.3.10 MCI Related Advertisements

MCI Related Advertisements are relevant only for SPIMCI (see section B.3).

The module can either advertise a fixed number of flow control dummy bytes, or it can advertise a required duration of the flow control phase, from which the number of dummy bytes sent during the flow control phase is computed, depending on the currently selected MCI clock speed.

Table 8-41 MCI Related Advertisements (Page 00h)

Byte	Bits	Register Name	Register Description	Type
213	7	MciFlowControlDuration Encoding	Defines the method to determine the duration of the Flow Control Phase 0b: static (given as number of bytes) 1b: speed-dependent (given as a duration)	RO Cnd.
213	6-0	MciFlowControlDuration	Used by both host and module to compute the number <i>N</i> of dummy bytes sent by the host during the Flow Control phase, in encoding dependent manner: MciFlowControlDurationEncoding static : $N = 2 + \text{MciFlowControlDuration}$ MciFlowControlDurationEncoding speed-dependent : $D = 0.2\mu\text{s} \cdot \text{MciFlowControlDuration}$ $N = \max(2, \text{ceil}(D \cdot \text{SpiSpeedInMHz}/8))$ where the 'ceiling' function <i>ceil</i> (.) returns the smallest integer that is not smaller than its argument and where the SPI speed SpiSpeedInMHz is to be determined from decoding MciSpeedConfiguration (0h:27:3-0)	RO Cnd.
214	7-0	-	Reserved	

8.3.11 Page 00h Page Checksum (required)

The page checksum is a one-byte code that can be used to verify that the read-only static data on Page 00h is valid. The page checksum value shall be the low order 8 bits of the arithmetic sum of all byte values from byte 128 to byte 221, inclusive.

8.3.12 Custom Info (non-volatile)

Bytes 223-255 are allocated in **non-volatile** storage for information provided by the original manufacturer of the module or the module reseller. This information persists across module reset and power cycle. The contents of this area are not defined by this specification.

8.4 Page 01h (Advertising)

Page 01h is an optional Page containing advertising fields for properties of paged memory modules.

The module advertises support of Page 01h in the MemoryModel Bit 00h:2.7.

Note: Page 01h is mandatory for paged memory modules.

All fields on Page 01h are read-only and static.

See the following subsections for detailed information on the subject areas listed in Table 8-42

Table 8-42 Page 01h Overview

Byte	Size (bytes)	Subject Area	Description
128-131	4	Inactive Firmware and Hardware revisions	Inactive FW revision and HW revision
132-137	6	Supported Link Length	Supported lengths of various fiber media
138-141	4	Wavelength Information	(for single wavelength modules)
142	1	Supported Pages	See subsection 8.4.4
143-144	2	Durations Advertisements	See subsection 8.4.5
145-154	10	Module Characteristics Advertisement	See subsection 8.4.6
155-156	2	Supported Controls	See subsection 8.4.7
157-158	2	Supported Flags	See subsection 8.4.8
159-160	2	Supported Monitors	See subsection 8.4.9
161-162	2	Supported Signal Integrity Controls	See subsection 8.4.10
163-166	4	Supported CDB Functionality	See subsection 8.4.11
167-169	3	Additional Durations Advertisements	See subsection 8.4.12
170-174	5	-	Reserved[5]
175	1	Normalized Application Descriptors Support	See subsection 8.4.13
176-190	15	Media Lane Advertising	See subsection 8.4.14
191-222	32	-	Custom[32]
223-250	28	Additional Application Descriptors	See subsection 8.4.15
251-254	4	Miscellaneous Advertisements	See subsection 8.4.16
255	1	Page Checksum	Page Checksum of bytes 130-254 ¹

Note 1: The firmware version bytes 128-129 are intentionally excluded from the Page Checksum to avoid requiring a Memory Map update when firmware is updated.

8.4.1 Inactive Firmware and Hardware Revisions

Table 8-43 describes the fields for reporting the module's inactive firmware revision (if any) and the module's hardware revision.

Note: The active firmware revision is reported in the Module Active Firmware Version fields (see section 8.2.9) and the overall module revision number is reported in the Vendor Revision Number field (see section 8.3.2.4).

The inactive firmware is that firmware stored that is not currently executing, when the module supports a second firmware image.

Note: This inactive firmware image may be an alternate or backup firmware image, or a new firmware image downloaded but not yet activated.

The ModuleInactiveFirmwareMajorRevision and ModuleInactiveFirmwareMinorRevision fields contain numbers with the same encoding as the ModuleActiveFirmwareMajorRevision and ModuleActiveFirmwareMinorRevision fields (see section 8.2.9): a module without inactive firmware clears these fields.

Bytes 01h:128-129 are not included in the Page 01h Page Checksum as these bytes may change dynamically for modules that support switching firmware version between multiple images during firmware updates.

The numeric ModuleHardwareMajorRevision and ModuleHardwareMinorRevision fields contain version numbers. These two fields are included in the Page 01h Page Checksum.

Table 8-43 Module Inactive Firmware and Hardware Revisions (Page 01h)

Byte	Bits	Register Name	Register Description	Type
128	7-0	ModuleInactiveFirmwareMajorRevision	U8 Numeric representation of module inactive firmware major revision	RO Rqd.
129	7-0	ModuleInactiveFirmwareMinorRevision	U8 Numeric representation of module inactive firmware minor revision	RO Rqd.
130	7-0	ModuleHardwareMajorRevision	U8 Numeric representation of module hardware major revision	RO Rqd.
131	7-0	ModuleHardwareMinorRevision	U8 Numeric representation of module hardware minor revision	RO Rqd.

Note: Modules that support a proprietary vendor specific firmware update method with two firmware images (instead of CDB as described later) shall still report inactive firmware version in bytes 01h:128-129.

8.4.2 Supported Link Length Advertisement

The Bytes described in Table 8-44 advertise the maximum supported fiber media length for each type of fiber media at the maximum module-supported bit rate for modules with a separable optical media interface. Unsupported media types shall be populated with zeroes.

Active optical cables shall populate the fields in this table with zeroes and instead report their actual length using the fields in Table 8-32.

Table 8-44 Supported Fiber Link Length (Page 01h)

Byte	Bits	Name	Description	Type
132	7-6	LengthMultiplierSMF	Link length multiplier for SMF fiber 00: 0.1 (0.1 to 6.3 km) 01: 1 (1 to 63 km) 10: 10 (10 to 630 km) 11: multiplier is specified in 01h:137	RO Rqd.
	5-0	BaseLengthSMF	Base link length for SMF fiber in km. Must be multiplied by multiplier defined in bits 7-6 to calculate actual link length.	RO Rqd.
133	7-0	LengthOM5	Link length supported for OM5 fiber, units of 2 m (2 to 510 m)	RO Rqd.
134	7-0	LengthOM4	Link length supported for OM4 fiber, units of 2 m (2 to 510 m)	RO Rqd.
135	7-0	LengthOM3	Link length supported for EBW 50/125 µm fiber (OM3), units of 2m (2 to 510 m)	RO Rqd.
136	7-0	LengthOM2	Link length supported for 50/125 µm fiber (OM2), units of 1m (1 to 255 m)	RO Rqd.
137	7-6	LengthMultiplierSMF2	Link length multiplier for SMF fiber (if 01h:132 = 11) 00: 50 (50 to 3150 km) 01: 100 (100 to 6300 km) 10: 200 (200 to 12600 km) 11: 500 (500 to 31500 km)	RO Rqd.
	5-0	-	Reserved	

The link length supported for SMF fiber specifies the link length that is supported by the device while operating in compliance with the applicable standards using single mode fiber. The supported link length is as specified in [12]. The value is in units of kilometers.

The link length supported for OM5 fiber specifies the link length that is supported by the device while operating in compliance with the applicable standards using 4700 MHz*km (850 nm) and 2470 MHz*km (953 nm) extended bandwidth 50 micron core multimode fiber. The value is in units of two meters.

The link length supported for OM4 fiber specifies the link length that is supported by the device while operating in compliance with the applicable standards using 4700 MHz*km (850 nm) extended bandwidth 50 micron core multimode fiber. The value is in units of two meters.

The link length supported for OM3 fiber specifies the link length that is supported by the device while operating in compliance with the applicable standards using 2000 MHz*km (850 nm) extended bandwidth 50 micron core multimode fiber. The value is in units of two meters.

The link length supported for OM2 fiber specifies the link length that is supported by the device while operating in compliance with the applicable standards using 500 MHz*km (850 nm and 1310 nm) 50 micron multi-mode fiber. The value is in units of one meter.

8.4.3 Wavelength Information

The **NominalWavelength** and **WavelengthTolerance** fields are defined for **single wavelength** modules.

Note: Nominal wavelength and wavelength tolerance are attributes of optical media lanes, whereas NominalWavelength and WavelengthTolerance are module level fields.

Note: All optical modules, whether single or multiple wavelengths, advertise their supported media specifications indirectly via the MediaInterfaceID of the relevant Application. In many of these Applications, fixed nominal wavelengths and maximum tolerances are known from the associated standard.

A single wavelength module with implicitly defined fixed wavelength and tolerance reports these specified values or more specific actual values (e.g. tighter tolerance).

A single wavelength module with (directly or indirectly) programmable wavelength reports actual nominal wavelength and actual wavelength tolerance.

A multi-wavelength module may optionally provide wavelength information, either for one of the wavelengths or for the entire wavelength range (as nominal center wavelength and overall tolerance). Since the interpretation is not uniquely defined, a host may ignore this field for multi-wavelength modules and use the advertised Application to determine module capabilities.

Table 8-45 Wavelength Information (Page 01h)

Byte	Bits	Register Name	Register Description	Type
138-139	7-0	NominalWavelength	U16 nominal transmitter output wavelength for a single wavelength module at room temperature in units of 0.05nm	RO Cnd.
140-141	7-0	WavelengthTolerance	U16 wavelength tolerance tol as the worst case +/- tol range around the NominalWavelength under all normal operating conditions in units of 0.005nm	RO Cnd.

Example 1 (Single Wavelength Module):

ITU-T Grid Wavelength = 1534.25 nm with 0.236 nm Tolerance

Nominal Wavelength = 1534.25 nm, represented as U16: (1534.25 nm * 20) = 30685

Wavelength Tolerance = 0.236 nm, represented as U16: (0.236 nm * 200) = 47

Example 2 (Multi-Wavelength Module):

100GBASE-LR4 Wavelength Range = 1294.53 to 1310.19 nm

Nominal Wavelength = 1302.36 nm, represented as U16: (1302.36 nm * 20) = 26047

Wavelength Tolerance = 7.83 nm, represented as U16: (7.83 nm * 200) = 1566

1 8.4.4 Supported Pages Advertisement

2 **Table 8-46 Supported Pages Advertising (Page 01h)**

Byte	Bit	Field Name	Field Description	Type
142	7	NetworkPathPagesSupported	Page 16h and NP-related parts of Page 17h supported	RO Rqd.
	6	VDMPagesSupported	VDM Pages 20h-2Fh (partially) supported (advertisement details in Page 2Fh)	
	5	DiagnosticPagesSupported	Banked Pages 13h-14h supported	
	4	CoherentPagesSupported	Banked Pages 30h-4Fh (partially) supported (advertisement details are specified in [7])	
	3	CmisFfSupported	Form Factor specific functionality as specified in the CMIS-FF supplement [8] is supported on Page 05h .	
	2	Page03hSupported	User Page 03h supported	
	1-0	BanksSupported ^{1, 2}	Banks supported for those Pages in range 10h-2Fh that are lane-banked by groups of 8 lanes 00b: Bank 0 supported (8 lanes) 01b: Banks 0 and 1 supported (16 lanes) 10b: Banks 0-3 supported (32 lanes) 11b: reserved	

3 Note 1: The response to a host using an invalid Bank Page combination is defined in section 8.2.15.

4 Note 2: Support of Pages may also be derived from other advertisements, see e.g. section 8.4.13.

5 8.4.5 Durations Advertisement

6 The **ModSelWaitTime** numerical value is represented in a special floating-point format by the two fields
ModSelWaitTimeMantissa and **ModSelWaitTimeExponent** and defines both the required **setup time**
7 (for the ModSel signal after the host asserts ModSel and before the start of an MCI bus transaction) and the
8 required **delay** (after completion of an MCI transaction before the host can deassert the ModSel signal).

9 For example, if the module wait time is 1.6 ms, the mantissa field (bits 4-0) value will be 11001b (25) and the
10 exponent field (bits 7-5) value will be 110b (6) for a result of $25 \cdot 2^6 = 1600$ us. Note that the representation of
11 a ModSelWaitTime value may be ambiguous, e.g. $8 = 1 \cdot 2^3 = 2 \cdot 2^2 = 4 \cdot 2^1 = 8 \cdot 2^0$.

12 The **MaxDurationDPInit** and **MaxDurationDPDeinit** fields are defined so that hosts can determine when
13 something failed in the module during these states, for example a module firmware hang up.

14 These maximum duration values represent worst-case durations across all advertised Applications and all
15 possible combinations of Data Paths.

16 See sections 6.3.3.5 and 6.3.3.7 for details of the DPInit and DPDeinit states, respectively.

17 See section 8.4.12 for other maximum duration advertisements.

18 **Table 8-47 Durations Advertising (Page 01h)**

Byte	Bit	Field Name	Field Description	Type
143	7-5	ModSelWaitTimeExponent e	ModSelWaitTime value represented as $m \cdot 2^e$ in μs 00h: no data available	RO Rqd.
	4-0	ModSelWaitTimeMantissa m		
144	7-4	MaxDurationDPDeinit	Maximum duration of the DPDeinit state (encoded as per Table 8-48)	RO Rqd.
	3-0	MaxDurationDPInit	Maximum duration of the DPInit state (encoded as per Table 8-48)	

1

Table 8-48 State Duration Encoding (Page 01h)

Encoding	Maximum State Duration T_{state}
0000b	$T_{state} < 1 \text{ ms}$
0001b	$1 \text{ ms} \leq T_{state} < 5 \text{ ms}$
0010b	$5 \text{ ms} \leq T_{state} < 10 \text{ ms}$
0011b	$10 \text{ ms} \leq T_{state} < 50 \text{ ms}$
0100b	$50 \text{ ms} \leq T_{state} < 100 \text{ ms}$
0101b	$100 \text{ ms} \leq T_{state} < 500 \text{ ms}$
0110b	$500 \text{ ms} \leq T_{state} < 1 \text{ s}$
0111b	$1 \text{ s} \leq T_{state} < 5 \text{ s}$
1000b	$5 \text{ s} \leq T_{state} < 10 \text{ s}$
1001b	$10 \text{ s} \leq T_{state} < 1 \text{ min}$
1010b	$1 \text{ min} \leq T_{state} < 5 \text{ min}$
1011b	$5 \text{ min} \leq T_{state} < 10 \text{ min}$
1100b	$10 \text{ min} \leq T_{state} < 50 \text{ min}$
1101b	$T_{state} \geq 50 \text{ min}$
1110b	Reserved
1111b	Reserved

2

8.4.6 Module Characteristics Advertisement

The fields in Table 8-49 describe the characteristics of certain **module** properties. Some features are optional. Advertisement of the implementation of optional features is described in sections 8.4.7 and beyond.

*Note: For versatile modules supporting multiple applications or application combinations, some characteristics may be **application dependent** rather than module dependent. In this case, application specific advertisement is necessary, while the following **module characteristics** advertisements describe either the characteristics of the default application or, when meaningful, the **highest capabilities** and **strongest restrictions** across all supported applications.*

A **Tx synchronous group** is defined as a Tx input lane or group of Tx input lanes sourced from the same clock domain. Two different Tx synchronous groups may be sourced from different clock domains. There may be a limit on the maximum permissible clock tolerance between two different Tx synchronous groups, as defined by the industry standard associated with a given Application. Refer to applicable industry standards.

A Tx synchronous group can contain one or more Data Paths, if the Tx lanes on all Data Paths are sourced from the same clock domain and the module takes measures to ensure that active Data Paths continue to operate undisturbed even as other Data Paths (and their associated Tx input lanes) are enabled/disabled by the host.

Table 8-49 Module Characteristics Advertisement (Page 01h)

Byte	Bit	Field Name	Field Description	Type
145	7	CoolingImplemented	0b: Uncooled transmitter device 1b: Cooled transmitter	RO Rqd.
	6-5	TxInputClockingCapabilities	Defines which Tx input lanes must be frequency synchronous 00b: Tx input lanes 1-8 01b: Tx input lanes 1-4 and 5-8 10b: Tx input lanes 1-2, 3-4, 5-6, 7-8 11b: Lanes may be asynchronous in frequency When more than 8 lanes are supported, the above grouping applies to each group of 8 lanes (each Bank)	RO Rqd.
	4	ePPSSupported	Support of the Enhanced Pulse Per Second timing signal [2] 0b/1b: ePPS signal processing not supported/supported	RO Rqd.
	3	TimingPage15hSupported	0b: Timing characteristics (Page 15h) not supported 1b: Timing characteristics (Page 15h) supported	RO Rqd.
	2	Aux3MonObservable	0b: Aux 3 monitor monitors Laser Temperature 1b: Aux 3 monitor monitors Vcc2	RO Adv.
	1	Aux2MonObservable	0b: Aux 2 monitor monitors Laser Temperature 1b: Aux 2 monitor monitors TEC current	RO Adv.
	0	Aux1MonObservable	0b: Aux 1 monitor is custom 1b: Aux 1 monitor monitors TEC current	RO Adv.
	7-0	ModuleTempMax	S8 Maximum allowed module temperature in 1 deg C increments. ModuleTempMax = ModuleTempMin = 0 indicates 'not specified'.	RO Cnd.
147	7-0	ModuleTempMin	S8 Minimum allowed module temperature in 1 deg C increments. ModuleTempMax = ModuleTempMin = 0 indicates 'not specified'.	RO Cnd.
148-149	7-0	PropagationDelay	U16 Propagation delay of a non-separable AOC in multiples of 10 ns rounded to the nearest 10 ns, or zero for 'not specified'.	RO Cnd.
150	7-0	OperatingVoltageMin	U8 Minimum supported module operating voltage, in 20 mV increments (0-5.1 V), or zero for 'not specified'.	RO Cnd.
151	7	OpticalDetectorType	0b: PIN detector 1b: APD detector	RO Rqd.
	6-5	RxOutputEqType	00b: Peak-to-peak (p-p) amplitude stays constant, or not implemented, or no information 01b: Steady-state amplitude stays constant 10b: Average of p-p and steady-state amplitude stays constant 11b: Reserved	
	4	RxPowerMeasurementType	0b: OMA 1b: average power	

Byte	Bit	Field Name	Field Description	Type
	3	RxLOSType	0b: Rx LOS responds to OMA 1b: Rx LOS responds to P_{av} <i>Note: LOS Type depends on interface standards supported</i>	
	2	RxLOSIsFast	0b: Module raises Rx LOS within regular timing limits 1b: Module raises Rx LOS within "fast mode" timing limits Refer to form factor hardware specification for regular and "fast mode" timing limit requirements	
	1	TxDisableIsFast	0b: Module responds to Tx Output Disable with regular timing 1b: Module responds to Tx Output Disable in "fast mode" timing limits Refer to form factor hardware specification for regular and "fast mode" timing limit requirements	
	0	TxDisableIsModuleWide	0b: Tx output disable is controlled per lane 1b: All Tx output lanes disabled when any OutputDisableTx set	
152	7-0	CDRPowerSavedPerLane	U8 Minimum power consumption saved per CDR per lane when placed in CDR bypass, in multiples of 0.01 W rounded up to the next whole multiple of 0.01 W	RO Cnd.
153	7	RxOutputLevel3Supported ¹	0b/1b: Amplitude Code 3 not supported/supported	RO Cnd.
	6	RxOutputLevel2Supported ¹	0b/1b: Amplitude Code 2 not supported/supported	
	5	RxOutputLevel1Supported ¹	0b/1b: Amplitude Code 1 not supported/supported	
	4	RxOutputLevel0Supported ¹	0b/1b: Amplitude Code 0 not supported/supported	
	3-0	TxInputEqMax	Maximum supported value of the Tx Input Equalization control for host-controlled equalizer operation (see section 6.2.5.1)	
154	7-4	RxOutputEqPostCursorMax	Maximum supported value of the Rx Output Eq Post-cursor control (see section 6.2.5.2)	RO Cnd.
	3-0	RxOutputEqPreCursorMax	Maximum supported value of the Rx Output Eq Pre-cursor control (see section 6.2.5.2)	

Note 1: See Table 6-8

1

2

8.4.7 Supported Controls Advertisement

Table 8-50 describes supported module and lane controls and related module functions (see section 8.9.2).

Table 8-50 Supported Controls Advertisement (Page 01h)

Byte	Bit	Field Name	Field Description	Type
155	7	WavelengthIsControllable	0b: No wavelength control 1b: Active wavelength control supported <i>Note: Active Control does not imply tunability</i>	RO Rqd.
	6	TransmitterIsTunable	0b: Transmitter not tunable 1b: Transmitter is tunable (Pages 04h & 12h supported)	
	5-4	SquelchMethodTx	00b: Tx output squelching function is not supported 01b: Tx output squelching function reduces OMA 10b: Tx output squelching function reduces P _{av} 11b: Host controls the method for Tx output squelching, reducing OMA or P _{av} (see Table 8-11) <i>Note: Support of the Tx output squelching function implies support of automatic Tx squelching control</i>	
	3	ForcedSquelchTxSupported	0b/1b: Host cannot/can force squelching of Tx outputs using OutputSquelchForceTx*	
	2	AutoSquelchDisableTxSupported	0b/1b: Host cannot/can disable automatic squelching of Tx outputs using AutoSquelchDisableTx*	
	1	OutputDisableTxSupported	0b/1b: Host cannot/can disable Tx outputs using the OutputDisableTx register	
	0	InputPolarityFlipTxSupported	0b/1b: InputPolarityFlipTx control not supported/supported	
	7	BankBroadcastSupported	0b/1b: The BankBroadcastEnable control is not supported/supported	
156	6-3	-	Reserved	RO Rqd.
	2	AutoSquelchDisableRxSupported	0b/1b: Host cannot/can disable automatic squelching of Rx outputs using the AutoSquelchDisableRx register <i>Note: Rx squelching and automatic Rx squelching control is not advertised and always assumed to be supported.</i>	
	1	OutputDisableRxSupported	0b/1b: Host cannot/can disable Rx outputs using the OutputDisableRx register	
	0	OutputPolarityFlipRxSupported	0b/1b: PolarityFlipRx not supported/supported	

8.4.8 Supported Flags Advertisement

Table 8-51 describes supported module and lane Flags (see section 8.10.3).

Table 8-51 Supported Flags Advertisement (Page 01h)

Byte	Bit	Field Name	Field Description	Type
157	7-4	-	Reserved	RO
	3	AdaptiveInputEqFailFlagTxSupported	0b: Tx Adaptive Input Eq Fail Flags not supported 1b: supported	RO Rqd.
	2	CDRLOLFlagTxSupported	0b: Tx CDR Loss of Lock Flags not supported 1b: supported	
	1	LOSFlagTxSupported	0b: Tx Loss of Signal Flags not supported 1b: supported	
	0	FailureFlagTxSupported	0b: Tx Fault Flags not supported 1b: supported	
158	7-3	-	Reserved	RO
	2	CDRLOLFlagRxSupported	0b: Rx CDR Loss of Lock Flags not supported 1b: Rx CDR Loss of Lock Flags supported	RO Rqd.
	1	LOSFlagRxSupported	0b: Rx Loss of Signal Flags not supported 1b: Rx Loss of Signal Flags supported	
	0	-	Reserved	

8.4.9 Supported Monitors Advertisement

Table 8-52 describes supported module and lane monitors.

Table 8-52 Supported Monitors Advertisement (Page 01h)

Byte	Bit	Field Name	Field Description	Type
159	7-6	-	Reserved	RO
	5	CustomMonSupported	0b: Custom monitor not supported 1b: Custom monitor supported	RO Rqd.
	4	Aux3MonSupported	0b: Aux 3 monitor not supported 1b: Aux 3 monitor supported	
	3	Aux2MonSupported	0b: Aux 2 monitor not supported 1b: Aux 2 monitor supported	
	2	Aux1MonSupported	0b: Aux 1 monitor not supported 1b: Aux 1 monitor supported	
	1	VccMonSupported	0b: Internal 3.3 V monitor not supported 1b: Internal 3.3 V monitor supported	
	0	TempMonSupported	0b: Temperature monitor not supported 1b: Temperature monitor supported	
160	7-5	-	Reserved	RO
	4-3	TxBiasCurrentScalingFactor	Multiplier for 2uA Bias current increment used in Tx Bias current monitor and threshold registers (see Table 8-63 and Table 8-89) 00b: multiply x1 01b: multiply x2 10b: multiply x4 11b: reserved	Rqd.
	2	RxOpticalPowerMonSupported	0b: Rx Optical Input Power monitor not supported 1b: Rx Optical Input Power monitor supported	
	1	TxOpticalPowerMonSupported	0b: Tx Output Optical Power monitor not supported 1b: Tx Output Optical Power monitor supported	
	0	TxBiasMonSupported	0b: Tx Bias monitor not supported 1b: Tx Bias monitor supported	

8.4.10 Supported Configuration and Signal Integrity Controls Advertisement

Table 8-53 describes the advertisement of signal integrity controls and related settings, and of CMIS-VCS [6].

Table 8-53 Supported Signal Integrity Controls Advertisement (Page 01h)

Byte	Bit	Field Name	Field Description	Type
161	7	-	Reserved	RO
	6-5	TxInputEqRecallBuffersSupported	00b: Tx Input Eq Store/Recall not supported 01b: Tx Input Eq Store/Recall buffer count=1 10b: Tx Input Eq Store/Recall buffer count=2 11b: reserved	RO Rqd.
	4	TxInputEqFreezeSupported	0b/1b: Tx Input Eq Freeze not supported / supported	
	3	TxInputAdaptiveEqSupported	0b/1b: Adaptive Tx Input Eq not supported / supported	
	2	TxInputEqHostControlSupported	0b: host-control of Tx Input Eq target not supported 1b: host-control of Tx Input Eq target supported	
	1	TxCDRBypassControlSupported	0b: If a Tx CDR is supported, it cannot be bypassed 1b: If a Tx CDR is supported, it can be bypassed	
	0	TxCDRSupported	0b/1b: Tx CDR not supported / supported	
162	7	VersatileControlSetSupported	0b/1b: CMIS-VCS [6] functionality is not supported / supported, for all supported Staged Control Sets. <i>Note: See specifications in [6] if the VCS parameter space extensions on Pages 18h and 19h are supported.</i>	RO
	6	UnidirReconfigSupported	0b/1b: ApplyImmediateTx/Rx on Page 10h and DPConfigTx/Rx on Page 19h are not supported / supported, for all supported Staged Control Sets	RO Rqd.
	5	StagedSet1Supported	Staged Control Set 1 supported on Page 10h	
	4-3	RxOutputEqControlSupported	Host control of Rx Output Eq targets 00b: not supported 01b: only Pre-cursor supported 10b: only Post-cursor supported 11b: both Pre- and Post-cursor supported	
	2	RxOutputAmplitudeControlSupported	0b: Rx Output Amplitude control not supported 1b: Rx Output Amplitude control supported	
	1	RxCDRBypassControlSupported	0b: Rx CDR Bypass control not supported (if a CDR is supported, it cannot be bypassed) 1b: Rx CDR Bypass control supported	
	0	RxCDRSupported	0b: Rx CDR not supported 1b: Rx CDR supported	

8.4.11 CDB Messaging Support Advertisement

Table 8-54 describes how the module advertises fundamental support of the Command Data Block (CDB) command and reply messaging functionality, as well as of some high-level CDB features and characteristics.

The host can query support for individual CDB commands through the CDB commands shown in Table 9-1.

See section 7.2 for conceptual information on CDB, section 8.23 for details of the messaging mechanism, and chapter 9 for the CDB Command Reference (message catalogue).

Table 8-54 CDB Advertisement (Page 01h)

Byte	Bit	Field Name	Field Description	Type																																				
163	7-6	CdbInstancesSupported	<p>00b: CDB functionality not supported 01b: One CDB instance supported 10b: Two CDB instances supported 11b: Reserved</p> <p>One CDB Instance Bank 0 of Pages 9Fh and the subset of Pages A0h-Afh advertised in CdbMaxPagesEPL (01h:163.3-0), CdbCmdCompleteFlag1 Flag (00h:8.6) and the associated Mask (00h:31.6) are supported.</p> <p>Two CDB Instances Banks 0 and 1 of Pages 9Fh and of the subset of Pages A0h-Afh advertised in CdbMaxPagesEPL (01h:163.3-0), CdbCmdCompleteFlag<i> (00h:8.7-6), <i> = 1,2, and the associated Masks (00h:31.7-6) are supported.</p>	RO Rqd.																																				
5		CdbBackgroundModeSupported	<p>0b: Background CDB operation not supported. 1b: Background CDB operation supported</p> <p><i>Note: In Background Mode, register access is possible while a CDB command is still being processed.</i></p>	RO Cnd.																																				
4		CdbAutoPagingSupported	<p>When the Management Memory current address pointer advances past the end of an Extended Payload (EPL) CDB Page (A0h-AFh), the page number in the PageSelect Byte will automatically increment and the Memory Map current address pointer automatically wraps to 128. When the last page number AFh is incremented, the page number wraps back to A0h.</p> <p>0b: Auto Paging not supported 1b: Auto Paging and Auto Page wrap supported</p>	RO Cnd.																																				
3-0		CdbMaxPagesEPL	<p>This field encodes the EPL Page range supported or, equivalently, the maximum length of extended payload:</p> <table> <thead> <tr> <th>Value</th> <th>Supported EPL Pages</th> <th>Total Number of Pages</th> <th>Total Number of EPL Bytes</th> </tr> </thead> <tbody> <tr> <td>0:</td> <td>(none)</td> <td>0</td> <td>0</td> </tr> <tr> <td>1:</td> <td>A0h</td> <td>1</td> <td>128</td> </tr> <tr> <td>2:</td> <td>A0h-A1h</td> <td>2</td> <td>256</td> </tr> <tr> <td>3</td> <td>A0h-A2h</td> <td>3</td> <td>384</td> </tr> <tr> <td>4</td> <td>A0h-A3h</td> <td>4</td> <td>512</td> </tr> <tr> <td>5</td> <td>A0h-A7h</td> <td>8</td> <td>1024</td> </tr> <tr> <td>6</td> <td>A0h-Abh</td> <td>12</td> <td>1536</td> </tr> <tr> <td>7</td> <td>A0h-Afh</td> <td>16</td> <td>2048</td> </tr> </tbody> </table> <p><i>Note: A host can access all supported EPL Pages and the EPL Page range is sufficient for all supported CDB commands. The required number of EPL pages may be CDB command specific.</i></p>	Value	Supported EPL Pages	Total Number of Pages	Total Number of EPL Bytes	0:	(none)	0	0	1:	A0h	1	128	2:	A0h-A1h	2	256	3	A0h-A2h	3	384	4	A0h-A3h	4	512	5	A0h-A7h	8	1024	6	A0h-Abh	12	1536	7	A0h-Afh	16	2048	RO Cnd.
Value	Supported EPL Pages	Total Number of Pages	Total Number of EPL Bytes																																					
0:	(none)	0	0																																					
1:	A0h	1	128																																					
2:	A0h-A1h	2	256																																					
3	A0h-A2h	3	384																																					
4	A0h-A3h	4	512																																					
5	A0h-A7h	8	1024																																					
6	A0h-Abh	12	1536																																					
7	A0h-Afh	16	2048																																					
164	7-0	CdbReadWriteLengthExtension	<p><i>Note: For READ and WRITE efficiency in CDB messaging, a module can support multi-byte ACCESS in the CDB Page range (9Fh-Afh) with more than 8 bytes.</i></p> <p>CdbReadWriteLengthExtension = i specifies i^8 allowable additional number of bytes in a WRITE or READ</p>	RO Cnd.																																				

Byte	Bit	Field Name	Field Description	Type
			<p>access to an EPL CDB Page (A0Fh-Afh), i.e. i is a length extension in units of byte octets (8 bytes).</p> <p>For page 9Fh (without auto paging support), the allowable length extension is $\min(i, 15) * 8 = 120$ Bytes.</p> <p>This leads to the maximum length of a READ or a WRITE</p> <p>Value Maximum number of bytes (EPL)</p> <p>0: 8 bytes (no extension of general length limit) i: $8 * (1+i)$ bytes ($0 \leq i \leq 255$) 255: $8 * 256 = 2048$ bytes max</p> <p>Value Maximum Number of Bytes (LPL)</p> <p>0: 8 bytes (no extension of general length limit) i: $8 * (1+i)$ bytes ($0 \leq i \leq 15$) i: $8 * 16 = 128$ bytes ($16 \leq i \leq 256$)</p> <p><i>Note: If the MCI transaction from the host is longer than the length allowed as per this advertisement, the module may ignore bytes written beyond the allowed length and not return more than so many bytes in a read.</i></p>	
165	7	CdbCommandTriggerMethod	<p>Determines how the host triggers CDB command processing in the module and when this occurs:</p> <p>1b: when the MCI transaction of a WRITE access including the CMDID register 9Fh:129 is properly terminated by the host (STOP).</p> <p>0b: when a single byte WRITE to 9Fh:129 or a two-byte WRITE to the CMDID register Byte 9Fh:128-129 is properly terminated by the host (STOP).</p> <p><i>Note: Preferred method 1b enables the host to WRITE a complete CMD message (header and body) in one go, whereas in Method 0b the host composes the CMD message body first and then triggers CMD processing in a second step. See also section 7.2.</i></p>	RO Cnd.
	6-5	-	Reserved	RO
	4-0	CdbExtMaxBusyTime	When CdbMaxBusySpecMethod=1b: CdbExtMaxBusyTime = X encodes the maximum CDB busy time T_{CDBB} as $\max(1, X) * 160$ ms in a range of 160 ms to 4960 ms. When CdbMaxBusySpecMethod=0b: don't care	RO Cnd.
166	7	CdbMaxBusySpecMethod	0b: Indicates that the maximum CDB busy time T_{CDBB} is specified via CdbMaxBusyTime (01h:166.6-0) 1b: Indicates that the maximum CDB busy time T_{CDBB} is specified via CdbExtMaxBusyTime (01h:165.4-0).	RO Cnd.
	6-0	CdbMaxBusyTime	When CdbMaxBusySpecMethod=0b: CdbMaxBusyTime=X encodes the maximum CDB busy time T_{CDBB} as $(80 - \min(80, X))$ ms in a range of 0 ms to 80 ms. When CdbMaxBusySpecMethod=1b: don't care	RO Cnd.

Editor's Note: The following text needs to be revised and cleared from redundancies in a future revision.

CDB instances, are supported by the module.

A value of CdbInstancesSupported = 0 indicates that CDB is not supported at all.

A CDB instance is identified by non-zero **CDB instance number**

The Bank Index of the Pages belonging to a CDB instance is the CDB instance number decremented by one.

All CDB instances behave identically and support the same set of CDB commands.

Note: Module support for multiple CDB instances can be useful when long-duration CDB commands operating in the background are supported, such as firmware update.

Background Operation

The **CdbBackgroundModeSupported** Bit 01h:163.5 defines if the host can ACCESS the addressable management memory while a CDB command is being processed by the module.

- If CdbBackgroundModeSupported is cleared, the module will hold-off ACCESS while a CDB command is being executed until the command is completed (see section 5.2.3 and section 10.2.2)
- If CdbBackgroundModeSupported is set, the module will possibly hold-off ACCESS only until the CDB command is parsed and captured or queued (see section 5.2.3 and section 10.2.2).

When CDB Background Operation is supported, the host can read the **CdbStatus** field to determine the status of in-progress CDB commands (see Table 8-13).

While a CDB Command is being executed in the background, the module ensures that the relevant internal background operations (such as flash EEPROM writes) do not affect other host interactions with the module that may concurrently occur in the foreground.

Auto Paging

The advertisements of Auto-Paging support (in **CdbAutoPagingSupported** Bit 01h:163.4), the number of supported EPL pages (in **CdbMaxPagesEPL** Field 01h:163.3-0), and the maximum write transaction length (in the **CdbReadWriteLengthExtension** Byte 01h:164) are interrelated, as described in Table 8-55, below.

Table 8-55 Overview of CDB advertising combinations

Auto Paging Supported? (01h:163.4)	# EPL Pages Supported? (01h:163.3-0)	Max # Bytes in Seq. Byte Write (01h:164)	Description
0	> 0	Any	Auto paging is not supported. Normal wrap of current address within page. The host should not write past the end of an EPL Page.
1	0	Any	Invalid (useless).
	> 0		Auto Paging is supported, a write past address 255 automatically increments the Page number and wraps the current address pointer to byte 128. If the Page number increment is past the last supported EPL Page, the Page number wraps back to A0h. <i>Note: The host may use the Auto Paging feature to write data in large chunks, without the overhead of explicitly programming Page changes.</i>

8.4.12 Additional Durations Advertisement

The fields described in Table 8-56 advertise maximum durations of module operations. See also section 8.4.5.

Table 8-56 Additional Durations Advertising (Page 01h)

Byte	Bit	Field Name	Field Description	Type
167	7-4	MaxDurationModulePwrDn	Encoded maximum duration of the ModulePwrDn state (see Table 8-48)	RO Rqd.
	3-0	MaxDurationModulePwrUp	Encoded maximum duration of the ModulePwrUp state (see Table 8-48)	
168	7-4	MaxDurationDPTxTurnOff	Encoded maximum duration of the DPTxTurnOff state (see Table 8-48)	RO Rqd.
	3-0	MaxDurationDPTxTurnOn	Encoded maximum duration of the DPTxTurnOn state (see Table 8-48)	
169	7-4	-	Reserved	
	3-0	MaxDurationBPC	U4: Encoded scaling of the maximum ACCESS hold-off duration tBPC, after a PageMapping register change (see Table 10-4). The maximum bank or page switch time given MaxDurationBPC = i is $tBPC / 2^i$	RO Rqd.

Note: The MaxDuration fields allow hosts to determine when something has gone wrong in the module during transient states, for example when a module firmware is hung up.*

The module shall report worst-case maximum state durations across all supported configurations.

See sections 6.3.2.5.5 and 6.3.2.5.7 for details of the ModulePwrUp and ModulePwrDn states and sections 6.3.3.8 and 6.3.3.10 for details of the DPTxTurnOn and DPTxTurnOff states.

8.4.13 Normalized Application Descriptor Support

The fields described in Table 8-57 advertise if Normalized Application Descriptors (**NAD**) are supported.

Application selection by **Application Numbers** in Control Sets requires NAD support. See section 6.2.1.4.2.

Table 8-57 Normalized Application Descriptors Support (Page 01h)

Byte	Bit	Field Name	Field Description	Type
175	7-4	-	Reserved	RO Rqd.
	3-0	NADBanksSupported	0: Normalized Application Descriptors not supported n>0: Normalized Application Descriptors supported on n Banks of Page 1Ch. Up to n·15 Applications can be advertised in n groups of 15 NADs each.	

8.4.14 Media Lane Assignment Options Advertisement

Each Application Descriptor comprises five bytes to advertise an Application, as described in section 6.2.1.4.

The first four bytes of the Application Descriptors are advertised in Bytes 00h:86-117 (see Table 8-23) or in Bytes 01h:223-250 (see Table 8-59). The fifth bytes are advertised in Bytes 01h:176-190 (see Table 8-58).

Both parts of the Application Descriptor are linked by the number known as AppSel Code. The AppSel code is simply the sequential number of the Application Descriptor in any of the memory locations storing (parts of) the Application Descriptor array.

Table 8-58 Media Lane Assignment Advertising (Page 01h)

Byte	Bits	Register Name	Register Description	Type
176	7-0	MediaLaneAssignmentOptionsApp1	MediaLaneAssignmentOptionsApp<i>	RO Rqd.
177	7-0	MediaLaneAssignmentOptionsApp2	Media Lane Assignment Options for the Application advertised in Application descriptor identified by AppSel <i>.	
178	7-0	MediaLaneAssignmentOptionsApp3	Bits 0-7 form a bit map corresponding to Media Lanes 1-8. A set bit indicates that a Data Path for the Application is allowed to begin on the corresponding Media Lane.	
179	7-0	MediaLaneAssignmentOptionsApp4		
180	7-0	MediaLaneAssignmentOptionsApp5		
181	7-0	MediaLaneAssignmentOptionsApp6		
182	7-0	MediaLaneAssignmentOptionsApp7		
183	7-0	MediaLaneAssignmentOptionsApp8		
184	7-0	MediaLaneAssignmentOptionsApp9		
185	7-0	MediaLaneAssignmentOptionsApp10		
186	7-0	MediaLaneAssignmentOptionsApp11		
187	7-0	MediaLaneAssignmentOptionsApp12		
188	7-0	MediaLaneAssignmentOptionsApp13		
189	7-0	MediaLaneAssignmentOptionsApp14		
190	7-0	MediaLaneAssignmentOptionsApp15	Each instance of an Application uses contiguous Media Lane numbers. If multiple instances of a single Application are allowed, each starting point is identified, and all instances must be supported concurrently. (See section 6.2.1)	

8.4.15 Additional Application Advertisement

Bytes 01h:223-250 (see Table 8-59) provide space for seven additional Application Descriptors (the first four bytes of each descriptor) in addition to the eight Application Descriptors in Bytes 86-177 (see Table 8-23).

See section 6.2.1.4 for information on Application Advertising and section 8.2.12 for information about the Application Descriptor and the array of Application Descriptors.

The HostInterfaceID field of the first unused descriptor in Table 8-59 shall have a value of FFh, indicating the end of the list of Application Descriptors.

Table 8-59 Additional Application Descriptor Registers (Page 01h)

Byte	Bits	Field Name	Register Description	Type
223	7-0	HostInterfaceIDApp9	AppDescriptor9	RO
224	7-0	MediaInterfaceIDApp9	AppSel 9 (1001b) See Table 8-22	Cnd.
225	7-4	HostLaneCountApp9		
	3-0	MediaLaneCountApp9		
226	7-0	HostLaneAssignmentOptionsApp9		
227	7-0	HostInterfaceIDApp10	AppDescriptor10	RO
228	7-0	MediaInterfaceIDApp10	AppSel 10 (1010b) See Table 8-22	Cnd.
229	7-4	HostLaneCountApp10		
	3-0	MediaLaneCountApp10		
230	7-0	HostLaneAssignmentOptionsApp10		
231	7-0	HostInterfaceIDApp11	AppDescriptor11	RO
232	7-0	MediaInterfaceIDApp11	AppSel 11 (1011b) See Table 8-22	Cnd.
233	7-4	HostLaneCountApp11		
	3-0	MediaLaneCountApp11		
234	7-0	HostLaneAssignmentOptionsApp11		
235	7-0	HostInterfaceIDApp12	AppDescriptor12	RO
236	7-0	MediaInterfaceIDApp12	AppSel 12 (1100b) See Table 8-22	Cnd.
237	7-4	HostLaneCountApp12		
	3-0	MediaLaneCountApp12		
238	7-0	HostLaneAssignmentOptionsApp12		
239	7-0	HostInterfaceIDApp13	AppDescriptor13	RO
240	7-0	MediaInterfaceIDApp13	AppSel 13 (1101b) See Table 8-22	Cnd.
241	7-4	HostLaneCountApp13		
	3-0	MediaLaneCountApp13		
242	7-0	HostLaneAssignmentOptionsApp13		
243	7-0	HostInterfaceIDApp14	AppDescriptor14	RO
244	7-0	MediaInterfaceIDApp14	AppSel 14 (1110b) See Table 8-22	Cnd.
245	7-4	HostLaneCountApp14		
	3-0	MediaLaneCountApp14		
246	7-0	HostLaneAssignmentOptionsApp14		
247	7-0	HostInterfaceIDApp15	AppDescriptor15	RO
248	7-0	MediaInterfaceIDApp15	AppSel 15 (1111b) See Table 8-22	Cnd.
249	7-4	HostLaneCountApp15		
	3-0	MediaLaneCountApp15		
250	7-0	HostLaneAssignmentOptionsApp15		

Note: The fifth bytes of these Application Descriptors are stored elsewhere (see Table 8-58)

1
2
8.4.16 Miscellaneous Advertisements

Table 8-60 Miscellaneous Advertisements (Page 01h)

Byte	Bit	Field Name	Field Description	Type
251	7-6	ScratchPadSupported	Host Scratchpad Area (13h:184-191) support: 00: unknown (only for CMIS 5.2 or earlier) 01: not supported 02: supported (recommended since CMIS 5.3) <i>Note: Support may become mandatory in a future major CMIS revision</i>	RO Rqd.
	5-4	PasswordEntrySupported [9]	00: unknown (only for CMIS 5.2 or earlier) 01: not supported 02: supported	
	3-2	PasswordEntryResultSupported	00: unknown (only for CMIS 5.2 or earlier) 01: not supported 02: supported	
	1-0	FullPageReadSupported	When supported, the maximum number of bytes that can be read with a READ access is 128. 00: unknown (only for CMIS 5.2 or earlier) 01: not supported 02: supported (recommended since CMIS 5.3) <i>Note: Support may become mandatory in a future major CMIS revision</i>	
252	7	HostLaneSwitchingSupported	0b/1b: Page 1Dh is not supported / supported	
	6	LinkTrainingSupported	0b/1b: Functionality specified in CMIS-LT [9] is not supported / supported on Pages 50h-53h . <i>Note: CMIS-LT-specific advertisements may be found on these pages.</i>	
	5-0	-	Reserved	
253	7-0	-	Reserved	
254	7-0	-	Reserved	

3
4
8.4.17 Page Checksum (Page 01h, Byte 255, RO RQD)

5
6
7
The Page Checksum is a one-byte code that can be used to verify that the read-only static data on Page 01h is valid. The checksum code shall be the low order 8 bits of the arithmetic sum of all byte values from byte 130 to byte 254, inclusive.

8
Note that the module firmware revision in bytes 128 and 129 is intentionally not included in the Page Checksum.

8.5 Page 02h (Module and Lane Thresholds)

Page 02h is an optional Page that informs about module-defined thresholds for module-level and lane-specific threshold crossing monitors. All fields on Page 02h are read-only.

The threshold values on this Page can depend on the commissioned set of Applications and therefore may change (including the checksum) whenever a new Application is commissioned.

Note: The module will update threshold values only for that reason. For Data Path Applications, the thresholds will be updated when the relevant Data Path reaches DPInitialized. For Network Path Applications, the Media Side thresholds will be updated when the relevant Network Path reaches NPInitialized, and the Host side thresholds will be updated when the relevant Host Path reaches DPInitialized.

Note: It is recommended that hosts interested in the thresholds read them after Data Path Initialization and/or Network Path Initialization.

The module advertises support of Page 02h in the MemoryModel Bit 00h:2.7.

Note: Page 02h is mandatory for paged memory modules.

Table 8-61 Page 02h Overview

Byte	Size (bytes)	Subject Area	Description
128-175	48	Module-level monitor thresholds	
176-199	24	Lane-specific monitor thresholds	
200-229	30	-	Reserved[30]
230-254	25	-	Custom[25]
255	1	Page Checksum	Covers bytes 128-254

8.5.1 Module-Level Monitor Thresholds

The following thresholds are provided by the module to inform the host of the monitored observable levels where alarms and warnings will be triggered.

Table 8-62 Module-Level Monitor Thresholds (Page 02h)

Byte	Bit	Name	Description	Type
128-129	7-0	TempMonHighAlarmThreshold	S16 Thresholds for internal temperature monitor: 1/256 degree Celsius increments	RO Cnd.
130-131	7-0	TempMonLowAlarmThreshold		
132-133	7-0	TempMonHighWarningThreshold		
134-135	7-0	TempMonLowWarningThreshold		
136-137	7-0	VccMonHighAlarmThreshold	U16 Thresholds for internal 3.3 volt input supply voltage monitor: 100 µV increments	RO Cnd.
138-139	7-0	VccMonLowAlarmThreshold		
140-141	7-0	VccMonHighWarningThreshold		
142-143	7-0	VccMonLowWarningThreshold		
144-145	7-0	Aux1MonHighAlarmThreshold	TEC Current: 100/32767% increments of maximum TEC current +32767 (100%) – Max Heating -32767 (-100%) – Max Cooling	RO Cnd.
146-147	7-0	Aux1MonLowAlarmThreshold		
148-149	7-0	Aux1MonHighWarningThreshold		
150-151	7-0	Aux1MonLowWarningThreshold		
152-153	7-0	Aux2MonHighAlarmThreshold	TEC Current: 100/32767% increments of maximum TEC current +32767 (100%) – Max Heating -32767 (-100%) – Max Cooling Laser Temperature: 1/256 degree Celsius increments	RO Cnd.
154-155	7-0	Aux2MonLowAlarmThreshold		
156-157	7-0	Aux2MonHighWarningThreshold		
158-159	7-0	Aux2MonLowWarningThreshold		
160-161	7-0	Aux3MonHighAlarmThreshold	Laser Temperature: 1/256 degree Celsius increments	RO Cnd.
162-163	7-0	Aux3MonLowAlarmThreshold		
164-165	7-0	Aux3MonHighWarningThreshold		
166-167	7-0	Aux3MonLowWarningThreshold		

Byte	Bit	Name	Description	Type
			<i>NOTE: Laser Temp can be below 0 if uncooled or in Tx Output Disable.</i> Additional supply voltage monitor: 100 μ V increments	
168-169	7-0	CustomMonHighAlarmThreshold	S16 or U16 Custom monitor	RO Cnd.
170-171	7-0	CustomMonLowAlarmThreshold		
172-173	7-0	CustomMonHighWarningThreshold		
174-175	7-0	CustomMonLowWarningThreshold		

1

8.5.2 Lane-Related Monitor Thresholds

The following thresholds are provided by the module to inform the host of the monitor levels where alarms and warnings will be triggered. These monitor thresholds apply to all lanes of the module.

Note: See section 8.10.4 for monitor details including accuracy.

5

Table 8-63 Lane-Related Monitor Thresholds (Page 02h)

Byte	Bit	Name	Description	Type
176-177	7-0	OpticalPowerTxHighAlarmThreshold	U16 Thresholds for Tx optical power monitor: in 0.1 uW increments Total measurement range of 0 to 6.5535 mW (~-40 dBm to +8.2 dBm for non-zero values)	RO Cnd.
178-179	7-0	OpticalPowerTxLowAlarmThreshold		
180-181	7-0	OpticalPowerTxHighWarningThreshold		
182-183	7-0	OpticalPowerTxLowWarningThreshold		
184-185	7-0	LaserBiasCurrentHighAlarmThreshold	U16 Thresholds for Tx laser bias monitor: 2 uA increments, times the multiplier encoded in 01h:160.4-3 (see Table 8-52)	RO Cnd.
186-187	7-0	LaserBiasCurrentLowAlarmThreshold		
188-189	7-0	LaserBiasCurrentHighWarningThreshold		
190-191	7-0	LaserBiasCurrentLowWarningThreshold		
192-193	7-0	OpticalPowerRxHighAlarmThreshold	U16 Thresholds for Rx optical power monitor: 0.1 uW increments Total measurement range of 0 to 6.5535 mW (~-40 dBm to +8.2 dBm for non-zero values)	RO Cnd.
194-195	7-0	OpticalPowerRxLowAlarmThreshold		
196-197	7-0	OpticalPowerRxHighWarningThreshold		
198-199	7-0	OpticalPowerRxLowWarningThreshold		

7

8.5.3 Page Checksum (Page 02h, Byte 255, RO RQD)

The Page Checksum code is a one-byte code that can be used to verify that the device property information in the module is valid. The Page Checksum code shall be the low order 8 bits of the arithmetic sum of all byte values from byte 128 to byte 254, inclusive.

8

9

10

11

8.6 Page 03h (User EEPROM)

Page 03h is an optional Page that allows the module to provide access to a host writeable EEPROM.

The module advertises support of Page 03h in Bit 01h:142.2 (see Table 8-46).

The host may read or write to this memory for any purpose.

Note: Actual usage of the EEPROM is not standardized by CMIS.

The maximum number of bytes in one WRITE access is 8 bytes.

The module **rejects** register ACCESS until a WRITE to EEPROM is completed internally. The maximum duration of the ACCESS hold-off period caused by internal write completion is tWRITENV as specified in section 10.2.2.

Note: In SFF-8636, a CLEI code may be present in the first 10 Bytes of Page 02h. In CMIS, this convention is aborted, because the CLEI code is allocated in Bytes 00h:190-199.

Table 8-64 Page 03h Overview

Byte	Size (bytes)	Subject Area	Description
128-255	128	User Data	Module user data stored in NV memory

8.7 Page 04h (Laser Capabilities Advertising)

Page 04h is an optional Page that allows the module to advertise tunable laser capabilities.

The module advertises support of Page 04h in Bit 01h:155.6 (see Table 8-50).

The tunable laser capabilities are defined at module level, for all media lanes.

Table 8-65 Page 04h Overview

Byte	Size (bytes)	Subject Area	Description
128-129	2	Wavelength grids	Supported grids and associated channel numbering
130-189	60	Channel number ranges	Channel number ranges per grid
190-197	8	Fine-Tuning Support	Wavelength fine tuning resolution and range
198-201	4	Programmable Output Power	Programmable range of output powers
202-254	53	-	Reserved[53]
255	1	Page Checksum	Covers bytes 128-254

Bytes 04h:128 advertises supported channel grids and 04h:129 advertises fine tuning support (see Table 8-66).

In the following Table 8-66 a local parameter **n** is used as a shorthand for a signed **channel offset number**, which is sometimes simply called a **channel number**.

This channel offset number effectively specifies a laser frequency in terms of a (signed) frequency offset from a nominal **reference** frequency at 193.1THz, in units of the grid resolution (except for the 75GHz and 150GHz grids, where the offset is defined in units of a third or a sixth of the grid resolution).

*Note: This **reference** channel is sometimes colloquially referred to as a center channel or center frequency, but this may be misleading. In fact, the reference channel should better be viewed merely as a **grid anchor** frequency for an equidistant grid which is not necessarily in the center of the supported band.*

The minimum and maximum channel offset numbers supported are used to define the tuning range on each grid in the **supported channel number range** registers 04h:130-161.

Note: The channel offset number range is large enough to cover the L Band range from 191.6THz to 184.5THz.

This channel offset number **n** is also used to specify the laser frequency in Page 12h.

Table 8-66 Laser capabilities for tunable lasers (Page 04h)

Byte	Bits	Name	Description	Type
128	7	GridSupported75GHz	Bool: Indicates whether the module supports channel-based tuning on the 75 GHz grid defined as: Frequency (THz) = 193.1 + n × 0.025 where n must be an integer multiple of 3. 0b: 75 GHz Grid not supported 1b: 75 GHz Grid supported n is the 16-bit signed channel number that is referred to in the highest/lowest supported channel numbers in Page 04h	RO Rqd.
	6	GridSupported33GHz	Bool: Indicates whether the module supports channel-based tuning on the 33 GHz grid defined as: Frequency (THz) = 193.1 + n × 0.1/3 0b: 33 GHz Grid not supported 1b: 33 GHz Grid supported	RO Rqd.
	5	GridSupported100GHz	Bool: Indicates whether the module supports channel-based tuning on the 100 GHz grid defined as: Frequency (THz) = 193.1 + n × 0.1 0b: 100 GHz Grid not supported 1b: 100 GHz Grid supported	RO Rqd.
	4	GridSupported50GHz	Bool: Indicates whether the module supports channel-based tuning on the 50 GHz grid defined as: Frequency (THz) = 193.1 + n × 0.05 0b: 50 GHz Grid not supported 1b: 50 GHz Grid supported	RO Rqd.

Byte	Bits	Name	Description	Type
	3	GridSupported25GHz	Bool: Indicates whether the module supports channel-based tuning on the 25 GHz grid defined as: Frequency (THz) = 193.1 + $n \times 0.025$ 0b: 25 GHz Grid not supported 1b: 25 GHz Grid supported	RO Rqd.
	2	GridSupported12p5GHz	Bool: Indicates whether the module supports channel-based tuning on the 12.5 GHz grid defined as: Frequency (THz) = 193.1 + $n \times 0.0125$ 0b: 12.5 GHz Grid not supported 1b: 12.5 GHz Grid supported	RO Rqd.
	1	GridSupported6p25GHz	Bool: Indicates whether the module supports channel-based tuning on the 6.25 GHz grid defined as: Frequency (THz) = 193.1 + $n \times 0.00625$ 0b: 6.25 GHz Grid not supported 1b: 6.25 GHz Grid supported	RO Rqd.
	0	GridSupported3p125GHz	Bool: Indicates whether the module supports channel-based tuning on 3.125 GHz grid defined as: Frequency (THz) = 193.1 + $n \times 0.003125$ 0b: 3.125 GHz Grid not supported 1b: 3.125 GHz Grid supported	RO Rqd.
129	7	FineTuningSupported	Bool: Indicates whether the module supports fine-tuning of laser frequency in the vicinity of an on-grid channel. 0b: module does not support fine-tuning 1b: module supports fine-tuning	RO Rqd.
	6	GridSupported150GHz	Bool: Indicates whether the module supports channel-based tuning on the 150 GHz grid defined as: Frequency (THz) = 193.1 + $(n+3) \times 0.025$ where n must be an integer multiple of 6. 0b: 150 GHz Grid not supported 1b: 150 GHz Grid supported	
	5-0	-	Reserved for future channel spacing advertisements	RO Rqd.
130-131	7-0	GridLowChannel3p125GHz	S16 Lowest supported n for 3.125 GHz spacing	RO Rqd.
132-133	7-0	GridHighChannel3p125GHz	S16 Highest supported n for 3.125 GHz spacing	RO Rqd.
134-135	7-0	GridLowChannel6p25GHz	S16 Lowest supported n for 6.25 GHz spacing	RO Rqd.
136-137	7-0	GridHighChannel6p25GHz	S16 Highest supported n for 6.25 GHz spacing	RO Rqd.
138-139	7-0	GridLowChannel12p5GHz	S16 Lowest supported n for 12.5 GHz spacing	RO Rqd.
140-141	7-0	GridHighChannel12p5GHz	S16 Highest supported n for 12.5 GHz spacing	RO Rqd.
142-143	7-0	GridLowChannel25GHz	S16 Lowest supported n for 25 GHz spacing	RO Rqd.
144-145	7-0	GridHighChannel25GHz	S16 Highest supported n for 25 GHz spacing	RO Rqd.
146-147	7-0	GridLowChannel50GHz	S16 Lowest supported n for 50 GHz spacing	RO Rqd.
148-149	7-0	GridHighChannel50GHz	S16 Highest supported n for 50 GHz spacing	RO Rqd.
150-151	7-0	GridLowChannel100GHz	S16 Lowest supported n for 100 GHz spacing	RO Rqd.
152-153	7-0	GridHighChannel100GHz	S16 Highest supported n for 100 GHz spacing	RO Rqd.
154-155	7-0	GridLowChannel33GHz	S16 Lowest supported n for 33 GHz spacing	RO Rqd.
156-157	7-0	GridHighChannel33GHz	S16 Highest supported n for 33 GHz spacing	RO Rqd.
158-159	7-0	GridLowChannel75GHz	S16 Lowest supported n for 75 GHz spacing	RO Rqd.
160-161	7-0	GridHighChannel75GHz	S16 Highest supported n for 75 GHz spacing	RO Rqd.
162-163	7-0	GridLowChannel150GHz	S16 Lowest supported n for 150 GHz spacing	RO Rqd.
164-165	7-0	GridHighChannel150GHz	S16 Highest supported n for 150 GHz spacing	RO Rqd.
166-189	7-0	-	Reserved[24] for future channel spacing support	RO Rqd.
190-191	7-0	FineTuningResolution	U16 Fine-tuning resolution, increments of 0.001 GHz	RO Rqd.
192-193	7-0	FineTuningLowOffset	S16 lowest fine-tuning offset with resolution of 0.001 GHz	RO Rqd.
194-195	7-0	FineTuningHighOffset	S16 highest fine-tuning offset in increments of 0.001 GHz	RO Rqd.

Byte	Bits	Name	Description	Type
196	7	ProgOutputPowerPerLaneSupported	0b: Programmable output power per lane not supported 1b: Programmable output power per lane supported	RO Rqd.
	6-0	-	Reserved	
197	7-0	-	Reserved[1]	RO
198-199	7-0	ProgOutputPowerMin	S16 Minimum Programmable Output Power in increments of 0.01 dBm	RO Rqd.
200-201	7-0	ProgOutputPowerMax	S16 Maximum Programmable Output Power in increments of 0.01 dBm	RO Rqd.
202-254	7-0	-	Reserved[53]	RO
255	7-0	PageChecksum	Page Checksum over bytes 128-254 (see section 8.1.4.3)	RO Rqd.

1

1 **8.8 Page 05h (Form Factor Specific Management Signals Management)**

2 Page 05h is an optional Page that is restricted for the management of form factor specific management signals,
3 i.e. to form factor specific extensions of the MSL (see section 5.1).

4 The actual specifications are defined in one or more separate OIF documents (see section 2.1.2).

5 The module advertises support of Page 05h in Bit 01h:142.3 (see Table 8-46).

6

8.9 Banked Page 10h (Lane Control and Data Path Control)

Page 10h is an optional Page that contains lane control bytes.

The module advertises support of Page 10h in the **MemoryModel** Bit 00h:2.7.

Note: Page 0:10h is mandatory for paged memory modules.

Page 10h may optionally be Banked. Each Bank of Page 10h refers to a group of 8 lanes.

Page 10h is subdivided into several areas as illustrated in the following table:

Table 8-67 Page 10h Overview

Byte	Size (bytes)	Subject Area	Description
128	1	Data Path Control	Data Path control bits for each lane, controlling associated Data Path State machines
129-142	14	Lane-Specific Control	Fields to control lane attributes independent of the Data Path State machine or control sets
143-177	35	Staged Control Set 0	Fields to select Applications and signal integrity settings
178-212	35	Staged Control Set 1	Fields to select Applications and signal integrity settings
213-232	20	Lane-Specific Masks	Masks to suppress Interrupts from lane-related Flags
233-239	7	-	Reserved[7]
240-255	16	-	Custom[16]

8.9.1 Data Path Initialization Control (DPDeinit Bits)

The **DPDeinit** byte controls the initialization of the lanes in all configured Data Paths that are associated with the 8 lanes represented in a Bank.

Note: See chapter 6 for general information on Data Paths and the Module State Machine, and section 6.3.3 for relevant information about the Data Path State Machine

The module evaluates this Byte only in Module State ModuleReady. When the Module State is ModuleReady, the Data Path associated with lanes whose DPDeinit bits are cleared will transition to the DPInit state and begin the initialization process.

Note: By default, all Data Paths will begin initializing when the Module State reaches ModuleReady. The host can prevent this auto-initialization behavior by setting all DPDeinit bits while the module is in the ModuleLowPwr state.

Note: Multiple Data Paths are mutually independent. They may be initialized or deinitialized at the same time or at different times.

Note: The number of lanes in any specific Data Path is a characteristic of the currently active Application that is selected by the AppSel code in the Active Control Set (see Table 8-93). Refer to section 6.2.1.4 for details on AppSel codes and section 6.2.3 for details on Control Sets.

Table 8-68 Data Path initialization control (Page 10h:128)

Byte	Bit	Field Name	Register Description	Type
128	7	DPDeinitLane8	DPDeinitLane<i> Data Path initialization control for host lane <i> 0b: Initialize the Data Path associated with host lane 1b: Deinitialize the Data Path associated with host lane All lanes of a Data Path must have the same value <i>Note: These bits represent static requests, not trigger events</i>	RW Rqd.
	6	DPDeinitLane7		
	5	DPDeinitLane6		
	4	DPDeinitLane5		
	3	DPDeinitLane4		
	2	DPDeinitLane3		
	1	DPDeinitLane2		
	0	DPDeinitLane1		

8.9.2 Lane-Specific Direct Effect Control Fields

Lane-specific direct effect control fields act on individual lanes in the module and are nominally independent of the Data Path (although all lanes of a Data Path are mostly treated identically, except for specific purposes).

Note: The lane-specific direct effect control field settings are not staged and not part of Control Sets.

A module advertises support for individual controls as described in section 8.4.7.

The direct effect control settings for muting a lane output, may remain temporarily ineffective as per Data Path State Machine specifications (e.g. clearing OutputDisableTx has no immediate effect in DPSM state DPInit).

*Note: Readers should carefully distinguish the muting **functions** (disabling or squelching an output) and the **control** of these functions that may be exercised by the module itself or by the external host. The effect of both host and module controlling these muting functions depends on advertised module capabilities and on host-controlled configuration.*

8.9.2.1 Tx Output Muting Functions and Their Control

A Tx output is muted when it is **disabled** or when it is **squelched**. Regarding the muting method applied and the resulting output signal characteristics, the **output disable takes precedence**.

Tx Output Disable Function

When an optical Tx output is **disabled**, its output signal has negligible optical average output power as defined by the relevant media interface standard (e.g. average power <-20dBm). The output is then **quiescent**.

Note: It is irrelevant if the implementation disables transmitter facilities (e.g. laser off) or just shuts down the output. That is why the function is named Tx Output Disable instead of Tx Disable as in earlier revisions.

When an electrical Tx output is **disabled**, it is **quiescent**.

Joint Control of Tx Output Disable Function

The Tx output disable function in the module is normally controlled by the host only, but the DPSM may override and keep an Tx output disabled when it is already nominally un-disabled by the host.

Tx Output Squelching Function

When a Tx output is **squelched**, either the optical modulation amplitude (OMA) or the average power (P_{av}) of the optical output is reduced such that the output is **quiescent**. The module advertises its **squelching method** or whether the host selects the squelching method (see Table 8-50 and Table 8-11). While a Tx output is disabled, the squelching function is masked or blocked (precedence of output disable).

Joint Control of Tx Output Squelch Function

The Tx output squelch function (if supported) in the module can be controlled both by the host and by a module-internal **squelch controller** that automatically activates the output squelch function when no suitable host side input signal is available for forwarding on the media side output. The output squelch function activation on host command is called **forced** output **squelching**, while the activation due to internal controller decision is called **automatic** output **squelching**.

For NP applications, **automatic** Tx output squelching is generally not supported.

The automatic Tx output squelch controller itself (if supported) can be enabled or disabled by the host.

8.9.2.2 Rx Output Muting Functions and Their Control

An Rx output is muted when it is **disabled** or when it is **squelched**. Regarding the muting method and the resulting output signal characteristics, there is no precedence defined.

Rx Output Disable Function

When an electrical Rx output is disabled, the output shall be **quiescent**.

Control of Rx Output Disable Function

The Rx output disable function in the module is controlled by the host.

Rx Output Squelching Function

When an electrical Rx output is squelched, the output shall be **quiescent**.

Control of Rx Output Squelching Function

The Rx output squelch function is controlled only by a module-internal **squelch controller** (if supported) that **automatically** activates the output squelch function when no suitable media side Rx input signal is available to be forwarded on the host side output.

Note: A single Rx input lane may feed more than one Rx output lane, and the data transmitted on one Rx output lane may originate from more than one Rx input lane. The internal controller squelches an Rx electrical output lane until all associated Rx input lanes have detected a valid input signal and all associated internal resources are fully initialized and capable of forwarding a valid stable signal, in order to avoid link flaps.

The automatic Rx output squelch controller itself (if supported) can be enabled or disabled by the host.

8.9.2.3 Lane-specific Tx and Rx Control Fields

This section lists the specific controls for the purposes discussed above.

Tx Output Controls

The host can disable and un-disable Tx output lane N using **OutputDisableTx<N>**.

The host can force and unforce squelching of Tx output lane N using **OutputSquelchForceTx<N>**.

The host can disable or enable the internal squelch controller for lane N using **AutoSquelchDisableTx<N>**.

Tx Input Controls

The host can switch the host side input signal polarity of lane N using **InputPolarityFlipTx<N>**.

The host can freeze the host side input equalizer adaptation for lane N using **AdaptiveInputEqFreezeTx<N>**.

The host can store the current equalizer setting of lane N for later recall using **AdaptiveInputEqStoreTx<N>**.

Rx Output Controls

The host can disable and un-disable Rx output lane N using **OutputDisableRx<N>**.

Note: There is no OutputSquelchForceRx<N> corresponding to OutputSquelchForceTx<N>.

The host can switch the output signal polarity of lane N using **OutputPolarityFlipRx<N>**.

The host can disable or enable the internal squelch controller for lane N using **AutoSquelchDisableRx<N>**.

Note: It is recommended that hosts configure the same auto-squelch disable settings for all lanes (of one direction) within a Data Path (media lanes for Tx and host lanes for Rx), otherwise behavior may be unexpected.

8.9.2.4 Register Lists

Table 8-69 Lane-specific Direct Effect Control Fields (Page 10h)

Byte	Bits	Field Name	Register Description	Type
129	7	InputPolarityFlipTx8	InputPolarityFlipTx<i> 0b: No Tx input polarity flip for lane <i> 1b: Tx input polarity flip for lane <i> Advertisement: 01h:155.0	RW Adv.
	6	InputPolarityFlipTx7		
	5	InputPolarityFlipTx6		
	4	InputPolarityFlipTx5		
	3	InputPolarityFlipTx4		
	2	InputPolarityFlipTx3		
	1	InputPolarityFlipTx2		
	0	InputPolarityFlipTx1		
130	7	OutputDisableTx8	OutputDisableTx<i> 0b: Tx output enabled for media lane <i> 1b: Tx output disabled for media lane <i> Advertisement: 01h:155.1	RW Adv.
	6	OutputDisableTx7		
	5	OutputDisableTx6		
	4	OutputDisableTx5		
	3	OutputDisableTx4		
	2	OutputDisableTx3		
	1	OutputDisableTx2		
	0	OutputDisableTx1		
131	7	AutoSquelchDisableTx8	AutoSquelchDisableTx<i> 0b: Automatic Tx output squelching control enabled for media lane <i>	RW Adv.
	6	AutoSquelchDisableTx7		
	5	AutoSquelchDisableTx6		
	4	AutoSquelchDisableTx5		

Byte	Bits	Field Name	Register Description	Type
	3	AutoSquelchDisableTx4	1b: Automatic Tx output squelching control disabled for media lane <i>	
	2	AutoSquelchDisableTx3		
	1	AutoSquelchDisableTx2		
	0	AutoSquelchDisableTx1	Advertisement: 01h:155.2	
132	7	OutputSquelchForceTx8	OutputSquelchForceTx< i>	RW Adv.
	6	OutputSquelchForceTx7	0b: No impact on Tx output for media lane <i>	
	5	OutputSquelchForceTx6	1b: Tx output squelched for media lane <i>	
	4	OutputSquelchForceTx5		
	3	OutputSquelchForceTx4		
	2	OutputSquelchForceTx3		
	1	OutputSquelchForceTx2		
	0	OutputSquelchForceTx1	Advertisement: 01h:155.3	
133	7-0	-	Reserved[1]	RO
134	7	AdaptiveInputEqFreezeTx8	AdaptiveInputEqFreezeTx< i>	RW Adv.
	6	AdaptiveInputEqFreezeTx7	0b: No impact on Tx input eq adaptation behavior for lane <i>	
	5	AdaptiveInputEqFreezeTx6	1b: Tx input equalizer adaptation frozen at last value for lane <i>	
	4	AdaptiveInputEqFreezeTx5		
	3	AdaptiveInputEqFreezeTx4	See section 6.2.5.4	
	2	AdaptiveInputEqFreezeTx3		
	1	AdaptiveInputEqFreezeTx2	Advertisement: 01h:161.4	
	0	AdaptiveInputEqFreezeTx1		
135	7-6	AdaptiveInputEqStoreTx4	AdaptiveInputEqStoreTx< i>	WO Adv.
	5-4	AdaptiveInputEqStoreTx3	Tx Input Equalizer Adaptation Store location for lane <i>	
	3-2	AdaptiveInputEqStoreTx2	00b: reserved	
	1-0	AdaptiveInputEqStoreTx1	01b: store to recall buffer 1	
136	7-6	AdaptiveInputEqStoreTx8	10b: store to recall buffer 2	WO Adv.
	5-4	AdaptiveInputEqStoreTx7	11b: reserved	
	3-2	AdaptiveInputEqStoreTx6	See section 6.2.5.4	
	1-0	AdaptiveInputEqStoreTx5	Advertisement: 01h:161.6-5	
137	7	OutputPolarityFlipRx8	OutputPolarityFlipRx< i>	RW Adv.
	6	OutputPolarityFlipRx7	0b: No Rx output polarity flip for lane <i>	
	5	OutputPolarityFlipRx6	1b: Rx output polarity flip for lane <i>	
	4	OutputPolarityFlipRx5		
	3	OutputPolarityFlipRx4		
	2	OutputPolarityFlipRx3		
	1	OutputPolarityFlipRx2		
	0	OutputPolarityFlipRx1	Advertisement: 01h:156.0	
138	7	OutputDisableRx8	OutputDisableRx< i>	RW Adv.
	6	OutputDisableRx7	0b: Rx output enabled for lane <i>	
	5	OutputDisableRx6	1b: Rx output disabled for lane <i>	
	4	OutputDisableRx5		
	3	OutputDisableRx4		
	2	OutputDisableRx3		
	1	OutputDisableRx2		
	0	OutputDisableRx1	Advertisement: 01h:156.1	
139	7	AutoSquelchDisableRx8	AutoSquelchDisableRx< i>	RW Adv.
	6	AutoSquelchDisableRx7	0b: Automatic Rx output squelching enabled for host lane <i>	
	5	AutoSquelchDisableRx6	1b: Automatic Rx output squelching disabled for host lane <i>	
	4	AutoSquelchDisableRx5	<i>Note: Automatic Rx output squelching control is not advertised, hence always assumed to be supported.</i>	
	3	AutoSquelchDisableRx4		
	2	AutoSquelchDisableRx3		
	1	AutoSquelchDisableRx2		
	0	AutoSquelchDisableRx1	Advertisement: 01h:156.2	
140-142	All	-	Reserved[3]	RO

8.9.3 Staged Control Set 0

Staged Control Set 0 is required for all paged memory modules.

Background on Control Sets and on the associated configuration and reconfiguration procedures (Provision and Provision-and-Commission) is described in sections 6.2.3. and 6.2.4

8.9.3.1 Apply Staged Control Set Triggers (Configuration Commands)

SCS0::ApplyDPInit and **SCS0::ApplyImmediate** are write-only (WO) trigger registers that allow the host to request execution of a configuration or reconfiguration procedure for one or more Data Paths selected by means of a lane bit mask in the value being written to the trigger register.

Note: Changes to per-lane configurations of a Staged Control Set have no effect on the behavior of a lane until the host has successfully triggered ApplyDPInit or ApplyImmediate for the lane (causing a copy of the lane configuration in the staged control set into the Active Control Set and, ultimately, into hardware or firmware).

Note: As described in section 6.2.3.3, an ApplyDPInit trigger may cause commissioning to hardware without host intervention by a series of DPSM state transitions through DPInit, initiated via the DPInitPending bits.

Note: Support for ApplyImmediate is advertised as described in Table 8-5. When ApplyImmediate is not supported, WRITE access to it is ignored.

A successful **Provision** configuration procedure copies settings from a Staged Control Set to the Active Control set and sets DPInitPending bits.

A successful **Provision-and-Commission** reconfiguration procedure copies settings from a Staged Control Set to the Active Control set and commits the Active Control set to hardware.

Note: The ApplyDPInit and ApplyImmediate registers are stateless trigger registers with write-only access type. This implies that the value read from the register is not specified. Modules may use the bits in these registers for any purpose, including to signal command execution or acceptance status, e.g. for debug purposes.

Command Handling Protocol

The module response to a WRITE onto one of the Apply registers follows a four-step command handling protocol with status reporting in the ConfigStatus register (see also section 8.10.5)

(1) Command Acceptance

When a previously triggered configuration procedure is still being executed on the lanes of a Data Path, the module (silently) ignores any new trigger bits for that Data Path and does not execute the following steps.

Otherwise the module sets ConfigStatusLane<i> = ConfigInProgress for all Data Path lanes <i>, immediately, to indicate that a configuration procedure is being executed.

(2) Parameter Validation

The module validates the defined configuration settings in Staged Control Set 0, and optionally checks if the command is admissible in the current state, before making any changes.

Note: Module hardware and Active Control Set remain unchanged when a trigger is ignored (not accepted) or when the settings in the Staged Control Set are not successfully validated.

(3) Command Execution

The module first copies the validated settings from **Staged Control Set 0** to the Active Control Set, for all lanes <i> selected and accepted in the value written to the Apply trigger register. Further processing then depends on the current state of the Data Path and on the particular Apply trigger register used, as follows:

Provision procedure: When the host WRITE was to ApplyDPInit the module sets the DPInitPendingLane<i> bits in the DPInitPending register.

Provision-and-Commission procedure: When the host WRITE was to ApplyImmediate for lanes indicating DPInitialized or DPActivated state, the module commits the configuration in the Active Control set to hardware.

(4) Result Feedback

Command processing terminates when the module eventually writes the command result status to the **ConfigStatus** fields: The same result status is written to all lanes of a Data Path.

Note: In case of multiple rejection reasons, the module chooses the most important rejection reason. The reporting priority is not specified.

Table 8-70 Staged Control Set 0, Apply Triggers (Page 10h)

Byte	Bits	Field Name	Register Description	Type
143	7	ApplyDPInitLane8	SCS0::ApplyDPInitLane<i> 0b: No action for host lane <i> 1b: Trigger the Provision procedure using the Staged Control Set 0 settings for host lane <i>, with feedback provided in the associated ConfigStatusLane<i> field Restriction: This byte must be written in a single-byte WRITE <i>Note: See Table 6-3 for preconditions and state dependencies</i>	WO Rqd.
144	7	ApplyImmediateLane8	SCS0::ApplyImmediate<i> 0b: No action for host lane <i> 1b: Trigger the Provision-and-Commission procedure using the Staged Control Set 0 settings for host lane <i>, with feedback provided in the associated ConfigStatusLane<i> field Restriction: This byte must be written in a single-byte WRITE Condition: SteppedConfigOnly (00h:2.6) = 0b <i>Note: See Table 6-3 for preconditions and state dependencies</i>	WO Cnd.

8.9.3.2 Data Path Configuration (Application Assignments)

The Data Path Configuration fields allow the host to **allocate** and configure the **Data Paths** consisting of one or more lanes for one or more of the **Applications** advertised by the module in the **Application Descriptor** registers (Table 8-23 and Table 8-59).

Table 8-72 describes the Data Path Configuration as a **DPConfigLane<i>** configuration register array.

Table 8-71 describes the fields in a DPConfigLane<i> configuration Byte, two of which are descriptive (AppSelCode and DataPathID), while one (ExplicitControl) optionally customizes standard lane configurations.

Table 8-71 Data Path Configuration per Lane (DPConfigLane<i>)

Lane	Bits	Field	Field Description	Type
<i>	7-4	AppSelCode	SCS<k>::AppSelCodeLane<i> If lane <i> is part of a Data Path for an Application instance, the AppSelCode field stores the AppSel code of the Descriptor of that Application (see Table 8-23 and Table 8-59) If lane <i> is not part of a Data Path for an Application and hence unused, the AppSelCode field is assigned the NULL value 0000b. All lanes of a Data Path for an Application that spans multiple lanes must have the same setting of AppSelCode.	RW Rqd.
	3-1	DataPathID	SCS<k>::DataPathIDLane<i> If lane <i> is part of a Data Path for an Application instance, this DataPathID field stores the DataPathID of that Data Path, i.e. the number of the first lane in the Data Path, decremented by one. If lane <i> is unused, the value of DataPathID is ignored. All lanes of a Data Path for an Application that spans multiple lanes must have the same setting of DataPathID. <i>Note: For example, the DataPathID of a Data Path including lane 1 is 0 (000b) and the DataPathID of a Data Path where lane 5 is the lowest lane number is 4 (100b).</i>	RW Rqd.

Lane	Bits	Field	Field Description	Type
	0	ExplicitControl	<p>SCS< k >::ExplicitControlLane< i ></p> <p>If lane < i > is part of a Data Path for an Application instance, the ExplicitControl field specifies if lane < i > is configured with host-defined signal integrity settings (provisioned in registers described in Table 8-73 and Table 8-74) or with Application-dependent settings known by the module, when the Staged Control Set is Applied.</p> <p>If lane < i > is unused, the field is ignored.</p> <p>0b: Use Application-dependent settings for lane < i > 1b: Use Staged Control Set 0 control values for lane < i ></p>	RW Rqd.

1

2

Table 8-72 Staged Control Set 0, Data Path Configuration (Page 10h)

Byte	Bits	Field Name	Register Description	Type
145	7-4	AppSelCodeLane1	SCS0::DPConfigLane1 See Table 8-71	RW Rqd.
	3-1	DataPathIDLane1		
	0	ExplicitControlLane1		
146	7-4	AppSelCodeLane2	SCS0::DPConfigLane2 See Table 8-71	RW Rqd.
	3-1	DataPathIDLane2		
	0	ExplicitControlLane2		
147	7-4	AppSelCodeLane3	SCS0::DPConfigLane3 See Table 8-71	RW Rqd.
	3-1	DataPathIDLane3		
	0	ExplicitControlLane3		
148	7-4	AppSelCodeLane4	SCS0::DPConfigLane4 See Table 8-71	RW Rqd.
	3-1	DataPathIDLane4		
	0	ExplicitControlLane4		
149	7-4	AppSelCodeLane5	SCS0::DPConfigLane5 See Table 8-71	RW Rqd.
	3-1	DataPathIDLane5		
	0	ExplicitControlLane5		
150	7-4	AppSelCodeLane6	SCS0::DPConfigLane6 See Table 8-71	RW Rqd.
	3-1	DataPathIDLane6		
	0	ExplicitControlLane6		
151	7-4	AppSelCodeLane7	SCS0::DPConfigLane7 See Table 8-71	RW Rqd.
	3-1	DataPathIDLane7		
	0	ExplicitControlLane7		
152	7-4	AppSelCodeLane8	SCS0::DPConfigLane8 See Table 8-71	RW Rqd.
	3-1	DataPathIDLane8		
	0	ExplicitControlLane8		

3

4

8.9.3.3 Tx and Rx Signal Integrity Controls

The following control fields allow the host to pre-program signal integrity settings per lane, which the module uses instead of default values that are associated with the Application configured on the Data Path of the lane.

Note: When Versatile Control Sets providing an extended set of SI settings are supported and used, please refer to [6] for a description of the management memory map registers described here in this subsection.

These fields are without effect unless the lane specific ExplicitControl bits are set, as described in section 6.2.3.

These fields have no effect until the staged control set is applied to the Active Control Set.

See section 6.2.3 for the dependency of these fields on the value of the ExplicitControl bit.

See section 6.2.5 for definitions of valid signal integrity control settings.

Table 8-73 Staged Control Set 0, Tx Controls (Page 10h)

Byte	Bits	Field Name	Register Description	Type
153	7	AdaptiveInputEqEnableTx8	SCS0::AdaptiveInputEqEnableTx< i > Adaptive Tx input equalizer for host lane < i > 1b: Enable 0b: Disable (allow host-controlled equalization) Advertisement: 01h:161.3	RW Adv.
	6	AdaptiveInputEqEnableTx7		
	5	AdaptiveInputEqEnableTx6		
	4	AdaptiveInputEqEnableTx5		
	3	AdaptiveInputEqEnableTx4		
	2	AdaptiveInputEqEnableTx3		
	1	AdaptiveInputEqEnableTx2		
	0	AdaptiveInputEqEnableTx1		
154	7-6	AdaptiveInputEqRecallTx4	SCS0::AdaptiveInputEqRecallTx< i > Recall stored Tx input equalizer adaptation settings for host lane < i > when Staged Control Set is copied to Active Control Set. See section 6.2.5.4 for Store/Recall	RW Adv.
	5-4	AdaptiveInputEqRecallTx3		
	3-2	AdaptiveInputEqRecallTx2		
	1-0	AdaptiveInputEqRecallTx1		
155	7-6	AdaptiveInputEqRecallTx8	00b: do not recall 01b: recall buffer 1 10b: recall buffer 2 11b: reserved Advertisement: 01h:161.6-5	RW Adv.
	5-4	AdaptiveInputEqRecallTx7		
	3-2	AdaptiveInputEqRecallTx6		
	1-0	AdaptiveInputEqRecallTx5		
156	7-4	HostControlledInputEqTargetTx2	SCS0:: HostControlledInputEqTargetTx< i > Host-controlled Tx input equalizer target for host lane < i >. Encoding as described in Table 6-6	RW Adv.
	3-0	HostControlledInputEqTargetTx1		
157	7-4	HostControlledInputEqTargetTx4	Advertisement: 01h:161.2	RW Adv.
	3-0	HostControlledInputEqTargetTx3		
158	7-4	HostControlledInputEqTargetTx6		RW Adv.
	3-0	HostControlledInputEqTargetTx5		
159	7-4	HostControlledInputEqTargetTx8		RW Adv.
	3-0	HostControlledInputEqTargetTx7		
160	7	CDREnableTx8	SCS0::CDREnableTx< i > 1b: CDR enabled 0b: CDR bypassed Advertisement: 01h:161.0-1	RW Adv.
	6	CDREnableTx7		
	5	CDREnableTx6		
	4	CDREnableTx5		
	3	CDREnableTx4		
	2	CDREnableTx3		
	1	CDREnableTx2		
	0	CDREnableTx1		

Table 8-74 Staged Control Set 0, Rx Controls (Page 10h)

Byte	Bits	Name	Register Description	Type
161	7	CDREnableRx8	SCS0::CDREnableRx< i > 1b: CDR enabled 0b: CDR bypassed Advertisement: 01h:162.0-1	RW Adv.
	6	CDREnableRx7		
	5	CDREnableRx6		
	4	CDREnableRx5		
	3	CDREnableRx4		
	2	CDREnableRx3		
	1	CDREnableRx2		
	0	CDREnableRx1		
162	7-4	OutputEqPreCursorTargetRx2	SCS0::OutputEqPreCursorTargetRx< i > Rx output equalization pre-cursor target. Used only when both pre- and post-cursor targets are supported. See Table 6-7 and section 6.2.5.2	RW Adv.
	3-0	OutputEqPreCursorTargetRx1		
163	7-4	OutputEqPreCursorTargetRx4		
	3-0	OutputEqPreCursorTargetRx3		
164	7-4	OutputEqPreCursorTargetRx6		
	3-0	OutputEqPreCursorTargetRx5		
165	7-4	OutputEqPreCursorTargetRx8	Advertisement: 01h:162.4-3	RW Adv.
	3-0	OutputEqPreCursorTargetRx7		
166	7-4	OutputEqPostCursorTargetRx2	SCS0::OutputEqPostCursorTargetRx< i > Rx output equalization post-cursor target, or pre-cursor target when only pre-cursor setting is advertised. See Table 6-7 and section 6.2.5.2	RW Adv.
	3-0	OutputEqPostCursorTargetRx1		
167	7-4	OutputEqPostCursorTargetRx4		
	3-0	OutputEqPostCursorTargetRx3		
168	7-4	OutputEqPostCursorTargetRx6		
	3-0	OutputEqPostCursorTargetRx5		
169	7-4	OutputEqPostCursorTargetRx8	Advertisement: 01h:162.4-3	RW Adv.
	3-0	OutputEqPostCursorTargetRx7		
170	7-4	OutputAmplitudeTargetRx2	SCS0::OutputAmplitudeTargetRx< i > Rx output amplitude target See Table 6-8 and section 6.2.5.3	RW Adv.
	3-0	OutputAmplitudeTargetRx1		
171	7-4	OutputAmplitudeTargetRx4		
	3-0	OutputAmplitudeTargetRx3		
172	7-4	OutputAmplitudeTargetRx6	Advertisement: 01h:162.2	RW Adv.
	3-0	OutputAmplitudeTargetRx5		
173	7-4	OutputAmplitudeTargetRx8		
	3-0	OutputAmplitudeTargetRx7		
174-175	All	-	Reserved[2]	RO

8.9.3.4 Unidirectional Apply Staged Control Set Triggers

The unidirectional command trigger registers are supported only when the module supports independent reconfiguration of the Tx and Rx directions (i.e. when UnidirReconfigSupported is set). See also section 7.7.

The host shall access these registers only in DPSM states DPInitialized or DPActivated.

Note: A common use of this feature is during Fibre Channel link speed negotiation.

Table 8-75 Staged Control Set 0, Unidirectional Apply Triggers (Page 10h)

Byte	Bits	Field Name	Register Description	Type
176	7	ApplyImmediateTx8	SCS0::ApplyImmediateTx<i> 0b: No action for host lane <i> 1b: Trigger the Provision-and-Commission procedure for Tx using the Staged Control Set 0 settings for host lane <i>, with feedback provided in the ConfigStatusLane<i> field Restriction: This byte must be written in a single-byte WRITE Condition: UnidirReconfigSupported (01h:162.6) = 1b <i>Note: See Table 6-3 for preconditions and state dependencies</i>	WO Cnd.
	6	ApplyImmediateTx7		
	5	ApplyImmediateTx6		
	4	ApplyImmediateTx5		
	3	ApplyImmediateTx4		
	2	ApplyImmediateTx3		
	1	ApplyImmediateTx2		
	0	ApplyImmediateTx1		
177	7	ApplyImmediateRx8	SCS0::ApplyImmediateRx<i> 0b: No action for host lane <i> 1b: Trigger the Provision-and-Commission procedure for Rx using the Staged Control Set 0 settings for host lane <i>, with feedback provided in the ConfigStatusLane<i> field Restriction: This byte must be written in a single-byte WRITE Condition: UnidirReconfigSupported (01h:162.6) = 1b <i>Note: See Table 6-3 for preconditions and state dependencies</i>	WO Cnd.
	6	ApplyImmediateRx7		
	5	ApplyImmediateRx6		
	4	ApplyImmediateRx5		
	3	ApplyImmediateRx4		
	2	ApplyImmediateRx3		
	1	ApplyImmediateRx2		
	0	ApplyImmediateRx1		

8.9.4 Staged Control Set 1

Staged Control Set 1 is optional. If supported, it behaves exactly like Staged Control Set 0 (see section 8.9.3), except that it is zero initialized.

The module advertises support for Staged Control Set 1 in StagedSet1Supported (01h:162.5, see Table 8-53).

Note: Supporting Staged Control Set 1 may be useful for Applications where speed negotiation is performed.

8.9.4.1 Apply Staged Control Set Triggers

See section 8.9.3.1 for description and rationale.

Table 8-76 Staged Control Set 1, Apply Triggers (Page 10h)

Byte	Bits	Name	Register Description	Type
178	7	ApplyDPInitLane8	SCS1::ApplyDPInitLane<i> 0b: No action for host lane <i>	WO Adv.
	6	ApplyDPInitLane7	1b: Trigger the Provision procedure using the Staged Control Set 1 settings for host lane <i>, with feedback provided in the associated ConfigStatusLane<i> field	
	5	ApplyDPInitLane6		
	4	ApplyDPInitLane5		
	3	ApplyDPInitLane4		
	2	ApplyDPInitLane3		
	1	ApplyDPInitLane2		
	0	ApplyDPInitLane1	Restriction: This byte must be written in a single-byte WRITE <i>Note: See Table 6-3 for preconditions and state dependencies</i>	
179	7	ApplyImmediateLane8	SCS1::ApplyImmediate<i> 0b: No action for host lane <i>	WO Cnd.
	6	ApplyImmediateLane7	1b: Trigger the Provision-and-Commission procedure using Staged Control Set 1 settings for host lane <i>, with feedback provided in the associated ConfigStatusLane<i> field	
	5	ApplyImmediateLane6		
	4	ApplyImmediateLane5		
	3	ApplyImmediateLane4		
	2	ApplyImmediateLane3	Restriction: This byte must be written in a single-byte WRITE	
	1	ApplyImmediateLane2	Condition: SteppedConfigOnly (00h:2.6) = 0b <i>Note: See Table 6-3 for preconditions and state dependencies</i>	
	0	ApplyImmediateLane1		

8.9.4.2 Data Path Configuration (Application Assignment)

See section 8.9.3.2 and Table 8-71 for description.

Table 8-77 Staged Control Set 1, Data Path Configuration (Page 10h)

Byte	Bits	Name	Register Description	Type
180	7-4	AppSelCodeLane1	SCS1::DPCConfigLane1	RW Adv.
	3-1	DataPathIDLane1	See Table 8-71	
	0	ExplicitControlLane1		
181	7-4	AppSelCodeLane2	SCS1::DPCConfigLane2	RW Adv.
	3-1	DataPathIDLane2	See Table 8-71	
	0	ExplicitControlLane2		
182	7-4	AppSelCodeLane3	SCS1::DPCConfigLane3	RW Adv.
	3-1	DataPathIDLane3	See Table 8-71	
	0	ExplicitControlLane3		
183	7-4	AppSelCodeLane4	SCS1::DPCConfigLane4	RW Adv.
	3-1	DataPathIDLane4	See Table 8-71	
	0	ExplicitControlLane4		
184	7-4	AppSelCodeLane5	SCS1::DPCConfigLane5	RW Adv.
	3-1	DataPathIDLane5	See Table 8-71	
	0	ExplicitControlLane5		
185	7-4	AppSelCodeLane6	SCS1::DPCConfigLane6	RW Adv.
	3-1	DataPathIDLane6	See Table 8-71	
	0	ExplicitControlLane6		
186	7-4	AppSelCodeLane7	SCS1::DPCConfigLane7	RW Adv.
	3-1	DataPathIDLane7	See Table 8-71	
	0	ExplicitControlLane7		
187	7-4	AppSelCodeLane8	SCS1::DPCConfigLane8	RW Adv.
	3-1	DataPathIDLane8	See Table 8-71	

Byte	Bits	Name	Register Description	Type
	0	ExplicitControlLane8		

8.9.4.3 Tx and Rx Signal Integrity Controls

See section 8.9.3.3 for a description of the fully analogous controls in Staged Control Set 0.

Note: When Versatile Control Sets providing an extended set of SI settings are supported and used, please refer to [6] for a description of the management memory map registers described here in this subsection.

Table 8-78 Staged Control Set 1, Tx Controls (Page 10h)

Byte	Bits	Name	Register Description	Type
188	7	AdaptiveInputEqEnableTx8	SCS1::AdaptiveInputEqEnableTx<i> Adaptive Tx input equalizer for host lane <i> 1b: Enable 0b: Disable (allow host-controlled equalization)	RW Adv.
	6	AdaptiveInputEqEnableTx7		
	5	AdaptiveInputEqEnableTx6		
	4	AdaptiveInputEqEnableTx5		
	3	AdaptiveInputEqEnableTx4		
	2	AdaptiveInputEqEnableTx3		
	1	AdaptiveInputEqEnableTx2		
	0	AdaptiveInputEqEnableTx1		
189	7-6	AdaptiveInputEqRecallTx4	SCS1::AdaptiveInputEqRecallTx<i> Recall stored Tx input equalizer adaptation settings for host lane <i> when Staged Control Set is copied to Active Control Set. See section 6.2.5.4 for Store/Recall mechanism	RW Adv.
	5-4	AdaptiveInputEqRecallTx3		
	3-2	AdaptiveInputEqRecallTx2		
	1-0	AdaptiveInputEqRecallTx1		
190	7-6	AdaptiveInputEqRecallTx8	00b: do not recall 01b: recall from recall buffer 1 10b: recall from recall buffer 2 11b: reserved	RW Adv.
	5-4	AdaptiveInputEqRecallTx7		
	3-2	AdaptiveInputEqRecallTx6		
	1-0	AdaptiveInputEqRecallTx5		
191	7-4	HostControlledInputEqTargetTx2	SCS1::HostControlledInputEqTargetTx<i> Host-controlled Tx input equalizer target for host lane <i>. Encoding as described in Table 6-6	RW Adv.
	3-0	HostControlledInputEqTargetTx1		
192	7-4	HostControlledInputEqTargetTx4		RW Adv.
	3-0	HostControlledInputEqTargetTx3		
193	7-4	HostControlledInputEqTargetTx6		RW Adv.
	3-0	HostControlledInputEqTargetTx5		
194	7-4	HostControlledInputEqTargetTx8		RW Adv.
	3-0	HostControlledInputEqTargetTx7		
195	7	CDREnableTx8	SCS1::CDREnableTx<i> 1b: CDR enabled 0b: CDR bypassed	RW Adv.
	6	CDREnableTx7		
	5	CDREnableTx6		
	4	CDREnableTx5		
	3	CDREnableTx4		
	2	CDREnableTx3		
	1	CDREnableTx2		
	0	CDREnableTx1		

Table 8-79 Staged Control Set 1, Rx Controls (Page 10h)

Byte	Bits	Name	Register Description	Type
196	7	CDREnableRx8	SCS1::CDREnableRx<i> 1b: CDR enabled 0b: CDR bypassed	RW Adv.
	6	CDREnableRx7		
	5	CDREnableRx6		
	4	CDREnableRx5		
	3	CDREnableRx4		
	2	CDREnableRx3		
	1	CDREnableRx2		
	0	CDREnableRx1		
197	7-4	OutputEqPreCursorTargetRx2	SCS1::OutputEqPreCursorTargetRx<i> Rx output equalization pre-cursor target.	RW Adv.
	3-0	OutputEqPreCursorTargetRx1		
198	7-4	OutputEqPreCursorTargetRx4		RW

Byte	Bits	Name	Register Description	Type
	3-0	OutputEqPreCursorTargetRx3	Used only when both pre- and post-cursor targets are supported.	Adv.
199	7-4	OutputEqPreCursorTargetRx6		RW
	3-0	OutputEqPreCursorTargetRx5		Adv.
200	7-4	OutputEqPreCursorTargetRx8	See Table 6-7 and section 6.2.5.2	RW
	3-0	OutputEqPreCursorTargetRx7		Adv.
201	7-4	OutputEqPostCursorTargetRx2	SCS1::OutputEqPostCursorTargetRx< i >	RW
	3-0	OutputEqPostCursorTargetRx1	Rx output equalization post-cursor, or pre-cursor target when only pre-cursor setting is advertised.	Adv.
202	7-4	OutputEqPostCursorTargetRx4		RW
	3-0	OutputEqPostCursorTargetRx3		Adv.
203	7-4	OutputEqPostCursorTargetRx6	See Table 6-7 and section 6.2.5.2	RW
	3-0	OutputEqPostCursorTargetRx5		Adv.
204	7-4	OutputEqPostCursorTargetRx8		RW
	3-0	OutputEqPostCursorTargetRx7		Adv.
205	7-4	OutputAmplitudeTargetRx2	SCS1::OutputAmplitudeTargetRx< i >	RW
	3-0	OutputAmplitudeTargetRx1	Rx output amplitude encoding	Adv.
206	7-4	OutputAmplitudeTargetRx4	See Table 6-8	RW
	3-0	OutputAmplitudeTargetRx3		Adv.
207	7-4	OutputAmplitudeTargetRx6		RW
	3-0	OutputAmplitudeTargetRx5		Adv.
208	7-4	OutputAmplitudeTargetRx8		RW
	3-0	OutputAmplitudeTargetRx7		Adv.
209-210	All	-	Reserved[2]	

1

8.9.4.4 Unidirectional Apply Staged Control Set Triggers

The unidirectional command trigger registers are supported only when the module supports independent reconfiguration of the Tx and Rx directions (i.e. when **UnidirReconfigSupported** is set). See also section 7.7.

The host shall access these registers only in DPSM states DPInitialized or DPActivated.

Note: A common use of this feature is during Fibre Channel link speed negotiation.

Table 8-80 Staged Control Set 1, Unidirectional Apply Triggers (Page 10h)

Byte	Bits	Field Name	Register Description	Type
211	7	ApplyImmediateTx8	SCS1::ApplyImmediateTx< i >	WO Cnd.
	6	ApplyImmediateTx7	0b: No action for host lane < i >	
	5	ApplyImmediateTx6	1b: Trigger the Provision-and-Commission procedure for Tx using the Staged Control Set 1 settings for host lane < i >, with feedback provided in the ConfigStatusLane< i > field	
	4	ApplyImmediateTx5	Restriction: This byte must be written in a single-byte WRITE	
	3	ApplyImmediateTx4	Condition: UnidirReconfigSupported (01h:162.6) = 1b	
	2	ApplyImmediateTx3	<i>Note: See Table 6-3 for preconditions and state dependencies</i>	
	1	ApplyImmediateTx2		
	0	ApplyImmediateTx1		
212	7	ApplyImmediateRx8	SCS1::ApplyImmediateRx< i >	WO Cnd.
	6	ApplyImmediateRx7	0b: No action for host lane < i >	
	5	ApplyImmediateRx6	1b: Trigger the Provision-and-Commission procedure for Rx using the Staged Control Set 1 settings for host lane < i >, with feedback provided in the ConfigStatusLane< i > field	
	4	ApplyImmediateRx5	Restriction: This byte must be written in a single-byte WRITE	
	3	ApplyImmediateRx4	Condition: UnidirReconfigSupported (01h:162.6) = 1b	
	2	ApplyImmediateRx3	<i>Note: See Table 6-3 for preconditions and state dependencies</i>	
	1	ApplyImmediateRx2		
	0	ApplyImmediateRx1		

8

8.9.5 Lane-Specific Masks

The host can control which lane-specific Flags contribute to hardware Interrupt request generation by setting Mask bits in the registers described in Table 8-81. A Mask bit is allocated for each Flag.

See section 8.1.4.2 for more information on Masks, Flags, and Interrupt requests.

Table 8-81 Lane-Specific Masks (Page 10h)

Byte	Bits	Name	Register Description	Type
213	7	DPStateChangedMask8	DPStateChangedMask<i><i></i> Mask for DPStateChangedFlag <i><i></i> for Data Path of host lane <i><i></i>	RW Rqd.
	6	DPStateChangedMask7		
	5	DPStateChangedMask6		
	4	DPStateChangedMask5		
	3	DPStateChangedMask4		
	2	DPStateChangedMask3		
	1	DPStateChangedMask2		
	0	DPStateChangedMask1		
214	7	FailureMaskTx8	FailureMaskTx<i><i></i> Mask for FailureFlagTx <i><i></i> , affecting media lane <i><i></i> Advertisement: 01h:157.0	RW Adv.
	6	FailureMaskTx7		
	5	FailureMaskTx6		
	4	FailureMaskTx5		
	3	FailureMaskTx4		
	2	FailureMaskTx3		
	1	FailureMaskTx2		
	0	FailureMaskTx1		
215	7	LOSMaskTx8	LOSMaskTx<i><i></i> Mask for LOSFlagTx, lane <i><i></i> Advertisement: 01h:157.1	RW Adv.
	6	LOSMaskTx7		
	5	LOSMaskTx6		
	4	LOSMaskTx5		
	3	LOSMaskTx4		
	2	LOSMaskTx3		
	1	LOSMaskTx2		
	0	LOSMaskTx1		
216	7	CDRLOLMaskTx8	CDRLOLMaskTx<i><i></i> Mask for CDRLOLFlagTx <i><i></i> , lane <i><i></i> Advertisement: 01h:157.2	RW Adv.
	6	CDRLOLMaskTx7		
	5	CDRLOLMaskTx6		
	4	CDRLOLMaskTx5		
	3	CDRLOLMaskTx4		
	2	CDRLOLMaskTx3		
	1	CDRLOLMaskTx2		
	0	CDRLOLMaskTx1		
217	7	AdaptiveInputEqFailMaskTx8	AdaptiveInputEqFailMaskTx<i><i></i> Mask for AdaptiveInputEqFailFlagTx <i><i></i> , lane <i><i></i> Advertisement: 01h:157.3	RW Adv.
	6	AdaptiveInputEqFailMaskTx7		
	5	AdaptiveInputEqFailMaskTx6		
	4	AdaptiveInputEqFailMaskTx5		
	3	AdaptiveInputEqFailMaskTx4		
	2	AdaptiveInputEqFailMaskTx3		
	1	AdaptiveInputEqFailMaskTx2		
	0	AdaptiveInputEqFailMaskTx1		
218	7	OpticalPowerHighAlarmMaskTx8	OpticalPowerHighAlarmMaskTx<i><i></i> Mask for OpticalPowerHighAlarmFlagTx <i><i></i> , media lane <i><i></i> Advertisement: 01h:160.1	RW Adv.
	6	OpticalPowerHighAlarmMaskTx7		
	5	OpticalPowerHighAlarmMaskTx6		
	4	OpticalPowerHighAlarmMaskTx5		
	3	OpticalPowerHighAlarmMaskTx4		
	2	OpticalPowerHighAlarmMaskTx3		
	1	OpticalPowerHighAlarmMaskTx2		
	0	OpticalPowerHighAlarmMaskTx1		
219	7	OpticalPowerLowAlarmMaskTx8	OpticalPowerLowAlarmMaskTx<i><i></i> Mask for OpticalPowerLowAlarmFlagTx <i><i></i> ,	RW Adv.
	6	OpticalPowerLowAlarmMaskTx7		

Byte	Bits	Name	Register Description	Type
	5	OpticalPowerLowAlarmMaskTx6	media lane <i>	
	4	OpticalPowerLowAlarmMaskTx5		
	3	OpticalPowerLowAlarmMaskTx4	Advertisement: 01h:160.1	
	2	OpticalPowerLowAlarmMaskTx3		
	1	OpticalPowerLowAlarmMaskTx2		
	0	OpticalPowerLowAlarmMaskTx1		
220	7	OpticalPowerHighWarningMaskTx8	OpticalPowerHighWarningMaskTx< i > Mask for OpticalPowerHighWarningFlagTx< i >, media lane <i>	RW Adv.
	6	OpticalPowerHighWarningMaskTx7		
	5	OpticalPowerHighWarningMaskTx6		
	4	OpticalPowerHighWarningMaskTx5		
	3	OpticalPowerHighWarningMaskTx4	Advertisement: 01h:160.1	
	2	OpticalPowerHighWarningMaskTx3		
	1	OpticalPowerHighWarningMaskTx2		
	0	OpticalPowerHighWarningMaskTx1		
221	7	OpticalPowerLowWarningMaskTx8	OpticalPowerLowWarningMaskTx< i > Mask for OpticalPowerLowWarningFlagTx< i >, media lane <i>	RW Adv.
	6	OpticalPowerLowWarningMaskTx7		
	5	OpticalPowerLowWarningMaskTx6		
	4	OpticalPowerLowWarningMaskTx5		
	3	OpticalPowerLowWarningMaskTx4	Advertisement: 01h:160.1	
	2	OpticalPowerLowWarningMaskTx3		
	1	OpticalPowerLowWarningMaskTx2		
	0	OpticalPowerLowWarningMaskTx1		
222	7	LaserBiasHighAlarmMaskTx8	LaserBiasHighAlarmMaskTx< i > Mask for LaserBiasHighAlarmFlagTx< i >, media lane <i>	RW Adv.
	6	LaserBiasHighAlarmMaskTx7		
	5	LaserBiasHighAlarmMaskTx6		
	4	LaserBiasHighAlarmMaskTx5		
	3	LaserBiasHighAlarmMaskTx4	Advertisement: 01h:160.0	
	2	LaserBiasHighAlarmMaskTx3		
	1	LaserBiasHighAlarmMaskTx2		
	0	LaserBiasHighAlarmMaskTx1		
223	7	LaserBiasLowAlarmMaskTx8	LaserBiasLowAlarmMaskTx< i > Mask for LaserBiasLowAlarmFlagTx< i >, media lane <i>	RW Adv.
	6	LaserBiasLowAlarmMaskTx7		
	5	LaserBiasLowAlarmMaskTx6		
	4	LaserBiasLowAlarmMaskTx5		
	3	LaserBiasLowAlarmMaskTx4	Advertisement: 01h:160.0	
	2	LaserBiasLowAlarmMaskTx3		
	1	LaserBiasLowAlarmMaskTx2		
	0	LaserBiasLowAlarmMaskTx1		
224	7	LaserBiasHighWarningMaskTx8	LaserBiasHighWarningMaskTx< i > Mask for LaserBiasHighWarningFlagTx< i >, media lane <i>	RW Adv.
	6	LaserBiasHighWarningMaskTx7		
	5	LaserBiasHighWarningMaskTx6		
	4	LaserBiasHighWarningMaskTx5		
	3	LaserBiasHighWarningMaskTx4	Advertisement: 01h:160.0	
	2	LaserBiasHighWarningMaskTx3		
	1	LaserBiasHighWarningMaskTx2		
	0	LaserBiasHighWarningMaskTx1		
225	7	LaserBiasLowWarningMaskTx8	LaserBiasLowWarningMaskTx< i > Mask for LaserBiasLowWarningFlagTx< i >, media lane <i>	RW Adv.
	6	LaserBiasLowWarningMaskTx7		
	5	LaserBiasLowWarningMaskTx6		
	4	LaserBiasLowWarningMaskTx5		
	3	LaserBiasLowWarningMaskTx4	Advertisement: 01h:160.0	
	2	LaserBiasLowWarningMaskTx3		
	1	LaserBiasLowWarningMaskTx2		
	0	LaserBiasLowWarningMaskTx1		
226	7	LOSMaskRx8	LOSMaskRx< i > Mask for LOSFlagRx< i >, media lane <i>	RW Adv.
	6	LOSMaskRx7		
	5	LOSMaskRx6		

Byte	Bits	Name	Register Description	Type
	4	LOSMaskRx5	Advertisement: 01h:158.1	
	3	LOSMaskRx4		
	2	LOSMaskRx3		
	1	LOSMaskRx2		
	0	LOSMaskRx1		
227	7	CDRLOLMaskRx8	CDRLOLMaskRx< i> Mask for CDRLOLFlagRx< i>, media lane < i> Advertisement: 01h:158.2	RW Adv.
	6	CDRLOLMaskRx7		
	5	CDRLOLMaskRx6		
	4	CDRLOLMaskRx5		
	3	CDRLOLMaskRx4		
	2	CDRLOLMaskRx3		
	1	CDRLOLMaskRx2		
	0	CDRLOLMaskRx1		
228	7	OpticalPowerHighAlarmMaskRx8	OpticalPowerHighAlarmMaskRx< i> Mask for OpticalPowerHighAlarmFlagRx< i>, media lane < i> Advertisement: 01h:160.2	RW Adv.
	6	OpticalPowerHighAlarmMaskRx7		
	5	OpticalPowerHighAlarmMaskRx6		
	4	OpticalPowerHighAlarmMaskRx5		
	3	OpticalPowerHighAlarmMaskRx4		
	2	OpticalPowerHighAlarmMaskRx3		
	1	OpticalPowerHighAlarmMaskRx2		
	0	OpticalPowerHighAlarmMaskRx1		
229	7	OpticalPowerLowAlarmMaskRx8		
	6	OpticalPowerLowAlarmMaskRx7		
	5	OpticalPowerLowAlarmMaskRx6		
	4	OpticalPowerLowAlarmMaskRx5		
	3	OpticalPowerLowAlarmMaskRx4		
	2	OpticalPowerLowAlarmMaskRx3		
	1	OpticalPowerLowAlarmMaskRx2		
	0	OpticalPowerLowAlarmMaskRx1		
230	7	OpticalPowerHighWarningMaskRx8	OpticalPowerHighWarningMaskRx< i> Mask for OpticalPowerHighWarningFlagRx< i>, media lane < i> Advertisement: 01h:160.2	RW Adv.
	6	OpticalPowerHighWarningMaskRx7		
	5	OpticalPowerHighWarningMaskRx6		
	4	OpticalPowerHighWarningMaskRx5		
	3	OpticalPowerHighWarningMaskRx4		
	2	OpticalPowerHighWarningMaskRx3		
	1	OpticalPowerHighWarningMaskRx2		
	0	OpticalPowerHighWarningMaskRx1		
231	7	OpticalPowerLowWarningMaskRx8		
	6	OpticalPowerLowWarningMaskRx7		
	5	OpticalPowerLowWarningMaskRx6		
	4	OpticalPowerLowWarningMaskRx5		
	3	OpticalPowerLowWarningMaskRx4		
	2	OpticalPowerLowWarningMaskRx3		
	1	OpticalPowerLowWarningMaskRx2		
	0	OpticalPowerLowWarningMaskRx1		
232	7	OutputStatusChangedMaskRx8	OutputStatusChangedMaskRx< i> Mask for OutputStatusChangedFlagRx< i>, media lane < i>	RW Rqd.
	6	OutputStatusChangedMaskRx7		
	5	OutputStatusChangedMaskRx6		
	4	OutputStatusChangedMaskRx5		
	3	OutputStatusChangedMaskRx4		
	2	OutputStatusChangedMaskRx3		
	1	OutputStatusChangedMaskRx2		
	0	OutputStatusChangedMaskRx1		

8.10 Banked Page 11h (Lane Status and Data Path Status)

Page 11h is an optional Page that contains lane dynamic status bytes. All fields on Page 11h are read-only.

The module advertises support of Page 11h in the MemoryModel Bit 00h:2.7.

Note: Page 0:11h is mandatory for paged memory modules.

Page 11h may optionally be Banked. Each Bank of Page 11h refers to a group of 8 lanes.

Page 11h is subdivided into several areas as illustrated in the following table:

Table 8-82 Page 11h Overview

Byte	Size (bytes)	Subject Area	Description
128-131	4	Data Path States	State of Data Path State Machine associated with each lane
132-133	2	Lane Output Status	Output signal validity status (per lane)
134-153	20	Lane-specific Flags	Flags per lane or per Data Path
154-201	48	Lane-specific Monitors	Generic media side monitors for optical power and laser bias
202-205	4	Configuration Status	Status of configuration commands (Apply register writes)
206-234	29	Active Control Set	Nominal or actual Data Path configuration. See section 6.2.3
235-239	5	Data Path Conditions	Dynamic condition indications
240-255	16	Media Lane to media wavelengths and fibers mapping	Indicates the mapping of Media Lanes to Media Wavelengths and Fibers

8.10.1 Data Path States

The following fields report the Data Path States of the Data Path State Machines associated with each host lane.

Data Path States apply to both the host and the media interfaces, but are reported by host lane.

For Data Paths with multiple lanes, the module reports the same state for the lanes of each Data Path.

For unused lanes that are not part of a Data Path, the module reports DPDeactivated

An indication of DPDeactivated means that no Data Path is initialized on that lane.

Table 8-84 defines the Data Path State encodings.

Table 8-83 Lane-associated Data Path States (Page 11h)

Byte	Bit	Field Name	Register Description (DPStateHostLane <i><i></i>)	Type
128	7-4	DPStateHostLane2	Data Path State of host lane 2 (see Table 8-84)	RO
	3-0	DPStateHostLane1	Data Path State of host lane 1 (see Table 8-84)	Rqd.
129	7-4	DPStateHostLane4	Data Path State of host lane 4 (see Table 8-84)	RO
	3-0	DPStateHostLane3	Data Path State of host lane 3 (see Table 8-84)	Rqd.
130	7-4	DPStateHostLane6	Data Path State of host lane 6 (see Table 8-84)	RO
	3-0	DPStateHostLane5	Data Path State of host lane 5 (see Table 8-84)	Rqd.
131	7-4	DPStateHostLane8	Data Path State of host lane 8 (see Table 8-84)	RO
	3-0	DPStateHostLane7	Data Path State of host lane 7 (see Table 8-84)	Rqd.

Table 8-84 Data Path State Encoding

Encoding	State
0h	Reserved
1h	DPDeactivated (or unused lane)
2h	DPIInit
3h	DPDInIt
4h	DPAActivated
5h	DPTxTurnOn
6h	DPTxTurnOff
7h	DPIInitialized
8h-Fh	Reserved

8.10.2 Lane Output Status Indications

The **OutputStatusRx** and **OutputStatusTx** output status indication registers described in Table 8-85 report the status of the high speed outputs of a module, independent of the state of the DPSM instances associated with those output lanes (see also section 6.3.3).

The signal on an **Rx** output host lane is declared valid in the **OutputStatusRx** register (11h:132) only while the module is actually sending a **valid** signal to the host (see definition in section 3.3).

Example: In simple transceiver modules the Rx output is usually valid only after the associated Rx Media lane input has detected an input signal sufficient to initialize and activate all associated internal module resources, and only after those resources have been initialized and activated. A detected valid input Rx input signal is then, after Application specific processing in the module, forwarded to the host as a valid Rx output signal.

For each Rx output host lane status in the OutputStatusRx register there is an associated status change Flag (in Byte 11h:153), which the module sets on a change in the Rx Output Status (in Byte 11h:132) of that lane. The Mask associated with each Flag is found in Byte 10h:232 (see Table 8-81).

Note: The intent here is to avoid link flaps, by ensuring that the module declares its Rx outputs valid only when it is capable of sending a valid and stable signal to the host (after internal resources are initialized and activated).

The signal on an **Tx** output media lane is declared valid in the **OutputStatusTx** register (11h:133) only while the module is actually sending a **valid** media lane signal (see definition in section 3.3).

Note: It should be obvious that the module can neither judge nor ensure signal validity of higher layers that are not processed by the module.

Note: Unlike for the Rx output, there is no Tx output status change reporting Flag defined. This is deliberate because a status change does not trigger any regular configuration activity by the host.

Table 8-85 Lane-Specific Output Status (Page 11h)

Byte	Bit	Field Name	Register Description	Type
132	7	OutputStatusRx8	OutputStatusRx<i> 0b: Rx<i> output signal invalid or muted 1b: Rx<i> output signal valid	RO Rqd.
	6	OutputStatusRx7		
	5	OutputStatusRx6		
	4	OutputStatusRx5		
	3	OutputStatusRx4		
	2	OutputStatusRx3		
	1	OutputStatusRx2		
	0	OutputStatusRx1		
133	7	OutputStatusTx8	OutputStatusTx<i> 0b: Tx<i> output signal muted or invalid 1b: Tx<i> output signal valid	RO Rqd.
	6	OutputStatusTx7		
	5	OutputStatusTx6		
	4	OutputStatusTx5		
	3	OutputStatusTx4		
	2	OutputStatusTx3		
	1	OutputStatusTx2		
	0	OutputStatusTx1		

8.10.3 Lane-Specific Flags

This section of the Memory Map contains lane-specific Flags. These Flags provide a mechanism for reporting lane-specific status changes, operating failures, alarms, and warnings for monitored observables, or event occurrences. Each lane-specific Flag has an associated Mask.

The general behavior of Flags, Masks, and Interrupt generation is described in section 8.1.4.2.

The lane-specific Flags are defined in Table 8-86, Table 8-87 and Table 8-88.

Table 8-86 Lane-Specific State Changed Flags (Page 11h)

Byte	Bit	Field Name	Register Description	Type
134	7	DPStateChangedFlag8	DPStateChangedFlag< i > Latched Data Path State Changed Flag, host lane < i >	RO/COR Rqd.
	6	DPStateChangedFlag7		
	5	DPStateChangedFlag6		
	4	DPStateChangedFlag5		
	3	DPStateChangedFlag4		
	2	DPStateChangedFlag3		
	1	DPStateChangedFlag2		
	0	DPStateChangedFlag1		

2

3

Table 8-87 Lane-Specific Tx Flags (Page 11h)

Byte	Bit	Field Name	Register Description	Type
135	7	FailureFlagTx8	FailureFlagTx< i > Latched Tx Failure Flag, affecting media lane < i > This Flag indicates an internal failure that causes an unspecified malfunction in the Tx facility used by media lane < i >. <i>Note: This Flag was formerly named Tx Fault. See glossary for definitions of Fault and Failure.</i> Advertisement: 01h:157.0	RO/COR Adv.
	6	FailureFlagTx7		
	5	FailureFlagTx6		
	4	FailureFlagTx5		
	3	FailureFlagTx4		
	2	FailureFlagTx3		
	1	FailureFlagTx2		
	0	FailureFlagTx1		
136	7	LOSFlagTx8	LOSFlagTx< i > Latched Tx LOS Flag, host lane < i > Advertisement: 01h:157.1	RO/COR Adv.
	6	LOSFlagTx7		
	5	LOSFlagTx6		
	4	LOSFlagTx5		
	3	LOSFlagTx4		
	2	LOSFlagTx3		
	1	LOSFlagTx2		
	0	LOSFlagTx1		
137	7	CDRLOLFlagTx8	CDRLOLFlagTx< i > Latched Tx CDR LOL Flag, host lane < i > Advertisement: 01h:157.2	RO/COR Adv.
	6	CDRLOLFlagTx7		
	5	CDRLOLFlagTx6		
	4	CDRLOLFlagTx5		
	3	CDRLOLFlagTx4		
	2	CDRLOLFlagTx3		
	1	CDRLOLFlagTx2		
	0	CDRLOLFlagTx1		
138	7	AdaptiveInputEqFailFlagTx8	AdaptiveInputEqFailFlagTx< i > Latched Tx Adaptive Input Eq Fail, host lane < i > Advertisement: 01h:157.3	RO/COR Adv.
	6	AdaptiveInputEqFailFlagTx7		
	5	AdaptiveInputEqFailFlagTx6		
	4	AdaptiveInputEqFailFlagTx5		
	3	AdaptiveInputEqFailFlagTx4		
	2	AdaptiveInputEqFailFlagTx3		
	1	AdaptiveInputEqFailFlagTx2		
	0	AdaptiveInputEqFailFlagTx1		
139	7	OpticalPowerHighAlarmFlagTx8	OpticalPowerHighAlarmFlagTx< i > Latched Tx output power High Alarm, media lane < i > Advertisement: 01h:160.1	RO/COR Adv.
	6	OpticalPowerHighAlarmFlagTx7		
	5	OpticalPowerHighAlarmFlagTx6		
	4	OpticalPowerHighAlarmFlagTx5		
	3	OpticalPowerHighAlarmFlagTx4		
	2	OpticalPowerHighAlarmFlagTx3		
	1	OpticalPowerHighAlarmFlagTx2		
	0	OpticalPowerHighAlarmFlagTx1		
140	7	OpticalPowerLowAlarmFlagTx8	OpticalPowerLowAlarmFlagTx< i > Latched Tx output power Low alarm, media lane < i >	RO/COR Adv.
	6	OpticalPowerLowAlarmFlagTx7		
	5	OpticalPowerLowAlarmFlagTx6		
	4	OpticalPowerLowAlarmFlagTx5		

Byte	Bit	Field Name	Register Description	Type	
	3	OpticalPowerLowAlarmFlagTx4	Advertisement: 01h:160.1		
	2	OpticalPowerLowAlarmFlagTx3			
	1	OpticalPowerLowAlarmFlagTx2			
	0	OpticalPowerLowAlarmFlagTx1			
141	7	OpticalPowerHighWarningFlagTx8	OpticalPowerHighWarningFlagTx< i >	RO/COR Adv.	
	6	OpticalPowerHighWarningFlagTx7	Latched Tx output power High warning, media lane <i>		
	5	OpticalPowerHighWarningFlagTx6	Advertisement: 01h:160.1		
	4	OpticalPowerHighWarningFlagTx5			
	3	OpticalPowerHighWarningFlagTx4			
	2	OpticalPowerHighWarningFlagTx3			
	1	OpticalPowerHighWarningFlagTx2			
	0	OpticalPowerHighWarningFlagTx1			
142	7	OpticalPowerLowWarningFlagTx8	OpticalPowerLowWarningFlagTx< i >	RO/COR Adv.	
	6	OpticalPowerLowWarningFlagTx7	Latched Tx output power Low warning, media lane <i>		
	5	OpticalPowerLowWarningFlagTx6	Advertisement: 01h:160.1		
	4	OpticalPowerLowWarningFlagTx5			
	3	OpticalPowerLowWarningFlagTx4			
	2	OpticalPowerLowWarningFlagTx3			
	1	OpticalPowerLowWarningFlagTx2			
	0	OpticalPowerLowWarningFlagTx1			
143	7	LaserBiasHighAlarmFlagTx8	LaserBiasHighAlarmFlagTx< i >	RO/COR Adv.	
	6	LaserBiasHighAlarmFlagTx7	Latched Tx Bias High Alarm, media lane <i>		
	5	LaserBiasHighAlarmFlagTx6	Advertisement: 01h:160.0		
	4	LaserBiasHighAlarmFlagTx5			
	3	LaserBiasHighAlarmFlagTx4			
	2	LaserBiasHighAlarmFlagTx3			
	1	LaserBiasHighAlarmFlagTx2			
	0	LaserBiasHighAlarmFlagTx1			
144	7	LaserBiasLowAlarmFlagTx8	LaserBiasLowAlarmFlagTx< i >	RO/COR Adv.	
	6	LaserBiasLowAlarmFlagTx7	Latched Tx Bias Low alarm, media lane <i>		
	5	LaserBiasLowAlarmFlagTx6	Advertisement: 01h:160.0		
	4	LaserBiasLowAlarmFlagTx5			
	3	LaserBiasLowAlarmFlagTx4			
	2	LaserBiasLowAlarmFlagTx3			
	1	LaserBiasLowAlarmFlagTx2			
	0	LaserBiasLowAlarmFlagTx1			
145	7	LaserBiasHighWarningFlagTx8	LaserBiasHighWarningFlagTx< i >	RO/COR Adv.	
	6	LaserBiasHighWarningFlagTx7	Latched Tx Bias High warning, media lane <i>		
	5	LaserBiasHighWarningFlagTx6	Advertisement: 01h:160.0		
	4	LaserBiasHighWarningFlagTx5			
	3	LaserBiasHighWarningFlagTx4			
	2	LaserBiasHighWarningFlagTx3			
	1	LaserBiasHighWarningFlagTx2			
	0	LaserBiasHighWarningFlagTx1			
146	7	LaserBiasLowWarningFlagTx8	LaserBiasLowWarningFlagTx< i >	RO/COR Adv.	
	6	LaserBiasLowWarningFlagTx7	Latched Tx Bias Low warning, media lane <i>		
	5	LaserBiasLowWarningFlagTx6	Advertisement: 01h:160.0		
	4	LaserBiasLowWarningFlagTx5			
	3	LaserBiasLowWarningFlagTx4			
	2	LaserBiasLowWarningFlagTx3			
	1	LaserBiasLowWarningFlagTx2			
	0	LaserBiasLowWarningFlagTx1			

Table 8-88 Rx Flags (Page 11h)

Byte	Bit	Field Name	Register Description	Type
147	7	LOSFlagRx8	LOSFlagRx< i > Latched Rx LOS Flag, media lane < i > Advertisement: 01h:158.1	RO/COR Adv.
	6	LOSFlagRx7		
	5	LOSFlagRx6		
	4	LOSFlagRx5		
	3	LOSFlagRx4		
	2	LOSFlagRx3		
	1	LOSFlagRx2		
	0	LOSFlagRx1		
148	7	CDRLOLFlagRx8	CDRLOLFlagRx< i > Latched Rx CDR LOL Flag, media lane < i > Advertisement: 01h:158.2	RO/COR Adv.
	6	CDRLOLFlagRx7		
	5	CDRLOLFlagRx6		
	4	CDRLOLFlagRx5		
	3	CDRLOLFlagRx4		
	2	CDRLOLFlagRx3		
	1	CDRLOLFlagRx2		
	0	CDRLOLFlagRx1		
149	7	OpticalPowerHighAlarmFlagRx8	OpticalPowerHighAlarmFlagRx< i > Latched Rx input power High alarm, media lane < i > Advertisement: 01h:160.2	RO/COR Adv.
	6	OpticalPowerHighAlarmFlagRx7		
	5	OpticalPowerHighAlarmFlagRx6		
	4	OpticalPowerHighAlarmFlagRx5		
	3	OpticalPowerHighAlarmFlagRx4		
	2	OpticalPowerHighAlarmFlagRx3		
	1	OpticalPowerHighAlarmFlagRx2		
	0	OpticalPowerHighAlarmFlagRx1		
150	7	OpticalPowerLowAlarmFlagRx8	OpticalPowerLowAlarmFlagRx< i > Latched Rx input power Low alarm, media lane < i > Advertisement: 01h:160.2	RO/COR Adv.
	6	OpticalPowerLowAlarmFlagRx7		
	5	OpticalPowerLowAlarmFlagRx6		
	4	OpticalPowerLowAlarmFlagRx5		
	3	OpticalPowerLowAlarmFlagRx4		
	2	OpticalPowerLowAlarmFlagRx3		
	1	OpticalPowerLowAlarmFlagRx2		
	0	OpticalPowerLowAlarmFlagRx1		
151	7	OpticalPowerHighWarningFlagRx8	OpticalPowerHighWarningFlagRx< i > Latched Rx input power High warning, media lane < i > Advertisement: 01h:160.2	RO/COR Adv.
	6	OpticalPowerHighWarningFlagRx7		
	5	OpticalPowerHighWarningFlagRx6		
	4	OpticalPowerHighWarningFlagRx5		
	3	OpticalPowerHighWarningFlagRx4		
	2	OpticalPowerHighWarningFlagRx3		
	1	OpticalPowerHighWarningFlagRx2		
	0	OpticalPowerHighWarningFlagRx1		
152	7	OpticalPowerLowWarningFlagRx8	OpticalPowerLowWarningFlagRx< i > Latched Rx input power Low warning, media lane < i > Advertisement: 01h:160.2	RO/COR Adv.
	6	OpticalPowerLowWarningFlagRx7		
	5	OpticalPowerLowWarningFlagRx6		
	4	OpticalPowerLowWarningFlagRx5		
	3	OpticalPowerLowWarningFlagRx4		
	2	OpticalPowerLowWarningFlagRx3		
	1	OpticalPowerLowWarningFlagRx2		
	0	OpticalPowerLowWarningFlagRx1		

153	7	OutputStatusChangedFlagRx8	OutputStatusChangedFlagRx< i > Latched Output Status Changed Flag for Rx host lane < i >	RO/COR Rqd.
	6	OutputStatusChangedFlagRx7		
	5	OutputStatusChangedFlagRx6		
	4	OutputStatusChangedFlagRx5		
	3	OutputStatusChangedFlagRx4		
	2	OutputStatusChangedFlagRx3		
	1	OutputStatusChangedFlagRx2		
	0	OutputStatusChangedFlagRx1		

8.10.4 Lane-Specific Monitors

Real time lane monitoring may be performed for each transmit and receive lane and includes Tx output optical power, Rx input optical power, and Tx bias current.

The monitored observables defined here all have associated alarm and/or warning thresholds with associated threshold crossing alarm and/or warning Flags and Flags.

Alarm threshold values and warning threshold values have the same numerical value representation as the associated monitor values for which they specify threshold values.

Monitor results of supported lane monitors shall be within the relevant accuracy requirements when the module is in the DPActivated state.

Measured Tx laser bias current is represented as a 16-bit unsigned integer with the current defined as the full 16-bit value (0 to 65535) with LSB equal to 2 uA times the multiplier from Byte 01h:160. For a multiplier of 1, this yields a total measurement range of 0 to 131 mA.

Accuracy is Vendor Specific but must be better than +/-10% of the manufacture's nominal value over specified operating temperature and voltage.

Measured Rx input optical power is in mW and can represent either average received power or OMA, depending on the Rx Optical Power Measurement type described in Table 8-49. The parameter is encoded as a 16-bit unsigned integer with the power defined as the full 16-bit value (0 to 65535) with LSB equal to 0.1 uW, yielding a total measurement range of 0 to 6.5535 mW (~-40 to +8.2 dBm for non-zero values).

Absolute accuracy is dependent upon the exact optical wavelength. For the vendor specified wavelength, accuracy shall be better than +/-3 dB over specified temperature and voltage. This accuracy shall be maintained for input power levels up to the lesser of maximum transmitted or maximum received optical power per the appropriate standard. It shall be maintained down to the minimum transmitted power minus cable plant loss (insertion loss or passive loss) per the appropriate standard. Absolute accuracy beyond this minimum required received input optical power range is Vendor Specific.

Measured Tx optical power is the average power represented in mW. The parameter is encoded as a 16-bit unsigned integer with the power defined as the full 16-bit value (0 to 65535) with LSB equal to 0.1 uW, yielding a total measurement range of 0 to 6.5535 mW (~-40 to +8.2 dBm).

Accuracy is Vendor Specific but shall be better than +/-3 dB over specified temperature and voltage for the vendor specified wavelength.

Table 8-89 Media Lane-Specific Monitors (Page 11h)

Byte	Bit	Field Name	Register Description	Type
154-155	7-0	OpticalPowerTx1	U16 OpticalPowerTx< i > Internally measured Tx output optical power: in 0.1 uW increments, yielding a total measurement range of 0 to 6.5535 mW (~-40 to +8.2 dBm for non-zero values) Advertisement: 01h:160.1	RO Adv.
156-157	7-0	OpticalPowerTx2		
158-159	7-0	OpticalPowerTx3		
160-161	7-0	OpticalPowerTx4		
162-163	7-0	OpticalPowerTx5		
164-165	7-0	OpticalPowerTx6		
166-167	7-0	OpticalPowerTx7		
168-169	7-0	OpticalPowerTx8		
170-171	7-0	LaserBiasTx1	U16 LaserBiasTx< i > Internally measured Tx bias current monitor: in 2 uA increments, times the multiplier from Table 8-52. Advertisement: 01h:160.0	RO Adv.
172-173	7-0	LaserBiasTx2		
174-175	7-0	LaserBiasTx3		
176-177	7-0	LaserBiasTx4		
178-179	7-0	LaserBiasTx5		

Byte	Bit	Field Name	Register Description	Type
180-181	7-0	LaserBiasTx6		
182-183	7-0	LaserBiasTx7		
184-185	7-0	LaserBiasTx8		
186-187	7-0	OpticalPowerRx1	U16 OpticalPowerRx<1>	RO Adv.
188-189	7-0	OpticalPowerRx2	Internally measured Rx input optical power: in 0.1 uW increments, yielding a total measurement range of 0 to 6.5535 mW (~-40 to +8.2 dBm for non-zero values)	
190-191	7-0	OpticalPowerRx3		
192-193	7-0	OpticalPowerRx4		
194-195	7-0	OpticalPowerRx5		
196-197	7-0	OpticalPowerRx6		
198-199	7-0	OpticalPowerRx7		
200-201	7-0	OpticalPowerRx8	Advertisement: 01h:160.2	

1

2 8.10.5 Configuration Command Execution and Result Status (ConfigStatus)

3 As described in sections 6.2.3.3, 6.2.4, and 8.9.3.1, the host can command a configuration procedure to be
 4 executed by the module, by writing a lane trigger bit mask to an Apply register in that Staged Control Set where
 5 the desired configuration has been defined.

6 Both the **Provisioning** procedure and the **Provisioning-and-Commissioning** procedure may take significant
 7 time to execute and may be rejected, for a variety of reasons. Therefore, a synchronization protocol about
 8 currently ongoing execution and eventual result feedback is implemented, by means of the Configuration
 9 Command Execution and Result Status (**ConfigStatus**) register (see Table 8-90).

10 The field **ConfigStatusLane<i>** informs the host both on the **current command handling status** and on
 11 the **final result status** of the most recently accepted configuration command affecting lane <i>.

12 Configuration Command Handling Protocol

- 13 (9) In response to a WRITE access to an Apply register, the module executes four generic command
 14 handling activities: (1) command acceptance, (2) parameter validation, (3) command execution, and
 15 (4) result feedback. In the **command acceptance** step the module **checks** for the precondition that
 16 none of the relevant fields **ConfigStatusLane<i>** indicate ConfigInProgress. If this check passes, the
 17 module immediately sets the ConfigStatusLane<i> field of all Data Path lanes <i> to ConfigInProgress.
 18 Otherwise, the module aborts all further command handling steps for the relevant Data Path silently
 19 (without feedback).

20 *Note: Busy signaling to the host is specified via the separate ConfigStatus register, but module implementations
 21 are free to use the undefined Apply register bits as busy status bits for their own purpose or for debugging.*

- 22 (2) In the **parameter validation** step the module validates the settings in the relevant Staged Control Set. If
 23 the module determines that the desired configuration is invalid, it skips the following command execution step.

24 *Note: It is assumed that the following command execution step cannot fail after successful validation.*

- 25 (3) In the **command execution** step, the module actually performs the desired procedure. See section 8.9.3.1.

- 26 (4) In the **result feedback** step, the module **updates** the fields **ConfigStatusLane<i>** of all lanes of the
 27 Data Path with the actual command result status, thereby clearing ConfigInProgress. The presence of a result
 28 status value in ConfigStatusLane<i> indicates to the host that the module will accept a new configuration
 29 command for the Data Path of lane <i>.

30 The module populates the ConfigStatusLane<i> fields for all triggered lanes of a Data Path with the same value.

31 *Note: Recall that all configuration procedures are defined to operate on entire Data Paths, with the exception
 32 of hot reconfiguration of SI attributes by ApplyImmediate in DPActivated or DPInitialized, where triggers on a
 33 subset of lanes are allowed.*

34 The reporting priority of the result status codes is not specified.

Table 8-90 Configuration Command Status registers (Page 11h)

Byte	Bit	Field Name	Field Description	Type
202	7-4	ConfigStatusLane2	ConfigStatusLane<i> Configuration Command Execution / Result Status for the Data Path of host lane <i>, during and after the most recent configuration command. See Table 8-91 for the encoding of values. <i>Note: There is no feedback to the host when an Apply trigger is ignored after failed readiness test (when another configuration is still in progress)</i>	RO
	3-0	ConfigStatusLane1		Rqd.
203	7-4	ConfigStatusLane4		RO
	3-0	ConfigStatusLane3		Rqd.
204	7-4	ConfigStatusLane6		RO
	3-0	ConfigStatusLane5		Rqd.
205	7-4	ConfigStatusLane8		RO
	3-0	ConfigStatusLane7		Rqd.

The codes in Table 8-91 represent both the current command handling status (**in-progress**, **ready**) and the result status information (**success**, **rejection** due to validation failure), whereby the **ready** command handling status is implicitly indicated by the presence of result status information.

Table 8-91 Configuration Command Execution and Result Status Codes (Page 11h)

Encoding	Name	Value Description
0h	ConfigUndefined	No status information available (initial register value)
1h	ConfigSuccess	Positive Result Status: The last accepted configuration command has been completed successfully
2h	ConfigRejected	Negative Result Status (2h-Bh. Dh-Fh): Configuration rejected: unspecific validation failure
3h	ConfigRejectedInvalidAppSel	Configuration rejected: invalid AppSel code
4h	ConfigRejectedInvalidDataPath	Configuration rejected: invalid set of lanes for AppSel
5h	ConfigRejectedInvalidSI	Configuration rejected: invalid SI control settings
6h	ConfigRejectedLanesInUse	Configuration rejected: some lanes not in DPDeactivated
7h	ConfigRejectedPartialDataPath	Configuration rejected: lanes are only subset of DataPath
8h	-	Reserved (other validation failures)
9h	-	
Ah	-	
Bh	-	
Ch	ConfigInProgress	Execution Status: A configuration command is still being processed by the module; a new configuration command is ignored for this lane while ConfigInProgress.
Dh	-	Custom Configuration rejected for custom reasons
Eh	-	
Fh	-	

8.10.6 Active Control Set

The Active Control Set is required for all paged memory modules. It provides information on provisioned Data Path Configuration settings (see Table 8-93) as well as provisioned Signal Integrity settings for the Tx direction (see Table 8-94) and for the Rx direction (see Table 8-95).

Refer to section 6.2.3 for background on Control Sets and section 6.2.4 for background on provisioning.

Note: Before exiting the MgmtInit state, the Active Control Set is populated with a default Application and with default signal integrity settings.

The module updates the Active Control Set during execution of a valid provisioning command that was previously triggered by the host, normally using the ApplyDPInit or ApplyImmediate trigger fields in either Staged Control Set 0 (see section 8.9.3) or Staged Control Set 1 (see section 8.9.4).

Should the module support independent reconfiguration of the Tx and Rx directions (i.e. when the advertisement bit UnidirReconfigSupported is set), then there are direction specific ApplyImmediateTx and ApplyImmediateRx trigger fields in each supported Staged Control set, and the provisioned Data Path Configuration is then shown separately for the Tx direction (see Table 8-159) and for the Rx direction (see Table 8-160). In this case, the normal Data Path Configuration reporting fields described in Table 8-93 represent the provisioned settings for the Tx direction, i.e. they are a mirrored copy of the Data Path Configuration Tx fields described in Table 8-159.

Note: Hosts who do not use direction specific reconfiguration can therefore ignore the direction specific Data Path Configuration fields.

8.10.6.1 Provisioned Data Path Configuration (Application Assignment)

The following fields allow the host to infer the current baud rate, modulation format, and Data Path width, as well as other interface related specification elements, for each lane in the module.

Table 8-92 Data Path Configuration per Lane (DPConfigLane<i> Field)

Lane	Bits	Field	Field Description	Type
<i>	7-4	AppSelCode	ACS::AppSelCodeLane<i> Defines the Application assigned to the Data Path containing lane <i> by reference to the Application Descriptor of that Application: 0000b: lane <i> is unused and unassigned (not part of a Data Path) 0001b, ..., 1111b: lane <i> is part of a Data Path for the Application described in Application Descriptor with AppSel code AppSelCodeLane<i>	RO Rqd.
3-1	DataPathID		ACS::DataPathIDLane<i> Index of first lane in the Data Path containing lane <i> 000b: Lane 1 001b: Lane 2 111b: Lane 8	
0	ExplicitControl		ACS::ExplicitControlLane<i> 0b: Lane <i> SI settings are Application dependent 1b: Lane <i> SI settings are host defined	

Table 8-93 Active Control Set, Provisioned Data Path Configuration (Page 11h)

Byte	Bits	Field Name	Register Description	Type
206	7-4	AppSelCodeLane1	ACS::DPConfigLane1 See Table 8-92	RO Rqd.
	3-1	DataPathIDLane1		
	0	ExplicitControlLane1		
207	7-4	AppSelCodeLane2	ACS::DPConfigLane2 See Table 8-92	RO Rqd.
	3-1	DataPathIDLane2		
	0	ExplicitControlLane2		
208	7-4	AppSelCodeLane3	ACS::DPConfigLane3 See Table 8-92	RO Rqd.
	3-1	DataPathIDLane3		

Byte	Bits	Field Name	Register Description	Type
	0	ExplicitControlLane3		
209	7-4	AppSelCodeLane4	ACS::DPConfigLane4 See Table 8-92	RO Rqd.
	3-1	DataPathIDLane4		
	0	ExplicitControlLane4		
210	7-4	AppSelCodeLane5	ACS::DPConfigLane5 See Table 8-92	RO Rqd.
	3-1	DataPathIDLane5		
	0	ExplicitControlLane5		
211	7-4	AppSelCodeLane6	ACS::DPConfigLane6 See Table 8-92	RO Rqd.
	3-1	DataPathIDLane6		
	0	ExplicitControlLane6		
212	7-4	AppSelCodeLane7	ACS::DPConfigLane7 See Table 8-92	RO Rqd.
	3-1	DataPathIDLane7		
	0	ExplicitControlLane7		
213	7-4	AppSelCodeLane8	ACS::DPConfigLane8 See Table 8-92	RO Rqd.
	3-1	DataPathIDLane8		
	0	ExplicitControlLane8		

8.10.6.2 Provisioned Tx and Rx Signal Integrity Settings

The fields described in Table 8-94 and Table 8-95 report the provisioned signal integrity settings for each Tx and Rx lane, respectively.

Note: When Versatile Control Sets providing an extended set of SI settings are supported and used, please refer to [6] for a description of the management memory map registers described here in this subsection.

If the ExplicitControl bit for a lane was set in the Data Path Configuration (see Table 8-93) when the provisioning was triggered in a Staged Control Set, the contents of the registers for that lane described in Table 8-94 and Table 8-95 originate from corresponding registers in that Staged Control Set.

If the ExplicitControl bit for a lane was cleared when the provisioning was triggered, the contents of the registers for that lane in Table 8-94 and Table 8-95 were determined by the module according to the selected Application.

See section 6.2.5 for definitions of valid signal integrity control settings.

Table 8-94 Active Control Set, Provisioned Tx Controls (Page 11h)

Byte	Bits	Field Name	Register Description	Type
214	7	AdaptiveInputEqEnableTx8	ACS::AdaptiveInputEqEnableTx< i >	RO Adv.
	6	AdaptiveInputEqEnableTx7	Adaptive Tx input equalization for lane < i >	
	5	AdaptiveInputEqEnableTx6	1b: Enabled 0b: Disabled	
	4	AdaptiveInputEqEnableTx5	(Host-controlled Tx input equalization for lane < i > as per HostControlledInputEqTargetTx< i >, if supported)	
	3	AdaptiveInputEqEnableTx4		
	2	AdaptiveInputEqEnableTx3		
	1	AdaptiveInputEqEnableTx2		
	0	AdaptiveInputEqEnableTx1	Advertisement: 01h:161.3	
215	7-6	AdaptiveInputEqRecalledTx4	ACS::AdaptiveInputEqRecalledTx< i >	RO Adv.
	5-4	AdaptiveInputEqRecalledTx3	Recalled Tx Eq settings status for lane < i >	
	3-2	AdaptiveInputEqRecalledTx2	00b: settings are not recalled	
	1-0	AdaptiveInputEqRecalledTx1	01b: settings have been recalled from recall buffer 1	
216	7-6	AdaptiveInputEqRecalledTx8	10b: settings have been recalled from recall buffer 2	RO Adv.
	5-4	AdaptiveInputEqRecalledTx7	11b: reserved	
	3-2	AdaptiveInputEqRecalledTx6	<i>Note: See Table 8-73 and Table 8-78 for the recall settings in the staged control sets.</i>	
	1-0	AdaptiveInputEqRecalledTx5	Advertisement: 01h:161.6-5	
217	7-4	HostControlledInputEqTargetTx2	ACS:: HostControlledInputEqTargetTx< i >	RO Adv.
	3-0	HostControlledInputEqTargetTx1	Host-controlled Tx input equalization target for lane < i >	
218	7-4	HostControlledInputEqTargetTx4	encoded as defined in Table 6-6	RO Adv.
	3-0	HostControlledInputEqTargetTx3		
219	7-4	HostControlledInputEqTargetTx6	Advertisement: 01h:161.2	RO Adv.
	3-0	HostControlledInputEqTargetTx5		

Byte	Bits	Field Name	Register Description	Type
220	7-4	HostControlledInputEqTargetTx8		RO
	3-0	HostControlledInputEqTargetTx7		Adv.
221	7	CDREnableTx8	ACS::CDREnableTx< i >	RO
	6	CDREnableTx7	1b: CDR enabled 0b: CDR bypassed	Adv.
	5	CDREnableTx6		
	4	CDREnableTx5		
	3	CDREnableTx4	Advertisement: 01h:161.1	
	2	CDREnableTx3		
	1	CDREnableTx2		
	0	CDREnableTx1		

1

2

Table 8-95 Active Control Set, Provisioned Rx Controls (Page 11h)

Byte	Bits	Name	Register Description	Type
222	7	CDREnabledRx8	ACS::CDREnableRx< i >	RO
	6	CDREnabledRx7	1b: CDR enabled 0b: CDR bypassed	Adv.
	5	CDREnabledRx6		
	4	CDREnabledRx5		
	3	CDREnabledRx4	Advertisement: 01h:162.1	
	2	CDREnabledRx3		
	1	CDREnabledRx2		
	0	CDREnabledRx1		
223	7-4	OutputEqPreCursorTargetRx2	ACS::OutputEqPreCursorTargetRx< i >	RO
	3-0	OutputEqPreCursorTargetRx1	Rx output pre-cursor equalization for lane < i > encoded as per Table 6-7	Adv.
224	7-4	OutputEqPreCursorTargetRx4		RO
	3-0	OutputEqPreCursorTargetRx3		Adv.
225	7-4	OutputEqPreCursorTargetRx6	Advertisement: 01h:162.4-3	RO
	3-0	OutputEqPreCursorTargetRx5		Adv.
226	7-4	OutputEqPreCursorTargetRx8		RO
	3-0	OutputEqPreCursorTargetRx7		Adv.
227	7-4	OutputEqPostCursorTargetRx2	ACS::OutputEqPostCursorTargetRx< i >	RO
	3-0	OutputEqPostCursorTargetRx1	Rx output post-cursor equalization for lane < i > encoded as per Table 6-7	Adv.
228	7-4	OutputEqPostCursorTargetRx4		RO
	3-0	OutputEqPostCursorTargetRx3		Adv.
229	7-4	OutputEqPostCursorTargetRx6	Advertisement: 01h:162.4-3	RO
	3-0	OutputEqPostCursorTargetRx5		Adv.
230	7-4	OutputEqPostCursorTargetRx8		RO
	3-0	OutputEqPostCursorTargetRx7		Adv.
231	7-4	OutputAmplitudeTargetRx2	ACS::OutputAmplitudeTargetRx< i >	RO
	3-0	OutputAmplitudeTargetRx1	Rx output amplitude level for lane < i > encoded as per Table 6-8	Adv.
232	7-4	OutputAmplitudeTargetRx4		RO
	3-0	OutputAmplitudeTargetRx3		Adv.
233	7-4	OutputAmplitudeTargetRx6	Advertisement: 01h:162.2	RO
	3-0	OutputAmplitudeTargetRx5		Adv.
234	7-4	OutputAmplitudeTargetRx8		RO
	3-0	OutputAmplitudeTargetRx7		Adv.

3

8.10.7 Data Path Conditions

After copying the settings for the lanes of a Data Path from a Staged Control Set to the Active Control Set during a successful **Provision** procedure triggered by ApplyDPIInit (see Table 6-3 and Table 6-4 in section 6.2.4.2) the module simultaneously sets the bits **DPIInitPendingLane**< i > for the lanes < i > of that Data Path.

This indicates to the host that the Data Path configuration associated with Lane< i > has been updated (not necessarily modified) in the Active Control Set, whereas the subsequent transit through the DPInit state that will eventually commit the nominal settings in the Active Control Set to hardware is still pending.

When **intervention-free** reconfiguration is supported (by default, SteppedConfigOnly=0), these bits may also cause **DPSM state transitions** without further host intervention, as described in section 6.3.3.1.

Table 8-96 Data Path Conditions (Page 11h)

Byte	Bits	Field Name	Register Description	Type
235	7	DPIInitPendingLane8	DPIInitPendingLane < i > 0b: DPInit not pending 1b: DPInit not yet executed after successful ApplyDPIInit, hence the Active Control Set content may deviate from the actual hardware configuration.	RO Rqd.
	6	DPIInitPendingLane7		
	5	DPIInitPendingLane6		
	4	DPIInitPendingLane5		
	3	DPIInitPendingLane4		
	2	DPIInitPendingLane3		
	1	DPIInitPendingLane2		
	0	DPIInitPendingLane1		
236-239			Reserved[4]	RO

8.10.8 Media Lane to Media Wavelength and Fiber Mapping

Table 8-97 describes the advertising fields to define the mapping of media lanes to media wavelengths and physical fibers for muxed or WDM implementations.

For WDM applications the shortest wavelength is always designated media wavelength 1 and starting from shortest wavelength through the longest all others are listed consecutively.

The fiber numbering and naming used in Table 8-97 is as defined in the appropriate hardware specification.

See the relevant hardware specification and Table 8-23, Table 8-38, and Table 8-40 for mapping constraints.

A mapping advertised in Table 8-97 that violates any of these constraints is invalid.

Table 8-97 Media Lane to Media Wavelength and Fiber mapping (Page 11h)

Byte	Bits	Name	Register Description	Type
240	7-4	MediaLaneToWavelengthMappingTx1	MediaLaneToWavelengthMappingTx < i > Mapping of media lane < i >	RO
	3-0	MediaLaneToFiberMappingTx1	0000: Mapping unknown or undefined 0001: Maps to media wavelength 1	Cnd.
241	7-4	MediaLaneToWavelengthMappingTx2	0010: Maps to media wavelength 2	RO
	3-0	MediaLaneToFiberMappingTx2	0011: Maps to media wavelength 3	Cnd.
242	7-4	MediaLaneToWavelengthMappingTx3	0100: Maps to media wavelength 4	RO
	3-0	MediaLaneToFiberMappingTx3	0101: Maps to media wavelength 5	Cnd.
243	7-4	MediaLaneToWavelengthMappingTx4	0110: Maps to media wavelength 6	RO
	3-0	MediaLaneToFiberMappingTx4	0111: Maps to media wavelength 7	Cnd.
244	7-4	MediaLaneToWavelengthMappingTx5	1000: Maps to media wavelength 8	RO
	3-0	MediaLaneToFiberMappingTx5	1001-1111: Reserved	Cnd.
245	7-4	MediaLaneToWavelengthMappingTx6	MediaLaneToFiberMappingTx < i > Mapping of media lane < i >	RO
	3-0	MediaLaneToFiberMappingTx6	0000: Mapping unknown or undefined 0001: Maps to media fiber 1 or TR1	Cnd.
246	7-4	MediaLaneToWavelengthMappingTx7	0010: Maps to media fiber 2 or RT1	RO
	3-0	MediaLaneToFiberMappingTx7	0011: Maps to media fiber 3 or TR2	Cnd.
247	7-4	MediaLaneToWavelengthMappingTx8	0100: Maps to media fiber 4 or RT2	RO
	3-0	MediaLaneToFiberMappingTx8		Cnd.
248	7-4	MediaLaneToWavelengthMappingRx1		RO
	3-0	MediaLaneToFiberMappingRx1		Cnd.
249	7-4	MediaLaneToWavelengthMappingRx2		RO
	3-0	MediaLaneToFiberMappingRx2		Cnd.

Byte	Bits	Name	Register Description	Type
250	7-4	MediaLaneToWavelengthMappingRx3	0101: Maps to media fiber 5 or TR3 0110: Maps to media fiber 6 or RT3 0111: Maps to media fiber 7 or TR4 1000: Maps to media fiber 8 or RT4 1001-1111: Reserved	RO
	3-0	MediaLaneToFiberMappingRx3		Cnd.
251	7-4	MediaLaneToWavelengthMappingRx4	0101: Maps to media fiber 5 or TR3 0110: Maps to media fiber 6 or RT3 0111: Maps to media fiber 7 or TR4 1000: Maps to media fiber 8 or RT4 1001-1111: Reserved	RO
	3-0	MediaLaneToFiberMappingRx4		Cnd.
252	7-4	MediaLaneToWavelengthMappingRx5	0101: Maps to media fiber 5 or TR3 0110: Maps to media fiber 6 or RT3 0111: Maps to media fiber 7 or TR4 1000: Maps to media fiber 8 or RT4 1001-1111: Reserved	RO
	3-0	MediaLaneToFiberMappingRx5		Cnd.
253	7-4	MediaLaneToWavelengthMappingRx6	0101: Maps to media fiber 5 or TR3 0110: Maps to media fiber 6 or RT3 0111: Maps to media fiber 7 or TR4 1000: Maps to media fiber 8 or RT4 1001-1111: Reserved	RO
	3-0	MediaLaneToFiberMappingRx6		Cnd.
254	7-4	MediaLaneToWavelengthMappingRx7	0101: Maps to media fiber 5 or TR3 0110: Maps to media fiber 6 or RT3 0111: Maps to media fiber 7 or TR4 1000: Maps to media fiber 8 or RT4 1001-1111: Reserved	RO
	3-0	MediaLaneToFiberMappingRx7		Cnd.
255	7-4	MediaLaneToWavelengthMappingRx8	0101: Maps to media fiber 5 or TR3 0110: Maps to media fiber 6 or RT3 0111: Maps to media fiber 7 or TR4 1000: Maps to media fiber 8 or RT4 1001-1111: Reserved	RO
	3-0	MediaLaneToFiberMappingRx8		Cnd.

8.11 Banked Page 12h (Tunable Laser Control and Status)

Page 12h is an optional Page for laser tuning control, status, and Flags, for transmitters with tunable technology.

The module advertises support of Page 12h in Bit 01h:155.6 (see Table 8-50).

Page 12h may optionally be Banked. Each Bank of Page 12h refers to 8 media lanes.

Table 8-98 Page 12h Overview

Byte	Size (bytes)	Subject Area	Description
128-135	8	Grid Spacings	array with one Byte per media lane
136-151	8 x 2	Channel Offset Numbers	array with one S16 Word per media lane
152-167	8 x 2	Fine Tuning Offsets	array with one S16 Word per media lane
168-199	8 x 4	Laser Frequencies	array with one U32 double word per media lane
200-215	8 x 2	Target Output Power	array with one S16 word per media lane
216-221	6	-	Reserved[6]
222-229	8	Status Indicators	array with one Byte per media lane
230	1	Flag Summary	one Bit per media lane
231-238	8	Flags	array with one Byte per media lane
239-246	8	Masks (default: 1)	array with one Byte per media lane, Masks all set by default
247-255	9	-	Reserved[9] Note: This page has no Page Checksum

In Byte register arrays with one byte per lane the lowest byte address represents lane 1.

In Bit arrays (within a Byte) with one bit per lane, the least significant bit 0 represents lane 1.

Table 8-99 Laser tuning, status, and Flags for tunable transmitters (Page 12h)

Byte	Bit	Field Name	Field Description	Type
128-135	7-4	GridSpacingTx<n>	Selected grid spacing of media lane <n>=1-8 0000b: 3.125 GHz 0001b: 6.25 GHz 0010b: 12.5 GHz 0011b: 25 GHz 0100b: 50 GHz 0101b: 100 GHz 0110b: 33 GHz 0111b: 75 GHz 1000b: 150GHz 9-14: Reserved 1111b: Not available	RW Rqd.
	3-1	-	Reserved	RO
	0	FineTuningEnableTx<n>	Bool: fine-tuning enabled for media lane <n>=1-8 0b: Fine-tuning disabled 1b: Fine-tuning enabled	RW Rqd.
136-151	7-0	ChannelNumberTx<n>	S16 Channel (offset) Number for media lane <n>=1-8 <i>The meaning of the signed channel (offset) number and its dependence on the selected grid spacing is defined in section 8.7.</i>	RW Rqd.
152-167	7-0	FineTuningOffsetTx<n>	S16 fine-tuning frequency offset for media lane <n>=1-8 in units of 0.001 GHz	RW Rqd.
168-199	7-0	CurrentLaserFrequencyTx<n>	U32 Current frequency for media lane <n>=1-8 in units of 0.001 GHz	RO Rqd.
200-215	7-0	TargetOutputPowerTx<n>	S16 Target programmable output power for media lane <n>=1-8 in units of 0.01 dBm	RW Rqd.
216-221	7-0	-	Reserved[6]	RO Rqd.
222-229	7-2	-	Reserved	RO Rqd.
	1	TuningInProgressTx<n>	Bool: Status indication for tuning in progress on media lane <n>=1-8	RO Rqd.

Byte	Bit	Field Name	Field Description	Type
	0	WavelengthUnlockStatusTx< n >	0b/1b: Tuning not in progress/in progress Tuning is in progress when the laser is tuning to an on-grid channel, or fine tuning to a frequency offset, or tuning to target output power.	
			Bool: Unlocked status indication for laser wavelength on media lane < n >=1-8 0b/1b: Wavelength locked/unlocked	RO Rqd.
230	7-0	LaserTuningFlagSummaryTx< n >	Laser tuning Flag summary for media lane < n >=1-8. The bit < n >-1 is set if and only if any of the Flags in Bytes 231-238 are 1 for the particular Lane < n >. <i>Note: The host should react to an interrupt by reading this Byte to indicate which Lane(s) are responsible for the interrupt. The host should then read the corresponding Lane Byte (231-238) to determine the specific interrupt and to clear the latch.</i>	RO Rqd.
231-238	7-6	-	Reserved	RO Rqd.
	5	TargetOutputPowerOORFlagTx< n >	Latched Flag indicating that a target output power value outside the allowed range was entered for media lane < n >=1-8	RO/COR Rqd.
	4	FineTuningOutOfRangeFlagTx< n >	Latched Flag indicating that a fine-tuning value outside the allowed range was given for media lane < n >=1-8	RO/COR Rqd.
	3	TuningNotAcceptedFlagTx< n >	Latched Flag indicating a failed tuning operation for media lane < n >=1-8: the module was temporarily unable to serve a tuning request in its current state, in response to the host programming grid or channel. The Flag indicates that the nominally programmed grid or channel do not match the actual condition of the laser. The Flag clears when channel or grid have been successfully re-programmed. <i>Note. This Flag may e.g. be asserted by the module if the host attempts to change the target output power or fine tune frequency while tuning in progress.</i>	RO/COR Rqd.
	2	InvalidChannelNumberFlagTx< n >	Latched Flag indicating that ChannelNumberTx< n > was not in advertised range of the channel spacing selected on media lane < n >=1-8	RO/COR Rqd.
	1	WavelengthUnlockedFlagTx< n >	Latched Flag version of WavelengthUnlockedTx< n >	RO/COR Rqd.
	0	TuningCompleteFlagTx< n >	Latched Flag indicating laser tuning has been completed for media lane < n >=1-8	RO/COR Rqd.
239-246	7-6	-	Reserved	RO Rqd.
	5	TargetOutputPowerOORMaskTx< n >	Mask for TargetOutputPowerOORFlagTx< n > Default: 1	RW Rqd.
	4	FineTuningOutOfRangeMaskTx< n >	Mask for FineTuningOutOfRangeFlagTx< n > Default: 1	RW Rqd.
	3	TuningNotAcceptedMaskTx< n >	Mask for TuningNotAcceptedFlagTx< n >=1-8 Default: 1	RW Rqd.
	2	InvalidChannelMaskTx< n >	Mask for InvalidChannelNumberFlagTx< n > Default: 1	RW Rqd.
	1	WavelengthUnlockedMaskTx< n >	Mask for WavelengthUnlockedFlagTx< n > Default: 1	RW Rqd.
	0	TuningCompleteMaskTx< n >	Mask for TuningCompleteFlagTx< n > Default: 1	RW Rqd.
247-255	7-0	-	Reserved[9]	RO

8.12 Banked Page 13h (Module Performance Diagnostics Control)

Pages 13h and 14h are optional Pages containing module diagnostic control and result status fields, respectively.

The module advertises support of Pages 13h and 14h jointly in Bit 01h:142.5 (see Table 8-46).

Page 13h may optionally be Banked. Each Bank of Page 13h refers to 8 lanes.

Module Performance Diagnostics allows the host to control and evaluate two types of measurements:

- **intrusive** measurements of **error performance metrics** using pattern generators and checkers
 - U64 bit error counts and total bits
 - F16 bit error ratio (BER)
- **non-intrusive** measurements or estimations of **physical channel metrics** (SNR)

For counting-based error performance metrics, two kinds of controlling the measurement interval are possible

- **host controlled** for on-demand single-shot or quasi-periodic measurements (with stop-restart gap)
- **module controlled** gating interval for single-shot or periodic measurements (gap-free diagnostics)

For both types of interval control, **progressive updates** may be available that also occur periodically using a subinterval of the overall single-shot or periodic measurement interval (gating period).

The host can read selected measurement results in the Diagnostics Data area of Page 14h (see Table 8-125). See Table 8-129 for the detailed list of selectable Diagnostics measurement results and Table 8-103 for the associated advertisement fields.

The actually available measurement methods depend on advertised module capabilities as defined in the following subsections and described in section 8.12.11.

Page 13h is subdivided into subject areas as illustrated in the following table:

Table 8-100 Page 13h Overview

Byte	Size (bytes)	Subject Area	Description
128	1	Loopback capabilities	Diagnostics capability advertisements
129	1	Diagnostics measurement capabilities	
130	1	Diagnostic reporting capabilities	
131	1	Pattern Generation and Checking locations	
132-142	11	Pattern Generation and Checking capabilities	
143	1	-	Reserved[1] for module advertisements
144-151	8	Pattern Generator controls, host side	Host controls
152-159	8	Pattern Generator controls, media side	
160-167	8	Pattern Checker controls, host side	
168-175	8	Pattern Checker controls, media side	
176-179	4	Clocking and Measurement controls	
180-183	4	Loopback controls	Auto-restart detection and state storage
184-191	8	Host Scratchpad Area	
192-195	4	-	Reserved[4]
196-205	10	-	Custom[10]
206-223	18	Masks for Diagnostics Flags	Masks for Flags in Bytes 14h:132-149
224-255	32	User Pattern	User defined 32-byte pattern

8.12.1 Loopback Capabilities Advertisement

Four different types of loopback are defined by this specification. Figure 8-3 illustrates each type characterized by the **location of the loopback** on Host or Media Side and the **direction of the signal** being looped-back.

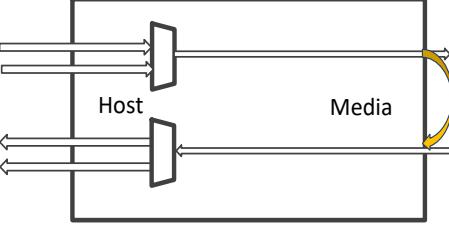
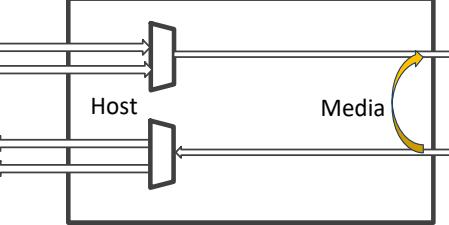
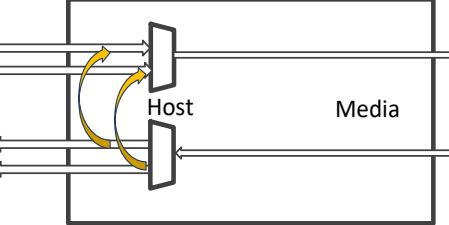
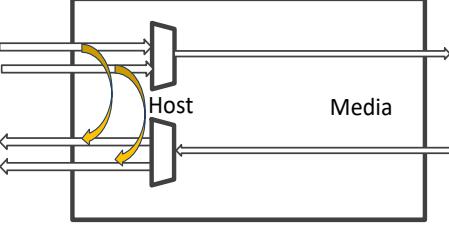
Name	Illustration
Media Side Output Loopback (only one media lane shown) advertised in 13h:128.0	
Media Side Input Loopback (only one media lane shown) advertised in 13h:128.1	
Host Side Output Loopback (only one media lane shown) advertised in 13h:128.2	
Host Side Input Loopback (only one media lane shown) advertised in 13h:128.3	

Figure 8-3 Loopback Type Illustrations

The loopback capabilities of the module are advertised as described in 13h:128 (see Table 8-101).

1

Table 8-101 Loopback Capabilities (Page 13h)

Byte	Bits	Field Name	Field Description	Type
128	7	-	Reserved	RO Rqd.
	6	SimultaneousHostAndMediaSideLoopbacks	0b: not supported 1b: supported	
	5	PerLaneMediaSideLoopbacks		
	4	PerLaneHostSideLoopbacks		
	3	HostSideInputLoopback		
	2	HostSideOutputLoopback		
	1	MediaSideInputLoopback		
	0	MediaSideOutputLoopback		

8.12.2 Diagnostics Measurement Capabilities Advertisement

2 Diagnostics measurement capabilities of the module are advertised in Table 8-102.

3

Table 8-102 Diagnostics Measurement Capabilities (Page 13h)

Byte	Bits	Field Name	Field Description	Type
129	7-6	GatingSupport	Gating (measurement over a given time interval) is 00b: Not supported (host defined measurement time) 01b: Supported with time accuracy <= 2 ms 10b: Supported with time accuracy <= 20 ms 11b: Supported with time accuracy > 20 ms <i>Note: The accuracy specifies the interval length error magnitude in the real-time time base of the module. Interval length may vary, e.g. by task scheduling jitter in module firmware.</i> <i>Note: The total bits counted and reported can be used as an accurate measure of the actual time interval length</i>	RO Rqd.
	5	GatingResultsSupported	Gating result statistics selectable by DiagnosticsSelector (14h:128) values 11h-15h are 0b: Not supported 1b: Supported	
	4	PeriodicUpdatesSupported	Real-time statistics selectable by DiagnosticsSelector (14h:128) values 01h-06h are periodically updated during measurement 0b: no periodic update during measurement 1b: periodic update during measurement <i>Note: The update rate of periodic update is selected by UpdatePeriodSelect (13h:177.0)</i> <i>Note: Single-shot measurements (un-gated non-periodic) can be started by enabling checkers and stopped by ResetErrorInformation (13h:177.5)</i>	
	3	PerLaneGatingTimersSupported	0b: Only two global gating timers are available for all lanes on all Banks, one for Host Side Measurements and one for Media Side Measurements. 1b: Per lane gating timers are supported in all Banks	
	2	AutoRestartGatingSupported	0b: AutoRestartGating control (13h:177.4) not supported 1b: AutoRestartGating control (13h:177.4) supported	
	1-0	-	Reserved	

4

1 **8.12.3 Diagnostic Reporting Capabilities Advertisement**

2 The diagnostic reporting capabilities of the module are advertised in Table 8-103.

3 **Table 8-103 Diagnostic Reporting Capabilities (Page 13h)**

Byte	Bits	Field Name	Field Description	Type
130	7	MediaSideFEC	1b: Supported (PRBS error information available) 0b: Not supported	RO Rqd.
	6	HostSideFEC	1b: Supported (PRBS error information available) 0b: Not supported	
	5	MediaSideInputSNRMeasurement	Indicates if media side SNR measurement reported via DiagnosticsSelector value 06h is supported (see Table 8-127) 1b: Supported 0b: Not supported	
	4	HostSideInputSNRMeasurement	Indicates if host side SNR measurement reported via Diagnostics Selection value 06h is supported (Byte 14h:128, see Table 8-127) 1b: Supported 0b: Not supported	
	3	-	Reserved	
	2	-	Reserved	
	1	BitsAndErrorsCountingSupported	Indicates if DiagnosticsSelector values 02h-05h are supported (Page 14h byte 128, Table 8-127) 1b: Supported 0b: Not supported <i>Some modules are unable to perform a 64-bit division to calculate and present BER. It is expected these types of module will present the BER as error counts and total bits elapsed for the error counts presented.</i>	
	0	BitErrorRatioResultsSupported	Indicates if DiagnosticsSelector value 01h is supported (See Table 8-127) 1b: Supported 0b: Not supported	

8.12.4 Pattern Generation and Checking Location Advertisement

The support options for pattern generator and pattern checkers are advertised in 13h:131 (see Table 8-104).

Note: Following common practice, the terms 'pattern' and 'PRBS' are used interchangeably.

Figure 8-4 shows a reference diagram of possible locations for PRBS generators, PRBS checkers, FEC encoders and FEC decoders in the conceptual data path of an Application. The encircled numbers indicate the associated advertisement bit in 13h:131.

Note: The terms Pre-FEC and Post-FEC refer to the position of a generator or checker in the data path, before or after the FEC encoder for pattern generators, and before or after the FEC decoder for pattern checkers. In modules without FEC, only Post-FEC generators and Pre-FEC checkers exist (if supported).

Note: FEC usage and type of FEC depends on the Applications supported by the module (see section 6.2).

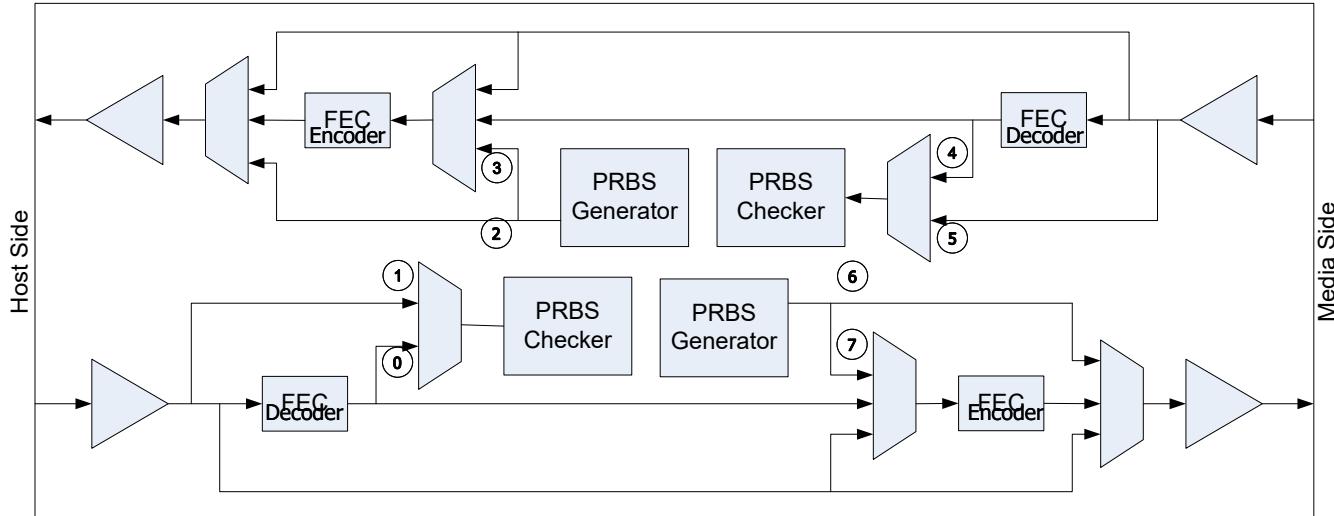


Figure 8-4 PRBS Paths Reference Diagram

Table 8-104 Pattern Generation and Checking Location (Page 13h)

Byte	Bits	Field Name	Field Description	Type
131	7	PRBSGeneratorMediaSidePreFEC	1b: Supported 0b: Not supported	RO Rqd.
	6	PRBSGeneratorMediaSidePostFEC		
	5	PRBSCheckerMediaSidePreFEC		
	4	PRBSCheckerMediaSidePostFEC		
	3	PRBSGeneratorHostSidePreFEC		
	2	PRBSGeneratorHostSidePostFEC		
	1	PRBSCheckerHostSidePreFEC		
	0	PRBSCheckerHostSidePostFEC		

8.12.5 Pattern Generation and Checking Capabilities Advertisement

This specification defines sixteen distinguished patterns that may be generated and/or checked by the module.

These patterns are identified by a Pattern ID according to Table 8-105.

All pattern names ending in 'Q' denote patterns for PAM4 modulation using Gray coding (see section 3.3).

Note: The signal rate for these patterns is determined by Application on the associated lane.

Table 8-105 Pattern IDs

Pattern ID	Pattern Name	Pattern Description
0	PRBS-31Q	As defined in IEEE 802.3-2018 clause 120.5.11.2.2
1	PRBS-31	
2	PRBS-23Q	ITU-T Recommendation O.172, 2005
3	PRBS-23	
4	PRBS-15Q	$x^{15} + x^{14} + 1$
5	PRBS-15	
6	PRBS-13Q	As defined in IEEE 802.3-2018 clause 120.5.11.2.1
7	PRBS-13	
8	PRBS-9Q	As defined in IEEE 802.3-2018 clause 120.5.11
9	PRBS-9	
10	PRBS-7Q	$x^7 + x^6 + 1$
11	PRBS-7	
12	SSPRQ	As defined in IEEE 802.3-2018 clause 120.5.11.2.3
13	-	Reserved
14	Custom	Vendor defined pattern
15	User Pattern	Programmable pattern provided in Bytes 13h:224-255

The **pattern generation capabilities** of the module are advertised in registers described in Table 8-106.

Table 8-106 PRBS Pattern Generation Capabilities (Page 13h)

Byte	Bits	Field Name	Register Description	Type
132	7	HostSideGeneratorSupportsPattern7	HostSideGeneratorSupportsPattern<i> Pattern with ID <i> (see Table 8-105) 1b: Supported 0b: Not supported (Little Endian)	RO Rqd.
	6	HostSideGeneratorSupportsPattern6		
	5	HostSideGeneratorSupportsPattern5		
	4	HostSideGeneratorSupportsPattern4		
	3	HostSideGeneratorSupportsPattern3		
	2	HostSideGeneratorSupportsPattern2		
	1	HostSideGeneratorSupportsPattern1		
	0	HostSideGeneratorSupportsPattern0		
133	7	HostSideGeneratorSupportsPattern15		RO Rqd.
	6	HostSideGeneratorSupportsPattern14		
	5	HostSideGeneratorSupportsPattern13		
	4	HostSideGeneratorSupportsPattern12		
	3	HostSideGeneratorSupportsPattern11		
	2	HostSideGeneratorSupportsPattern10		
	1	HostSideGeneratorSupportsPattern9		
	0	HostSideGeneratorSupportsPattern8		
134	7	MediaSideGeneratorSupportsPattern7	MediaSideGeneratorSupportsPattern<i> Pattern <i> (see Table 8-105 for Pattern IDs) 1b: Supported 0b: Not supported (Little Endian)	RO Rqd.
	6	MediaSideGeneratorSupportsPattern6		
	5	MediaSideGeneratorSupportsPattern5		
	4	MediaSideGeneratorSupportsPattern4		
	3	MediaSideGeneratorSupportsPattern3		
	2	MediaSideGeneratorSupportsPattern2		
	1	MediaSideGeneratorSupportsPattern1		
	0	MediaSideGeneratorSupportsPattern0		
135	7	MediaSideGeneratorSupportsPattern15		RO Rqd.
	6	MediaSideGeneratorSupportsPattern14		

Byte	Bits	Field Name	Register Description	Type
	5	MediaSideGeneratorSupportsPattern13		
	4	MediaSideGeneratorSupportsPattern12		
	3	MediaSideGeneratorSupportsPattern11		
	2	MediaSideGeneratorSupportsPattern10		
	1	MediaSideGeneratorSupportsPattern9		
	0	MediaSideGeneratorSupportsPattern8		

1

2 The **pattern checking capabilities** of the module are advertised in registers described in Table 8-107.

3

Table 8-107 Pattern Checking Capabilities (Page 13h)

Byte	Bits	Name	Description	Type
136	7	HostSideCheckerSupportsPattern7	HostSideCheckerSupportsPattern<i> Pattern <i> (see Table 8-105 for Pattern IDs) 1b: Supported 0b: Not supported (Little Endian)	RO Rqd.
	6	HostSideCheckerSupportsPattern6		
	5	HostSideCheckerSupportsPattern5		
	4	HostSideCheckerSupportsPattern4		
	3	HostSideCheckerSupportsPattern3		
	2	HostSideCheckerSupportsPattern2		
	1	HostSideCheckerSupportsPattern1		
	0	HostSideCheckerSupportsPattern0		
137	7	HostSideCheckerSupportsPattern15		RO Rqd.
	6	HostSideCheckerSupportsPattern14		
	5	HostSideCheckerSupportsPattern13		
	4	HostSideCheckerSupportsPattern12		
	3	HostSideCheckerSupportsPattern11		
	2	HostSideCheckerSupportsPattern10		
	1	HostSideCheckerSupportsPattern9		
	0	HostSideCheckerSupportsPattern8		
138	7	MediaSideCheckerSupportsPattern7	MediaSideCheckerSupportsPattern<i> Pattern <i> (see Table 8-105 for Pattern IDs) 1b: Supported 0b: Not supported (Little Endian)	RO Rqd.
	6	MediaSideCheckerSupportsPattern6		
	5	MediaSideCheckerSupportsPattern5		
	4	MediaSideCheckerSupportsPattern4		
	3	MediaSideCheckerSupportsPattern3		
	2	MediaSideCheckerSupportsPattern2		
	1	MediaSideCheckerSupportsPattern1		
	0	MediaSideCheckerSupportsPattern0		
139	7	MediaSideCheckerSupportsPattern15		RO Rqd.
	6	MediaSideCheckerSupportsPattern14		
	5	MediaSideCheckerSupportsPattern13		
	4	MediaSideCheckerSupportsPattern12		
	3	MediaSideCheckerSupportsPattern11		
	2	MediaSideCheckerSupportsPattern10		
	1	MediaSideCheckerSupportsPattern9		
	0	MediaSideCheckerSupportsPattern8		

4

5 Additional pattern capabilities are advertised in registers described in Table 8-108.

6

Table 8-108 Pattern Generator and Checker swap and invert Capabilities (Page 13h)

Byte	Bits	Field Name	Field Description	Type
140	7-6	RecoveredClockForGeneratorOptions	Options to use recovered clock for contra-directional pattern generator on the same module side 00b: not supported 01b: supported without loopback 10b: supported with loopback 11b: supported with and without loopback	RO Rqd.
	5	ReferenceClockForPatternsSupported		
			Option to use reference clock for pattern generation 1b/0b: supported/not supported	

Byte	Bits	Field Name	Field Description	Type
	4		Reserved	
	3-0	UserPatternLengthSupported	U4 Maximum length L of the user defined pattern, where the field value n encodes L as L=2(n+1), i.e. 0000b: 2 bytes, ..., 1111b: 32 bytes	
141	7	MediaSideCheckerSupportsDataSwap	0b/1b: Byte 13h:170 not supported/supported	RO Rqd.
	6	MediaSideCheckerSupportsDataInvert	0b/1b: Byte 13h:169 not supported/supported	
	5	MediaSideGeneratorSupportsDataSwap	0b/1b: Byte 13h:154 not supported/supported	
	4	MediaSideGeneratorSupportsDataInvert	0b/1b: Byte 13h:153 not supported/supported	
	3	HostSideCheckerSupportsDataSwap	0b/1b: Byte 13h:162 not supported/supported	
	2	HostSideCheckerSupportsDataInvert	0b/1b: Byte 13h:161 not supported/supported	
	1	HostSideGeneratorSupportsDataSwap	0b/1b: Byte 13h:146 not supported/supported	
	0	HostSideGeneratorSupportsDataInvert	0b/1b: Byte 13h:145 not supported/supported	
142	7	MediaCheckerSupportsPerLaneEnable	Media side pattern checker for lane i enabled in 13h:168 enables lane i (or all lanes of the Bank) 0b/1b: per lane enable not supported/supported	RO Rqd.
	6	MediaCheckerSupportsPerLanePattern	Media side pattern selection for checker 0b: Lane 1 pattern 13h:172.3-0 is used for all lanes 1b: Per lane pattern selection in 13h:172-175	
	5	MediaGeneratorSupportsPerLaneEnable	Media side pattern generator for lane i enabled in 13h:152 enables lane i (or all lanes of the Bank) 0b/1b: per lane enable not supported/supported	
	4	MediaGeneratorSupportsPerLanePattern	Media side pattern selection for generator 0b: Lane 1 pattern 13h:156.3-0 is used for all lanes 1b: Per lane pattern selection in 13h:156-159	
	3	HostCheckerSupportsPerLaneEnable	Host side pattern checker for lane i enabled in 13h:160 enables lane i (or all lanes of the Bank) 0b/1b: per lane enable not supported/supported	
	2	HostCheckerSupportsPerLanePattern	Host side pattern selection for checker 0b: Lane 1 pattern 13h:164.3-0 is used for all lanes 1b: Per lane pattern selection in 13h:164-167	
	1	HostGeneratorSupportsPerLaneEnable	Host side pattern generator for lane i enabled in 13h:144 enables lane i (or all lanes of the Bank) 0b/1b: per lane enable not supported/supported	
	0	HostGeneratorSupportsPerLanePattern	Host side pattern selection for generator 0b: Lane 1 pattern 13h:148.3-0 is used for all lanes 1b: Per lane pattern selection in 13h:148-151	
143	7-0	-	Reserved[1]	RO

8.12.6 Host Side Pattern Generator Controls

The controls in this section control pattern generation on the host side of the module in the direction of the Rx electrical output.

Table 8-109 defines the host side pattern generator controls and Table 8-110 defines the host side pattern generator selection controls.

Table 8-109 Host Side Pattern Generator Controls (Page 13h)

Byte	Bits	Field Name	Register Description	Type
144	7	HostSideGeneratorEnableLane8	1b: Enable generator with configuration defined in Bytes 145-151 0b: Disable pattern generator Advertisement: 13h:131.3-2	RW Adv.
	6	HostSideGeneratorEnableLane7		
	5	HostSideGeneratorEnableLane6		
	4	HostSideGeneratorEnableLane5		
	3	HostSideGeneratorEnableLane4		
	2	HostSideGeneratorEnableLane3		
	1	HostSideGeneratorEnableLane2		
	0	HostSideGeneratorEnableLane1		
145	7	HostSideGeneratorDataInvertLane8	1b: Invert the selected pattern 0b: Do not invert the selected pattern <i>Note: This control inverts the pattern; it does not swap the P and N signals (see polarity controls in Table 8-69.)</i>	RW Adv.
	6	HostSideGeneratorDataInvertLane7		
	5	HostSideGeneratorDataInvertLane6		
	4	HostSideGeneratorDataInvertLane5		
	3	HostSideGeneratorDataInvertLane4		
	2	HostSideGeneratorDataInvertLane3		
	1	HostSideGeneratorDataInvertLane2		
	0	HostSideGeneratorDataInvertLane1		
146	7	HostSideGeneratorSwapSymbolBitsLane8	1b: Swap MSB and LSB for PAM4 patterns 0b: Do not swap MSB and LSB	RW Adv.
	6	HostSideGeneratorSwapSymbolBitsLane7		
	5	HostSideGeneratorSwapSymbolBitsLane6		
	4	HostSideGeneratorSwapSymbolBitsLane5		
	3	HostSideGeneratorSwapSymbolBitsLane4		
	2	HostSideGeneratorSwapSymbolBitsLane3		
	1	HostSideGeneratorSwapSymbolBitsLane2		
	0	HostSideGeneratorSwapSymbolBitsLane1		
147	7	HostSideGeneratorPreFECEnableLane8	1b: Encoded pattern: Generate the selected pattern at the input to the internal host side FEC encoder 0b: Unencoded pattern: Generate the selected pattern at a location after the internal host side FEC encoder block	RW Adv.
	6	HostSideGeneratorPreFECEnableLane7		
	5	HostSideGeneratorPreFECEnableLane6		
	4	HostSideGeneratorPreFECEnableLane5		
	3	HostSideGeneratorPreFECEnableLane4		
	2	HostSideGeneratorPreFECEnableLane3		
	1	HostSideGeneratorPreFECEnableLane2		
	0	HostSideGeneratorPreFECEnableLane1		

Table 8-110 Host Side Pattern Generator Pattern Select Controls (Page 13h)

Byte	Bits	Field Name	Field Description	Type
148	7-4	HostSideGeneratorPatternSelectLane2	Selected pattern to be generated on each lane. See Table 8-105 for pattern coding.	RW Adv.
	3-0	HostSideGeneratorPatternSelectLane1		
149	7-4	HostSideGeneratorPatternSelectLane4		
	3-0	HostSideGeneratorPatternSelectLane3		
150	7-4	HostSideGeneratorPatternSelectLane6		
	3-0	HostSideGeneratorPatternSelectLane5		
151	7-4	HostSideGeneratorPatternSelectLane8		
	3-0	HostSideGeneratorPatternSelectLane7		

8.12.7 Media Side Pattern Generator Controls

The controls in this section control pattern generation on the media side of the module in the direction of the Tx electrical or optical output.

Table 8-111 defines the media side pattern generator controls and Table 8-112 defines the media side pattern generator selection controls.

Table 8-111 Media Side Pattern Generator Controls (Page 13h)

Byte	Bits	Field Name	Register Description	Type
152	7	MediaSideGeneratorEnableLane8	1b: Enable pattern generator with configuration defined in Bytes 153-159 0b: Disable pattern generator Advertisement: 13h:131.7-6	RW Adv.
	6	MediaSideGeneratorEnableLane7		
	5	MediaSideGeneratorEnableLane6		
	4	MediaSideGeneratorEnableLane5		
	3	MediaSideGeneratorEnableLane4		
	2	MediaSideGeneratorEnableLane3		
	1	MediaSideGeneratorEnableLane2		
	0	MediaSideGeneratorEnableLane1		
153	7	MediaSideGeneratorDataInvertLane8	1b: Invert the selected pattern 0b: Do not invert the selected pattern <i>Note: This control inverts the pattern; it does not swap the P and N signals (see polarity controls in Table 8-69)</i>	RW Adv.
	6	MediaSideGeneratorDataInvertLane7		
	5	MediaSideGeneratorDataInvertLane6		
	4	MediaSideGeneratorDataInvertLane5		
	3	MediaSideGeneratorDataInvertLane4		
	2	MediaSideGeneratorDataInvertLane3		
	1	MediaSideGeneratorDataInvertLane2		
	0	MediaSideGeneratorDataInvertLane1		
154	7	MediaSideGeneratorSwapSymbolBitsLane8	1b: Swap MSB and LSB for PAM4 patterns 0b: Do not swap MSB and LSB	RW Adv.
	6	MediaSideGeneratorSwapSymbolBitsLane7		
	5	MediaSideGeneratorSwapSymbolBitsLane6		
	4	MediaSideGeneratorSwapSymbolBitsLane5		
	3	MediaSideGeneratorSwapSymbolBitsLane4		
	2	MediaSideGeneratorSwapSymbolBitsLane3		
	1	MediaSideGeneratorSwapSymbolBitsLane2		
	0	MediaSideGeneratorSwapSymbolBitsLane1		
155	7	MediaSideGeneratorPreFECEnableLane8	1b: Encoded pattern: Generate the selected pattern at the input to the internal media side FEC encoder 0b: Unencoded pattern: Generate the selected pattern at a location after the internal media side FEC encoder block	RW Adv.
	6	MediaSideGeneratorPreFECEnableLane7		
	5	MediaSideGeneratorPreFECEnableLane6		
	4	MediaSideGeneratorPreFECEnableLane5		
	3	MediaSideGeneratorPreFECEnableLane4		
	2	MediaSideGeneratorPreFECEnableLane3		
	1	MediaSideGeneratorPreFECEnableLane2		
	0	MediaSideGeneratorPreFECEnableLane1		

Table 8-112 Media Side Pattern Generator Pattern Select Controls (Page 13h)

Byte	Bits	Field Name	Register Description	Type
156	7-4	MediaSideGeneratorPatternSelectLane2	Selected pattern to be generated on each lane. See Table 8-105 for pattern coding.	RW Adv.
	3-0	MediaSideGeneratorPatternSelectLane1		
157	7-4	MediaSideGeneratorPatternSelectLane4		
	3-0	MediaSideGeneratorPatternSelectLane3		
158	7-4	MediaSideGeneratorPatternSelectLane6		
	3-0	MediaSideGeneratorPatternSelectLane5		
159	7-4	MediaSideGeneratorPatternSelectLane8		
	3-0	MediaSideGeneratorPatternSelectLane7		

8.12.8 Host Side Pattern Checker Controls

The controls in this section control pattern checking on the host side of the module for data arriving in the direction of the Tx electrical input.

Table 8-113 defines the host side pattern checker controls and Table 8-114 defines the host side pattern checker selection controls.

Table 8-113 Host Side Pattern Checker Controls (Page 13h)

Byte	Bits	Field Name	Register Description	Type
160	7	HostSideCheckerEnableLane8	1b: Enable pattern checker with configuration defined in Bytes 161-167 0b: Disable pattern checker Advertisement: 13h:131.1-0	RW Adv.
	6	HostSideCheckerEnableLane7		
	5	HostSideCheckerEnableLane6		
	4	HostSideCheckerEnableLane5		
	3	HostSideCheckerEnableLane4		
	2	HostSideCheckerEnableLane3		
	1	HostSideCheckerEnableLane2		
	0	HostSideCheckerEnableLane1		
161	7	HostSideCheckerDataInvertLane8	1b: Invert the selected pattern 0b: Do not invert the selected pattern <i>Note: This control inverts the pattern; it does not swap the P and N signals (P/N swap for input signals is not currently supported by this specification).</i>	RW Adv.
	6	HostSideCheckerDataInvertLane7		
	5	HostSideCheckerDataInvertLane6		
	4	HostSideCheckerDataInvertLane5		
	3	HostSideCheckerDataInvertLane4		
	2	HostSideCheckerDataInvertLane3		
	1	HostSideCheckerDataInvertLane2		
	0	HostSideCheckerDataInvertLane1		
162	7	HostSideCheckerSwapSymbolBitsLane8	1b: Swap MSB and LSB for PAM4 patterns 0b: Do not swap MSB and LSB	RW Adv.
	6	HostSideCheckerSwapSymbolBitsLane7		
	5	HostSideCheckerSwapSymbolBitsLane6		
	4	HostSideCheckerSwapSymbolBitsLane5		
	3	HostSideCheckerSwapSymbolBitsLane4		
	2	HostSideCheckerSwapSymbolBitsLane3		
	1	HostSideCheckerSwapSymbolBitsLane2		
	0	HostSideCheckerSwapSymbolBitsLane1		
163	7	HostSideCheckerPostFECEnableLane8	1b: Check the selected encoded pattern at the output from the internal FEC decoder 0b: Check the selected unencoded pattern at the input to the internal FEC decoder block	RW Adv.
	6	HostSideCheckerPostFECEnableLane7		
	5	HostSideCheckerPostFECEnableLane6		
	4	HostSideCheckerPostFECEnableLane5		
	3	HostSideCheckerPostFECEnableLane4		
	2	HostSideCheckerPostFECEnableLane3		
	1	HostSideCheckerPostFECEnableLane2		
	0	HostSideCheckerPostFECEnableLane1		

Table 8-114 Host Side Pattern Checker Pattern Select Controls (Page 13h)

Byte	Bits	Name	Register Description	Type
164	7-4	HostSideCheckerPatternSelectLane2	Selected pattern to be generated on each lane. See Table 8-105 for pattern coding.	RW Adv.
	3-0	HostSideCheckerPatternSelectLane1		
165	7-4	HostSideCheckerPatternSelectLane4		
	3-0	HostSideCheckerPatternSelectLane3		
166	7-4	HostSideCheckerPatternSelectLane6		
	3-0	HostSideCheckerPatternSelectLane5		
167	7-4	HostSideCheckerPatternSelectLane8		
	3-0	HostSideCheckerPatternSelectLane7		

8.12.9 Media Side Pattern Checker Controls

The controls in this section control pattern checking on the media side of the module for data arriving in the direction of the Rx electrical or optical input.

Table 8-115 defines the media side pattern checker controls and Table 8-116 defines the media side pattern checker selection controls.

Table 8-115 Media Side Pattern Checker Controls (Page 13h)

Byte	Bits	Field Name	Register Description	Type
168	7	MediaSideCheckerEnableLane8	1b: Enable pattern checker with configuration defined in Bytes 169-175 0b: Disable pattern checker Advertisement: 13h:131.5-4	RW Adv.
	6	MediaSideCheckerEnableLane7		
	5	MediaSideCheckerEnableLane6		
	4	MediaSideCheckerEnableLane5		
	3	MediaSideCheckerEnableLane4		
	2	MediaSideCheckerEnableLane3		
	1	MediaSideCheckerEnableLane2		
	0	MediaSideCheckerEnableLane1		
169	7	MediaSideCheckerDataInvertLane8	1b: Invert the selected pattern 0b: Do not invert the selected pattern	RW Adv.
	6	MediaSideCheckerDataInvertLane7		
	5	MediaSideCheckerDataInvertLane6		
	4	MediaSideCheckerDataInvertLane5		
	3	MediaSideCheckerDataInvertLane4		
	2	MediaSideCheckerDataInvertLane3		
	1	MediaSideCheckerDataInvertLane2		
	0	MediaSideCheckerDataInvertLane1		
170	7	MediaSideCheckerSwapSymbolBitsLane8	1b: Swap MSB and LSB for PAM4 patterns 0b: Do not swap MSB and LSB	RW Adv.
	6	MediaSideCheckerSwapSymbolBitsLane7		
	5	MediaSideCheckerSwapSymbolBitsLane6		
	4	MediaSideCheckerSwapSymbolBitsLane5		
	3	MediaSideCheckerSwapSymbolBitsLane4		
	2	MediaSideCheckerSwapSymbolBitsLane3		
	1	MediaSideCheckerSwapSymbolBitsLane2		
	0	MediaSideCheckerSwapSymbolBitsLane1		
171	7	MediaSideCheckerPostFECEnableLane8	1b: Check the selected encoded pattern at the output from the internal FEC decoder 0b: Check the selected unencoded pattern at the input to the internal FEC decoder block	RW Adv.
	6	MediaSideCheckerPostFECEnableLane7		
	5	MediaSideCheckerPostFECEnableLane6		
	4	MediaSideCheckerPostFECEnableLane5		
	3	MediaSideCheckerPostFECEnableLane4		
	2	MediaSideCheckerPostFECEnableLane3		
	1	MediaSideCheckerPostFECEnableLane2		
	0	MediaSideCheckerPostFECEnableLane1		

Table 8-116 Media Side Pattern Checker Select Controls (Page 13h)

Byte	Bits	Field Name	Register Description	Type
172	7-4	MediaSideCheckerPatternSelectLane2	Selected pattern to be generated on each lane. See Table 8-105 for pattern coding.	RW Adv.
	3-0	MediaSideCheckerPatternSelectLane1		
173	7-4	MediaSideCheckerPatternSelectLane4		
	3-0	MediaSideCheckerPatternSelectLane3		
174	7-4	MediaSideCheckerPatternSelectLane6		
	3-0	MediaSideCheckerPatternSelectLane5		
175	7-4	MediaSideCheckerPatternSelectLane8		
	3-0	MediaSideCheckerPatternSelectLane7		

8.12.10 Clocking and Measurement Controls

Table 8-117 describes general controls for the pattern generator and checker features.

Examples showing the usage of the Pattern Generator/Checker controls are described in Appendix E.

Table 8-117 Clocking and Measurement Controls (Page 13h)

Byte	Bits	Field Name	Field Description	Type
176	7-4	HostPRBSGeneratorClockSource	Clock source for Host Side PRBS Pattern Generation: 0: All lanes use Internal Clock i=1-8: All lanes use Reference Clock Media Lane i 9-14: Reserved 15: Recovered clock per Media Lane or Data Path	RW Opt.
	3-0	MediaPRBSGeneratorClockSource	Clock source for Media Side PRBS Pattern Generation: 0: All lanes use Internal Clock 1: All lanes use Reference Clock i= 2-9: All lanes use Reference Clock Host Lane i-1 10-14: Reserved 15: Recovered clock per Host Lane or Data Path	
177	7	StartStopIsGlobal	Bool StartStopIsGlobal controls whether writing a measurement Start/Stop control (see below for a list) on one Bank globally starts/stops measurements across all Banks as if the same control value change had occurred in all supported Banks synchronously (which is impossible): 0b: A start/stop control change acts on current Bank 1b: A start/stop control change acts on all Banks Start/Stop controls (instances) Locations ResetErrorInformation (4) (0-3:13h:177.5) HostSideCheckerEnable (32) (0-3:13h:160.0-7) MediaSideCheckerEnable (32) (0-3:13h:168.0-7) <i>Note: The intended effect of StartStopIsGlobal is that starting or stopping error counting measurement in one supported Bank occurs on all affected lanes in all supported Banks synchronously.</i> <i>Note: The actually affected lanes depend on per Bank capabilities like *CheckerSupportsPerLaneEnable.</i>	RW Opt.
6	-		Reserved	
5		ResetErrorInformation	This bit has effects on error information, and on gating timers when the value changes. Effect on Error Information 0b→1b: Freeze. Currently accumulating error statistics identified by DiagnosticsSelector 01h to 05h are frozen and, if supported, gated results identified by Selectors 11h-15h are updated with the frozen current error statistics. 1b→0b: Reset. Error statistics identified by Selectors 01h-05h are reset to 0, whereas results identified by Selectors 11h-15h are unaffected. <i>Whenever a ResetErrorInformation operation on an individual lane or all lanes is triggered, it should always clear associated Error Information (both PRBS BER and Error Counters) in the Diagnostics Data memory at the same time.</i> The effect of ResetErrorInformation depends also on - PerLaneGatingTimersSupported (13h:129.3) - StartStopIsGlobal (13h:177.7) Effect on global gating timers (13h:129.3=0b)	

Byte	Bits	Field Name	Field Description	Type
			<p>0b→1b: single gate timer is stopped 1b→0b: starts the single gating timer Effect on per lane gating timers (13h:129.3=1b) 1b: keeps the gate timers stopped for the lanes in this Bank (13h:177.7=0) or on all. 0b: starts the individual gating timers for all enabled lanes in this Bank.</p>	
	4	AutoRestartGating	<p>0b: When Gate Time expires, the module will set the Gate Complete Flag in Bytes 14h:134-135. The module will update the error counter result and stop error detection. 1b: When Gate Time expires, the module will set the Gate Complete Flag in Bytes 14h:134-135. The module will update the error counter results, then immediately clear internal error counters and restart the Gate Timer, while continuing to count errors. <i>Note: With Auto-restart the host must read the error counter results before the Gate Time expires to avoid the result values to be overwritten with new results.</i></p>	
	3-1	MeasurementTime	<p>Measurement (gating) time for one complete result over a defined measurement period. 000b: ungated, counters accrue indefinitely (infinite gate time) 001b: 5 sec gate time 010b: 10 sec gate time 011b: 30 sec gate time 100b: 60 sec gate time 101b: 120 sec gate time 110b: 300 sec gate time 111b: Custom</p>	
	0	UpdatePeriodSelect	<p>Time between incremental updates to intermediate error counting results during a longer gating period 0b: 1 sec update interval 1b: 5 sec update interval This period is relevant only if PeriodicUpdatesSupported (13h:129.4=1b) advertises that the module will update the BER or Error counters while gating is in progress. Then two update rates at which the module will process the internal error counters and update the error counter fields in the Memory Map are possible. The host can then poll "real-time" error counts, up until the gating results are completed.</p>	
178	7-4	-	Reserved	RW
	3-2	HostPRBSCheckerClockSource	<p>The clock source used for the Host PRBS Pattern Checker can be configured using this control register: 0h: Recovered clocks from Host Lane/Data paths 1h: All lanes use Internal Clock 2h: All lanes use Reference Clock 3h: Reserved</p>	RW
	1-0	MediaPRBSCheckerClockSource	<p>The clock source used for the Media PRBS Pattern Checker can be configured using this control register: 0h: Recovered clocks from Media Lane/Data paths 1h: All lanes use Internal Clock 2h: All lanes use Reference Clock 3h: Reserved.</p>	RW
179	7-0	-	Reserved[1]	RW

8.12.11 Diagnostics Measurement Behavior

Editor's Note: Section should be revised for better clarity in the next revision of this specification

8.12.11.1 Un-Gated Measurements

This section describes error performance measurements when gating is not used (13h:177.3-1= 000b).

When intermediate periodic update of error information is supported (PeriodicUpdatesSupported=1), the retrievable measurement results are cumulative and progressively updated until the host eventually stops the measurement. These intermediate updates are not notified by a Flag, so the host needs to poll the updates.

When intermediate periodic update of error information is not supported, measurements without gating are single-shot and results are available only when the measurement has been stopped by the host.

Table 8-118 PRBS Checker Behavior Un-Gated Mode

Configuration	13h:177 (D7-D0)	Description
13h:177.3-1=000b Not Gated 13h:177.4 ignored AutoRestartGating = x 13h:177.7 StartStopIsGlobal = 0	000x 000y	<p>In this mode, the error metrics measurement runs continuously. When the host enables disabled PRBS checkers (in 13h:160 or in 13h:168) all error counters for the enabled lanes are cleared and then start accumulating.</p> <p>When the host disables enabled PRBS checkers (in 13h:160 or in 13h:168) error counting is stopped, and error counting results will be available both via Selector 01-05h and 11h-15h (if supported). If 13h:129.4=0, real time error information is not updated and error information is only available when the error counting is stopped by checker disable.</p> <p>If 13h:129.4 = 1, real time error information is available with Selectors 01h-05h, updated period configured by y. Accurate error counter accumulation times can be derived from the total bit counters in the error information.</p> <p>A write to 13h:177. 5=1 also cause the error information registers to reset to 0 and restart accumulation on the enabled lanes.</p>
13h:177.3-1=000b Not Gated 13h:177.4 ignored AutoRestartGating = x 177.7 StartStopIsGlobal = 1	100x 000y	<p>The behavior is the same as above with the exception that toggling 13h:177. 5 causes the error information registers of all lanes in all Banks to reset and restart accumulation.</p> <p>As described above just prior to reset of error information of all lanes in all Banks, the previous error information should be copied to both 01-05h (and 11h-15h if supported). Selector 01-05h restarts accumulation and 11h-15h (if supported) contains the error counting results of the period prior to reset.</p>
ResetErrorInformation 13h:177.5=1b (freeze) 13h:177.5=0b (restart) PerLaneGatingTimersSupported 13h:129.3 =0 13h:177.7 = x StartStopIsGlobal irrelevant for per lane gating	x01x 000y x00x 000y	<p>The ResetErrorInformation (13h:177.5) bit is toggled to reset error information and start a new measurement for all enabled lanes.</p> <p>When ResetErrorInformation (13h:177.5) is raised, error counters are frozen for all enabled lane checkers and can be read via Selectors 01h-05h or 11h-15h (if supported).</p> <p>When ResetErrorInformation (13h:177.5) is ceased, error counters of all enabled checkers of both host and media lanes are reset and started.</p> <p>The host may also individually toggle host (13h:160) or media (13h:168) lane checker enable bits to restart error counting of specific host or media lanes individually.</p>

8.12.11.2 Gated Measurements with Global Gate Timer

This section describes error performance measurements with gating based on a single global Gate Timer for all lanes and all Banks (13h:129.3 = 0) on Host Side, or on Media Side, respectively.

Bit 13h:129.4 defines if real-time error information is available while the gating timer has not expired.

When gating is enabled, the Flags 14h:134-135 are raised at the expiry of the gate timers to indicate that new error information results are available.

If the host fails to read the error information registers prior to expiry of subsequent gate time expiry, only the latest error information will be available.

Table 8-119 PRBS Checker Behavior Single Gate Timer

Configuration	13h:177 (D7-D0)	Description
Gated, <i>nnn</i> configured secs. Since 13h:129.3=0 the control 13h:177.7 is ignored.	X000 <i>nnny</i> <i>nnn</i> ≠ 000	<p>PRBS Error Counters are gated. Whenever PRBS checkers are enabled in 13h:160 or 13h:168, the respective Host or Media single gate timer resets to 0, all error counters for the enable lanes reset to 0 and start accumulating errors.</p> <p>When the gate timer expires, the error counters stop counting. At this time the module shall be able to present the error information collected within this gating period on Page 14h by using the Selectors 01h-05h (and 11h-15h if supported).</p> <p>If 13h:129.4 = 0, real time error information is not updated while gating is in progress. The error information will only be available at the end of the gate.</p> <p>If 13h:129.4 = 1, real time error information is available by using Page 14h using the Selectors 01h-05h. This real time error information will be updated by the module every configure "y" seconds. This mode is useful if the gating period is long. A host may choose to periodically read the real-time error information during gating and take any necessary action if a bad BER is detected.</p> <p>To restart gating, the host toggles the PRBS enable registers 160 and 168.</p>
Gated, <i>nnn</i> configured secs. Since 13h:129.3=0 the control 13h:177.7 is ignored. If 13h:129.2 is set, and the host sets 13h:177.4 AutoRestartGating , the gating timer will automatically restart.	Xx01 <i>nnny</i> <i>nnn</i> ≠ 000	<p>This behavior is the same as the above row, except for the auto restart behavior on gate timer expiry or exceeding the configured gate time. Module should support Selectors 11h-15h to use this feature, otherwise error information from the previous gating period will be lost.</p> <p>When the gate timer expires and exceeds the configured "nnn" elapsed time, the error information is presented in the error information results via Selectors 11h-15h. The current error information will reset, and the gate timer will be reset to 0 and restart accumulating errors for a new gating period. The host must read the error information from the previous gated time using the error information results Selectors 11h-15h.</p>
Restart gate timer (Bit 5) aka ResetErrorInformation Since 13h:129.3=0 there is only one timer for host and one timer for media lanes and 13h:177.7 is irrelevant.	X01x <i>nnny</i> <i>nnn</i> ≠ 000	<p>This bit is used to restart both the host timer and the media timer:</p> <p>When Byte177.5 is set, all enabled lanes error counters are frozen and the gate timers are stopped.</p> <p>When Byte177.5 is cleared, both the host and media enabled the gate timers are restarted from 0, the error information is reset, and a new error accumulation gate period is restarted.</p> <p>The host may also individually toggle enable bits for host lanes (13h:160) or media lanes (13h:168) to restart the relevant gate timer individually.</p>

8.12.11.3 Gated Measurements with Per Lane Gate Timer

This section describes error performance measurements with gating based on individual Gate Timers for all lanes and on all Banks.

Byte 13h:129.4 defines if real-time error information is available while the gating timer has not expired.

When gating is enabled, the Flags 14h:134-135 are raised at gate timer expiry, and Interrupt may be asserted to indicate that new error information is available. If the host then fails to read error information prior to the next gate timer expiry, the previous error information is lost and only the latest error information is available.

Table 8-120 PRBS Checker Behavior Per Lane Gate Timer

Configuration	Byte 177 (D7-D0)	Description
Gated, <i>nnn</i> configured secs. Since 13h:129.3=0 the control 13h:177.7 is ignored.	X000 <i>nnny</i> <i>nnn</i> !=000	<p>PRBS Error Counters are gated. When a PRBS checker is enabled in 13h:160 (or in 13h:168), the relevant Host (or Media) lane gate timer resets to 0, error counters for the enabled lanes are reset and then start accumulating errors. Since the gate timers are individual per lane, these timers reset independently to provide the most accurate per lane gated time as possible.</p> <p>When the configured gate time expires, the error counters stop counting and the host can inspect or collect the results of the finished gating period in the Diagnostics Area on Page 14h, using the Selectors 01h-05h and 11h-15h (if supported).</p> <p>If 13h:129.4 = 0, real time error information is not updated while gating is in progress, but only at the end of the gate.</p> <p>If 13h:129.4 = 1, real time error information is available on Page 14h using the Selectors 01h-05h, whereby real time error information is updated with period configured by y.</p> <p>This mode is useful if the gating period is long. A host may then periodically read the real-time error information during gating and, e.g., react if bad performance is detected.</p> <p>To restart gating, the host toggles the PRBS checker enable registers 13h:160 and 13h:168 or sets and ceases 13h:177.5.</p>
Gated, <i>nnn</i> configured secs. Since 13h:129.3=0 the control 13h:177.7 is ignored. If the Page 13h:129.2 is set, and the host writes a 1 to Bit4 of Byte 177, the gating timer will automatically restart.	Xx01 <i>nnny</i> <i>nnn</i> !=000	<p>This behavior is the same as the above row, except for the behavior when the gate timer expires and exceeds the configured gate time. Module shall support Selectors 11h-15h to use this feature.</p> <p>Here at the end of the gate, that is when the individual per lane gate timer expires and exceeds the configured “<i>nnn</i>” elapsed time, the PRBS error information will be presented in the error information results via Selectors 11h-15h. The current error information will reset, and the gate timer will be reset to 0 and restart accumulating errors for a new gating period. The host will have to read the error information from the previous gated time using the error information results Selectors 11h-15h.</p>
Restart gate timer (Bit 5) aka Reset Error Information	x01x <i>nnny</i> <i>nnn</i> !=000	<p>This bit is used to restart all enabled Bank and lanes gate timers. When 13h:177.5 is raised, all enabled lane error counters are frozen, and the gate timers are stopped.</p> <p>If 13h:177.7=0 when 13h:177.5 is ceased, all enabled gate timers of the current Bank are restarted from 0, error information is reset, and a new error accumulation gate period is started.</p> <p>If 13h:177.7=1 when 13h:177.5 is ceased, all enabled gate timers of all Banks with 13h:177.7=1 are restarted from 0, error information is reset, and a new accumulation gate period is started.</p> <p>The host may also individually toggle host Byte 160 or media Byte 168 enable bits to restart the gate timer for the host and media lanes independently. In this case, only the error information of the enable or disabled lane will reset and start count or freeze in its last value respectively.</p>

8.12.12 Loopback Controls

Host and Media side loopback control registers and module behaviors depend on the advertised loopback capabilities described in Table 8-101.

Controls for loopback features are described in Table 8-121.

The module may reject unsupported host-written loopback settings (no change in affected register bits).

For example, if the module advertises that it can only perform host side or media side loopback one at a time and not simultaneously, the module may reject the command such that the rejected values in the loopback controls do not change.

Table 8-121 Loopback Controls (Page 13h)

Byte	Bits	Field Name	Register Description	Type
180	7	MediaSideOutputLoopbackEnableLane8	MediaSideOutputLoopbackEnableLane< i > 0b: normal non-loopback operation 1b: loopback enabled.	RWW Adv.
	6	MediaSideOutputLoopbackEnableLane7		
	5	MediaSideOutputLoopbackEnableLane6		
	4	MediaSideOutputLoopbackEnableLane5	If the Per-lane Media Side Loopback Supported field=1, loopback control is per lane. Otherwise, if any loopback enable bit is set to 1, all Media side lanes are in output loopback. Advertisement: 13h:128.	
	3	MediaSideOutputLoopbackEnableLane4		
	2	MediaSideOutputLoopbackEnableLane3		
	1	MediaSideOutputLoopbackEnableLane2		
	0	MediaSideOutputLoopbackEnableLane1		
181	7	MediaSideInputLoopbackEnableLane8	MediaSideInputLoopbackEnableLane< i > 0b: normal non-loopback operation 1b: loopback enabled.	RWW Adv.
	6	MediaSideInputLoopbackEnableLane7		
	5	MediaSideInputLoopbackEnableLane6		
	4	MediaSideInputLoopbackEnableLane5		
	3	MediaSideInputLoopbackEnableLane4		
	2	MediaSideInputLoopbackEnableLane3		
	1	MediaSideInputLoopbackEnableLane2		
	0	MediaSideInputLoopbackEnableLane1		
182	7	HostSideOutputLoopbackEnableLane8	HostSideOutputLoopbackEnableLane< i > 0b: normal non-loopback operation 1b: loopback enabled.	RWW Adv.
	6	HostSideOutputLoopbackEnableLane7		
	5	HostSideOutputLoopbackEnableLane6		
	4	HostSideOutputLoopbackEnableLane5		
	3	HostSideOutputLoopbackEnableLane4		
	2	HostSideOutputLoopbackEnableLane3		
	1	HostSideOutputLoopbackEnableLane2		
	0	HostSideOutputLoopbackEnableLane1		
183	7	HostSideInputLoopbackEnableLane8	HostSideInputLoopbackEnableLane< i > 0b: normal non-loopback operation 1b: loopback enabled.	RWW Adv.
	6	HostSideInputLoopbackEnableLane7		
	5	HostSideInputLoopbackEnableLane6		
	4	HostSideInputLoopbackEnableLane5		
	3	HostSideInputLoopbackEnableLane4		
	2	HostSideInputLoopbackEnableLane3		
	1	HostSideInputLoopbackEnableLane2		
	0	HostSideInputLoopbackEnableLane1		

8.12.13 Host Scratchpad Area

Support of the HostScratchPad registers is advertised in 01h:251.7.

The module clears the Scratchpad area on each Firmware restart, including auto-recovery reboots.

Otherwise, it ignores the Scratchpad area.

Table 8-122 Host Scratchpad Area (Page 13h)

Byte	Bits	Field Name	Register Description	Type
184	7-0	HostScratchPad0		
185	7-0	HostScratchPad1		
186	7-0	HostScratchPad2		
187	7-0	HostScratchPad3		
188	7-0	HostScratchPad4		
189	7-0	HostScratchPad5		
190	7-0	HostScratchPad6		
191	7-0	HostScratchPad7	HostScratchPad <i>, <i> = 0, ..., 7 These 8 bytes can be used by the host to store state in the module. The scratch pad clears in every firmware restart to its default zero values. It can therefore also be used by the host to detect a module-initiated recovery reboot. Advertisement: 01h:251.7	RW Adv.

6

8.12.14 Diagnostics Masks

Table 8-123 provides Mask bits for all diagnostics Flags.

The default value for all Mask bits on this page is 1 (masked).

Diagnostics Flags are located on Page 14h (see Table 8-128).

Table 8-123 Diagnostics Masks (Page 13h)

Byte	Bits	Field Name	Field/Register Description	Type
206	7	LossOfReferenceClockMask	Loss of reference clock Mask for the module	RW Opt.
	6-0	-	Reserved	
207	7-0	-	Reserved[1]	RW
			PatternCheckGatingCompleteMaskHostLane< i >	
208	7	PatternCheckGatingCompleteMaskHostLane8	Per-host lane gating complete Mask. Default:1	RW Adv.
	6	PatternCheckGatingCompleteMaskHostLane7		
	5	PatternCheckGatingCompleteMaskHostLane6		
	4	PatternCheckGatingCompleteMaskHostLane5		
	3	PatternCheckGatingCompleteMaskHostLane4		
	2	PatternCheckGatingCompleteMaskHostLane3		
	1	PatternCheckGatingCompleteMaskHostLane2		
	0	PatternCheckGatingCompleteMaskHostLane1		
			PatternCheckGatingCompleteMaskMediaLane< i >	
209	7	PatternCheckGatingCompleteMaskMediaLane8	Per-media lane gating complete Mask. Default:1	RW Adv.
	6	PatternCheckGatingCompleteMaskMediaLane7		
	5	PatternCheckGatingCompleteMaskMediaLane6		
	4	PatternCheckGatingCompleteMaskMediaLane5		
	3	PatternCheckGatingCompleteMaskMediaLane4		
	2	PatternCheckGatingCompleteMaskMediaLane3		
	1	PatternCheckGatingCompleteMaskMediaLane2		
	0	PatternCheckGatingCompleteMaskMediaLane1		
			PatternGeneratorLOLMaskHostLane< i >	
210	7	PatternGeneratorLOLMaskHostLane8	Per-host lane pattern generator loss of lock Mask Default:1	RW Adv.
	6	PatternGeneratorLOLMaskHostLane7		
	5	PatternGeneratorLOLMaskHostLane6		
	4	PatternGeneratorLOLMaskHostLane5		
	3	PatternGeneratorLOLMaskHostLane4		
	2	PatternGeneratorLOLMaskHostLane3		
	1	PatternGeneratorLOLMaskHostLane2		
	0	PatternGeneratorLOLMaskHostLane1		
			PatternGeneratorLOLMaskMediaLane< i >	
211	7	PatternGeneratorLOLMaskMediaLane8	Per-media lane pattern generator loss of lock Mask Default:1	RW Adv.
	6	PatternGeneratorLOLMaskMediaLane7		
	5	PatternGeneratorLOLMaskMediaLane6		
	4	PatternGeneratorLOLMaskMediaLane5		
	3	PatternGeneratorLOLMaskMediaLane4		
	2	PatternGeneratorLOLMaskMediaLane3		
	1	PatternGeneratorLOLMaskMediaLane2		
	0	PatternGeneratorLOLMaskMediaLane1		
			PatternCheckerLOLMaskHostLane< i >	
212	7	PatternCheckerLOLMaskHostLane8	Per-host lane pattern checker loss of lock Mask Default:1	RW Adv.
	6	PatternCheckerLOLMaskHostLane7		
	5	PatternCheckerLOLMaskHostLane6		
	4	PatternCheckerLOLMaskHostLane5		
	3	PatternCheckerLOLMaskHostLane4		
	2	PatternCheckerLOLMaskHostLane3		
	1	PatternCheckerLOLMaskHostLane2		
	0	PatternCheckerLOLMaskHostLane1		
			PatternCheckerLOLMaskMediaLane< i >	
213	7	PatternCheckerLOLMaskMediaLane8	Per-media lane pattern checker loss of lock Mask Default:1	RW Adv.
	6	PatternCheckerLOLMaskMediaLane7		
	5	PatternCheckerLOLMaskMediaLane6		
	4	PatternCheckerLOLMaskMediaLane5		
	3	PatternCheckerLOLMaskMediaLane4		
	2	PatternCheckerLOLMaskMediaLane3		
	1	PatternCheckerLOLMaskMediaLane2		

Byte	Bits	Field Name	Field/Register Description	Type
	0	PatternCheckerLOLMaskMediaLane1		
214-223	7-0	-	Reserved[10]	RW

1

8.12.15 User Pattern

2

Table 8-124 provides space for the host to define a user pattern of up to 32 bytes in length. The module may not support the full 32-byte length. Refer to Table 8-108 for the maximum supported user pattern length.

3

Table 8-124 User Pattern (Page 13h)

Byte	Bits	Register Name	Register Description	Type
224-255	7-0	UserPattern	Host defined user pattern (32 Bytes)	RW Opt.

8.13 Banked Page 14h (Module Performance Diagnostics Results)

Pages 13h and 14h are optional Pages and contain module diagnostic control and status fields.

The module advertises support of Pages 13h and 14h jointly in Bit 01h:142.5 (see Table 8-46).

Page 14h may optionally be Banked. Each Bank of Page 14h refers to 8 lanes.

Page 14h is subdivided into subject areas as illustrated in the following table:

Table 8-125 Page 14h Overview

Byte	Size (bytes)	Subject Area	Description
128	1	DiagnosticsSelector	Selects the content of Diagnostics Data (Bytes 192-255)
129	1	-	Reserved[1]
130-131	2	-	Custom[2]
132-139	18	Diagnostics Flags	Latched diagnostics Flags
140-149	10	-	Reserved[10]
192-255	64	Diagnostics Data	Diagnostics Data selected by Diagnostics Selector (a host visible set of Error Information Registers)

8.13.1 Diagnostics Selection

The **DiagnosticsSelector** value (Byte 14h:128) selects the type of diagnostics information that the module makes available to the host in the **Diagnostics Data** area (Bytes 14h:192-255).

Table 8-126 Diagnostics Selection Register (Page 14h)

Byte	Bits	Register Name	Register Description	Type
128	7-0	DiagnosticsSelector	Select content of Diagnostics Data, see Table 8-127. Reverts to 0 if value not supported	RWW Rqd.
129	7-0	-	Reserved[1]	
130	7-0	-	Custom[1]	
131	7-0	-	Custom[1]	

The possible types of diagnostics information (detection error performance information and channel metrics) and their associated DiagnosticsSelector values are listed in Table 8-127.

The DiagnosticsSelector values 01-05h allow the host to read **intermediate results** while gating is in progress.

The DiagnosticsSelector values 11-15h allow the host to read the **results** of the last completed gating period.

The module updates the values in **Diagnostics Data** area (Bytes 14h:192-255) with the selected diagnostic data within the Diagnostics Data Content Switch time (**tDDCS**) after the **DiagnosticsSelector** has been written.

Note that this is a pure timing specification; the module does not reject further register accesses. There is no confirmation by the module when the new data are available, so host should wait for a tDDCS guard period.

Table 8-127 Diagnostics Selector Options

Diagnostics Selector	Selection	Diagnostics Data Contents (Bytes 14h:192-255)
00h	None	All zeroes
Real-Time Results		
01h	Host/Media Input Lane 1-8 BER	F16 BER values (Pre- or Post-FEC)
02h	Host Lane 1-4 errors and bits counters	U64 little endian counters with PSL indicator
03h	Host Lane 5-8 errors and bits counters	U64 little endian counters with PSL indicator
04h	Media Lane 1-4 errors and bits counters	U64 little endian counters with PSL indicator
05h	Media Lane 5-8 errors and bits counters	U64 little endian counters with PSL indicator
06h	Host/Media Input Lane 1-8 SNR(dB)	U16 little endian in units of 1/256 dB
07h-10h	-	Reserved
Results over most recently completed Gating Period (stable for Gating Period)		
11h	Gated Host/Media Input Lane 1-8 BER	F16 BER values

Diagnostics Selector	Selection	Diagnostics Data Contents (Bytes 14h:192-255)
12h	Gated Host Lane 1-4 errors and bits counters	U64 little endian counters with PSL indicator
13h	Gated Host Lane 5-8 errors and bits counters	U64 little endian counters with PSL indicator
14h	Gated Media Lane 1-4 errors and bits counters	U64 little endian counters with PSL indicator
15h	Gated Media Lane 5-8 errors and bits counters	U64 little endian counters with PSL indicator
16h-BFh	-	Reserved
C0h-FFh	-	Custom

8.13.2 Diagnostics Flags

The diagnostics Pages contain Flags that are specific to diagnostics features.

The Masks associated with these diagnostics Flags are located on Page 13h (see Table 8-123).

Table 8-128 Latched Diagnostics Flags (Page 14h)

Byte	Bits	Field Name	Field/ Register Description	Type
132	7	LossOfReferenceClockFlag	Latched loss of reference clock Flag.	RO/COR Opt.
	6-0	-	Reserved	
133	7-0	-	Reserved[1]	RO/COR Opt.
134	7	PatternCheckGatingCompleteFlagHostLane8	PatternCheckGatingCompleteFlag HostLane< i > Latched per-host lane gating complete Flag. When gating is complete, this bit will be set.	RO/COR Adv.
	6	PatternCheckGatingCompleteFlagHostLane7		
	5	PatternCheckGatingCompleteFlagHostLane6		
	4	PatternCheckGatingCompleteFlagHostLane5		
	3	PatternCheckGatingCompleteFlagHostLane4		
	2	PatternCheckGatingCompleteFlagHostLane3		
	1	PatternCheckGatingCompleteFlagHostLane2		
	0	PatternCheckGatingCompleteFlagHostLane1		
135	7	PatternCheckGatingCompleteFlagMediaLane8	PatternCheckGatingCompleteFlag MediaLane< i > Latched per-media lane gating complete Flag. When gating is complete, this bit will be set.	RO/COR Adv.
	6	PatternCheckGatingCompleteFlagMediaLane7		
	5	PatternCheckGatingCompleteFlagMediaLane6		
	4	PatternCheckGatingCompleteFlagMediaLane5		
	3	PatternCheckGatingCompleteFlagMediaLane4		
	2	PatternCheckGatingCompleteFlagMediaLane3		
	1	PatternCheckGatingCompleteFlagMediaLane2		
	0	PatternCheckGatingCompleteFlagMediaLane1		
136	7	PatternGeneratorLOLFlagHostLane8	PatternGeneratorLOLFlagHostLane< i > Latched per-host lane pattern generator failure Flag. Indicates that the generator has failed to generate a valid pattern.	RO/COR Adv.
	6	PatternGeneratorLOLFlagHostLane7		
	5	PatternGeneratorLOLFlagHostLane6		
	4	PatternGeneratorLOLFlagHostLane5		
	3	PatternGeneratorLOLFlagHostLane4		
	2	PatternGeneratorLOLFlagHostLane3		
	1	PatternGeneratorLOLFlagHostLane2		
	0	PatternGeneratorLOLFlagHostLane1		
137	7	PatternGeneratorLOLFlagMediaLane8	PatternGeneratorLOLFlagMediaLane< i > Latched per-media lane pattern generator failure Flag. Indicates that the generator has failed to generate a valid pattern.	RO/COR Adv.
	6	PatternGeneratorLOLFlagMediaLane7		
	5	PatternGeneratorLOLFlagMediaLane6		
	4	PatternGeneratorLOLFlagMediaLane5		
	3	PatternGeneratorLOLFlagMediaLane4		
	2	PatternGeneratorLOLFlagMediaLane3		
	1	PatternGeneratorLOLFlagMediaLane2		
	0	PatternGeneratorLOLFlagMediaLane1		
138	7	PatternCheckerLOLFlagHostLane8	PatternCheckerLOLFlagHostLane< i > Latched per-host lane pattern checker loss of pattern lock Flag.	RO/COR Adv.
	6	PatternCheckerLOLFlagHostLane7		
	5	PatternCheckerLOLFlagHostLane6		
	4	PatternCheckerLOLFlagHostLane5		
	3	PatternCheckerLOLFlagHostLane4		
	2	PatternCheckerLOLFlagHostLane3		

Byte	Bits	Field Name	Field/ Register Description	Type
	1	PatternCheckerLOLFlagHostLane2		
	0	PatternCheckerLOLFlagHostLane1		
139	7	PatternCheckerLOLFlagMediaLane8	PatternCheckerLOLFlagMediaLane <i><i></i> Latched per-media lane pattern checker loss of pattern lock Flag.	RO/COR Adv.
	6	PatternCheckerLOLFlagMediaLane7		
	5	PatternCheckerLOLFlagMediaLane6		
	4	PatternCheckerLOLFlagMediaLane5		
	3	PatternCheckerLOLFlagMediaLane4		
	2	PatternCheckerLOLFlagMediaLane3		
	1	PatternCheckerLOLFlagMediaLane2		
	0	PatternCheckerLOLFlagMediaLane1		
140-149	7-0	-	Reserved[10]	

8.13.3 Diagnostics Data

The **Diagnostics Selector** field value (Byte 14h:128) determines which diagnostics information is reported in the **Diagnostics Data** area (Bytes 192-255). The meaning of the selector field value is specified in Table 8-129.

Note: The following required functionality for a Pattern Synchronization Loss indication in a counter should be noted especially by module vendors.

Embedded Pattern Checker Sync Loss Indication in Total Bits Counters

The least significant bit of any ***TotalBitsCount*** register indicates if the pattern checker was always in pattern sync lock during accumulation since total bits counting had started. This bit serves as a latched pattern sync loss Flag.

An **even** total bits counter value encodes that there was not a single pattern checker sync loss while counting, so BER can be calculated reliably from counted errors and counted total bits received.

An **odd** total bits counter value encodes that there was at least one pattern checker sync loss while counting, hence the counted errors are suspicious.

Note: Regular pattern checker loss of synchronization lock (LOL) Flags are also reported in Byte 14h:136. However, embedding a latched pattern lock indication in total bits counter allows the host to read a single 2x64 byte block to measure BER from bit error counts and an associated total bits counter. Note that the error rate distortion due to this encoding is negligible.

Table 8-129 Diagnostics Data (Bytes 192-255) Contents per Diagnostics Selector (Page 14h)

Diag. Selector	Bytes	Size	Diagnostics Field	Diagnostics Data Description
00h	192-255	64	-	Reserved
01h	192-193	2	HostSideBERLane1	F16 BER of host and media lanes 1-8 (periodically updated while gating)
	194-195	2	HostSideBERLane2	
	196-197	2	HostSideBERLane3	
	198-199	2	HostSideBERLane4	
	200-201	2	HostSideBERLane5	
	202-203	2	HostSideBERLane6	
	204-205	2	HostSideBERLane7	
	206-207	2	HostSideBERLane8	
	208-209	2	MediaSideBERLane1	
	210-211	2	MediaSideBERLane2	
	212-213	2	MediaSideBERLane3	
	214-215	2	MediaSideBERLane4	
	216-217	2	MediaSideBERLane5	
02h	218-219	2	MediaSideBERLane6	
	220-221	2	MediaSideBERLane7	
	222-223	2	MediaSideBERLane8	
	192-199	8	HostSideErrorCountLane1	U64 Little-endian (LSB first) Counters for host lanes 1-4 (periodically updated while gating)
	200-207	8	HostTotalBitsCountLane1	
	208-215	8	HostSideErrorCountLane2	
	216-223	8	HostTotalBitsCountLane2	

	224-231	8	HostSideErrorCountLane3	<p><i>Note: Odd TotalBits counter values indicate transient or permanent pattern sync loss during the accumulation period.</i></p>
	232-239	8	HostTotalBitsCountLane3	
	240-247	8	HostSideErrorCountLane4	
	248-255	8	HostTotalBitsCountLane4	
03h	192-199	8	HostSideErrorCountLane5	<p>U64 Little-endian (LSB first) Counters for host lanes 5-8 (periodically updated while gating)</p> <p><i>Note: Odd TotalBits counter values indicate transient or permanent pattern sync loss during the accumulation period.</i></p>
	200-207	8	HostTotalBitsCountLane5	
	208-215	8	HostSideErrorCountLane6	
	216-223	8	HostTotalBitsCountLane6	
	224-231	8	HostSideErrorCountLane7	
	232-239	8	HostTotalBitsCountLane7	
	240-247	8	HostSideErrorCountLane8	
	248-255	8	HostTotalBitsCountLane8	
04h	192-199	8	MediaSideErrorCountLane1	<p>U64 Little-endian (LSB first) Counters for media lanes 1-4 (periodically updated while gating)</p> <p><i>Note: Odd TotalBits counter values indicate transient or permanent pattern sync loss during the accumulation period.</i></p>
	200-207	8	MediaSideTotalBitsCountLane1	
	208-215	8	MediaSideErrorCountLane2	
	216-223	8	MediaSideTotalBitsCountLane2	
	224-231	8	MediaSideErrorCountLane3	
	232-239	8	MediaSideTotalBitsCountLane3	
	240-247	8	MediaSideErrorCountLane4	
	248-255	8	MediaSideTotalBitsCountLane4	
05h	192-199	8	MediaSideErrorCountLane5	<p>U64 Little-endian (LSB first) Counters for media lanes 5-8 (periodically updated while gating)</p> <p><i>Note: Odd TotalBits counter values indicate transient or permanent pattern sync loss during the accumulation period.</i></p>
	200-207	8	MediaSideTotalBitsCountLane5	
	208-215	8	MediaSideErrorCountLane6	
	216-223	8	MediaSideTotalBitsCountLane6	
	224-231	8	MediaSideErrorCountLane7	
	232-239	8	MediaSideTotalBitsCountLane7	
	240-247	8	MediaSideErrorCountLane8	
	248-255	8	MediaSideTotalBitsCountLane8	
06h	192-193	2	-	<p>Reserved</p> <p>See section 7.1.4 for a definition.</p> <p>Reserved</p> <p>See section 7.1.4 for a definition.</p>
	194-195	2	-	
	196-197	2	-	
	198-199	2	-	
	200-201	2	-	
	202-203	2	-	
	204-205	2	-	
	206-207	2	-	
	208-209	2	HostSideSNRLane1	
	210-211	2	HostSideSNRLane2	
	212-213	2	HostSideSNRLane3	
	214-215	2	HostSideSNRLane4	
	216-217	2	HostSideSNRLane5	
	218-219	2	HostSideSNRLane6	
	220-221	2	HostSideSNRLane7	
	222-223	2	HostSideSNRLane8	
	224-225	2	-	
	226-227	2	-	
	228-229	2	-	
	230-231	2	-	
	232-233	2	-	
	234-235	2	-	
	236-237	2	-	
	238-239	2	-	
	240-241	2	MediaSideSNRLane1	<p>U16 Little-endian in units of 1/256dB SNR(dB) for host lanes 1-8 (real time status)</p> <p>See section 7.1.4 for a definition.</p>
	242-243	2	MediaSideSNRLane2	
	244-245	2	MediaSideSNRLane3	
	246-247	2	MediaSideSNRLane4	
	248-249	2	MediaSideSNRLane5	
	250-251	2	MediaSideSNRLane6	See section 7.1.4 for a definition.

	252-253	2	MediaSideSNRLane7	
	254-255	2	MediaSideSNRLane8	
11h	192-193	2	GatedHostSideBERLane1	F16 BER BER of host and media lanes 1-8 (stable gating results)
	194-195	2	GatedHostSideBERLane2	
	196-197	2	GatedHostSideBERLane3	
	198-199	2	GatedHostSideBERLane4	
	200-201	2	GatedHostSideBERLane5	
	202-203	2	GatedHostSideBERLane6	
	204-205	2	GatedHostSideBERLane7	
	206-207	2	GatedHostSideBERLane8	
	208-209	2	GatedMediaSideBERLane1	
	210-211	2	GatedMediaSideBERLane2	
	212-213	2	GatedMediaSideBERLane3	
	214-215	2	GatedMediaSideBERLane4	
	216-217	2	GatedMediaSideBERLane5	
	218-219	2	GatedMediaSideBERLane6	
	220-221	2	GatedMediaSideBERLane7	
	222-223	2	GatedMediaSideBERLane8	
12h	192-199	8	GatedHostSideErrorCountLane1	U64 Little-endian (LSB first) Counters for host lanes 1-4 (stable gating results)
	200-207	8	GatedHostSideTotalBitsCountLane1	
	208-215	8	GatedHostSideErrorCountLane2	
	216-223	8	GatedHostSideTotalBitsCountLane2	
	224-231	8	GatedHostSideErrorCountLane3	
	232-239	8	GatedHostSideTotalBitsCountLane3	
	240-247	8	GatedHostSideErrorCountLane4	
	248-255	8	GatedHostSideTotalBitsCountLane4	
13h	192-199	8	GatedHostSideErrorCountLane5	U64 Little-endian (LSB first) Counters for host lanes 5-8 (stable gating results)
	200-207	8	GatedHostSideTotalBitsCountLane5	
	208-215	8	GatedHostSideErrorCountLane6	
	216-223	8	GatedHostSideTotalBitsCountLane6	
	224-231	8	GatedHostSideErrorCountLane7	
	232-239	8	GatedHostSideTotalBitsCountLane7	
	240-247	8	GatedHostSideErrorCountLane8	
	248-255	8	GatedHostSideTotalBitsCountLane8	
14h	192-199	8	GatedMediaSideErrorCountLane1	U64 Little-endian (LSB first) Counters for media lanes 1-4 (stable gating results)
	200-207	8	GatedMediaSideTotalBitsCountLane1	
	208-215	8	GatedMediaSideErrorCountLane2	
	216-223	8	GatedMediaSideTotalBitsCountLane2	
	224-231	8	GatedMediaSideErrorCountLane3	
	232-239	8	GatedMediaSideTotalBitsCountLane3	
	240-247	8	GatedMediaSideErrorCountLane4	
	248-255	8	GatedMediaSideTotalBitsCountLane4	
15h	192-199	8	GatedMediaSideErrorCountLane5	U64 Little-endian (LSB first) Counters for media lanes 5-8 (stable gating results)
	200-207	8	GatedMediaSideTotalBitsCountLane5	
	208-215	8	GatedMediaSideErrorCountLane6	
	216-223	8	GatedMediaSideTotalBitsCountLane6	
	224-231	8	GatedMediaSideErrorCountLane7	
	232-239	8	GatedMediaSideTotalBitsCountLane7	
	240-247	8	GatedMediaSideErrorCountLane8	
	248-255	8	GatedMediaSideTotalBitsCountLane8	

8.14 Banked Page 15h (Timing Characteristics)

Page 15h is an optional Page containing per-lane timing characteristics.

The module advertises support of Page 15h in Bit 01h:145.3 (see Table 8-49).

Page 15h may optionally be Banked. Each Bank of Page 15h refers to 8 lanes.

All fields in Page 15h are read-only reporting registers (not necessarily static).

Page 15h is subdivided in subject areas as illustrated in the following table:

Table 8-130 Page 15h Overview

Byte	Size (Bytes)	Subject Area	Description
128 – 223	96	-	Reserved
224 – 239	8 x 2	Data Path Rx Latency	Total Rx delay thru the module reported by host lane
240 – 255	8 x 2	Data Path Tx Latency	Total Tx delay thru the module reported by host lane

Data Path Rx and Tx Latency

The following fields report the Data Path Rx and Tx Latency associated with each host lane in the module.

Data Path Rx Latency and Data Path Tx Latency convey the total delay thru the module, in nanoseconds, and are reported by host lane. For Data Paths with multiple lanes, all lanes shall report the same latency.

The accuracy of the reported latency values is not specified in this version of CMIS.

Note: These registers may be used to either convey static average or max latencies, per-module values that have been characterized by the manufacturer, or to convey dynamic/per-reset values that are changed by the module after each initialization. There may be other possibilities.

Note: For modules which perform dynamic updates of these registers, it is currently undefined as to when the values in these registers are guaranteed to be valid.

Table 8-131 Data Path Rx and Tx Latency, per lane (Page 15h)

Byte	Bit	Register Name	Register Array Description	Type
224-225	7-0	DataPathRxLatencyHostLane1	DataPathRxLatencyHostLane<i> U16 lane <i> Rx delay in ns	RO Rqd.
226-227	7-0	DataPathRxLatencyHostLane2		RO Rqd.
228-229	7-0	DataPathRxLatencyHostLane3		RO Rqd.
230-231	7-0	DataPathRxLatencyHostLane4		RO Rqd.
232-233	7-0	DataPathRxLatencyHostLane5		RO Rqd.
234-235	7-0	DataPathRxLatencyHostLane6		RO Rqd.
236-237	7-0	DataPathRxLatencyHostLane7		RO Rqd.
238-239	7-0	DataPathRxLatencyHostLane8		RO Rqd.
240-241	7-0	DataPathTxLatencyHostLane1	DataPathTxLatencyHostLane<i> U16 lane <i> Tx delay in ns	RO Rqd.
242-243	7-0	DataPathTxLatencyHostLane2		RO Rqd.
244-245	7-0	DataPathTxLatencyHostLane3		RO Rqd.
246-247	7-0	DataPathTxLatencyHostLane4		RO Rqd.
248-249	7-0	DataPathTxLatencyHostLane5		RO Rqd.
250-251	7-0	DataPathTxLatencyHostLane6		RO Rqd.
252-253	7-0	DataPathTxLatencyHostLane7		RO Rqd.
254-255	7-0	DataPathTxLatencyHostLane8		RO Rqd.

8.15 Banked Page 16h (Network Path Functionality)

Page 16h is an optional Page supporting the optional Network Path (**NP**) functionality that is required for multiplex or uniplex client encapsulation applications.

Concept and functionality of Network Paths and NP Applications is described in section 7.6.

The module advertises support of Page 16h (and Page 17h) in Bit 01h:142.7 (see Table 8-46).

Page 16h may optionally be Banked. Each Bank of Page 16h refers to 8 lanes.

Page 16h is subdivided in subject areas as illustrated in the following table:

Table 8-132 Page 16h Overview

Byte	Size (bytes)	Subject Area	Description
128-159	32	Provisioning	
128-135	8	NP Staged Control Set 0	Lane to NP Assignments Provisioning – Staged Control Set 0
136-143	8	NP Staged Control Set 1	Lane to NP Assignments Provisioning – Staged Control Set 1
144-159	16	-	Reserved[16]
160-175	16	Control	
160	1	NP Control	Network Path initialization control
161	1	-	Reserved[1]
162-163	2	NP Source Selectors	Signal source selection at the Network Path connection points
164-175	12	-	Reserved[12]
176-191	16	Command & Response	
176	1	NP Apply SCS 0	Apply command for NP Staged Control Set 0
177	1	NP Apply SCS 1	Apply command for NP Staged Control Set 1
178-181	4	Configuration Status	Status of most recent Network Path configuration command
182-191	10	-	Reserved[10]
192-223	32	Status	
192-199	8	NP Active Control Set	Provisioned Network Path Configuration
200-203	4	Network Path Status	Network Path State Machine state of each NP media lane
204	1	NPIInitPending Condition	Commissioning status (NPIInitPending condition)
205-223	19	-	Reserved[19]
224-255	32	Advertisement	
224-225	2	NPSM Max Durations	Maximum durations for all NPSM transient states
226	1	Options	Miscellaneous options
227	1	-	Reserved[1]
228-247	20	Mixed Multiplex Support	Advertising for mixed HP multiplexing support
248-249	2	Application Advertisement	Application Advertisement Extensions
250-255	6	-	Reserved[4]

8.15.1 Network Path Provisioning

Network Path Provisioning fields allow the host to provision Network Paths into NP Active Control Sets, for subsequent commissioning into module hardware when the NP transits through the NPInit state of the NPSM.

Note: See sections 6.2.3 and 6.2.4 for the core concepts and procedures of provisioning and commissioning Data Paths, which are very similar to those defined here (and in section 7.6.5), for Network Paths.

There are two NP Staged Control Sets (see Table 8-134 and Table 8-135), both offering the same provisioning fields, allowing the host to prepare two different configurations to be provisioned on demand.

Parallel Network Paths can be provisioned when there are no host lane conflicts and when each NP carries one of the multiplex or uniplex NP Applications advertised by the module in the Application Descriptor registers (see Table 8-23, Table 8-58, and Table 8-59) and its extension (see Table 8-148), and in the Mixed Multiplex Advertisement registers (see Table 8-150 and Table 8-151).

Note: See sections 7.6.4 and 8.15.5. for more information about NP Application advertising.

Provisioning the Data Path of an NP Application with one Network Path serving a number N of Host Paths also requires provisioning the N Host Paths (using the Data Path configuration mechanisms). The AppSel fields of all Host Paths must refer to (partial) Application Descriptors that all advertise the same media interface ID, with throughput larger than the sum of the N host interface IDs associated with the N Data Paths.

The host provisions the NP configuration prepared in an NP Staged Control Set by writing to the Apply Trigger register of the relevant NP Staged Control Set (see Table 8-138 and Table 8-139), which triggers execution of a command to update the NP Active Control Set (see Table 8-142).

Note: The NPID of all lanes belonging to host paths of the NP Application Data Path must be identical because they are all part of the same NP Application Data Path, and the NPInUse field of all those lanes must be set.

The mechanism and command handling protocol to trigger execution of a provisioning command is fully analogous to the configuration of Data Paths or Host Paths (see section 6.2.3.3). However, intentionally only step by step configuration is supported (there is no NP ApplyImmediate trigger register)

The actual commissioning of a provisioned NP configuration into hardware occurs after the host initiated state transition of the NPSM from NPDeactivated to NPInit, in the NPInit state.

Table 8-133 Network Path Provisioning per Lane (NPConfigLane<i>)

Lane	Bits	Field Name	Register Description	Type
<i>	7-4	-	Reserved	RW Rqd.
	3-1	NPID	SCS<k>::NPIDLane<i> If host lane <i> feeds the Network Path of an NP Application instance, the NPIDLane<i> field stores the Network Path ID of that Network Path, which is defined as the number of the first host lane feeding the Network Path, decremented by one. If lane<i> is unused (NPInUseLane<i> = 0), the value of NPIDLane<i> is ignored. <i>Note: All lanes of the Network Path of an Application that spans multiple host lanes have the same NPID.</i> <i>Note: For example, the NPID of a Network Path carrying the HP of host lane 1 is 0 and the NPID of a NP where host lane 5 is the lowest lane number feeding the NP is 4.</i>	
	0	NPInUse	SCS<k>::NPInUseLane<i> 0b: host lane <i> is either part of the Data Path of a DP Application or it is unused 1b: host lane <i> is part of a Host Path that feeds the Network Path identified by the NPID field in a NP Application	

Table 8-134 Staged Control Set 0, Network Path Configuration (Page 16h)

Byte	Bits	Field Name	Register Description	Type
128	7-4	-	SCS0::NPConfigLane1 See Table 8-133	RW Rqd.
	3-1	NPIDLane1		
	0	NPInUseLane1		
129	7-4	-	SCS0::NPConfigLane2 See Table 8-133	RW Rqd.
	3-1	NPIDLane2		
	0	NPInUseLane2		
130	7-4	-	SCS0::NPConfigLane3 See Table 8-133	RW Rqd.
	3-1	NPIDLane3		
	0	NPInUseLane3		
131	7-4	-	SCS0::NPConfigLane4 See Table 8-133	RW Rqd.
	3-1	NPIDLane4		
	0	NPInUseLane4		
132	7-4	-	SCS0::NPConfigLane5 See Table 8-133	RW Rqd.
	3-1	NPIDLane5		
	0	NPInUseLane5		
133	7-4	-	SCS0::NPConfigLane6 See Table 8-133	RW Rqd.
	3-1	NPIDLane6		
	0	NPInUseLane6		
134	7-4	-	SCS0::NPConfigLane7 See Table 8-133	RW Rqd.
	3-1	NPIDLane7		
	0	NPInUseLane7		
135	7-4	-	SCS0::NPConfigLane8 See Table 8-133	RW Rqd.
	3-1	NPIDLane8		
	0	NPInUseLane8		

Table 8-135 Staged Control Set 1, Network Path Configuration (Page 16h)

Byte	Bits	Field Name	Register Description	Type
136	7-4	-	SCS1::NPConfigLane1 See Table 8-133	RW Rqd.
	3-1	NPIDLane1		
	0	NPInUseLane1		
137	7-4	-	SCS1::NPConfigLane2 See Table 8-133	RW Rqd.
	3-1	NPIDLane2		
	0	NPInUseLane2		
138	7-4	-	SCS1::NPConfigLane3 See Table 8-133	RW Rqd.
	3-1	NPIDLane3		
	0	NPInUseLane3		
139	7-4	-	SCS1::NPConfigLane4 See Table 8-133	RW Rqd.
	3-1	NPIDLane4		
	0	NPInUseLane4		
140	7-4	-	SCS1::NPConfigLane5 See Table 8-133	RW Rqd.
	3-1	NPIDLane5		
	0	NPInUseLane5		
141	7-4	-	SCS1::NPConfigLane6 See Table 8-133	RW Rqd.
	3-1	NPIDLane6		
	0	NPInUseLane6		
142	7-4	-	SCS1::NPConfigLane7 See Table 8-133	RW Rqd.
	3-1	NPIDLane7		
	0	NPInUseLane7		
143	7-4	-	SCS1::NPConfigLane8 See Table 8-133	RW Rqd.
	3-1	NPIDLane8		
	0	NPInUseLane8		

8.15.2 Network Path Control

The operational status of the provisioned Network Paths (see the NP Active Control Set described in section 8.15.4) is determined by the content of the **NPDeinit** register.

The NPDeinit register controls the initialization of the media lanes in all configured Network Paths that are associated with the 8 lanes represented in a Bank.

The module evaluates the NPDeinit register only in the fully operational Module State ModuleReady. When the Module State is ModuleReady, the Network Path associated with lanes whose NPDeinit bits are cleared will transition to the NPIInit state and begin the media lane initialization process.

Note: By default, all configured Network Paths will begin initializing when the Module State reaches ModuleReady. The host can prevent this auto-initialization behavior by setting the NPDeinit bits while the module is in the ModuleLowPwr state.

At time of initialization (commissioning), the multiplex structure of the NP application must be provisioned in the Active Control set, in terms of HP definitions, but the HPs can be in any achievable DPSM state.

Parallel Network Paths are mutually independent. They may be initialized or deinitialized by one command or by a sequence of commands.

Table 8-136 Network Path Initialization Control (Page 16h)

Byte	Bits	Field Name	Register Description	Type
160	7	NPDeinitLane8	NPDeinitLane< i > Initialization control for the Network Path fed by host lane < i > 0b: Initialize the Network Path associated with host lane < i > 1b: Deinitialize the Network Path associated with host lane < i >	RW Rqd.
	6	NPDeinitLane7		
	5	NPDeinitLane6		
	4	NPDeinitLane5		
	3	NPDeinitLane4		
	2	NPDeinitLane3		
	1	NPDeinitLane2		
	0	NPDeinitLane1	All host lanes feeding one Network Path must have the same NPDeinitLane< i > value set <i>Note: These bits represent static requests, not trigger events</i>	

Supporting configurable client replacement signal data insertion in Tx direction at NP inputs from HP outputs, replacing received host signal data to be mapped into NP inputs, is optional and hence advertised.

Supporting configurable client replacement signal insertion in Rx direction at HP inputs from NP outputs, replacing received and demapped host signal data of NP outputs, is optional and hence advertised.

Note: It is assumed that the pertinent specifications of an NP application specify well defined consequent actions for the case when no valid host signal is available to be forwarded. The default configuration is therefore to forward the received signal or the signal resulting from a consequent action, as this allows for intervention-free power up. See section 7.6.1.2 and 8.15.5.2 for more information on client replacement signals

Table 8-137 Network and Host Path Signal Source Selection (Page 16h)

Byte	Bits	Field Name	Register Description	Type
162	7	HPSourceRx8	HPSourceRx< i > Controls the signal source feeding host lane < i > in Rx direction 0b: Signal received from Network Path 1b: Internally generated client replacement signal All host lanes belonging to the same HP must have the same HPSourceRx< i > value set Advertisement: 16h:226.0	RW Cond.
	6	HPSourceRx7		
	5	HPSourceRx6		
	4	HPSourceRx5		
	3	HPSourceRx4		
	2	HPSourceRx3		
	1	HPSourceRx2		
	0	HPSourceRx1		
163	7	NPSourceTx8	NPSourceTx< i > Controls the signal source feeding the Network Path input related to the HP containing host lane < i > 0b: Signal received from Host Path 1b: Internally generated client replacement signal All host lanes belonging to the same HP must have the same NPSourceTx< i > value set Advertisement: 16h:226.1	RW Cond.
	6	NPSourceTx7		
	5	NPSourceTx6		
	4	NPSourceTx5		
	3	NPSourceTx4		
	2	NPSourceTx3		
	1	NPSourceTx2		
	0	NPSourceTx1		

8.15.3 Network Path Commands

Triggering the execution of a command to **provision** a configuration prepared in an NP Staged Control Set into the NP Active Control Set by writing to an **ApplyNPInit** trigger register and the associated command handling protocol using the **NPConfigStatus** status register is fully analogous to DP command handling for Data Paths described in section 6.2.4.

The ApplyNPInit trigger register allows the host to trigger the execution of the Network Path provisioning or re-provisioning commands for the lanes selected via the lane bit mask in the value written to the trigger register.

Note: Unlike the name might suggest, writing to ApplyNPInit causes a provisioning step to be executed but does not itself cause the NPInit state to be entered, as this step is governed by the NPSM.

Note: The ApplyNPInit register is a stateless trigger registers with write-only access type. This implies that the value read from the register is not specified. Modules may use the bits in these registers for any purpose, including to signal command execution or acceptance status, e.g. for debug purposes.

Command acceptance: The module (silently) ignores a set trigger bit for lanes where execution of a previously triggered provisioning command is still in progress, as indicated in the associated NPConfigStatus field. Conversely, when the module accepts a triggered command for a lane, it immediately sets the NPConfigStatus field of that lane to ConfigInProgress.

Command validation: After setting ConfigInProgress, the module first validates the configuration to be provisioned.

Note: Full multiplex structure validation at time of provisioning may be difficult if not impossible for the module because Host Path provisioning and Network Path provisioning are separate steps with undefined order.

Command execution: After successful validation and after subsequent successful copy from the relevant NP Staged Control Set to the NP Active Control Set, the module sets the bits of the provisioned lanes in the **NPInitPending** register, indicating that the commissioning of the NP Active Control Set during the NPInit state is still outstanding.

Command termination: Finally the module updates the NPConfigStatus.

Table 8-138 Staged Control Set 0, Apply Triggers (Page 16h)

Byte	Bits	Field Name	Register Description	Type
176	7	ApplyNPInitLane8	SCS0::ApplyNPInitLane<i> 0b: No action for host lane <i>	
	6	ApplyNPInitLane7	1b: Trigger the Provision procedure using the NP Staged Control Set 0 settings for host lane <i>, with feedback provided in the associated NPConfigStatusLane<i> field	
	5	ApplyNPInitLane6		
	4	ApplyNPInitLane5		
	3	ApplyNPInitLane4		
	2	ApplyNPInitLane3		
	1	ApplyNPInitLane2		
	0	ApplyNPInitLane1	Restriction: This byte must be written in a single-byte WRITE	

Table 8-139 Staged Control Set 1, Apply Triggers (Page 16h)

Byte	Bits	Field Name	Register Description	Type
177	7	ApplyNPInitLane8	SCS1::ApplyNPInitLane<i> 0b: No action for host lane <i>	
	6	ApplyNPInitLane7	1b: Trigger the Provision procedure using the NP Staged Control Set 1 settings for host lane <i>, with feedback provided in the associated NPConfigStatusLane<i> field	
	5	ApplyNPInitLane6		
	4	ApplyNPInitLane5		
	3	ApplyNPInitLane4		
	2	ApplyNPInitLane3		
	1	ApplyNPInitLane2		
	0	ApplyNPInitLane1	Restriction: This byte must be written in a single-byte WRITE	

The NPConfigStatus register provides feedback on the command handling status (**in-progress**, **ready**) and the result status (**success**, **rejection** due to validation failure)

Table 8-140 NP Configuration Command Status registers (Page 16h)

Byte	Bit	Field Name	Field Description	Type
178	7-4	NPConfigStatusLane2	NPConfigStatusLane<i> Provisioning Command Execution / Result Status for the Network Path of host lane <i>, during and after the most recent configuration command.	RO
	3-0	NPConfigStatusLane1		Rqd.
179	7-4	NPConfigStatusLane4	See Table 8-141 for the encoding of values. <i>Note: There is no feedback to the host when an Apply trigger is ignored after failed readiness test (when another configuration is still in progress)</i>	RO
	3-0	NPConfigStatusLane3		Rqd.
180	7-4	NPConfigStatusLane6		RO
	3-0	NPConfigStatusLane5		Rqd.
181	7-4	NPConfigStatusLane8		RO
	3-0	NPConfigStatusLane7		Rqd.

The status codes in Table 8-141 represent both the current command handling status (**in-progress**, **ready**) and the result status information (**success**, **rejection** due to validation failure), whereby the **ready** command handling status is implicitly indicated by the presence of result status information.

Table 8-141 NP Configuration Command Execution and Result Status Codes (Page 16h)

Encoding	Name	Value Description
0h	ConfigUndefined	No status information available (initial register value)
1h	ConfigSuccess	Positive Result Status: The last accepted configuration command has been completed successfully
2h	ConfigRejected	Negative Result Status (2h-Bh, Dh-Fh): Configuration rejected: unspecific validation failure
3h	ConfigRejectedInvalidAppSel	Configuration rejected: invalid AppSel codes
4h	ConfigRejectedInvalidNetworkPath	Configuration rejected: invalid set of lanes for AppSel
5h	-	Reserved
6h	ConfigRejectedLanesInUse	Configuration rejected: some lanes not in NPDeactivated
7h	ConfigRejectedPartialNetworkPath	Configuration rejected: lanes are only subset of Network Path
8h	-	Reserved (other validation failures)
9h	-	
Ah	-	
Bh	-	
Ch	ConfigInProgress	Execution Status: A configuration command is still being processed by the module; a new configuration command is ignored for this lane while ConfigInProgress.
Dh	-	Custom Configuration rejected for custom reasons
Eh	-	
Fh	-	

1 8.15.4 Network Path Status

2 The Active Control Set represents the currently commissioned NP configuration, except transiently during a
 3 reconfiguration when the NPInitPending status is active.

4 **Table 8-142 NP Active Control Set, Network Path Configuration (Page 16h)**

Byte	Bits	Field Name	Register Description	Type
192	7-4	-	ACS::NPConfigLane1 See Table 8-133	RO Rqd.
	3-1	NPIDLane1		
	0	NPInUseLane1		
193	7-4	-	ACS::NPConfigLane2 See Table 8-133	RO Rqd.
	3-1	NPIDLane2		
	0	NPInUseLane2		
194	7-4	-	ACS::NPConfigLane3 See Table 8-133	RO Rqd.
	3-1	NPIDLane3		
	0	NPInUseLane3		
195	7-4	-	ACS::NPConfigLane4 See Table 8-133	RO Rqd.
	3-1	NPIDLane4		
	0	NPInUseLane4		
196	7-4	-	ACS::NPConfigLane5 See Table 8-133	RO Rqd.
	3-1	NPIDLane5		
	0	NPInUseLane5		
197	7-4	-	ACS::NPConfigLane6 See Table 8-133	RO Rqd.
	3-1	NPIDLane6		
	0	NPInUseLane6		
198	7-4	-	ACS::NPConfigLane7 See Table 8-133	RO Rqd.
	3-1	NPIDLane7		
	0	NPInUseLane7		
199	7-4	-	ACS::NPConfigLane8 See Table 8-133	RO Rqd.
	3-1	NPIDLane8		
	0	NPInUseLane8		

5
 6 The current state of the NPSM associated with a host lane (if any) is indicated as follows, using the encoding
 7 defined in Table 8-144

8 **Table 8-143 Lane-associated Network Path States (Page 16h)**

Byte	Bit	Field Name	Register Description (NPStateHostLane<i><i></i>)	Type
200	7-4	NPStateHostLane2	Network Path State of host lane 2 (see Table 8-144)	RO Rqd.
	3-0	NPStateHostLane1	Network Path State of host lane 1 (see Table 8-144)	
201	7-4	NPStateHostLane4	Network Path State of host lane 4 (see Table 8-144)	RO Rqd.
	3-0	NPStateHostLane3	Network Path State of host lane 3 (see Table 8-144)	
202	7-4	NPStateHostLane6	Network Path State of host lane 6 (see Table 8-144)	RO Rqd.
	3-0	NPStateHostLane5	Network Path State of host lane 5 (see Table 8-144)	
203	7-4	NPStateHostLane8	Network Path State of host lane 8 (see Table 8-144)	RO Rqd.
	3-0	NPStateHostLane7	Network Path State of host lane 7 (see Table 8-144)	

9 **Table 8-144 Network Path State Encoding**

Encoding	State
0h	Reserved
1h	NPDeactivated (or unused lane)
2h	NPInit
3h	NPDeinit
4h	NPAActivated
5h	NPTxTurnOn
6h	NPTxTurnOff
7h	NPIinitialized
8h-Fh	Reserved

1

Table 8-145 Network Path Conditions (Page 16h)

Byte	Bits	Field Name	Register Description	Type
204	7	NPIInitPendingLane8	NPIInitPendingLane<i> 0b: NPIInit not pending 1b: Commissioning the NP Active Control Set during NPIInit has not yet been executed after a successful ApplyNPIInit, hence the NP Active Control Set content may still deviate from the actual hardware configuration.	RO Rqd.

2

8.15.5 Network Path Related Advertisements (Capabilities and Restrictions)

8.15.5.1 Maximum Durations Advertisement

The maximum duration of transient NPSM states is advertised by the module as follows

Table 8-146 NPSM Durations Advertising (Page 16h)

Byte	Bit	Field Name	Field Description	Type
224	7-4	MaxDurationNPDeinit	Maximum duration of the NPDeinit state (encoded as per Table 8-48)	RO Rqd.
	3-0	MaxDurationNPInit	Maximum duration of the NPInit state (encoded as per Table 8-48)	
225	7-4	MaxDurationNPTxTurnOff	Encoded maximum duration of the NPTxTurnOff state (see Table 8-48)	RO Rqd.
	3-0	MaxDurationNPTxTurnOn	Encoded maximum duration of the NPTxTurnOn state (see Table 8-48)	

8.15.5.2 Miscellaneous Options

A module advertises if the host can configure a host replacement signal to be forwarded in downstream direction, instead of the (possibly but not necessarily failed) host signal received from an upstream source.

Note: Replacement signal insertion is intended to be used when no host signal is connected to a multiplex connection point of the NP (Tx direction) or when no host signal is expected from a multiplex connection point of the NP (Rx direction), i.e. always when a host signal is known to be intentionally missing or "unequipped".

Note: It is assumed that the pertinent transmission specifications of an NP application provide transmit and forwarding specifications in case of an upstream host signal failure, such as automatic AIS or LF insertion. Reactive modifications of the data stream, automatically performed by the module, prior to forwarding the data stream are often called "consequent actions".

Note: When configurable replacement signals are not supported, a deliberately unconnected host signal would be treated by the module as a host signal failure, while supporting configurable replacement signals may allow to distinguish an intentionally unconnected signal from a signal failure leading to consequent actions.

Table 8-147 Miscellaneous Options (Page 16h)

Byte	Bit	Field Name	Field Description	Type
226	7-2	-	Reserved	RO Rqd.
	1	ReplaceHPSignalTxSupported	0b: not supported 1b: Tx replacement signals for NP inputs are supported	
	0	ReplaceHPSignalRxSupported	0b: not supported 1b: Rx replacement signals for NP outputs are supported	

8.15.5.3 Application Advertisement Extensions

As an extension to the basic application advertisement described in section 8.2.12 a module supporting NP Applications advertises for each Application Descriptor (identified by its AppSel code) if the advertised application is a DP Application or a NP Application.

This application advertisement extension is necessary for robust distinction of system interface applications (DP Applications) and uniplex applications (NP Applications).

Note: An Application Descriptor for a genuine homogeneous multiplex application can also be recognized from the data rates associated with the MediaInterfaceID and the HostInterfaceID: The number of homogeneously multiplexed Host Paths is the quotient of the larger information rate of the media interface (MediaInterfaceID) and the smaller information rate of the host interface (HostInterfaceID) in the Application Descriptor.

Note: The extension to the Application Descriptor is required (for modules supporting NP Applications) because the distinction of a uniplex NP Application and a DP Application cannot always be derived from the meaning of the advertised MediaInterfaceID, and hence the distinction must be indicated explicitly.

Table 8-148 NP Extended Application Advertisement (Page 16h)

Byte	Bits	Field Name	Register Description	Type	
248	7	ExtAppDescriptor15	ExtAppDescriptor< i > The Application Descriptor identified by AppSel< i >: 0b: describes a supported DP Application 1b: (partially) describes a supported NP Application	RO Rqd.	
	6	ExtAppDescriptor14			
	5	ExtAppDescriptor13			
	4	ExtAppDescriptor12			
	3	ExtAppDescriptor11			
	2	ExtAppDescriptor10			
	1	ExtAppDescriptor9			
	0	ExtAppDescriptor8			
249	7	ExtAppDescriptor7	<i>Note: The value is irrelevant when the Application Descriptor identified by AppSel< i > is unused.</i>		
	6	ExtAppDescriptor6			
	5	ExtAppDescriptor5			
	4	ExtAppDescriptor4			
	3	ExtAppDescriptor3			
	2	ExtAppDescriptor2			
	1	ExtAppDescriptor1			
	0	-			

8.15.5.4 Multiplex and Uniplex Application Advertisement

NP Applications are advertised by one or more Application Descriptors (see Table 8-23, Table 8-58, and Table 8-59), each of which indicates a HP type (HostInterfaceID) and the NP type (MediaInterfaceID), together with the Application Descriptor Extensions (see Table 8-148) distinguishing NP Applications and DP Applications.

Uniplex NP and Homogeneous Multiplex Application Advertisement

A uniplex or homogeneous multiplex NP Application is advertised in one Application Descriptor with AppSel=< i > where the associated ExtAppDescriptor< i > bit is set.

Note: One characteristic resulting from a single Application Descriptor being used is that the host lane data rates of all tributaries are identical, simply because they all use the same HostInterfaceID.

Mixed Multiplex NP Application Advertisement

A mixed (heterogenous) multiplex application is advertised via a **set** of mutually **consistent** Application Descriptors with all the associated ExtAppDescriptor< i > bits set, together with the additional Mixed Multiplex Descriptor registers (see Table 8-150 and Table 8-151 and section 8.15.5.5).

Application Descriptors are mutually **consistent** when they all advertise the same MediaInterfaceID with at least one common option in the MediaLaneAssignmentOptions, and when there is at least one combination of HostLaneAssignmentOptions in the Application Descriptors that assigns different host lanes to each of them.

A non-trivial **uniform lane data rate restriction** for a **mixed multiplex** application requires that the host lane data rates of all tributaries must be identical. See also subsection 7.6.1.5.

Note: In other words, all tributaries of a multiplex application are multi-lane signals with possibly different number of lanes but with uniform lane data rate. As specified below in section 8.15.5.5, this lane data uniformity requirement extends even across parallel multiplex applications.

Note: The rationale for this apparently undesired restriction is to keep accurate advertising both of capabilities and of restrictions at a reasonable level of advertisement complexity.

8.15.5.5 Constraints and Advertisements for Parallel NP Applications

The **uniform lane data rate restriction** applies also across **parallel** multiplexing applications (of any kind).

Note: The rationale for this apparently undesired restriction is to keep accurate advertising both of capabilities and of restrictions, as defined below, at a reasonable level of advertisement complexity.

The ordered list of Host Path lane groups feeding one single Network Path, in lane number order, is called the **multiplex structure** of the Network Path.

The ordered list of all Host Path lane groups (feeding any of possibly several parallel Network Paths) is called the **global multiplex structure**.

The **multiplexing granularity** of a multiplex structure is defined by a HostInterfaceID value that indicates both the common host lane data rate and the multiplexing rate granularity of a mixed multiplex structure.

A particular global multiplex structure is supported when

- The data rate of each Host Path is a **power-of-two multiple** of one selected multiplexing granularity
- The information rate of each Network Path is the **sum of the** information **rates** of all its Host Paths
- The **lane grouping** of the host lanes into Host Paths is advertised as **supported** by the module

Table 8-149 lists all possible lane groupings (global multiplex structures) together with their global multiplex structure encodings (Multiplex Structure IDs) for mixed Host Path widths in a multiplex application.

Table 8-149 also shows the Host Path DPIDs (HP DPIDs) for the Host Paths of a particular global multiplex structure, independent of any NP they belong to.

When parallel Network Paths are supported, the first NP is associated with one or more HPs using a first group of host lanes, and the n^{th} NP is associated with one or more HPs using the n^{th} group of host lanes.

Note. This table has the same hand-crafted structure as Table 8-38, which is not scalable to fewer or more lanes. It is reused here because a look-up table is expected to be used in implementations. The ID numbering (0-25) is different (off by one) because here the multiplex structure ID selects a bit position in a bit mask.

Table 8-149 Multiplex Lane Grouping Advertisement

Multiplex Structure			HP DPID per Host Lane #							
ID	# of HPs	HP Widths	1	2	3	4	5	6	7	8
0	8	1, 1, 1, 1, 1, 1, 1, 1	1	2	3	4	5	6	7	8
1	1	8	1	1	1	1	1	1	1	1
2	2	4, 4	1	1	1	1	5	5	5	5
3	4	2, 2, 2, 2	1	1	3	3	5	5	7	7
4	3	4, 2, 2	1	1	1	1	5	5	7	7
5	4	4, 2, 1, 1	1	1	1	1	5	5	7	8
6	4	4, 1, 1, 2	1	1	1	1	5	6	7	7
7	5	4, 1, 1, 1, 1	1	1	1	1	5	6	7	8
8	3	2, 2, 4	1	1	3	3	5	5	5	5
9	4	2, 1, 1, 4	1	1	3	4	5	5	5	5
10	4	1, 1, 2, 4	1	2	3	3	5	5	5	5
11	4	1, 1, 1, 1, 4	1	2	3	4	5	5	5	5
12	5	2, 2, 2, 1, 1	1	1	3	3	5	5	7	8
13	5	2, 2, 1, 1, 2	1	1	3	3	5	6	7	7
14	5	2, 1, 1, 2, 2	1	1	3		5	5	7	7
15	5	1, 1, 2, 2, 2	1	2	3	3	5	5	7	7
16	6	2, 2, 1, 1, 1, 1	1	1	3	3	5	6	7	8
17	6	2, 1, 1, 2, 1, 1	1	1	3	4	5	5	7	8
18	6	2, 1, 1, 1, 1, 2	1	1	3	4	5	6	7	7
19	6	1, 1, 2, 2, 1, 1	1	2	3	3	5	5	7	8
20	6	1, 1, 2, 1, 1, 2	1	2	3	3	5	6	7	7
21	6	1, 1, 1, 1, 2, 2	1	2	3	4	5	5	7	7
22	6	2, 1, 1, 1, 1, 1	1	1	3	4	5	6	7	8
23	7	1, 1, 2, 1, 1, 1, 1	1	2	3	3	5	6	7	8
24	7	1, 1, 1, 1, 2, 1, 1	1	2	3	4	5	5	7	8
25	7	1, 1, 1, 1, 1, 1, 2	1	2	3	4	5	6	7	7

Based on this coded enumeration of global multiplex structures, the module advertises the supported global multiplex options by a list of up to four **Multiplex Descriptors** as follows:

1 Each Multiplex Descriptor advertises a multiplex granularity **MuxGranularity** (bandwidth) and a set of
 2 supported global multiplex structures **MuxStructsSupported** for that granularity.

3 For reasons of access efficiency, the two parts of a Multiplex Descriptor (granularity and list of global multiplex
 4 structures) are stored in two separate register arrays, as follows:

5 **Table 8-150 Multiplex Granularities Advertisement (Page 16h)**

Byte	Bits	Field Name	Register Description	Type
228	7-0	MuxGranularity1	U8 MuxGranularity<i> 0: Not supported (end of granularity list: after a zero value, all following MuxGranularity<i> fields are zero as well) >0: HostInterfaceID indicating the lane data rate and the multiplex rate granularity: all multiplexed signals have the same lane data rate and a data rate that is a power of two multiple of the multiplex rate granularity	RO Rqd.
229	7-0	MuxGranularity2		
230	7-0	MuxGranularity3		
231	7-0	MuxGranularity4		

6 **Table 8-151 Global Multiplex Structures Advertisement (Page 16h)**

Byte	Bits	Field Name	Register Description	Type
232-235	31-0	MuxStructsSupported1	U32 MuxStructsSupported<i> contains a bit mask where each bit <j> set indicates support for the multiplex structure ID <j> as defined in Table 8-149	RO Rqd.
236-239	31-0	MuxStructsSupported2		
240-243	31-0	MuxStructsSupported3		
244-247	31-0	MuxStructsSupported4		

8.16 Banked Page 17h (Flags and Masks)

Page 17h is a conditionally required Page that provides Flags and Masks for various optional features.

Note: Collecting Flags for different features on one Page allows hosts to efficiently search for the source of an Interrupt, without feature related Page changes.

The module supports this page when one of the relevant features is advertised as supported.

Page 17h may optionally be Banked. Each Bank of Page 17h refers to 8 lanes.

Page 17h is subdivided into equally sized Flags and Masks areas as illustrated in the following table:

Table 8-152 Page 17h Overview

Byte	Size (bytes)	Subject Area	Description
128	1	Network Path Flags	Network Path related Flags
129-191	63	-	Reserved[63]
192	1	Network Path Masks	Network Path related Masks
193-255	63	-	Reserved[63]

8.16.1 Flags

Table 8-153 Network Path Related Flags (Page 17h)

Byte	Bit	Field Name	Register Description	Type
128	7	NPStateChangedFlag8	NPStateChangedFlag<i>	
	6	NPStateChangedFlag7	Latched Network Path State Changed Flag, for Network Path of host lane <i>	
	5	NPStateChangedFlag6		
	4	NPStateChangedFlag5	Condition: Page 16h supported	
	3	NPStateChangedFlag4		
	2	NPStateChangedFlag3		
	1	NPStateChangedFlag2		
	0	NPStateChangedFlag1		

8.16.2 Masks

Table 8-154 Network Path Related Masks (Page 17h)

Byte	Bit	Field Name	Register Description	Type
192	7	NPStateChangedMask8	NPStateChangedMask<i>	
	6	NPStateChangedMask7	Mask for NPStateChangedFlag<i> for Network Path of host lane <i>	
	5	NPStateChangedMask6		
	4	NPStateChangedMask5	Condition: Page 16h supported	
	3	NPStateChangedMask4		
	2	NPStateChangedMask3		
	1	NPStateChangedMask2		
	0	NPStateChangedMask1		

8.17 Banked Page 18h (Lane Control and Data Path Control Part 2)

Page 18h may optionally be Banked. Each Bank of Page 18h refers to a group of 8 lanes.

Page 18h is supported when page 1Ch is supported or when usage of the VCS parameter space is required as per specifications in [6].

Page 18h is subdivided into areas as illustrated in the following table:

Table 8-155 Page 18h Overview

Bytes	Size (bytes)	Subject Area	Description
128-135	8	SCS0 Extension	Per Lane NAD Block Numbers for Normalized Application Descriptors in Staged Control Set 0
136-143	8	SCS 1 Extension	Per Lane NAD Block Numbers for Normalized Application Descriptors in Staged Control Set 1
144-199	56	SCS 0 Versatile Control Set Parameter Space	Restricted – defined in CMIS-VCS [6] Support of CMIS-VCS is advertised in: 01h:162.7.
200-255	56	SCS 1 Versatile Control Set Parameter Space	

8.17.1 Staged Control Set Extension – Application Descriptors Block Indices

The classical Control Sets identify an Application by an AppSel code ranging from 1 to 15.

When Normalized Application Descriptors (NAD) are used (e.g. to support more than 15 Applications), the relevant Normalized Application Descriptor Block must be identified. See also sections 6.2.1.4.2 and 8.20.

Table 8-156 Staged Control Set 0 – Normalized Application Descriptor Block Indices (Page 18h)

Byte	Bits	Field Name	Register Description	Type
128	7-4	-	SCS0:: NADBlockIndex <lane>	RW Cnd.
	3-0	NADBlockIndex1		
129	7-4	-	Condition: Page 1Ch supported	RW Cnd.
	3-0	NADBlockIndex2		
130	7-4	-	Bits 7-4 are Reserved to allow future expansion	RW Cnd.
	3-0	NADBlockIndex3		
131	7-4	-		RW Cnd.
	3-0	NADBlockIndex4		
132	7-4	-		RW Cnd.
	3-0	NADBlockIndex5		
133	7-4	-		RW Cnd.
	3-0	NADBlockIndex6		
134	7-4	-		RW Cnd.
	3-0	NADBlockIndex7		
135	7-4	-		RW Cnd.
	3-0	NADBlockIndex8		

Table 8-157 Staged Control Set 1 – Normalized Application Descriptor Block Indices (Page 18h)

Byte	Bits	Field Name	Register Description	Type
136	7-4	-	SCS1:: NADBlockIndex <lane>	RW Cnd.
	3-0	NADBlockIndex1		
137	7-4	-	Condition: Page 1Ch supported and SCS1 supported	RW Cnd.
	3-0	NADBlockIndex2		
138	7-4	-	Bits 7-4 are Reserved to allow future expansion	RW Cnd.
	3-0	NADBlockIndex3		
139	7-4	-		RW Cnd.
	3-0	NADBlockIndex4		
140	7-4	-		RW Cnd.
	3-0	NADBlockIndex5		
141	7-4	-		RW Cnd.
	3-0	NADBlockIndex6		

Byte	Bits	Field Name	Register Description	Type
142	7-4	-		RW
	3-0	NADBlockIndex7		Cnd.
143	7-4	-		RW
	3-0	NADBlockIndex8		Cnd.

8.18 Banked Page 19h (Lane Status and Data Path Status Part 2)

Page 19h is an optional Page that is present when any of the associated advertised features require its presence.

All fields on Page 19h are read-only.

Page 19h is supported when unidirectional reconfiguration is supported or when usage of the ACS parameter space as per specifications in [6].

Page 19h may optionally be Banked. Each Bank of Page 19h refers to a group of 8 lanes.

Page 19h is subdivided into areas as illustrated in the following table:

Table 8-158 Page 19h Overview

Bytes	Size (bytes)	Subject Area	Description
128-135	8	Active Control Set Data Path Configuration Tx	Provisioned DP settings for Tx , mirrored in 11h:206-213 Advertisement: 01h:162.6
136-143	8	Active Control Set Data Path Configuration Rx	Provisioned DP settings for Rx Advertisement: 01h:162.6
144-151	8	ACS Extension	Per Lane NAD Block Indices of Normalized Application Descriptors in Active Control Set
152-207	56	ACS Versatile Control Set Parameter Space	Restricted – defined in CMIS-VCS [6] CMIS-VCS support is advertised in: 01h:162.7 Condition: See [6]
208-255	48	-	Reserved

8.18.1 Direction-specific Provisioned Data Path Configuration

When the module supports independent reconfiguration of the Tx and Rx directions (when the advertisement bit UnidirReconfigSupported is set), the provisioned DataPath Configuration settings of Rx and Tx directions may become different when the host uses the ApplyDPIInitTx and ApplyDPIInitRx triggers. See section 8.10.6.

In this case (when the UnidirReconfigSupported bit is set) the DPConfigTx registers (see Table 8-159) and the DPConfigRx registers (see Table 8-160) display the provisioned Data Path Configuration per direction.

Table 8-159 Active Control Set, Provisioned Tx Data Path Configuration (Page 19h)

Byte	Bits	Field Name	Register Description	Type
128	7-4	AppSelCodeTx1	ACS::DPConfigTx1 See Table 8-92	RO Cnd.
	3-1	DataPathIDTx1		
	0	ExplicitControlTx1		
129	7-4	AppSelCodeTx2	ACS::DPConfigTx2 See Table 8-92	RO Cnd.
	3-1	DataPathIDTx2		
	0	ExplicitControlTx2		
130	7-4	AppSelCodeTx3	ACS::DPConfigTx3 See Table 8-92	RO Cnd.
	3-1	DataPathIDTx3		
	0	ExplicitControlTx3		
131	7-4	AppSelCodeTx4	ACS::DPConfigTx4 See Table 8-92	RO Cnd.
	3-1	DataPathIDTx4		
	0	ExplicitControlTx4		
132	7-4	AppSelCodeTx5	ACS::DPConfigTx5 See Table 8-92	RO Cnd.
	3-1	DataPathIDTx5		
	0	ExplicitControlTx5		
133	7-4	AppSelCodeTx6	ACS::DPConfigTx6 See Table 8-92	RO Cnd.
	3-1	DataPathIDTx6		
	0	ExplicitControlTx6		
134	7-4	AppSelCodeTx7	ACS::DPConfigTx7 See Table 8-92	RO Cnd.
	3-1	DataPathIDTx7		
	0	ExplicitControlTx7		
135	7-4	AppSelCodeTx8	ACS::DPConfigTx8 See Table 8-92	RO Cnd.
	3-1	DataPathIDTx8		
	0	ExplicitControlTx8		

1
2**Table 8-160 Active Control Set, Provisioned Rx Data Path Configuration (Page 19h)**

Byte	Bits	Field Name	Register Description	Type
136	7-4	AppSelCodeRx1	ACS::DPConfigRx1 See Table 8-92	RO Cnd.
	3-1	DataPathIDRx1		
	0	ExplicitControlRx1		
137	7-4	AppSelCodeRx2	ACS::DPConfigRx2 See Table 8-92	RO Cnd.
	3-1	DataPathIDRx2		
	0	ExplicitControlRx2		
138	7-4	AppSelCodeRx3	ACS::DPConfigRx3 See Table 8-92	RO Cnd.
	3-1	DataPathIDRx3		
	0	ExplicitControlRx3		
139	7-4	AppSelCodeRx4	ACS::DPConfigRx4 See Table 8-92	RO Cnd.
	3-1	DataPathIDRx4		
	0	ExplicitControlRx4		
140	7-4	AppSelCodeRx5	ACS::DPConfigRx5 See Table 8-92	RO Cnd.
	3-1	DataPathIDRx5		
	0	ExplicitControlRx5		
141	7-4	AppSelCodeRx6	ACS::DPConfigRx6 See Table 8-92	RO Cnd.
	3-1	DataPathIDRx6		
	0	ExplicitControlRx6		
142	7-4	AppSelCodeRx7	ACS::DPConfigRx7 See Table 8-92	RO Cnd.
	3-1	DataPathIDRx7		
	0	ExplicitControlRx7		
143	7-4	AppSelCodeRx8	ACS::DPConfigRx8 See Table 8-92	RO Cnd.
	3-1	DataPathIDRx8		
	0	ExplicitControlRx8		

3

4

8.18.2 Active Control Set Extension – Applications Block Select5
6
7

The classical Control Sets identify an Application by an AppSel code ranging from 1 to 15, which refers to the original (0'th) Application Descriptors Segment. When more than 15 Applications are advertised, the relevant Normalized Application Descriptor Block (see section 6.2.1.4.2) must be identified.

8

Table 8-161 Active Control Set – Normalized Application Descriptor Block Indices (Page 19h)

Byte	Bits	Field Name	Register Description	Type
144	7-4	-	ACS::NADBlockIndex<lane> Bits 7-4 are Reserved to allow future expansion	RO Cnd.
	3-0	NADBlockIndex1		
145	7-4	-		RO Cnd.
	3-0	NADBlockIndex2		
146	7-4	-		RO Cnd.
	3-0	NADBlockIndex3		
147	7-4	-		RO Cnd.
	3-0	NADBlockIndex4		
148	7-4	-		RO Cnd.
	3-0	NADBlockIndex5		
149	7-4	-		RO Cnd.
	3-0	NADBlockIndex6		
150	7-4	-		RO Cnd.
	3-0	NADBlockIndex7		
151	7-4	-		RO Cnd.
	3-0	NADBlockIndex8		

9

10

8.19 Banked Pages Range 1Ah-1Bh (Resource Modules)

1 These Banked Pages are defined in Resource Module specifications such as [10].

2 The Resource Module Type (if any) is advertised in the ModuleFunctionType field 00h:57 (see Table 8-18)

8.20 Banked Page 1Ch (Normalized Application Descriptors)

Page 1Ch is an optional Banked Page that is present when **NADBanksSupported** is non-zero (see Table 8-57).

All fields on Page 1Ch are read-only and static.

Each of the NADBanksSupported Banks provides space for 15 **Normalized Application Descriptors**

The concept of Normalized Application Descriptors is introduced in section 6.2.1.4.2.

The first **NAD Block** located in Bank 0 **mirrors** the content of the always available list of basic Application Descriptors that are allocated in Lower Memory (see section 8.2.12) and on Page 1h (see section 8.4.15).

Unlike the basic Application Descriptors, which are scattered in memory, each Normalized Application Descriptor (**NAD**) is a contiguous 8-byte data structure (see Table 8-163) and allows systematic and convenient access.

Page 1Ch is subdivided into areas as illustrated in the following table:

Table 8-162 Page 1Ch Overview

Bytes	Size (bytes)	Subject Area	Description
128-247	120	Normalized Application Descriptor Block	NAD Block containing an array of 15 normalized application descriptor (NAD) instances
248-255	8	-	Reserved [8]

8.20.1 Normalized Application Descriptor Format

Table 8-163 Normalized Application Descriptor (NAD) Structure (Page 1Ch)

Byte Offset	Bits	Field Name	Register Description	Type
0	7-0	HostInterfaceID	See section 0	
1	7-0	MediaInterfaceID		
2	7-4	HostLaneCount		
	3-0	MediaLaneCount		
3	7-0	HostLaneAssignmentOptions		
4	7-0	MediaLaneAssignmentOptions		
5	7	NetworkPathIndicator	Extension for Network Paths, see section 8.15.5.3 0b: NAD describes a Data Path 1b: NAD describes a Network Path	RO Rqd.
	6-0	-	Reserved	
6	7-0	-		
7	7-0	-		

8.20.2 Normalized Application Descriptor Block

Table 8-164 Normalized Application Descriptor Block (Page 1Ch)

Byte	Bits	Field Name	Register Description	Type
128-135	7-0	NAD1	Date structure NAD<i> , with AppSel code = <i> = 1..15, accessed on Bank Index <j>, j = 0..15	RO Rqd.
136-143	7-0	NAD2		
144-151	7-0	NAD3		
152-159	7-0	NAD4		
160-167	7-0	NAD5		
168-175	7-0	NAD6		
176-183	7-0	NAD7		
184-191	7-0	NAD8		
192-199	7-0	NAD9		
200-207	7-0	NAD10		
208-215	7-0	NAD11		
216-223	7-0	NAD12		
224-231	7-0	NAD13		
232-239	7-0	NAD14		
240-247	7-0	NAD15	<i>Note: To select the Application with this AN, e.g. in a Control Set, the settings AppSelCode = <i> and NADBlockIndex = <j> must be used (see e.g. Table 8-72, Table 8-77, Table 8-93, as well as Table 8-156, Table 8-157, and Table 8-161 for the relevant fields)</i>	

8.21 Banked Page 1Dh (Host Lane Switching)

Page 1Dh is an optional Banked Page that is present when advertised in 01h:252.7 (see section 8.4.16)

Each supported Bank provides space for host lane switching functionality within a group of 8 lanes.

See section 7.8 for a description of this optional feature.

Page 1Dh is subdivided into areas as illustrated in the following table:

Table 8-165 Page 1Dh Overview

Bytes	Size (bytes)	Subject Area	Description
128-135	8	Advertisement	Host Lane Switching related features
136-151	16	Provisioning	Provisioned Configuration of Host Lane Switch
152-159	8	Configuration	Enable Host Lane Switching
160-167	8	Commands	Triggers for command execution
168-183	16	Results	Command execution status and results
184-199	16	Status	Committed Configuration of Host Lane Switch
200-255	56	-	Reserved [56]

Table 8-166 Host Lane Switching (Page 1Dh)

Byte	Bits	Field Name	Register Description	Type
Advertisement				
128	7-4	MaxRedirectionCommitDuration	Maximum duration of a CommitRedirection command execution being in progress, encoded as defined in Table 8-48.	RO Rqd.
128	3-0	-	Reserved	
129-135	7-0	-	Reserved [7]	RO
Provisioning				
136		RedirectionOfLane1	RedirectionOfLane <i> contains the nominal host lane number j in {1, ..., 8} to which the electrical host lane with lane number <i> is to be redirected. A valid redirection is consistent with the granularity of active Data Paths and with the lane ordering within the Data Paths, and it is a permutation of the lanes.	RW Rqd.
137		RedirectionOfLane2		
138		RedirectionOfLane3		
139		RedirectionOfLane4		
140		RedirectionOfLane5		
141		RedirectionOfLane6		
142		RedirectionOfLane7		
143		RedirectionOfLane8		
144-151		-	Reserved [8]	
Configuration				
152	7-1	-	Reserved	
152	0	EnableHostLaneRedirection	0b: disabled: commit command is without effect 1b: enabled: commit command is effective When enabled initially, for the first time, the RedirectionOfLane register shows the trivial unpermuted host lane switching, i.e. the commit command is not executed automatically. When disabled, the current switch configuration remains in effect.	RW Rqd.
153-159		-	Reserved [7]	
Commands				
160	7-1	-	Reserved	
160	0	CommitRedirection	1b: Commit RedirectionOfLane register to HW. Traffic integrity on lanes with unchanged redirection shall not be affected. Like elsewhere in CMIS the command shall be validated prior to execution, and nothing shall	WO/ SC Rqd.

			be changed in case of a validation failure.	
161-167	-		Reserved [7]	
Command Results				
168		RedirectionCommitResult	RedirectionCommitResult < <i>i</i> > encodes the result of the last (or current) host lane switching commit command on lane < <i>i</i> > 0: No status (initial value) 1: Command execution successful 2: Command execution in progress 3: Rejected, unspecific validation failure 4: Rejected, invalid configuration (no perm.) 5: Rejected, inconsistent with active Data Path 6: Rejected, lane ordering not supported >6 Reserved	RO Rqd.
169-183	-		Reserved [15]	
Configuration Status				
184		RedirectStatusOfLane1	RedirectStatusOfLane < <i>i</i> > contains the nominal lane number <i>j</i> in {1, ..., 8} to which the electrical host lane < <i>i</i> > is redirected Default value: <i>j</i> = <i>i</i>	RO Rqd.
185		RedirectStatusOfLane2		
186		RedirectStatusOfLane3		
187		RedirectStatusOfLane4		
188		RedirectStatusOfLane5		
189		RedirectStatusOfLane6		
190		RedirectStatusOfLane7		
191		RedirectStatusOfLane8		
192-199	-		Reserved [8]	

8.22 Banked Pages Range 20h-2Fh (VDM)

Versatile diagnostics monitoring (VDM) is an optional feature allowing extension of the types of observables that can be monitored beyond those defined in Lower Memory and in Page 11h.

Basic monitoring functionality includes both access to a recent sample and threshold crossing detection.

Advanced monitoring functionality provides interval statistics (min, max, average) for certain observables.

The module advertises support of the VDM feature in bit 01h:142.6 (see Table 8-46)

A module may offer up to 256 so called VDM instances ("slots") for up to 256 observables.

The data of VDM instances are available in the VDM Page range 20h-2Fh in groups of 64 VDM instances.

Additional Banks of Pages can provide VDM instances for modules with more than 8 lanes.

See also section 7.1 for more information and an introduction to VDM functionality.

Introduction and Overview

A **VDM instance** is a numbered "slot" in the VDM Pages which is either active and associated with a basic observable or with a statistic of a basic observable (min, max, average), or passive and unused.

The observable associated with a VDM instance can be related to a lane, to a Data Path, or to the entire module.

The association of VDM instances and basic observables or statistics is module-defined: a module describes each supported VDM instance in an array of **VDM instance descriptors** which is distributed over up to four Pages (in groups of 64 descriptors per Page).

The descriptor of a VDM instance either identifies the associated observable (basic or statistic) from a list of supported observable types or marks the VDM instance as unused (passive). For active VDM instances, the descriptor points to an associated set of four module-defined thresholds that the module uses to detect and flag threshold crossings. The host can set a Mask for each threshold crossing Flag to suppress contribution of that Flag to Interrupt request generation.

Note: Different modules will often support monitoring the same observables, but the VDM instance used for each observable is not pre-defined. The host will have to search and map desired monitors from the VDM instance descriptor array.

In case of a **basic observable**, VDM periodically reports "current value" samples of that observable via the sample register of its VDM instance (in Pages 24h-27h) and moreover it provides high/low alarm/warning threshold crossing supervision for that observable by setting the appropriate Flags on threshold crossing.

In case of a **statistics observable** (min, max, average) the host can read current values of the statistics "live" as updated by the module, but it can also request the module to stop (freeze) the current value update in order to ensure consistent reading of statistics results across one or more VDM instances (see section 8.22.6). The module performs threshold crossing detection also for statistics observables.

An important side effect of (globally) freezing the statistics reporting registers is that the module internally starts a new statistics interval by resetting its internal statistics variables at precisely the time when it performs the last update on a then frozen register.

This freeze-and-reset mechanism enables gap-free **interval statistics** over host-defined intervals.

Overview of VDM Pages

The following table provides an overview of the VDM Pages and their purpose:

Table 8-167 Summary of Page Definitions for Page 20h-2Fh

Page	Subject Area	Description
20h	Descriptors for VDM Instances 1-64 (Group 1)	Every 2-byte value describes a VDM instance offered by the module.
21h	Descriptors for VDM Instances 65-128 (Group 2)	
22h	Descriptors for VDM Instances 129-192 (Group 3)	
23h	Descriptors for VDM Instances 193-256 (Group 4)	(RO access)
24h	Samples of VDM Instances 1-64 (Group 1)	Every 2-byte value is a (possibly frozen) sample of the observable monitored by a VDM instance.
25h	Samples of VDM Instances 65-128 (Group 2)	
26h	Samples of VDM Instances 129-192 (Group 3)	
27h	Samples of VDM Instances 193-256 (Group 4)	(RO access)

Page	Subject Area	Description
28h	Thresholds 1-16 (Group 1)	Every set of four 2-byte values describes a threshold set of a threshold crossing detector, possibly shared by several VDM instances.
29h	Thresholds 17-32 (Group 2)	
2Ah	Thresholds 33-48 (Group 3)	
2Bh	Thresholds 49-64 (Group 4)	The association of these threshold sets and VDM instances is described in the VDM descriptors. (RO access)
2Ch	VDM Flags (Groups 1-4)	Every Byte contains the latched threshold crossing Flags of 2 VDM instances (4 bits each) (RO/COR access)
2Dh	VDM Masks (Groups 1-4)	Every Byte contains the Masks for the threshold crossing Flags of 2 VDM instances (4 bits each) (RW access) The default of each Mask is 1 (masked)
2Eh	-	Reserved
2Fh	Advertisement and Dynamic Controls	VDM support details and dynamic controls (Mixed RO and RW access)

8.22.1 Pages 20h-23h (VDM Descriptor Groups Pages)

There can be up to 256 VDM instances to monitor up to 256 VDM observables.

The module describes the list of VDM instances it supports in VDM instance descriptor registers.

Each Page contains one of four VDM instance groups.

Table 8-168 VDM Configuration (Page 20h-23h)

Page	Byte	Register Name	Description	Type
20h	128-129	VDMDescriptor1	Descriptor register for VDM instance 1, in group 1	RO
	130-131	VDMDescriptor2	Descriptor register for VDM instance 2, in group 1	RO
	132-251	...	Descriptor registers for VDM instance 3 to 62 , in group 1	RO
	252-253	VDMDescriptor63	Descriptor register for VDM instance 63, in group 1	RO
	254-255	VDMDescriptor64	Descriptor register for VDM instance 64, in group 1	RO
21h	128-129	VDMDescriptor65	Descriptor register for VDM instance 65, in group 2	RO
	130-131	VDMDescriptor66	Descriptor register for VDM instance 66, in group 2	RO
	132-251	...	Descriptor registers for VDM instance 67 to 126 , in group 2	RO
	252-253	VDMDescriptor127	Descriptor register for VDM instance 127, in group 2	RO
	254-255	VDMDescriptor128	Descriptor register for VDM instance 128, in group 2	RO
22h	128-129	VDMDescriptor129	Descriptor register for VDM instance 129, in group 3	RO
	130-131	VDMDescriptor130	Descriptor register for VDM instance 130, in group 3	RO
	132-251	...	Descriptor registers for VDM instance 131 to 190 , in group 3	RO
	252-253	VDMDescriptor191	Descriptor register for VDM instance 191, in group 3	RO
	254-255	VDMDescriptor192	Descriptor register for VDM instance 192, in group 3	RO
23h	128-129	VDMDescriptor193	Descriptor register for VDM instance 193, in group 4	RO
	130-131	VDMDescriptor194	Descriptor register for VDM instance 194, in group 4	RO
	132-251	...	Descriptor registers for VDM instance 195 to 254 , in group 4	RO
	252-253	VDMDescriptor255	Descriptor register for VDM instance 255, in group 4	RO
	254-255	VDMDescriptor256	Descriptor register for VDM instance 256, in group 4	RO

8.22.1.1 VDM Instance Descriptors

The module describes each available VDM instance by a two-byte descriptor field.

Table 8-169 Definition of 2-byte VDM Instance Descriptor

Byte	Bits	Field Name and Description
Even Address	7-4	LocalThresholdSetID This number determines which threshold set will be used for this observable. The threshold set is in the same group as the observable descriptor. The global ThresholdSetID (1-64) is defined as follows: Page 20h: (group 1): ThresholdSetID = 1 + LocalThresholdSetID Page 21h: (group 2): ThresholdSetID = 17 + LocalThresholdSetID Page 22h: (group 3): ThresholdSetID = 33 + LocalThresholdSetID Page 23h: (group 4): ThresholdSetID = 49 + LocalThresholdSetID
	3-0	Monitored Resource 0: Lane 1 or Data Path starting on lane 1 1: Lane 2 or Data Path starting on lane 2 2: Lane 3 or Data Path starting on lane 3 3: Lane 4 or Data Path starting on lane 4 4: Lane 5 or Data Path starting on lane 5 5: Lane 6 or Data Path starting on lane 6 6: Lane 7 or Data Path starting on lane 7 7: Lane 8 or Data Path starting on lane 8 15: Module (not associated with a lane or Data Path)
Odd Address	7-0	Observable Type (see Table 8-170 for observables and encodings)

8.22.1.2 VDM Observable Types

The list of defined VDM observable types, an associated Type ID, and the data type of samples of the observable type is defined in Table 8-170.

Observables related to the entire module are presented on Bank index 0 zero only.

All VDM data types are **Big Endian**, unlike *the regular Module Diagnostics observables (on Page 13h-14h)*.

Table 8-170 VDM Observable Types (Type Coding)

Type ID	Observable Type	Instance Type	Data Type	Unit Scale	Unit
0	Not Used indicator ¹	N/A	N/A		
1	Laser Age (0% at BOL, 100% EOL) (Media Lane)	Basic	U16	1	%
2	TEC Current (Module)	Basic	S16	100/32767	%
3	Laser Frequency Error (Media Lane)	Basic	S16	10	MHz
4	Laser Temperature (Media Lane)	Basic	S16	1/256	C
5	SNR (dB) Media Input (Media Lane) -- see section 7.1.4	Basic	U16	1/256	dB
6	SNR (dB) Host Input (Lane) -- see section 7.1.4	Basic	U16	1/256	dB
7	PAM4 Level Transition Parameter Media Input (Media Lane)	Basic	U16	1/256	dB
8	PAM4 Level Transition Parameter Host Input (Lane)	Basic	U16	1/256	dB
9	Pre-FEC BER Minimum Sample Media Input (Data Path)	Statistic	F16	N/A	
10	Pre-FEC BER Minimum Sample Host Input (Data Path)	Statistic	F16	N/A	
11	Pre-FEC BER Maximum Sample Media Input (Data Path)	Statistic	F16	N/A	
12	Pre-FEC BER Maximum Sample Host Input (Data Path)	Statistic	F16	N/A	
13	Pre-FEC BER Sample Average Media Input (Data Path)	Statistic	F16	N/A	
14	Pre-FEC BER Sample Average Host Input (Data Path)	Statistic	F16	N/A	
15	Pre-FEC BER Current Sample Media Input (Data Path)	Basic	F16	N/A	
16	Pre-FEC BER Current Sample Host Input (Data Path)	Basic	F16	N/A	
17	FERC Minimum Sample Value Media Input (Data Path)	Statistic	F16	N/A	
18	FERC Minimum Sample Value Host Input (Data Path)	Statistic	F16	N/A	
19	FERC Maximum Sample Value Media Input (Data Path)	Statistic	F16	N/A	
20	FERC Maximum Sample Value Host Input (Data Path)	Statistic	F16	N/A	
21	FERC Sample Average Value Media Input (Data Path)	Statistic	F16	N/A	
22	FERC Sample Average Value Host Input (Data Path)	Statistic	F16	N/A	
23	FERC Current Sample Value Media Input (Data Path)	Basic	F16	N/A	
24	FERC Current Sample Value Host Input (Data Path)	Basic	F16	N/A	
25	FERC Total Accumulated Media Input (Data Path)	Statistic	F16	N/A	
26	FERC Total Accumulated Host Input (Data Path)	Statistic	F16	N/A	
27	SEWmax Minimum Sample Value Media Input (Data Path)	Statistic	U16	N/A	
28	SEWmax Minimum Sample Value Host Input (Data Path)	Statistic	U16	N/A	
29	SEWmax Maximum Sample Value Media Input (Data Path)	Statistic	U16	N/A	
30	SEWmax Maximum Sample Value Host Input (Data Path)	Statistic	U16	N/A	
31	SEWmax Sample Average Value Media Input (Data Path) ¹	Statistic	U16	N/A	
32	SEWmax Sample Average Value Host Input (Data Path)	Statistic	U16	N/A	
33	SEWmax Current Sample Value Media Input (Data Path)	Basic	U16	N/A	
34	SEWmax Current Sample Value Host Input (Data Path)	Basic	U16	N/A	
35-76	-	Reserved [42]			
77	Vcc2p6 Voltage Monitor (Module)	Basic	U16	100	uV
78	Vcc1p8 Voltage Monitor (Module)	Basic	U16	100	uV
79	Vcc1p2 Voltage Monitor (Module)	Basic	U16	100	uV
80	Vcc0p9 Voltage Monitor (Module)	Basic	U16	100	uV
81	Vcc0p7A Voltage Monitor (Module)	Basic	U16	100	uV
82	Vcc0p7B Voltage Monitor (Module)	Basic	U16	100	uV
83	Vcc12 Voltage Monitor (Module)	Basic	U16	250	uV

¹ A non-integer SEWmax average should be rounded up to the next integer.

Type ID	Observable Type	Instance Type	Data Type	Unit Scale	Unit
84	ELS Input Power (Lane=Laser ID) ¹	Basic	S16	0.01	dBm
85-99	-	Reserved for CPO Observables [15]			
100-127	-	Custom Observables [28]			
128-255	-	Restricted OIF [128]			

¹ Not Used means that the module does not present data for the VDM Instance described by the descriptor

8.22.2 Pages 24h-27h (VDM Sample Groups Pages)

There are up to 256 read-only sample registers (real time value registers) for current values of basic observables or statistics associated with up to 256 VDM instances.

All sample values are stored in **BigEndian** order (MSB in lower address).

The observable actually reported in the VDM< i >Sample register of VDM instance < i > is defined in the descriptor VDM< i >Descriptor of that VDM instance in Pages 20h-23h.

Table 8-171 VDM Real-Time Values (Page 24h-27h)

Page	Byte	Register Name	Register Description	Type
24h	128-129	VDMsample1	X16 Real-time value of VDM instance 1 in group 1	RO
	130-131	VDMsample2	X16 Real-time value for VDM instance 2 in group 1	RO
	132-251	VDM 3..62	X16 Real-time values for VDM instance 3 to 62 in group 1	RO
	252-253	VDMsample63	X16 Real-time value for VDM instance 63 in group 1	RO
	254-255	VDMsample64	X16 Real-time value for VDM instance 64 in group 1	RO
25h	128-129	VDMsample65	X16 Real-time value of VDM instance 65 in group 2	RO
	130-131	VDMsample66	X16 Real-time value for VDM instance 66 in group 2	RO
	132-251	VDM 67..126	X16 Real-time values for VDM instance 67 to 126 in group 2	RO
	252-253	VDMsample127	X16 Real-time value for VDM instance 127 in group 2	RO
	254-255	VDMsample128	X16 Real-time value for VDM instance 128 in group 2	RO
26h	128-129	VDMsample129	X16 Real-time value of VDM instance 129 in group 3	RO
	130-131	VDMsample130	X16 Real-time value for VDM instance 130 in group 3	RO
	132-251	VDM 131..190	X16 Real-time values for VDM instance 131 to 190 in group 3	RO
	252-253	VDMsample191	X16 Real-time value for VDM instance 191 in group 3	RO
	254-255	VDMsample192	X16 Real-time value for VDM instance 192 in group 3	RO
27h	128-129	VDMsample193	X16 Real-time value of VDM instance 193 in group 4	RO
	130-131	VDMsample194	X16 Real-time value for VDM instance 194 in group 4	RO
	132-251	VDM 195..254	X16 Real-time values for VDM instance 195 to 254 in group 4	RO
	252-253	VDMsample255	X16 Real-time value for VDM instance 255 in group 4	RO
	254-255	VDMsample256	X16 Real-time value for VDM instance 256 in group 4	RO

10

¹ The numbering scheme for lasers in a CPO module is defined elsewhere.

8.22.3 Pages 28h-2Bh (VDM Threshold Set Groups Pages)

Pages 28h-2Bh contain a read only array of threshold sets, stored in groups of 16 threshold sets per Page.

Each used threshold set (quad) contains four threshold values for threshold crossing supervision at high/low alarm/warning thresholds.

The actual threshold values can depend on the commissioned set of Applications and therefore may change whenever a new Application is commissioned.

Note: The module will update threshold values only for that reason. For Data Path Applications, the thresholds will be updated when the relevant Data Path reaches DPInitialized. For Network Path Applications, the Media Side thresholds will be updated when the relevant Network Path reaches NPInitialized, and the Host side thresholds will be updated when the relevant Host Path reaches DPInitialized.

Note: It is recommended that hosts interested in the thresholds read them after Data Path Initialization and/or Network Path Initialization.

The usage of a threshold sets for one or more observables is defined in the VDM instance descriptors of those observables on Pages 20h-23h.

Table 8-172 VDM Alarm/Warning Thresholds (Page 28h-2Bh)

Page	Byte	Register Name	Register Description	Type
28h	128-129	HighAlarmThreshold1	X16 Thresholds for group 1 monitored observables. Threshold type and storage same as supervised real-time value.	RO
	130-131	LowAlarmThreshold1		RO
	132-133	HighWarningThreshold1		RO
	134-135	LowWarningThreshold1		RO
	136-137	HighAlarmThreshold2		RO
	138-139	LowAlarmThreshold2		RO
	140-141	HighWarningThreshold2		RO
	142-143	LowWarningThreshold2		RO
	144-247	... 3 to 15		RO
	248-249	HighAlarmThreshold16		RO
	250-251	LowAlarmThreshold16		RO
	252-253	HighWarningThreshold16		RO
	254-255	LowWarningThreshold16		RO
29h	128-129	HighAlarmThreshold17	X16 Thresholds for group 2 monitored observables. Threshold type and storage same as supervised real-time value.	RO
	130-131	LowAlarmThreshold17		RO
	132-133	HighWarningThreshold17		RO
	134-135	LowWarningThreshold17		RO
	136-137	HighAlarmThreshold18		RO
	138-139	LowAlarmThreshold18		RO
	140-141	HighWarningThreshold18		RO
	142-143	LowWarningThreshold18		RO
	144-247	... 19 to 31		RO
	248-249	HighAlarmThreshold32		RO
	250-251	LowAlarmThreshold32		RO
	252-253	HighWarningThreshold32		RO
	254-255	LowWarningThreshold32		RO
2Ah	128-129	HighAlarmThreshold33	X16 Thresholds for group 3 monitored observables. Threshold type and storage same as supervised real-time value.	RO
	130-131	LowAlarmThreshold33		RO
	132-133	HighWarningThreshold33		RO
	134-135	LowWarningThreshold33		RO
	136-247	... 34 to 47		RO
	248-249	HighAlarmThreshold48		RO
	250-251	LowAlarmThreshold48		RO
	252-253	HighWarningThreshold48		RO
	254-255	LowWarningThreshold48		RO
2Bh	128-129	HighAlarmThreshold49	X16 Thresholds for group 4 monitored observables.	RO
	130-131	LowAlarmThreshold49		RO
	132-133	HighWarningThreshold49		RO

Page	Byte	Register Name	Register Description	Type
	134-135	LowWarningThreshold49	Threshold type and storage same as supervised real-time value.	RO
	144-247	... 50 to 63		RO
	248-249	HighAlarmThreshold64		RO
	250-251	LowAlarmThreshold64		RO
	252-253	HighWarningThreshold64		RO
	254-255	LowWarningThreshold64		RO

1

2 8.22.4 Page 2Ch (VDM Flags Page)

3 Each VDM instance has four Flags for threshold crossing events: alarms and warnings for crossing high and low
4 thresholds.

5 Page 2Ch contains the Flags for these threshold crossing alarms and warnings, for all active VDM instances.

6 **Table 8-173 VDM Threshold Crossing (TC) Flags Byte**

Byte	Bit	Field Name	Field Description	Type
<i>Any</i>	0 or 4	HighTCAlarmFlag	Latched when monitored observable > high alarm threshold	RO/COR
	1 or 5	LowTCAlarmFlag	Latched when monitored observable < low alarm threshold	RO/COR
	2 or 6	HighTCWarningFlag	Latched when monitored observable > high warning threshold	RO/COR
	3 or 7	LowTCWarningFlag	Latched when monitored observable < low warning threshold	RO/COR

7

8 **Table 8-174 VDM Alarm and Warning Configuration (Page 2Ch)**

Bytes	Bits	Field Name	Field Description	Type
128	7-4	VDMFlags2	High/Low Alarm/Warning TC Flags for VDM instance 2, group 1	RO/COR
	3-0	VDMFlags1	High/Low Alarm/Warning TC Flags for VDM instance 1, group 1	RO/COR
129-158	7-0	VDMFlags3-62	High/Low Alarm/Warning TC Flags as described in Table 8-172, for VDM instances 3-62, group 1	RO/COR
159	7-4	VDMFlags64	High/Low Alarm/Warning TC Flags for VDM instance 64, group 1	RO/COR
	3-0	VDMFlags63	High/Low Alarm/Warning TC Flags for VDM instance 63, group 1	RO/COR
160	7-4	VDMFlags66	High/Low Alarm/Warning TC Flags for VDM instance 66, group 2	RO/COR
	3-0	VDMFlags65	High/Low Alarm/Warning TC Flags for VDM instance 65, group 2	RO/COR
161-190	7-0	VDMFlags 67-126	High/Low Alarm/Warning TC Flags as described in Table 8-172, for VDM instances 67-126, group 2	RO/COR
191	7-4	VDMFlags128	High/Low Alarm/Warning TC Flags for VDM instance 128, group 2	RO/COR
	3-0	VDMFlags127	High/Low Alarm/Warning TC Flags for VDM instance 127, group 2	RO/COR
192	7-4	VDMFlags130	High/Low Alarm/Warning TC Flags for VDM instance 130, group 3	RO/COR
	3-0	VDMFlags129	High/Low Alarm/Warning TC Flags for VDM instance 129, group 3	RO/COR
193-222	7-0	VDMFlags 131-190	High/Low Alarm/Warning TC Flags as described in Table 8-172, for VDM instances 131-190, group 3	RO/COR
223	7-4	VDMFlags192	High/Low Alarm/Warning TC Flags for VDM instance 192, group 3	RO/COR
	3-0	VDMFlags191	High/Low Alarm/Warning TC Flags for VDM instance 191, group 3	RO/COR
224	7-4	VDMFlags194	High/Low Alarm/Warning TC Flags for VDM instance 194, group 4	RO/COR
	3-0	VDMFlags193	High/Low Alarm/Warning TC Flags for VDM instance 193, group 4	RO/COR
224-254	7-0	VDMFlags 195-254	High/Low Alarm/Warning Flags as described in Table 8-172, for VDM instances 195-254, group 4,	RO/COR
255	7-4	VDMFlags256	High/Low Alarm/Warning TC Flags for VDM instance 256, group 4	RO/COR
	3-0	VDMFlags255	High/Low Alarm/Warning TC Flags for VDM instance 255, group 4	RO/COR

9

8.22.5 Page 2Dh (VDM Masks Page)

1 Each VDM instance can set Flags as described in Page 2Ch.

2 Each set Flag asserts the Interrupt request signal unless the corresponding Mask bit is set.

3 The default value for all VDM Mask bits is 1 (masked).

4 Page 2Dh defines a writeable Page of alarm Mask bits for the corresponding Flags defined in Page 2Ch.

5 **Table 8-175 VDM ThresholdSet0 to 15 Alarm and Warning Configuration (Page 2Dh)**

Byte	Bit	Field Name	Field Description	Type
128	7-4	VDM Masks2	High/Low Alarm/Warning Masks for VDM instance 2, group 1	RW
	3-0	VDM Masks1	High/Low Alarm/Warning Masks for VDM instance 1, group 1	RW
129-158	7-4	VDM Masks3-62	High/Low Alarm/Warning Masks as described in Table 8-172, for VDM instances 3-62, group 1	RW
	3-0			
159	7-4	VDM Masks64	High/Low Alarm/Warning Masks for VDM instance 64, group 1	RW
	3-0	VDM Masks63	High/Low Alarm/Warning Masks for VDM instance 63, group 1	RW
160	7-4	VDM Masks66	High/Low Alarm/Warning Masks for VDM instance 66, group 2	RW
	3-0	VDM Masks65	High/Low Alarm/Warning Masks for VDM instance 65, group 2	RW
161-190	7-4	VDM Masks67-126	High/Low Alarm/Warning Masks as described in Table 8-172, for VDM instances 67-126, group 2	RW
	3-0			
191	7-4	VDM Masks128	High/Low Alarm/Warning Masks for VDM instance 128, group 2	RW
	3-0	VDM Masks127	High/Low Alarm/Warning Masks for VDM instance 127, group 2	RW
192	7-4	VDM Masks130	High/Low Alarm/Warning Masks for VDM instance 130, group 3	RW
	3-0	VDM Masks129	High/Low Alarm/Warning Masks for VDM instance 129, group 3	RW
193-222	7-4	VDM Masks131-190	High/Low Alarm/Warning Masks as described in Table 8-172, for VDM instances 131-190, group 3	RW
	3-0			
223	7-4	VDM Masks192	High/Low Alarm/Warning Masks for VDM instance 192, group 3	RW
	3-0	VDM Masks191	High/Low Alarm/Warning Masks for VDM instance 191, group 3	RW
224	7-4	VDM Masks194	High/Low Alarm/Warning Masks for VDM instance 194, group 4	RW
	3-0	VDM Masks193	High/Low Alarm/Warning Masks for VDM instance 193, group 4	RW
224-254	7-4	VDM Masks195-254	High/Low Alarm/Warning Masks as described in Table 8-172, for VDM instances 195-254, group 4,	RW
	3-0			
255	7-4	VDM Masks256	High/Low Alarm/Warning Masks for VDM instance 256, group 4	RW
	3-0	VDM Masks255	High/Low Alarm/Warning Masks for VDM instance 255, group 4	RW

8.22.6 Page 2Fh (VDM Advertisement and Dynamic Controls)

Page 2Fh is used for two purposes

- **advertisement** of the number of supported VDM groups (i.e. availability of VDM Pages)
- **dynamic control** (mainly) for freezing and unfreezing statistics reporting registers

VDM Groups Advertisement

The VDMSSupport Byte 2Fh:128 advertises how many VDM instance groups (of 64 VDM instances each) are supported, i.e. which of the Pages 20h-2Bh, and which parts of Pages 2Ch and 2Dh are available.

The individual VDM instances that the module actually provides to report samples of basic observables or of statistics of basic observables are described in the VDM Descriptor pages (see section 8.22.1).

Freezing and Unfreezing Statistics Reporting Registers

To ease the following description, the sample registers of VDM instances (in Pages 24h-27h) that report a statistic (minimum, maximum, or average) of a basic observable are called **statistics reporting registers**.

Note: Please see the beginning of section 8.19 for an introduction to statistics support in VDM

During statistics collection, a module always updates internal statistics variables (e.g. counters) but it updates the associated statistics reporting registers only when the statistics reporting is not frozen as described below.

Globally freezing the statistics reporting implies that the module terminates the current statistics collection interval (providing the results in the now frozen statistics reporting registers) and atomically starts a new statistics collection interval internally (without losing any sample of the basic observables).

Note: Raising the FreezeRequest bit effectively defines a host-controlled Performance Monitoring Interval.

Note: The initial behavior of statistics collection is not specified; it is unknown if statistics collection is initially halted or running. To achieve a well-defined starting point, hosts are advised to initially raise a FreezeRequest and then to start statistics collection at a suitable point in time after the module confirmed that the Freeze has been executed.

Dynamic Control of Statistics Intervals and Freezing

The host raises the FreezeRequest Bit 2Fh:144.7 to request that the module shall stop updating the statistics reporting registers (in order to allow consistent results readout) and at the same to start a new statistics interval.

When finished reading, the host may clear the FreezeRequest bit to signal that the module can resume updating the Interval Statistics registers. As a preparation for requesting the next freeze, the Host must eventually clear the FreezeRequest bit.

When freezing the reporting registers, the module starts a new statistics interval and ensures that no sample of any observable is lost. To this end the module atomically resets and restarts its internal statistics variables when it performs the last update on a then frozen register and then just stops updating the statistics reporting registers until the host releases the reporting registers again by terminating the freeze period.

Scenario

A scenario of host-module interactions for (repeatedly) obtaining statistics results over contiguous statistics intervals defined by the host (i.e. over host-defined Performance Monitoring Intervals) is as follows:

1. Host ceases a **FreezeRequest** (if still set) in order to request unfreezing the statistics reporting registers
2. Module clears **UnfreezeDone** (within **tVDMF** time) to prepare for a later Unfreeze Done indication
3. Module copies current values of its internally collected statistics (internally collected during the Freeze) to the statistics reporting registers and then raises **UnfreezeDone**
4. Host may now read live statistics from the statistics reporting registers to get incrementally updated statistics results. The module now continues to update the statistics reporting registers.
5. To request termination of the current statistics interval and prepare for statistics results readout, the host raises **FreezeRequest**
6. Module clears **FreezeDone** (within **tVDMF** time) to prepare for a later FreezeDone indication
7. Module stops updating statistics reporting registers after freeze, resets its internal statistics variable, and immediately continues to update the internal statistics variables (just without copying updates to the frozen statistics reporting registers) and then sets **FreezeDone**
8. The statistics reporting registers are now frozen and the host is free to read the frozen statistics reporting registers with the results of the statistics interval just ended.
9. While the statistics results registers are frozen, the module updates its internal statistics variables, but does not update the statistics reporting registers.
10. Host may cease the still active **FreezeRequest** at any time before ending the current statistics interval,

either immediately after having read the statistics of the previous interval, or just prior to ending the current interval, and the scenario repeats.

Power Saving

A power saving mode (PowerSavingMode) can be supported by a module, where the monitoring of observables and derived functions (e.g. alarm threshold detection) are suspended for power saving.

Note: This feature is intended for modules with a tight power envelope such that monitoring should only be enabled transiently for diagnostics. It is not expected to be used otherwise.

When the host makes changes to the PowerSavingMode the following procedures should be employed to ensure the module reports consistent information, and transient statistics are ignored:

To enable PowerSavingMode:

1. Host issues a FreezeRequest
2. Host waits for the module to acknowledge
3. Host changes the PowerSavingMode to Enable

To disable PowerSavingsMode, if desired, and assuming FreezeRequest from above is still engaged:

1. Host changes the PowerSavingMode to Disabled
2. Host issues UnFreezeRequest
3. Host waits for the module to acknowledge
4. Host issues FreezeRequest
5. Host waits for the module to acknowledge
6. Host continues as normal

Table 8-176 VDM Advertisement and Control Registers Summary (Page 2Fh)

Byte	Bit	Field Name	Field Description	Type
128	7-3	-	Reserved	RO
	2	PowerSavingSupport	1b: PowerSavingMode is supported	RO
	1-0	VDMSupport	Advertisement 0: Group 1 (Page 20h, 24h, 28h, first 1/4 of 2Ch, 2Dh) 1: Groups 1-2 (Page 20h-21h, 24h-25h, 28h-29h, first 1/2 of 2Ch, 2Dh) 2: Groups 1-3 (Page 20h-22h, 24h-26h, 28h-2Ah, first 3/4 of 2Ch, 2Dh) 3: Groups 1-4 (Page 20h-23h, 24h-27h, 28h-2Bh, 2Ch, 2Dh)	RO
129-130	7-0	FineIntervalLength	U16 Length of fine interval (measurement time for one sample) used for internal BER/FERC monitoring, in units of 0.1 ms	RO
131-143	7-0	-	Reserved[13]	RO
144	7	FreezeRequest	When raised by the host, causes the module to freeze and hold all reported statistics reporting registers (minimum, maximum and average values) in Pages 24h-27h. When ceased by the host, releases the freeze request, allowing the reported minimum, maximum and average values to update again. Multi-bank behavior: Freezing or unfreezing in one supported bank occurs in all supported banks synchronously, independent of whether BankBroadcastEnable is supported and enabled. <i>Note. The module internally records new statistics while the reporting registers are frozen, so that no data is lost.</i>	RW
	6	PowerSavingMode	0b: VDM monitoring functions are enabled (ON) 1b: VDM monitoring functions are disabled (OFF) for power saving Advertisement: 2Fh:128.2 (PowerSavingSupport)	RW Adv.
	5-0	-	Reserved	RO
145	7	FreezeDone	Raised by the module when it has finished freezing the reporting registers (such that the host can now safely read) and simultaneously restarted all supported statistics internally. Ceased by the module (within T_{VDMF} time from when Freeze Request is raised by the host). The FreezeDone status in any of the supported banks reflects the completion of a FreezeRequest across all supported banks. Until a first FreezeRequest is served, the module reports 0b.	RO

	6	UnfreezeDone	Raised by the module when real-time updates of the reporting registers are performed again Ceased by the module (within T_{VDMF} time from when Freeze Request is ceased by the host). The UnfreezeDone status in any of the supported banks reflects the completion of unfreezing across all supported banks. Until a first request to unfreeze is served, the module reports 0b.	RO
	5-0	-	Reserved	RO
146-255	7-0	-	Reserved[110]	RO

8.23 Banked Page 9Fh (CDB Message)

Page 9Fh is the main Page of the optional Command Data Block (CDB) messaging feature (see section 7.2).

The module advertises support of Page 9Fh in the **CdbInstancesSupported** field (see Table 8-54).

Each supported Bank of Page 9Fh (with the same Bank supported for the supported subset of pages A0h-AFh) implements one **CDB instance** as a separate memory mapped command/reply message exchange facility

Note: Concurrent messaging over different CDB instances is possible only when the module supports background mode message processing.

Table 8-177 Page 9Fh Overview (CDB Message)

Byte	Size (Bytes)	Subject Area	Description
128 – 133	6	Command Message Header	Host-written command message header
134 – 135	2	Reply Message Header	Module-written header additions in reply message
136 – 255	120	Local Payload (LPL)	Area to store command and/or reply message body (or parts)

The **CDB Message Page 9Fh** contains fields for the **message header** and for (parts of) the **message body**, both for **command messages** (composed by host) and for **reply messages** (composed by module).

The **CDB Command Message Header** Bytes (9Fh:128-133) described in Table 8-178 are filled by the host to specify the CDB command to be executed by the module, the length of the used local command payload (LPLLength), of the used extended command payload (EPLLength), and a Message Checksum (CdbChkCode). These fields are filled by the host when composing the command message and evaluated by the module when receiving the command message. The module does not change the command header fields in a reply.

Table 8-178 CDB Command Message Header (Page 9Fh)

Byte	Bits	Register Name	Register Description	Type
128-129	7-0	CMDID	U16 CDB Command Code (CMDID) identifies a CDB command to be executed and writing this field also "sends" the CMD message from host to module for processing. <i>Note: See Table 9-7 for a list of commands and their CMD IDs.</i> <i>Note: The module advertises the actual trigger condition for the "message transfer". Ignoring some details, it may either be a write to 9Fh:128-129 or 9Fh:129, or it may be end of a multibyte write that includes Byte 9Fh:129. The detailed specification is found in the main text.</i> <i>Note: This means that the host must compose header and body either before writing to Byte 9Fh:129, or during a multi-byte write operation that includes Byte 9Fh:129.</i>	RW Rqd.
130-131	7-0	EPLLength	U16 Extended Payload Length (EPLLength) specifies the host-written number of command message body bytes in EPL, in Pages A0h-AFh. Valid lengths are 0-2048. <i>Note: The EPLLength field is included in the calculation of the CdbChkCode field value.</i>	RW Cnd.
132	7-0	LPLLength	U8 Local Payload Length (LPLLength) specifies the host-written number of command message body bytes in LPL, on this Page 9Fh. Valid lengths are 0-- 120 <i>Note: The LPLLength field is included in the calculation of the CdbChkCode field value.</i>	RW Rqd.
133	7-0	CdbChkCode	U8 CDB Check Code (CdbChkCode) is computed by the host as the one's complement of the arithmetic sum of Bytes 9Fh:128 to 9Fh:(136+LPLLength-1) excluding Bytes 9Fh:133-135 . Integrity checks of EPL Pages (if any) are defined individually for each CMD as defined in chapter 9. <i>Note: Neither the reply header fields 9Fh:134-135 nor any EPL data used by the command are included in the check computation.</i> <i>Note: Modules can compare a correctly recomputed CdbChkCode against the received value to check message integrity.</i>	RW Rqd.

The **CDB Reply Message Header** Bytes (9Fh:134-135) described in Table 8-179 are filled by the module to provide information about the **reply message body** (if any) carried in one or both of the payload areas, LPL or EPL. These reply header fields are ignored by the module when receiving a command; instead they are determined and filled by the module when composing the reply.

Note: To verify later that the module has actually filled the reply header fields properly, the host may set all reply header fields to values not expected in a reply, when composing the command message.

Table 8-179 CDB Reply Message Header (Page 9Fh)

Byte	Bits	Register Name	Register Description	Type
134	7-0	RPLLength	<p>U8 REPLY Payload Length (RPLLength) is computed by the module and encodes the length of REPLY data returned:</p> <p>0-120: Number of LPL bytes used in the LPL area</p> <p>240≤i≤255: Number of EPL Pages used, encoded as i-239 (Pages A0 to A0h+i-240)</p> <p>121-239: Reserved</p> <p>The exact number of bytes returned in EPL Pages is determined in a CMDID specific way.</p> <p>RPLLength is computed in a CMDID specific way when data are returned in both LPL and EPL areas.</p> <p><i>Note: This REPLY header field is not included in the CdbChkCode computation. Hosts may want to set this byte to a suitable value when composing a CMD, in order to later check if the module has actually set this field in the REPLY.</i></p>	RW Cnd.
135	7-0	RPLChkCode	<p>U8 REPLY Payload Check Code (RPLChkCode) is computed by the module as follows:</p> <p>No REPLY message body: RPLChkCode = 0</p> <p>REPLY message body in Local Payload: RPLChkCode is the one's complement of the sum of LPL Bytes 136 to (136+RPLLength-1).</p> <p>REPLY message body in Extended Payload: RPLChkCode is the one's complement of the sum of all used bytes in EPL Pages A0h to (A0h+RPLLength-240).</p> <p>REPLY message body in LPL and EPL: RPLChkCode is computed as specified for the particular CMD.</p> <p><i>Note: This REPLY header field is not included in the CdbChkCode computation. Hosts may want to set this byte to a suitable value when composing a CMD, in order to later check if the module has actually set this field in the REPLY.</i></p>	RW Rqd.

The **Local Payload (LPL)** area 9Fh:136-255 described in Table 8-180 can carry a message body in the same Page as the message header, for message body lengths not exceeding 120 bytes.

The LPL area is used both for CMD arguments and for REPLY data (i.e. REPLY LPL may overwrite CMD LPL).

*Note: The **CDB EPL Pages** range (A0h-AFh) is called **Extended Payload (EPL)** and allows for large message body lengths of up to 2048 bytes, as described in the next section.*

Table 8-180 CDB Message Body (Page 9Fh)

Byte	Bits	Register Name	Register Description	Type
136-255	7-0	LPL	Local Payload (LPL): Message body area sufficient for lengths not exceeding 120 bytes for host-written CMD data or module-written REPLY data (possibly overwriting CMD data) as specified individually for each CDB Command in chapter 9.	RW Cnd.

Message Exchange Protocol

The host must write the length, Message Checksum, payload (if applicable), and CMDID.

A host WRITE either ending at or including the LSB of the CMDID field "sends the message" and triggers the module to execute the command.

Note: If the host writes a new CMDID before a current CDB command completes the behavior is unpredictable.

While processing a CDB command, the module updates command execution status in **CdbStatus** 00h:38 or 00h:37 and signals completion in the CDB **CMD Completion Flags** 00h:8.7 or 00h:8.6.

Note: The command message header fields EPLLength and LPLLength are always included in the command message checksum CdbChkCode (even when they are not used), while the reply header fields RPLLength and RPLChkCode are never included.

Editor's Note: In future revisions of this specification, the following text should be consolidated.

Triggering CDB Command Execution

The event triggering the module to start CDB command processing depends on the advertising bit 01h:165.7.

Assuming that the current Page is 9Fh, CDB processing in the module is triggered as follows:

- When 01h:165.7=0b: when the I2CMCI STOP condition of the I2CMCI transaction for a two-byte or one-byte WRITE ending at Byte 9Fh:129 is received
- When 01h:165.7=1b: when the I2CMCI STOP condition of the I2CMCI transaction for a multi-byte WRITE that includes Byte 9Fh:129 is received

Note: In the first method, at least two WRITES are required to invoke a CDB command, unless the host writes all 128 Bytes of Page 9Fh with wrapping. In the second method a single WRITE may be sufficient to invoke a CDB command represented fully in Page 9Fh.

Note: The host must use the advertised method.

Note: The maximum number of bytes per single ACCESS for CDB is advertised in Page 01h Byte 164.

The LPLLength field and/or the EPLLength field is populated by the host and is CDB command dependent.

A command may use either LPL or EPL or both LPL and EPL.

The length of the EPL is specified in the EPLLength field. Any checks to ensure validity of EPL data are not defined generally. If additional checks are required, they are to be specified independently for each type or group of command codes.

For example, firmware download assumes that the CDB data written in the EPL block come directly from a binary byte stream representing the vendor provided file and it is up to the vendor to add any additional checks to any data embedded within the file.

8.23.1 Triggering CDB Command on WRITE ending at 9Fh:129

The module advertises 01h:165.7=0b when CDB command execution is triggered by WRITE to Byte 9Fh:129 or to 9Fh:128-129.

Note: In this case the recommended method to invoke a CDB command with LPL consists of at least two WRITES:

1. Write command parameters to bytes 9Fh:130-(135+LPLLength).
 - a. These writes may have to be broken into multiple WRITE accesses.
 - b. The maximum allowed number bytes per WRITE is advertised in 01h:164.
2. Write the 16-bit CDBID field one two-byte WRITE to 9Fh:128-129 or in two one-byte WRITES to 9Fh:128 followed by 9Fh:129

Note: The outcome of writing bytes 9Fh:128-(135+LPLLength) in one WRITE are unpredictable in this mode.

8.23.2 Triggering CDB Command on WRITE including 9Fh:129

The module advertises 01h:165.7=1b when CDB command execution is triggered by the STOP condition of an MCI transaction for a WRITE that includes Byte 129, i.e. after the MCI write transaction has been completed.

Note: With this option the recommended method to invoke a CDB command with LPL involves one WRITE to 9Fh:128-(135+LPLLength) if the maximum WRITE length in CDB pages allows.

Note: Although Byte 129 is sent as the 2nd byte, the command is not processed until the STOP bit is received.

1 *Note: If the maximum bytes per WRITE is at least 8, then commands without LPL (LPLLength of 0) can be*
2 *invoked with a single WRITE.*

3 *Note: If the maximum bytes per WRITE is 128, then any CDB command that uses LPL only may be triggered in*
4 *one WRITE.*

5 *Note: The host may still use multiple WRITES or byte by byte WRITES, as a single byte write to Byte 129 ends*
6 *with a STOP condition and will also invoke the CDB command.*

8.24 Banked Pages Range A0h-AFh (CDB Extended Payload Pages)

Pages A0h-AFh contain the optional **Extended Payload (EPL)** area for the optional Command Data Block (CDB) messaging feature (see section 7.2), providing space for long message body data beyond the capacity of the Local Payload (LPL) area on **CDB Message Page** 9Fh (see previous section).

Each Bank of Pages A0h-AFh refers to one CDB messaging instance.

The actually supported set of EPL Pages is advertised in the **CdbMaxPagesEPL** field and is conditional on CDB support advertised in the general **CdbInstancesSupported** field (see Table 8-54).

For efficient READ or WRITE access to the multi-page EPL area modules may support Auto Paging and READ and WRITE length extensions.

Auto Paging

If the CdbAutoPagingSupported bit is set (see Table 8-54), except for the last **supported** EPL Page, host READ or WRITE accesses past the end of a Page in the EPL Page Range A0h-AEh cause the module to automatically increment the Page number and wrap the current address pointer to the beginning of the next Page (Byte 128). At the end of the last **supported** EPL Page, Auto Paging wraps the Page number back to A0h and the current address pointer to Byte A0h:128.

Otherwise, if the CdbAutoPagingSupported bit is cleared, host READ or WRITE accesses past the end of a Page will wrap around to Byte 128 of the same Page, as described in section B.1 and section B.2.5.2.1 respectively.

Length Extension

The allowed length of READ and WRITE operations is determined from the **CdbReadWriteLengthExtension** advertisement field (see Table 8-54).

Table 8-181 EPL Segments (Pages A0h-AFh)

Byte	Bits	Page Name	Page Description	Type
128-255	7-0	EPLSegment<i>	<p>EPLSegment<i> is the <i>th 128 bytes segment of the overall Extended Payload (EPL) area, where <i> is a sequential segment number, starting with EPLSegment1 (on Page A0h) and counting up to EPLSegment16 (on Page AFh).</p> <p>The maximum usable EPL length depends on the number of EPL segments supported and does not exceed 2048 Bytes.</p> <p>Advertisement: 01h:163</p> <p><i>Note: The EPL area is used both for host-written CMD data and for module-written REPLY data (possibly overwriting CMD data), as specified for each CDB Command in chapter 9.</i></p>	RW Adv.

9 CDB Command Reference

CDB is an optional feature for command-reply message exchange between host and module.

Support of CDB is advertised in 01h:163 (section 8.4.11).

When a module supports CDB, a subset of all defined CDB commands, such as e.g. feature queries, is (conditionally) required. Information about conditionally required command groups can be found in Table 9-1. Within each conditionally required command group there is then further per-command information whether a command is conditionally required.

All other commands are optional.

As advertised in 01h:163.7-6 a module may offer more than one CDB message exchange instance.

CDB messaging instances are distinguished by the Bank Address of the CDB message exchange Pages.

9.1 CDB Command Group Summary

A single CDB data exchange between host and module follows a command-reply pattern.

The individual message interactions are therefore referred to as **CDB commands** (for short, despite the fact that there are both commands and associated responses).

Each type of CDB interaction (CDB command) is identified by a CDB command identifier (**CMD ID**) that determines purpose, structure, and semantics of a CDB data exchange.

Table 9-1 CDB Command Groups

CMD IDs		Command Group	Description	Group	See
from	to				
0000h	003Fh	Module Commands	CDB module level commands.	Rqd.	9.3
0040h	005Fh	Capabilities Inquiry	Query advertised CDB features and capabilities	Rqd.	9.4
0060h	006Fh	-	Reserved , for Bulk Read of Banks and Pages.	-	9.4.6
0070h	007Fh	-	Reserved , for Bulk Write of Banks and Pages.	-	9.6
0080h	00FFh	-	Reserved , for module level commands	-	-
0100h	011Fh	Firmware Management	CDB Firmware Management	Adv.	9.7
0120h	01FFh	-	Reserved	-	-
0200h	027Fh	Performance Monitoring	CDB Performance Monitoring	Adv.	9.8
0280h	02FFh	Data Recording	Non-volatile Data Recording and Monitoring	Adv.	9.9
0300h	037Fh	-	Reserved , for BERT functionality		9.10
0380h	03FFh	Diagnostics and Debug	Reserved , for Diagnostics & Debug commands	Adv.	9.11
0400h	04FFh	Security Features	Security related CDB commands	Adv.	9.12
0500h	3FFFh	-	Reserved	-	-
4000h	40FFh	Versatile Control Set	Module-defined and possibly application-specific SI parameters in the control sets are defined in the CMIS-VCS supplement [6]	Adv.	[6]
4100h	7FFFh	-	Restricted for use by OIF. This CMD ID range allows the OIF to define new groups of messages specific for the application in separate supplement documents. <i>Note: The CMD ID ranges of each group may provide additional sub-ranges restricted for use by OIF.</i>		
8000h	FFFFh	-	Custom		

9.2 General Messaging Rules

9.2.1 Command and Reply

One message interaction consisting of a host command (CMD) and the associated module response (REPLY).

The host prepares the CMD parameters and triggers the module to execute the CMD.

The module processes the requested command and prepares REPLY header and REPLY data.

When processing is complete, the module sets the command completion Flag and the command status Byte, depending on the CDB instance used, as follows:

- CdbCmdCompleteFlag1 (00h:8.6) and CdbStatus1 (00h:37) for CDB instance 1
- CdbCmdCompleteFlag2 (00h:8.7) and CdbStatus2 (00h:38) for CDB instance 2

These rules apply to all CDB commands and with both background mode and foreground mode processing.

9.2.2 Use of Multiple CDB Instances

One command-reply message interaction always uses the same CDB instance (i.e. the same Bank address).

For protocols involving sequences of CDB commands (such as Firmware download), each protocol execution is bound to one of the possibly several CDB messaging instances. Sending protocol messages in parallel or sequentially over different CDB instances may result in undefined module behavior.

9.2.3 Preparing for Command Extensions

Future CMIS revisions may add optional parameters (with backwards compatible defaults) to existing CDB commands, by redefining previously reserved portions of the message body in a backwards compatible way.

It has not been specified yet if such changes to a message definition may include changing the message length.

Modules wishing to be tolerant against such future command extensions would have to avoid checking the length fields against expectations.

Such modules are therefore advised to react to commands with unexpected length fields in the following way:

- When receiving shorter message content (compatible with field structure), process only this content and assume zero values in the remaining expected fields.
- When receiving longer message content, process the content only up to the length known
- In all cases, the check code RPLChckCode needs to be calculated over the received length

9.2.4 Preparing for Query Reply Extensions

It is expected that message content may be added to CDB query command **replies** in future CMIS versions.

Example: A CDB command to retrieve object attributes may return more attributes in future revisions.

To be backwards compatible, such extensions will always be appended to the existing message reply body, and therefore the content of the length field changes accordingly.

Note: This is possible because almost all replies define the remaining space in LPL or EPL as unspecified.

Hosts are therefore advised to react to message replies with unexpected length fields in the following way:

- When receiving shorter message content (compatible with field structure), process only this content
- When receiving longer message content, process the content only up to the length known
- In all cases, the check code RPLChckCode needs to be calculated over the received length

1 **9.3 CDB Module Commands**

2

Table 9-2 CDB Module Commands Summary

ID	Command Title	Description	Type	Section
0000h	Query Status	The module returns a CdbStatus of success (01h) as well as basic status information of the module, such as password unlock status. <i>Note: This can also be used to check if CDB is supported.</i>	Rqd.	9.3.1
0001h	Enter Password	This command allows the host to enter a password. The module reports success or failure, especially if the password was accepted or rejected. <i>Note: Entering passwords is also possible by writing to Bytes 00h:122-125, but this method lacks feedback if the password was accepted or rejected.</i>	Adv.	9.3.2
0002h	Change Password	This command allows the host to change the Host Password. The module's reply reports success or failure, i.e. if the new Host Password has been accepted and stored successfully in non-volatile memory, or if there were any errors. <i>Note: Changing the Host Password persistently is also possible by writing to Bytes 00h:118-121, but without feedback if the operation completed successfully.</i>	Adv.	9.3.3
0003h	-	Reserved		
0004h	Abort Processing	Abort Current Background Operation	Adv.	9.3.4
0005h-001Fh	-	Reserved		

3

9.3.1 CMD 0000h: Query Status

This Query Status command may be used to retrieve the password acceptance status (if any) and to perform a test of the CDB messaging.

The **ResponseDelay** time parameter in the message defines a host specified delay before the module is to return with the response to this command by asserting the CDB complete Flag.

Note: This delay insertion can be used as a test especially if the module advertises via 01h:163.5 that CDB commands are processed in the background. If the MCI is enabled during the delay time, CdbStatus reports "In Progress" status until the delay time has expired. If MCI is disabled, then the MCI will NACK until the delay time has elapsed and the CDB response is ready.

The return message shows the current module unlock level.

Table 9-3 CDB Command 0000h: Query Status

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Query Status CMD ID	0000h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	2
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-137	ResponseDelay	U16. Programmable delay in ms for module responding to this command. A value of 0 asks for module response as fast as possible	
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy capturing command 10 000010b: Busy checking/validating command 10 000011b: Busy executing command On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	Encoded Length and Check Code for REPLY message body in LPL or EPL. See Table 8-179.	2
9Fh	135	RPLChkCode		comp.
REPLY Data (LPL)				
9Fh	136	Length	Length of this message payload (including this byte)	2
9Fh	137	Status	0000 0000b: Module Boot Up. 0000 0001b: Host Password Accepted. 1xxx xxxx b: Module Password accepted. Bits 'x' may contain custom information. Module passwords are passwords with Bit 31 set in the password field (see section 8.2.14).	
9Fh	138-255	-	No data returned. Content not modified by the module ¹	

12

¹ Bytes that are not part of the returned reply message are usually 'not specified'. Here the stronger specification 'not modified' is used, for historical reasons (maintain backwards compatibility with previous wording).

1 **9.3.2 CMD 0001h: Enter Password**

2 The Enter Password command allows the host to enter a host password

3 **Table 9-4 CDB Command 0001h: Enter Password**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Enter Password CMD ID	0001h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	4
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-139	Password	Password to be entered	
9Fh	140-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error 01 000110b: Password error – not accepted	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

9.3.3 CMD 0002h: Change Password

The Change Password command allows the host to change the Host Password.

Table 9-5 CDB Command 0002h: Change Password

Page	Byte	Field Name	Description	Value
CMD Header Fields				
9Fh	128-129	CMDID	Change Password CMD ID	0002h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	4
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-139	New password		
9Fh	140-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error (e.g. Bit 31 is set). 01 000101b: CdbChkCode error 01 000110b: Insufficient privilege to change password	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

1 **9.3.4 CMD 0004h: Abort Processing**

2 The Abort command allows the host to abort any current background operation.

3 **Table 9-6 CDB Command 0004h: Abort**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Abort Processing CMD ID	0004h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	FBh
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

1 **9.4 CDB Features and Capabilities Inquiry Commands**

2
3

Table 9-7 CDB Feature and Capabilities Commands Overview

ID	Command Title	Description	Type	Section
0040h	Module Features	Identify which module level commands are supported in the CMD ID space 0000h to 00FFh . <i>Note that this includes the commands in this group.</i>	Rqd.	9.4.1
0041h	Firmware Management Features	Identify which firmware management features are supported in the CMD ID space 0100h to 011Fh .	Rqd.	9.4.2
0042h	Performance Monitoring Features	Identify which Performance Monitoring and Data Monitoring and Recording commands are supported in CMD ID space 0200h to 02FFh	Rqd.	9.4.3
0043h	BERT and Diagnostic Features	Identify which BERT and Diagnostic features are supported in the CMD ID space 0300h to 03FFh	Rqd.	9.4.4
0044h	Security Features	Identify which Security related commands and features are supported in the CMD ID space 0400h to 04FFh	Adv.	9.4.5
0045h	Externally Defined Features	Identify which CMIS supplement specifications that define CDB messages in a dedicated part of the restricted CMD ID space from 4000h to 7FFFh are supported	Adv.	9.4.6
0046h-005Fh	-	Reserved		

4

9.4.1 CMD 0040h: Module Features

This command is used to query which CDB commands are supported.

Table 9-8 CDB Command 0040h: Module Features

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Module Features CMD ID	0040h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	Bfh
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	36
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136	CDB Flags	Reserved for additional CDB Flags.	00h
9Fh	137	-	Reserved	00h
9Fh	138	CMDs 0000h-0007h	This array of 32 bytes indicates support of CDB commands CMD <i>, with identifiers $0 \leq < i > \leq 255$, as follows: CMD <i> is supported when bit <j>=<i> mod 8 of byte <k> = $138 + \text{floor}(<i>/8)$ is set. <i>Note: Previous CMIS revisions were ambiguous if required commands are also "advertised" here. To honor existing implementations, not indicating a required command is allowable (but deprecated).</i>	
9Fh	139	CMDs 0008h-000Fh		
9Fh	140	CMDs 0010h-0017h		
9Fh	141	CMDs 0018h-001Fh		
9Fh	142	CMDs 0020h-0027h		
9Fh	143	CMDs 0028h-002Fh		
9Fh	144	CMDs 0030h-0037h		
9Fh	145	CMDs 0038h-003Fh		
9Fh	146	CMDs 0040h-0047h		
9Fh	147	CMDs 0048h-004Fh		
9Fh	148	CMDs 0050h-0057h		
9Fh	149	CMDs 0058h-005Fh		
9Fh	150	CMDs 0060h-0067h		
9Fh	151	CMDs 0068h-006Fh		
9Fh	152	CMDs 0070h-0077h		
9Fh	153	CMDs 0078h-007Fh		
9Fh	154-169	CMDs 0080h-00FFh		
9Fh	170-171	MaxCompletionTime	U16 Maximum CDB command execution time in ms, of all supported CDB commands. <i>Note: If exceeded, the host may send the CDB Abort Command.</i> <i>Note: The maximum possible MaxCompletionTime is about one minute (65.535 seconds).</i>	
9Fh	172-255	-	No data returned. Content not specified.	undef.

9.4.2 CMD 0041h: Firmware Management Features

This command is used to query which firmware management features are supported and to query command performance attributes.

Table 9-9 CDB Command 0041h: Firmware Management Features

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Firmware Management Features CMD ID	0041h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	BEh
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	18
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136	-	Reserved	0
9Fh	137.7	ImageReadback	0b = Full Image Readback Not Supported 1b = Full Image Readback Supported (see section 7.3.1)	
	137.6-4	-	Reserved	000b
	137.3	MaxDurationCoding	0b = max duration multiplier M is 1 1b = max duration multiplier M is 10 This bit encodes a multiplier value M which governs the interpretation of values found in the U16 array of advertised max durations in Bytes 144-153 of this message: These advertised values are multiplied by M .	
	137.2	SkippingErasedBlocks	0b = Skipping erased blocks Not Supported 1b = Skipping erased blocks Supported	
	137.1	CopyCmd	0b = CMD 0108h (Copy image) Not Supported 1b = CMD 0108h (Copy image) Supported	
	137.0	AbortCmd	0b = CMD 0102h (Abort) Not Supported 1b = CMD 0102h (Abort) Supported	
	138	StartCmdPayloadSize	This defines the number of bytes that the host must extract from the beginning of the vendor-delivered binary firmware image file and send to the module in CMD 0101h (Start).	
9Fh	139	ErasedByte	This is the value representing an erased byte. The purpose of advertising this byte is to optionally reduce download time by allowing the host to skip sending blocks of the image containing ErasedByte values only. <i>Note: Typically for NAND flash the erased state is FFh, for other flash or EEPROM technology it is 00h.</i>	
9Fh	140	ReadWriteLengthExt	ReadWriteLengthExt = i specifies the allowable additional number of byte octets in a READ or a WRITE, specifically for Firmware Management Commands (IDs 0100h-01FFh) as follows	

Page	Byte	Field Name	Description	Value
			<p>EPL: For accessing the multi-page EPL field, the allowable length extension is i byte octets (8 bytes).</p> <p>LPL: For accessing the LPL field on page 9Fh, the allowable length extension is min(i, 15) byte octets.</p> <p>This leads to the maximum length of a READ or a WRITE</p> <p>Value Maximum Number of Bytes (EPL)</p> <p>0: 8 bytes (no extension of general length limit) i: 8 * (1+i) bytes (0 ≤ i ≤ 255) 255: 8 * 256 = 2048 bytes</p> <p>Value Maximum Number of Bytes (LPL)</p> <p>0: 8 bytes (no extension of general length limit) i: 8 * (1+i) bytes (0 ≤ i ≤ 15) i: 8 * 16 = 128 bytes (16 ≤ i ≤ 256)</p> <p><i>Note: See also the CdbReadWriteLengthExtension Byte 01h:164 which defines a similar length extension for arbitrary CDB commands.</i></p>	
9Fh	141	WriteMechanism	Firmware update supported mechanism. 00h: None Supported. 01h: Write to LPL supported. 10h: Write to EPL supported. 11h: Both Write to LPL and EPL supported.	
9Fh	142	ReadMechanism	Firmware read / readback support mechanism. 00h: None Supported. 01h: Read via LPL supported. 10h: Read via EPL supported. 11h: Both Read via LPL and EPL supported.	
9Fh	143	HitlessRestart	0: CMD Run Image causes a reset. Traffic is affected. 1: CMD Run Image may reset but module will do its best to maintain traffic and management states. Data path functions are not reset.	
9Fh	144-145	MaxDurationStart	U16 Maximum time in M ms for a CDB Start command to complete execution, where M is defined by bit 137.3	
9Fh	146-147	MaxDurationAbort	U16 Maximum time in M ms for a CDB Abort command to complete execution, where M is defined by bit 137.3	
9Fh	148-149	MaxDurationWrite	U16 Maximum time in M ms for a CDB Write command to complete execution, where M is defined by bit 137.3	
9Fh	150-151	MaxDurationComplete	U16 Maximum time in M ms for a CDB Complete command to complete execution, M defined by bit 137.3	
9Fh	152-153	MaxDurationCopy	U16 Maximum time in M ms for a CDB Copy command to complete execution, where M is defined by bit 137.3	
9Fh	154-255	-	No data returned. Content not specified.	undef.

1 **9.4.3 CMD 0042h: Performance Monitoring Features**

2
3

Table 9-10 CDB Command 0042h: Performance Monitoring Features

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Performance Monitoring Features CMD ID	0042h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	BDh
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	32
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136	CMDs 0200h-0207h support	Each bit represents a mask. If bit is "1" then command is supported. D0: CMD 0200h is supported. .. D7: CMD 0207h is supported.	
9Fh	137	CMDs 0208h-020Fh	Each bit represents a mask. If bit is "1" then command is supported. D0: CMD 0208h is supported. .. D7: CMD 020Fh is supported.	
9Fh	138-151	CMDs 0210h-027Fh	Bitmask defines command supported.	
9Fh	152	CMDs 0280h-0287h	Data Monitoring Command Supported. D0: CMD 0280h supported. .. D7: CMD 0287h is supported.	
9Fh	153-167	CMDs 0288h-02FFh		
9Fh	168-255	-	No data returned. Content not specified.	undef.

4

1 **9.4.4 CMD 0043h: BERT and Diagnostics Features**

2

3

Table 9-11 CDB Command 0043h: BERT and Diagnostics Features

Page	Byte	Field Name	Description	Value
CMD Header Fields				
9Fh	128-129	CMDID	BERT and Diagnostics Features CMD ID	0043h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	BCh
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	32
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136	CMDs 0300h-0307h support	Each bit represents a mask. If bit is "1" then command is supported D0: CMD 0300h is supported. .. D7: CMD 0307h is supported.	
9Fh	137	CMDs 0308h-030Fh	Each bit represents a mask. If bit is "1" then command is supported D0: CMD 0308h is supported. .. D7: CMD 030Fh is supported.	
9Fh	138-151	CMDs 0310h-037Fh	Bitmask defines command supported.	
9Fh	152	CMDs 0380h-0387h	Data Monitoring Command Supported. D0: CMD 0380h supported. .. D7: CMD 0387h is supported.	
9Fh	153-167	CMDs 0388h-03FFh		
9Fh	168-255	-	No data returned. Content not specified.	undef.

1 **9.4.5 CMD 0044h: Security Features and Capabilities**

2

3

Table 9-12 CDB Command 0044h: Security Features and Capabilities

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Security Features and Capabilities CMD ID	0044h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	50
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136	CMD 0400h-0407h support	Each bit represents a mask. If a bit is set, the corresponding command is supported D0: CMD 0400h is supported. .. D7: CMD 0407h is supported	
9Fh	137-167	CMD 0408h-04FFh	Bitmask defining commands supported.	0
9Fh	168	NumCertificates	Number of public certificates the host may obtain from the module. The device must contain a single leaf certificate and it may optionally contain one or more intermediate certificates optionally followed by a root certificate. For X.509 certificates, intermediate certificates are not self-signed, and the root cert is self-signed. NumCertificates <= 4.	
9Fh	169	CertChainSupported	0: Certificate chain is not supported. Module contains leaf certificate instance i = 0 only. 1: Module supports certificate chain and host must specify the instance when downloading a certificate. Instance i = 0 is the start of the chain, i.e. the leaf certificate, and any instance i+1 is another certificate used to sign the certificate instance i , where i < NumCertificates <= 4	
9Fh	170	CertificateFormat	IDevID certificate format: 0: Not supported 1: Custom 2: X509v3 DER encoding 3-255: Reserved	
9Fh	171	-	Reserved	0
9Fh	172-173	CertificateLength1	Length of leaf certificate i = 0	
9Fh	174-175	CertificateLength2	Length of certificate i = 1 or 0 when not supported	
9Fh	176-177	CertificateLength3	Length of certificate i = 2 or 0 when not supported	
9Fh	178-179	CertificateLength4	Length of certificate i = 3 or 0 when not supported	

Page	Byte	Field Name	Description	Value
9Fh	180	DigestLength	Required message hash digest length (in bytes) 0: Not supported 1: 28 bytes (SHA224) 2: 32 bytes (SHA256) 3: 48 bytes (SHA384) 4: 64 bytes (SHA512) 5-255: Reserved	
9Fh	181	-	Reserved	0
9Fh	182-183	SignatureTime	Maximum time (in milliseconds) for signature generation. After a "Set Signature Digest" command, the host must wait prior to "Get Signature".	
9Fh	184-185	SignatureLength	Length (in bytes) of the encoded/padded (if applicable) digest signature	
9Fh	186	SignatureFormat	Signature Format: 0: Not supported 1: Custom , vendor specific encoding 2: Raw binary byte stream 3: DER encoding 4: ECDSA (R,S) integer pair, integers prefixed with length 5-255: Reserved	
9Fh	187	SignaturePadScheme	Signature Padding Scheme: 0: None 1: Custom 2: PKCS#1 v1.5 3: OAEP 4: PSS 5-255: Reserved. <i>Note: Randomized padding is recommended</i>	
9Fh	188-255	-	No data returned. Content not specified.	undef.

1 **9.4.6 CMD 0045h: Externally Defined Features**

2
3

Table 9-13 CDB Command 0045h: Externally Defined Features

Page	Byte	Field Name	Description	Value
CMD Header Fields				
9Fh	128-129	CMDID	Externally Defined Features CMD ID	0043h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	1
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136	SupplementSupport	Bit 0 = 0/1: CMIS-VCS not supported/supported	
9Fh	137-255	-	No data returned. Content not specified.	undef.

9.4.7 CMD 0050h: Get Application Attributes

Some module advertisements are Application dependent, especially in versatile programmable modules.

For instance, a host may not be able to accommodate the advertised worst case power dissipation of a module, but there may be less power-hungry Applications that could be used.

To allow the host to understand Application-dependent properties, the following command can be used.

Table 9-14 CDB Command 0050h: Get Application Attributes

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Application Attributes CMD ID	0050h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL	2
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-137	ApplicationNumber	U16 Application Number 15-8: reserved (0) 7-4: NADBlockIndex (0-15) or 0 3-0: AppSelCode (1-15)	
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	20
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136-137	ApplicationNumber	U16: Application Number 15-8: reserved (0) 7-4: NADBlockIndex (0-15) or 0 3-0: AppSelCode (1-15)	
9Fh	138-139	MaxModulePower	U16: Worst case module power dissipation when this Application is instantiated homogeneously as often as possible in parallel (when applicable) with worst case configuration options. ¹ Unit: 0.25 W	
9Fh	140-141	ProgOutputPowerMin	S16: Minimum Programmable Output Power, Unit: 0.01 dBm	
9Fh	142-143	ProgOutputPowerMax	S16: Maximum Programmable Output Power, Unit: 0.01 dBm	
9Fh	144-145	PreFECBERThreshold	F16: Pre FEC BER VDM high alarm threshold	
9Fh	146-147	RxLOSOpticalPowerThr	S16: Optical power threshold for RxLOS alarm Unit: 0.01dBm	
9Fh	148-149	RxPowerHighAlarmThr	U16: OpticalPowerRxHighAlarmThreshold Unit: 0.1uW	

¹ When hybrid parallel Data Paths carrying different Application are possible and envisaged, hosts are advised to assume the overall module power dissipation of the worst-case Application.

Page	Byte	Field Name	Description	Value
9Fh	150-151	RxPowerLowAlarmThr	U16: OpticalPowerRxLowAlarmThreshold Unit: 0.1uW	
9Fh	152-153	RxPowerHighWarnThr	U16: OpticalPowerRxHighWarningThreshold Unit: 0.1uW	
9Fh	154-155	RxPowerLowWarnThr	U16: OpticalPowerRxLowWarningThreshold Unit: 0.1uW	
9Fh	156-255	-	No data returned. Content not specified.	undef.

9.4.8 CMD 0051h: Get Interface Code Description

SFF_8024 [5] provides the Interface ID codes for host and media interfaces together with descriptive attributes, which are well-known for standardized interfaces but unknown for vendor-defined custom interfaces.

To allow hosts building descriptive pick lists for all Applications including those using custom interfaces, this command can be used to query those descriptive attributes from the module.

Modules supporting this command are encouraged to return information for all supported interface types, while they are required to do that for all supported custom interfaces.

Table 9-15 CDB Command 0051h: Get Interface Code Description

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Interface Code Description CMD ID	0051h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL	3
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	
9Fh	134	RPLLength	<i>Note: Initiator may fill those Ascii fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-137	InterfaceID	U16: HostInterfaceID or MediaInterfaceID 15-8: reserved (0) 7-0: InterfaceID	
9Fh	138	InterfaceLocation	0: media side 1: host side	
9Fh	139-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	92
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136-137	InterfaceID	Unmodified CMD Data	
9Fh	138	InterfaceLocation	Unmodified CMD Data	
9Fh	139	-	Reserved (for alignment of X16 fields below)	0
9Fh	139-154	InterfaceName	ASCII[16]: Short Name	
9Fh	155-202	InterfaceDescription	ASCII[48]: Description and Specification Reference	
9Fh	204-205	InterfaceDataRate	F16: Application Bit Rate in Gb/s	
9Fh	206-207	InterfaceLaneCount	U16: Number of parallel lanes	
9Fh	208-209	LaneSignalingRate	F16: Lane Signaling Rate in GBd	
9Fh	210-225	Modulation	ASCII[16]: Lane Modulation Format	
9Fh	226-227	BitsPerSymbol	U16: Bits per Lane Modulation Symbol	
9Fh	228-255	-	No data returned. Content not specified.	undef.

9.5 CDB Bulk Read Commands

Table 9-16 CDB Bulk Read Commands Overview

ID	Command Group	Description	Type	Section
0060h-006Fh	-	Reserved		

9.6 CDB Bulk Write Commands

Table 9-17 CDB Bulk Write Commands Overview

ID	Command Group	Description	Type	Section
0070h-007Fh	-	Reserved		

9.7 CDB Firmware Management Commands

Firmware management using CDB is an optional feature.

See section 8.2.9 for background on module firmware as a potential aggregate of several firmware components and on aggregate firmware version identification.

Table 9-18 CDB Firmware Download Commands Overview

ID	Command Title	Description	Type	Section
0100h	Get Firmware Info	When the host issues this command, the module returns the requested FW information of all field-updateable firmware in the module.	Rqd.	9.7.1
0101h	Start Firmware Download	The host issues this command to initiate a firmware update. The module may completely erase the target or simply acknowledge and prepare the module firmware for any appropriate update or dynamically erase as each data block arrives. On success, the host may begin sending the firmware image using command codes 0103h-0107h. Condition: Any of CMD 103h, 104h, 105h, 106h supported	Cnd.	9.7.2
0102h	Abort Firmware Download	Aborts the FW Download if a FW Start has been issued. Advertisement: CMD 0041h, 9Fh.137.0	Adv.	9.7.3
0103h	Write Firmware Block LPL	With this command, the host downloads a firmware image block previously stored in the LPL area. The module transfers that firmware image block into non-volatile storage. <i>Note: Each image block transfer from host to module is covered by the CDB command block checksum; this checksum does not ensure that the image has been properly transferred to non-volatile storage.</i> Advertisement: CMD 0041h, 9Fh.141	Adv.	9.7.4
0104h	Write Firmware Block EPL	With this command, the host downloads a firmware image block previously stored in the EPL Page(s). The module transfers that firmware image block into non-volatile storage. Advertisement: CMD 0041h, 9Fh.141	Adv.	9.7.5
0105h	Read Firmware Block LPL	The host may use this command to read back the firmware image that was most recently written to non-volatile storage. The module copies the image from non-volatile storage to the LPL Page. The firmware image transfer from the module to the host is covered by the CDB command block checksum, but this checksum does not ensure that the image has been properly transferred from non-volatile storage to the LPL Page. Advertisement: CMD 0041h, 9Fh.142	Adv.	9.7.6
0106h	Read Firmware Block EPL	The host may use this command to read back the firmware image that was most recently written to non-volatile storage. The module copies the image from non-volatile storage to the EPL Page(s). The firmware image transfer from the module to host is covered by a block checksum, but this checksum does not ensure that the image has been properly transferred from non-volatile storage to the EPL Page(s). Advertisement: CMD 0041h, 9Fh.142	Adv.	9.7.7

ID	Command Title	Description	Type	Section
0107h	Complete Firmware Download	<p>The host issues this command when the entire firmware image has been written via LPL or via EPL Pages, or to stop the download after failure. If this command is not issued, the firmware cannot be run even if the image is properly loaded to non-volatile storage. The module validates the checksum associated with the image when the host issues this command.</p> <p>Condition: Any of CMD 103h, 104h, 105h, 106h supported</p>	Cnd.	9.7.8
0108h	Copy Firmware Image	<p>When multiple images are supported for a given subsystem in the module, this command causes the module to copy an image from one non-volatile storage location to another one. It is assumed that the copy firmware image command includes a validation process by the module firmware to ensure the copied image is valid. The CDB complete firmware image command 0107h does not need to be called after a copy firmware image command.</p> <p>Advertisement: CMD 0041h, 9Fh.137.1</p>	Adv.	9.7.9
0109h	Run Firmware Image	<p>This command is used to start and run a selected image. This command transfers control from the currently running firmware to a selected firmware that is started. It can be used to switch between firmware versions, or to perform a restart of the currently running firmware.</p> <p>Condition: CMD 103h or CMD 104h supported</p>	Cnd.	9.7.10
010Ah	Commit Firmware Image	<p>The host uses this command to commit the running image so that the module will boot from it on future boots.</p> <p><i>The assumption is that the host has a time period where it runs the new firmware and "accepts" the new firmware. During this time, if a reset occurs, the previously committed image will be run. When the host issues this command, the module will mark a non-volatile storage location to be used after future module resets. Only the running image can be committed. This is to avoid committing a "bad" image.</i></p> <p>Condition: CMD 103h or CMD 104h supported</p>	Cnd.	9.7.11
010Bh - 011Fh	-	Reserved		

9.7.1 CMD 0100h: Get Firmware Info

This command returns the firmware versions and firmware default running images that reside in the module. Firmware images A, B and either a factory or boot firmware image version.

Note: See section 8.2.9 for semantical expectations about unique identification of the entire firmware aggregate by the version triple (major revision, minor revision, build number).

Table 9-19 CDB Command 0100h: Get Firmware Info

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Firmware Info CMD ID	0100h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	Feh
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-178	110
9Fh	135	RPLChkCode	See Table 8-178	comp.
REPLY Data (LPL)				
9Fh	136	FirmwareStatus	Bitmask to indicate FW Status. Image in Bank A : Bit 0: Operational Status Bit 1: Administrative Status Bit 2: Validity Status Bit 3: Reserved Image in Bank B : Bit 4: Operational Status Bit 5: Administrative Status Bit 6: Validity Status Bit 7: Reserved Encoding as follows: Operational Status: 1 = running, 0 = not running Administrative Status: 1=committed, 0=uncommitted Validity Status: 1 = invalid, 0 = valid <i>Note: Zero-encoding of valid maintains backwards compatibility with CMIS 4.0</i> As per Table 7-4, codes 00h, 04h, 40h, 44h indicate that a boot loader or a factory image is running (if supported). See section 7.3.1.4 for a detailed description.	
9Fh	137	ImageInformation	Bit 0: Firmware image A information in 9Fh:138-173 Bit 1: Firmware image B information in 9Fh:174-209 Bit 2: Factory or Boot image information in 9Fh:201-245	
9Fh	138	ImageAMajor	Image A firmware major revision	
9Fh	139	ImageAMinor	Image A firmware minor revision	

Page	Byte	Field Name	Description	Value
9Fh	140-141	ImageABuild	Image A firmware build number	
9Fh	142-173	ImageAExtraString	Additional information	
9Fh	174	ImageBMajor	Image B firmware major revision	
9Fh	175	ImageBMinor	Image B firmware minor revision	
9Fh	176-177	ImageBBuild	Image B firmware build number	
9Fh	178-209	ImageBExtraString	Additional information	
9Fh	210	FactoryBootMajor	Factory or Boot image firmware major revision	
9Fh	211	FactoryBootMinor	Factory or Boot image firmware minor revision	
9Fh	212-213	FactoryBootBuild	Factory or Boot image firmware build number	
9Fh	214-245	FactoryBootExtraStr	Additional information	
9Fh	246-255	-	No data returned. Content unspecified	undef.

1

9.7.2 CMD 0101h: Start Firmware Download

The module may erase the whole image or simply prepare for the firmware to be updated. The duration of this command depends on the host-written payload of the command. A complete erase of an image may take several seconds.

Table 9-20 CDB Command 0101h: Start Firmware Download

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Start firmware download process CMD ID	0101h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	comp.
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-139	ImageSize	U32 Size of firmware image to download into the module. This should be the file size including the LPL bytes sent as vendor data in this message.	Var.
9Fh	140-143	-	Reserved	0
9Fh	144-255	VendorData	U8[112] The vendor may send up to 112 bytes of information in the Start Firmware Download command. It is recommended that the binary file delivered has up to 112 bytes of header that is sent to the module. The information within this field can e.g. be used by a vendor to reject an incorrect file (binary file) and prevent firmware loading of an incorrect file presented to the module.	
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	

1 **9.7.3 CMD 0102h: Abort Firmware Download**

2 Aborts the firmware update process.

3 Modules supporting only a single image should not implement this command and advertise accordingly.

4 **Table 9-21 CDB Command 0102h: Abort Firmware Download**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Abort Firmware Download CMD ID	0102h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	FCh
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

1 **9.7.4 CMD 0103h: Write Firmware Block LPL**

2 Download one block of the firmware image via LPL

3 **Table 9-22 CDB Command 0103h: Write Firmware Block LPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Write Firmware Block LPL CMD ID	0103h
9Fh	130-131	EPLLength	EPL is not used	0000h
9Fh	132	LPLLength	The actual length of the firmware block in the FirmwareBlock field + 4.	Comp.
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-139	BlockAddress	U32 Starting byte address of this block of data within the supplied image file minus the size of the "Start Command Payload Size". See section 7.3.1.2.	
9Fh	140-255	FirmwareBlock	U8[116] One block of the firmware image. The actually needed length may be shorter than the available FirmwareBlock field size. This actual length of the block is defined in Byte 132 (LPLLength), see above.	
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

1 **9.7.5 CMD 0104h: Write Firmware Block EPL**

2 Download one block of the firmware image via EPL

3 **Table 9-23 CDB Command 0104h: Write Firmware Block EPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Write Firmware Block EPL CMD ID	0104h
9Fh	130-131	EPLLength	Length of the firmware block in the FirmwareBlock field, which may span over multiple Pages.	Comp.
9Fh	132	LPLLength	LPL length	4
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-139	BlockAddress	U32 Starting byte address of this block of data within the supplied image file minus the size of the "Start Command Payload Size". See section 7.3.1.2.	
9Fh	140-255	-	No data passed. Content not specified.	undef.
CMD Data (EPL)				
A0h-Afh	128-255	FirmwareBlock	Up to 2048 Bytes. Actual Length specified in EPLLength	
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.
REPLY Data (EPL)				
A0h-Afh	128-255	-	No data returned. Content not specified.	undef.

1 **9.7.6 CMD 0105h: Read Firmware Block LPL**

2 Upload one block of firmware image from the indicated block in non-volatile storage via LPL

3 *Note: The image being read is the most recently written, implicitly.*

4 **Table 9-24 CDB Command 0105h: Read Firmware Block LPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Read Firmware Block LPL CMD ID	0105h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	6
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-139	BlockAddress	U32 Starting byte address of this block of data within the supplied image file minus the size of the size of the "Start Command Payload Size". See section 7.3.1.2.	
9Fh	140-141	Length	U16 Number of bytes to read back to the LPL in this command, starting at the indicated address.	
9Fh	142-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	var.
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136-139	Address of block	U32 Base address of the data block within the firmware image	
9Fh	140-255	ImageData	U8[116]: Up to 116 Bytes	

1 **9.7.7 CMD 0106h: Read Firmware Block EPL**

2 Upload one block of firmware image from the indicated block in non-volatile storage via EPL

3 Note: The image being read is the most recently written, implicitly.

4 **Table 9-25 CDB Command 0106h: Read Firmware Block EPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Read Firmware Block EPL CMD ID	0106h
9Fh	130-131	EPLLength	Length of EPL	0
9Fh	132	LPLLength	LPL length	6
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-139	BlockAddress	U32 Starting byte address of this block of data within the supplied image file minus the size of the size of the "Start Command Payload Size". See section 7.3.1.2.	
9Fh	140-141	Length	U16 Number of bytes to read back to the EPL in this command, starting at the indicated address.	
9Fh	142-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	Variable
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136-255	-	LPL unused. Contents unspecified	undef.
REPLY Data (EPL)				
A0h-Afh	128-255	ImageData	U8[128]: Up to 128 Bytes. Actual Length specified in RPLLength	

9.7.8 CMD 0107h: Complete Firmware Download

When the host issues this command, the module validates the complete firmware image and then return success or failure (could be checksum failure).

Note: This command can also be used to abort a firmware download. The module detects an incomplete download and is expected to return a failure.

Table 9-26 CDB Command 0107h: Complete Firmware Download

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Complete Firmware Download CMD ID	0107h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	F7h
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

9.7.9 CMD 0108h: Copy Firmware Image

Copy Firmware Image is an optional command within the firmware download commands that may be used in the firmware update process.

Note: This command is typically used in a system where both images that are written to flash are identical, and thus if the host desires to have both images A and B be identical, it can simply tell the module to copy from the running image which is committed to the uncommitted backup image.

Table 9-27 CDB Command 0108h: Copy Firmware Image

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Copy Firmware Image CMD ID	0108h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	1
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	6
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136-139	Length	U32 Number of bytes copied	
9Fh	140	CopyDirection	ABh: Image A was copied into Image B BAh: Image B was copied into Image A	
9Fh	141	CopyStatus	00h : Copy Successful 01h : Copy Failed	
9Fh	142-255	-	No data returned. Content not specified.	undef.

9.7.10 CMD 0109h: Run Firmware Image

The host uses this command to run a selected image from module internal firmware banks.

For example, after the firmware has been updated, this command may be used to switch to the new firmware version. The host may use CMD 0100h to determine both the active and inactive firmware versions.

Executing the Run Image command may potentially be hitless, non-disruptive or minimally disruptive to high-speed traffic.

Behavior of the module and its control loop transients during resets are vendor and technology dependent.

The maximum duration of a successful execution of the Run Firmware Image command is the same as that of a regular reboot using the ResetS transition signal.

Note: This duration is mainly governed by the tMgmtInit timing parameter. The transition time to the Reset state is assumed to be negligible.

Table 9-28 CDB Command 0109h: Run Firmware Image

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Run FW Image CMD ID	0109h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	4
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	-	Reserved	0
9Fh	137	ImageToRun	00h = Traffic affecting Reset to Inactive Image. 01h = Attempt Hitless Reset to Inactive Image 02h = Traffic affecting Reset to Running Image. 03h = Attempt Hitless Reset to Running Image	
9Fh	138-139	DelayToReset	U16 Indicates the delay in ms after receiving this command before a reset will occur, starting from the time the CDB complete Flag is set (or NACK clearing if the CDB background mode is not set). <i>Note: When DelayToReset is 0, the module may reset before the host has read the CdbStatus message.</i>	
9Fh	140-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	
9Fh	135	RPLChkCode	See Table 8-179	
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified	undef.

9.7.11 CMD 010Ah: Commit Firmware Image

A Commit is the process where the “running” image is set to be the image to be used on exit from module reset. In other words, a committed image is the image that will run and is expected to be a ‘good’ firmware version to run upon any resets (including watch dog).

This command is used to switch the committed image after the firmware update process, when the new firmware is running and when the host has determined that the new firmware is working properly.

Note: Commit Image commits only the image that it is currently running. If it was possible to commit a non-running, a bad version may be committed and attempted to run after the next module reset.

Note: There is no specific advertisement of the maximum duration for this command which is assumed to be insignificant and in the order of writing administrative information to nonvolatile memory (tWRITE).

Table 9-29 CDB Command 010Ah: Commit Image

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Commit Image CMD ID	010Ah
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified	undef.

1 9.8 CDB Performance Monitoring Commands

2
3 **Table 9-30 CDB Performance Monitoring Commands Overview**

ID	Command Title	Description	Type	Section
0200h	Control PM	General Performance Monitoring Controls	Rqd.	9.8.1
0201h	Get PM Feature Information	Advertisement on optional PM is supported.	Rqd.	9.8.2
0202h-020Fh	-	Reserved		
0210h	Get PM Module LPL	Get Module-level X16 PM using LPL	Adv.	9.8.3
0211h	Get PM Module EPL	Get Module-level X16 PM using EPL	Adv.	9.8.3
0212h	Get PM Host Side LPL	Get Lane-specific host side X16 PM using LPL	Adv.	9.8.4
0213h	Get PM Host Side EPL	Get Lane-specific host side X16 PM using EPL	Adv.	9.8.4
0214h	Get PM Media Side LPL	Get Lane-specific media side X16PM using LPL	Adv.	9.8.5
0215h	Get PM Media Side EPL	Get Lane-specific media side X16 PM using EPL	Adv.	9.8.5
0216h	Get PM Data Path LPL	Get Lane-specific Data Path X16 PM using LPL	Adv.	9.8.6
0217h	Get PM Data Path EPL	Get Lane-specific Data Path X16 PM using EPL	Adv.	9.8.6
0218h-0219h	-	Reserved		
0220h	Get Data Path RMON Statistics	Get Data Path RMON Statistics using LPL	Adv.	9.8.7
0221h-0229h	-	Reserved		
0230h	Control Data Path SEW Histogram	Control FEC Symbol Error Weight Histogram	Adv.	9.8.8
0231h	Get Data Path SEW Histogram	Get FEC Symbol Error Weight Histogram	Adv.	9.8.9
0232h	Control Data Path SEW _{max} Stats	Control Max FEC Symbol Error Weight Statistics	Adv.	9.8.10
0233h	Get Data Path SEW _{max} Stats	Get Max FEC Symbol Error Weight Statistics	Adv.	9.8.11
0234h-027Fh	-	Reserved		

4 The following Table provides an Overview of the PM Observables available for retrieval by CDB commands.

5 See Table 8-170 for VDM Observable IDs

6 **Table 9-31 CDB Performance Monitoring Observables Cross Reference**

PM Group	Instances	CDB ID: Type and Observable	Unit	Corresponds to Observables
Module	1	0: S16 Module Temperature 1: S16 Vcc 2: S16 Aux1Mon 3: S16 Aux2Mon 4: S16 Aux3Mon	1/256 degC 100 uV see → see → see →	0:00h:14-15 0:00h:16-17 0:00h:18-19 0:00h:20-21 0:00h:22-23
Host Side	Per Lane (across Banks)	0: U16 Host Side Lane SNR 1: F16 Host Side PAM4 LTP 2: F16 Host Side Pre-FEC BER	1/256 dB 1/256 dB	VDM ID 6 VDM ID 8 n/a
Media Side	Per Media Lane (across Banks)	0: U16 Tx Laser Bias 1: U16 Tx Optical Power 2: U16 Rx Optical Power 3: S16 Per-Lane Laser Temperature	0.1 uW 2uA * x 0.1 uW 1/256 degC	0-3:11h:170-185 0-3:11h:154-169 0-3:11h:186-201 VDM ID 4
Data Path	Per Data Path (across Banks)	0: F16 Media Side Frame Errors (FERC) 1: F16 Media Side Pre-FEC BER --: U16 Media Side SEW _{max} --: F16 Host Side Frame Errors (FERC) --: F16 Host Side Pre-FEC BER --: U16 Host Side SEW _{max}		VDM ID 23 VDM ID 15 VDM ID 33 VDM ID 24 VDM ID 16 VDM ID 34

9.8.1 CMD 0200h: Control PM

Controls the behavior of CDB Performance Monitoring and its behavior with respect to the Versatile Diagnostic Monitoring in Page 20h-2Fh.

This section details the messages used to extract PM data records such as minimum, average, maximum values.

Note: Unless otherwise specified, a 2-byte, 4-byte, or 8-byte value is encoded in Big Endian format, i.e. the lowest byte address stores the most significant byte of the word.

Table 9-32 CDB Command 0200h: Control PM

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Control PM CMD ID	0200h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	4
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136.1-7	-	Reserved	0
	136.0	LinkMode	0b: PM Objects are Independent 1b: PM Objects are the same as 20h-2Fh (linked) When PM objects are linked, this means that PM data in Page 20-2Fh are based on the same objects as the PM records that are returned by CDB Get PM commands. It also means that clearing statistics using either method in Page 2Fh will clear the statistics in CDB, too.	
9Fh	137	-	Reserved	0
9Fh	138.1-7	-	Reserved . Set to 0.	0
	138.0	ClearAllStatistics	0b: No operation 1b: Clear all statistics (minimum, average, maximum) for all observables for all lanes at the same time, across all Banks, in a best-effort manner	
9Fh	139	-	Reserved	0
9Fh	140-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	0
9Fh	135	RPLChkCode	See Table 8-179	0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

9.8.2 CMD 0201h: Get PM Feature Information

Identifies which of the PM monitors defined in CMD 0210h to 0217h is supported by the module.

Table 9-33 CDB Command 0201h: Get PM Feature Information

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get PM Feature Information CMD ID	0201h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	FCh
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	4
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136	HostSideMonitors	Bit 0: Bool: Host Side SNR monitor available Bit 1: Bool: Host Side LTP monitor available	X
9Fh	137	MediaSideMonitors	Bit 0: Bool: Media Side SNR monitor available Bit 1: Bool: Media Side LTP monitor available	X
9Fh	138	-	Reserved	0
9Fh	139	-	Reserved	0
9Fh	140-255	-	No data returned. Content not specified.	undef.

9.8.3 CMD 0210h/0211h: Get Module PM LPL/EPL

This command collects performance monitoring observables related to the module as a whole.

Table 9-34 CDB Command 0210h/0211h: Get Module PM LPL/EPL

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Module PM using LPL CMD ID Get Module PM using EPL CMD ID	0210h 0211h
9Fh	130-131	EPLLenght	EPL is not used	0
9Fh	132	LPLLenght	LPL length	5
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLenght	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136.7	ClearOnRead	0b: return selected PM data 1b: return selected PM data and then reset their statistics	
	136.1-6	-	Reserved	
	136.0	RecordType	0b: Return 6-byte PM record (min, mean, max) 1b: Return 8-byte PM record (append "current" value)	
9Fh	137	Observables	Bit 0: Module Temperature Bit 1: Vcc Bit 2: Aux1 Bit 3: Aux2 Bit 4: Aux3 Bit 5-7: Reserved	
9Fh	138	-	Reserved	
9Fh	139	-	Restricted (OIF)	
9Fh	140	-	Custom	
9Fh	141-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLenght	See Table 8-179	comp.
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL) (CMD 0210h)				
9Fh	136-255	LPL PM data (up to 120 bytes)	PM record of one observable consists of 6 or 8 bytes: X16 minimum value X16 average (mean) value X16 maximum value X16 current value (if requested) <i>Note: A maximum of 15 to 20 records can be returned using LPL, depending on the requested PM record length. The sequence of PM records is the same as the sequence of set bits in the Observables field, in order of ascending significance.</i>	

Page	Byte	Field Name	Description	Value
REPLY Data (EPL) (CMD 0211h)				
A0h to AFh	128-255	EPL PM data (maximum number of bytes depends on available EPL Pages)	PM record of one observable consists of 6 or 8 bytes: X16 minimum value X16 average (mean) value X16 maximum value X16 current value (if requested) <i>Note: The maximum number of records depends on the size of the EPL and the requested PM record length.</i> Data is contiguous across EPL Pages, and the sequence of PM records corresponds to the sequence of set bits in the Observables field, in order of ascending significance.	

9.8.4 CMD 0212h/0213h: Get PM Host Side LPL/EPL

This command collects performance monitoring observables related to individual host lanes.

Table 9-35 CDB Command 0212h/0213h: Get PM Host Side LPL/EPL

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get PM Host Side LPL CMD ID Get PM Host Side EPL CMD ID	0212h 0213h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	20
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136.7	ClearOnRead	0b: return selected PM data 1b: return selected PM data and then reset their statistics	
	136.1-6	-	Reserved	
	136.0	RecordType	0b: Return 6-byte PM record (min, mean, max) 1b: Return 8-byte PM record (append "current" value)	
9Fh	137-139	-	Reserved <i>In future, we could define start/stop/increment lanes here if needed for modules with more than 32 lanes.</i>	
9Fh	140-143	Lanes	U32 Lanes is a bitmask indicating which host lane is present, where bit i represents lane i+1	
9Fh	144	Observables	Bit 0: Host Side Lane SNR Bit 1: Host Side PAM4 LTP Bit 2: Host Side Pre-FEC BER ¹ Bits 3-7: Reserved	
9Fh	145-147	-	Reserved	
9Fh	148-151	-	Restricted (OIF)	
9Fh	152-155	-	Custom	
9Fh	156-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	comp.
9Fh	135	RPLChkCode	See Table 8-179	comp.

¹ As FEC decoding usually straddles lanes of a Data Path, a module may either chose to assume that the estimated pre-FEC BER estimated is identical on all lanes of a Data Path or reject the request as not supported.

Page	Byte	Field Name	Description	Value
REPLY Data (LPL) (CMD 0212h)				
9Fh	136-255	LPL PM data (up to 120 bytes)	<p>PM record of one PM observable consists of 6 or 8 bytes:</p> <p>X16 minimum value</p> <p>X16 average (mean) value</p> <p>X16 maximum value</p> <p>X16 current value (if requested)</p> <p><i>Note: A maximum of 15 to 20 records can be returned using LPL depending on the requested PM record length.</i></p> <p>The sequence of PM records is the same as the sequence of set bits in the Observables field, in order of ascending significance.</p>	
REPLY Data (EPL) (CMD 0213h)				
A0h to AFh	128-255	EPL PM data (number of bytes depends on available Pages)	<p>PM record of one observable consists of 6 or 8 bytes:</p> <p>X16 minimum value</p> <p>X16 average (mean) value</p> <p>X16 maximum value</p> <p>X16 current value (if requested)</p> <p><i>Note: The maximum number of records depends on the size of the EPL and the requested PM record length.</i></p> <p>Data is contiguous across EPL Pages, and the sequence of PM records corresponds to the sequence of set bits in the Observables field, in order of ascending significance.</p>	

9.8.5 CMD 0214h/0215h: Get PM Media Side LPL/EPL

This command collects performance monitoring observables related to individual media lanes.

Table 9-36 CDB Command 0214h/0215h: Get PM Media Side LPL/EPL

Page	Byte	Field Name (Type)	Description	Value
CMD Header				
9Fh	128-129	CMDID (U16)	Get PM Media Side LPL CMD ID	0214h
			Get PM Media Side EPL CMD ID	0215h
9Fh	130-131	EPLLlength (U16)	EPL is not used	0
9Fh	132	LPLLlength	LPL length	20
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLlength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136.7	ClearOnRead	0b: return selected PM data 1b: return selected PM data and then reset their statistics	
	136.1-6	-	Reserved	
	136.0	RecordType	0b: Return 6-byte PM record (min, mean, max) 1b: Return 8-byte PM record (append "current" value)	
9Fh	137-139	-	Reserved for modules with more than 32 lanes.	
9Fh	140-143	Lanes	U32 Lanes is a bitmask indicating which media lane is present, where bit i represents lane i+1.	
9Fh	144	Selected PM data	Bit 0: Media Side SNR Bit 1: Media Side PAM4 LTP Bits 2-7: Reserved	
9Fh	145	Selected PM data	Bit 0: Tx Laser Bias Bit 1: Tx Power Bit 2: Rx Power Bit 3: Per-Lane Laser Temperature Bits 4-7: Reserved	
9Fh	146-147		Reserved	
9Fh	148-151		Restricted (OIF)	
9Fh	152-155		Custom	
9Fh	156-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLlength	See Table 8-179	comp.
9Fh	135	RPLChkCode	See Table 8-179	comp.

Page	Byte	Field Name (Type)	Description	Value
REPLY Data (LPL) (CMD 0214h)				
9Fh	136-255	LPL PM data (up to 120 bytes)	PM record of one observable consists of 6 or 8 bytes: X16 minimum value X16 average (mean) value X16 maximum value X16 current value (if requested) <i>Note: A maximum of 15 to 20 records can be returned using LPL, depending on the requested PM record length.</i> The sequence of PM records is the same as the sequence of set bits in the Observables field, in order of ascending significance.	
REPLY Data (EPL) (CMD 0215h)				
A0h to AFh	128-255	EPL PM data (maximum number of bytes depends on available EPL Pages)	PM record of one PM parameter consists of 6 or 8 bytes: X16 minimum value X16 average (mean) value X16 maximum value X16 current value (if requested) <i>Note: The maximum number of records depends on the size of the EPL and the requested PM record length.</i> Data is contiguous across EPL Pages, and the sequence of PM records corresponds to the sequence of set bits in the Observables field, in order of ascending significance.	

9.8.6 CMD 0216h/0217h: Get Data Path PM LPL/EPL

This command collects performance monitoring observables related Data Paths.

Table 9-37 CDB Command 0216/0217h: Get Data Path PM LPL/EPL

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get PM Data Path LPL CMD ID Get PM Data Path EPL CMD ID	0216h 0217h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	20
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136.7	ClearOnRead	0b: return selected PM data 1b: return selected PM data and then reset their statistics	
	136.1-6	-	Reserved	
	136.0	RecordType	0b: Return 6-byte PM record (min, mean, max) 1b: Return 8-byte PM record (append "current" value)	
9Fh	137-139	-	Reserved for modules with more than 32 lanes.	
9Fh	140-143	DataPaths	U32 DataPaths is a mask indicating which Data Path is present, where bit i represents the Data Path with DataPathID (lowest lane number) i+1	
9Fh	144	Observables	Bit 0: Media Side Frame Error Count (FERC) Bit 1: Media Side Pre-FEC BER Bits 2-7: Reserved	
9Fh	145-147	-	Reserved	
9Fh	148-151	-	Restricted (OIF)	
9Fh	152-155	-	Custom	
9Fh	156-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	comp.
9Fh	135	RPLChkCode	See Table 8-179	comp.
Returned Data Path Lane PM Using LPL (0216h)				
9Fh	136-255	LPL PM data (up to 120 bytes)	PM record of one observable consists of 6 or 8 bytes: X16 minimum value X16 average (mean) value X16 maximum value X16 current value (if requested) <i>Note: A maximum of 15 to 20 records can be returned using LPL depending on the requested PM record length.</i> The sequence of PM records data is the same as the sequence of set bits in the Observables field, in order of ascending significance.	

Page	Byte	Field Name	Description	Value
Returned Data Path Lane PM Using EPL (0217h)				
A0h to AFh	128-255	EPL PM data (maximum number of bytes depends on available EPL Pages)	PM record of one observable consists of 6 or 8 bytes: X16 minimum value X16 average (mean) value X16 maximum value X16 current value (if requested) <i>Note: The maximum number of records depends on the size of the EPL and the requested PM record length.</i> Data is contiguous across EPL Pages and the sequence of PM records corresponds to the sequence of set bits in the Observables field, in order of ascending significance.	

1 **9.8.7 CMD 0220h: Get Data Path RMON Statistics**

2 **Table 9-38 CDB Command 0220h: Get Data Path RMON Statistics**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Data Path RMON Statistics CMD ID	0220h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	2
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136.7	ClearOnRead	0b: return selected PM data 1b: return selected PM data and then reset their statistics <i>Note: atomic read-reset-resume (i.e. gap-free) counting is strongly preferred.</i>	
9Fh	136.7-1	-	Reserved	0
9Fh	136.0	MonitorLocation	0: media side, Rx direction 1: host side, Tx direction	
9Fh	137	DPID	U8: Data Path ID for the Data Path of interest.	
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	See Table 8-179	120
9Fh	135	RPLChkCode	See Table 8-179	comp.
REPLY Data (LPL)				
9Fh	136-141	FrameCount	U48: Saturating counters	
9Fh	142-147	OctetCount		
9Fh	148-153	BadFrameCount		
9Fh	154-159	BadOctetCount		
9Fh	160-165	MulticastCount		
9Fh	166-171	BroadcastCount		
9Fh	172-177	Packets64B		
9Fh	178-183	Packets64to127B		
9Fh	184-189	Packets128to255B		
9Fh	190-195	Packets256to511B		
9Fh	196-201	Packets512to1023B		
9Fh	202-207	Packets1024to1518B		
9Fh	208-211	PacketsLargeNonJumbo		
9Fh	212-215	PacketsJumbo		
9Fh	216-219	BadMulticastCount		
9Fh	220-223	BadBroadcastCount		
9Fh	224-227	BadPackets64B		
9Fh	228-231	BadPackets64to127B		
9Fh	232-235	BadPackets128to255B		
9Fh	236-239	BadPackets256to511B		
9Fh	240-243	BadPackets512to1023B		
9Fh	244-247	BadPackets1024to1518B		
9Fh	248-251	BadPacketsLargeNonJumbo		
9Fh	252-255	BadPacketsJumbo		

9.8.8 CMD 0230h: Control FEC Symbol Error Weight Histogram

This command resets, starts, or stops a given type of FEC Error Correction Statistics for a given FEC location for a given Data Path.

Table 9-39 CDB Command 0230h: Control FEC Symbol Error Weight Histogram

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Control FEC Symbol Error Weight Histogram CMD ID	0230h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	2
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	DPID	U8: Data Path ID of the Data Path to be monitored.	
9Fh	137.0	DecoderLocation	0: media side 1: host side	
9Fh	137.2-1	Command	0: reset (clear stats) 1: start (clear stats and start collecting) 2: stop (stop collecting)	
9Fh	137.7-3	-	Reserved: 0	
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header and Data (LPL)				
9Fh	134	RPLLength	Encoded Length and Check Code for REPLY message body in LPL or EPL. See Table 8-179	0
9Fh	135	RPLChkCode		comp.
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

9.8.9 CMD 0231h: Get FEC Symbol Error Weight Histogram

This command provides the histogram of Symbol Error Weights (symbol error corrections per FEC frame) for a given Data Path, with FEC decoder either on media side or on host side.

The statistics collection should have been started before, to get valid current data. The command can be run without having stopped the statistics.

Table 9-40 CDB Command 0231h: Get FEC Symbol Error Weight Histogram

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get FEC Symbol Error Weight Histogram CMD ID	0231h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	2
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	DPID	U8: Data Path ID of the Data Path of interest	
9Fh	137.0	DecoderLocation	0: media side 1: host side	
9Fh	137.7-1	-	Reserved: 0	0
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	Encoded Length and Check Code for REPLY message body in LPL or EPL. See Table 8-179	120
9Fh	135	RPLChkCode		comp.
REPLY Data (LPL)				
9Fh	136	-	Unmodified CMD parameter	
9Fh	137	-	Unmodified CMD parameter	
9Fh	138	CorrectionCapability	U8 Number of FEC correctable symbols per code word	
9Fh	139	-	Reserved (pad for alignment)	
9Fh	140-145	UncorrectableCount	U48: Number of uncorrectable FEC frames	
9Fh	146-151	ErrorWeight0Count	U48: Number of error free FEC frames	
9Fh	152-157	ErrorWeight1Count	U48: Number of FEC frames with 1 corrected symbol	
9Fh	158-163	ErrorWeight2Count	U48: Number of FEC frames with 2 corrected symbols	
9Fh	164-169	ErrorWeight3Count	U48: Number of FEC frames with 3 corrected symbols	
9Fh	170-175	ErrorWeight4Count	U48: Number of FEC frames with 4 corrected symbols	
9Fh	176-181	ErrorWeight5Count	U48: Number of FEC frames with 5 corrected symbols	
9Fh	182-187	ErrorWeight6Count	U48: Number of FEC frames with 6 corrected symbols	
9Fh	188-193	ErrorWeight7Count	U48: Number of FEC frames with 7 corrected symbols	
9Fh	194-199	ErrorWeight8Count	U48: Number of FEC frames with 8 corrected symbols	
9Fh	200-205	ErrorWeight9Count	U48: Number of FEC frames with 9 corrected symbols	
9Fh	206-211	ErrorWeight10Count	U48: Number of FEC frames with 10 corrected symbols	
9Fh	212-217	ErrorWeight11Count	U48: Number of FEC frames with 11 corrected symbols	
9Fh	218-223	ErrorWeight12Count	U48: Number of FEC frames with 12 corrected symbols	
9Fh	224-229	ErrorWeight13Count	U48: Number of FEC frames with 13 corrected symbols	
9Fh	230-235	ErrorWeight14Count	U48: Number of FEC frames with 14 corrected symbols	
9Fh	236-241	ErrorWeight15Count	U48: Number of FEC frames with 15 corrected symbols	

Page	Byte	Field Name	Description	Value
9Fh	242-247	ErrorWeight16Count	U48: Number of FEC frames with 16 corrected symbols	
9Fh	248-253	HighErrorWeightCount	U48: FEC frames with 17 or more corrected symbols <i>Note: This collector bin is for (limited) support of FEC schemes with higher symbol error correction capabilities.</i>	
9Fh	254-255	BER	F16: Pre-FEC BER determined by counting (or estimating) the bit errors in the symbol errors that have been counted into the error weight histogram, excluding the unknown bit errors that occur in uncorrectable FEC frames.	

9.8.10 CMD 0232h: Control Max FEC Symbol Error Weight

This command resets, starts, or stops a given type of FEC Error Correction Statistics for a given FEC location for a given Data Path.

Table 9-41 CDB Command 0232h: Control Max FEC Symbol Error Weight

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Control Max FEC Symbol Error Weight CMD ID	0232h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	2
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	DPID	U8: Data Path ID of the Data Path to be monitored.	
9Fh	137.0	DecoderLocation	0: media side 1: host side	
9Fh	137.2-1	Command	0: reset (clear stats) 1: start (clear stats and start collecting) 2: stop (stop collecting)	
9Fh	137.7-3	-	Reserved	0
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header and Data (LPL)				
9Fh	134	RPLLength	Encoded Length and Check Code for REPLY message body in LPL or EPL. See Table 8-179	0
9Fh	135	RPLChkCode		0
REPLY Data (LPL)				
9Fh	136-255	-	No data returned. Content not specified.	undef.

5

6

1 **9.8.11 CMD 0233h: Get Max FEC Symbol Error Weight**

2 This command provides the Maximum FEC Symbol Corrections per Frame high water mark for all Data Paths,
 3 reported against the first lane of each Data Path.

4 The statistics collection should have been started before, to get valid current data. The command can be run
 5 without having stopped the statistics.

6 **Table 9-42 CDB Command 0233h: Get Max FEC Symbol Error Weight**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Max FEC Symbol Error Weight CMD ID	0233h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	2
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	-	Reserved	0
9Fh	137.0	DecoderLocation	0: media side 1: host side	
9Fh	137.7-1	-	Reserved	0
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header and Data (LPL)				
9Fh	134	RPLLength	Encoded Length and Check Code for REPLY message body in LPL or EPL. See Table 8-179	20
9Fh	135	RPLChkCode		comp.
REPLY Data (LPL)				
9Fh	136	-	Unmodified	
9Fh	137	-	Unmodified CMD parameter	
9Fh	138-139	CorrectionCapability	U16 Number of correctable FEC symbols per FEC frame	
9Fh	140-141	MaxErrorWeightDPID1	U16 MaxErrorWeightDPID<DPID> := Maximum FEC symbol error weight observed on the Data Path with DPID = <DPID> or zero if no such Data Path exists. When there were uncorrectable FEC frames, a value exceeding CorrectionCapability should be reported.	
9Fh	142-143	MaxErrorWeightDPID2		
9Fh	144-145	MaxErrorWeightDPID3		
9Fh	146-147	MaxErrorWeightDPID4		
9Fh	148-149	MaxErrorWeightDPID5		
9Fh	150-151	MaxErrorWeightDPID6		
9Fh	152-153	MaxErrorWeightDPID7		
9Fh	154-155	MaxErrorWeightDPID8		
9Fh	156-255	-	No data returned. Content not specified.	undef.

1 9.9 CDB Data Monitoring and Recording Commands

2 Table 9-43 CDB Data Monitoring and Recording Commands Overview

ID	Command Title	Description	Type	Section
0280h	Data Monitoring and Recording Controls		Rqd.	9.9.1
0281h	-	Reserved (for Advertising)		9.9.2
0282h-02FFh	-	Reserved		
0290h	Temperature Histogram		Opt.	9.9.3
0291h-02FFh	-	Reserved		

3

4

5 9.9.1 CMD 0280h: Data Monitoring and Recording Controls

6 Table 9-44 CDB Command 0280h: Data Monitoring and Recording Controls

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Data Monitoring Controls CMD ID	0280h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	4
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	-	Reserved	0
9Fh	137	-	Reserved	0
9Fh	138.1-7	-	Reserved. Set to 0	0
	138.0	ClearAllStatistics	0b: No operation. 1b: Clear all statistics (minimum, average, maximum) for all monitored observables for all lanes, at the same time (best-effort).	
9Fh	139	-	Reserved	0
9Fh	140-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header and Data (LPL)				
9Fh	134	RPLLength	See Table 8-179	
9Fh	135	RPLChkCode	See Table 8-179	
9Fh	136-255	-	No data returned. Content not specified.	undef.

7

9.9.2 CMD 0281h: Data Monitoring and Recording Advertisement

Reserved for advertisement of the features defined in CMDs 0280-029Fh

Table 9-45 CDB Command 0281h: Data Monitoring and Recording Advertisements

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID		0281h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	7Ch
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	-		
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header and Data (LPL)				
9Fh	134	RPLLength	See Table 8-179	
9Fh	135	RPLChkCode	See Table 8-179	
REPLY Data (LPL)				
9Fh	136-255	-		

9.9.3 CMD 0290h: Temperature Histogram

This is an optional feature of a module to track the number of seconds in nonvolatile memory, in which the module has operated within the defined temperature bins, since the last time the host or user has cleared the temperature histogram.

The sampling interval and the interval at which the module writes to nonvolatile memory is module dependent. For the host not to lose histogram data, the host may send a command to save the current histogram to NVR before decommissioning or unplugging a module from its slot.

The timing accuracy of the internal clock is not defined here.

Table 9-46 CDB Command 0290h: Temperature Histogram

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Temperature Histogram CMD ID	0290h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL length	1
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	SubCommands	Bit 0: Reserved Bit 1: Save current histogram to NVR (host triggered) Bit 2: Return histogram (10degC bins) Bit 3-6: Reserved Bit 7: Clear Temperature Histogram (should be NV)	
9Fh	136-255	-	No data passed. Content not specified.	undef.

Page	Byte	Field Name	Description	Value
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	<p>In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution</p> <p>On Success 00 000001b: Success</p> <p>On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error</p>	
REPLY Header and Data (LPL)				
9Fh	134	RPLLength	Encoded Length and Check Code for REPLY message body in LPL or EPL. See Table 8-179	52
9Fh	135	RPLChkCode		comp.
REPLY Data (LPL)				
9Fh	136	SubCommands	Echo of host-written values	
9Fh	137	NHoursToNextWrite	0: Module does not indicate hours to next write to NVR. 1-254: Temperature Histogram will be written to NVR within this specified number of hours. <i>Note: Depends on device technology use, NVR writes may need to be triggered by host.</i> 255: NVR write needs to be triggered by host. If the module were to write NVR autonomously, some operations may share EEPROM or flash device and may cause smaller microcontrollers to stop responding to MCI while the autonomous write is in progress.	
9Fh	138-139	-	Reserved	0
9Fh	140-143	TotalSeconds	U32 Total time of temperature histogram accumulation (s)	
9Fh	144-147	TenDegBinBelowM5	U32 Total seconds operated with temp. below -5 degC.	
9Fh	148-151	TenDegBin0	U32 Total seconds operated in temp. interval [-5, 5] degC	
9Fh	152-155	TenDegBin10	U32 Total seconds operated in temp. interval [5, 15] degC	
9Fh	156-159	TenDegBin20	U32 Total seconds operated in temp. interval [15, 25] degC	
9Fh	160-163	TenDegBin30	U32 Total seconds operated in temp. interval [25, 35] degC	
9Fh	164-167	TenDegBin40	U32 Total seconds operated in temp. interval [35, 45] degC	
9Fh	168-171	TenDegBin50	U32 Total seconds operated in temp. interval [45, 55] degC	
9Fh	172-175	TenDegBin60	U32 Total seconds operated in temp. interval [55, 65] degC	
9Fh	176-179	TenDegBin70	U32 Total seconds operated in temp. interval [65, 75] degC	
9Fh	180-183	TenDegBin80	U32 Total seconds operated in temp. interval [75, 85] degC	
9Fh	184-187	TenDegBinAbove85	U32 Total seconds operated with temp. above 85 degC	
9Fh	188-255	-	No data returned. Content not specified.	undef.

9.10 CDB PRBS BERT Commands

See also features available using Pages 13h-14h.

Table 9-47 CDB BERT Commands Overview

ID	Command Title	Description	Type	Section
0300h	-	Reserved for PRBS related capabilities		
030xh	-	Reserved for PRBS Generator Config, Enable and Disable		
030xh	-	Reserved for PRBS Error Injector		
030xh	-	Reserved for PRBS Detector (Checker) Config, Enable and Disable		
030xh	-	Reserved for PRBS Error Counts		
030xh	-	Reserved for PRBS BER		
030xh-037Fh	-	Reserved		

4

1 9.11 CDB Diagnostics and Debug Commands

2 3 Table 9-48 CDB Diagnostics and Debug Commands Overview

ID	Command Title	Description	Type	Section
0380h	Loopbacks	Reserved for Loopbacks	Adv.	9.11.1
0390h	-	Reserved for PAM4 Histogram		
03A0h	-	Reserved for Eye Monitors		

7 9.11.1 CMD 0380h: Loopbacks

8 Basic loopback modes are available using features supported via Page 13h and Page 14h.

9 This CDB command is **reserved** for additional loopback capabilities.

10 Table 9-49 CDB Command 0380h: Loopbacks

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Loopbacks CMD ID	0380h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	7Ch
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	In Progress 10 000001b: Busy processing command, CMD captured 10 000010b: Busy processing command, CMD checking 10 000011b: Busy processing command, CMD execution On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000010b: Parameter range error or not supported 01 000101b: CdbChkCode error	
REPLY Header and Data (LPL)				
9Fh	134	RPLLength	Empty reply message body. See Table 8-179	0
9Fh	135	RPLChkCode		0
REPLY Data (LPL)				
9Fh	136-255	-		undef.

9.12 CDB Security Related Commands

Module capabilities related to security commands are advertised in the reply to CMD 0044h (see section 9.4.5)

Security related commands may carry larger payloads, such as certificates or signatures. Depending on advertised module capabilities, security related command and reply messages are exchanged either via **LPL** (segmented into several commands if necessary), or via **EPL** (complete data transferred in one CDB interaction).

When data need to be segmented, several command invocations with incremented segment index need to be issued to transmit or collect and reassemble all segments. The host can derive the number of segments needed from the advertised or given length of the data and the known maximum length of segment.

Since the module may host several certificates in a certificate issuance chain, the certificate instances in the chain need to be distinguished. This is done using an instance index, where the leaf certificate is identified by an index value of 0.

Table 9-50 CDB Diagnostics and Debug Commands Overview

ID	Command Title	Description	Type	Section
0400h	Get IDevID Certificate in LPL	<p>Read (segment of) the module's Initial Device Id (IDevID) certificate, or (segment of) a selected element of its certificate chain, in the advertised format (X.509 v3 or custom).</p> <p>The IDevID certificate may be an end point issued directly by a host-trusted Root CA, or a leaf certificate in an issuance chain of at most 4 certificates, the highest one of which is issued by a host-trusted Root CA.</p> <p><i>Note that RSA 3072-bit leaf certificates are ~1200 bytes and RSA 4096-bit certificates are ~1500 bytes without extensions. Smaller keys (and signatures) used in other algorithms result in smaller certificates. The CDB EPL 2048 byte is not expected to be a problem unless large keys with cert chains and/or extensions are utilized.</i></p>	Adv.	9.12.3
0401h	Get IDevID Certificate in EPL	This command variant retrieves an entire certificate in a single REPLY (no segmentation). The number of EPL pages must accommodate the largest certificate.	Adv.	
0402h	Set Digest To Sign given in LPL	<p>Provide a digest (fingerprint) for a random message to be signed by the module.</p> <p><i>Note: This implements the challenge portion of a Challenge-Response interaction, where the host creates a random message and then applies the hash function that is characteristic for the advertised DigestLength (as returned by CMD 0044h) to create a message digest (fingerprint) that the module must sign.</i></p> <p>CMD 0044h SignatureTime indicates the maximum time required to complete the signing operation.</p>	Adv.	
0403h	Set Digest To Sign given in EPL	This command variant transmits the message digest to be signed in a single CMD (no segmentation). The number of EPL pages must be sufficient to accommodate the largest message digest.	Adv.	
0404h	Get Digest Signature in LPL	<p>Read the module's signature of the most recently provided digest to be signed</p> <p><i>Note: This implements the response portion of the Challenge-Response interaction, which allows the module to prove to the host that it owns the private key associated with the provided module certificate.</i></p> <p><i>The encoding of the signature is defined by CMD 0044h SignatureFormat. The signature algorithm is specified in the module certificate.</i></p> <p>Repeated Get Digest Signature commands will return the same signature until a new Set Digest To Sign command is issued.</p>	Adv.	

ID	Command Title	Description	Type	Section
0405h	Get Digest Signature in EPL	This command variant retrieves the entire Digest Signature in a single CMD. The number of EPL pages must be sufficient to accommodate the largest signature.	Adv.	
0406h -04FFh	-	Reserved for additional Security Related commands	Adv.	

1

2

Table 9-51 Detailed Return Status Codes for Security Commands

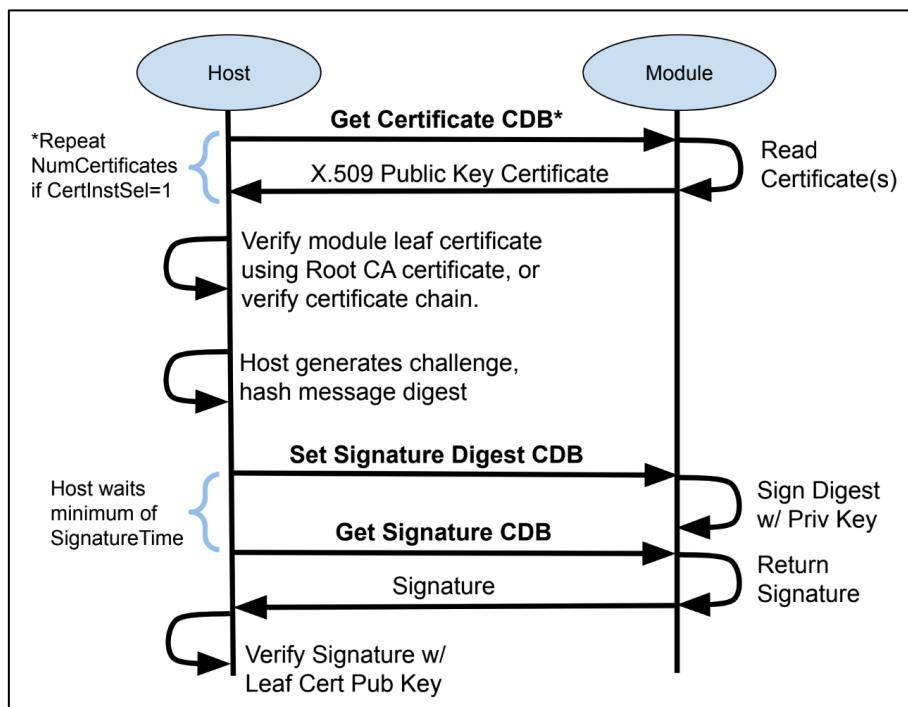
Value	Symbolic	Description
0000h	OK	Successful command completion
0001h	COMMUNICATION_ERROR	General internal communication error
0002h	UNEXPECTED_ERROR	General unspecific error
0010h	USER_COMMAND_DATA_ERROR	General command data error
0011h	USER_NO_DIGEST_ERROR	Signature request without digest supplied before
0012h	USER_NOT_READY_ERROR	Signature request during processing
0020h	SEC_INVALID_PRIVATE_KEY	Invalid private key
0021h	SEC_INVALID_PUBLIC_KEY	Invalid public key in certificate
0022h	SEC_DEVICE_NOT_FOUND	Security device not present or not operational
0023h	SEC_DEVICE_READ_ERROR	Security device read error
0024h	SEC_DEVICE_WRITE_ERROR	Security device write error
1000h-1FFFh ...	-	Custom return codes

3

4

9.12.1 Module Authentication Overview

Here is a high-level message exchange diagram for module authentication.

**Figure 9-1 Module Authentication Flow**

6

7

8

1 **9.12.2CMD 0400h: Get Initial Device ID Certificate in LPL**

2 **Table 9-52 CDB Command 0400h: Get Initial Device ID Certificate in LPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Initial Device ID Certificate CMD ID in LPL	0400h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL Length	2
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	CertificateIndex	U8: Identifies the certificate instance index to obtain the leaf certificate (CertificateIndex = 0) or another certificate in the chain (0 < CertificateIndex < 4).	
9Fh	137	SegmentIndex	U8: The host always starts reading a certificate with SegmentIndex=0 and issues further commands with incremented SegmentIndex if the module reply indicates in RPLLength that repeated calls are necessary. <i>Note: The total certificate length in bytes is known from the CertificateLen* fields retrieved via CMD 044h</i>	
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	Number of LPL bytes returned, as per Table 8-179.	
9Fh	135	RPLChkCode	Check code as specified in Table 8-179.	
REPLY Data (LPL)				
9Fh	136-137	ReturnStatus	U16: Detailed return status (see Table 9-51)	
9Fh	138	SegmentIndex	U8: The index of the segment returned; should be the same value as in the CMD Data	
9Fh	139-255	CertificateSegment	U8[117]: Certificate segment requested. The last segment may be shorter (RPLLength-3 < 117 bytes), with the remaining bytes undefined.	

1 9.12.3 CMD 0401h: Get Initial Device ID Certificate in EPL

2 Note: The reply to this command uses both LPL and EPL.

3 **Table 9-53 CDB Command 0401h: Get Initial Device ID Certificate in EPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Initial Device ID Certificate CMD ID in EPL	0401h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL Length	1
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	CertificateIndex	U8: Identifies the instance index of the certificate to obtain the leaf certificate (CertificateIndex = 0) or a certificate in the chain (0 < CertificateIndex < 4).	
9Fh	137-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	The number of EPL pages used for the certificate is encoded as specified in Table 8-179.	
9Fh	135	RPLChkCode	Check code as specified in Table 8-179.	
REPLY Data (LPL)				
9Fh	136-137	ReturnStatus	U16: Detailed return status (see Table 9-51)	
9Fh	138-139	CertificateLength	U16: Actual length (in bytes) of the requested Certificate as written to EPL, or 0 if there are insufficient EPL pages. <i>Note: The expected certificate lengths in bytes are known from the CertificateLen* fields retrieved via CMD 044h</i>	
REPLY Data (EPL)				
A0h-AFh	128-255	CertificateData	U8[128]: Complete certificate extending over one or more EPL pages (as needed). The last EPL page may be used only partially (CertificateLength mod 128 bytes), with the remaining bytes undefined.	

1 **9.12.4 CMD 0402h: Set Digest To Sign given in LPL**

2 **Table 9-54 CDB Command 0402h: Set Digest To Sign given in LPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Set Digest To Sign CMD ID given in LPL	0402h
9Fh	130-131	EPLLength	EPL is not used.	0.
9Fh	132	LPLLength	The given signature digest length in bytes (see also CMD 0044h DigestLength).	comp.
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136-255	Digest Data	U8[120]: The digest to sign by the module. The actual digest length depends on the hash function used for creating the digest and is given in LPLLength	
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	Number of LPL bytes returned, as per Table 8-179.	2
9Fh	135	RPLChkCode	Check code as specified in Table 8-179.	
REPLY Data (LPL)				
9Fh	136-137	ResultStatus	U16: Detailed return status (see Table 9-51)	
9Fh	138-255	-	No data returned. Content not specified.	undef.

3

1 **9.12.5 CMD 0403h: Set Digest To Sign given in EPL**

2 Note: The reply to this command uses LPL while the message digest to be signed is provided in EPL.

3 **Table 9-55 CDB Command 0403h: Set Digest To Sign given in EPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Set Digest To Sign CMD ID given in EPL	0403h
9Fh	130-131	EPLLength	Signature (message) digest length in bytes (see CMD 0044h DigestLength).	comp.
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (EPL)				
A0h	128-255	Digest Data	U8[128]: The digest to sign by the module. The actual digest length depends on the hash function used for creating the digest and is given in LPLLength	
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	Number of LPL bytes returned, as per Table 8-179.	2
9Fh	135	RPLChkCode	Check code as specified in Table 8-179.	
REPLY Data (LPL)				
9Fh	136-137	ResultStatus	U16: Detailed return status (see Table 9-51)	
9Fh	138-255	-	No data returned. Content not specified.	undef.

1 **9.12.6 CMD 0404h: Get Digest Signature in LPL**

2 **Table 9-56 CDB Command 0404h: Get Digest Signature in LPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Digest Signature CMD ID in LPL	0404h
9Fh	130-131	EPLLength	EPL is not used	0000h
9Fh	132	LPLLength	LPL length	1
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	7Ch
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply.</i> See Table 8-178	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	136	SegmentIndex	U8: The host always starts reading a signature with SegmentIndex=0 and issues further commands with incremented SegmentIndex. <i>Note: The total signature length in bytes is known from CMD 044h SignatureLength, so the host knows how many segments to request</i>	undef.
9Fh	137-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	Number of LPL bytes returned. See Table 8-179.	
9Fh	135	RPLChkCode	Check code as specified in Table 8-179.	comp.
REPLY Data (LPL)				
9Fh	136-137	ReturnStatus	U16: Detailed return status (see Table 9-51)	
9Fh	138	SegmentIndex	U8: The index of the segment returned; should be the same value as in the CMD Data	
9Fh	139-255	SignatureSegment	U8[117]: Signature segment requested. The last segment may be shorter (RPLLength-3 < 117 bytes), with the remaining bytes undefined.	

3

1 **9.12.7 CMD 0405h: Get Digest Signature in EPL**

2 Note: The reply to this command uses both LPL and EPL.

3 **Table 9-57 CDB Command 0405h: Get Digest Signature in EPL**

Page	Byte	Field Name	Description	Value
CMD Header				
9Fh	128-129	CMDID	Get Digest Signature CMD ID in EPL	0405h
9Fh	130-131	EPLLength	EPL is not used	0
9Fh	132	LPLLength	LPL is not used	0
9Fh	133	CdbChkCode	Check Code over 9Fh:128-132 and LPL. See Table 8-178	Comp.
9Fh	134	RPLLength	<i>Note: Initiator may fill those reply fields, to later verify field updates by the target in the reply. See Table 8-178</i>	undef.
9Fh	135	RPLChkCode		undef.
CMD Data (LPL)				
9Fh	138-255	-	No data passed. Content not specified.	undef.
REPLY Status				
00h	8.6 or 8.7	CdbCmdCompleteFlag	Set by module when the CDB command is complete.	1
00h	37 or 38	CdbStatus	On Success 00 000001b: Success On Failure 01 000000b: Failed, no specific failure 01 000101b: CdbChkCode error	
REPLY Header				
9Fh	134	RPLLength	The number of EPL pages used for the returned signature is encoded as specified in Table 8-179.	4
9Fh	135	RPLChkCode	Check code as specified in Table 8-179.	comp.
REPLY Data (LPL)				
9Fh	136-137	ReturnStatus	U16: Detailed return status (see Table 9-51)	
9Fh	138-139	SignatureLength	U16: length (in bytes) of the requested Signature as written to EPL, or 0 if there are insufficient EPL pages. <i>Note: The expected signature length in bytes is known from the SignatureLength field retrieved via CMD 044h</i>	
9Fh	140-255	-	No data returned. Content not specified.	undef.
REPLY Data (EPL)				
A0h-AFh	128-255	SignatureData	U8[128]: Returned signature data extending over one or more EPL pages (as needed). The last EPL page may be used only partially (SignatureLength mod 128 bytes), with the remaining bytes undefined.	

10 Management Timing Specifications

2 This chapter provides fixed, unadvertised timing specifications for

- 3 • management control signals (Management Signal Layer)
- 4 • register access and register interdependencies (Register Access Layer)
- 5 • interdependencies between signal conditions and the register map
- 6 • interdependencies between high-speed signal conditions

7 In certain cases, advertisement registers may exist that allow a module to advertise tighter timing specifications.

8 *Note: The timing specifications in this chapter are intentionally limited to timing specifications related to module
9 management, management registers, or management control signals. Timing specifications for the underlying
10 Management Communication Interface (MCI) are defined in Appendix B and must be consistent with timings
11 specified here.*

12 *Note: The timings between high-speed signals conditions are listed here because they are relevant for a host
13 managing the module.*

14 10.1 Timings for Management Control Signals (MSL)

15 Table 10-1 defines timing parameters and requirements for management control signal waveforms.

16 Table 10-2 defines timing parameters and requirements related to the latency of effects caused by management
17 control signals or by other host or operator actions (power supply, hot plug).

18 **Table 10-1 Signal Waveform Timings**

Parameter	Symbol	Min	Unit	Conditions
Min Reset Assert Duration	t_reset_init	10	us	Minimum pulse duration of the Reset signal to initiate a module reset

19 **Table 10-2 Effect Latency Timings**

Parameter	Symbol	Max	Unit	Conditions
Max Time to reach manageability in ModuleLowPwr	tMgmtInit	2000	ms	Time from power-on, hot plug, or Reset release until the START condition of a READ retrieving the default register value of an arbitrary register (including a page switch).

21 Note: Power on and hot plug events are defined as the instant when supply voltages reach and remain at or
22 above the minimum electrical level specified in the form factor-dependent specification. The Reset release event
23 is defined as the time of the most recent DEASSERTING edge of a correctly ASSERTED Reset signal.

10.2 Timings for Register Access (RAL)

2 Register ACCESS related timings are defined with reference to observable events at the Management
3 Communication Interface port of the module.

4 *Note: Future CMIS versions should seek specification methods that are independent of the MCI variant used.*

5 10.2.1 Single ACCESS Timings

6 Table 10-3 defines timing parameters and requirements related to one individual register ACCESS.

7 *Note: Refer to Appendix B for MCI implementation dependent timings related to register access*

8 **Table 10-3 Timings for Register Access**

Parameter	Symbol	Max	Unit	Conditions
Max READ Latency Variation	tREAD	500	μs	Allowable Latency Variation, i.e. the maximum time the module may internally delay returning the requested data in an MCI transaction for a READ, beyond the fastest transaction execution. <i>Note: This RAL parameter is an abstraction of the I2CMCI clock stretch time</i>

9 10.2.2 Consecutive ACCESS Timings

10 This subsection defines timing requirements involving two adjacent register accesses with the possibility of an
11 ACCESS hold-off period after the first ACCESS.

12 An **ACCESS hold-off period** is a period of time during which a module **rejects** READ or WRITE accesses for
13 reasons specified in the relevant sections of this specification (see sections 5.2.3 and 5.2.4).

14 The **actual duration** of an ACCESS hold-off period is defined as the **time interval between** the event when
15 the module receives the **I2CMCI STOP** token of the MCI transaction representing the WRITE causing the begin
16 of the ACCESS hold-off period **and** the event when the module sends the **I2CMCI ACK** token for the I2CMCI
17 command byte in the first MCI transaction representing a non-rejected ACCESS.

18 *Note: A host not willing to wait for specified maximum durations can retry a rejected ACCESS or use the TEST
19 primitive (defined in section 5.2.2.3) to poll if earlier ACCESS is possible.*

20 Table 10-4 defines the **maximum duration** of each type of ACCESS hold-off period.

21 **Table 10-4 Maximum ACCESS Hold-Off Durations**

Parameter	Symbol	Max	Unit	Conditions
Max WRITE Latency for regular (volatile) register or memory <i>Note: Almost all CMIS memory locations are classified as volatile.</i>	tWRITE	10	ms	Maximum time for the result of a WRITE to be completed (written value retrieved in a READ), expressed as the longest ACCESS hold-off period after a WRITE to an address representing volatile memory or management register. <i>Note: This RAL parameter is an abstraction of the I2CMCI parameter tNACK.</i>
Max WRITE Latency for NV (persistent) register or memory	tWRITENV	80	ms	Maximum time for the result of a WRITE to be completed (written value retrieved in a READ), expressed as the longest ACCESS hold-off period after a WRITE to an address representing non-volatile memory or management register. <i>Note: This RAL parameter is an abstraction of the I2CMCI parameter tWR.</i>
Maximum Bank/Page Change Time	tBPC	10	ms	Maximum duration of the ACCESS hold-off after a PageMapping register change (see section 8.2.15)
Maximum CDB command foreground processing completion time	tCDBF	4960	ms	Maximum advertised length of the ACCESS hold-off period while module processes a CDB command in foreground processing mode. <i>Note: The maximum CDB busy time defaults to 80 ms but the module may advertise a different</i>

				<i>value ranging from 0 to 4960 ms (see section 8.4.11)</i>
Maximum CDB command capture time	tCDBC	80	ms	Maximum advertised length of the ACCESS hold-off period while module captures a CDB command for subsequent execution.

10.2.3 Register Content Dependencies

This section defines timings for cause-to-effect latencies of changes in certain registers to become visible in dependent registers, especially for bulk data.

The module does not prevent access to stale data in these cases (i.e. ACCESS that is too early is not **rejected**).

Measurement condition: The effect latency of a WRITE to a register causing an effect in specified dependent registers is the **time interval between** the event when the module receives the **I2CMCI STOP** token of the MCI transaction representing the WRITE **and** the event when the module sends the **I2CMCI ACK** token for the I2CMCI command byte in the first MCI transaction representing a READ of an arbitrary byte in the dependent registers that retrieves non-stale, correct data.

Table 10-5 Content Dependency Timings

Parameter	Symbol	Max	Unit	Conditions
Maximum Diagnostics Data Content Switch Time	tDDCS	10	ms	Maximum time from STOP of the WRITE access to the Diagnostics Data Selector (see section 8.13.) to the START of a READ access retrieving correct data
Maximum VDM Statistics Freeze Time	tVDMF	10	ms	Maximum time required by the module to freeze all VDM statistics reporting registers after the STOP of the Freeze request WRITE (see section 8.22.6)

Note: In future versions of this specification, timings may be added for other interdependencies such as Flag summary registers (dependent on Flag updates), statis registers dependent on State Machine state changes, or register effects depending on hardware reconfigurations being completed, etc.

10.3 Timings between Conditions and Management Registers or Signals

This section describes timing requirements for interactions between signals and the management register map.

10.3.1 Interrupt and Flag Related Timings

Note: The relative timing behavior of different but correlated management information representations such as Interrupt request signal, and Flag and Flag bits, is largely unspecified and can vary significantly across modules.

The condition interrupt signal assertion or de-assertion refers to the voltage of the interrupt signal crossing the voltage threshold associated with assertion or de-assertion, respectively.

Table 10-6 Condition to Interrupt Timings

Parameter	Symbol	Max	Unit	Conditions
Max Interrupt Assert Time	ton_Int	200	ms	Time from onset of condition or occurrence of event until associated unmasked Interrupt asserted
Max Rx LOS Interrupt Assert Time	ton_los	100	ms	Time from onset of Rx LOS condition to Rx LOS Flag raised and Interrupt asserted
Max Rx LOS Interrupt Assert Time (optional fast mode)	ton_losf	1	ms	
Max Tx Failure Interrupt Assert Time	ton_Txfail	200	ms	Time from Tx Failure state to Tx Failure Flag raised and Interrupt asserted
Max Flag Assert Time	ton_flag	200	ms	Time from onset of condition or occurrence of event to associated Flag bit raised and Interrupt asserted

Table 10-7 Register to Interrupt Timings

Parameter	Symbol	Max	Unit	Conditions
Max Interrupt Deassert Time	toff_Int	0.5	ms	Time from reading (clear on read) last Flag set until Interrupt deasserted, assuming that no conditions nor events causing a Flag setting are present
Max Mask Assert Time	ton_mask	100	ms	Time from Mask bit raised while associated Flag bit is set until Interrupt deasserted
Max Mask Deassert Time	toff_mask	100	ms	Time from Mask bit ceased while associated Flag bit is set until Interrupt asserted

10.3.2 High Speed Signal Related Timings

Table 10-8 defines causal timing requirements between register access and high-speed signals.

Table 10-8 Register to High-Speed Signal Timings

Parameter	Symbol	Max	Unit	Conditions
Max Tx Output Disable Latency ¹	ton_txdis	100	ms	Time from setting a OutputDisableTx<i> bit until optical output falls below 10% of nominal.
Max Fast Mode Tx Output Disable Latency (optional) ¹	ton_txdisf	3	ms	For I2CMCI the time interval begins with the STOP token ending the MCI write transaction
Max Tx Output UnDisable Latency ¹	toff_txdis	400	ms	Time from clearing a OutputDisableTx<i> bit until an enabled optical output rises above 90% of nominal.
Max Fast Mode Tx Output UnDisable Latency (optional) ¹	toff_txdisf	10	ms	For I2CMCI the time interval begins with the STOP token ending the MCI write transaction
Max Rx Output Disable Assert Time	ton_rxdis	100	ms	Time from the termination of the Rx Output Disable register write sequence until an enabled Rx output falls below 10% of nominal
Max Rx Output Disable Deassert Time	toff_rxdis	100	ms	Time from the termination of the Rx Output Disable register write sequence until an enable Rx output rises above 90% of nominal
Max Auto Squelch Disable Assert Time	ton_sqdis	100	ms	This applies to Rx and Tx Auto Squelch and is the time from the termination of the register write sequence until auto squelch functionality is disabled.
Max Auto Squelch Disable Deassert Time	toff_sqdis	100	ms	This applies to Rx and Tx Auto Squelch and is the time from the termination of the register write sequence

Parameter	Symbol	Max	Unit	Conditions
				until auto squelch functionality is enabled.

Note 1: Values specified here place an upper limit on advertised timings like MaxDurationDPTxTurnOff and MaxDurationDPTxTurnOn (01h:168).

10.4 Timings between High-Speed Signal Conditions

Table 10-9 defines causal timing requirements between conditions of high-speed signals.

Table 10-9 Signal Condition to Signal Condition Timings

Parameter	Symbol	Max	Unit	Conditions
Max Rx Auto Squelch Reaction Latency	ton_Rxsq	15	ms	Time from onset of loss of Rx input signal condition until the squelched output condition is reached
Max Rx Auto UnSquelch Reaction Latency	toff_Rxsq	-	ms	Obsoleted: time for normal Rx output signal condition restored when loss of Rx input signal condition clears
Max Tx Auto Squelch Reaction Latency	ton_Txsq	400	ms	Time from onset of loss of Tx input signal condition until the squelched output condition is reached
Max Tx Auto UnSquelch Reaction Latency	toff_Txsq	-	ms	Obsoleted: time for normal Tx output signal condition restored when loss of Tx input signal condition clears

Note: These timings are not directly related to management registers or to management signals. They provide an upper bound for module timing behavior as may be assumed by a host managing a CMIS module. Form factor specific timing specifications may be tighter.

Appendix A Form Factor Specific MSL or MCI Signal Names

Table A-1 associates the generic names for management related signals that are used in this specification with form factor-specific signal names, for five form factor **examples**. These generic signals form the CMIS Management Signaling Layer (MSL) core (see section 4.2 and section 5.1) or they are part of a CMIS specified Management Communication Interface (MCI) variant (see section 4.2 and Appendix B).

Other specifications may define form factor specific extensions to the CMIS MSL core (see section 2.1.2), while management memory map extensions for the management (advertisement, administration, reporting) of these signals will be found on Page 05h (see section 8.8), which is restricted to this purpose.

Table A-1 Form Factor Dependent Signal Name Associations

CMIS Signal Name	Layer	QSFP-DD	QSFP	SFP112	OSFP	COBO
Reset	MSL	ResetL	ResetL	ResetL	RSTn	ResetL
Interrupt		IntL	IntL	IntL	INT	IntL
LowPwrRequestHW		LPMode	LPMode	- ¹	LPWn	LPMode
ModSel		I2CMCI	ModSell	ModSell	ModSell	ModSell

These generic CMIS hardware signals assume two **symbolic** signal levels, ASSERTED and DEASSERTED (section 2.3.2), in order to address realizations both in active-low logic and in active-high logic.

Note: Active-high voltage signals are sometimes equated, silently, as positive logic bits with a representation of TRUE=1 and FALSE=0, while active-low voltage signals are equated as negative or inverted logic bits, with TRUE=0 and FALSE=1. These silent associations can be confusing and are therefore avoided by introducing symbolic signal levels with a fixed association to the truth values TRUE and FALSE.

The correspondence between the symbolic signal levels ASSERTED and DEASSERTED, and the voltage levels of a given logic convention used in the hardware specification of a form factor is defined in the following table:

Table A-2 Symbolic Logical Signal Values

CMIS Symbolic Value Signal Level	Active-High Logic Voltage Level	Active-Low Logic Voltage Level	Positive Bit Logic CMIS Bool Type	Boolean Logic Truth Values
ASSERTED	High	Low	1	TRUE
DEASSERTED	Low	High	0	FALSE

*Note: Boolean logic expressions and equations in CMIS are mostly formulated in terms of numerical **bit values** 0 and 1, thereby assuming a positive logic representation (TRUE=1 and FALSE=0). Mapping conventions defined in section 2.3.2 allow the specification to use the abstract signal levels ASSERTED and DEASSERTED for hardware signals together with bit values in those pseudo-Boolean expressions.*

¹ Note: Modules without a LowPwrRequestHW signal input behave as if the missing signal was ASSERTED.

1 Appendix B Management Communication Interface Definitions

2 This appendix contains specification variants for the Management Communication Interface (MCI) for use by
 3 reference in module hardware (form factor) specifications.

4 Relationship to Hardware Specifications

5 The MCI variant used by a particular module is generally assumed to be specified in a form-factor-dependent
 6 hardware specification. Please refer to the relevant form factor hardware specification to determine which MCI
 7 variant is employed for a given case.

8 *Note: This documentation structure inverts the referential dependency of previous CMIS revisions. In the new
 9 structure the hardware specifications refer to CMIS when CMIS management is to be supported.*

10 Relationship to other parts of the CMIS Specification

11 Each MCI variant specified in this appendix provides the necessary data transfer services for the basic READ,
 12 WRITE, and TEST register access operations specified in chapter 5.

13 *Note: This separation between the register access layer (see section 5.2) and the underlying management
 14 communication interface in this specification does not need to be visible in an implementation.*

15 *Note: The main purpose of separating data transfer mechanism across the MCI and the basic CMIS management
 16 operations (READ from or WRITE to management addressable memory in a module) is to allow the CMIS
 17 management application layer to be easily migrated to other physical communication interfaces¹, in the future.*

18 MCI Variants Available

19 In this CMIS revision, the following MCI variants are available:

20 **Table B-1 Management Communication Interface Variants**

MCI Variant	Description	Specification
I2CMCI	The initial I2C-based MCI specification for CMIS <i>Note: This was called TWI in previous versions</i>	Section B.2
SPIMCI	An SPI based MCI with addressing improvements	Section B.3

21
 22 *Note: New MCI variants may be added to future CMIS revisions on request, for form factors that intend to
 23 implement CMIS based management over a new, e.g. higher bandwidth, communication infrastructure.*

24 B.1 Generic Definitions

25 A **management communication interface** provides transactional data transfer services to implement the
 26 basic management register access operations specified in chapter 5: READ, WRITE, TEST.

27 *Note: These transactional data transfer services are sometimes called bus transactions, for short.*

28 The data transfer occurs between a managing host and a managed module, which take the roles of initiator
 29 and target, respectively, in the neutral data communication perspective of the MCI.

30 B.1.1 Communication Roles and Transactions

31 The management communication interface distinguishes two roles.

32 An **initiator** initiates and terminates a **transaction** for data transfer.

33 A **target** responds to the initiator within a transaction.

34 A transaction to implement a READ access is called a read transaction.

35 A transaction to implement a WRITE access is called a write transaction.

36 Within a transaction data can be transferred from the initiator to the target and from the target to the initiator.

¹ Editor's note: It might be argued that the low abstraction level of the CMIS management operations may not be ideal for the management of future modules with increasingly complex functionality. Even if that is, hypothetically, assumed to be true, there is certainly value in allowing higher bandwidth CMIS implementations, both for porting existing management applications and for expanding the use of CDB messaging.

1 *Note: In CMIS the host is always the initiator, and the module is always the target.*

2 *Note: Completion of a data transfer transaction and associated register or memory modifications in the module
3 are not necessarily synchronous: internal modifications may occur only after the transaction is terminated by
4 the initiator. Any necessary synchronization conditions are specified in section 5.2.4.*

5 B.1.2 Current Byte Address

6 *Note: This section, strictly speaking, speaks about local and **byte-sized** byte addresses, i.e. addresses just
7 within the small 128 byte address space of a (possibly banked) page. These byte-sized byte addresses are
8 unique only locally on a given or assumed (banked) page, while the globally unique address of a given or
9 assumed byte in management memory requires all three addressing components: bank index, page index, byte
10 index. This distinction becomes important with MCI variants that address bytes globally and uniquely rather
11 than locally, within the context of a previously set current (banked) page.*

12 The target maintains one internal **current byte address** variable containing the Byte address of the Byte to
13 be accessed in the next access in case that no address is provided in that next access. When an address is
14 specified in the access, the current byte address is set to that address.

15 The target increments the current byte address by one after each byte read or byte written and rolls over
16 (wraps around) after incrementing past the end of the current 128-byte memory area (see section 8.1).

17 Roll-over in Lower Memory sets the current byte address to 0 when the byte-sized byte address to be
18 incremented is 127.

19 Roll-over in Upper Memory sets the current byte address to 128 when the byte-sized byte address to be
20 incremented is 255.

21 *Note: The current byte address roll-over may be combined with other mechanisms in certain (optional) address
22 ranges, such as auto paging in the CDB page range described in Section 7.2.*

23 The current byte address remains valid between two register accesses as long as power to the target is
24 maintained and as long as no MCI protocol violations occur.

25 The current byte address is indeterminate upon loss of power to or reset of the module (i.e. after initialization),
26 or upon an access failure (e.g. I2CMCI transactions not terminated by a STOP condition).

B.2 I2C-Based Management Communication Interface (I2CMCI)

The I2CMCI specification for communication between an initiator and a target is a tailored and simplified variant of the general I2C-bus [1] specification.

B.2.1 Communication Topology

The I2CMCI protocol for data transfer appears as a point-to-point interaction between initiator and target because the protocol frame does not provide for address information (see section B.2.4).

The physical management communication topology in a host system might still be a multi-point bus when host and module support an optional generic hardware control signal for module selection (ModSel).

B.2.2 I2CMCI Control Signals

The following discrete hardware control signal is not a required part of the I2CMCI but needs to be accounted for when present

- a module selection signal (**ModSel**) instructing the module to activate or passivate its management interface, in order to effectively allow MCI bus sharing (when such precautions are needed)

Note: This is motivated by the fact that the I2CMCI data transfer protocol described in Appendix B does not employ endpoint addressing and therefore inherently allows only point to point communication at any one time.

When module selection is used in a host system, the host must provide setup time on the ModSel line of all modules on the same bus and must not change the ModSel line of any module until the management communication is complete and hold time requirements are satisfied, as specified in section B.2.7.1.

B.2.3 Physical Layer Signals

The I2CMCI electrical layer consists of a clock signal (**SCL**) and a data signal (**SDA**).

The initiator drives the SCL signal to clock data and control information on SDA onto the I2CMCI bus.

Both initiator and target latch the state of SDA on the positive transitioning edge of SCL.

The initiator shall initially use a 0-400 kHz SCL clock speed. If a higher management interface speed is supported by the target (see appropriate Hardware Specification) the initiator may later switch to a faster 0-1 MHz SCL clock speed.

The SDA signal is bi-directional.

During binary data transfer, the SDA signal transitions when SCL is low.

Between binary data transfers, SDA transitions when SCL is high are used to mark either the beginning (**START**) or ending (**STOP**) of a data transfer.

Note: An I2CMCI data transfer from initiator to target can be viewed as a serial stream of four tokens, 1, 0, START, STOP, while the reverse data stream is binary.

All bytes (control bytes, address bytes, and data bytes) are transferred bit-serially over the SDA, with the most significant bit (MSB) of each byte transferred first.

A LOW voltage encodes 0. A HIGH voltage encodes 1.

Other electrical signal specifications are found in the appropriate hardware/module specification.

The serial data communication protocol for the data transfer associated with CMIS register access operations over the I2CMCI bit-serial communication link is described in detail in the remainder of this section.

B.2.4 Serial Communication Protocol

The I2CMCI serial data communication protocol governs the I2CMCI bus transactions used for the data transfer associated with CMIS register access operations.

Note: In this section a subtle distinction needs to be made between a I2CMCI bus transaction (a data transfer which begin with START and end either with a STOP or with a new START) and the CMIS register access operation (READ or WRITE) that is conveyed from the host to the module by means of that bus transaction. As will be seen below, register modifications caused by bus transactions occur after the transaction is terminated.

1 B.2.4.1. Basic Definitions and Protocol Elements

2 B.2.4.1.1. Start Condition (START)

3 A high-to-low transition of SDA with SCL high is a **START** condition.

4 All I2CMCI bus transactions begin with a START condition generated by the initiator.

5 *Note: To distinguish from binary data 0 and 1, a START condition can be represented by a backslash \ character*

6 B.2.4.1.2. Stop Condition (STOP)

7 A low-to-high transition of SDA with SCL high is a **STOP** condition.

8 All regular I2CMCI bus transactions end with a STOP condition generated by the initiator.

9 *Note: To distinguish from binary data 0 and 1, a STOP condition can be represented by a slash / character*

10 B.2.4.1.3. Word Size (Byte) and Bit Serial Transmission Order

11 The I2CMCI word size is 8-bits.

12 I2CMCI words are transferred bit-serially with the most significant bit (MSB) first.

13 *Note: The following text uses the term byte to refer to 8-bit words.*

14 B.2.4.1.4. Basic Operation Encoding (Control Byte)

15 The address byte of [1] is used in I2CMCI as a control byte indicating the type of basic management operation
16 that is conveyed in the bus transaction.

17 In I2CMCI the address 10100001b (A1h) indicates a READ access and is called the **read control byte**.

18 In I2CMCI the address 10100000b (A0h) indicates a WRITE access and is called the **write control byte**.

19 B.2.4.1.5. Acknowledge (ACK and NACK)

20 After sending a byte, the side driving the I2CMCI bus releases the SDA line for one-bit time. During this period
21 the receiving side of the I2CMCI bus pulls SDA low (zero) in order to acknowledge (**ACK**) that it has received
22 the byte.

23 Not pulling the SDA low (one) in this period is interpreted as a negative acknowledge (**NACK**).

24 The target sends ACK after each individual control byte received during a read or write transaction, and after
25 each individual data byte received during a write transaction.

26 The initiator sends ACK after each individual data byte received during a read transaction, except for the last
27 data byte, where it terminates the read transaction by sending NACK and then STOP.

28 B.2.4.2. Protocol Reset and Recovery

29 B.2.4.2.1. Power-On Reset

30 The I2CMCI adapter (interface circuitry) enters a reset state upon application of power.

31 B.2.4.2.2. Protocol Reset and Recovery

32 Synchronization issues may cause the initiator and target to disagree on the specific bit location currently being
33 transferred, the type of transaction, or even if a transaction is in progress.

34 The I2CMCI protocol has no explicitly defined reset mechanism. The following procedure may force completion
35 of a currently pending transaction and cause the target to release SDA.

- 36 1. The initiator shall provide up to nine SCL clock cycles (drive low, then high) to the target
- 37 2. The initiator shall monitor SDA while SCL is high on each cycle.
- 38 3. On any clock cycle, if the target releases SDA, the initiator shall initiate a STOP condition. The initiator
39 is then free to send a START condition for the next transaction
- 40 4. If SDA remains low after nine clock cycles the I2CMCI protocol reset has failed

B.2.5 Serial MCI Transactions for READ/WRITE/TEST Access

This section describes the I2CMCI transactions between initiator and target, for READ, WRITE, TEST.

I2CMCI provides defined data transfer transactions to implement the following basic register access operations

Table B-2 I2CMCI Transactions

I2CMCI Transaction	Register Access Primitive	Section
Read one byte from current Byte address	VALUE = READ(1)	B.2.5.1.1
Read n>1 bytes sequentially, starting at current Byte address	VALUES = READ(n)	B.2.5.1.2
Read one byte from given Byte address	VALUE=READ(ByteAddress, 1)	B.2.5.2.1
Read n>1 bytes sequentially, starting at given Byte address	VALUES=READ(ByteAddress, n)	B.2.5.2.2
Write one byte to a given address	WRITE(ByteAddress, Value)	B.2.5.3
Write N>1 bytes sequentially, starting at given Byte address	WRITE(ByteAddress, Val1, ..., ValN)	B.2.5.4
Test if target is ready to accept transaction (ACK polling)	TEST()	B.2.5.5

Note: Standard length limitations for reading or writing multiple bytes are specified in section 5.2 with possible exceptions explicitly specified elsewhere.

Read/Write Control Byte and Response

In the first part of a transaction, after the START condition, the first byte of a I2CMCI bus transaction is a control byte consisting of a fixed 7-bit part 1010000b followed by a bit indicating the type of transaction: A **read** transaction is requested if this bit is 1 (high). A **write** transaction is requested if this bit is 0 (low).

Upon reception of the control byte the target asserts the SDA signal low to acknowledge (ACK) receipt unless a transaction **rejection** is required (see section B.2.6.2).

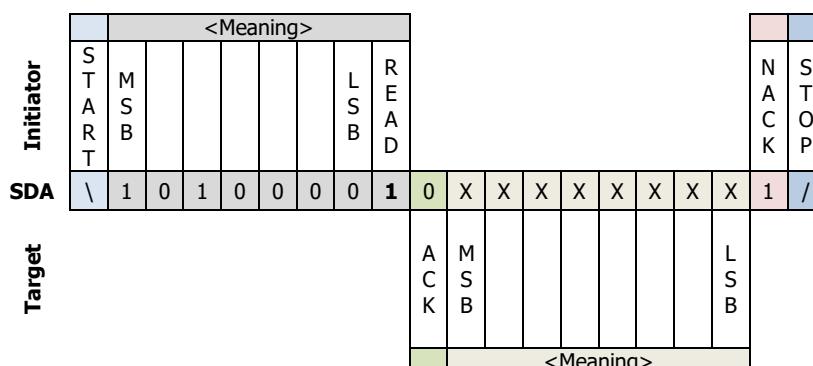
Byte Address and Data Bytes

When the first part of a transaction was acknowledged, addresses and/or data are transmitted in units of bytes, in the second part of a transaction. Format and interpretation of this transferred byte sequence is specific for each basic transaction (as described section B.2.4.2).

In transactions transferring more than one data byte, the data bytes are transmitted in the order of increasing byte addresses (cyclically increasing in case of roll-over described in section B.1.2).

Illustrations

To visually illustrate the bus transactions, figures like the following are used to show the bidirectional token stream on the SDA line, consisting of 0, 1, \ (START), and / (STOP) tokens. Colors and structure around the SDA line are used to show when Initiator or Target drive the SDA wire, and for what purpose.



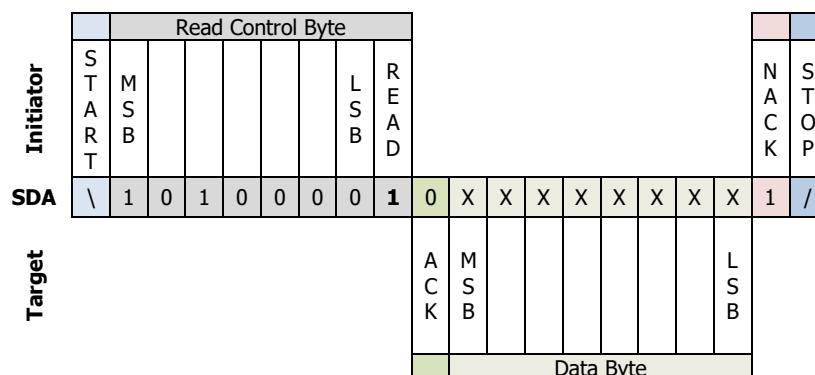
B.2.5.1. Transactions for a Byte Read Operation**B.2.5.1.1. Transaction for a Current Address Read Operation**

The initiator generates a START condition and sends the read control byte (10100001b).

The target sends ACK followed by the data byte retrieved from the address in the current byte address.

The initiator terminates the transaction with NACK and STOP.

The target increments the current byte address after orderly termination of the transaction (STOP) and before accepting a new transaction.

**Figure B-1 Current Address Read**

B.2.5.1.2. Transaction for a Random Read Operation

A random read transaction is implemented as an aborted dummy write transaction, followed by a current address read transaction.

The dummy write transaction is used to load (but not increment) the target byte address into the current byte address, from which the subsequent current address read transaction then reads. The procedure is as follows:

The initiator generates a START condition and sends the target byte address after the write control byte (10100000b).

The initiator then generates another START condition (aborting the dummy write) and begins a current address read transaction by sending a read control byte (10100001b).

The target acknowledges each byte received.

When the byte address of the dummy write is received, the target updates the current byte address. When the target then sees the write transaction aborted (i.e. when it receives a START instead of a byte to be written), the current byte address is not incremented and so contains the address to be read in the subsequent current address read transaction.

The initiator terminates the transaction with NACK and STOP when it has received the requested byte.

The target increments the current byte address after orderly termination of the transaction (STOP) and before accepting a new transaction.

Figure B-2 Random Read

B.2.5.2. Transaction for a Sequential Bytes Read Operation

A sequential bytes read transaction is implemented as a continuation within either a current address read (see Figure B-3) or a random address read (see Figure B-4).

The initiator indicates the continuation of a reading sequence to the target by sending an ACK after a data byte received (instead of terminating the transaction with NACK and STOP).

When the target receives an ACK after sending a byte to the initiator, it increments the current byte address and sends the byte that is stored at this address.

The sequential bytes read is terminated when the initiator sends a NACK and a STOP (instead of ACK).

The target increments the current byte address after orderly termination of the transaction (STOP) and before accepting a new transaction.

B.2.5.2.1. Transaction for a Sequential Bytes Read from Current Start Address

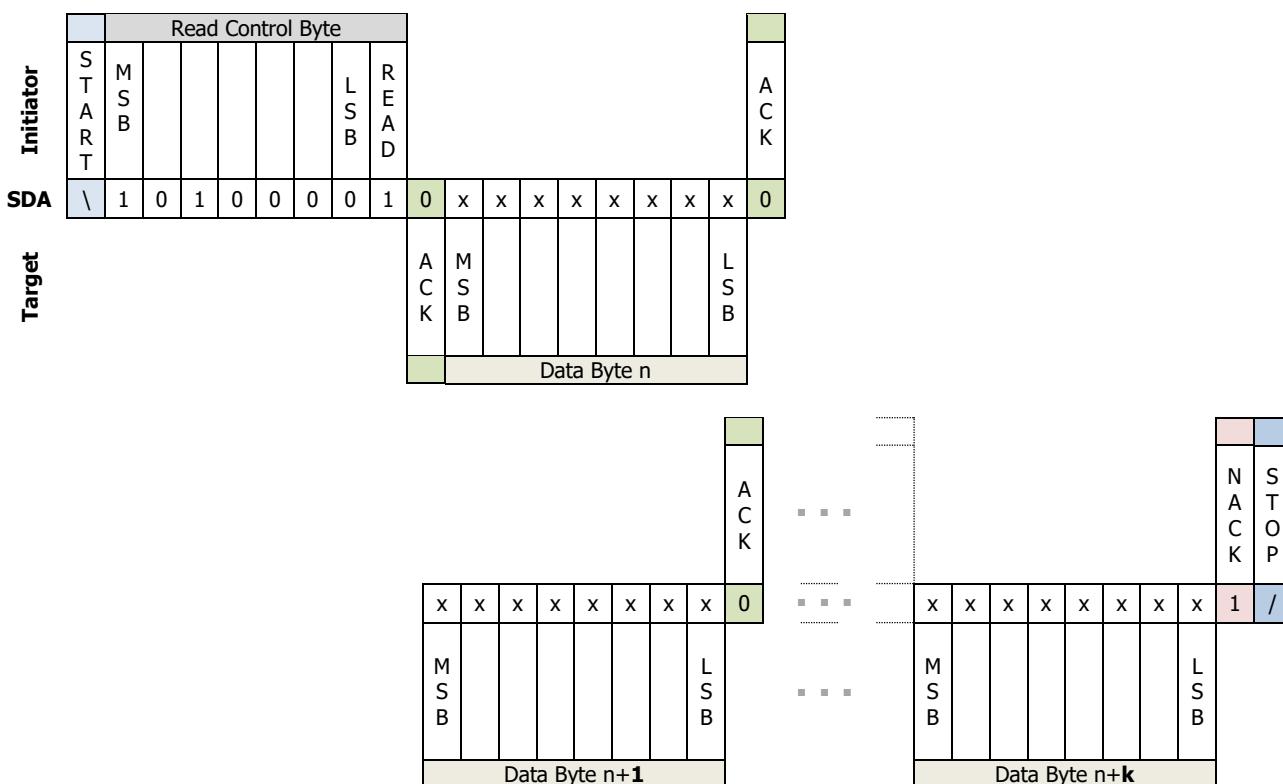


Figure B-3 Sequential Bytes Read Starting at Current Address

B.2.5.2.2. Transaction for a Sequential Bytes Read from Random Start Address

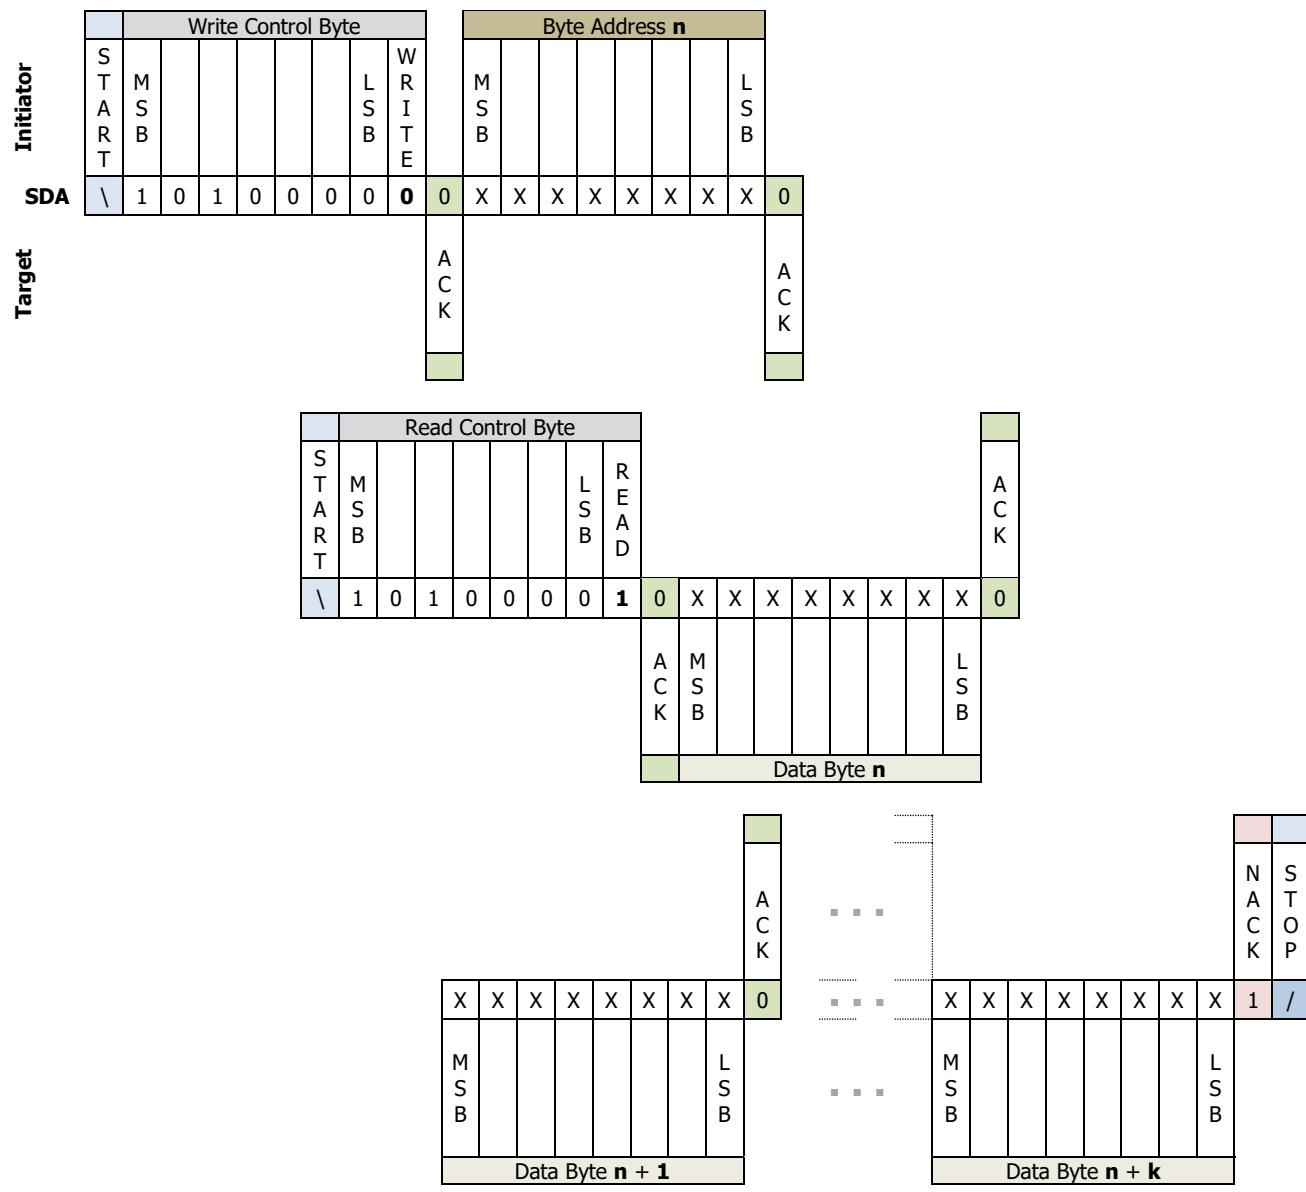


Figure B-4 Sequential Bytes Read Starting with Random Read

B.2.5.3. Transaction for a Single Byte Write Operation

The initiator generates a START condition and sends the write control byte (10100000b).

When the target has responded to the write control byte with ACK, the initiator sends the target byte address.

When the target has responded to the target byte address with ACK, the initiator sends the data byte value.

When the target has responded to the data byte value with ACK, the initiator sends STOP to terminate the write transaction. Otherwise the write transaction is either continued (see section B.2.5.4) or aborted.

On receipt of the target byte address, the target immediately updates its current byte address.

However, the target begins its internal memory write cycle of the received data byte to the address in the current byte address not before the write transaction is properly terminated by STOP.

Note: In a future revision, these register access related specifications may be moved to the appropriate section.

After receiving STOP, the target increments the current byte address before accepting the next transaction.

If a START condition is received in place of a STOP condition the target discards the data byte received and does not increment the current byte address.

For writes to non-volatile memory, upon receipt of the proper STOP condition, the target must enter and complete an internally timed write cycle before the next transaction can be accepted. If needed, the target must hold-off a subsequent transaction for a maximum duration tWR.

For writes to volatile memory, upon receipt of the proper STOP condition, the target must enter and complete an internally timed write cycle before the next transaction can be accepted. If needed, the target must hold-off subsequent transaction for a maximum duration tNACK.

Note: To hold-off a subsequent transaction the target may e.g. disable its management interface and not respond or acknowledge subsequent commands until the internal memory write cycle is complete.

Note: the 'Combined Format' [1] using repeated START conditions is not supported on write transactions.

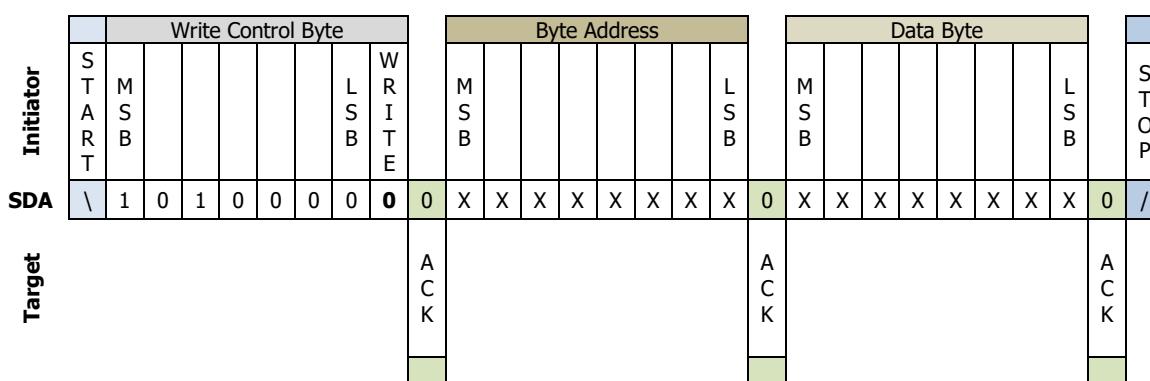


Figure B-5 Write Byte Transaction

B.2.5.4. Transaction for a Sequential Bytes Write Operation

The initiator initiates a sequential write transaction for several bytes in the same way as a single byte write, but it then does not send a STOP condition after the first byte. Instead, after the target acknowledges (ACK) receipt of the first data byte, the initiator transmits the additional data bytes without transmitting new explicit byte address information or control bytes.

The initiator terminates the sequential write sequence with a STOP condition, or otherwise, the transaction is aborted, and the results of the sequential write are undetermined.

The target shall ACK each byte received.

The target may either store a data byte after sending ACK or it may decide to buffer all bytes of the sequential write transaction until the transaction is terminated by STOP.

After a properly terminated sequential write transaction (STOP received) the target ensures that the current byte address contains the address of the next byte after the last byte written, before accepting the next transaction. Otherwise, the value of the current byte address is undetermined.

At the end of each 128-byte Page, the current byte address counter rolls over to the first byte of that Page.¹

For writes to **non-volatile** memory, upon receipt of the proper STOP condition the target must complete its internally timed write cycle before the next basic register access operation can be accepted. If needed, the target must hold-off subsequent transaction for a maximum duration **tWR**, to internal memory.

For writes to **volatile** memory, upon receipt of the proper STOP condition the target must complete its internally timed write cycle before the next basic register access operation can be accepted. If needed, the target must hold-off subsequent transaction for a maximum duration **tNACK**, to internal memory.

Note: To hold-off a subsequent transaction the target may e.g. disable its management interface input and not respond or acknowledge subsequent commands until the internal memory write cycle is complete.

Note: the 'combined format' using repeated START conditions is not supported.

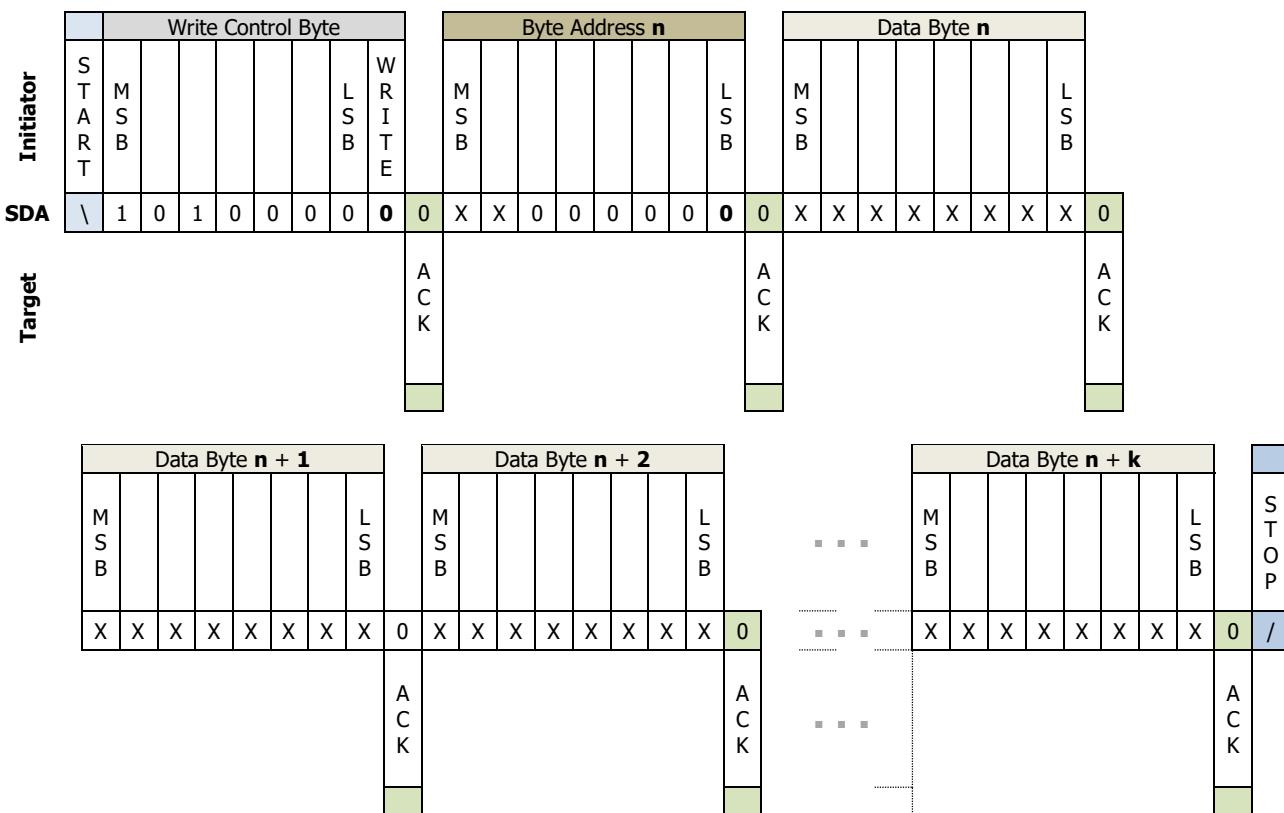


Figure B-6 Sequential Bytes Write Transaction

¹ There may be exceptions to this basic rule for advanced management features described in chapter 7. Such exceptions are described where those features are specified.

B.2.5.5. Transaction for a Test Operation

To test if the target is ready to accept a new transaction, the initiator can issue a **dummy** write transaction (START, **write** control byte, STOP) to see if the module ACKs the write control byte.

Note: This method of explicit Acknowledge Polling is the most lightweight method to synchronize to the target becoming ready without side effects on the target.

Note: A dummy read transaction (START, read control byte, STOP) cannot be used for Acknowledge Polling as this may have unintended side effects, when a target interprets this as a current address read transaction (see section B.2.5.1.1) and starts to drive SDA at the same time the initiator is driving SDA for the STOP.

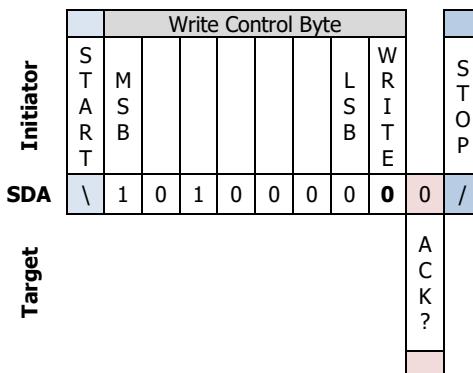


Figure B-7 Test Readiness Transaction

1 B.2.6 Transaction Flow Control Mechanisms

2 The MCI offers two flow control mechanisms to the local CMIS register access layer (RAL) instance.

3 B.2.6.1. Delaying Current Transaction (Clock Stretching)

4 Clock stretching is a mechanism for the target to delay completion of the current I2CMCI transaction.

5 The target exercises clock stretching by pulling the clock signal SCL low for a certain duration.

6 The duration while the target pulls the SCL signal low is the clock stretching duration.

7 The target is allowed to initiate clock stretching (SCL pull down) only while SCL is low.

8 The maximum total allowed clock stretching duration during one transaction is specified in section B.2.7.3.

9 *Note: The target may use the clock stretching mechanism whenever needed to implement the register access
10 associated with the transfer, e.g. to fetch a register value in a read transaction or, theoretically, to complete a
11 register update from a preceding write transaction.*

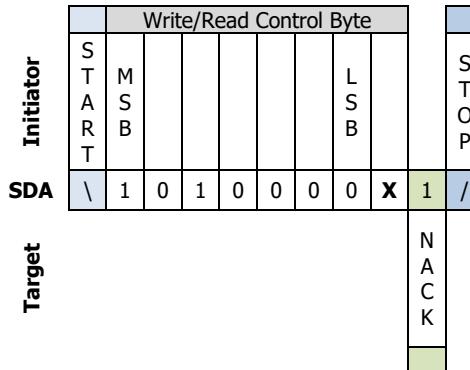
12 B.2.6.2. Rejecting Subsequent Transaction (Transaction Hold-Off)

13 Transaction hold-off is a temporary phase when new transactions are **rejected** by signaling NACK (instead of
14 ACK) after having received the control byte of a new transaction (see figure below).

15 *To implement hold-off the target may choose to simply disable its I2CMCI inputs, as this results in the host
16 receiving a NACK after the first byte of a transaction following the START condition.*

17 The maximum allowed transaction hold-off duration is specified in sections B.2.7.3.

18 *Note: The target exercises transaction hold-off when it is unable or not allowed to accept a new transaction in
19 its current state or condition. For instance, after a properly terminated write transaction, the target may need
20 to reject subsequent transactions in order to internally finalize the WRITE operation represented by the previous
21 transaction, due to synchronization requirements specified in section 5.2.4.*



23 24 **Figure B-8 Test Readiness Transaction**

B.2.7 Timing Specifications

Note: The I2CMCI timing specifications must be compatible with the timing specifications of the register access layer specified in chapter 10.

B.2.7.1. Module Select Timings

The following timing parameters apply to host-module interactions over the MCI in form factors implementing the ModSel module select signal.

Table 10-10 Module Select Timings

Parameter	Symbol	Min	Max	Unit	Conditions
Max Aborted sequence – bus release	Deselect_Abort		2	ms	Delay from a host de-asserting ModSelL (at any point in a bus sequence) to the module releasing SCL and SDA
Min ModSelL Setup Time ¹	tSU_ModSelL	2		ms	ModSelL Setup Time is the setup time on the select line before the start of a host initiated serial bus sequence.
Max ModSelL Hold Time ¹	tHD_ModSelL		500	us	ModSelL Hold Time is the delay from completion of a serial bus sequence to changes of module select status.

Note 1: Support for these setup and hold times are required for all modules. Once the module has initialized the management interface, the host may read advertised ModSelL setup and hold times from the register map.

B.2.7.2. Waveform Timings

An example I₂C/MCI waveform is shown in Figure B-9, and the timing parameters are specified in Table B-3. The default clock rate is a maximum of 400 kHz with an option to support up to a maximum of 1 MHz.

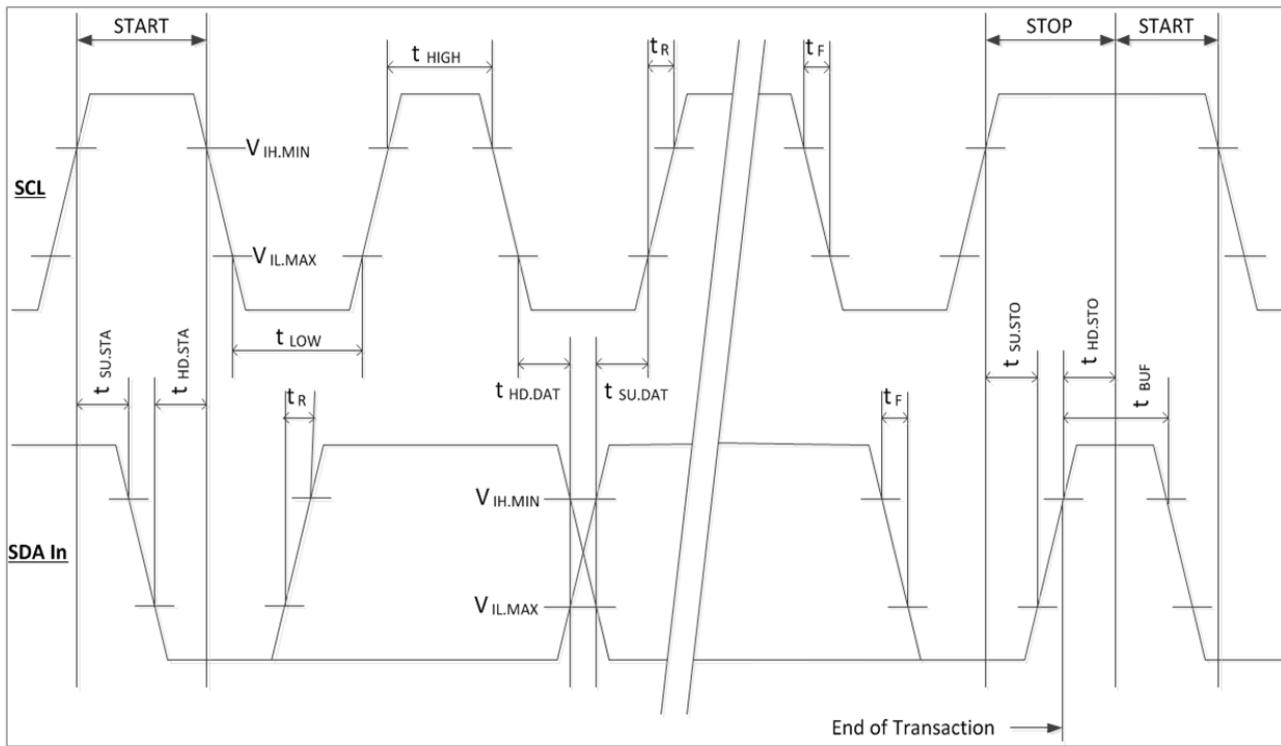


Figure B-9 I2CMCI Waveform Timing Diagram

Table B-3 I2CMCI Waveform Timing Parameters

Parameter	Symbol	0.4 MHz		1 MHz		Unit	Conditions
		Min	Max	Min	Max		
Max Clock Frequency	fSCL	0	400	0	1000	kHz	
Min Clock Pulse Width Low	tLOW	1.3		0.50		μs	
Min Clock Pulse Width High	tHIGH	0.6		0.26		μs	
Min START Hold Time	tHD.STA	0.6		0.26		μs	Delay required between SDA becoming low and SCL starting to go low in a START
Min START Setup Time	tSU.STA	0.6		0.26		μs	Delay required between SCL becoming high and SDA starting to go low in a START
Min Data In Hold Time	tHD.DAT	0		0		μs	
Min Data In Setup Time	tSU.DAT	0.1		0.1		μs	
Max Input Rise Time	tR		300		120	ns	From (VIL,MAX=0.3*Vcc) to (VIH, MIN=0.7*Vcc), see Figure B-9.
Max Input Fall Time	tF		300		120	ns	From (VIH,MIN=0.7*Vcc) to (VIL,MAX=0.3*Vcc), see Figure B-9.
Min STOP Setup Time	tSU.STO	0.6		0.26		μs	
Min STOP Hold Time	tHD.STO	0.6		0.26		μs	

2

3

B.2.7.3. Transaction Timings

4

Table B-4 Flow Control Timings

Parameter	Symbol	Min	Max	Unit	Conditions
Max I2CMCI READ Delay (Clock Stretch) (formerly: T_clock_hold)	tRD		500	μs	Maximum time the module may delay returning the requested data in an MCI Read transaction (e.g. by clock stretching)
Max time to complete a standard WRITE to volatile memory or register	tNACK		10	ms	Maximum transaction hold-off time after a write transaction to an address representing volatile memory or a volatile management register
Max time to complete a standard WRITE access to non-volatile memory or register	tWR		80	ms	Maximum transaction hold-off time after a write transaction to an address representing non-volatile memory or a non-volatile management register
Min inter-transaction idle time	tBUF	20		μs	Min time bus free before new transaction can start, between STOP and START and between ACK and START

6

B.2.7.3.1. tWR Timing

The timing attribute tWR is the maximum time allowed for a module to complete its internally timed write cycle after a single or sequential write to non-volatile memory before the next basic management operation can be accepted.

The write cycle completion time is measured from the low to high SDA edge of the STOP condition of the Write transaction to the high to low SDA edge of the START condition for the next transaction.

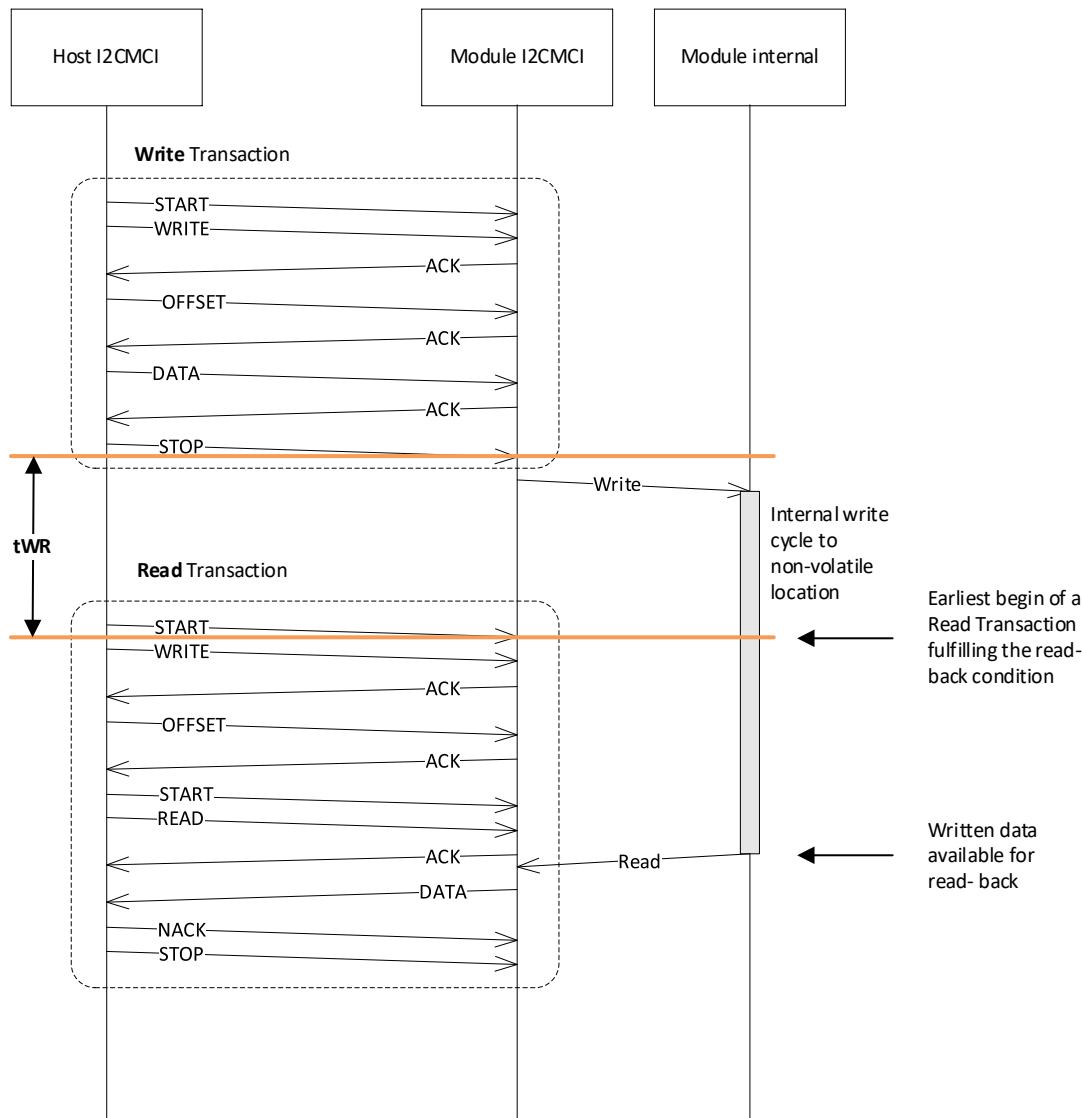


Figure B-10 tWR bus timing

B.2.7.3.2. tNACK Timing

The timing attribute tNACK is the maximum time allowed for a module to complete its internally timed write cycle after a single or sequential write to a volatile memory location before the next basic management operation can be accepted.

The write cycle completion time is measured from the low to high SDA edge of the STOP condition of the Write transaction to the high to low SDA edge of the START condition for the next transaction.

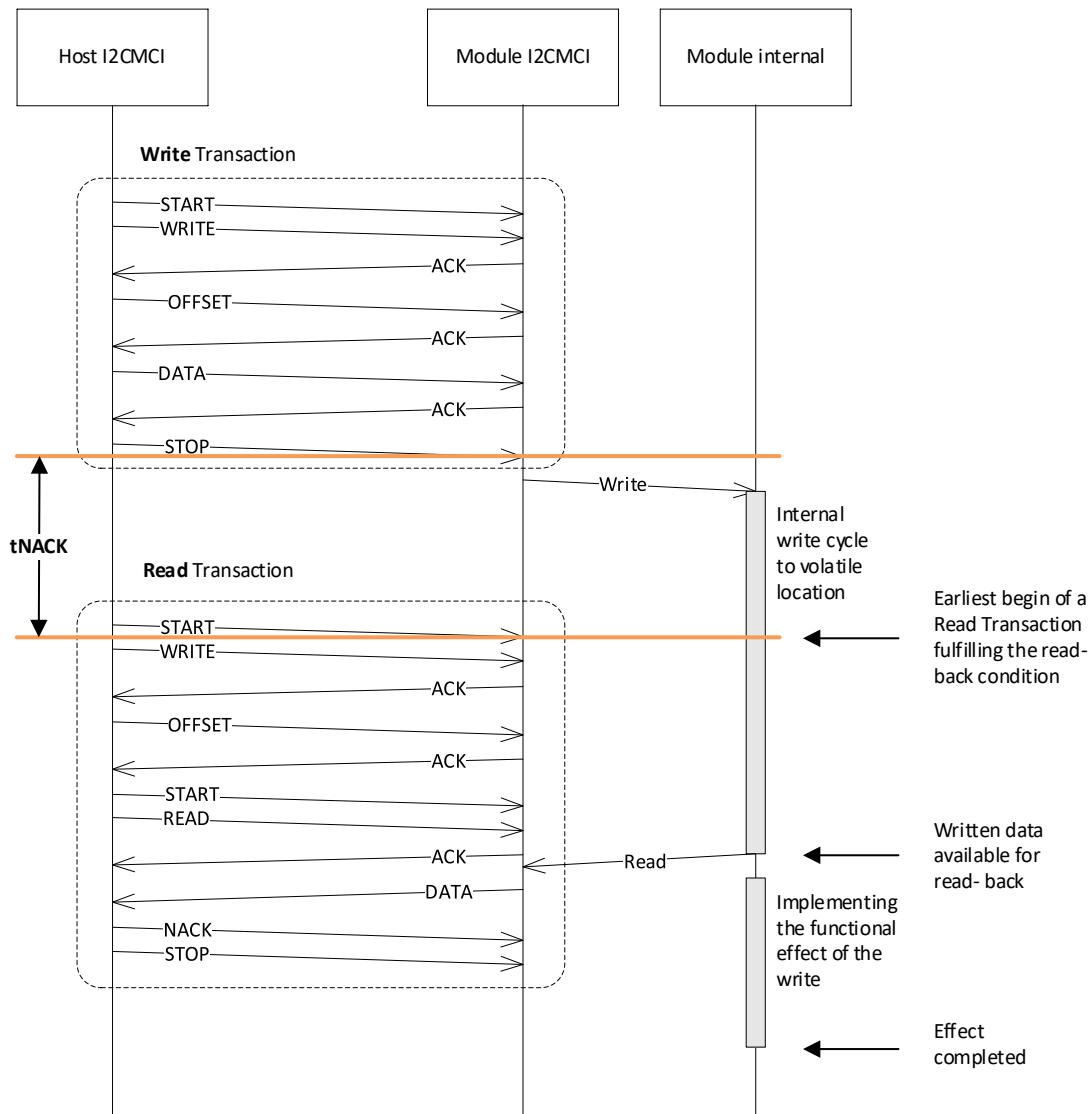


Figure B-11 tNACK bus timing

1 **B.3 SPI-Based Management Communication Interface (SPIMCI)**

2 **B.3.1 Introduction**

3 This SPIMCI specification is a specialization of the industry-standard SPI bus, for the purpose of management
4 communication between an Initiator (host) and one or more Targets (modules).

5 SPIMCI specifies how CMIS register access operations are implemented by SPIMCI bus transactions.

6 SPIMCI requires a four-wire electrical interface when connecting one host and one module; for each additional
7 module there is one additional dedicated chip select wire.

8 **Relation to SPI Bus**

9 The original SPI bus is used to connect peripherals to a controller. It allows synchronous bit-serial data
10 communication in both directions (full duplex), with all serially transmitted binary data synchronized to the rising
11 or falling edge of a controller supplied clock.

12 *Note: Due to undesired negative social connotations, the historical SPI signal names MISO and MOSI are
13 replaced by socially neutral terms, IOTT and IITO, see below).*

14 **Relation to I2CMCI**

15 Compared to I2CMCI described in section B.2, SPIMCI provides an important addressing enhancement.

16 Whereas I2CMCI transactions carry only a page-local byte-sized byte address, creating the need for special
17 page and bank switching registers, SPIMCI transactions always carry a full 3D (bank, page, byte) address, with
18 implicit bank and page switching as needed. This improvement enhances throughput.

19 *Note: A similar global addressing variant could easily be defined for I2CMCI.*

20 **B.3.2 Communication Topology**

21 On an SPIMCI bus there can be one Initiator and one or more Targets.

22 The Initiator can communicate with one Target at a time, by asserting at most one CSn signal at any one time.

23 **B.3.3 SPIMCI Control Signals**

24 A chip select signal (CSn) controlled by the Initiator is used both to select a Target for communication and to
25 delimit a bus transaction.

26 The CSn signal is therefore considered to be part of the SPIMCI physical layer signals and not a separate control
27 signal.

28 When multiple Targets are used, the Initiator must provide one individual chip select signal for each Target and
29 ensure that at most one such signal is asserted.

B.3.4 Physical Layer Signals

SPIMCI uses four physical layer signals to implement bus transactions, three of which can be used as a bus.

The following figure illustrates this physical interface between the Initiator and one Target.

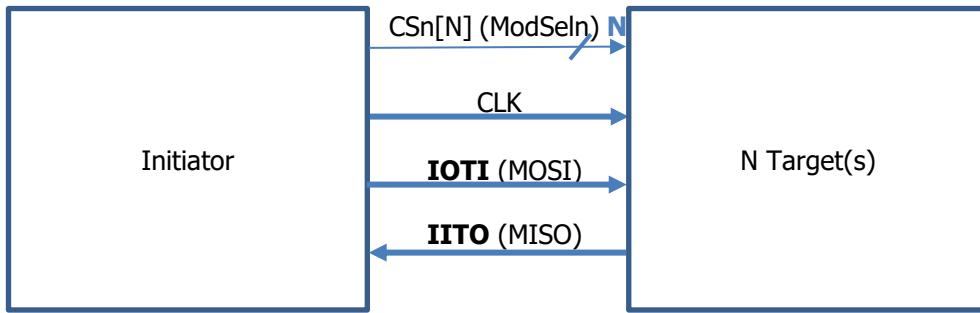


Figure B-12 SPIMCI Physical Layer Signals

- **CSn** (Chip Select): a dedicated active low signal used select a Target and to start and stop a transaction
- **CLK** (SPI Clock): to synchronize binary data exchanged between Initiator and selected Target
- **IOTI** (Initiator Out, Target In) or MOSI: serial binary data (also PICO, peripheral in, controller out)
- **IITO** (Initiator In, Target Out) or MISO: serial binary data (also POCI, peripheral out, controller in)

The device that generates the clock signal (CLK) is called the Initiator.

The non-shared chip select signal (CSn) from the Initiator to a Target is used to select that Target and at the same time to delineate a transaction.

Data transmitted between Initiator and Target in any direction is synchronized to the CLK signal.

CSn is pulled low to initiate communication with a Target.

CSn is pulled high to disconnect the Target from the SPI bus (usually at the end of a transaction).

IOTI and IITO are the data lines.

IOTI transmits data from the Initiator to the Target.

IITO transmits data from the Target to the Initiator.

B.3.5 Communication Speed

SPIMCI supports clock rates of 1, 2, 4, 8, 12, 16, 20 MHz.

The Initiator should initially use the lowest speed to query Target capabilities and switch to higher speeds (MciSpeedConfiguration) once the Target capabilities have been determined from module advertisements (MciMaxSpeed, MciFlowControlDurationEncoding, and MciFlowControlDuration).

B.3.6 Physical Encoding

A LOW voltage encodes a digital bit value of 0.

A HIGH voltage encodes a digital bit value of 1.

B.3.7 Serial Communication Protocol

The SPIMCI serial data communication protocol governs the SPI bus transactions used for the data transfer associated with CMIS register access operations.

Since SPIMCI is a full-duplex interface, the Initiator simultaneously transmits data (serially shifting bits out onto the IOTI bus) and receives data (serially sampling bits in from the IITO bus), and likewise the Target receives data (sampling bits from the IOTI bus) and transmits data (serially shifting bits out onto the IITO bus).

The serial clock CLK edge synchronizes the shifting-out and sampling-in of the data.

During a transaction, both Initiator and Target send data continuously and simultaneously in opposite directions via the IOTI and IITO lines respectively.

The meaning and use of segments of these contra-directional data streams at any point in time during a bus transaction is defined for each management transaction, as described in sections below.

B.3.7.1. Bus Transactions

To enable communication the Initiator must generate the clock signal (CLK).

To begin a bus transaction, the Initiator selects the desired Target by asserting the CS signal (CSn).

To terminate a bus transaction, the Initiator deselects the desired Target by deasserting the CS signal (CSn).

Note: Recall that chip select (CSn) is an active low signal.

The **rising** edge of the CLK signal is used to

- sample the IITO signal and shift out IOTI data at the Initiator
- shift out IITO data and sample the IOTI signal at the Target

The figure below shows the general form of an SPIMCI bus transaction. The START condition (described below) determines the start of the transaction.

Once the START condition is determined, the protocol distinguishes three phases:

1. Transaction Control Phase (**4** bytes)
2. Flow Control Phase (**N** \geq 1 bytes)
3. Data Transfer Phase (**M** = EncodedSize + 1 Bytes)

Byte	Bit	IOTI	IITO
0	7	Command Type	Undefined
	6:3	Bank Index	
	2:0	EncodedSize [10:8]	
1		EncodedSize [7:0]	Undefined
2		Page Index	Undefined
3		Byte Index	Undefined
(3+1) ... (3+N-1)		00h	Undefined
(3+N)		00h	00h (ACK) FFh (NACK)
(4+N) ... (4+N+M-1)		00h (read) M Data Bytes (write)	M Data Bytes (read) Undefined (write)

The diagram illustrates the three phases of an SPIMCI bus transaction using curly braces on the right side of the table. The first three bytes (0, 1, 2) are grouped under 'Transaction Control Phase'. The fourth byte (3) and all subsequent bytes up to (4+N-1) are grouped under 'Flow Control Phase'. All bytes from (3+N) onwards are grouped under 'Data Transfer Phase'.

Figure B-13 SPIMCI Bus Transactions (general)

If the Initiator asserts CSn longer than announced in the Transaction Control Header, the Target ignores the IOTI data and sends undefined content on IITO.

If the Initiator asserts CSn shorter than announced in the Transaction Control Header, the Target stops processing the current transaction with undefined effects of a WRITE transaction or undefined side effects of a READ transaction.

B.3.7.1.1. Transaction Control Phase

The Transaction Control phase covers the first four bytes (32 bits) transmitted by the Initiator to the Target.

The Transaction Control bytes are divided into several bit fields.

Read/Write (R/Wn) (transmitted bit index 0:0)

The first bit in the stream is the read/write indicator bit (R/W).

When this bit is 1, a read is being requested. Otherwise, a write is requested.

At successful completion of the transaction control phase (the first 32 bits), the Target uses the information provided to decode the internal address to be read or written.

Bank (transmitted bit indices 1:4)

The next 4 bits in the bit stream represent the bank index in management memory to be accessed.

Note: This format does not allow for arbitrary bank index values greater than 15!

Size (transmitted bit indices 5:15)

The next 11 bits encode the number of bytes to transfer for either read or write.

If the numerical value represented by unsigned size field is n, then (n+1) bytes of data are transferred.

Note: If the value represented by the size field is 0, one byte of data is transferred. If the value represented by size field is 1, two bytes of data are transferred. A maximum of 2048 bytes can be transferred in suitable regions of the memory map, where auto-page increment overrides page wrap around.

Page Index (transmitted bit indices 16:23)

The next 8 bits in the bit stream represent the page index of the management memory to be accessed.

Byte Index (transmitted bit indices 24:31)

The final 8 bits represent the starting byte address of the data sent or received. If more than one word is being sent, sequential addressing is used, starting with the address specified.

Note that the address increment logic is governed by generic rules in section B.1.2. For instance, addresses in regular pages wrap around. E.g. if the end of page (or of Lower Memory) is reached during a transaction on a page without auto page incrementing, a byte address wrap around occurs within the current 128-byte memory segment and address processing continues at 00h or 80h of the current 128-byte memory segment, depending on whether Lower Memory or Upper Memory is being accessed.

B.3.7.1.2. Flow Control Phase (N Bytes)

The Flow Control Phase follows the Transaction Control Phase.

During the Flow Control Phase, the Initiator sends $\mathbf{N} \geq 2$ zero-valued dummy bytes, which are to be ignored by the Target, while the Target evaluates its ability to service the requested transaction, returning ACK or NACK in the last byte of the Flow Control Phase (called the flow control bits below).

The duration of the Flow Control Phase is $\mathbf{N} \cdot 8$ bit durations at the current bus speed.

The actual value of $\mathbf{N} \geq 2$ is determined from advertisements (MciFlowControlDuration) and from the currently configured SPI bus speed (MciSpeedConfiguration): Depending on the advertised encoding method (MciFlowControlDurationEncoding), \mathbf{N} is either a fixed number or a speed-dependent number which is determined from a given time duration and the current SPI speed. Details are defined in section 8.3.10

Note: Both Initiator and Target must determine the value of \mathbf{N} consistent with the module advertisement and with the host configured speed, in the same way.

The Target uses the Flow Control Phase to decode the transaction control header and to determine if it is ready to process the command encoded in the transaction.

1 During a READ transaction, the flow control phase provides time for the read data to be retrieved internally.
2 If the Target is busy processing a previous WRITE command, the Target will send a logic 1 (busy, NACK) on
3 the flow control bits on its IITO line.

4 If the Target can accept the requested command, it will send a logic 0 (idle, ACK) on the flow control bits on
5 its IITO line.

6 The Target guarantees to send the flow control result (ACK or NACK) in the last byte of the Flow Control Phase.

7 **B.3.7.1.3. Data Transfer Phase (M Bytes)**

8 The Data Transfer Phase follows the Flow Control phase. The amount (in number of bytes **M**) of data announced
9 to be sent or requested to be received is determined from the EncodedSize field as **M** = EncodedSize + 1.

10 **B.3.7.2. Basic Definitions and Protocol Elements**

11 **B.3.7.2.1. Start Condition (START):**

12 A high-to-low transition of CSn is a START condition.

13 All SPIMCI bus transactions begin with a START condition generated by the Initiator.

14 **B.3.7.2.2. Stop Condition (STOP):**

15 A low-to-high transition of CSn is a STOP condition.

16 All SPIMCI bus transactions must be terminated with a STOP condition generated by the Initiator.

17 A regular STOP condition always occurs on an 8-bit boundary at an expected point in time.

18 An irregular STOP condition (a.k.a. an ABORT condition) occurs on a non-8-bit boundary or earlier than
19 expected. An ABORT terminates the transaction, and any incompletely received byte must be discarded or
20 ignored by the receiving side.

21 *Note: Completely received data bytes, however, may or may not be processed normally.*

22 **B.3.7.2.3. Bus Word Size (Byte) and Bit Serial Transmission Order**

23 The SPIMCI bus word size is 8-bits (one byte).

24 SPIMCI bus words (i.e., bytes) are transferred bit-serially with the most significant bit (MSB) first.

25 **B.3.7.2.4. Acknowledge (ACK and NACK)**

26 As described above, SPIMCI uses flow control bits on its IITO line to determine whether the Target is ready to
27 accept a command (ACK) or is busy (NACK).

28 **B.3.7.2.5. TEST**

29 The Initiator can use the dedicated TEST transaction to check if the Target is busy or idle. See section B.3.8.5

30 **B.3.7.3. Protocol Reset and Recovery**

31 **B.3.7.3.1. Power-On Reset**

32 The SPIMCI adapter (interface circuitry) enters a reset state upon application of power.

33 **B.3.7.3.2. Protocol Violation**

34 There are only a few detectable protocol violations:

35 An Initiator not sending STOP after specified length. Module behavior has been partially specified.

36 A Target not responding with one of the expected patterns (all zero, all ones in the last Flow Control byte).

37 **B.3.7.3.3. Protocol Reset and Recovery**

38 The transactional data transmission protocol is reset after each CSn deassert.

39 The current address pointer state in the module is of no interest with SPIMCI, because it is set to an explicitly
40 transmitted address in every transaction.

B.3.8 SPI MCI Transactions for READ/WRITE/TEST Access

SPIMCI provides defined data transfer transactions to implement the following basic register access operations.

SPIMCI Transaction	Register Access Primitive	Section
Not supported: Read one byte from current Byte address	VALUE = READ(1)	
Not supported: Read n bytes starting at current Byte address	VALUE = READ(n)	
Read one byte from given Byte address	VALUE = READ(ByteAddress, 1)	B.3.8.1
Read n>1 bytes sequentially, starting at given Byte address	VALUES = READ(ByteAddress, n)	B.3.8.2
Write one byte to a given address	WRITE(ByteAddress, Value)	B.3.8.3
Write N>1 bytes sequentially, starting at given Byte address	WRITE(ByteAddress, Val1, ..., ValN)	B.3.8.4
Test if target is ready to accept transaction (ACK polling)	TEST()	B.3.8.5

B.3.8.1. Read One Byte from Given Byte Address

The Initiator generates a START condition.

After the START condition, the Initiator sends 32 bits of Transaction Control with the Read bit set to 0. Transaction Control contains information on the bank, page, and address to be read. The Size field of the instruction data is set to 0 indicating a single byte read.

Then the Initiator sends N zero valued bytes to allow the Target to decode the internal address to be read.

The Initiator can read its IITO line to determine if the Target is Busy (FFh) or Idle (00h).

If the target indicates it is idle, then the Initiator can clock out the data to be read from the IITO line.

The Initiator generates a STOP condition indicating end of a command.

Byte	Bit	IOTI	IITO
0	7	0 (Read)	Undefined
	6:3	Bank#	
	2:0	Encoded Size = 0	
1		Undefined	
2		Page#	Undefined
3		Address#	Undefined
(3+1) ... (3+N-1)		00h	Undefined
(3+N)		00h	00h (ACK) FFh (NACK)
(4+N)		00h (Ignored)	Data Byte (read)

Figure B-14 SPIMCI READ Transaction

1 B.3.8.2. Read n Bytes from Given Byte Address

2 The Initiator generates a START condition.

3 After the START condition, the Initiator sends 32 bits of Transaction Control with the Read bit set to 0.
4 Transaction Control contains information on the bank, page, and address to be read. The Size field of the
5 instruction data is set to (n-1) indicating a multiple byte read of M = n bytes.

6 Then the Initiator sends N zero valued bytes to allow the Target to decode the internal address to be read.

7 The Initiator can read its IITO line to determine if the Target is Busy (FFh) or Idle (00h).

8 If the target indicates it is idle, then the Initiator can clock out the data to be read from the IITO line. The
9 Initiator clocks out the number bytes indicated by the size field.

10 The Initiator generates a STOP condition indicating end of a command.

Byte	Bit	IOTI	IITO
0	7	0 (Read)	Undefined
	6:3	Bank#	
	2:0	EncodedSize = n - 1	
1			Undefined
2		Page#	Undefined
3		Address#	Undefined
(3+1) ... (3+N-1)		00h	Undefined
(3+N)		00h	00h (ACK) FFh (NACK)
(4+N) ... (4+N+n-2)		00h (Ignored) ... 00h (Ignored)	Data Byte 1 (read) ... Data Byte n-1 (read)
(4+N+n-1)		00h (Ignored)	Data Byte n (read)

11 **Figure B-15 SPIMCI READ N Transaction**

12

B.3.8.3. Write 1 Byte to a Given Byte Address

The Initiator generates a START condition.

After the START condition, the Initiator sends 32 bits of Transaction Control with the Write bit set to 1. Transaction Control contains information on the bank, page, and address to be read. The Size field of the instruction data is set to 0 indicating a single byte write.

Then the Initiator sends 16 bits of flow control data with a logic 0. During these 16 bits, the Initiator can read its IITO line to determine if the Target is Busy (FFh) or Idle (00h).

If the target indicates it is idle, then the Initiator can clock out the data to be written on the IOTI line.

The Initiator generates a STOP condition indicating end of a command.

Byte	Bit	IOTI	IITO
0	7	1 (Write)	Undefined
	6:3	Bank#	
	2:0	0	
1		0	Undefined
2		Page#	Undefined
3		Address#	Undefined
(3+1) ... (3+N-1)		00h	Undefined
(3+N)		00h	00h (ACK) FFh (NACK)
(4+N)		Data Byte	Undefined

Figure B-16 SPIMCI WRITE Transaction

10

11

1 B.3.8.4. Write n Bytes to a Given Byte Address

2 The Initiator generates a START condition.

3 After the START condition, the Initiator sends 32 bits of Transaction Control with the Write bit set to 1.
 4 Transaction Control contains information on the bank, page, and address to be read. The Size field of the
 5 instruction data is set to (n-1) indicating a sequential multi byte write from the given address.

6 Then the Initiator sends 16 bits of flow control data with a logic 0. During these 16 bits, the Initiator can read
 7 its IITO line to determine if the Target is Busy (FFh) or Idle (00h).

8 If the target indicates it is idle, then the Initiator can clock out the data to be written on the IOTI line.

9 The Initiator generates a STOP condition indicating end of a command.

Byte	Bit	IOTI	IITO
0	7	1 (Write)	Undefined
	6:3	Bank#	
	2:0		
1			Undefined
2		Page#	Undefined
3		Address#	Undefined
(3+1) ... (3+N-1)		00h	Undefined
(3+N)		00h	00h (ACK) FFh (NACK)
(4+N) ... (4+N+n-2)		Data Byte 1 ... Data Byte n-1	Undefined
(4+N+n-1)		Data Byte n	Undefined

10 **Figure B-17 SPIMCI WRITE N Transaction**

11

B.3.8.5. Test Command

There is no specific TEST command. Instead, the Initiator checks the target response when trying to read one byte from byte address 0 in Low Memory, where the target responds with ACK (idle) or NACK (busy).

In more detail:

- After the START condition, the Initiator sends 32 bits of Transaction Control with all 32 bits set to 0.
- Then the Initiator sends 8·N zero valued bits, allowing the target to decode the address to be read.
- The Initiator then reads the flow control response byte from the IITO line to determine if the target is Busy (FFh) or Idle (00h).
- The Initiator eventually generates a STOP condition indicating the end of the command.

Byte	Bit	IOTI	IITO
0	7	0 (Read)	Undefined
	6:3	0 (Bank 0)	
	2:0	0 (Read 1 byte)	
1			Undefined
2		00h (Page 0)	Undefined
3		00h (Byte 0)	Undefined
(3+1) ... (3+N-1)		00h	Undefined
(3+N)		00h	00h (ACK) FFh (NACK)
(4+N)		Undefined	Don't care

Figure B-18 SPIMCI TEST Transaction

B.3.9 Transaction Flow Control Mechanism

B.3.9.1. Stretching Current Transaction

Note: Unlike in I2CMCI, there is no option for the target to delay completion of a current transaction.

B.3.9.2. Rejecting Subsequent Transaction (Transaction Hold-Off)

Transaction hold-off is a temporary phase when the target rejects new transactions by signaling NACK in the last byte of the Flow Control phase.

The maximum allowed transaction hold-off duration is specified in section B.3.10. The target exercises transaction hold-off when it is unable or not allowed to accept a new transaction in its current state or condition. For instance, after a properly terminated write transaction, the target may need to reject subsequent transactions in order to internally finalize the WRITE operation represented by the previous transaction, due to synchronization requirements specified in section 5.2.4).

B.3.10 Timing Specification

For a complete management stack specification, timing specification should be added here.

1 Appendix C Examples of Application Advertisements

2 This appendix presents Application advertisement examples to illustrate how fixed or programmable modules
 3 will advertise the Applications they support (as described in chapter 6). The following cases are described:

- 4 Programmable 400GBASE-DR4 module that can also operate as 4x100GBASE-DR
- 5 Fixed 400BASE-SR8 module
- 6 Programmable 400G-SR8 module that can also operate as a 2x200G, 4x100G, or 8x50G breakout.
- 7 Dual Use Active Optical Cable for 1x400G or 2x200G Applications

9 The first example in Table C-1 400GBASE-DR4 Transceiver with Dual Application Advertising illustrates the
 10 Application advertisement of a module that operates as a 400GBASE-DR4 transceiver but supports alternative
 11 operation as four integrated parallel 100GBASE-DR transceivers.

12 With this example, the host can select Application 1 or Application 2 (by programming the AppSel value
 13 AppSel=1 or AppSel=2, respectively) in one of the Staged Control Sets. AppSel=1 (400GAUI-8 to 400GBASE-
 14 DR4) is the default Application populated in the Active Control Set at power-on.

15 **Table C-1 400GBASE-DR4 Transceiver with Dual Application Advertising**

Byte	Bits	AppSel Code	Name	Value	Description
85	7-0	N/A	MediaType	02h	Optical Interfaces: SMF
86	7-0		HostInterfaceID	11h	400GAUI-8 C2M
87	7-0		MediaInterfaceID	1Ch	400GBASE-DR4
88	7-4		HostLaneCount	8	8 host lanes
	3-0		MediaLaneCount	4	4 media lanes
89	7-0		HostLaneAssignmentOptions	01h	Permissible first host lane number for Application: lane 1
01h:176	7-0		MediaLaneAssignmentOptions	01h	Permissible first media lane number for Application: lane 1
90	7-0	0010b	HostInterfaceID	0Dh	100GAUI-2 C2M
91	7-0		MediaInterfaceID	14h	100GBASE-DR
92	7-4		HostLaneCount	2	2 host lanes
	3-0		MediaLaneCount	1	1 media lane
93	7-0		HostLaneAssignmentOptions	55h	Permissible first host lane number for Application: lanes 1, 3, 5, and 7
01h:177	7-0		MediaLaneAssignmentOptions	0Fh	Permissible first media lane number for Application: lanes 1, 2, 3, 4
94	7-0	0011b	HostInterfaceID	FFh	End of list of supported Applications
95	7-0		MediaInterfaceID	00h	
96	7-4		HostLaneCount	0	
	3-0		MediaLaneCount	0	
97	7-0		HostLaneAssignmentOptions	00h	
01h:178	7-0		MediaLaneAssignmentOptions	00h	

The second example in Table C-2 400GBASE-SR8 Fixed Transceiver Application Advertising illustrates the Application advertisement for a fixed 400GBASE-SR8 transceiver that does not support other Applications.

With this example, the host can set only AppSel=1 in one of the Staged Control Sets

Table C-2 400GBASE-SR8 Fixed Transceiver Application Advertising

Byte	Bits	AppSel Code	Name	Value	Description
85	7-0	N/A 0001b	MediaType	01h	Optical Interfaces: MMF
86	7-0		HostInterfaceID	11h	400GAUI-8 C2M
87	7-0		MediaInterfaceID	10h	400GBASE-SR8
88	7-4		HostLaneCount	8	8 host lanes
	3-0		MediaLaneCount	8	8 media lanes
89	7-0		HostLaneAssignmentOptions	01h	Permissible first host lane number for Application: lane 1
01h:176	7-0		MediaLaneAssignmentOptions	01h	Permissible first media lane number for Application: lane 1
90	7-0	0010b	HostInterfaceID	FFh	End of list of supported Applications
91	7-0		MediaInterfaceID	00h	
92	7-4		HostLaneCount	0	
	3-0		MediaLaneCount	0	
93	7-0		HostLaneAssignmentOptions	00h	
01h:177	7-0		MediaLaneAssignmentOptions	00h	

The third example in Table C-3 400GBASE-SR8 Transceiver supporting 200GBASE-SR4, 100GBASE-SR2 and 50GBASE-SR illustrates the Application advertisement of a module that operates as a 400GBASE-SR8 transceiver by default but supports breakout to integrated parallel transceivers at a variety of speeds (2x200G, 4x100G, 8x50G).

With this example, the host can write either AppSel 1, 2, 3, or 4 in one of the Staged Control Sets, where AppSel=1 (400GAUI-8 to 400G-SR8) is the default Application populated in the Active Control Set at power-on.

Table C-3 400GBASE-SR8 Transceiver supporting 200GBASE-SR4, 100GBASE-SR2 and 50GBASE-SR

Byte	Bits	AppSel Code	Name	Value	Description
85	7-0	N/A	MediaType	01h	Optical Interfaces: MMF
86	7-0		HostInterfaceID	11h	400GAUI-8 C2M
87	7-0		MediaInterfaceID	10h	400GBASE-SR8
88	7-4		HostLaneCount	8	8 host lanes
	3-0		MediaLaneCount	8	8 media lanes
89	7-0		HostLaneAssignmentOptions	01h	Permissible first host lane number for Application: lane 1
01h:176	7-0		MediaLaneAssignmentOptions	01h	Permissible first media lane number for Application: lane 1
90	7-0	0001b	HostInterfaceID	0Fh	200GAUI-4 C2M
91	7-0		MediaInterfaceID	0Eh	200GBASE-SR4
92	7-4		HostLaneCount	4	4 host lanes
	3-0		MediaLaneCount	4	4 media lanes
93	7-0		HostLaneAssignmentOptions	11h	Permissible first host lane number for Application: lanes 1 and 5
01h:177	7-0		MediaLaneAssignmentOptions	11h	Permissible first media lane number for Application: lanes 1, 5
94	7-0		HostInterfaceID	0Dh	100GAUI-2 C2M
95	7-0	0010b	MediaInterfaceID	0Ch	100GBASE-SR2
96	7-4		HostLaneCount	2	2 host lanes
	3-0		MediaLaneCount	2	2 media lanes
97	7-0		HostLaneAssignmentOptions	55h	Permissible first host lane number for Application: lanes 1, 3, 5, and 7
01h:178	7-0		MediaLaneAssignmentOptions	55h	Permissible first media lane number for Application: lanes 1, 3, 5, 7
98	7-0	0011b	HostInterfaceID	0Ah	50GAUI-1 C2M
99	7-0		MediaInterfaceID	07h	50GBASE-SR
100	7-4		HostLaneCount	1	1 host lane
	3-0		MediaLaneCount	1	1 media lane
101	7-0		HostLaneAssignmentOptions	FFh	Permissible first host lane number for Application: lanes 1, 2, 3, 4, 5, 6, 7, 8
01h:179	7-0		MediaLaneAssignmentOptions	FFh	Permissible first media lane number for Application: lanes 1, 2, 3, 4, 5, 6, 7, 8
102	7-0		HostInterfaceID	FFh	End of list of supported Applications
103	7-0		MediaInterfaceID	00h	
104	7-4		HostLaneCount	0	
	3-0		MediaLaneCount	0	
105	7-0		HostLaneAssignmentOptions	00h	
01h:180	7-0		MediaLaneAssignmentOptions	00h	

The fourth example in Table C-4 8x50G AOC Application Advertising Example illustrates the Application advertisement for an 8x50G Active Optical Cable (AOC) that supports one 8x50G host interface or two parallel 4x50G host interfaces.

With this example, the host could write either AppSel=1 or AppSel=2 in one of the Staged Control Sets. AppSel=1 (400GAUI-8) is the default Application populated in the Active Control Set at power-on.

Table C-4 8x50G AOC Application Advertising Example

Byte	Bits	AppSel Code	Name	Value	Description
85	7-0	N/A 0001b	MediaType	04h	Active cables
86	7-0		HostInterfaceID	11h	400GAUI-8 C2M
87	7-0		MediaInterfaceID	03h	AOC with BER < 2.4e-4
88	7-4		HostLaneCount	8	8 host lanes
	3-0		MediaLaneCount	8	8 media lanes
89	7-0		HostLaneAssignmentOptions	01h	Permissible first host lane number for Application: lane 1
01h:176	7-0		MediaLaneAssignmentOptions	01h	Permissible first media lane number for Application: lane 1
90	7-0	0010b	HostInterfaceID	0Fh	200GAUI-4 C2M
91	7-0		MediaInterfaceID	03h	AOC with BER < 2.4e-4
92	7-4		HostLaneCount	4	4 host lanes
	3-0		MediaLaneCount	4	4 media lanes
93	7-0		HostLaneAssignmentOptions	11h	Permissible first host lane number for Application: lanes 1 and 5
01h:177	7-0		MediaLaneAssignmentOptions	11h	Permissible first media lane number for Application: lanes 1, 5
94	7-0	0011b	HostInterfaceID	FFh	End of list of supported Applications
95	7-0		MediaInterfaceID	00h	
96	7-4		HostLaneCount	0	
	3-0		MediaLaneCount	0	
97	7-0		HostLaneAssignmentOptions	00h	
01h:178	7-0		MediaLaneAssignmentOptions	00h	

Appendix D Examples of Initialization and Deinitialization

This appendix includes example scenarios illustrating possible sequences of events and interactions between host and module during module and Data Path initialization or deinitialization.

Refer to section 6.3.2 for information on the Module State Machine and to section 6.3.3 for information on Data Path State Machines, Applications, and Control Sets.

The following example scenarios illustrate how a host could power up and initialize, or deinitialize and power down a module:

Scenario	Description	Section
Quick hardware initialization	Power up and initialization without host software interaction	0
Quick software initialization	Power up and initialization with minimal host software interaction	D.1.2
Software configuration and initialization	Power up and initialization with module configuration by host	D.1.3
Hardware deinitialization	Power-down sequence using hardware control	D.2.1
Software deinitialization	Power-down sequence using software control	D.2.2

Note: For simplicity, the examples may deliberately omit minor behavioral variations depending on module advertisement or configuration (for instance it is assumed that all state change notifications are flagged)

D.1 Initialization Examples

The power-up initialization behavior of pausing in the ModuleLowPwr state (for host software control), or of passing to full operation through the ModuleLowPwr state (without host control) is determined by the settings LowPwrRequestSW and LowPwrAllowRequestHW and by the LowPwrRequestHW hardware signal (Table 6-12).

The default register settings LowPwrAllowRequestHW=1 and LowPwrRequestHW=0 (see Table 8-11) require the LowPwrRequestHW hardware signal to be ASSERTED when host software controlled start-up is desired.

D.1.1 Quick Hardware Initialization

This section describes an example of a simple module power-up sequence where the module fully powers up under hardware control, without host software intervention.

This scenario has the following characteristics:

- Module is powered-up from an un-powered state (LowPwrAllowRequestHW=1 and LowPwrRequestSW=0)
- Module is fully powered-up under hardware control (LowPwrRequestHW=**DEASSERTED**)
- Host uses the default settings for the selected Application.
- Host does not use custom signal integrity settings (i.e. ExplicitControl indicator)
- Host does not perform speed negotiation

#	Host Action	Module Action	Module State (M) Data Path State (D)
0	Host applies Vcc, deasserts LowPwrRequestHW, asserts ModSel (if supported), and deasserts Reset. Host ensures that host transmitters are configured and enabled to output a valid signal for the default Application in the module, prior to step 1 below.		
1	Hot Plug		
2	Host detects module presence, waits for Interrupt assertion	Module powers up and initializes management interface, sets power on defaults, such as LowPwrRequestSW=0, LowPwrAllowRequestHW=1, DPDeinit=00h, OutputDisableTx = 0 and OutputSquelchForceTx = 0, and writes the power on default Data Path configurations into the Active Control Set and Staged Control Set 0	M=MgmtInit D=DPDeactivated

#	Host Action	Module Action	Module State (M) Data Path State (D)
3		Module sees LowPwrS transition signal is FALSE on entry into ModuleLowPwr and transitions to ModulePwrUp	M=ModuleLowPwr D=DPDeactivated
4		Module powers up to High Power Mode	M=ModulePwrUp D=DPDeactivated
5		Module sets ModuleStateChangedFlag on entry into ModuleReady	M=ModuleReady D=DPDeactivated
6	Host detects assertion of Interrupt and reads all Flag registers, which deasserts Interrupt	Module sees DPDeinitS transition signal is FALSE and transitions all Data Path states to DPInit	
7	Host waits for second Interrupt assertion to indicate completion of Data Path initialization	Module initializes all Data Paths according to the configuration in the Active Control Set.	M=ModuleReady D=DPInit
8		Module sees DPDeactivateS transition signal is FALSE and transitions all Data Path states to DPTxTurnOn	M=ModuleReady D=DPInitialized
9		Modules enables all Tx outputs	M=ModuleReady D=DPTxTurnOn
10		Module sets the Data Path State Changed Flags to 1 on entry into DPActivated	M=ModuleReady D=DPActivated
11	Host detects assertion of Interrupt and reads all Flag registers, which deasserts Interrupt	Module waits for host action	
12	Host uses the activated Data Path to carry live traffic		



D.1.2 Quick Software Initialization

This section describes an example of a simple module power-up sequence where the module gets under host software control in ModuleLowPwr, such that the host can validate module capabilities and power dissipation before letting the module complete its initialization, using the default Application and Data Path configuration.

This scenario has the following characteristics:

- Module is powered-up from an un-powered state (LowPwrAllowRequestHW=1 and LowPwrRequestSW=0)
- Module function is initialized under software control (LowPwrRequestHW=**ASSERTED**)
- Host uses the default settings for the selected Application.
- Host does not use custom signal integrity settings (i.e. ExplicitControl indicator)
- Host does not perform speed negotiation

#	Host Action	Module Action	Module State (M) Data Path State (D)
0	Host applies Vcc, asserts LowPwrRequestHW and ModSel (if supported), and deasserts Reset		
1	Hot Plug		
2	Host detects module presence, waits for Interrupt assertion	Module powers up and initializes management interface, sets power on defaults, such as LowPwrRequestSW=0, LowPwrAllowRequestHW=1, DPDeinit=00h, OutputDisableTx = 0 and OutputSquelchForceTx = 0, and writes the power on default Data Path configurations into the Active Control Set and Staged Control Set 0	M=MgmtInit D=DPDeactivated

#	Host Action	Module Action	Module State (M) Data Path State (D)
3		Module sets ModuleStateChangedFlag on entry into ModuleLowPwr	M=ModuleLowPwr D=DPDeactivated
4	Host detects assertion of Interrupt and reads the Flag registers, which deasserts Interrupt	Module waits for host action	
5	Host reads module power requirements and Data Path configuration information		
6	Host configures and enables host transmitters such that they output a valid signal for the default Application towards the module		
7	Host clears LowPwrAllowRequestHW ¹ bit, which effectively initiates a module transition to High Power Mode (ModuleReady)		
8	Host waits for Interrupt assertion to indicate completion of transition to High Power Mode	Module sees LowPwrS transition signal become FALSE and transitions to ModulePwrUp	M=ModulePwrUp D=DPDeactivated
9		Module powers up to High Power Mode	
10		Module sets ModuleStateChangedFlag on entry into ModuleReady	M=ModuleReady D=DPDeactivated
11	Host detects assertion of Interrupt and reads all Flag registers, which deasserts Interrupt	Module sees DPDeinitS transition signal is FALSE and transitions all Data Path states to DPInit	
12	Host waits for second Interrupt assertion to indicate completion of Data Path initialization	Module initializes all Data Paths according to the configuration in the Active Control Set.	M=ModuleReady D=DPInit
13		Module sees DPDeactivateS transition signal is FALSE and transitions all Data Path states to DPTxTurnOn	M=ModuleReady D=DPInitialized
14		Modules enables all Tx outputs	M=ModuleReady D=DPTxTurnOn
15		Module sets the DPStateChangedFlag bits to 1 on entry into DPActivated	M=ModuleReady D=DPActivated
16	Host detects assertion of Interrupt and reads all Flag registers, which deasserts Interrupt	Module waits for host action	
17	Host uses the activated Data Path to carry live traffic		

1

D.1.3 Software Configuration and Initialization

2

This section describes an example of a simple module power-up sequence where the module gets under host software control in ModuleLowPwr state, such that the host can validate module capabilities and its maximum power dissipation and subsequently configure the module function, before letting the configured module complete its functional initialization. This method may be the most common initialization method used for Ethernet applications.

3

In this example, host software powers up the module and the Data Paths in two separate steps.

4

This scenario has the following characteristics:

5

- a. Module is powered-up from an un-powered state (LowPwrAllowRequestHW=1 and LowPwrRequestSW=0)

6

¹ To allow full SW control from now on, the host disables the LowPwrRequestHW signal. Alternatively the host could also deassert LowPwrRequestHW signal.

- 1 b. Module function is initialized under host software control (LowPwrRequestHW=**ASSERTED**)
- 2 c. Host selects one of the Applications advertised by the module
- 3 d. Host initializes one instance of this Application (i.e. one Data Path).
- 4 e. Host uses the default settings for the selected Application.
- 5 f. Host does not use custom signal integrity settings (i.e. ExplicitControl indicator is 0)
- 6 g. Host uses only staged Control Set 0 to configure the module (optional Staged Control Set 1 is not used).
- 7 h. Host does not perform speed negotiation

8
9 Steps 1-11 provide an example of how to power-up a module into the ModuleReady state, starting from an un-
10 powered state. Steps 12-28 provide an example for the initialization and activation of a single Data Path within
11 a module (these steps can also be used for the initialization and activation of subsequent Data Paths, e.g. in
12 the case of breakout, etc.).

#	Host Action	Module Action	Module State (M) Data Path State (D)
0	Host applies Vcc, asserts LowPwrRequestHW and ModSel (if supported), and deasserts Reset		
1	Hot Plug		
2	Host detects module presence, waits for Interrupt assertion	Module powers up and initializes management interface, sets power on defaults, such as LowPwrRequestSW=0, LowPwrAllowRequestHW=1, DPDeinit=00h, OutputDisableTx = 0 and OutputSquelchForceTx = 0, and writes the power on default Data Path configurations into the Active Control Set and Staged Control Set 0.	M=MgmtInit D=DPDeactivated
3		Module sets ModuleStateChangedFlag on entry into ModuleLowPwr	M=ModuleLowPwr D=DPDeactivated
4	Host detects assertion of Interrupt and reads all Flag registers, which deasserts Interrupt	Module waits for host action	
5	Host reads module power requirements		
6	Host writes FFh to the DPDeinit register to prevent automatic Data Path initialization when the module reaches ModuleReady, and writes FFh into the OutputDisableTx register to prevent automatic Data Path activation when the Data Path state reaches DPInitialized.		
7	Host clears LowPwrAllowRequestHW ¹ bit, which effectively initiates a module transition to High Power Mode (ModuleReady)		
8	Host waits for Interrupt assertion to indicate completion of transition to High Power Mode	Module sees LowPwrS transition signal become FALSE and transitions to ModulePwrUp	
9		Module powers up to High Power Mode	M=ModulePwrUp D=DPDeactivated
10		Module sets ModuleStateChangedFlag on entry into ModuleReady	M=ModuleReady D=DPDeactivated

¹ To allow full SW control from now on, the host disables the LowPwrRequestHW signal. Alternatively the host could also deassert LowPwrRequestHW signal.

#	Host Action	Module Action	Module State (M) Data Path State (D)
11	Host detects assertion of Interrupt and reads all Flag registers, which deasserts Interrupt	Module waits for host action	
12	Host reads Application advertising registers		
13	Host writes desired AppSel code into applicable Data Path Configuration Control registers in Staged Control Set 0		
14	Host configures and enables host transmitters such that they output a valid signal for the selected Application towards the module		
15	Host writes a 1 to the bits in the ApplyDPInit register in Staged Control Set 0, corresponding to host lanes of the Data Path the host wants to configure. The host lanes that are not part of the Data Path are masked (bits cleared) in the register write. The host performs this step with a single write to the ApplyDPInit register.		
16	Host waits for module action	Module validates the configuration requested in Staged Control Set 0. If the configuration was found to be valid, the module copies the contents to the Active Control Set, updates DPInitPending, and finally updates the ConfigStatus fields.	
17	Host reads the ConfigStatus fields to confirm that the requested configuration was validated and accepted by the module for all lanes of the selected Data Path.	Module waits for host action	
18	Host requests initialization of the newly configured Data Path by clearing the DPDeinit bits representing the host lanes of the Data Path being initialized. The host preserves the existing bit values for the other lanes in the DPDeinit register (using read-modify-write). The new contents of the DPDeinit register are written by the host in a single write transaction.		
19	Host waits for Interrupt assertion to indicate completion of Data Path initialization	Module sees DPDeinitS transition signal is FALSE and transitions the Data Path state to DPInit	
20		Module initializes the Data Path and clears its DPInitPending bits	M=ModuleReady D=DPInit
21		Once the Data Path is initialized the module sets the DPStateChangedFlag bits, on entry into DPInitialized	M=ModuleReady D=DPInitialized
22	Host detects assertion of Interrupt and reads all Flags to clear the interrupt	Module waits for host action	

#	Host Action	Module Action	Module State (M) Data Path State (D)
23	Host requests activation of the Data Path by clearing the bits representing its media lanes in the OutputDisableTx register. The host preserves the existing bit values for the other media lanes in the OutputDisableTx register (using read-modify-write). Host writes the new contents of the OutputDisableTx register in a single write transaction.		
24	Host waits for Interrupt assertion to indicate completion of Data Path activation	Module sees DPDeactivateS transition signal is FALSE and transitions the Data Path state to DPTxTurnOn. Modules enables all Tx outputs for the Data Path being activated	M=ModuleReady D=DPTxTurnOn
25		Once the Data Path is activated the module sets the DPStateChangedFlag bits, on entry into DPActivated	M=ModuleReady D=DPActivated
26		Module waits for host action	
27	Host detects assertion of Interrupt and reads all Flag registers, which deasserts Interrupt.		
28	Host uses the activated Data Path to carry live traffic		

1 D.2 Deinitialization Examples

2 D.2.1 Hardware Deinitialization

3 This section describes an example of a simple module power-down sequence where the module powers down
 4 under hardware control (LowPwrRequestHW hardware signal transitions from DEASSERTED to ASSERTED)
 5 without software interaction.

6 This scenario has the following characteristics:

- 7 a. Module was previously powered up under hardware control (LowPwrRequestHW=DEASSERTED)
 8 b. At least one Data Path is in the DPActivated state

#	Host Action	Module Action	Module State (M) Data Path State (D)
0	Module is powered up with at least one Data Path activated. LowPwrRequestHW = DEASSERTED, Reset = DEASSERTED, ModSel=ASSERTED (if supported).	Initial conditions: Module fully configured and powered, and with at least one Data Path in the DPActivated state.	M=ModuleReady D=DPActivated
1	Host asserts the LowPwrRequestHW signal	Module sees DPDeactivateS transition signal is TRUE and transitions all Data Path states to DPTxTurnOff	M=ModuleReady D=DPTxTurnOff
2		All Data Path states transition to DPInitialized without any further host action.	M=ModuleReady D=DPInitialized
3		Module sees DPReDeinitS is TRUE and transitions all Data Path states to DPDeinit	M=ModuleReady D=DPDeinit
4		Module disables all Tx outputs and deinitializes Data Path resources	M=ModuleReady D=DPDeactivated
5		Module sets the Data Path State Changed Flags to 1 on entry into DPDeactivated	M=ModuleReady D=DPDeactivated
6	Host detects assertion of Interrupt and reads all Flag registers, which deasserts Interrupt	Module sees LowPwrExS is TRUE and transitions the module state to ModulePwrDn	M=ModulePwrDn D=DPDeactivated
7	Host waits for second Interrupt assertion to indicate completion of module power down	Module reduces the module power to low power mode levels	M=ModulePwrDn D=DPDeactivated
8		Module sets ModuleStateChangedFlag on entry into ModuleLowPwr	M=ModuleLowPwr D=DPDeactivated
9	Host detects assertion of Interrupt and reads all Flags to clear the interrupt	Module waits for host action	

D.2.2 Software Deinitialization

This section describes an example of a simple module power-down sequence where the module powers down under software control (LowPwrRequestHW hardware signal = ASSERTED). This example uses a two-step power down process where all Data Paths are first deactivated and then the module is powered down using the LowPwrAllowRequestHW bit. One scenario where this might be used is where a host has deactivated all Data Paths but wants to go one step further and also transition the module to low power mode. Note that this approach does require a two-step power up process if the host wants to reactivate one of the Data Paths at a later point in time, where the host has to first transition the module to high power mode before activating the Data Path.

This scenario has the following characteristics:

- a. Module was previously powered up under software control (LowPwrRequestHW=ASSERTED)
- b. At least one Data Path is in the DPActivated state

#	Host Action	Module Action	Module State (M) Data Path State (D)
0	Module is powered up with at least one Data Path activated. LowPwrRequestHW = ASSERTED, ResetL = DEASSERTED, ModSel=ASSERTED (if supported).	Initial condition: Module fully configured and powered, and with at least one Data Path in the DPActivated state.	M=ModuleReady D=DPActivated
1	The host sets the DPDeinit bits for all host lanes in the applicable Data Path to 1.	Module sees DPDeactivateS transition signal is TRUE and transitions the Data Path state to DPTxTurnOff	M=ModuleReady D=DPTxTurnOff
2		All Data Path states transition to DPInitialized without any further host action.	M=ModuleReady D=DPInitialized
3		Module sees DPReDeinitS is TRUE and transitions all Data Path states to DPDeinit	M=ModuleReady D=DPDeinit
4		Module disables the Tx outputs and deinitializes Data Path resources for all Data Paths.	M=ModuleReady D=DPDeactivated
5		Module sets the Data Path State Changed Flags to 1 on entry into DPDeactivated	M=ModuleReady D=DPDeactivated
6		Module waits for host action	M=ModuleReady D=DPDeactivated
7		Module sees LowPwrExS is TRUE and transitions the module state to ModulePwrDn	M=ModulePwrDn D=DPDeactivated
8		Module reduces the module power to low power mode levels	M=ModulePwrDn D=DPDeactivated
9		Module sets ModuleStateChangedFlag on entry into ModuleLowPwr	M=ModuleLowPwr D=DPDeactivated
10		Module waits for host action	M=ModuleLowPwr D=DPDeactivated

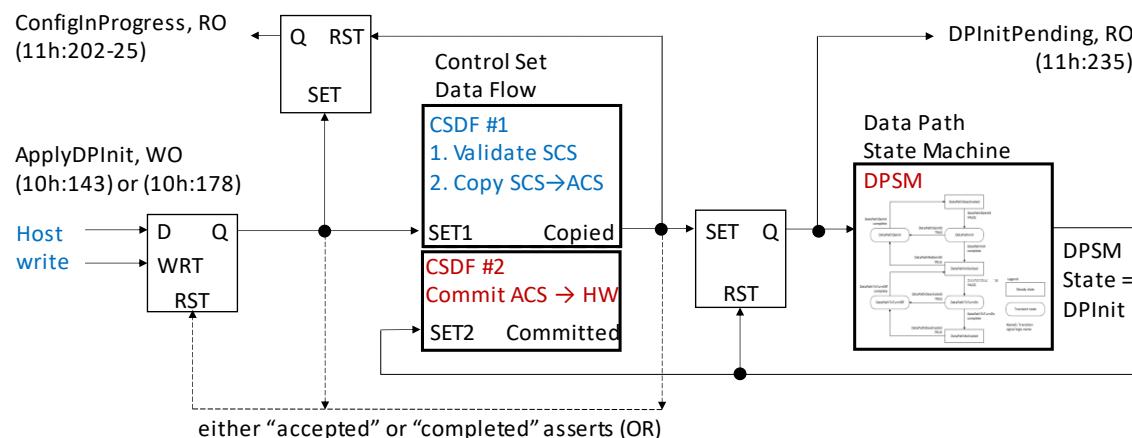
Appendix E Illustration of Applying Control Sets

This appendix illustrates different module behaviors after a host writing to one of the Apply register, depending on the advertisement SteppedConfigOnly (00h:2.6)

E.1 Default Behavior (SteppedConfigOnly = 0)

CMIS 5.0 Default ApplyDPIInit Processing

SteppedConfigOnly = 0, RO (00h:2.6) – **DPIInitPending** affects **DPSM** and may initiate automatic commissioning



Not shown: Module must ignore host write to ApplyDPIInit while ConfigStatus = ConfigInProgress
 Module is free to use ApplyDPIInit and to clear on command completion or on acceptance
 Module ignores writes in transient DPSM states

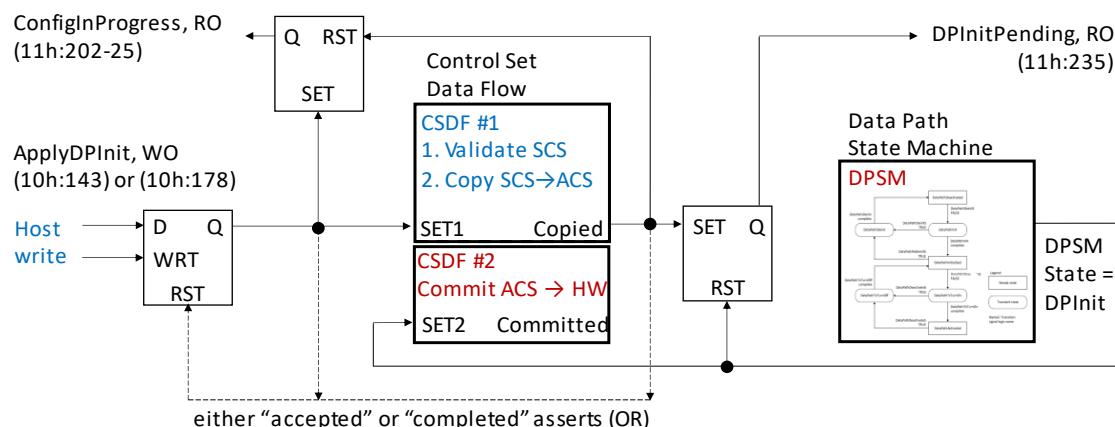
Figure E-19 ApplyDPIInit (default: SteppedConfigOnly=0)

Note that ApplyImmediate with state dependent behavior is supported in steady DPSM states as per Table 6-3

E.2 Restricted Behavior (SteppedConfigOnly = 1)

CMIS 5.0 Restricted ApplyDPIInit Processing

SteppedConfigOnly = 1, RO (00h:2.6) – **DPIInitPending** is pure provisioning and does **not** affect DPSM



Not shown: Module must ignore host write to ApplyDPIInit while ConfigStatus = ConfigInProgress
 Module is free to use ApplyDPIInit and to clear on command completion or on acceptance

Figure E-20 ApplyDPIInit (restricted: SteppedConfigOnly=1)

Note that ApplyImmediate is not supported in this configuration.

1 Appendix F Examples of Diagnostic Features Usage

2 This Appendix contains usage recommendations and examples for the diagnostic features available on Pages
3 13h and 14h.

4 *Note: For use of these diagnostic features the module should be in the ModuleReady state. An Application
5 should be selected but otherwise no specific relation to Data Path states is defined for these diagnostic features.*

6 F.1 Enabling and Disabling Pattern Generator (Host or Media Side)

7 The following procedure contains a recommended sequence of register accesses to set the module into Host
8 side (or Media side) pattern generation mode. Host side registers are used in the examples, with media side
9 registers shown in parentheses.

- 10 1. Write Bank Select and Page Select registers 00h:127-128 to select Page 13h on the appropriate Bank
- 11 2. Configure desired mode of operation in Byte 13h:177
- 12 3. Configure desired pattern generator configuration and lane pattern in Bytes 13h:145-151 (13h:153-
13 159)
- 14 4. Configure the desired control options in Bytes 13h:176-179
- 15 5. Enable the pattern generator on the selected lanes by writing to 13h:144.7-0 (13h:152.7-0)

16 After the above sequence of commands, the host side electrical (media side electrical or optical) output will be
17 generating the selected pattern for the enabled lanes.

19 *Note: The pattern generation feature may vary by module. On some modules, when in pattern generation
20 mode, per lane control is not provided and all lanes may be generating the pattern. On other modules, such as
21 modules design to support break-out, the selected lanes may be in pattern generation mode while other lanes
22 are either disabled or in normal mission mode.*

23 To disable the pattern generator on selected host side (media side) lanes:

- 24 6. Write to 13h:144.7-0 (13h:152.7-0) with the relevant bits representing the selected lanes cleared

25 When pattern generation is disabled on selected lanes, those lanes are expected to revert to mission mode if
26 possible. In some modules, mission mode can only be achieved if none of the module interfaces, host side or
27 media side are in pattern generation mode. A module reset may be used to guarantee that the module reverts
28 back to mission mode. Otherwise, the host has to ensure that all pattern generation modes on all lanes for both
29 the host and media sides are disabled for these types of modules.

30 F.2 Enabling and Disabling Pattern Checker (Host or Media Side)

31 The following procedure contains the recommended sequence of register accesses to set the module into Host
32 side (or Media side) pattern checking mode. Host side registers are used in the example, with media side
33 registers shown in parentheses.

- 34 1. Write Bank Select and Page Select registers 00h:127-128 to select Page 13h on the appropriate Bank.
- 35 2. Configure desired mode of operation in Byte 13h:177
- 36 3. Configure desired pattern checker configuration and lane pattern in 13h:161-167 (13h:169-175)
- 37 4. Configure desired control options in 13h:176-179
- 38 5. Enable the pattern checker on the selected lanes by writing to 13h:160.7-0 (13h:168.7-0)

40 When the host enables the PRBS checker, the module is expected to reset the error counters and enable the
41 error counters to begin counting. The behavior of the error counters is defined by Byte 13h:177.

42 To disable the pattern checker (at any time) including in the middle of a gated error count operation

- 43 6. Clear the bits representing the desired lanes in byte 13h:160.7-0 (13h:168.7-0)

44 If the pattern checker is disabled in the middle of a gated operation, all of the error counters are undefined.

45 Disabling in the middle of a gated count is considered an abort operation by the module.

46 The following section details some of the error counter configurations and their expected behavior.

47 F.3 Reading Pattern Checker Error Counters

48 There are many scenarios in which pattern checkers can be used, based on the configuration in Byte 13h:177.

1 The following sections describe recommended host write sequences for some commonly used scenarios.

2 These recommended sequences are provided for information. Media side registers are provided in these
3 examples, but the same sequences can also be applied to host side registers.

4 The example Application for these sequences has 8 host side lanes and 4 media side lanes.

5 F.3.1 Not Gated (Continuous) Error Counters, Individual Lanes

6 In this scenario the error statistics are collected as cumulative (non-periodic, open-end) statistics

7 Host Selected Mode of Operation

- 8 a. 13h:177.3-1 = 000b (not gated)
- 9 b. 13h:177.5 = 0 (do not hold checkers in reset)
- 10 c. 13h:177.7 = 0 (reset error counts per lane)
- 11 d. 13h:129.4 = 0 (periodic update disabled)
- 12 e. 13h:177.0 = X (don't care update interval)

13 Host Write Sequence

- 14 1. Write 13h:160.7-0 (13h:168.7-0) to enable the pattern checker on selected media side lanes.
 - 15 a. The module will apply the configuration and control options to the enabled lanes.
 - 16 b. Since the configuration is not gated and periodic update is disabled, the latest Error counters
 - 17 are available "on demand" when the host reads the current pattern checker data.
- 18 2. Write 14h to the Page select register. Since error information is on demand, module may take additional
- 19 time to provide the selected diagnostics data.
- 20 3. Write 02h (04h) into the Diagnostics Selector 14h:128 to select lanes 1-4 for the host side (media side)
- 21 or write a 03h (05h) to select lanes 5-8 (e.g. for DR4, FR4, LR4)
- 22 4. Module will perform a read of the pattern checker error counters when byte 128 is written. (Module will
- 23 not clear the error counters.)
- 24 5. Read Byte 14h:138 (14h:139) to ensure the pattern checker has not lost lock.
- 25 6. Read Bytes 14h:192-255 to obtain error counters and total bits.

28 F.3.2 Not Gated (Continuous) Error Counters, Individual Lanes, Reset Error Counter

29 In this scenario error statistics are collected as interval statistics. The interval length per lane is under host
30 control

31 Configuration Assumptions

- 32 a. 13h:177.3-1 = 000b (not gated)
- 33 b. 13h:177.5 = 0 (do not hold checkers in reset)
- 34 c. 13h:177.7 = 0 (reset error counts per lane)
- 35 d. 13h:129.4 = 1 (periodic update enabled)
- 36 e. 13h:177.0 = 0 (**1 second update interval**)

37 Host Write Sequence

- 38 1. Write bits 13h:160.7-0 (13h:168.7-0) to enable the pattern checker on selected media side lanes.
 - 39 a. The module will apply the configuration and control options to the enabled lanes.
 - 40 b. Since the configuration is not gated and periodic update is enabled, the module may provide
 - 41 the error information from the last update period.
 - 42 c. Error counters may also be available "on demand" when the host reads the current pattern
 - 43 checker data (this behavior is deliberately left to be vendor dependent).
- 44 2. Write 14h to the Page select register. To provide better response to host, the module may return the
- 45 last updated error information.
- 46 3. Write 02h (or 04h) into the Diagnostics Selector 14h:128 to select lanes 1-4 for the host side (media
- 47 side) or write a 03h (05h) to select host (media) lanes 5-8 (e.g. for DR8, FR8, LR8)
- 48 4. Read byte 14h:138 (14h:139) to ensure the pattern checker has not lost lock.
- 49 5. To reset the error information the host may set Bit 13h:177.5. When set:
 - 50 a. Module (may also read from the data path chips and then) freezes the current error information.
 - 51 b. Module copies current internal pattern checker counters into the counter results.
- 52 6. Host may write 11h-15h to the Diagnostics Selector 14h:128 to select the collected error information
- 53 results, and then read bytes 14h:192-255 to obtain the results of BER or of error counters and total
- 54 bits.

- 1 7. **To restart error counting** the host may clear Bit 13h:177.5. This will reset the current error
2 information (selector 01h-05h) without affecting the error information results (selector 11h-15h).

3 F.3.3 Not Gated (Continuous) Error Counters, All Lanes, all Banks

4 In this scenario the error statistics are operated as host-controlled interval statistics

5 Configuration Assumptions

- 6 a. 13h:177.3-1 = 000b (not gated)
7 b. 13h:177.5 = 0 (do not hold checkers in reset)
8 c. 13h:177.7 = 1 (**reset error counts on all Banks all enabled lanes**)
9 d. 13h:129.4 = 0 (periodic update enabled)
10 e. 13h:177.0 = 1 (**5 second update interval**)

12 Host Write Sequence

- 13 1. Write bits 13h:160.7-0 (13h:168.7-0) to enable the pattern checker on selected media lanes.
14 a. The module will apply the configuration and control options to the enabled lanes.
15 b. Since the configuration is not gated and periodic update is enabled, the module may provide
16 the error information from the last update period.
17 c. Error counters may also available "on demand" when the host reads the current pattern checker
18 data, but this behavior is deliberately left as vendor dependent.
19 2. Write 14h to the Page select register. To provide better response to host, the module may return the
20 last polled error information.
21 3. Write 02h (04h) to the Diagnostics Selector 14h:128 to select lanes 1-4 for the host side (media side)
22 or write 03h (05h) to select host (media) lanes 5-8 (e.g. for DR8, FR8, LR8)
23 4. Host should read byte 14h:138 (14h:139) to ensure the pattern checker has not lost lock.
24 5. If the host wants to reset error information, the host shall set Bit 13h:177.5. When set:
25 a. Module (may also read from the data path chips and then) freezes the current error information.
26 Since this Bank's Bit 13h:177.7 is set, the module will also freeze the current error information
27 of all enabled Banks with Bit 13h:177.7 set.
28 b. Module copies current internal pattern checker counters into counting results. The module will
29 copy error information of all enabled lanes and Banks with Bit 13h:177.7 set to error information
30 results.
31 6. Host may write 11h-15h to the Diagnostics Selector 14h:128 to select the collected error information
32 results of all Banks, and then read Bytes 14h:192-255 to obtain the results of BER or of error counters
33 and total bits for the lanes represented in the respective Bank.
34 7. **To restart error counting** the host may clear Bit 13h:177.5. This will reset the current error
35 information (selector 01h-05h) without affecting the error information results (selector 11h-15h).

1 Appendix G Specification Evolution and Maintenance Notes

2 This appendix documents considerations for future evolution of the specification.

3 The audience of this appendix is the editor and the group involved in maintaining and evolving CMIS.

4 The desired property of “backwards compatibility” or “cross-version interoperability” of a new CMIS revision
5 leads to subtle constraints both for allowable specification changes and for later extensions of the specification.

6 The purpose of this appendix is to make those subtle constraints explicit and clearly visible to the CMIS audience
7 and the specification evolution group in particular.

8 As will be shown below, the following guideline should be used whenever possible:

CMIS Evolution Guideline: New CMIS functionality for modules should be specified in a way that it either works silently and unattended in the default configuration of the module, or that it is passivated by default until explicitly enabled by a host who is aware of the new functionality.

10 G.1 Definitions

11 A CMIS version is indicated by a version number with two components called **major revision number** and
12 **minor revision number**.

- 13 Backwards compatibility from a **host perspective** means that a host can, in principle, always interwork
14 with older modules (i.e. a host can manage modules using an earlier version of CMIS than the host).
- 15 Backwards compatibility from a **module perspective** means that a module can interwork with older hosts
16 (i.e. the module can be managed by hosts implementing an earlier version of CMIS than the module).

17 G.2 Cross-Version Compatibility

18 A host can query the CMIS version of a module, while the module does not know about the CMIS version
19 implemented in the host.

20 When the CMIS version supported by a module is **older** than the CMIS version run by the host, the host can,
21 at least in principle, adapt to a known older CMIS version and manage the module.

22 When the CMIS version supported by a module is **newer** and hence unknown to the host, it depends on the
23 nature of the changes between the host’s CMIS version and the newer CMIS version: If the new version provides
24 only **compatible** extensions, or **incompatible but passivated** extensions that work well with their default
25 settings and behavior, the host can still manage the module using the functionality supported by the host’s
26 older CMIS version.

27 G.3 Interpretation of CMIS Version Numbers

28 The CMIS version of a new CMIS revision will be defined to enable the following rules for hosts:

- 29 If a module reports the **same or a smaller CMIS major revision number** than the host, the module
30 can be managed (in principle, possibly depending on the host dynamically adapting to an older version)
- 31 If the module reports a **higher CMIS major revision number** than the host, the module may not behave
32 as per host expectations and therefore cannot be managed

Appendix H Examples for Network Path Applications

This appendix presents examples of NP Application advertisement and NP Application provisioning, in order to illustrate the specifications of section 7.6

Note: This appendix is relevant only to modules supporting NP Applications.

H.1 Advertisement Examples

The following advertisement examples explore

- homogeneous versus **mixed** multiplex application
- one versus several **alternative** multiplex applications
- one versus several **parallel** multiplex applications
- system interface (DP) versus uniplex (NP) application

See sections 0 and 7.6.4 for a description of the advertisement concepts.

See sections 8.2.12, 8.4.13, and 8.15.5.5 for the specification of the advertisement registers used.

Note: The shown sequence of Application Descriptors is not pre-determined and may vary across modules.

Note: In the following examples, symbolic values of Interface IDs like <400ZR> are shown when several suitable values or variants exist in [5], such as 62 and 63 for 400ZR DWDM and for 400ZR single channel, respectively.

H.1.1 400G Module for 400ZR DP and NP Application supporting Homogeneous Multiplex

The following advertisement describes a module supporting simple (homogeneous) multiplexing, in addition to a classical 400ZR system interface and a 400ZR uniplex application.

Note: The non-multiplexing <400ZR> system interface is actually a DP Application, by default, not a uniplex NP Application. The example assumes that the multiplex NP applications use the same Media Interface ID.

Note: The module supports also the uniplex NP Application that enables changing the multiplex structure without disrupting the network signal; this may be most useful in the 400ZR DWDM application.

Table H-1 400ZR NP Application Advertisement Example

Page	Byte	Bits	Field Name	Value	Remark
00h	86	7-0	HostInterfaceIDApp1	<400GAUI-8>	ID for "400GAUI-8"
	87	7-0	MediaInterfaceIDApp1	63	ID for "400ZR" (DP Application)
	88	7-4	HostLaneCountApp1	8	
		3-0	MediaLaneCountApp1	1	
	89	7-0	HostLaneAssignmentOptionsApp1	0000 0001b	
01h	176	7-0	MediaLaneAssignmentOptionsApp1	0000 0001b	
00h	90	7-0	HostInterfaceIDApp2	<200GAUI-4>	ID for "200GAUI-4"
	91	7-0	MediaInterfaceIDApp2	62	ID for "400ZR" (NP Application)
	92	7-4	HostLaneCountApp2	4	
		3-0	MediaLaneCountApp2	1	
	93	7-0	HostLaneAssignmentOptionsApp2	0001 0001b	
01h	177	7-0	MediaLaneAssignmentOptionsApp2	0000 0001b	
00h	94	7-0	HostInterfaceIDApp3	0Dh	ID for "100GAUI-2"
	95	7-0	MediaInterfaceIDApp3	62	ID for "400ZR" (NP Application)
	96	7-4	HostLaneCountApp3	2	
		3-0	MediaLaneCountApp3	1	
	97	7-0	HostLaneAssignmentOptionsApp3	0101 0101b	
01h	178	7-0	MediaLaneAssignmentOptionsApp3	0000 0001b	
00h	98	7-0	HostInterfaceIDApp4	<400GAUI-8>	ID for "400GAUI-8"
	99	7-0	MediaInterfaceIDApp4	62	"400ZR DWDM" (NP Application)
	100	7-4	HostLaneCountApp4	8	
		3-0	MediaLaneCountApp4	1	
	101	7-0	HostLaneAssignmentOptionsApp4	0000 0001b	
01h	179	7-0	MediaLaneAssignmentOptionsApp4	0000 0001b	
16h	228	7-0	MuxGranularity1	0	homogeneous multiplexing only
	248	7-0	ExtAppDescriptor<15-8>	xxxx xxxx b	
	249	7-0	ExtAppDescriptor<7-1>	xxx1 1100b	AppSel 1 is DP, 2, 3, 4 is NP

1 H.1.2 400G Module for 400ZR NP Application supporting Mixed Multiplex

2 The following advertisement describes a module supporting mixed multiplexing using the native 50G lane
 3 granularity of a QSFP-DD 400ZR module (as indicated by Host Interface IDs 100GAUI-2, 200GAUI-4, or
 4 400GAUI-8, any of which could be used to specify the granularity).

5 *Note: For reading convenience, a copy of Table 8-149 is provided below to show the meaning of the multiplex
 6 structure encoding.*

7 **Table H-2 400ZR NP Application Advertisement Mixed Multiplex Example**

Page	Byte	Bits	Field Name	Value	Remark
00h	86	7-0	HostInterfaceIDApp1	<400GAUI-8>	ID for "400GAUI-8"
	87	7-0	MediaInterfaceIDApp1	<400ZR>	ID for "400ZR"
	88	7-4	HostLaneCountApp1	8	
		3-0	MediaLaneCountApp1	1	
	89	7-0	HostLaneAssignmentOptionsApp1	0000 0001b	
01h	176	7-0	MediaLaneAssignmentOptionsApp1	0000 0001b	
00h	90	7-0	HostInterfaceIDApp2	<200GAUI-4>	ID for "200GAUI-4"
	91	7-0	MediaInterfaceIDApp2	<400ZR>	ID for "400ZR"
	92	7-4	HostLaneCountApp2	4	
		3-0	MediaLaneCountApp2	1	
	93	7-0	HostLaneAssignmentOptionsApp2	0001 0001b	
01h	177	7-0	MediaLaneAssignmentOptionsApp2	0000 0001b	
00h	94	7-0	HostInterfaceIDApp3	0Dh	ID for "100GAUI-2"
	95	7-0	MediaInterfaceIDApp3	<400ZR>	ID for "400ZR"
	96	7-4	HostLaneCountApp3	2	
		3-0	MediaLaneCountApp3	1	
	97	7-0	HostLaneAssignmentOptionsApp3	0101 0101b	
01h	178	7-0	MediaLaneAssignmentOptionsApp3	0000 0001b	
16h	228	7-0	MuxGranularity1	0Dh	ID for "100GAUI-2": 50G lanes
	229	7-0	MuxGranularity2	0	end of granularity list
	230	7-0	MuxGranularity3	-	don't care
	231	7-0	MuxGranularity4	-	don't care
	232	31-0	MuxStructsSupported1	0000 0000	Multiplex Structure IDs 1, 2, 3 for homogeneous and 4, 8 for some mixed multiplex (see also table below)
				0000 0000	
				0000 0001	
				0001 1110b	
	236	31-0	MuxStructsSupported2	-	don't care
	240	31-0	MuxStructsSupported3	-	don't care
	244	31-0	MuxStructsSupported4	-	don't care
	248	7-0	ExtAppDescriptor<15-8>	xxxx xxxx	
	249	7-0	ExtAppDescriptor<7-1>	xxxx 1110b	AppSel 1, 2, 3 is NP

8 **Table H-3 Global Multiplex Structure Advertisement**

Multiplex Structure			HP DPID per Host Lane #							
ID	# of HPs	HP Widths [lanes]	1	2	3	4	5	6	7	8
0	8	1, 1, 1, 1, 1, 1, 1, 1	1	2	3	4	5	6	7	8
1	1	8	1	1	1	1	1	1	1	1
2	2	4, 4	1	1	1	1	5	5	5	5
3	4	2, 2, 2, 2	1	1	3	3	5	5	7	7
4	3	4, 2, 2	1	1	1	1	5	5	7	7
5	4	4, 2, 1, 1	1	1	1	1	5	5	7	8
6	4	4, 1, 1, 2	1	1	1	1	5	6	7	7
7	5	4, 1, 1, 1, 1	1	1	1	1	5	6	7	8
8	3	2, 2, 4	1	1	3	3	5	5	5	5
9	4	2, 1, 1, 4	1	1	3	4	5	5	5	5
10	4	1, 1, 2, 4	1	2	3	3	5	5	5	5
11	4	1, 1, 1, 1, 1	1	2	3	4	5	5	5	5

12	5	2, 2, 2, 1, 1	1	1	3	3	5	5	7	8
13	5	2, 2, 1, 1, 2	1	1	3	3	5	6	7	7
14	5	2, 1, 1, 2, 2	1	1	3		5	5	7	7
15	5	1, 1, 2, 2, 2	1	2	3	3	5	5	7	7
16	6	2, 2, 1, 1, 1, 1	1	1	3	3	5	6	7	8
17	6	2, 1, 1, 2, 1, 1	1	1	3	4	5	5	7	8
18	6	2, 1, 1, 1, 1, 2	1	1	3	4	5	6	7	7
19	6	1, 1, 2, 2, 1, 1	1	2	3	3	5	5	7	8
20	6	1, 1, 2, 1, 1, 2	1	2	3	3	5	6	7	7
21	6	1, 1, 1, 1, 2, 2	1	2	3	4	5	5	7	7
22	6	2, 1, 1, 1, 1, 1, 1	1	1	3	4	5	6	7	8
23	7	1, 1, 2, 1, 1, 1, 1	1	2	3	3	5	6	7	8
24	7	1, 1, 1, 1, 2, 1, 1	1	2	3	4	5	5	7	8
25	7	1, 1, 1, 1, 1, 1, 2	1	2	3	4	5	6	7	7

1

H.1.3 400G Module with Alternative Support of 400G or 200G NP Application

If the module described in H.1.2 would **alternatively** support a 200G NP on its single media lane that just uses 25G host lane granularity instead of 50G granularity, additional multiplex Application Descriptors for this type of 200G NP would be needed, and the advertising in page 16h might look as follows

6

Table H-4 Multiple Multiplex Granularities Advertisement Example

Page	Byte	Bits	Field Name	Value	Remark
16h	228	7-0	MuxGranularity1	0Dh	ID for "100GAUI-2": 50G lanes
	229	7-0	MuxGranularity2	<50GAUI-2>	ID for "50GAUI-2": 25G lanes
	230	7-0	MuxGranularity3	0	end of granularity list
	231	7-0	MuxGranularity4	-	don't care
	232	31-0	MuxStructsSupported1	0000 0000 0000 0000 0000 0001 0001 1110b	Multiplex Structure IDs 1, 2, 3 for homogeneous and 4, 8 for mixed multiplex of 50G lanes (see also table below)
	236	31-0	MuxStructsSupported2	0000 0000 0000 0000 0000 0001 0001 1110b	Same multiplex structures for 25G lanes as for 50G lanes
	240	31-0	MuxStructsSupported3	-	don't care
	244	31-0	MuxStructsSupported4	-	don't care

7

H.1.4 800G Module with Parallel 400ZR NP or DP Applications

The following configuration advertises **simultaneous** (rather than alternative) support of two identical **parallel** 400G NP Applications (each with one media lane) or 400G DP Applications, or a mix thereof.

8

Table H-5 2x400ZR NP Application Advertisement Example

Page	Byte	Bits	Field Name	Value	Remark
00h	86	7-0	HostInterfaceIDApp1	<400GAUI-4>	ID for "400GAUI-4"
	87	7-0	MediaInterfaceIDApp1	<400ZR>	ID for "400ZR"
	88	7-4	HostLaneCountApp1	4	
		3-0	MediaLaneCountApp1	1	
	89	7-0	HostLaneAssignmentOptionsApp1	0001 0001b	
01h	176	7-0	MediaLaneAssignmentOptionsApp1	0000 0011b	
00h	90	7-0	HostInterfaceIDApp2	<200GAUI-2>	ID for "200GAUI-2"
	91	7-0	MediaInterfaceIDApp2	<400ZR>	ID for "400ZR"
	92	7-4	HostLaneCountApp2	2	
		3-0	MediaLaneCountApp2	1	
	93	7-0	HostLaneAssignmentOptionsApp2	0101 0101b	
01h	177	7-0	MediaLaneAssignmentOptionsApp2	0000 0011b	
00h	94	7-0	HostInterfaceIDApp3	<100GAUI-1>	ID for "100GAUI-1"

	95	7-0	MediaInterfaceIDApp3	<400ZR>	ID for "400ZR"
	96	7-4	HostLaneCountApp3	1	
		3-0	MediaLaneCountApp3	1	
	97	7-0	HostLaneAssignmentOptionsApp3	1111 1111b	
01h	178	7-0	MediaLaneAssignmentOptionsApp3	0000 0011b	
00h	98	7-0	HostInterfaceIDApp4	<400GAUI-4>	ID for "400GAUI-4"
	99	7-0	MediaInterfaceIDApp4	<400ZR>	ID for "400ZR DWDM"
	100	7-4	HostLaneCountApp4	4	
		3-0	MediaLaneCountApp4	1	
	101	7-0	HostLaneAssignmentOptionsApp4	0001 0001b	
01h	179	7-0	MediaLaneAssignmentOptionsApp4	0000 0011b	
16h	228	7-0	MuxGranularity1	0	homogeneous multiplexing only
	248	7-0	ExtAppDescriptor<15-8>	xxxx xxxx b	
	249	7-0	ExtAppDescriptor<7-1>	xxx0 1110b	AppSel 1,2,3 is NP, 4 is DP

1

H.1.5 800G Module for 400ZR NP Application and Parallel DP Applications

The following configuration advertises **simultaneous** (rather than alternative) support of a 400G NP Application (with one media lane) and in **parallel** either one 400G DP Application (with four media lanes) or four DP Applications (with one media lane each).

2

Table H-6 400ZR + 400G-DR4 or 4x100G-DR1 Application Advertisement Example

Page	Byte	Bits	Field Name	Value	Remark
00h	86	7-0	HostInterfaceIDApp1	<400GAUI-4>	ID for "400GAUI-4"
	87	7-0	MediaInterfaceIDApp1	<400ZR>	ID for "400ZR"
	88	7-4	HostLaneCountApp1	4	
		3-0	MediaLaneCountApp1	1	
	89	7-0	HostLaneAssignmentOptionsApp1	0000 0001b	
01h	176	7-0	MediaLaneAssignmentOptionsApp1	0000 0001b	
00h	90	7-0	HostInterfaceIDApp2	<200GAUI-2>	ID for "200GAUI-2"
	91	7-0	MediaInterfaceIDApp2	<400ZR>	ID for "400ZR"
	92	7-4	HostLaneCountApp2	2	
		3-0	MediaLaneCountApp2	1	
	93	7-0	HostLaneAssignmentOptionsApp2	0000 0101b	
01h	177	7-0	MediaLaneAssignmentOptionsApp2	0000 0001b	
00h	94	7-0	HostInterfaceIDApp3	<100GAUI-1>	ID for "100GAUI-1"
	95	7-0	MediaInterfaceIDApp3	<400ZR>	ID for "400ZR"
	96	7-4	HostLaneCountApp3	1	
		3-0	MediaLaneCountApp3	1	
	97	7-0	HostLaneAssignmentOptionsApp3	0000 1111b	
01h	178	7-0	MediaLaneAssignmentOptionsApp3	0000 0001b	
00h	98	7-0	HostInterfaceIDApp4	<400GAUI-4>	ID for "400GAUI-4"
	99	7-0	MediaInterfaceIDApp4	<400G-DR4>	ID for "400G-DR4"
	100	7-4	HostLaneCountApp4	4	
		3-0	MediaLaneCountApp4	4	
	101	7-0	HostLaneAssignmentOptionsApp4	0001 0000b	
01h	179	7-0	MediaLaneAssignmentOptionsApp4	0000 0010b	
00h	102	7-0	HostInterfaceIDApp5	<100GAUI-1>	ID for "100GAUI-1"
	103	7-0	MediaInterfaceIDApp5	<100G-DR1>	ID for "100G-DR1"
	104	7-4	HostLaneCountApp5	1	
		3-0	MediaLaneCountApp5	1	
	105	7-0	HostLaneAssignmentOptionsApp5	1111 0000b	
01h	180	7-0	MediaLaneAssignmentOptionsApp5	0001 1110b	
16h	228	7-0	MuxGranularity1	0	homogeneous multiplexing only
	248	7-0	ExtAppDescriptor<15-8>	xxxx xxxx b	
	249	7-0	ExtAppDescriptor<7-1>	xx00 1110b	AppSel 1,2,3 is NP, 4,5 is DP

H.2 Provisioning Examples

The group of Host Paths (HP) and the Network Path (NP) participating in an N:1 NP Application instance are defined, **provisioned**, and **commissioned** in separate steps. Also the dynamic configuration states of all commissioned HPs and of the NP are **observed** and **controlled** separately.

See section 7.6.5 for more information.

Note: The examples below arbitrarily use the first staged control set instance

Note: HPs are managed through the DP registers (Page 10h), NPs through the NP registers (Page 16h)

H.2.1 4x100G NP Application with Unused Multiplexing Slots

Assuming the advertisement described in section H.1.2, the following registers provision and then commission the NP Application instance for 4 x 100G ZR.

Table H-7 400ZR NP Provisioning Example

Page	Byte	Bits	Register Name	Value	Remark
10h	145	7-0	SCS0::DPConfigLane1	0011 000 0b	(AppSel, DPID, EC) = (3, 0, 0)
	146	7-0	SCS0::DPConfigLane2	0011 000 0b	(AppSel, DPID, EC) = (3, 0, 0)
	147	7-0	SCS0::DPConfigLane3	0011 010 0b	(AppSel, DPID, EC) = (3, 2, 0)
	148	7-0	SCS0::DPConfigLane4	0011 010 0b	(AppSel, DPID, EC) = (3, 2, 0)
	149	7-0	SCS0::DPConfigLane5	0011 100 0b	(AppSel, DPID, EC) = (3, 4, 0)
	150	7-0	SCS0::DPConfigLane6	0011 100 0b	(AppSel, DPID, EC) = (3, 4, 0)
	151	7-0	SCS0::DPConfigLane7	0011 110 0b	(AppSel, DPID, EC) = (3, 6, 0)
	152	7-0	SCS0::DPConfigLane8	0011 110 0b	(AppSel, DPID, EC) = (3, 6, 0)
16h	128	7-0	SCS0::NPConfigLane1	0000 000 1b	(NPID, NPInUse) = (0, 1)
	129	7-0	SCS0::NPConfigLane2	0000 000 1b	(NPID, NPInUse) = (0, 1)
	130	7-0	SCS0::NPConfigLane3	0000 000 1b	(NPID, NPInUse) = (0, 1)
	131	7-0	SCS0::NPConfigLane4	0000 000 1b	(NPID, NPInUse) = (0, 1)
	132	7-0	SCS0::NPConfigLane5	0000 000 1b	(NPID, NPInUse) = (0, 1)
	133	7-0	SCS0::NPConfigLane6	0000 000 1b	(NPID, NPInUse) = (0, 1)
	134	7-0	SCS0::NPConfigLane7	0000 000 1b	(NPID, NPInUse) = (0, 1)
	135	7-0	SCS0::NPConfigLane8	0000 000 1b	(NPID, NPInUse) = (0, 1)
10h	128	7	DPDeinitLane8	1b	Keep 4 th HP in DPDeactivated
		6	DPDeinitLane7	1b	
		5	DPDeinitLane6	0b	Initialize 3 rd HP
		4	DPDeinitLane5	0b	
		3	DPDeinitLane4	0b	Initialize 2 nd HP
		2	DPDeinitLane3	0b	
		1	DPDeinitLane2	0b	Initialize 1 st HP
		0	DPDeinitLane1	0b	
16h	160	7	NPDeinitLane8	0b	Initialize the NP entirely (all 8 host lanes feed into that NP) <i>Note: the multiplex structure of the NP must be provisioned in the Active Control Set, in terms of HPs, while the DPSM state of each HP is irrelevant.</i>
		6	NPDeinitLane7	0b	
		5	NPDeinitLane6	0b	
		4	NPDeinitLane5	0b	
		3	NPDeinitLane4	0b	
		2	NPDeinitLane3	0b	
		1	NPDeinitLane2	0b	
		0	NPDeinitLane1	0b	

H.2.2 800G Module with Parallel 400ZR NP Applications

Assuming the advertisement described in section H.1.4, the following registers provision one 400ZR NP Application instance and one parallel 400ZR DP Application instance.

Table H-8 2 x 400ZR NPs Provisioning Example

Page	Byte	Bits	Register Name	Value	Remark
10h	145	7-0	SCS0::DPConfigLane1	0001 000 0b	(AppSel, DPID, EC) = (1, 0, 0)
	146	7-0	SCS0::DPConfigLane2	0001 000 0b	(AppSel, DPID, EC) = (1, 0, 0)
	147	7-0	SCS0::DPConfigLane3	0001 000 0b	(AppSel, DPID, EC) = (1, 0, 0)
	148	7-0	SCS0::DPConfigLane4	0001 000 0b	(AppSel, DPID, EC) = (1, 0, 0)
	149	7-0	SCS0::DPConfigLane5	0100 100 0b	(AppSel, DPID, EC) = (4, 4, 0)
	150	7-0	SCS0::DPConfigLane6	0100 100 0b	(AppSel, DPID, EC) = (4, 4, 0)
	151	7-0	SCS0::DPConfigLane7	0100 100 0b	(AppSel, DPID, EC) = (4, 4, 0)
	152	7-0	SCS0::DPConfigLane8	0100 100 0b	(AppSel, DPID, EC) = (4, 4, 0)
16h	128	7-0	SCS0::NPConfigLane1	0000 000 1b	(NPID, NPInUse) = (0, 1)
	129	7-0	SCS0::NPConfigLane2	0000 000 1b	(NPID, NPInUse) = (0, 1)
	130	7-0	SCS0::NPConfigLane3	0000 000 1b	(NPID, NPInUse) = (0, 1)
	131	7-0	SCS0::NPConfigLane4	0000 000 1b	(NPID, NPInUse) = (0, 1)
	132	7-0	SCS0::NPConfigLane5	0000 100 1b	(NPID, NPInUse) = (X, 0)
	133	7-0	SCS0::NPConfigLane6	0000 100 1b	(NPID, NPInUse) = (X, 0)
	134	7-0	SCS0::NPConfigLane7	0000 100 1b	(NPID, NPInUse) = (X, 0)
	135	7-0	SCS0::NPConfigLane8	0000 100 1b	(NPID, NPInUse) = (X, 0)

H.2.3 800G Module for 400ZR NP Application and Other Parallel DP Applications

Assuming the advertisement described in section H.1.5, the following registers provision the five NP Application instances for 400G ZR + 4 x 100G DR1

Table H-9 400ZR + 4x100G-DR1 NP Provisioning Example

Page	Byte	Bits	Register Name	Value	Remark
10h	145	7-0	SCS0::DPConfigLane1	0010 000 0b	(AppSel, DPID, EC) = (2, 0, 0)
	146	7-0	SCS0::DPConfigLane2	0010 000 0b	(AppSel, DPID, EC) = (2, 0, 0)
	147	7-0	SCS0::DPConfigLane3	0011 010 0b	(AppSel, DPID, EC) = (3, 2, 0)
	148	7-0	SCS0::DPConfigLane4	0011 011 0b	(AppSel, DPID, EC) = (3, 3, 0)
	149	7-0	SCS0::DPConfigLane5	0101 100 0b	(AppSel, DPID, EC) = (5, 4, 0)
	150	7-0	SCS0::DPConfigLane6	0101 101 0b	(AppSel, DPID, EC) = (5, 5, 0)
	151	7-0	SCS0::DPConfigLane7	0101 110 0b	(AppSel, DPID, EC) = (5, 6, 0)
	152	7-0	SCS0::DPConfigLane8	0101 111 0b	(AppSel, DPID, EC) = (5, 7, 0)
16h	128	7-0	SCS0::NPConfigLane1	0000 000 1b	(NPID, NPInUse) = (0, 1)
	129	7-0	SCS0::NPConfigLane2	0000 000 1b	(NPID, NPInUse) = (0, 1)
	130	7-0	SCS0::NPConfigLane3	0000 000 1b	(NPID, NPInUse) = (0, 1)
	131	7-0	SCS0::NPConfigLane4	0000 000 1b	(NPID, NPInUse) = (0, 1)
	132	7-0	SCS0::NPConfigLane5	0000 xxx 0b	(NPID, NPInUse) = (X, 0)
	133	7-0	SCS0::NPConfigLane6	0000 xxx 0b	(NPID, NPInUse) = (X, 0)
	134	7-0	SCS0::NPConfigLane7	0000 xxx 0b	(NPID, NPInUse) = (X, 0)
	135	7-0	SCS0::NPConfigLane8	0000 xxx 0b	(NPID, NPInUse) = (X, 0)

1 Appendix I Companies Belonging to OIF at Time of Approval

2	Accelight Technologies, Inc.	Google
	Accton Technology Corporation	H3C Technologies Co., Ltd.
	Adtran Networks SE	Hakusan Inc
	Advanced Fiber Resources (AFR)	Hewlett Packard Enterprise (HPE)
	Advanced Micro Devices, Inc.	HGGenuine Optics Tech Company
	AIO Core Co., Ltd	Hirose Electric Co. Ltd.
	Alibaba	Hisense Broadband Multimedia Technologies Co., LTD
	Alphawave Semi	Huawei Technologies Co., Ltd.
	Amazon	Infinera Corporation
	Amphenol Corp.	InfiniLink
	Anritsu	InnoLight Technology Limited
	Applied Optoelectronics, Inc.	Integrated Device Technology
	Arista Networks	Intel
	Astera Labs	Juniper Networks
	Ayar Labs	Kandou Bus
	BitifEye Digital Test Solutions GmbH	KDDI Research, Inc.
	BizLink Technology, Inc.	Keysight Technologies, Inc.
	Broadcom Inc.	KYOCERA Corporation
	Cadence Design Systems	Lessengers Inc.
	Casela Technologies USA	Lightmatter
	Celestica	Linktel Technologies Co., Ltd.
	China Telecom	Lumentum
	CICT	Lumiphase AG
	Ciena Corporation	LUXIC Technology Co
	Cisco Systems	Luxshare Technologies International, Inc.
	Coherent	MACOM Technology Solutions
	ColorChip LTD	Marvell Semiconductor, Inc.
	Cornelis Networks, Inc.	MaxLinear Inc.
	Corning	MediaTek
	Credo Semiconductor (HK) LTD	Meta Platforms
	Dai Nippon Printing Co., Ltd.	Microchip Technology Incorporated
	Dell, Inc.	Microsoft Corporation
	Dexerials Corporation	Mitsubishi Electric US, Inc.
	EFFECT Photonics B.V.	Molex
	Eoptolink Technology	Multilane Inc.
	Epson Electronics America, Inc.	NEC Corporation
	Ericsson	Nokia
	EXFO	NTT Corporation
	Foxconn Interconnect Technology Ltd	Nubis Communications, Inc.
	Fujikura	NVIDIA
	Fujitsu	O-Net Technologies (Shenzhen) Group Co., Limited
	Furukawa Electric Co., Ltd.	Omniva LLC
	Global Foundries	

Optomind Inc.
Orange
PETRA
Point2 Technology
Precision Optical Technologies
Quantifi Photonics USA Inc.
Quintessent Inc.
Ranovus
Retym
Rosenberger Hochfrequenztechnik
GmbH & Co. KG
Samsung Electronics Co. Ltd.
Samtec Inc.
SCINTIL Photonics
Semtech Canada Corporation
Senko Advanced Components
SerialLink Systems Ltd.
Sicoya GmbH
SiFotonics Technologies Inc.
Silith Technology
Socionext Inc.
Source Photonics, Inc.
Spirent Communications
Sumitomo Electric Industries, Ltd.
Sumitomo Osaka Cement
Synopsys, Inc.
TE Connectivity
Tektronix
Telefonica S.A.
TELUS Communications, Inc.
Teramount
TeraSignal, LLC.
US Conec
Viavi Solutions Deutschland GmbH
Wilder Technologies, LLC
Wistron Corporation
Xphor Ltd.
Yamaichi Electronics Ltd.
ZTE Corporation

1

2

3

End of Document